



Attributes



Devi Parikh
Virginia Tech



Ali Farhadi
UW



Kristen Grauman
UT Austin



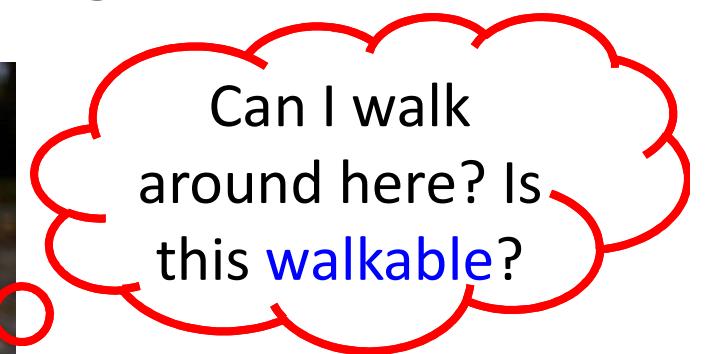
Tamara Berg
UNC (Fall 13)



Abhinav Gupta
CMU

Why attributes?

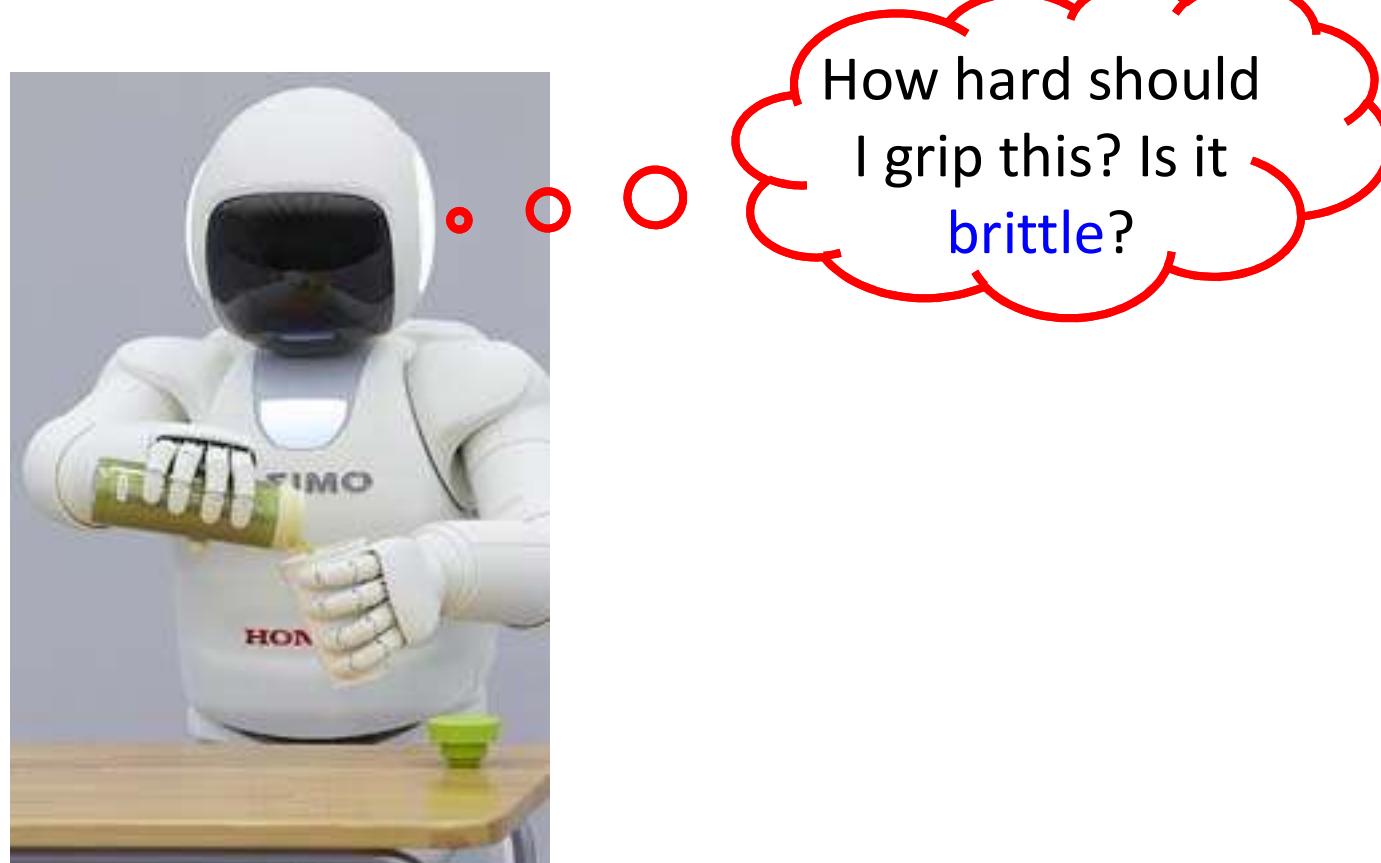
- Why recognition?
- Why would a robot need to recognize a scene?



Slide credit: Devi Parikh

Why attributes?

- Why would a robot need to recognize an object?



Slide credit: Devi Parikh

Why attributes?

- How humans naturally describe visual concepts
- Image search



I want elegant
silver sandals
with high heels

Why attributes?

- How humans think of concepts
- Domain knowledge helpful to build visual models

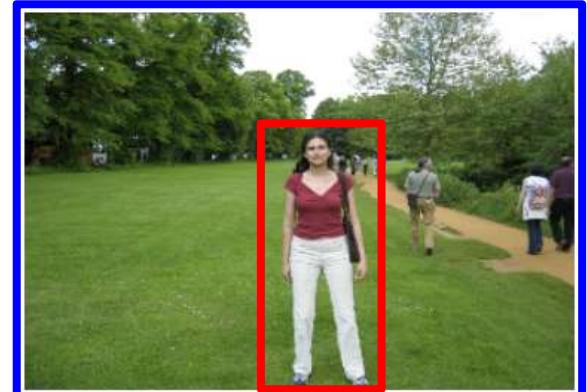


Slide credit: Devi Parikh

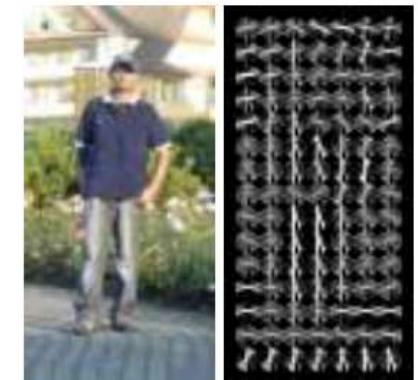
What are attributes?

- Mid-level concepts
 - Higher than low-level features
 - Lower than high-level categories
- Shared across categories
- Human-understandable (semantic)
- Machine-detectable (visual)

Outdoor, Peaceful, Natural, Green,...



Female, Long-hair, Young ...



Dalal and Triggs, 2005

Slide credit: Devi Parikh



Ladicky et al, 2010 6

What are attributes?

- Material, Appearance, Function, ...
- Any adjective
- Statements *about* visual concepts
- Objects, scene: Nouns
- Actions: Verbs
- Distinctions less critical in how to predict attributes
- Distinctions important in uses of attributes for improving computer vision
- Parts often semantic, shared, mid-level
- Can be used as attributes

1at·tri·bute noun \ə-trə-,byüt\

Definition of ATTRIBUTE

1 : an inherent characteristic; *also* : an accidental quality

2 : an object closely associated with or belonging to a specific person, thing, or office <a scepter is the *attribute* of power>; *especially* : such an object used for identification in painting or sculpture

3 : a word ascribing a quality; *especially* : **ADJECTIVE**



Example Attributes

Face Tracer Image Search

“Smiling Asian Men With Glasses”



Kumar et al. 2008

Slide credit: Devi Parikh

Example Attributes



Farhadi et al. 2009

Example Attributes

otter

black: yes
white: no
brown: yes
stripes: no
water: yes
eats fish: yes



polar bear

black: no
white: yes
brown: no
stripes: no
water: yes
eats fish: yes



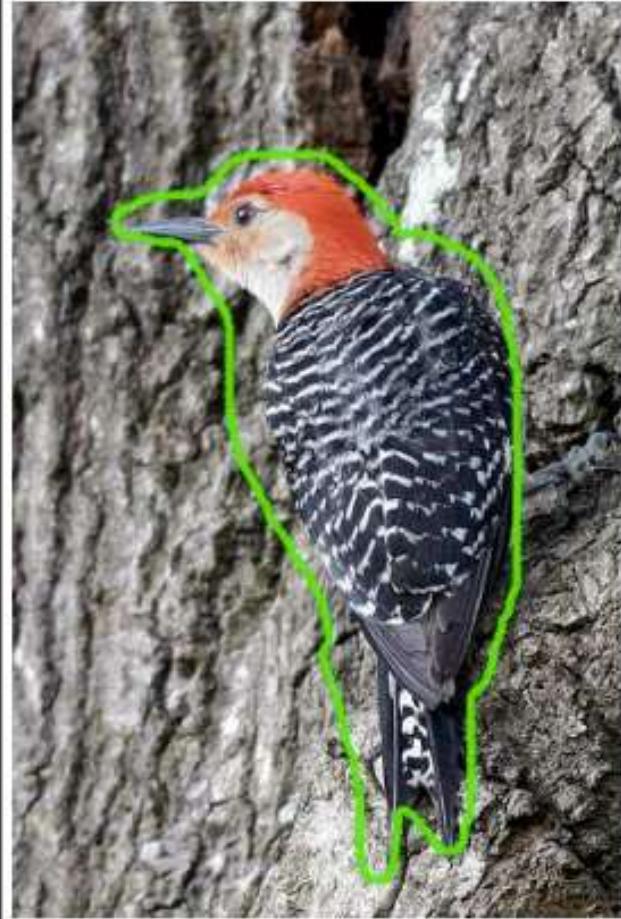
zebra

black: yes
white: yes
brown: no
stripes: yes
water: no
eats fish: no



Lampert et al. 2009

Example Attributes

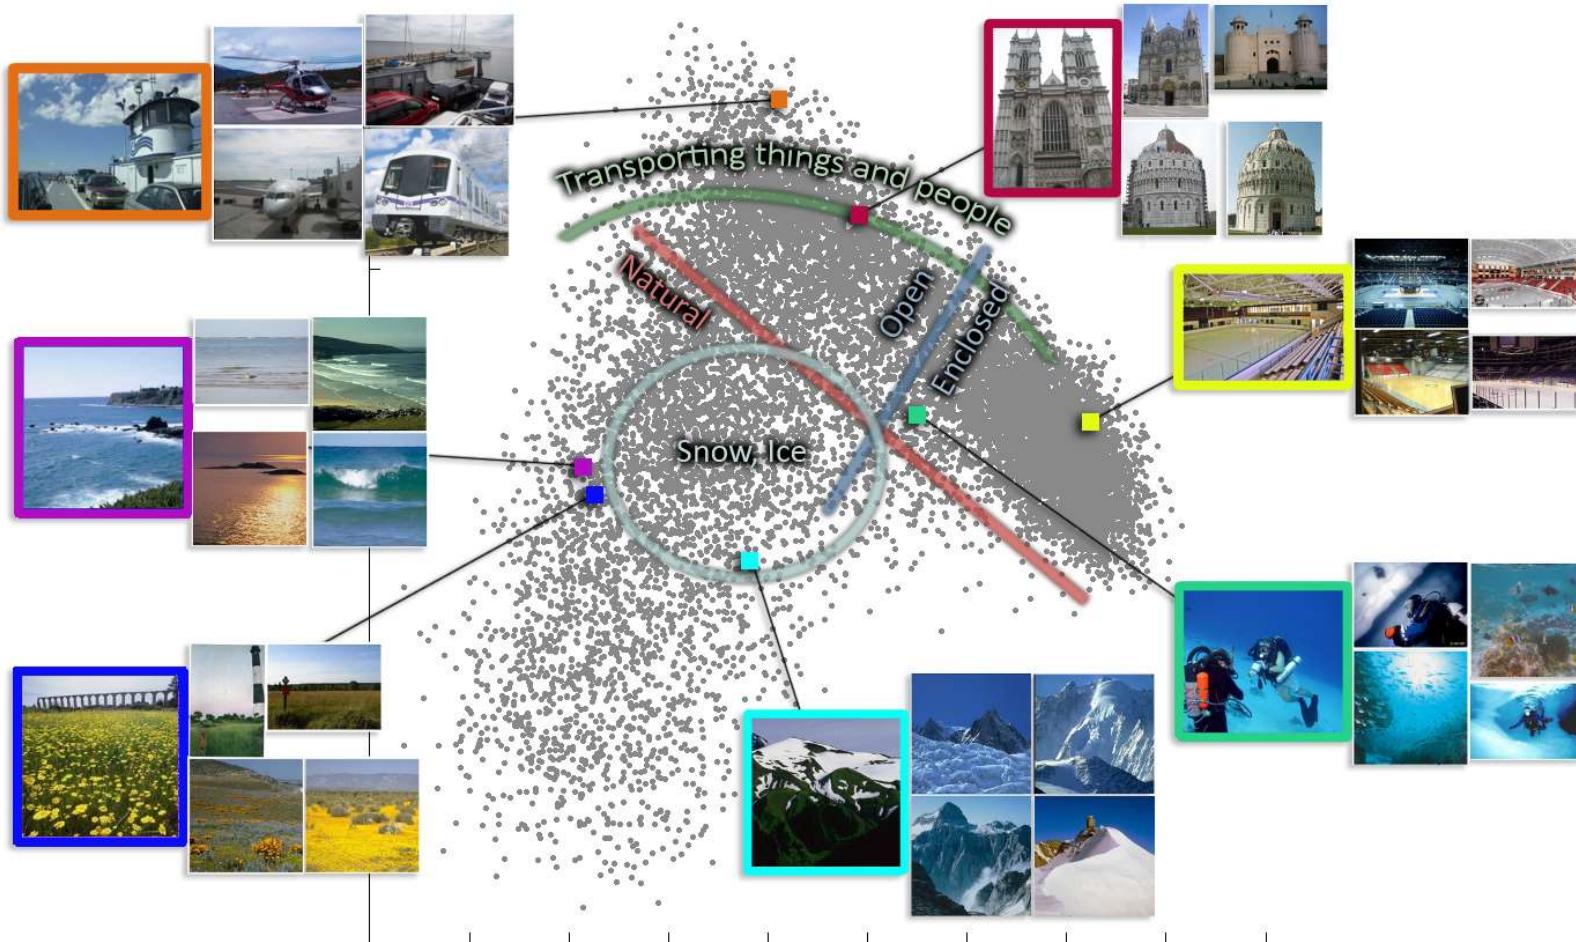


forehead_color	red	red	red
breast_pattern	multi-colored	solid	solid
breast_color	white	white/red	white
head_pattern	capped	capped	capped
back_color	white/black	white/black	white/black
wing_color	white/black	white/black	white/black
leg_color	buff	black	black
size	small	medium	medium
bill_shape	all-purpose	dagger	all-purpose
wing_shape	pointed	tapered	pointed
...
primary_color	black, red	white, black	white, black

Welinder et al. 2010

Slide credit: Devi Parikh

Example Attributes



Patterson and Hays 2011

Slide credit: Devi Parikh

Example Attributes



Berg et al. 2010

Slide credit: Devi Parikh

Relative Attributes



>
natural



<
smiling



Parikh and Grauman 2011

Slide credit: Devi Parikh

Some Notation

- Vocabulary of attributes:

$$A = \{a_m\}, m \in \{1, \dots, M\}$$

- Image features:

$$\{\mathbf{x}_i\}, i \in \{1, \dots, N\}$$

Attribute Models

- Classifiers for binary attributes

Attribute	Positive Examples			Negative Examples		
Asian						
Blond Hair						
Child						
Male						

$$\mathbf{x}_i \rightarrow \{+1, -1\}$$

(Or confidence)

Kumar et al. 2010

Slide credit: Devi Parikh

Attribute Models

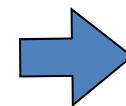
Weakly supervised learning

Noisy labels !

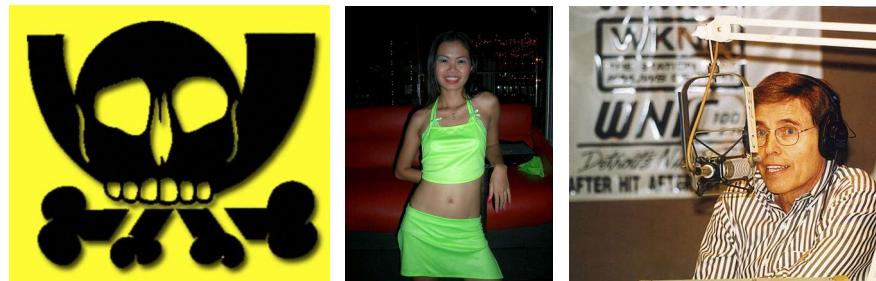
Google™
"stripes"
"things"



Positive training images
(unsegmented)



Attribute
model



Negative training images

Ferrari et al. 2007

Slide adapted from Vittorio Ferrari by Devi Parikh

Attribute Models

- Ranking functions for relative attributes

For each attribute a_m , open

Supervision is

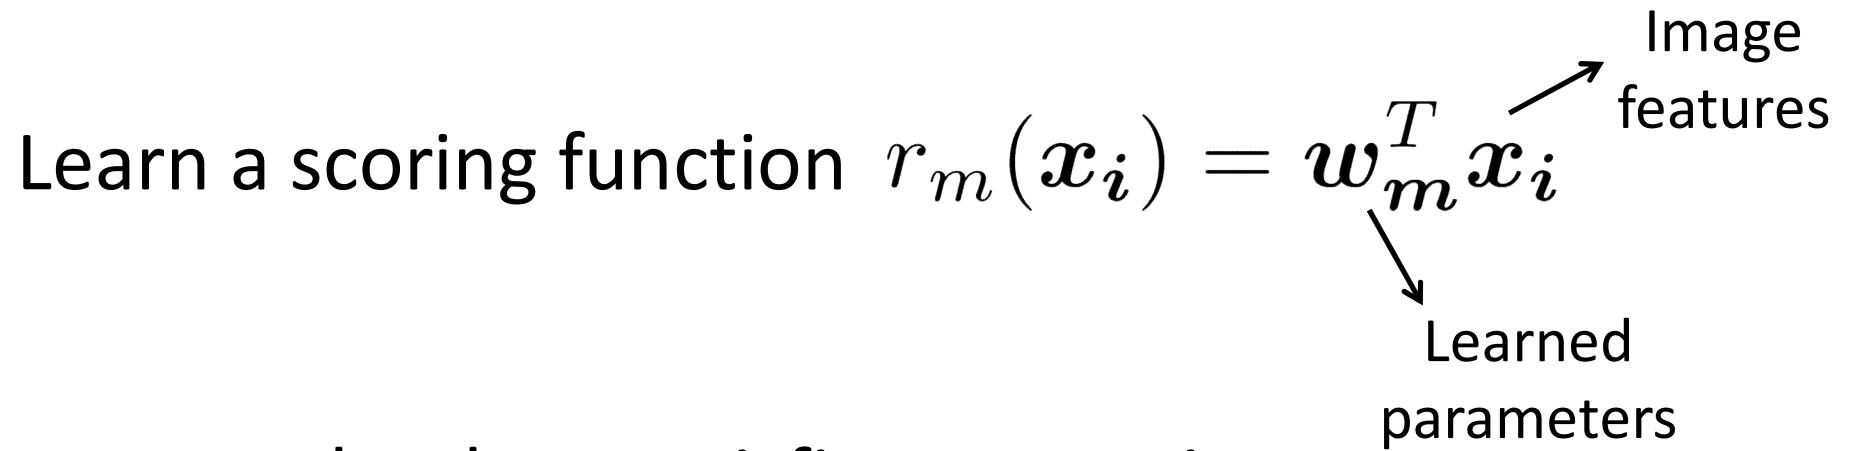
$$O_m: \left\{ \left(\begin{array}{c} \text{[Image of Hallgrímskirkja]} \\ \sim \\ \text{[Image of skyscrapers]} \end{array} \right), \dots \right\},$$

$$S_m: \left\{ \left\{ \begin{array}{c} \text{[Image of beach]} \\ \sim \\ \text{[Image of landscape]} \end{array} \right\}, \dots \right\}$$

Slide credit: Devi Parikh

Attribute Models

Learn a scoring function $r_m(\mathbf{x}_i) = \mathbf{w}_m^T \mathbf{x}_i$



that best satisfies constraints:

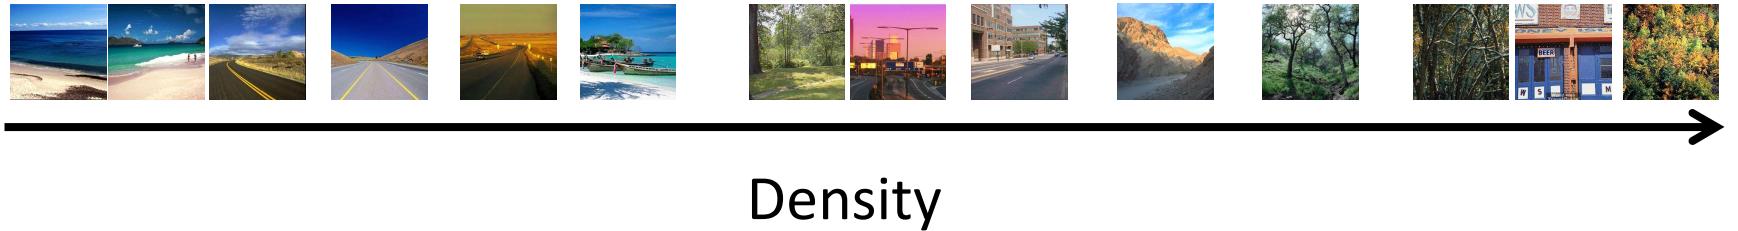
$$\forall (i, j) \in O_m : \mathbf{w}_m^T \mathbf{x}_i > \mathbf{w}_m^T \mathbf{x}_j$$

$$\forall (i, j) \in S_m : \mathbf{w}_m^T \mathbf{x}_i = \mathbf{w}_m^T \mathbf{x}_j$$

Max-margin learning to rank formulation of Joachims 2002

Attribute Models

$x_i \rightarrow$ Real value
(strength of attribute presence)



“I am 60% sure this person is smiling”
(Binary Classifier Confidence)

“This person is smiling 60%”
(Attribute Strength)

“Person A is smiling more than Person B”
(Relative Attribute)

Offline Uses of Attributes

Zero-shot Learning

- Aye-ayes
 - Are nocturnal
 - Live in trees
 - Have large eyes
 - Have long middle fingers

Which one of these is an aye-aye?



Humans can learn from descriptions (zero examples).

Zero-shot Learning

- Seen categories with labeled images
 - Train attribute predictors
- Unseen categories
 - No examples, only description

	bea	turtle	rabbit
furry	○	●	○
big	○	●	●
...

Zero-shot Learning

- Test image: \mathbf{x}
- Test class: z
- Classification: $\operatorname{argmax}_z p(z|\mathbf{x})$

Zero-shot Learning

Lampert et al. 2009

$$p(z|x)$$

$z =$ zebra

$$a =$$

striped	1
4-legs	1
large	
0	
...	

z is independent of x given a

0 for all other a

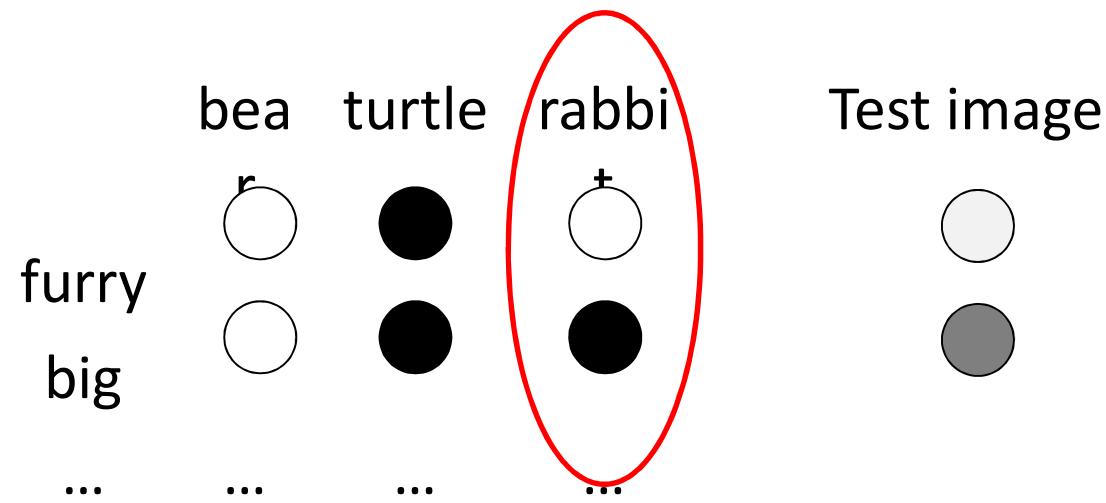
→ prior (e.g. uniform)

estimate from data or uniform

(assuming uniform priors above)

outputs of binary classifiers

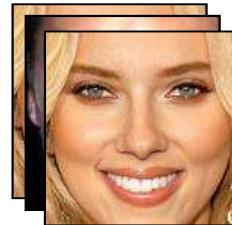
Zero-shot Learning



Farhadi et al. 2009

Relative Zero-shot Learning

Training: Images from **S seen** categories and
Descriptions of **U unseen** categories



Age: **Hugh** \succ **Clive** \succ **Scarlett**

Jared \succ **Miley**



Miley \succ **Jared**

Smiling:

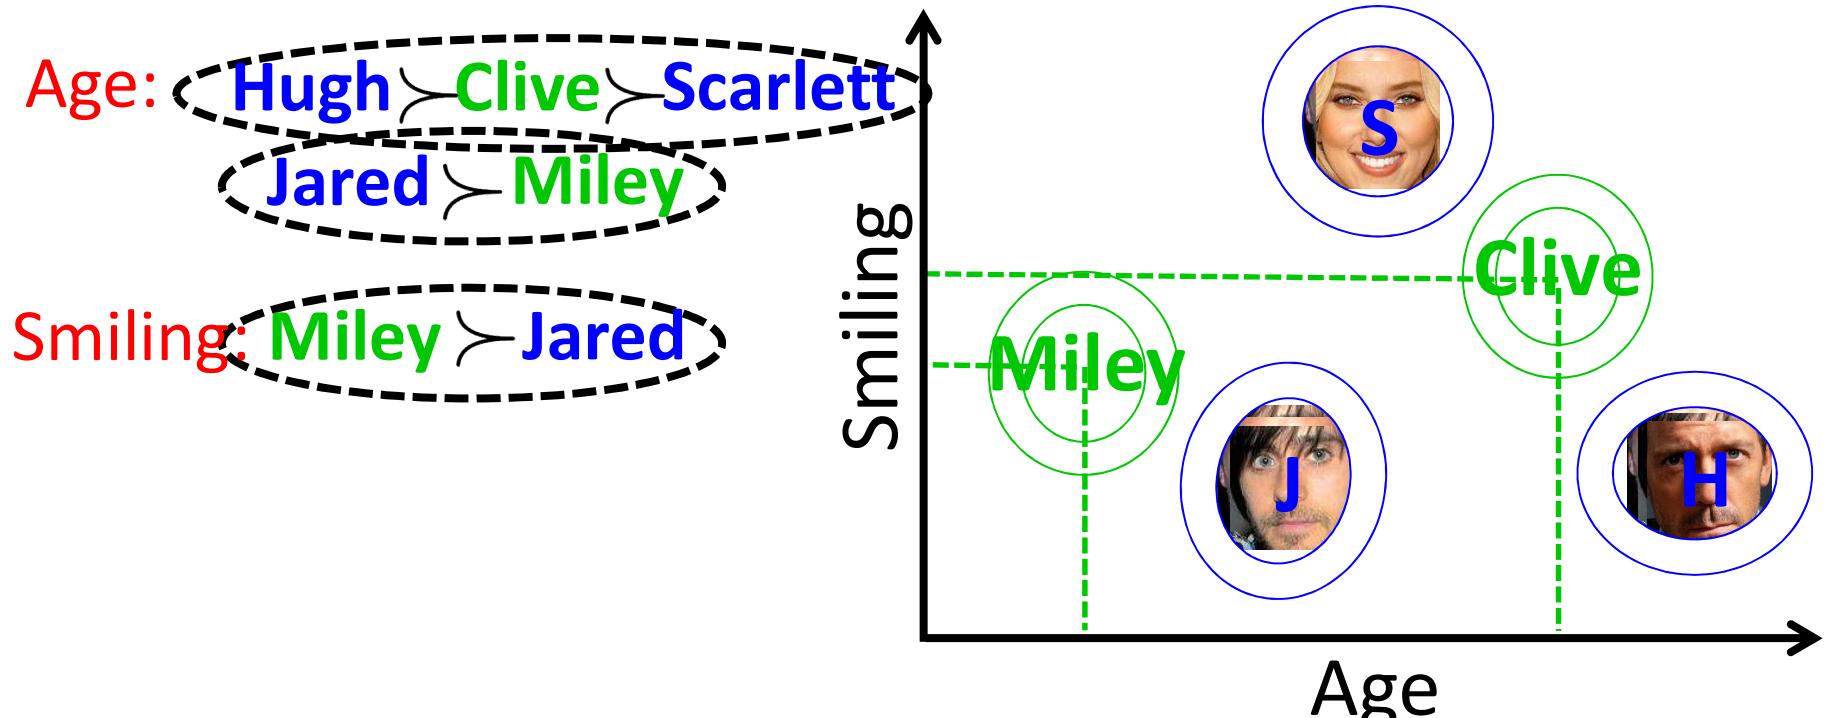
Need not use all attributes, or all seen categories

Testing: Categorize image into one of **S+U** categories

Relative Zero-shot Learning

Parikh and Grauman 2011

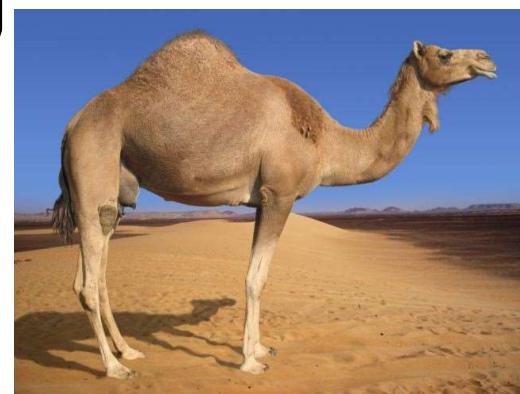
Can predict new classes based on their relationships to existing classes – without training images



Infer image category using max-likelihood

Learning with a few examples

- Active Learning



Current belief

Focused feedback Knowledge of the world

I think this is a giraffe. What do you think?



No, its neck is too short for it to be a giraffe.

3

- Learner learns better from its mistakes
- Accelerated discriminative learning with few examples

not be giraffes either then.



.....

Feedback on one, transferred to many

Attributes-based Feedback

Attribute predictors:

a_1 a_2 ... a_M



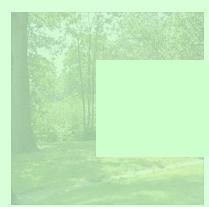
Forest



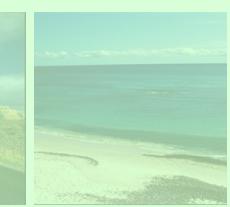
Unlabeled pool
of images

No,

It is too open to be a forest



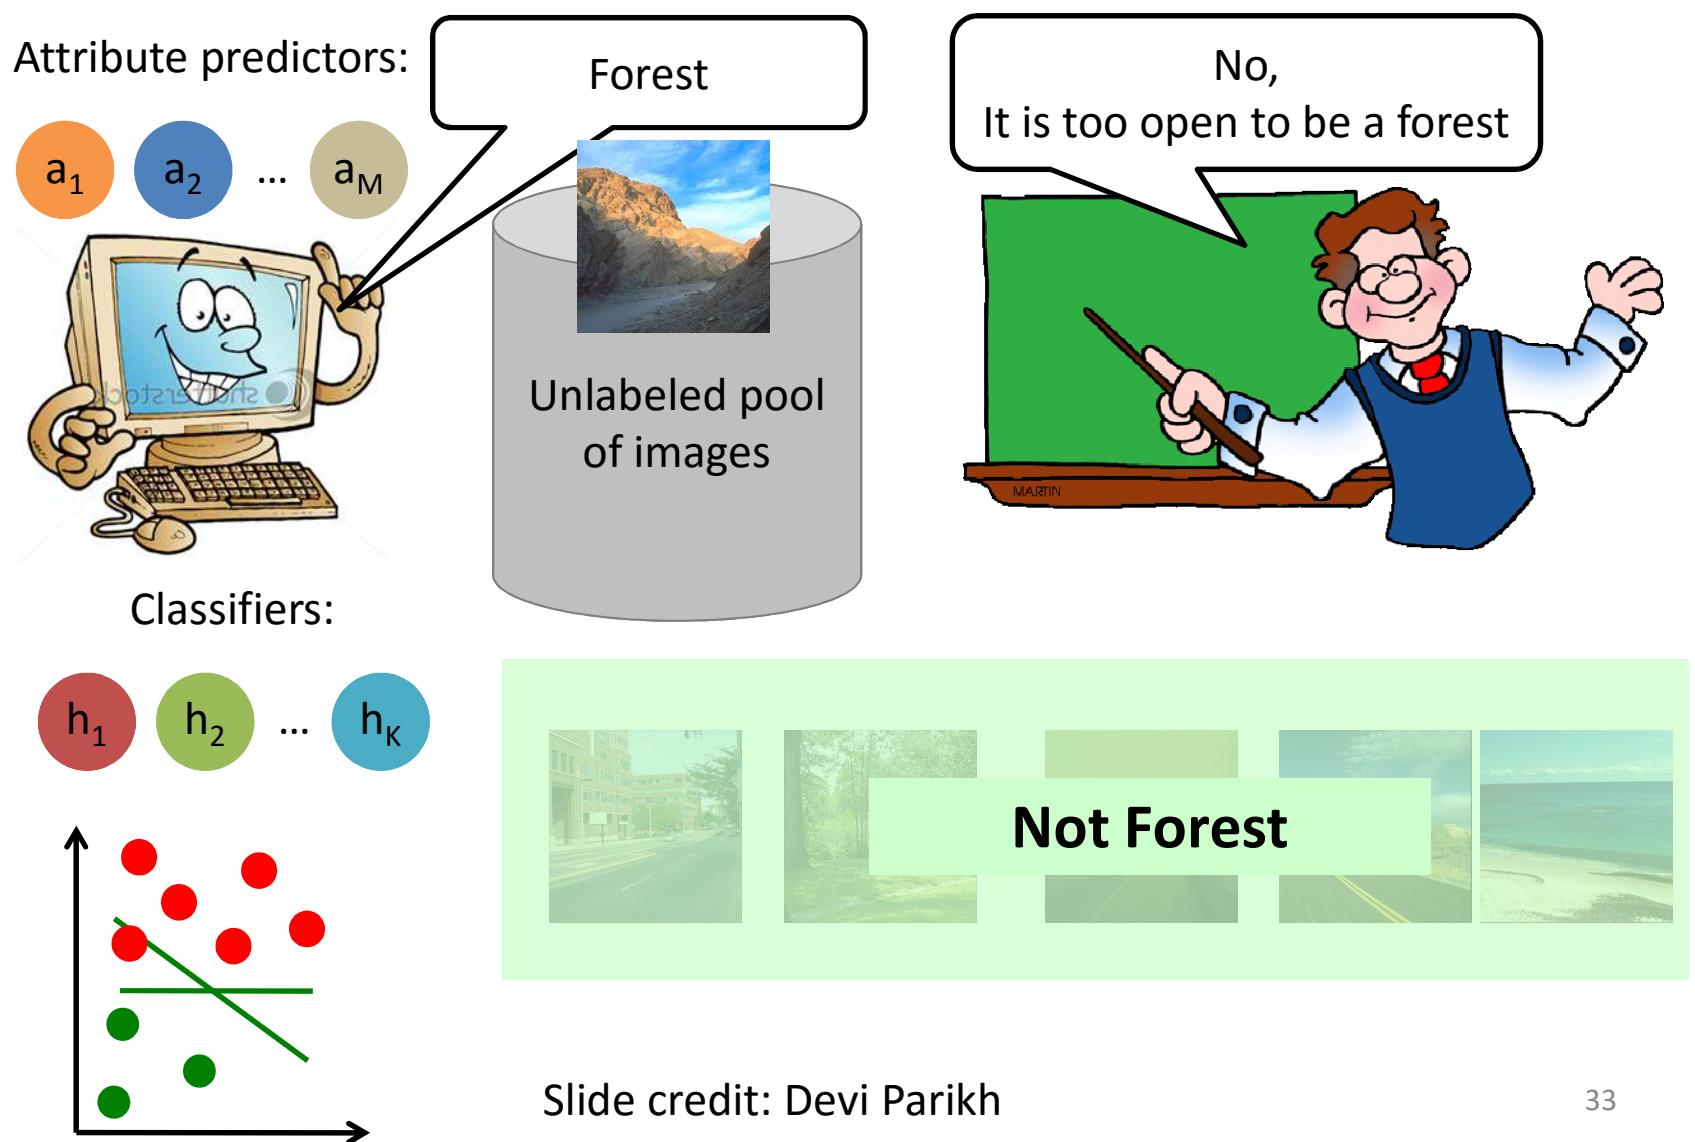
Not Forest



Openness

Slide credit: Devi Parikh

Attributes-based Feedback



I think this is a giraffe. What do you think?

No, its neck is too short for it to be a giraffe.

Ah! These must not be giraffes either then.

Ah! These must have longer necks than this image.

[Animals with even shorter necks]

[Giraffes]

MARTIN

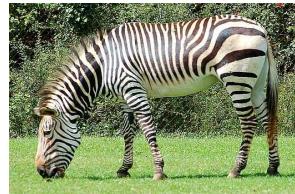
.....

.....



.....

Slide credit: Devi Parikh



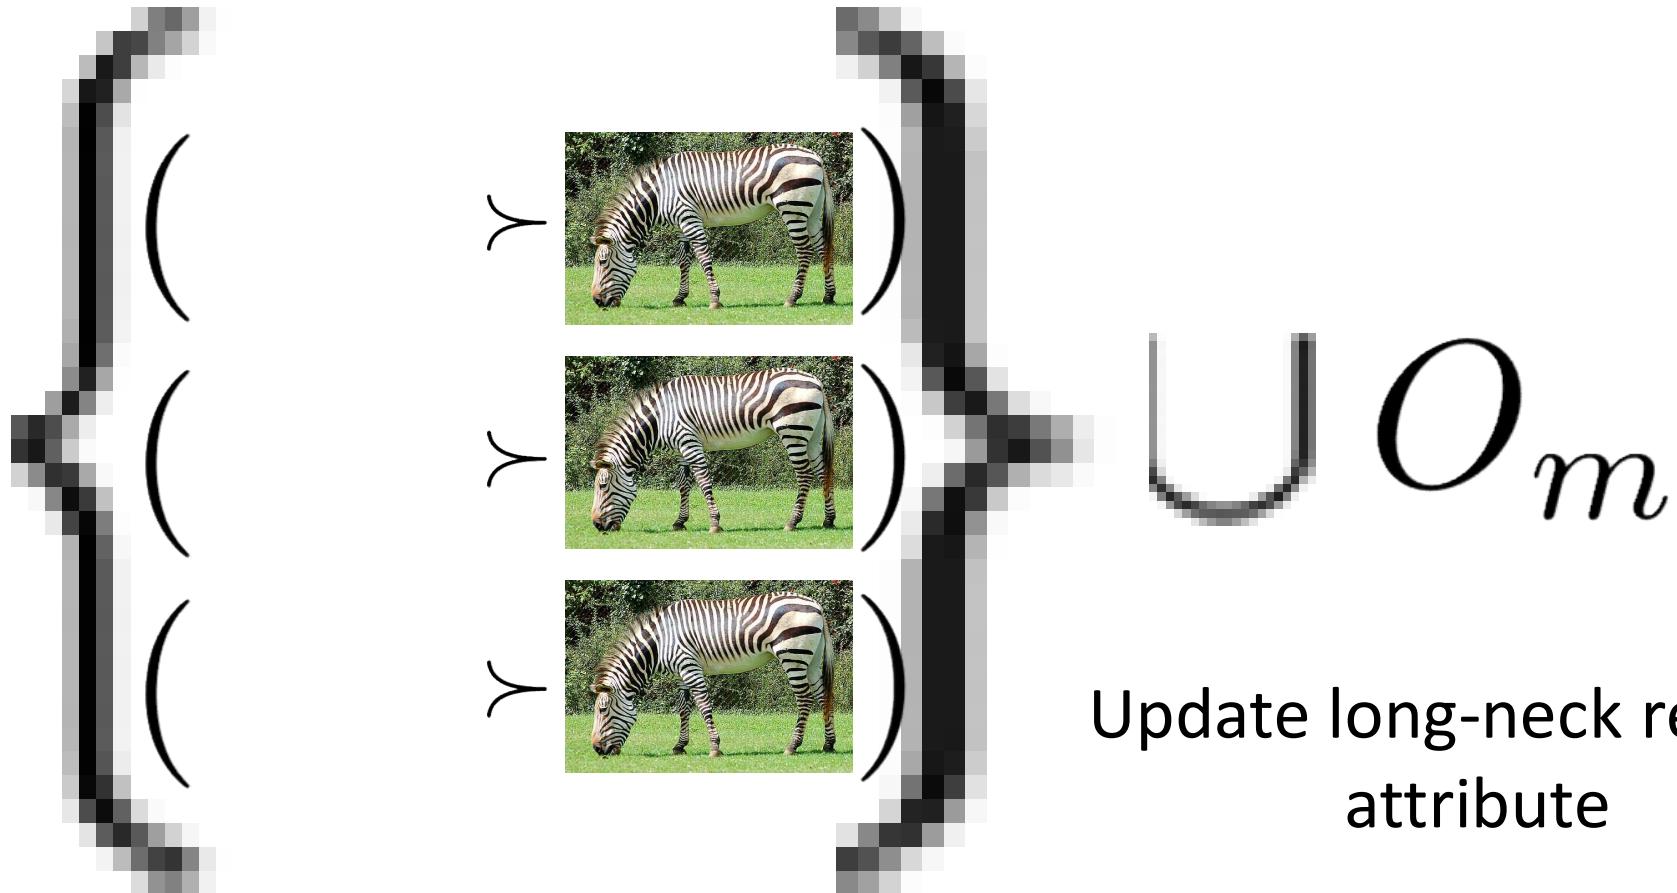
.....

Slide credit: Devi Parikh

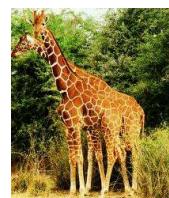


.....

Slide credit: Devi Parikh

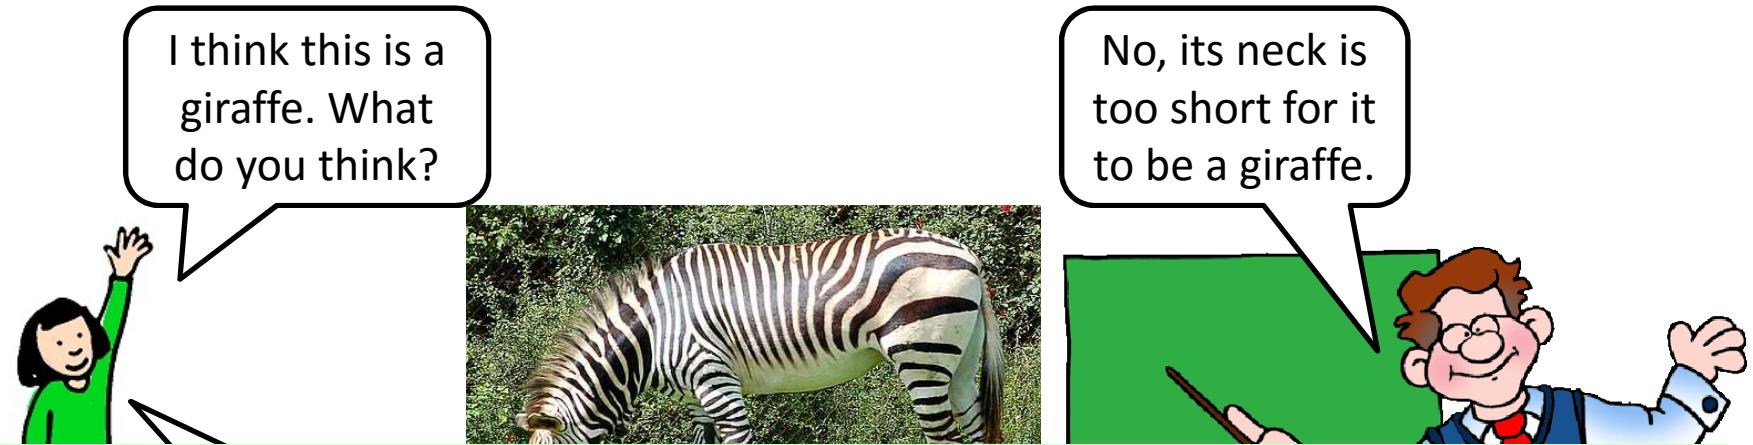


Update long-neck relative
attribute



.....

Slide credit: Devi Parikh



- Learn or update attribute models on-the-fly
- Start with an unlabeled pool of image and learn categories and attributes from scratch
- Actively select image for this form of feedback

Ah! These must have longer necks than this image.



.....

Questions?