

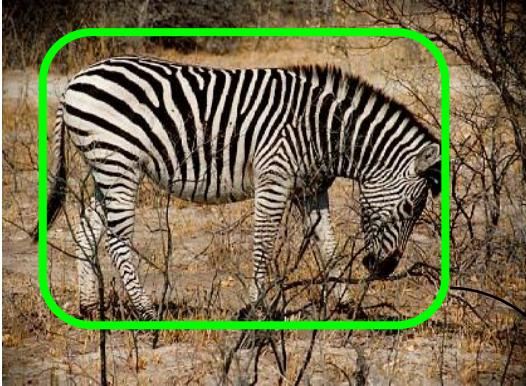
# Classifier based methods for Object Recognition

CMP 719– Computer Vision

Pinar Duygulu

(Slide credits:

Kristen Grauman, Fei fei Li, Antonio Torralba, Hames Hays)



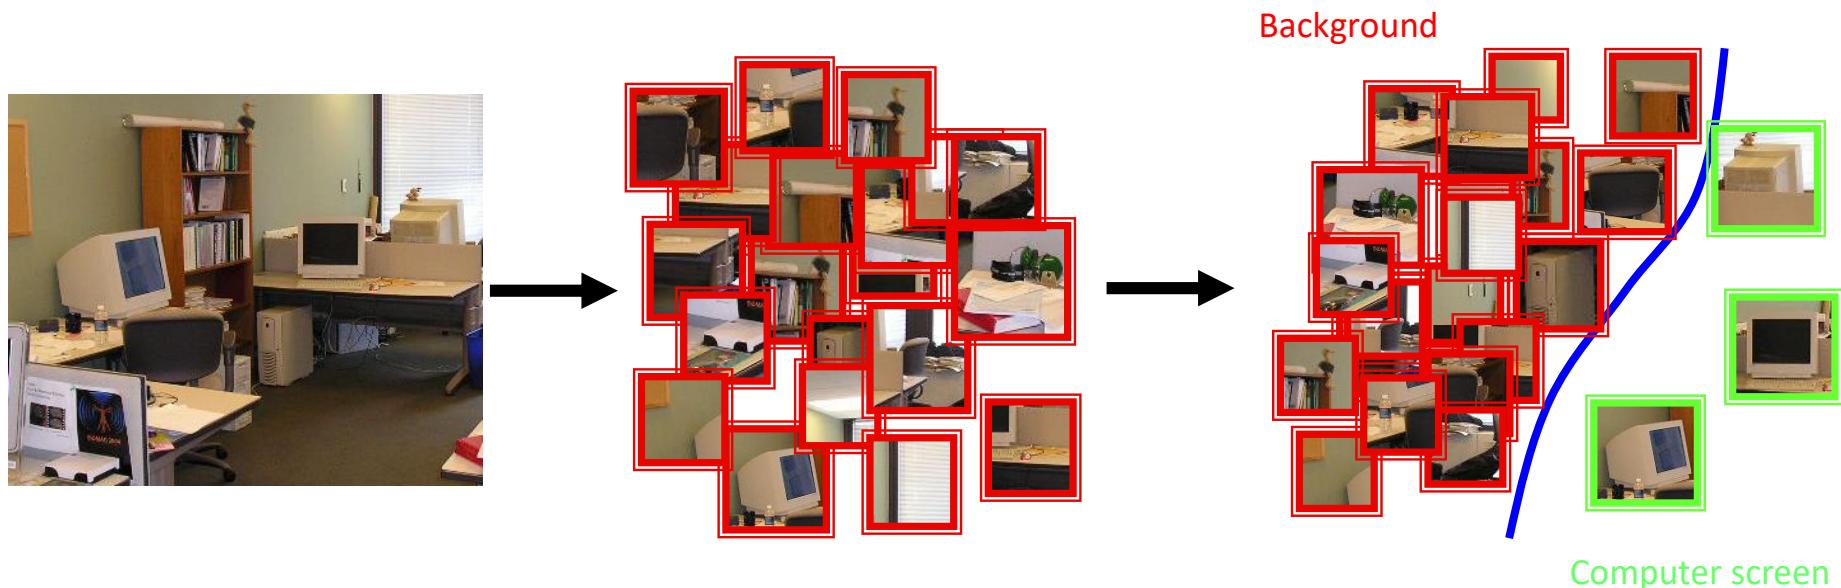
# Classifier-based methods

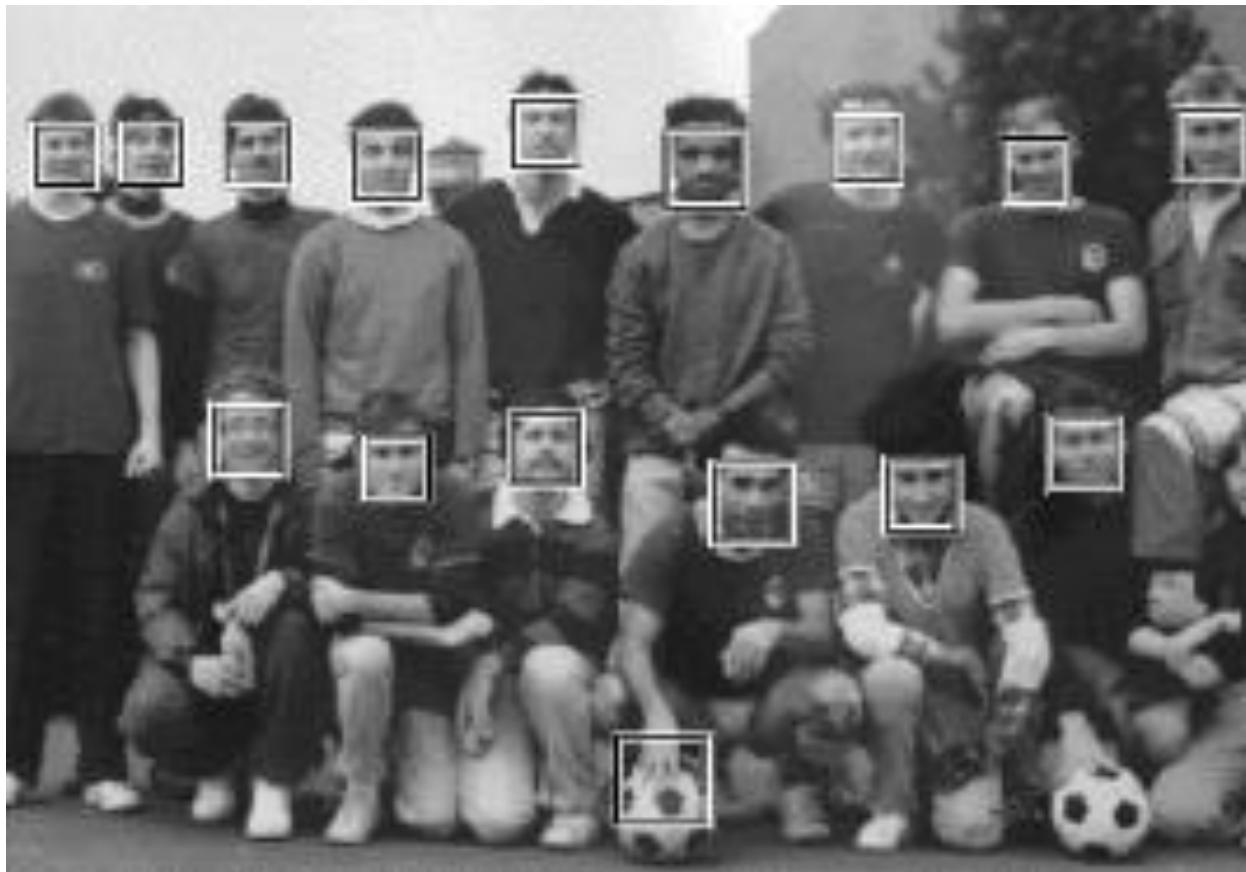
# Classifier based methods

Object detection and recognition is formulated as a classification problem.

The image is partitioned into a set of overlapping windows

... and a decision is taken at each window about if it contains a target object or not.





# Learning Models

## Training

Training  
Images

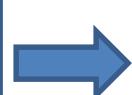


Image  
Features



Training  
Labels



Training



Learned  
model

## Testing



Image  
Features



Test Image

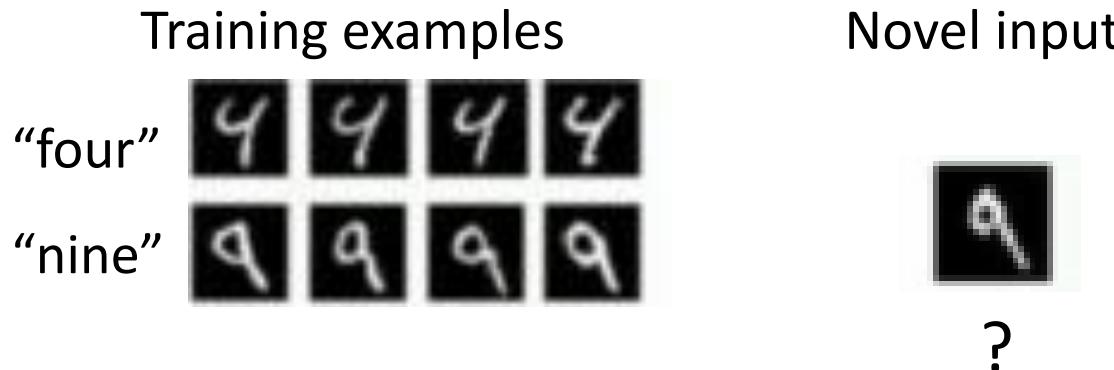
Learned  
model



Prediction

# Supervised classification

- Given a collection of *labeled* examples, come up with a function that will predict the labels of new examples.



- How good is some function we come up with to do the classification?
- Depends on
  - Mistakes made
  - Cost associated with the mistakes

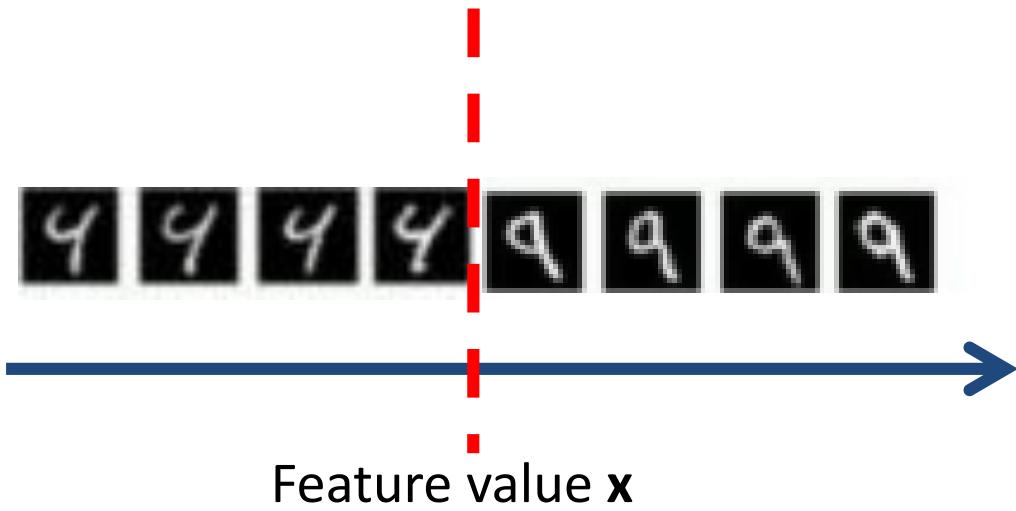
# Supervised classification

- Given a collection of *labeled* examples, come up with a function that will predict the labels of new examples.
- Consider the two-class (binary) decision problem
  - $L(4 \rightarrow 9)$ : Loss of classifying a 4 as a 9
  - $L(9 \rightarrow 4)$ : Loss of classifying a 9 as a 4
- **Risk** of a classifier  $s$  is expected loss:

$$R(s) = \Pr(4 \rightarrow 9 \mid \text{using } s)L(4 \rightarrow 9) + \Pr(9 \rightarrow 4 \mid \text{using } s)L(9 \rightarrow 4)$$

- We want to choose a classifier so as to minimize this total risk

# Supervised classification



Optimal classifier will minimize total risk.

At decision boundary, either choice of label yields same expected loss.

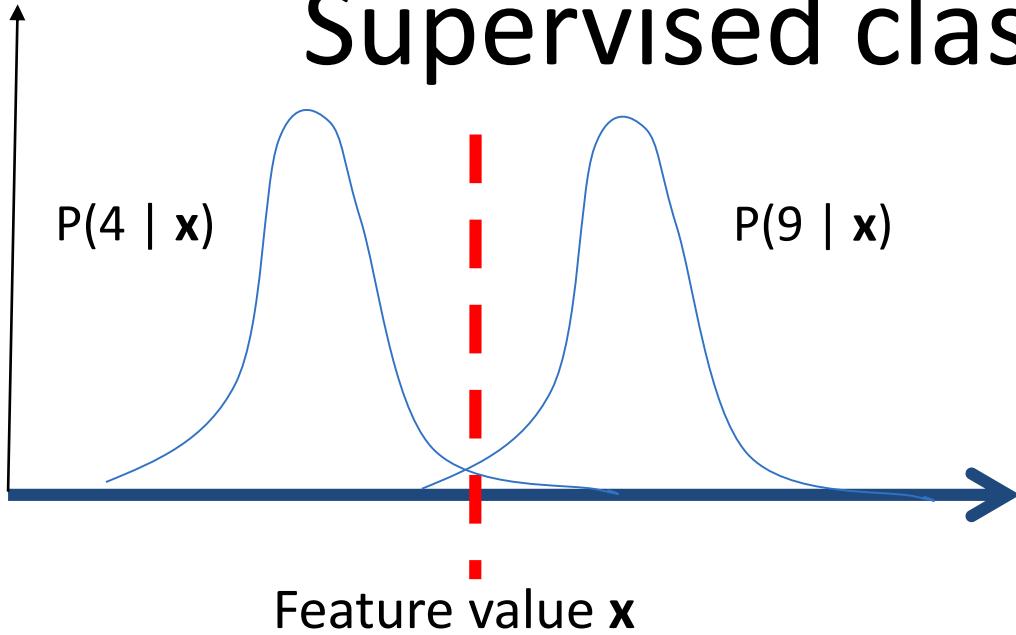
So, best decision boundary is at point  $x$  where

$$P(\text{class is } 9 \mid \mathbf{x}) L(9 \rightarrow 4) = P(\text{class is } 4 \mid \mathbf{x}) L(4 \rightarrow 9)$$

To classify a new point, choose class with lowest expected loss;  
i.e., choose “four” if

$$P(4 \mid \mathbf{x}) L(4 \rightarrow 9) > P(9 \mid \mathbf{x}) L(9 \rightarrow 4)$$

# Supervised classification



Optimal classifier will minimize total risk.

At decision boundary, either choice of label yields same expected loss.

So, best decision boundary is at point  $\mathbf{x}$  where

$$P(\text{class is } 9 \mid \mathbf{x}) L(9 \rightarrow 4) = P(\text{class is } 4 \mid \mathbf{x}) L(4 \rightarrow 9)$$

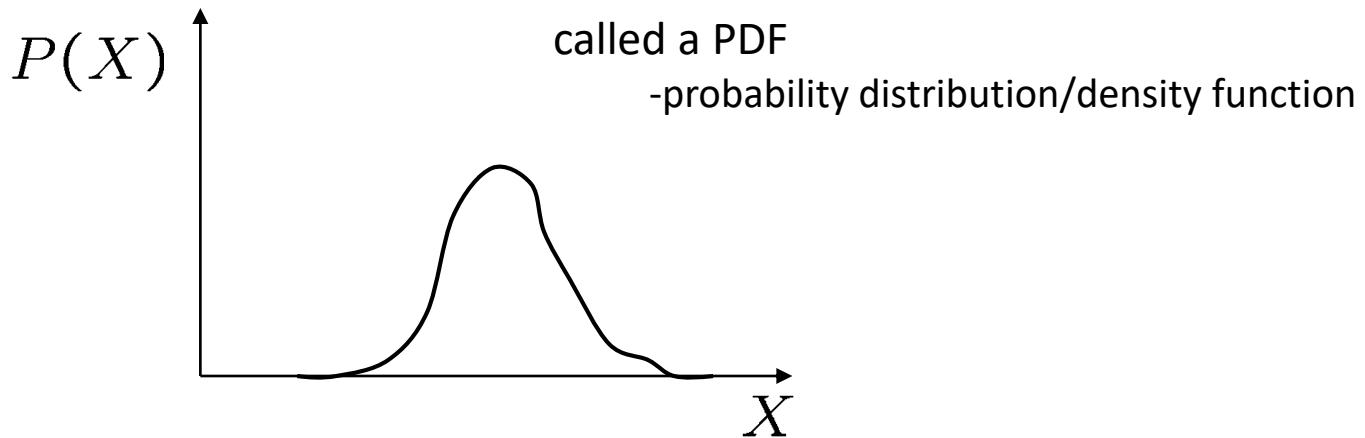
To classify a new point, choose class with lowest expected loss;  
i.e., choose “four” if

$$P(4 \mid \mathbf{x}) L(4 \rightarrow 9) > P(9 \mid \mathbf{x}) L(9 \rightarrow 4)$$

*How to evaluate these probabilities?*

# Probability

- Basic probability
  - $X$  is a random variable
  - $P(X)$  is the probability that  $X$  achieves a certain value



- $0 \leq P(X) \leq 1$
- $\int_{-\infty}^{\infty} P(X)dX = 1$     or     $\sum P(X) = 1$ 
  - continuous  $X$
  - discrete  $X$

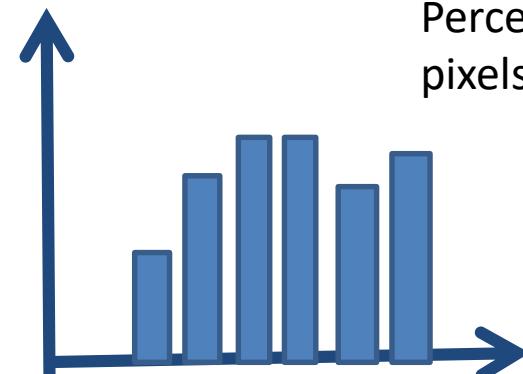
- Conditional probability:  $P(X | Y)$ 
  - probability of  $X$  given that we already know  $Y$

# Example: learning skin colors

- We can represent a class-conditional density using a histogram (a “non-parametric” distribution)

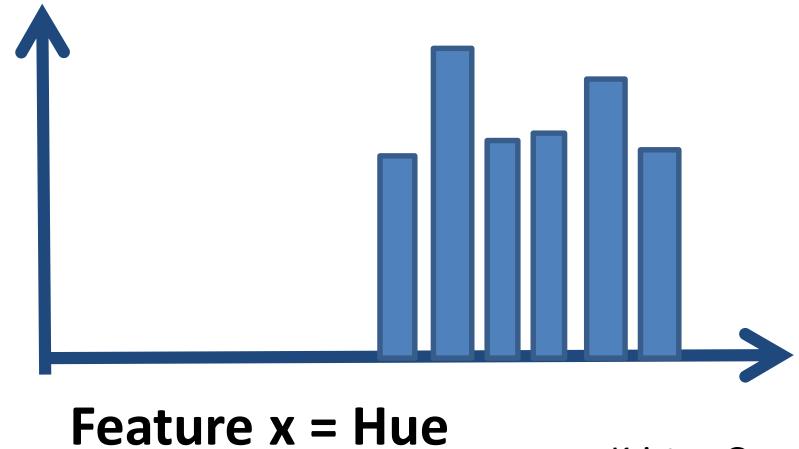


$P(x|\text{skin})$



Percentage of skin pixels in each bin

$P(x|\text{not skin})$



Feature x = Hue

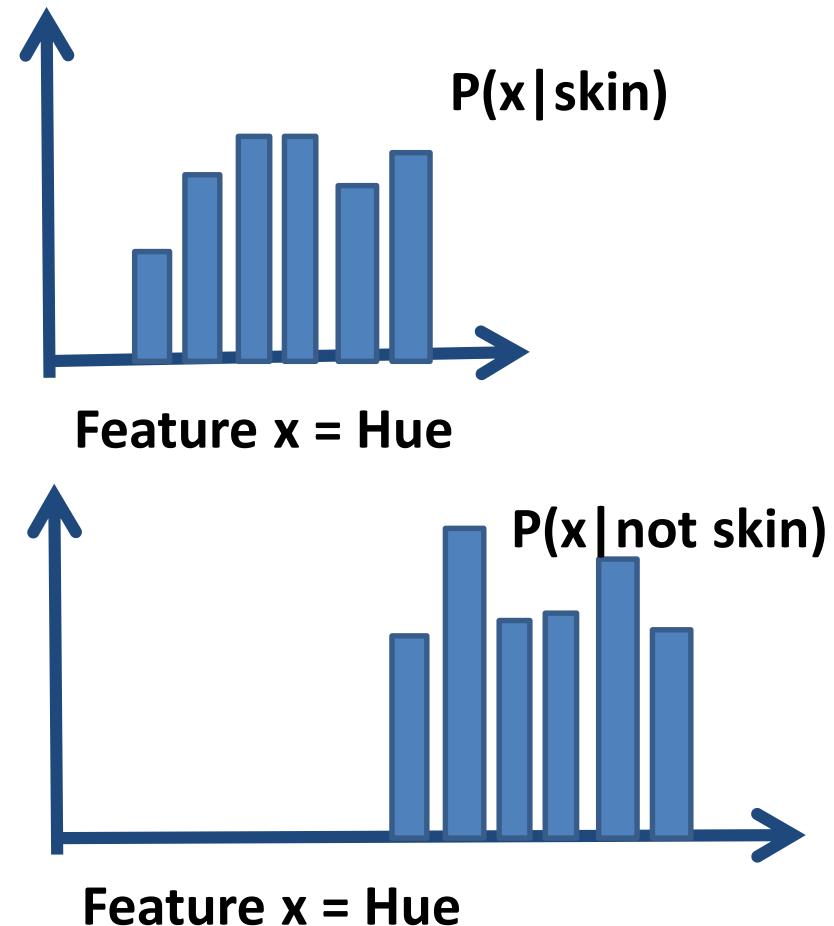
# Example: learning skin colors

- We can represent a class-conditional density using a histogram (a “non-parametric” distribution)



Now we get a new image, and want to label each pixel as skin or non-skin.

*What's the probability we care about to do skin detection?*



# Bayes rule

$$P(\text{skin} | x) = \frac{P(x | \text{skin}) P(\text{skin})}{P(x)}$$

posterior                      likelihood                      prior

$$P(\text{skin} | x) \propto P(x | \text{skin}) P(\text{skin})$$

# Example: classifying skin pixels

Now for every pixel in a new image, we can estimate probability that it is generated by skin.



Brighter pixels →  
higher probability  
of being skin

Classify pixels based on these probabilities

- if  $p(\text{skin}|\mathbf{x}) > \theta$ , classify as skin
- if  $p(\text{skin}|\mathbf{x}) < \theta$ , classify as not skin

# Example: classifying skin pixels

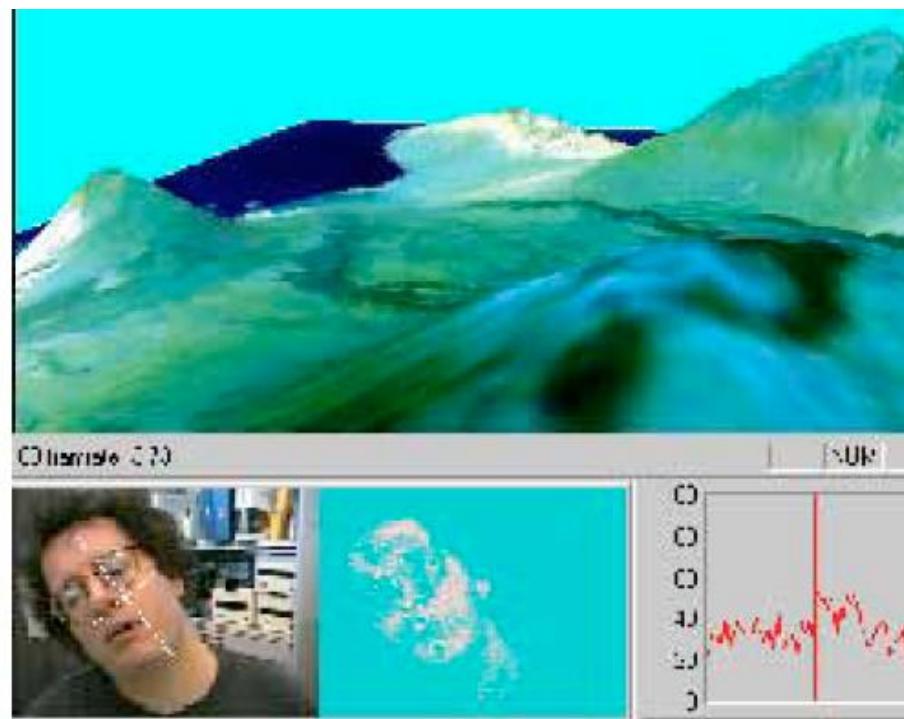


**Figure 6:** A video image and its flesh probability image

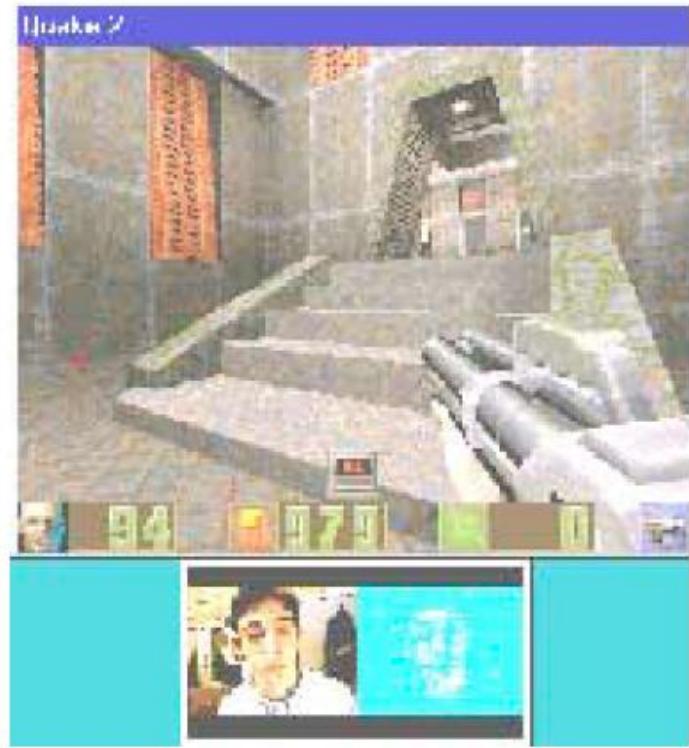


**Figure 7:** Orientation of the flesh probability distribution marked on the source video image

# Example: classifying skin pixels



**Figure 13:** CAMSHIFT-based face tracker used to track a person's face over a 3D graphic's model of Hawaii



**Figure 12:** CAMSHIFT-based face tracker used to play Quake 2 hands free by inserting control variables into the mouse queue

Using skin color-based face detection and pose estimation as a video-based interface

# Supervised classification

- Want to minimize the expected misclassification
- Two general strategies
  - Use the training data to build representative probability model; separately model class-conditional densities and priors (*generative*)
  - Directly construct a good decision boundary, model the posterior (*discriminative*)

# Discriminative classifiers for image recognition

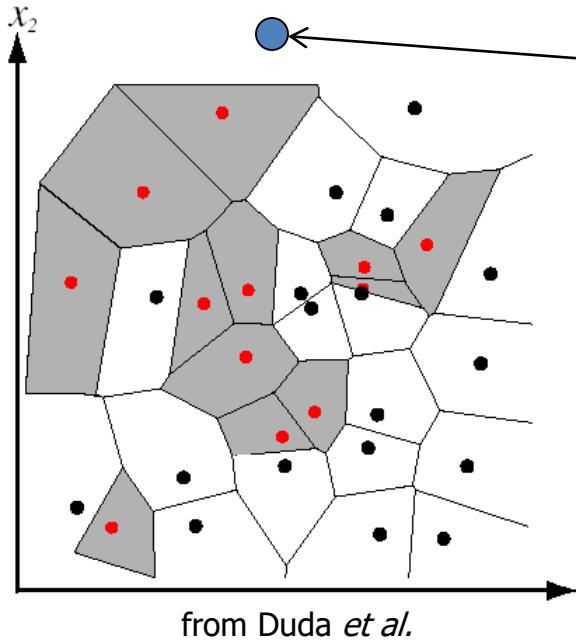
- nearest neighbors (+ scene match app)
- support vector machines (+ gender, person app)

# Nearest Neighbor classification

- Assign label of nearest training data point to each test data point

Black = negative

Red = positive



Novel test example

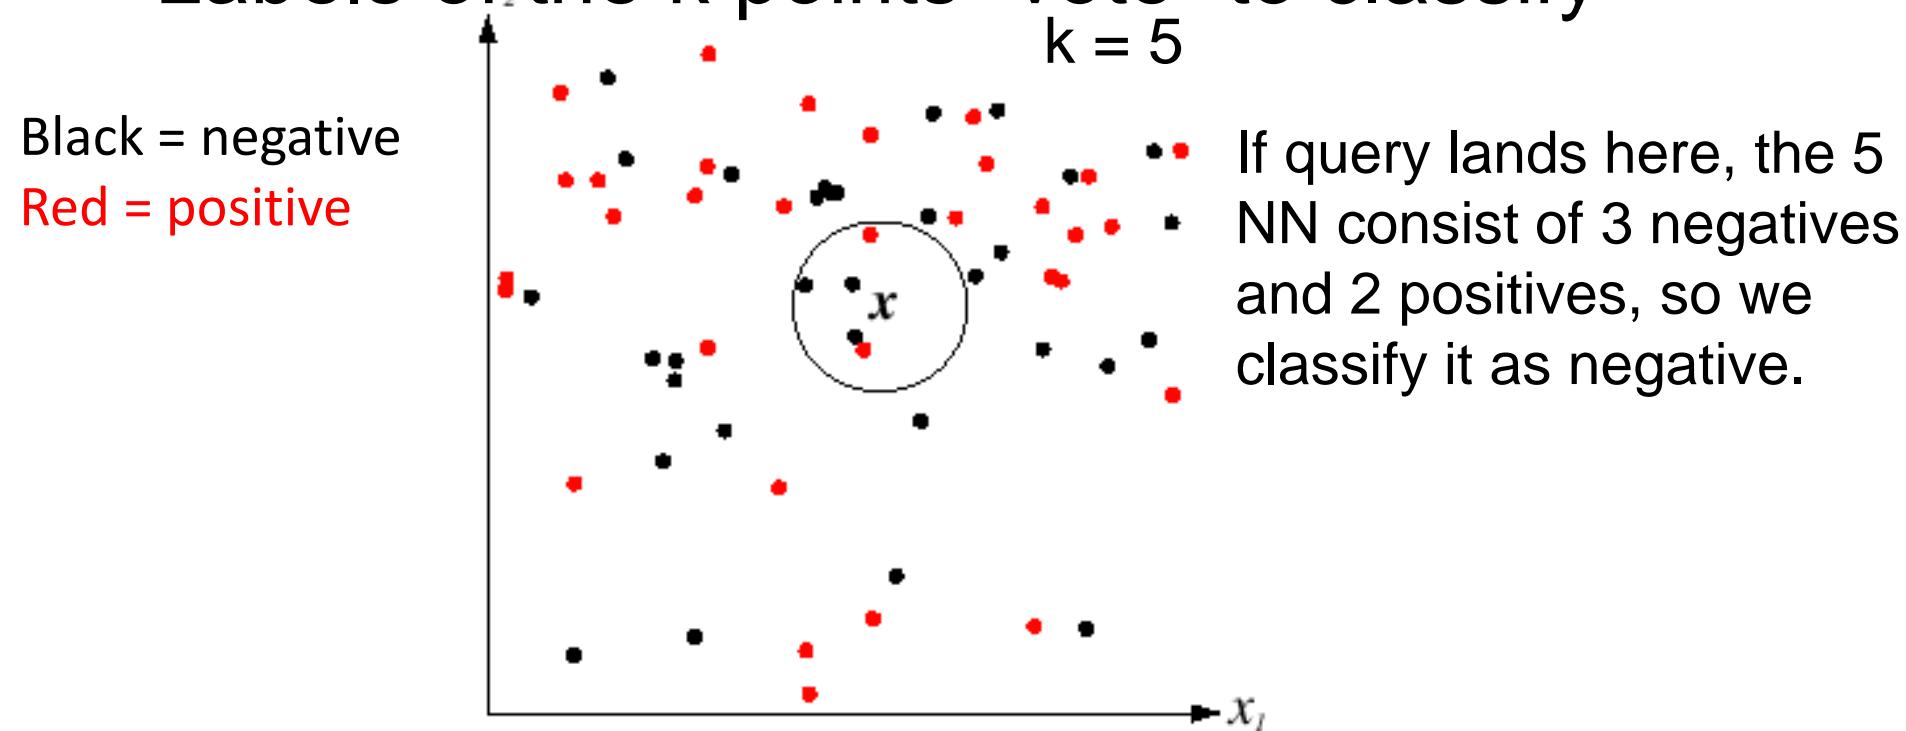
Closest to a  
**positive** example  
from the training  
set, so classify it  
as positive.

from Duda *et al.*

Voronoi partitioning of feature space  
for 2-category 2D data

# K-Nearest Neighbors classification

- For a new point, find the  $k$  closest points from training data
- Labels of the  $k$  points “vote” to classify



A nearest neighbor  
recognition example

# Where in the World?



[Hays and Efros. **im2gps**: Estimating Geographic Information from a Single Image.  
CVPR 2008.]

Slides: James Hays

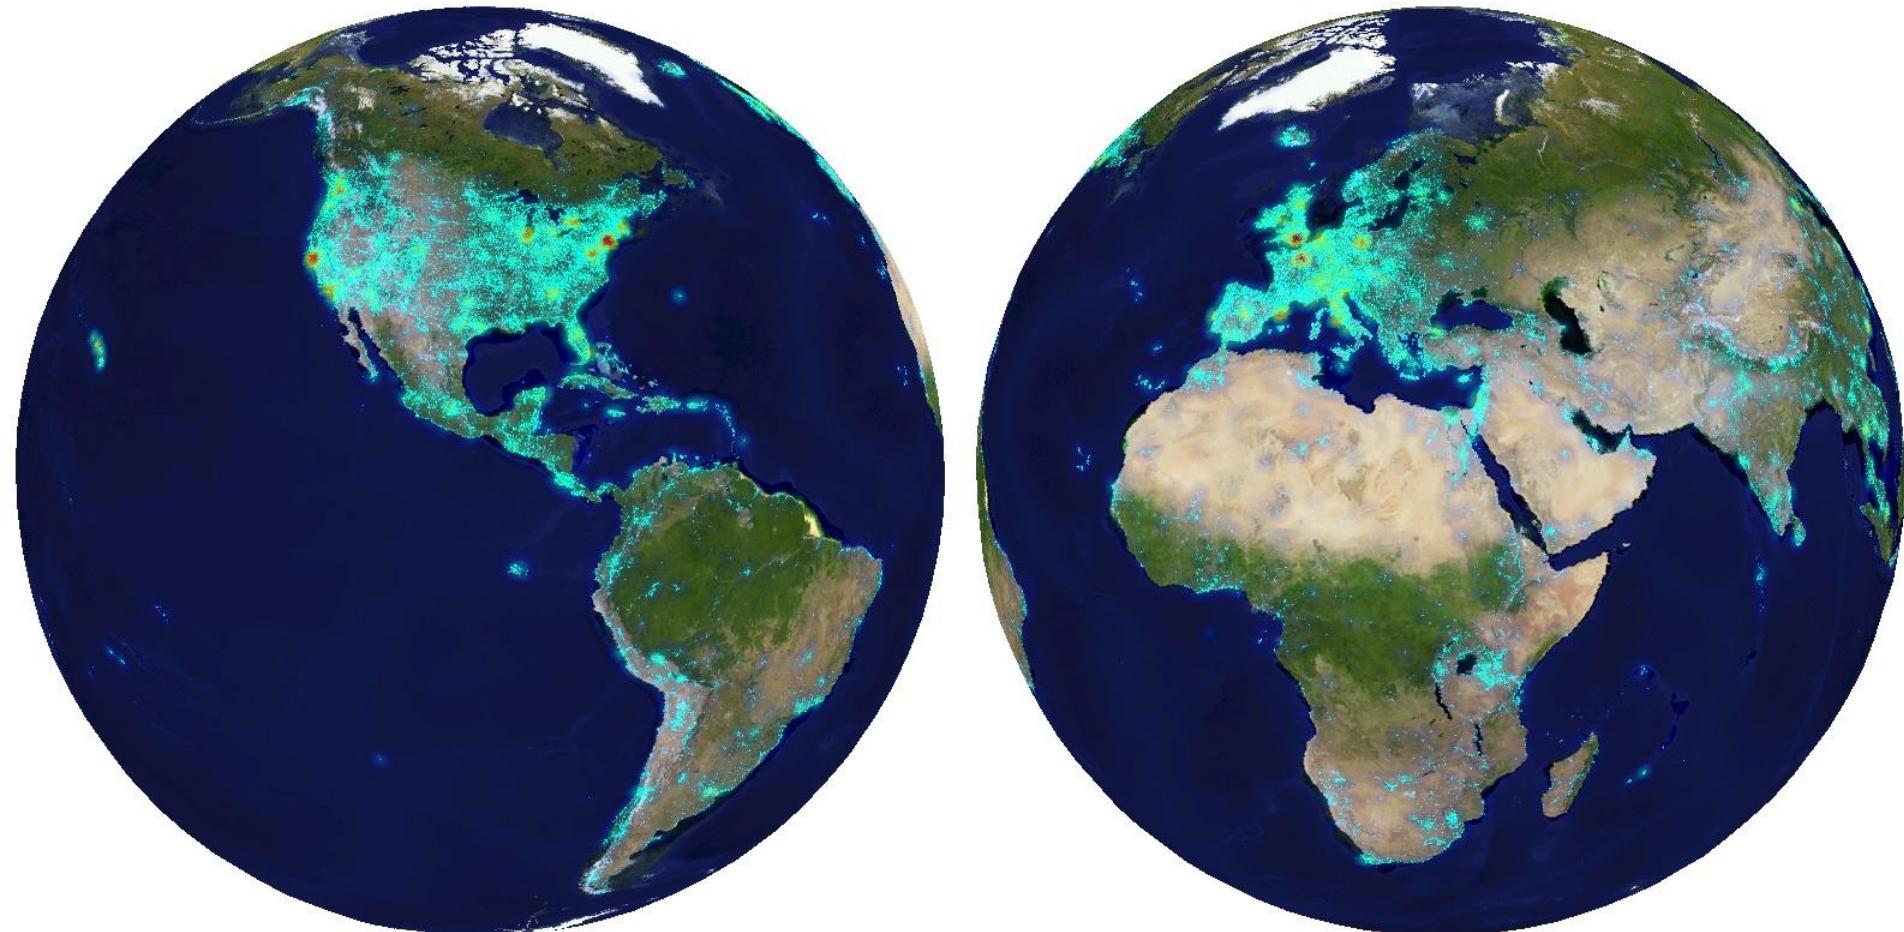
# Where in the World?



# Where in the World?



6+ million geotagged photos  
by 109,788 photographers



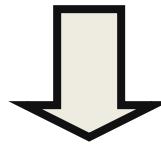
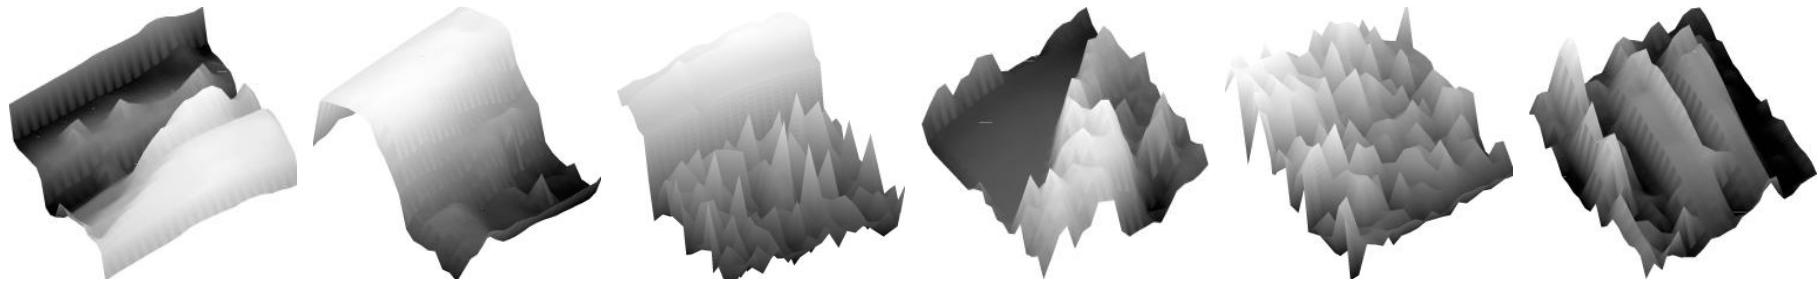
Annotated by Flickr users

Slides: James Hays

# Which scene properties are relevant?

# Spatial Envelope Theory of Scene Representation

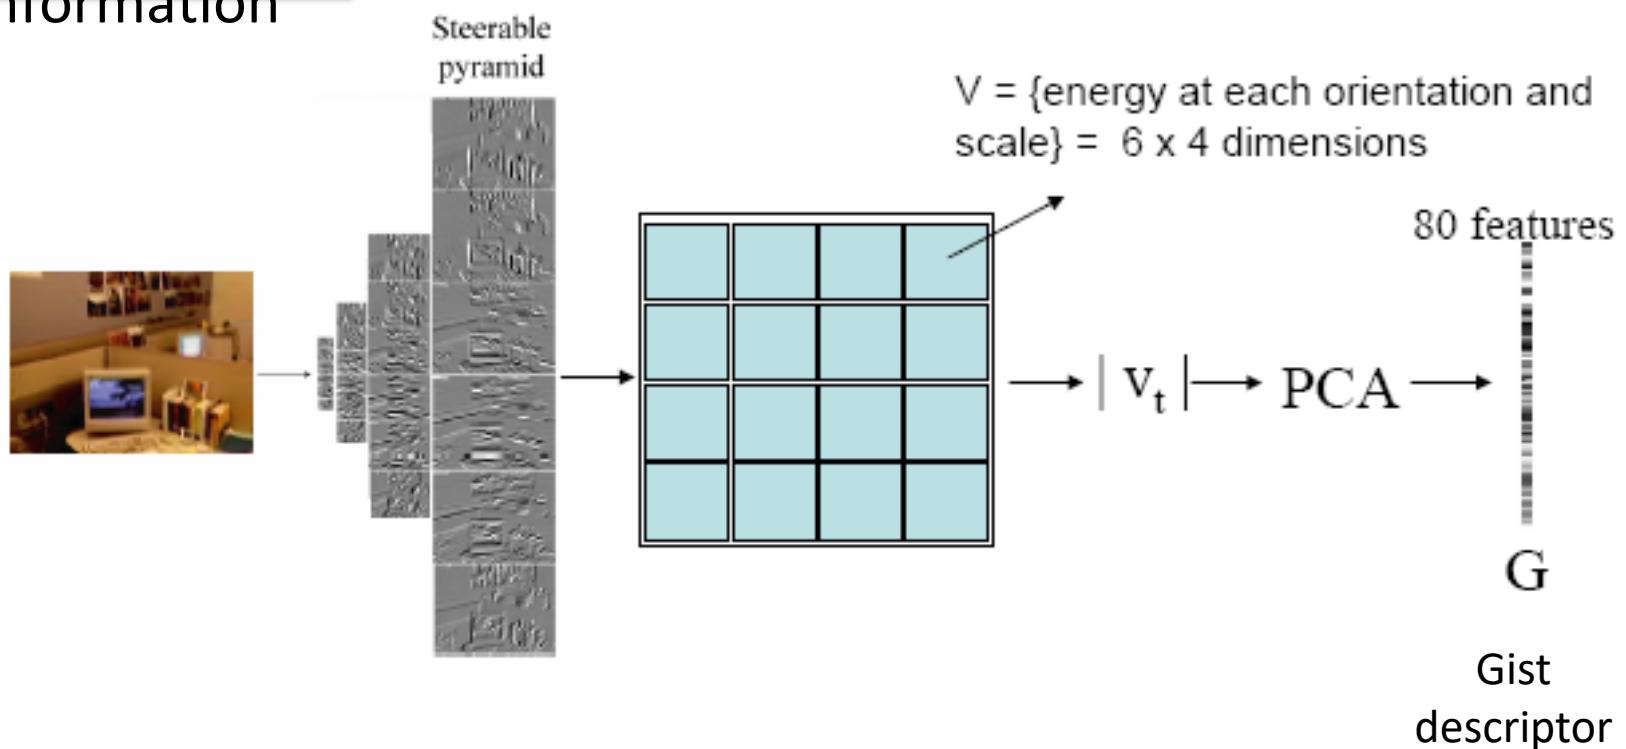
Oliva & Torralba (2001)



A scene is a single surface that can be represented by global (statistical) descriptors

# Global texture: capturing the “Gist” of the scene

Capture global image properties while keeping some spatial information

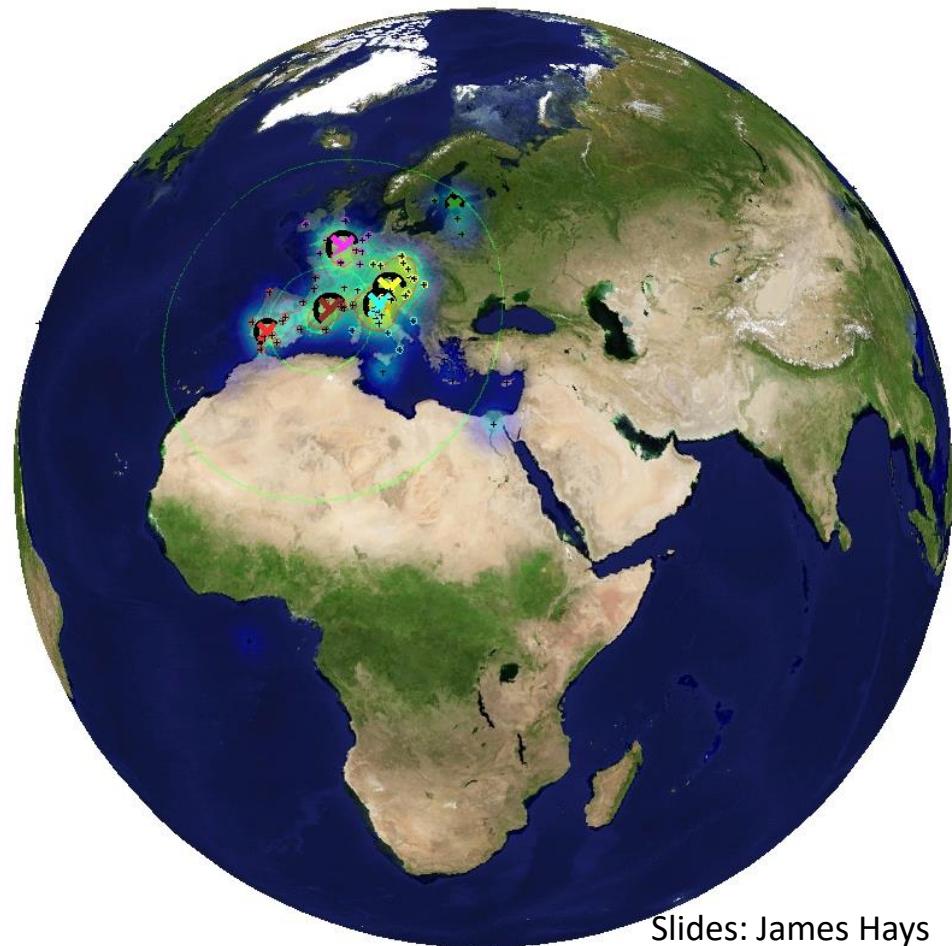


# Which scene properties are relevant?

- **Gist scene descriptor**
- **Color Histograms** - L\*A\*B\* 4x14x14 histograms
- **Texton Histograms** – 512 entry, filter bank based
- **Line Features** – Histograms of straight line stats

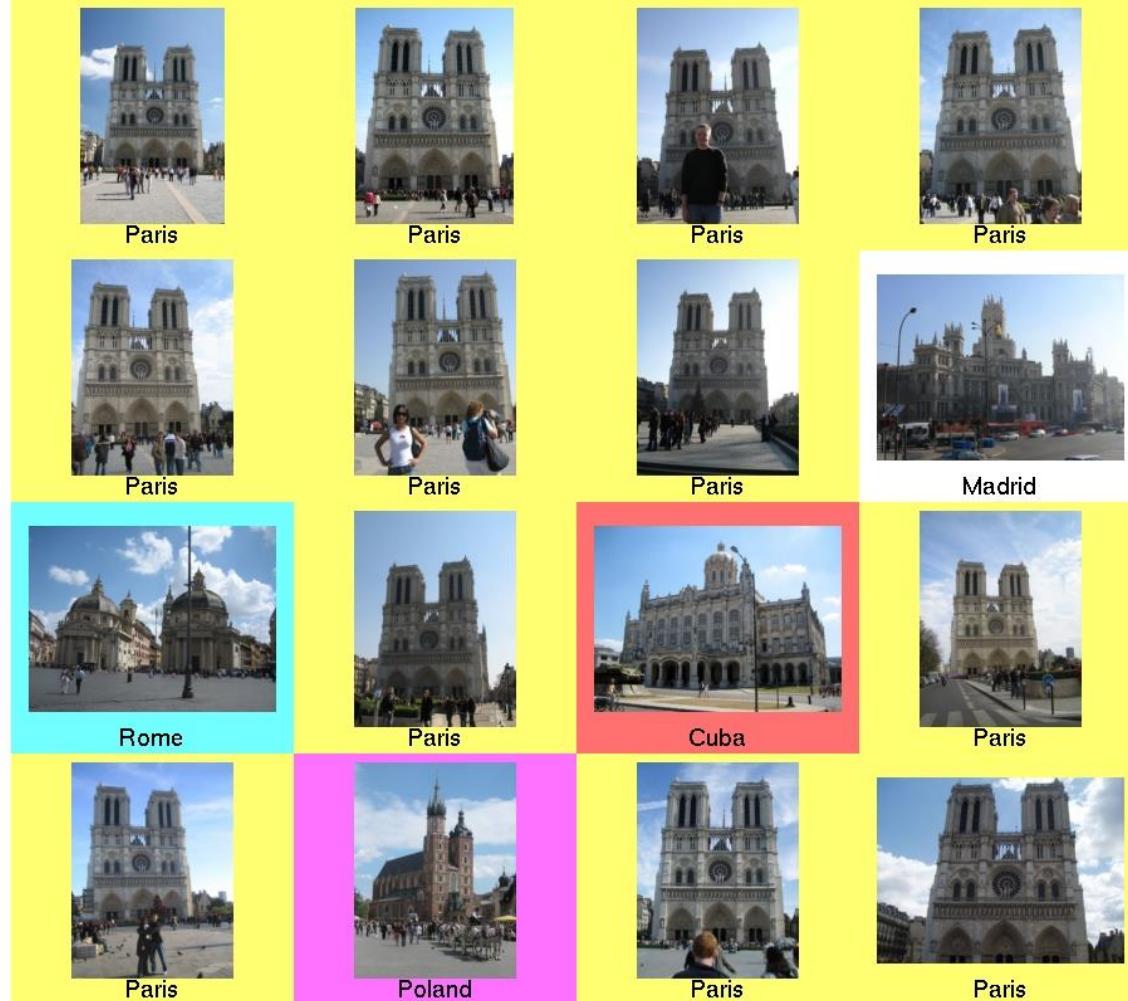
# Scene Matches

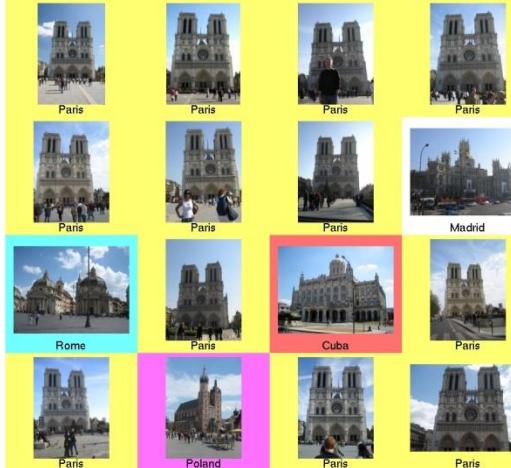




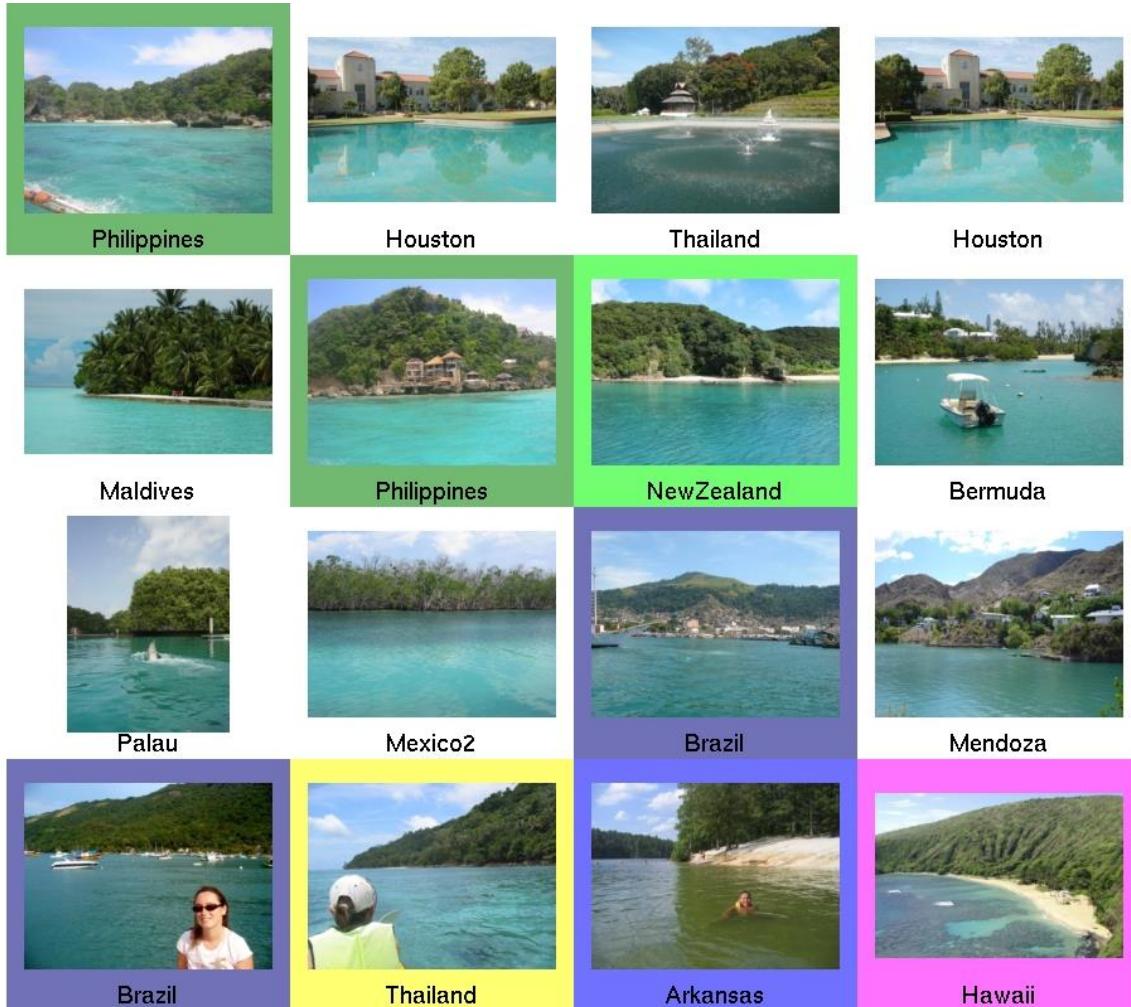
Slides: James Hays

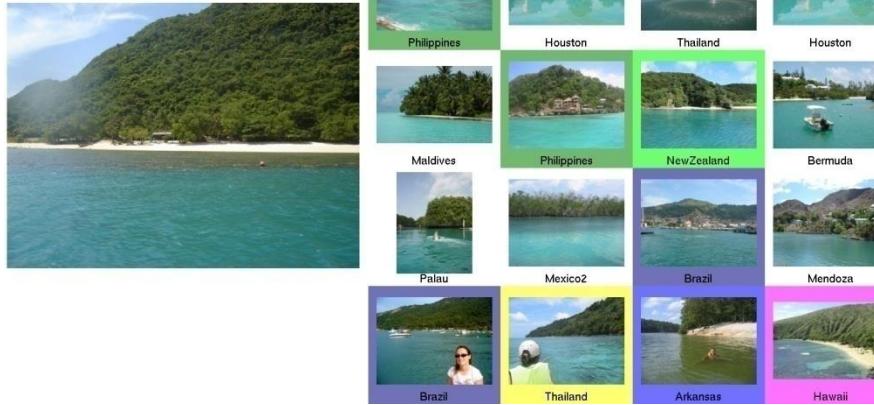
# Scene Matches



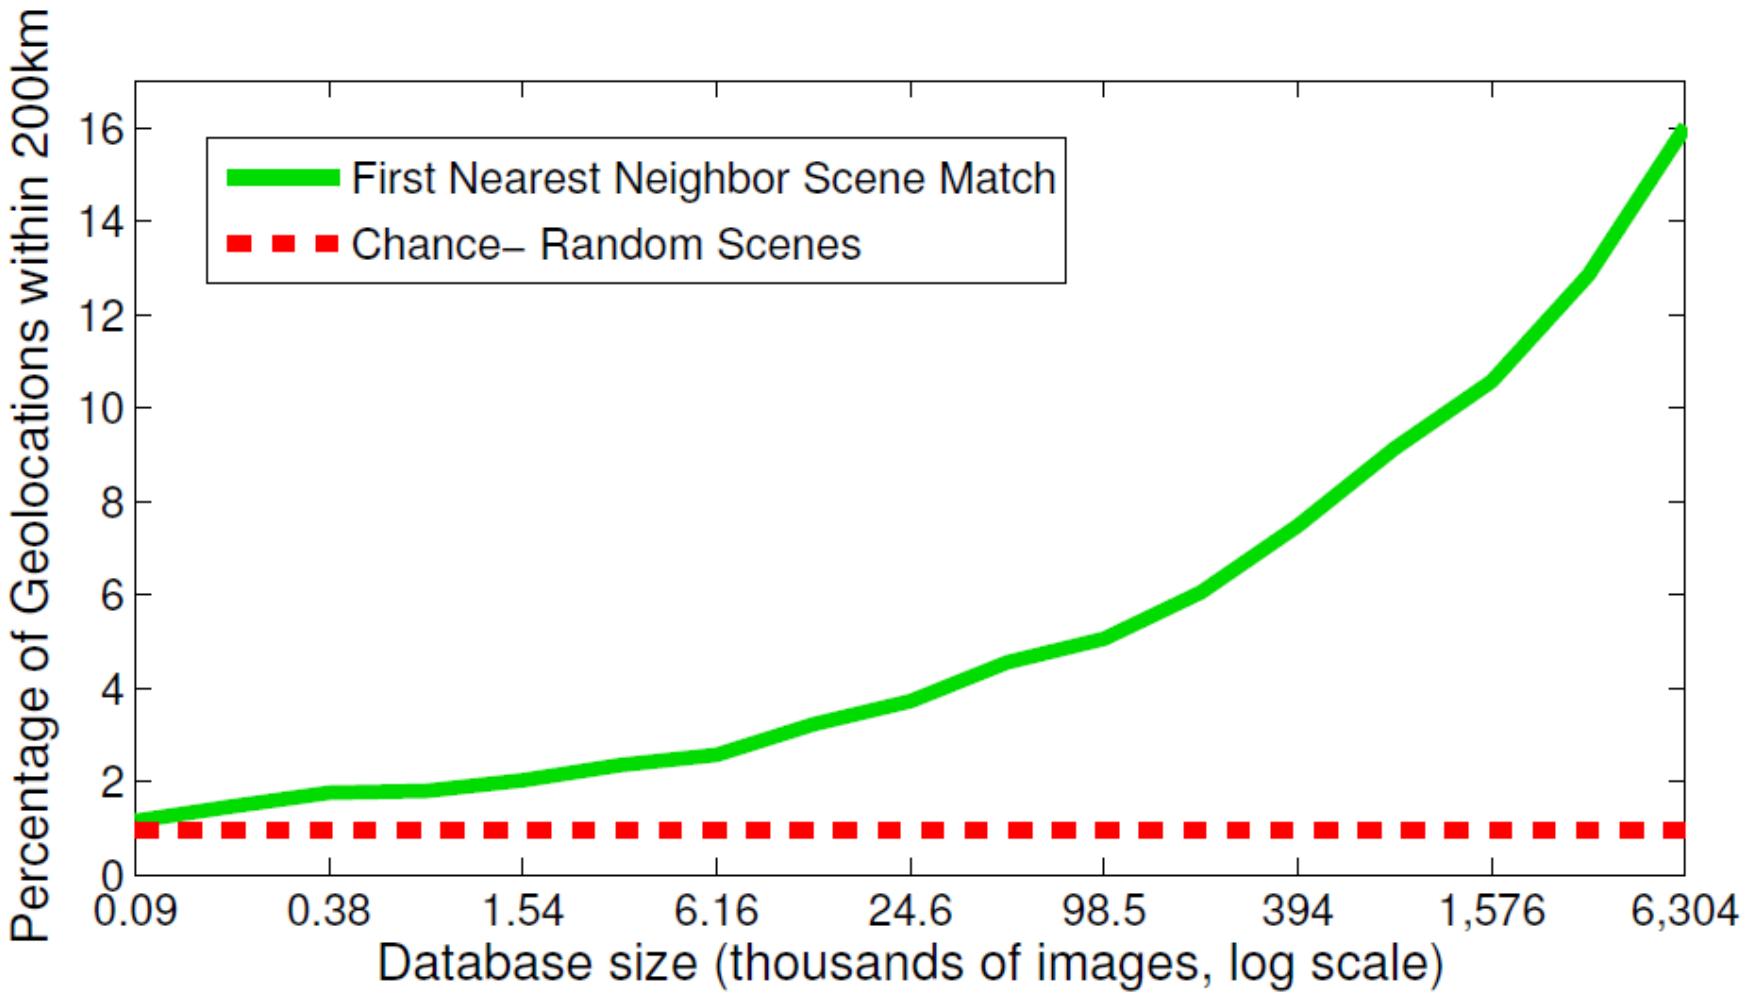


# Scene Matches





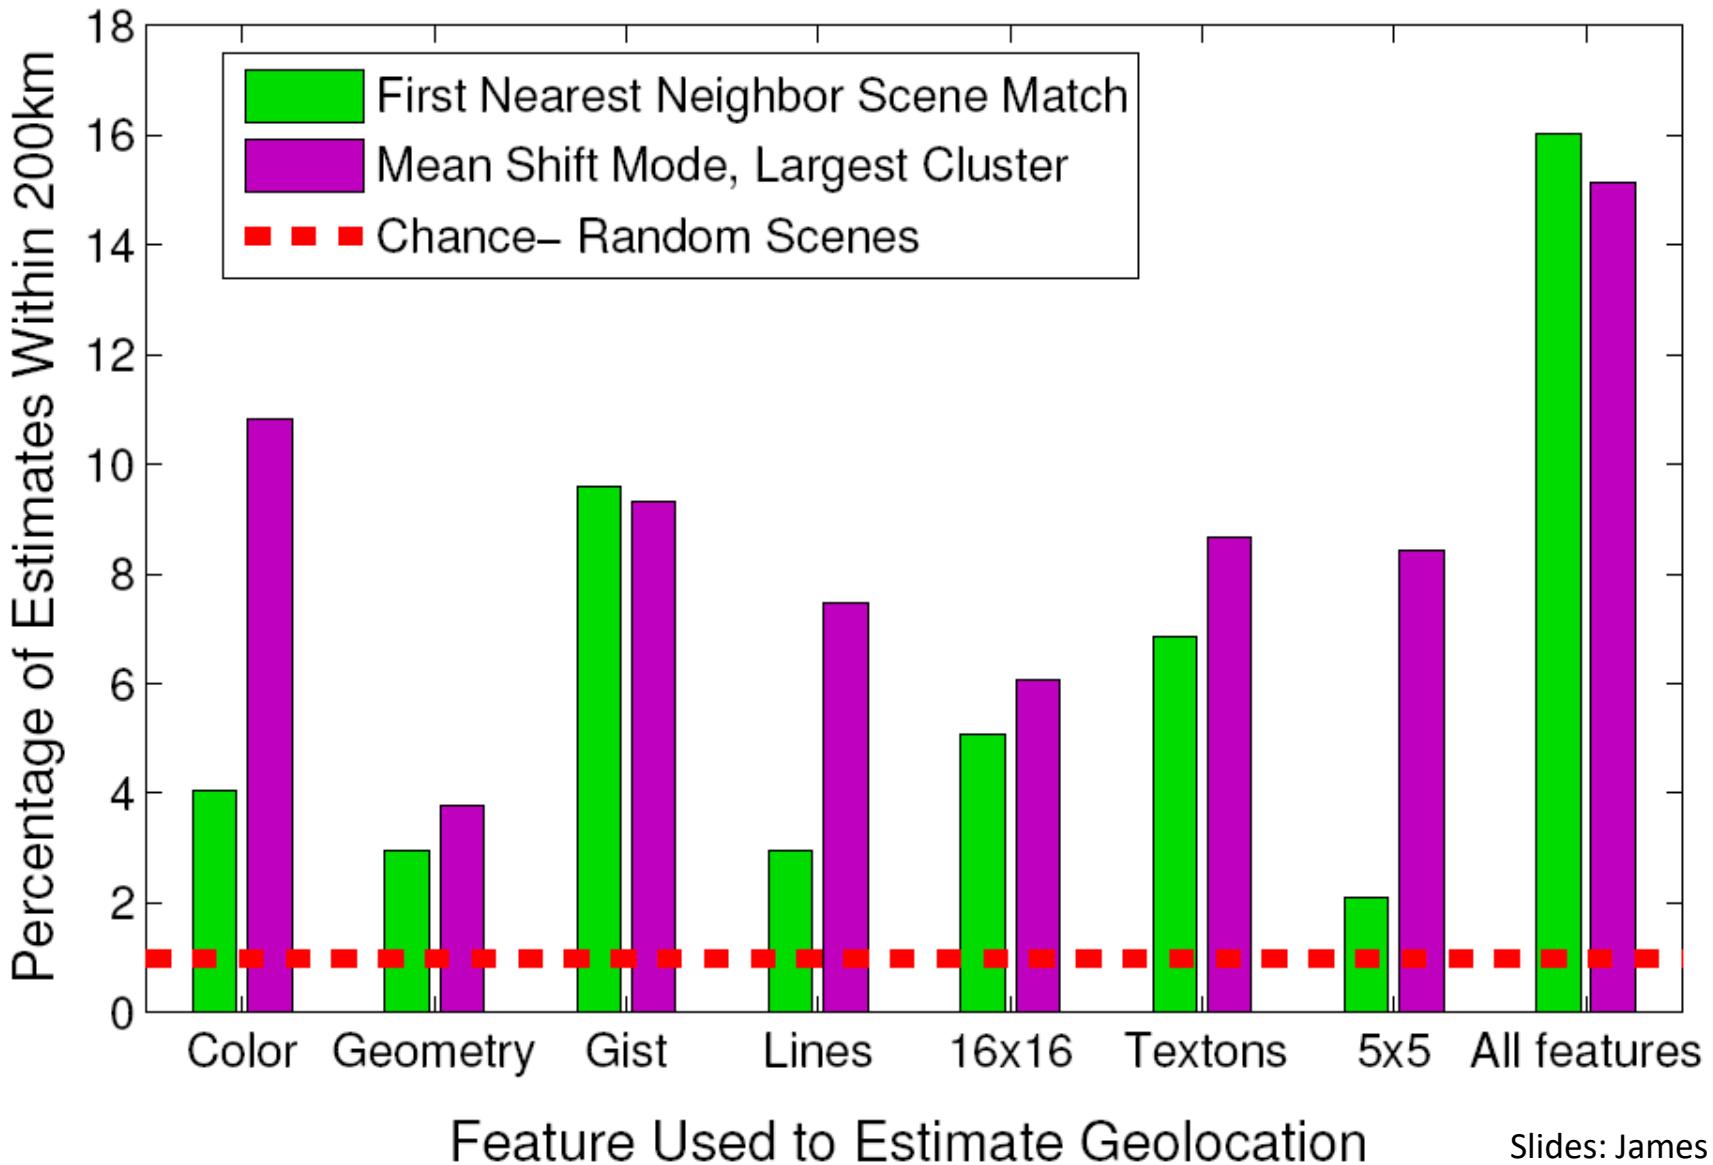
# The Importance of Data



[Hays and Efros. **im2gps**: Estimating Geographic Information from a Single Image. CVPR 2008.]

Slides: James Hays

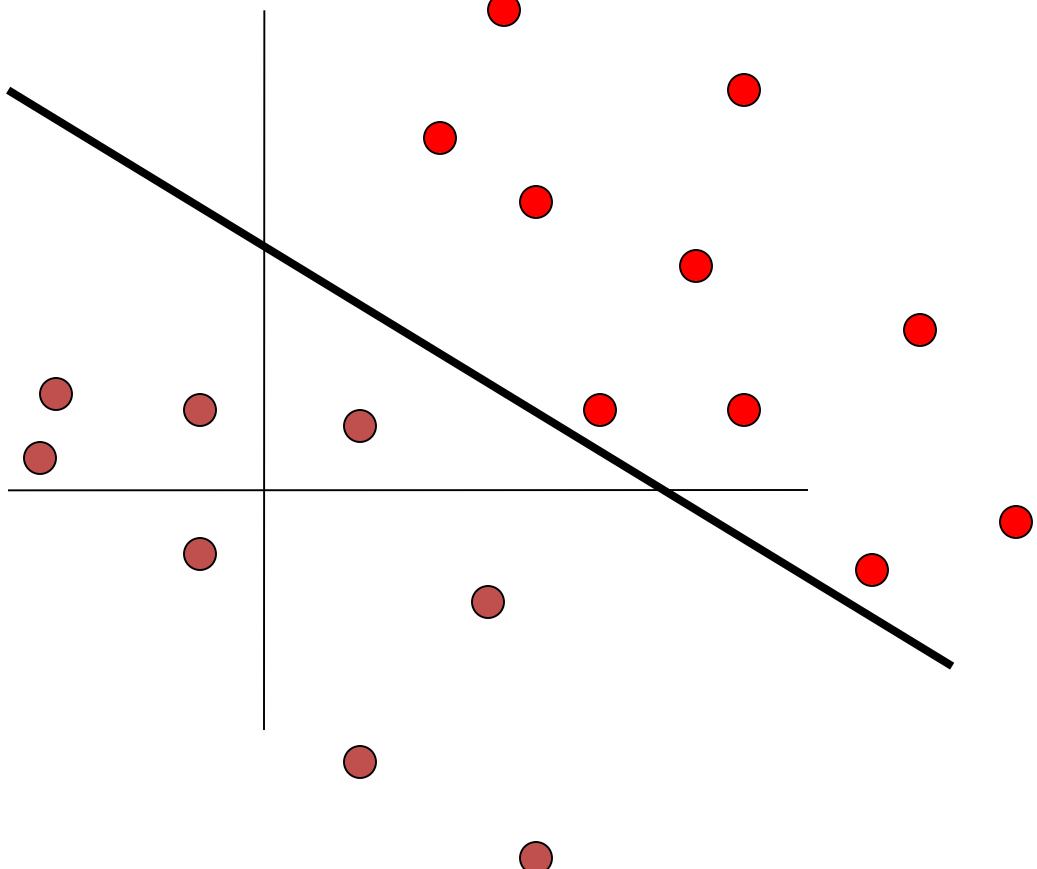
# Feature Performance



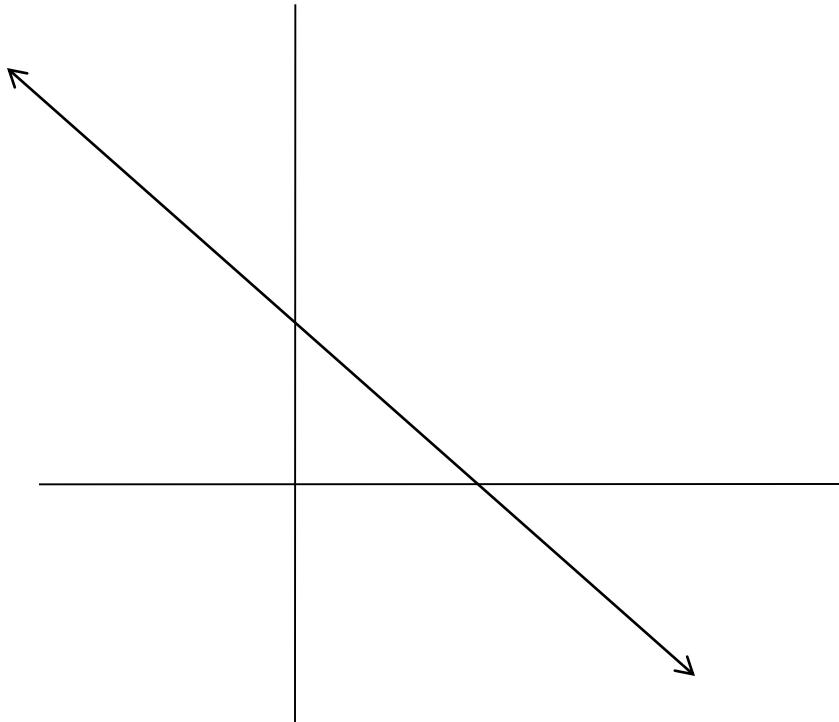
# Nearest neighbors: pros and cons

- **Pros:**
  - Simple to implement
  - Flexible to feature / distance choices
  - Naturally handles multi-class cases
  - Can do well in practice with enough representative data
- **Cons:**
  - Large search problem to find nearest neighbors
  - Storage of data
  - Must know we have a meaningful distance function

# Linear classifiers



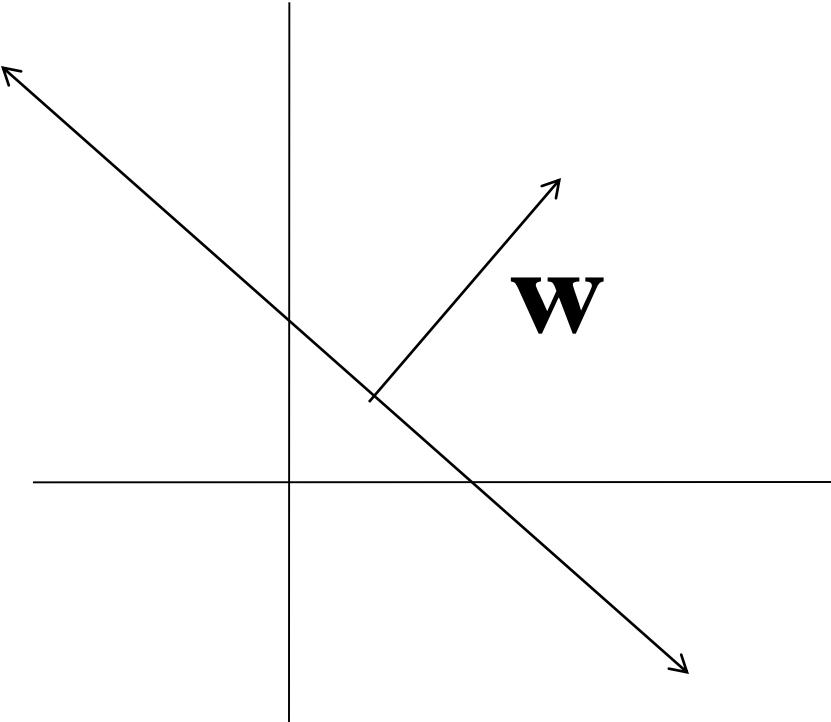
# Lines in $\mathbb{R}^2$



Let  $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$   $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$

# Lines in $\mathbb{R}^2$

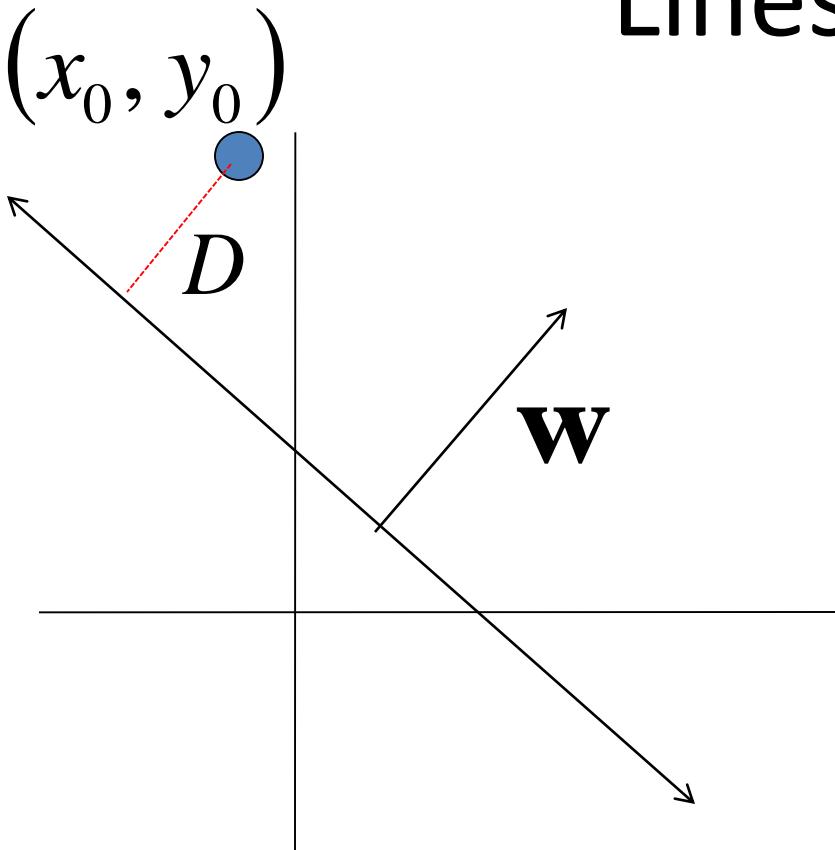


Let  $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$   $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$

$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

# Lines in $\mathbb{R}^2$

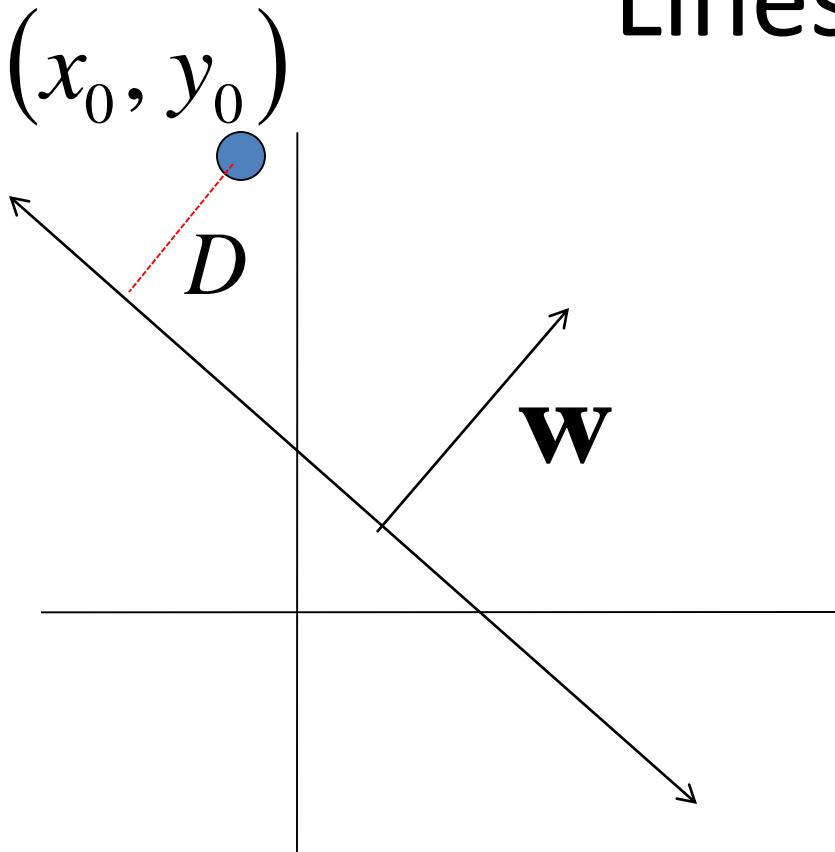


Let  $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$   $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$

$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

# Lines in $\mathbb{R}^2$



Let  $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$   $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$

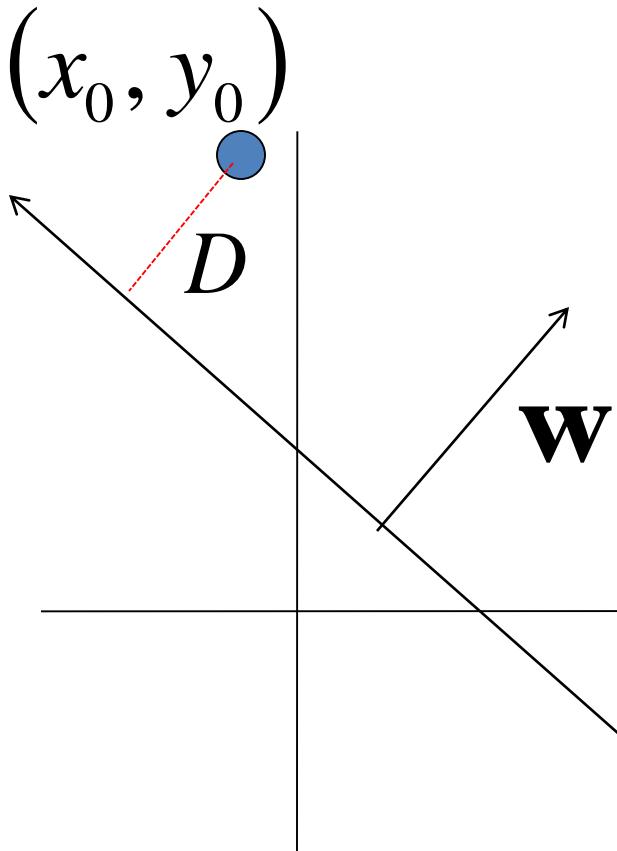


$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

$$D = \frac{|ax_0 + cy_0 + b|}{\sqrt{a^2 + c^2}}$$

distance from  
point to line

# Lines in $\mathbb{R}^2$



Let  $\mathbf{w} = \begin{bmatrix} a \\ c \end{bmatrix}$   $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

$$ax + cy + b = 0$$



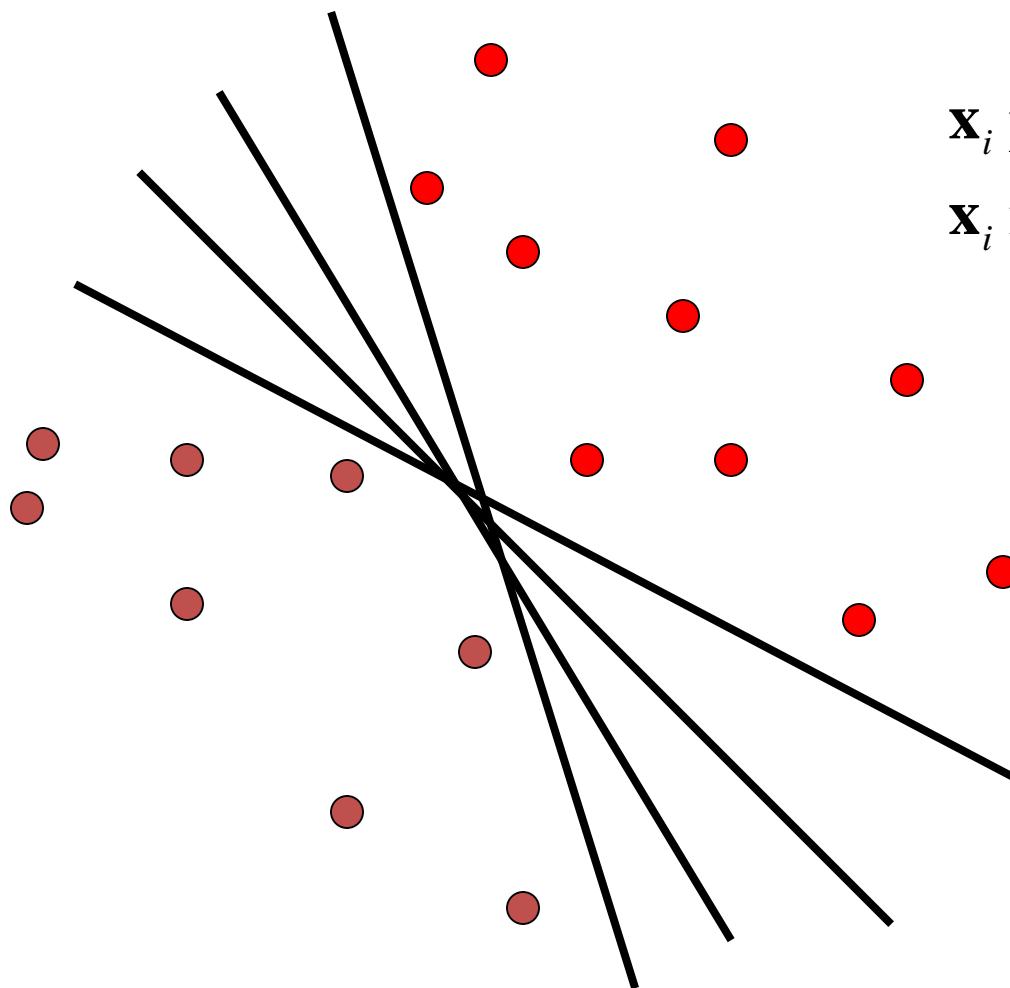
$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

$$D = \frac{|ax_0 + cy_0 + b|}{\sqrt{a^2 + c^2}} = \frac{\mathbf{w}^\top \mathbf{x} + b}{\|\mathbf{w}\|}$$

} distance from  
point to line

# Linear classifiers

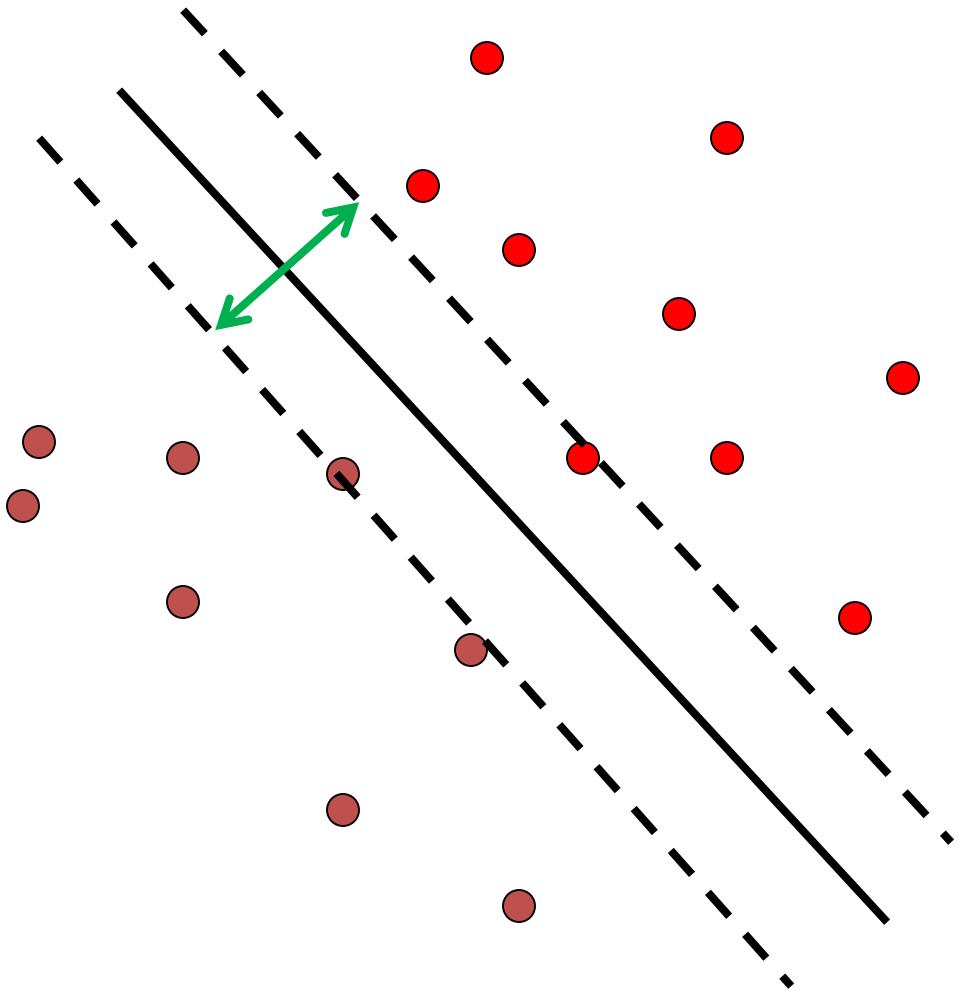
- Find linear function to separate positive and negative examples



$\mathbf{x}_i$  positive :  $\mathbf{x}_i \cdot \mathbf{w} + b \geq 0$   
 $\mathbf{x}_i$  negative :  $\mathbf{x}_i \cdot \mathbf{w} + b < 0$

Which line  
is best?

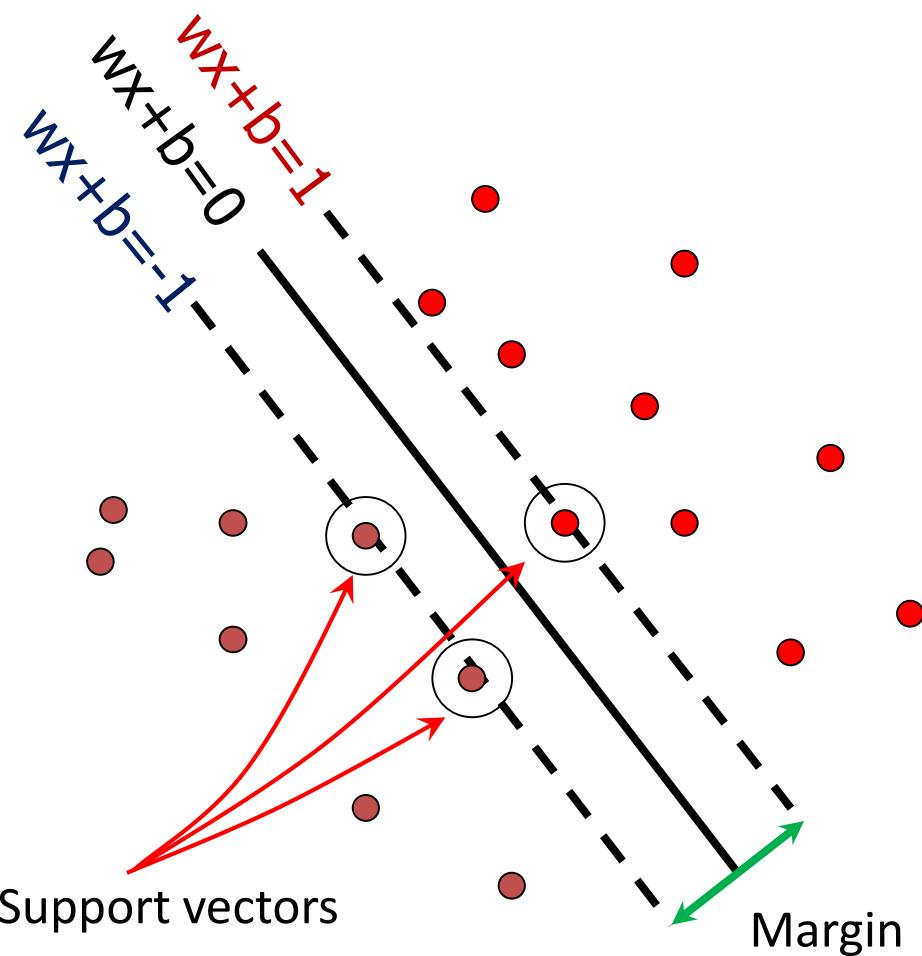
# Support Vector Machines (SVMs)



- Discriminative classifier based on *optimal separating line* (for 2d case)
- Maximize the *margin* between the positive and negative training examples

# Support vector machines

- Want line that maximizes the margin.



$$\mathbf{x}_i \text{ positive } (y_i = 1): \quad \mathbf{x}_i \cdot \mathbf{w} + b \geq 1$$

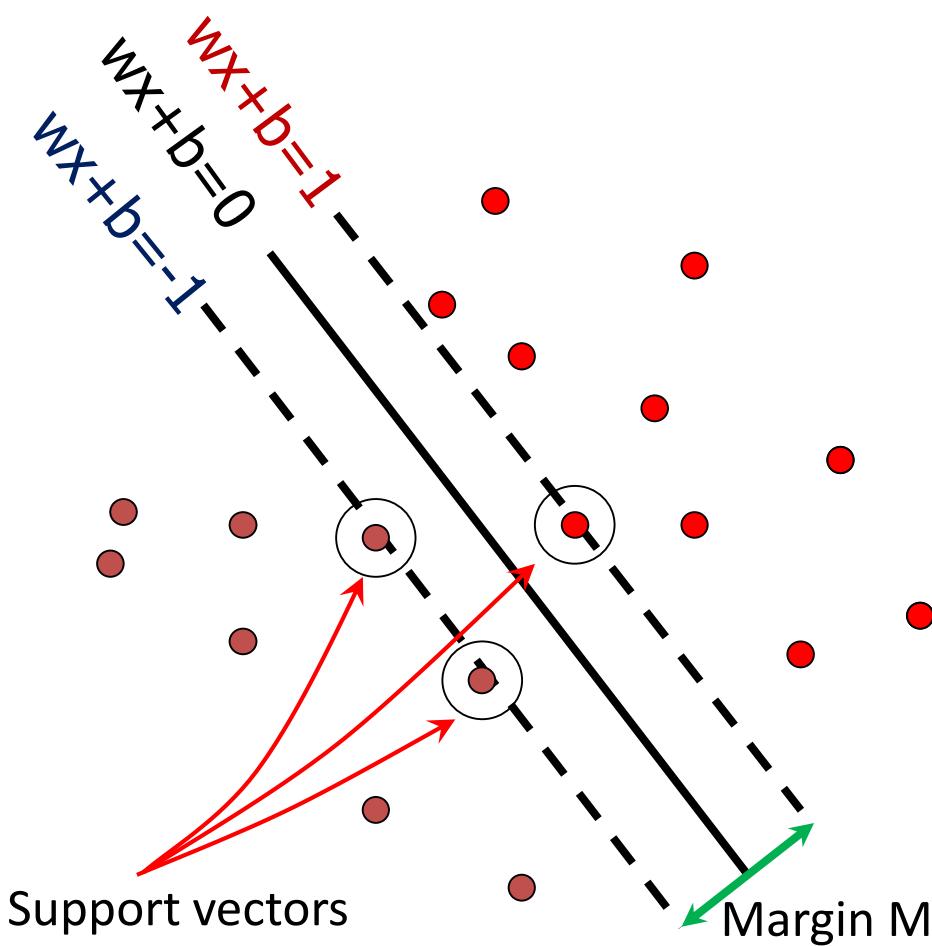
$$\mathbf{x}_i \text{ negative } (y_i = -1): \quad \mathbf{x}_i \cdot \mathbf{w} + b \leq -1$$

For support, vectors,

$$\mathbf{x}_i \cdot \mathbf{w} + b = \pm 1$$

# Support vector machines

- Want line that maximizes the margin.



$$\mathbf{x}_i \text{ positive } (y_i = 1): \quad \mathbf{x}_i \cdot \mathbf{w} + b \geq 1$$

$$\mathbf{x}_i \text{ negative } (y_i = -1): \quad \mathbf{x}_i \cdot \mathbf{w} + b \leq -1$$

$$\text{For support, vectors,} \quad \mathbf{x}_i \cdot \mathbf{w} + b = \pm 1$$

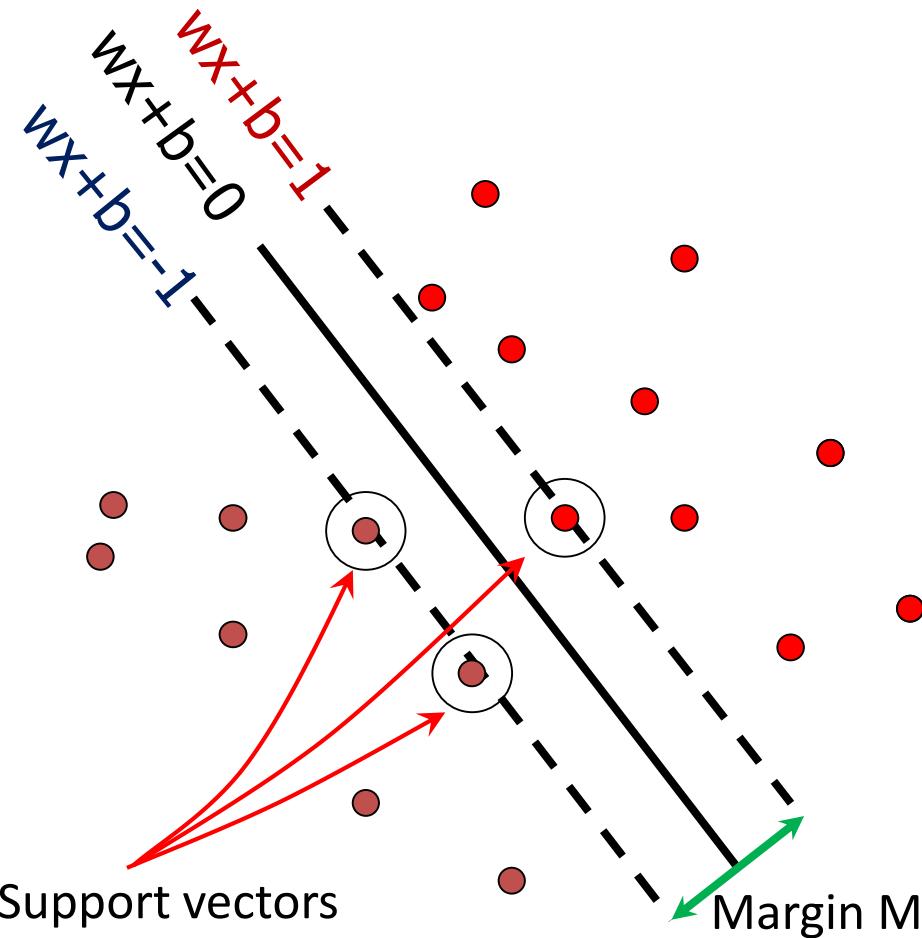
$$\text{Distance between point and line:} \quad \frac{|\mathbf{x}_i \cdot \mathbf{w} + b|}{\|\mathbf{w}\|}$$

$$\text{For support vectors:}$$

$$\frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|} = \frac{\pm 1}{\|\mathbf{w}\|} \quad M = \left| \frac{1}{\|\mathbf{w}\|} - \frac{-1}{\|\mathbf{w}\|} \right| = \frac{2}{\|\mathbf{w}\|}$$

# Support vector machines

- Want line that maximizes the margin.



$\mathbf{x}_i$  positive ( $y_i = 1$ ):  $\mathbf{x}_i \cdot \mathbf{w} + b \geq 1$

$\mathbf{x}_i$  negative ( $y_i = -1$ ):  $\mathbf{x}_i \cdot \mathbf{w} + b \leq -1$

For support, vectors,  $\mathbf{x}_i \cdot \mathbf{w} + b = \pm 1$

Distance between point and line: 
$$\frac{|\mathbf{x}_i \cdot \mathbf{w} + b|}{\|\mathbf{w}\|}$$

Therefore, the margin is  $2 / \|\mathbf{w}\|$

# Finding the maximum margin line

1. Maximize margin  $2/\|\mathbf{w}\|$
2. Correctly classify all training data points:

$$\mathbf{x}_i \text{ positive } (y_i = 1): \quad \mathbf{x}_i \cdot \mathbf{w} + b \geq 1$$

$$\mathbf{x}_i \text{ negative } (y_i = -1): \quad \mathbf{x}_i \cdot \mathbf{w} + b \leq -1$$

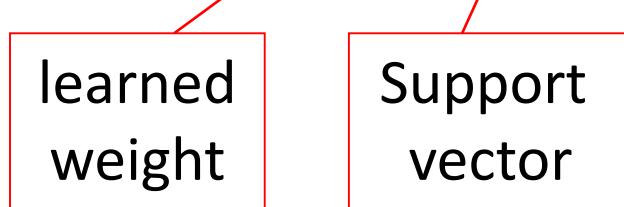
- *Quadratic optimization problem:*

- $$\text{Minimize } \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

$$\text{Subject to } y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$$

# Finding the maximum margin line

- Solution:  $\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$



# Finding the maximum margin line

- Solution:  $\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$   
 $b = y_i - \mathbf{w} \cdot \mathbf{x}_i$  (for any support vector)  
 $\mathbf{w} \cdot \mathbf{x} + b = \sum_i \alpha_i y_i \mathbf{x}_i \cdot \mathbf{x} + b$

- Classification function:



If  $f(x) < 0$ , classify as negative,  
if  $f(x) > 0$ , classify as positive

$$\begin{aligned}f(x) &= \text{sign} (\mathbf{w} \cdot \mathbf{x} + b) \\&= \text{sign}\left(\sum_i \alpha_i \mathbf{x}_i \cdot \mathbf{x} + b\right)\end{aligned}$$

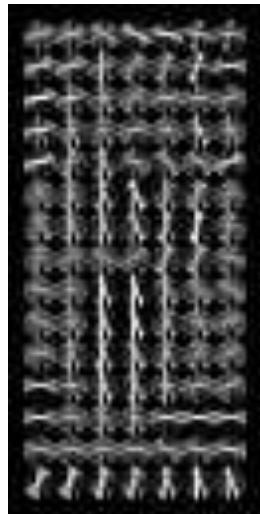
# Questions

- **What if the features are not 2d?**
- What if the data is not linearly separable?
- What if we have more than just two categories?

# Questions

- What if the features are not 2d?
  - Generalizes to d-dimensions – replace line with “hyperplane”
- What if the data is not linearly separable?
- What if we have more than just two categories?

# Person detection with HoG's & linear SVM's



- Map each grid cell in the input window to a histogram counting the gradients per orientation.
- Train a linear SVM using training set of pedestrian vs. non-pedestrian windows.

Code available:  
<http://pascal.inrialpes.fr/soft/olt/>

# Person detection with HoG's & linear SVM's

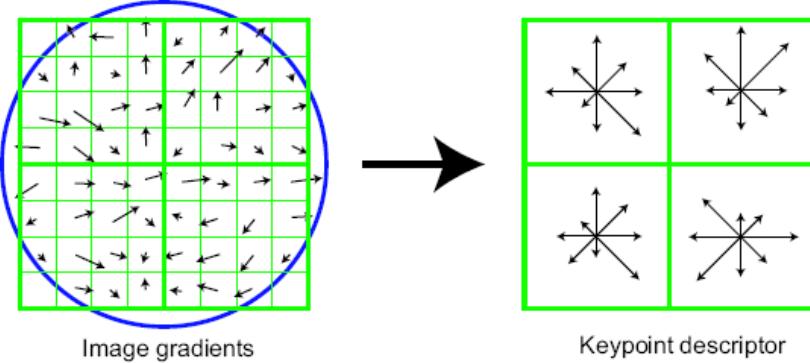


- Histograms of Oriented Gradients for Human Detection, [Navneet Dalal](#), [Bill Triggs](#), International Conference on Computer Vision & Pattern Recognition - June 2005
- <http://lear.inrialpes.fr/pubs/2005/DT05/>

# Histograms of oriented gradients

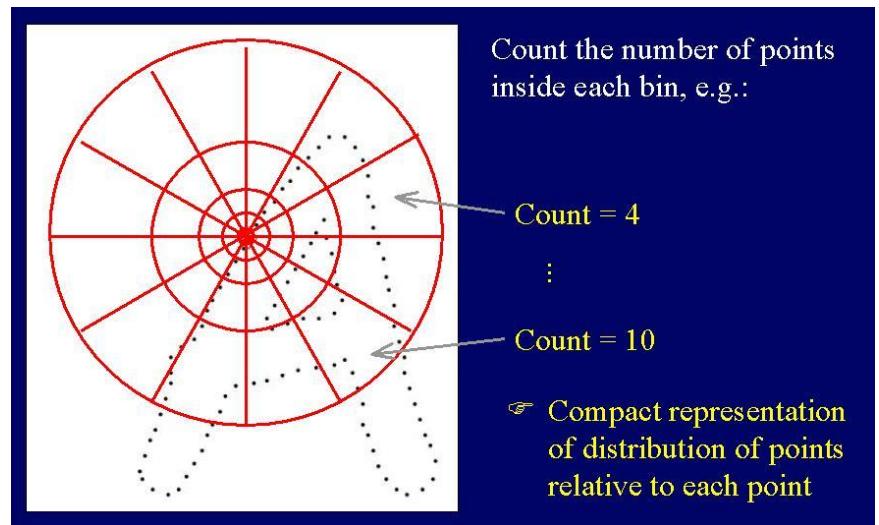
# Histograms of oriented gradients

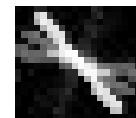
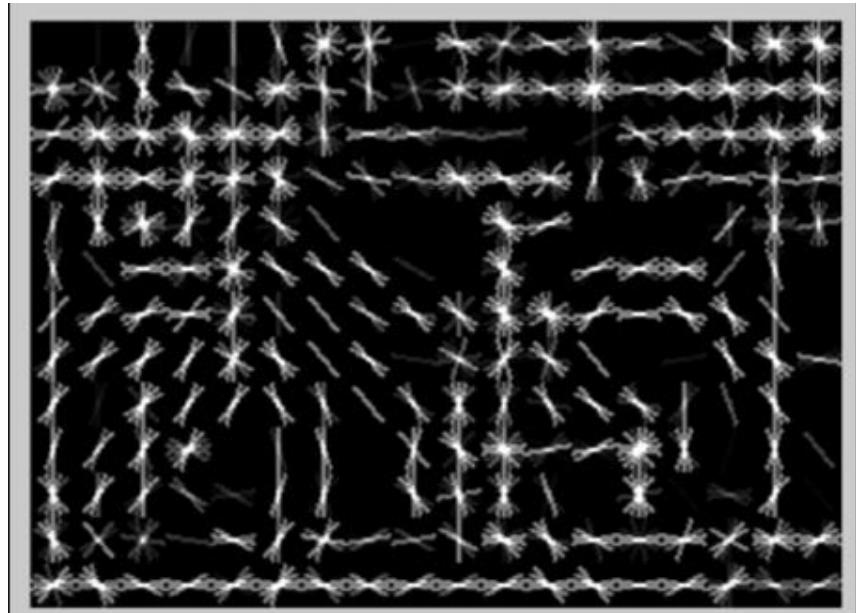
SIFT, D. Lowe, ICCV 1999



Shape context

Belongie, Malik, Puzicha, NIPS 2000





Source: Deva Ramanan

# Histograms of Oriented Gradients for Human Detection

Navneet Dalal and Bill Triggs

INRIA Rhône-Alps, 655 avenue de l'Europe, Montbonnot 38334, France  
{Navneet.Dalal,Bill.Triggs}@inrialpes.fr, <http://lear.inrialpes.fr>



Figure 1. An overview of our feature extraction and object detection chain. The detector window is tiled with a grid of overlapping blocks in which Histogram of Oriented Gradient feature vectors are extracted. The combined vectors are fed to a linear SVM for object/non-object classification. The detection window is scanned across the image at all positions and scales, and conventional non-maximum suppression is run on the output pyramid to detect object instances, but this paper concentrates on the feature extraction process.

# Histograms of Oriented Gradients for Human Detection

Navneet Dalal and Bill Triggs

INRIA Rhône-Alps, 655 avenue de l'Europe, Montbonnot 38334, France  
{Navneet.Dalal,Bill.Triggs}@inrialpes.fr, http://lear.inrialpes.fr

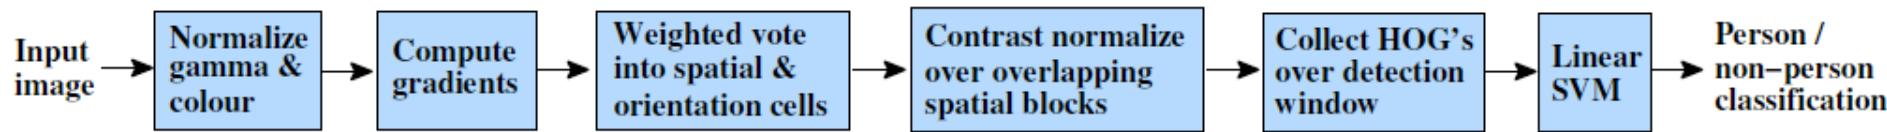
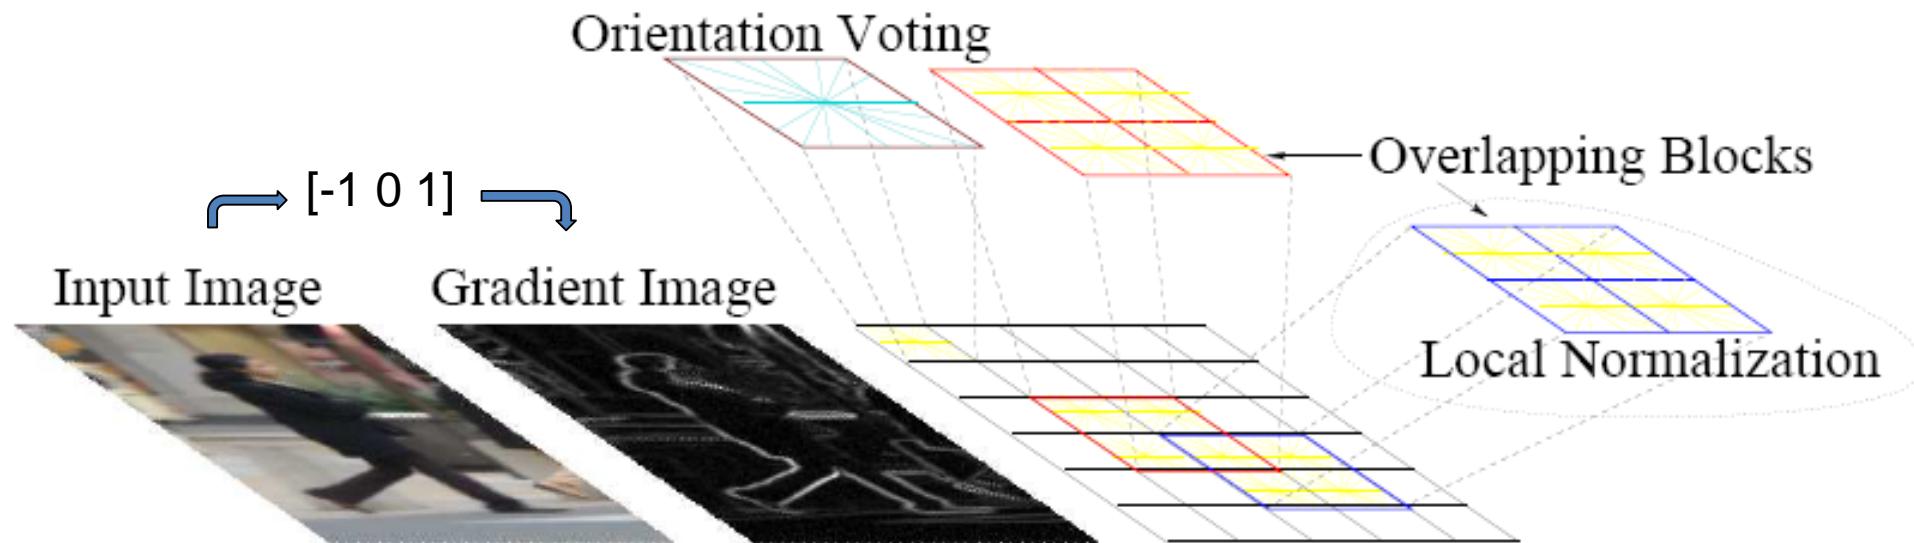


Figure 1. An overview of our feature extraction and object detection chain. The detector window is tiled with a grid of overlapping blocks in which Histogram of Oriented Gradient feature vectors are extracted. The combined vectors are fed to a linear SVM for object/non-object classification. The detection window is scanned across the image at all positions and scales, and conventional non-maximum suppression is run on the output pyramid to detect object instances, but this paper concentrates on the feature extraction process.



# SVM

A Support Vector Machine (SVM) learns a classifier with the form:

$$H(x) = \sum_{m=1}^M a_m y_m k(x, x_m)$$

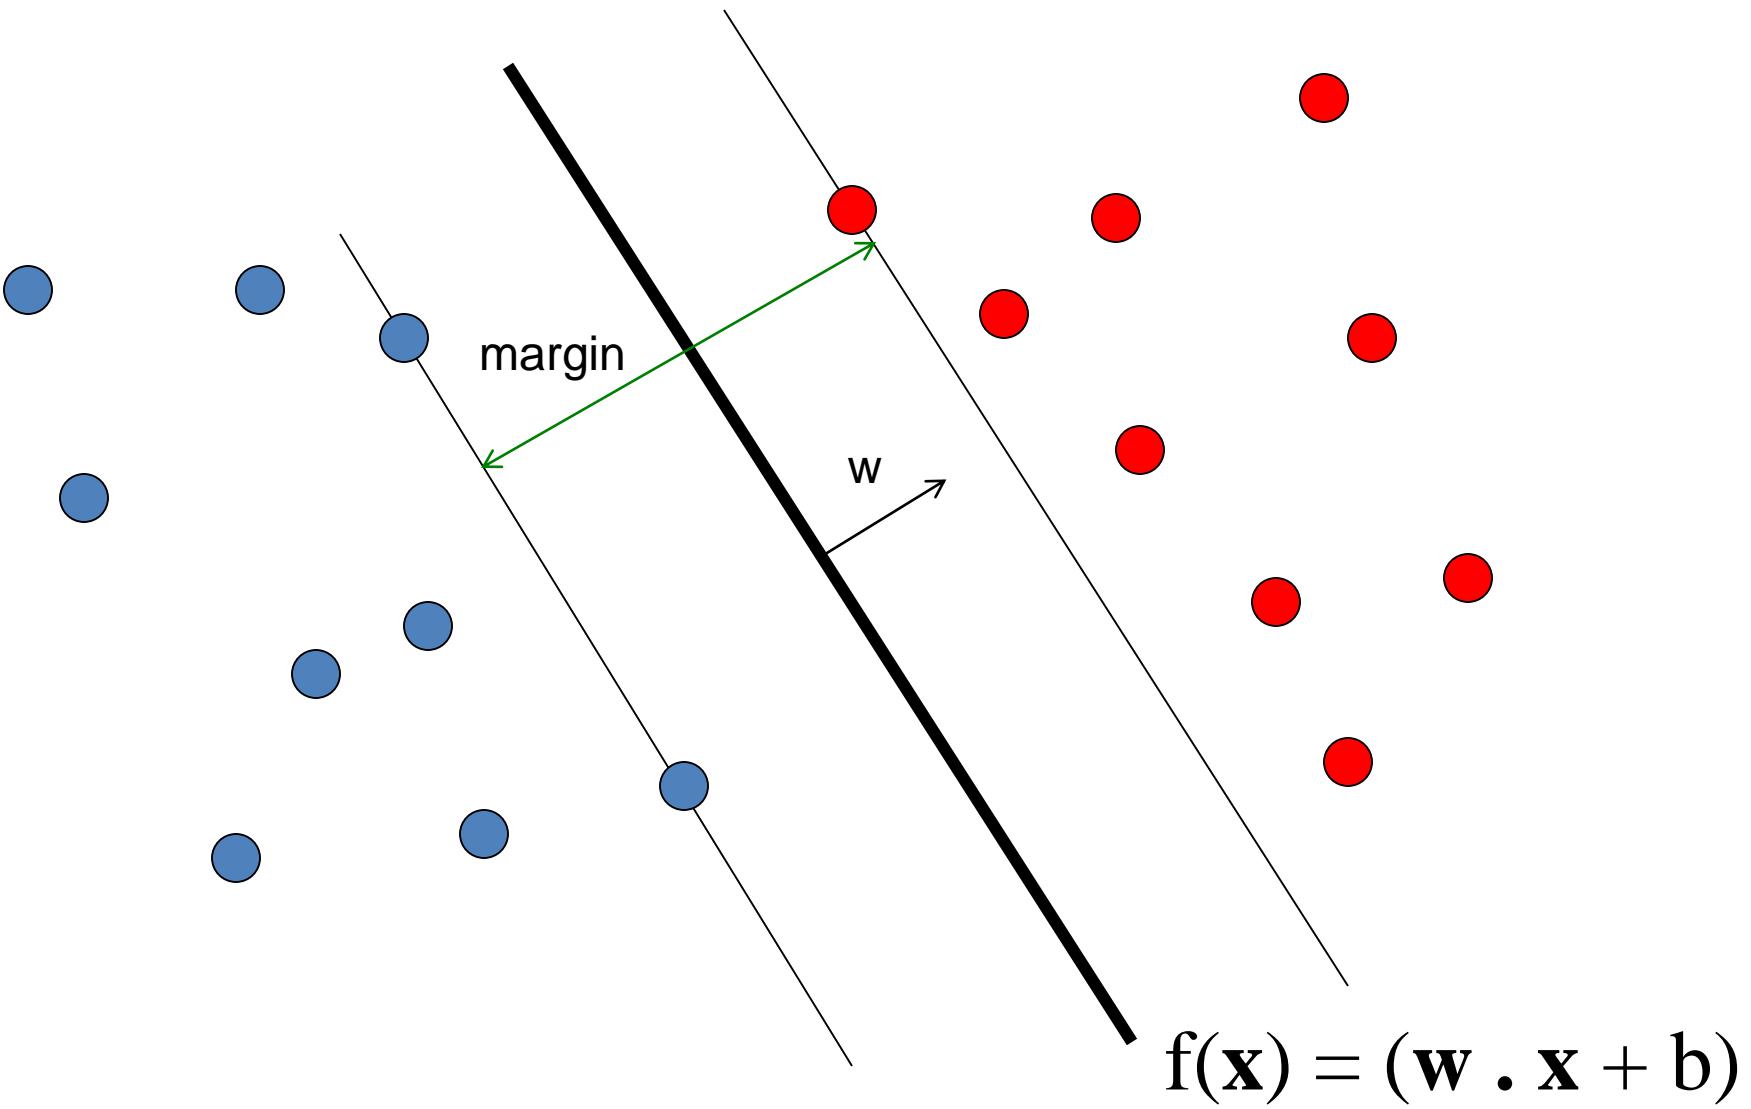
Where  $\{x_m, y_m\}$ , for  $m = 1 \dots M$ , are the training data with  $x_m$  being the input feature vector and  $y_m = +1, -1$  the class label.  $k(x, x_m)$  is the kernel and it can be any symmetric function satisfying the Mercer Theorem.

The classification is obtained by thresholding the value of  $H(x)$ .

There is a large number of possible kernels, each yielding a different family of decision boundaries:

- Linear kernel:  $k(x, x_m) = x^T x_m$
- Radial basis function:  $k(x, x_m) = \exp(-|x - x_m|^2/\sigma^2)$ .
- Histogram intersection:  $k(x, x_m) = \sum_i (\min(x(i), x_m(i)))$

# Linear SVM



# Scanning-window templates

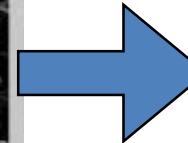
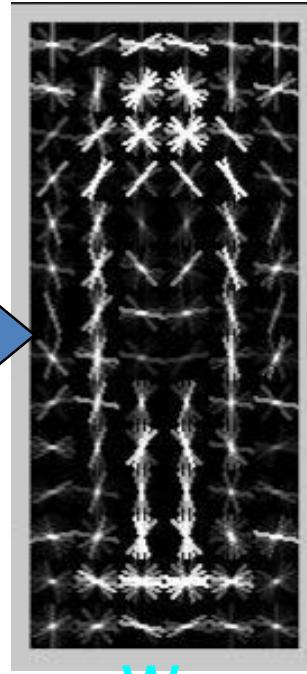
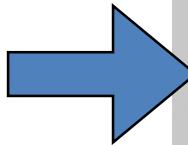
Dalal and Triggs CVPR05 (HOG)

Papageorgiou and Poggio ICI99 (wavelets)

pos



neg



$W$

$W$  = weights for orientation and spatial bins

$$w \cdot x > 0$$

Train with a linear classifier (perceptron, logistic regression, SVMs...)



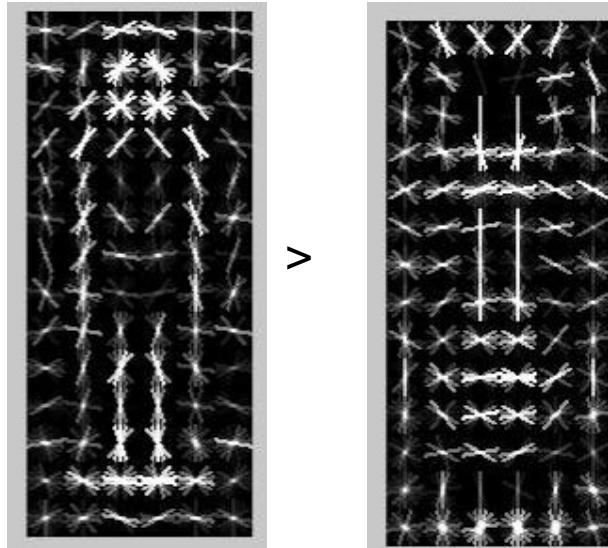
# How to interpret positive and **negative** weights?

$$w \cdot x > 0$$

$$(w_{\text{pos}} - w_{\text{neg}}) \cdot x > 0$$

$$w_{\text{pos}} \cdot x > w_{\text{neg}} \cdot x$$

Pedestrian template



Pedestrian  
background  
template

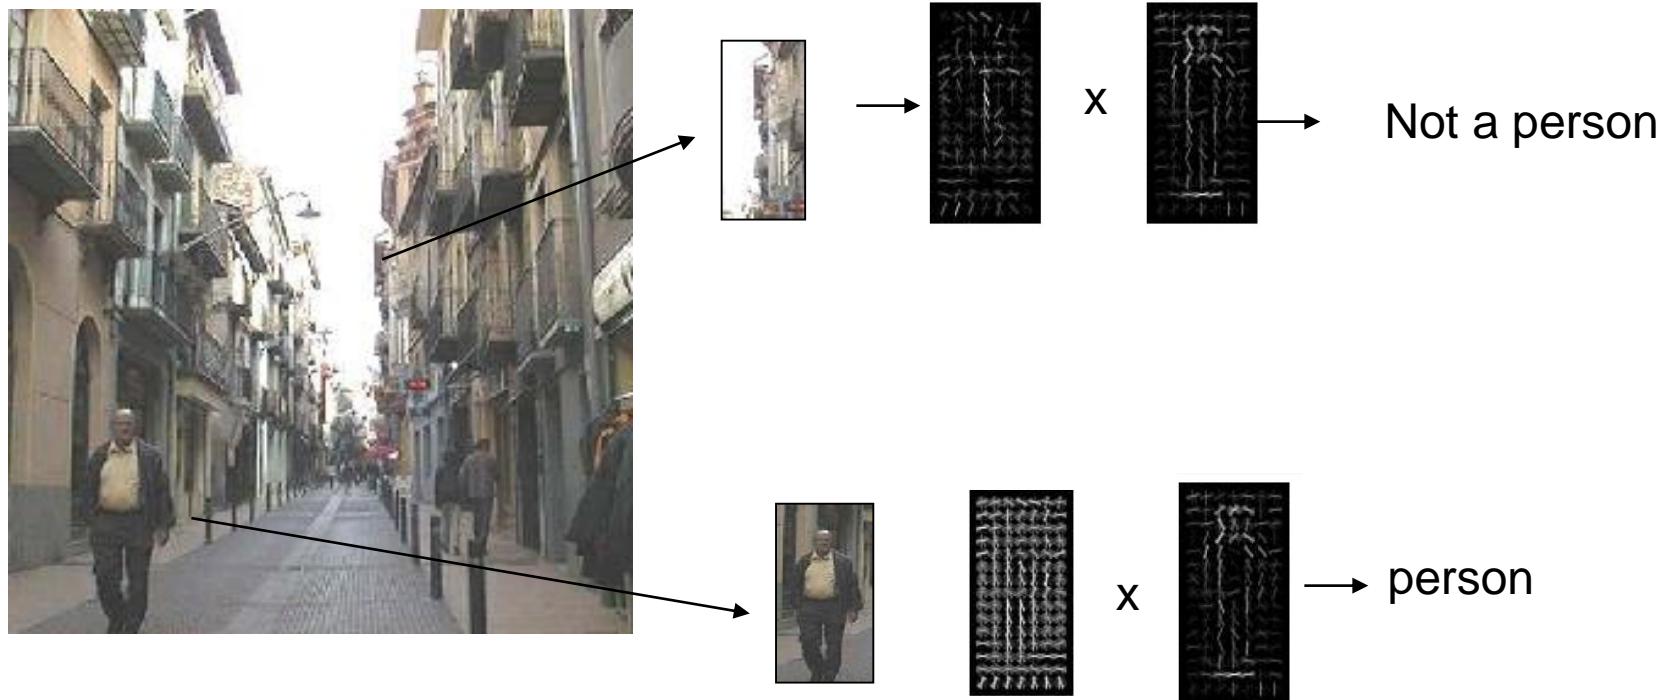
$w_{\text{pos}}, w_{\text{neg}}$  = weighted average of positive, negative support vectors

Right approach is to **compete** pedestrian, pillar, doorway... models

Background class is hard to model - easier to penalize particular vertical edges

# Histograms of oriented gradients

Dalal & Trigs, 2006



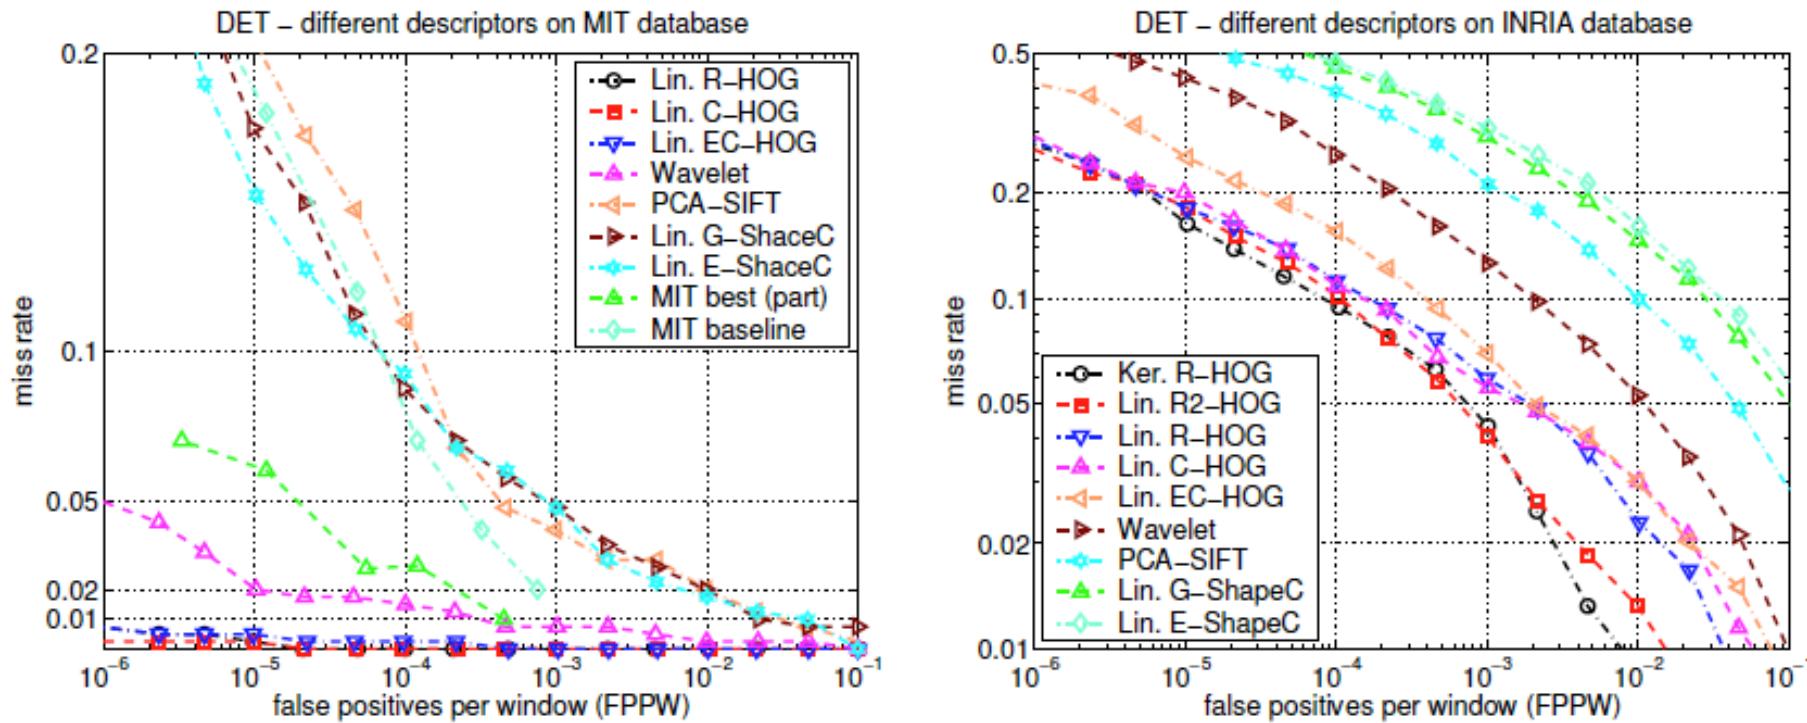


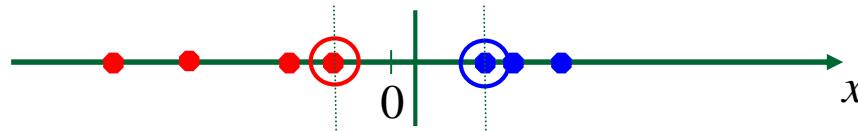
Figure 3. The performance of selected detectors on (left) MIT and (right) INRIA data sets. See the text for details.

# Questions

- What if the features are not 2d?
- **What if the data is not linearly separable?**
- What if we have more than just two categories?

# Non-linear SVMs

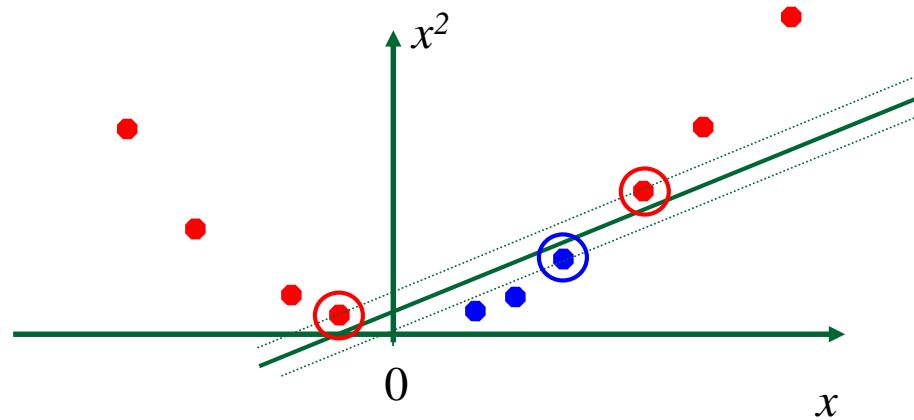
- Datasets that are linearly separable with some noise work out great:



- But what are we going to do if the dataset is just too hard?

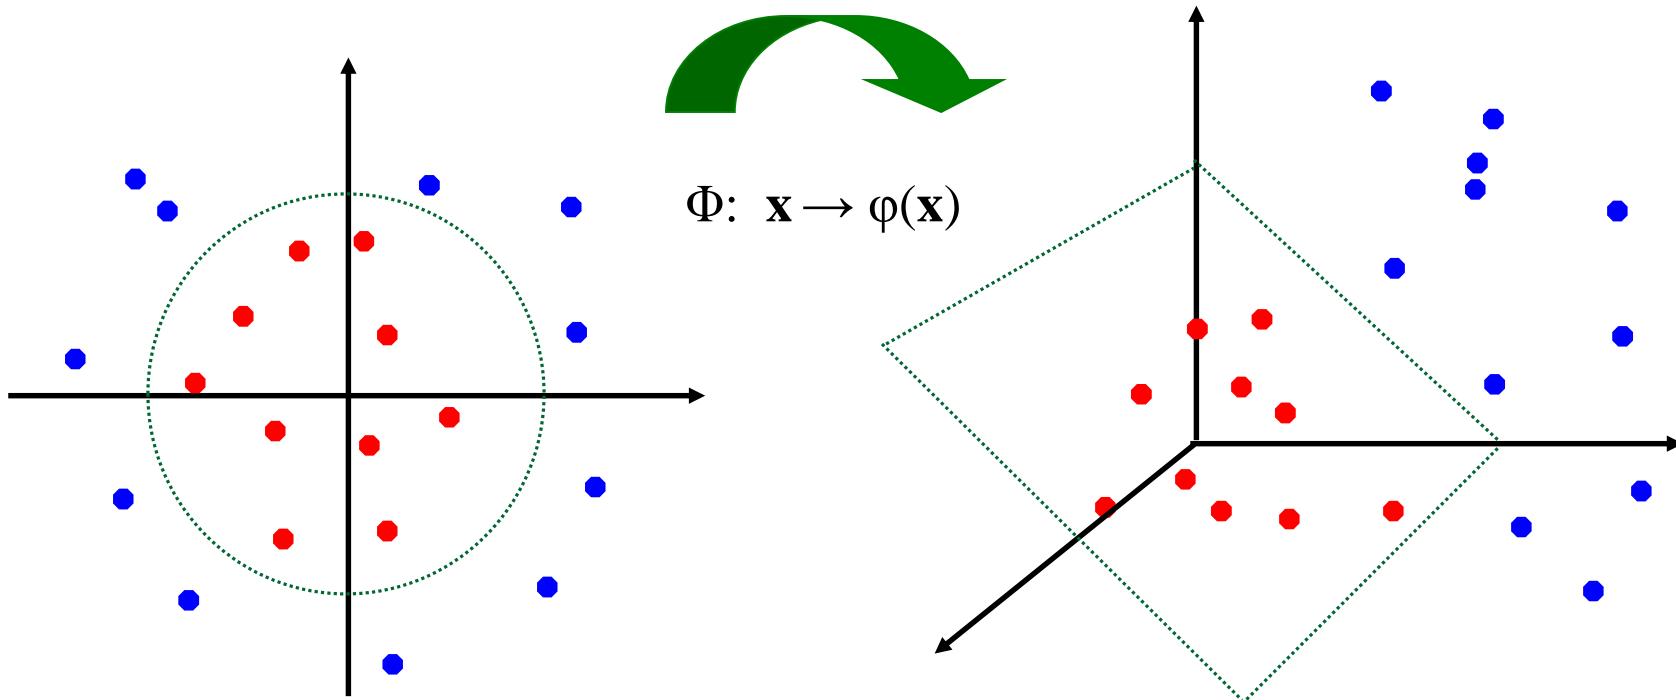


- How about... mapping data to a higher-dimensional space:



# Non-linear SVMs: feature spaces

- General idea: the original input space can be mapped to some higher-dimensional feature space where the training set is separable:



# The “Kernel Trick”

- The linear classifier relies on dot product between vectors  $K(x_i, x_j) = x_i^T x_j$
- If every data point is mapped into high-dimensional space via some transformation  $\Phi: x \rightarrow \phi(x)$ , the dot product becomes:

$$K(x_i, x_j) = \phi(x_i)^T \phi(x_j)$$

- A *kernel function* is similarity function that corresponds to an inner product in some expanded feature space.

# Example

2-dimensional vectors  $\mathbf{x} = [x_1 \ x_2]$ ;

$$\text{let } K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^2$$

Need to show that  $K(\mathbf{x}_i, \mathbf{x}_j) = \varphi(\mathbf{x}_i)^T \varphi(\mathbf{x}_j)$ :

$$K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^2,$$

$$= 1 + x_{i1}^2 x_{j1}^2 + 2 x_{i1} x_{j1} x_{i2} x_{j2} + x_{i2}^2 x_{j2}^2 + 2 x_{i1} x_{j1} + 2 x_{i2} x_{j2}$$

$$= [1 \ x_{i1}^2 \ \sqrt{2} x_{i1} x_{i2} \ x_{i2}^2 \ \sqrt{2} x_{i1} \ \sqrt{2} x_{i2}]^T$$

$$[1 \ x_{j1}^2 \ \sqrt{2} x_{j1} x_{j2} \ x_{j2}^2 \ \sqrt{2} x_{j1} \ \sqrt{2} x_{j2}]$$

$$= \varphi(\mathbf{x}_i)^T \varphi(\mathbf{x}_j),$$

$$\text{where } \varphi(\mathbf{x}) = [1 \ x_1^2 \ \sqrt{2} x_1 x_2 \ x_2^2 \ \sqrt{2} x_1 \ \sqrt{2} x_2]$$

# Nonlinear SVMs

- *The kernel trick:* instead of explicitly computing the lifting transformation  $\varphi(\mathbf{x})$ , define a kernel function  $K$  such that

$$K(\mathbf{x}_i, \mathbf{x}_j) = \varphi(\mathbf{x}_i) \cdot \varphi(\mathbf{x}_j)$$

- This gives a nonlinear decision boundary in the original feature space:

$$\sum_i \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b$$

# Examples of kernel functions

$$K(x_i, x_j) = x_i^T x_j$$

- Linear:

$$K(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

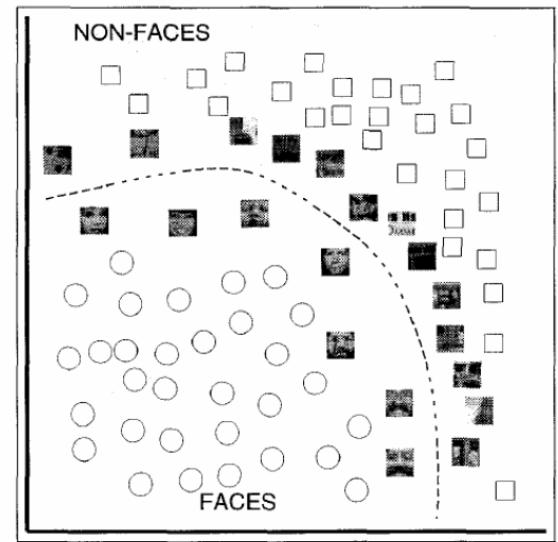
- Gaussian RBF:

$$K(x_i, x_j) = \sum_k \min(x_i(k), x_j(k))$$

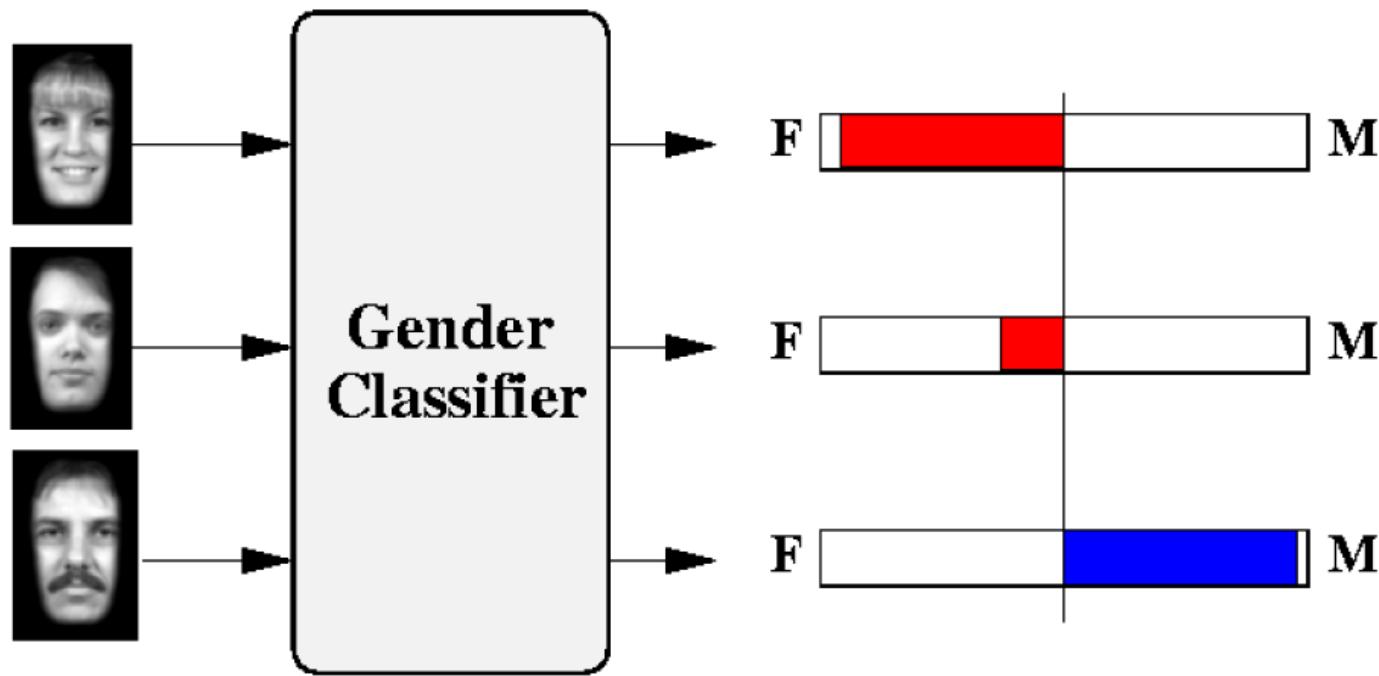
- Histogram intersection:

# SVMs for recognition

1. Define your representation for each example.
2. Select a kernel function.
3. Compute pairwise kernel values between labeled examples
4. Use this “kernel matrix” to solve for SVM support vectors & weights.
5. To classify a new example: compute kernel values between new input and support vectors, apply weights, check sign of output.



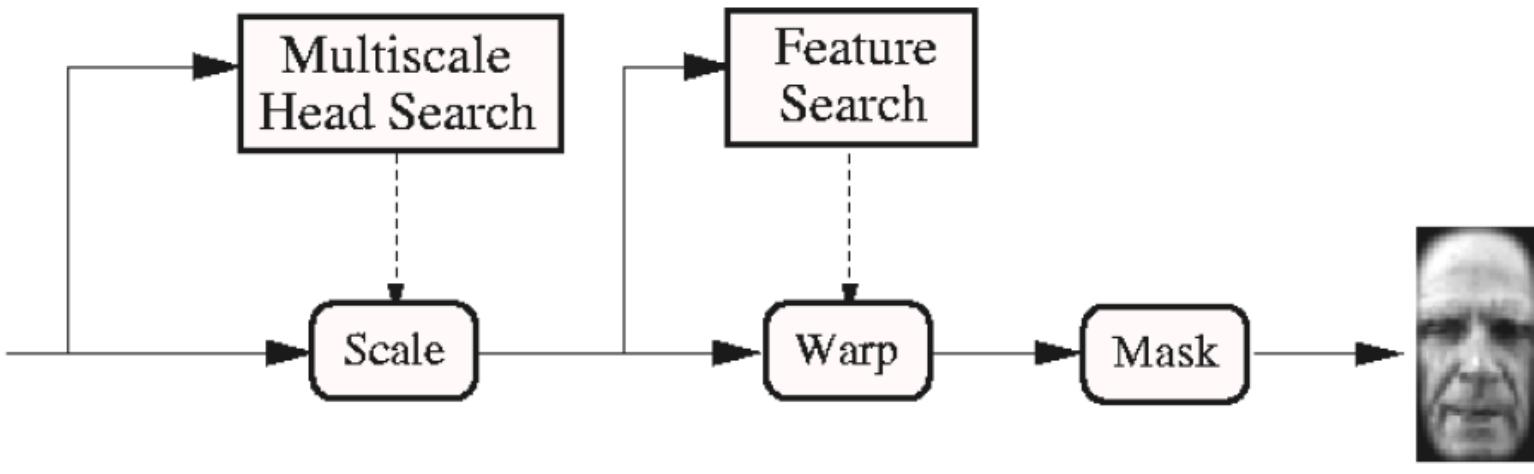
# Example: learning gender with SVMs



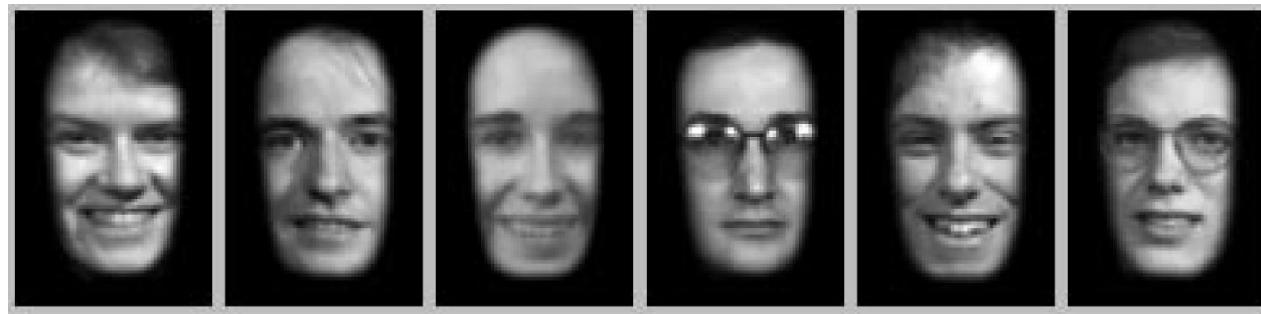
Moghaddam and Yang, Learning Gender with Support Faces, TPAMI 2002.

Moghaddam and Yang, Face & Gesture 2000.

Face alignment  
processing



Processed faces

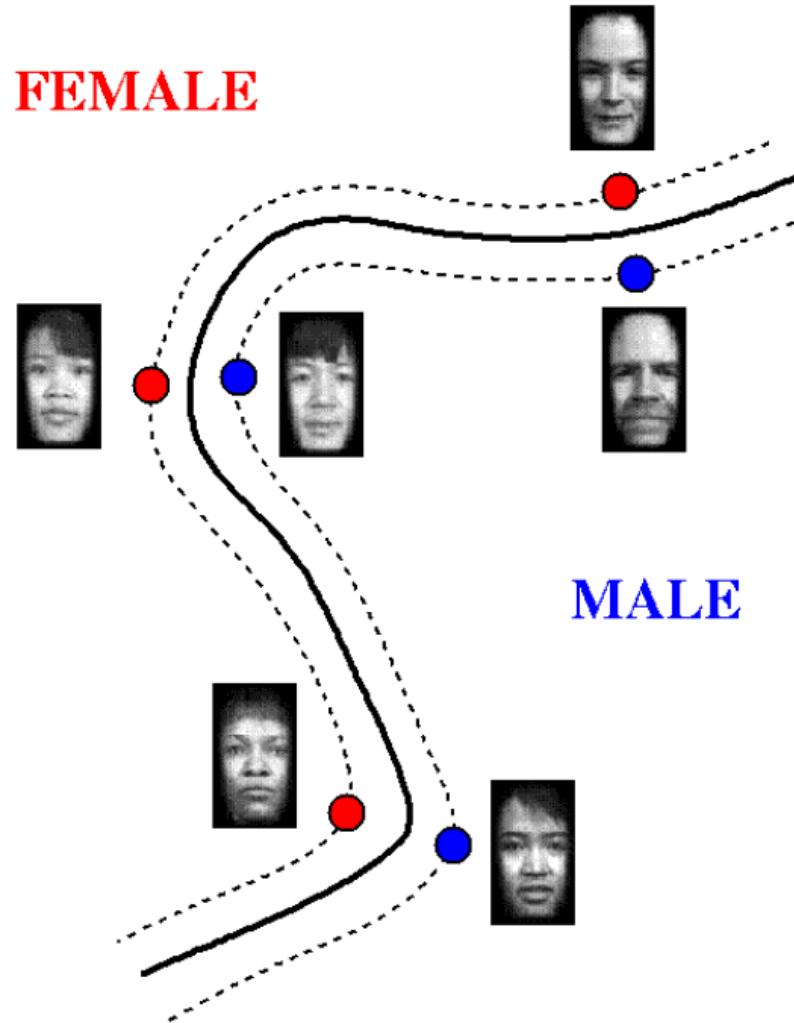


# Learning gender with SVMs

- Training examples:
  - 1044 males
  - 713 females
- Experiment with various kernels, select Gaussian RBF

$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$

# Support Faces



# Classifier Performance

Classifier	Error Rate		
	Overall	Male	Female
SVM with RBF kernel	3.38%	2.05%	4.79%
SVM with cubic polynomial kernel	4.88%	4.21%	5.59%
Large Ensemble of RBF	5.54%	4.59%	6.55%
Classical RBF	7.79%	6.89%	8.75%
Quadratic classifier	10.63%	9.44%	11.88%
Fisher linear discriminant	13.03%	12.31%	13.78%
Nearest neighbor	27.16%	26.53%	28.04%
Linear classifier	58.95%	58.47%	59.45%

# Gender perception experiment: How well can humans do?

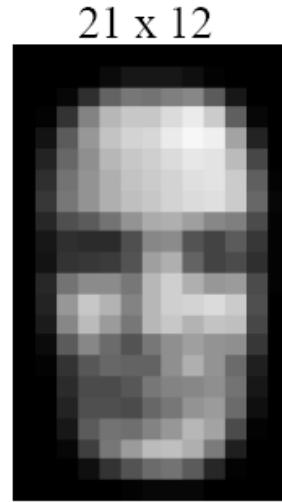
- Subjects:
  - 30 people (22 male, 8 female)
  - Ages mid-20's to mid-40's
- Test data:
  - 254 face images (6 males, 4 females)
  - Low res and high res versions
- Task:
  - Classify as male or female, forced choice
  - No time limit

# Gender perception experiment: How well can humans do?

Stimuli →



N = 4032



N = 252

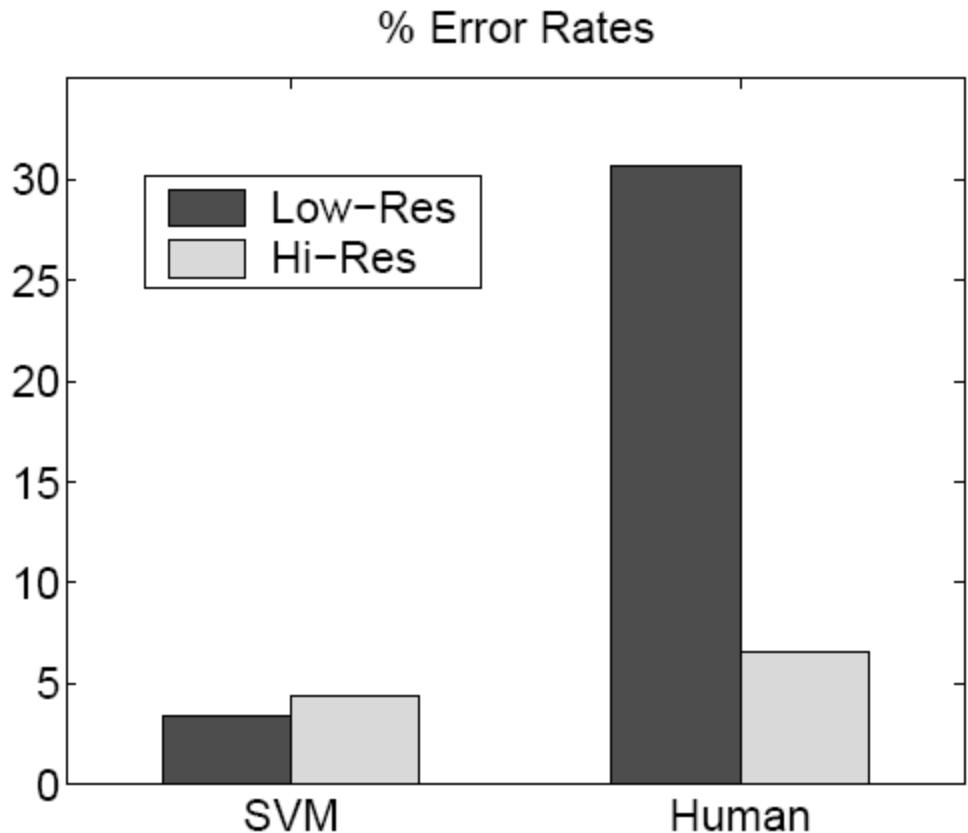
Results →

High-Res	Low-Res
6.54%	30.7%

Error      Error

$\sigma = 3.7\%$

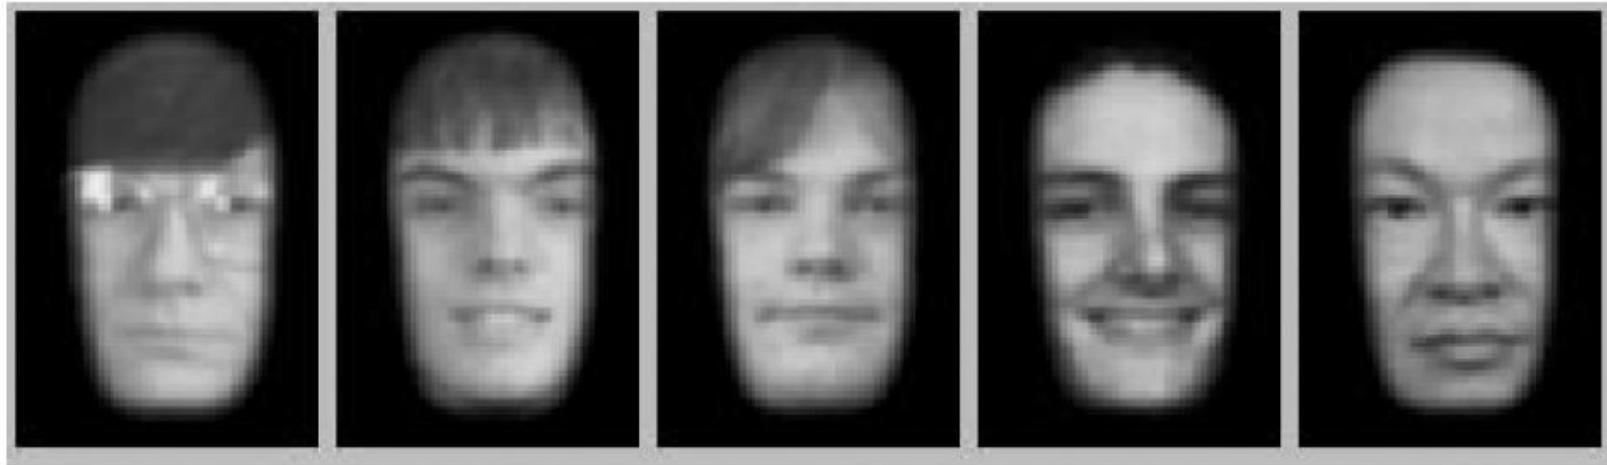
# Human vs. Machine



- SVMs performed better than any single human test subject, at either resolution

**Figure 6. SVM vs. Human performance**

# Hardest examples for humans



**Top five human misclassifications**

# Questions

- What if the features are not 2d?
- What if the data is not linearly separable?
- **What if we have more than just two categories?**

# Multi-class SVMs

- Achieve multi-class classifier by combining a number of binary classifiers
- **One vs. all**
  - Training: learn an SVM for each class vs. the rest
  - Testing: apply each SVM to test example and assign to it the class of the SVM that returns the highest decision value
- **One vs. one**
  - Training: learn an SVM for each pair of classes
  - Testing: each learned SVM “votes” for a class to assign to the test example

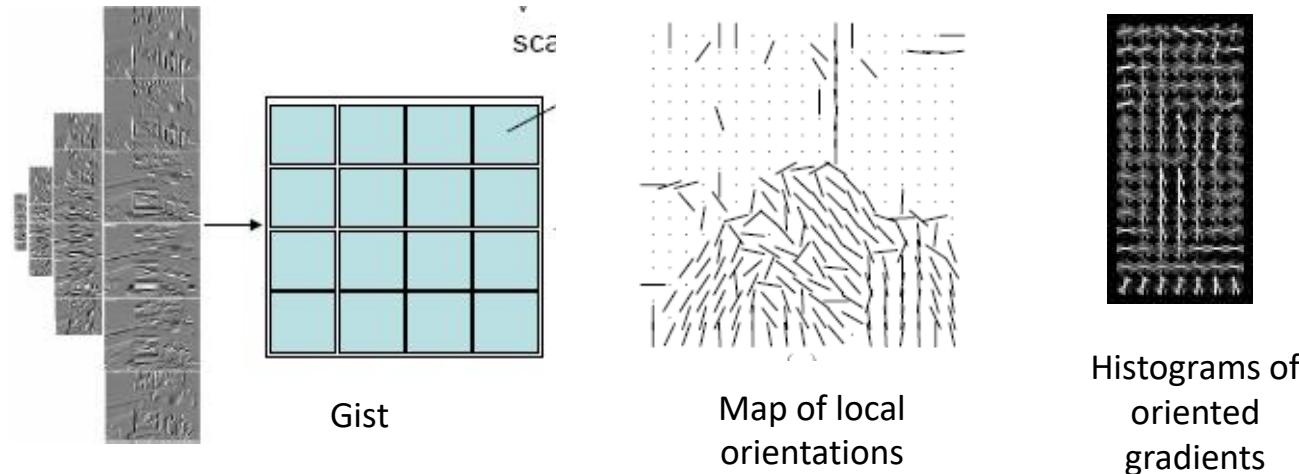
# SVMs: Pros and cons

- Pros
  - Many publicly available SVM packages:  
<http://www.kernel-machines.org/software>
  - <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
  - Kernel-based framework is very powerful, flexible
  - Often a sparse set of support vectors – compact at test time
  - Work very well in practice, even with very small training sample sizes
- Cons
  - No “direct” multi-class SVM, must combine two-class SVMs
  - Can be tricky to select best kernel function for a problem
  - Computation, memory
    - During training time, must compute matrix of kernel values for every pair of examples
    - Learning can take a very long time for large-scale problems

# Summary

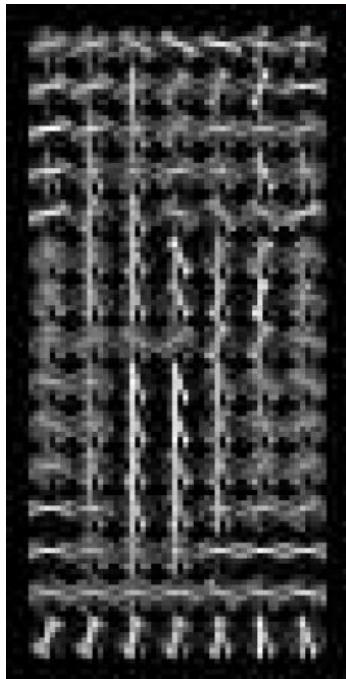
- Discriminative classifiers
  - Boosting
  - Nearest neighbors
  - Support vector machines
- Useful for object recognition when combined with “window-based” or holistic appearance descriptors

# Global window-based appearance representations

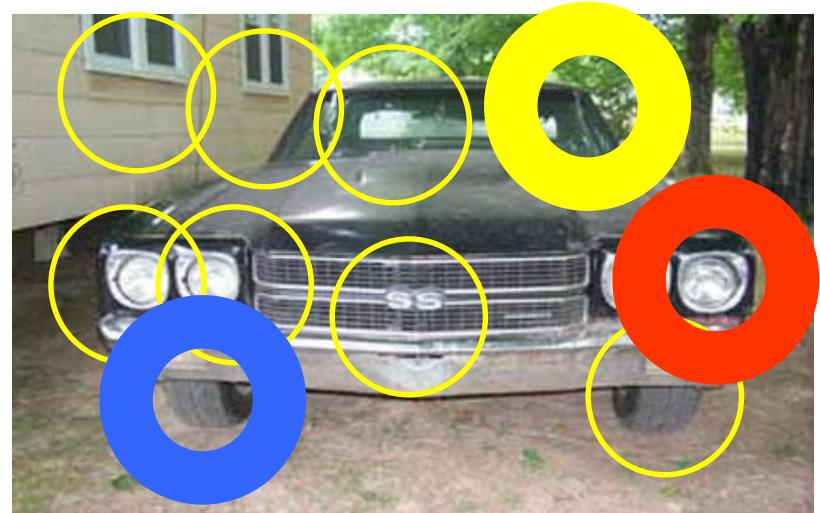


- These examples are truly global; each pixel in the window contributes to the representation.
- Classifier can account for relative relevance...
- *When might this not be ideal?*

# Generic category recognition: representation choice



Window-based



Part-based