

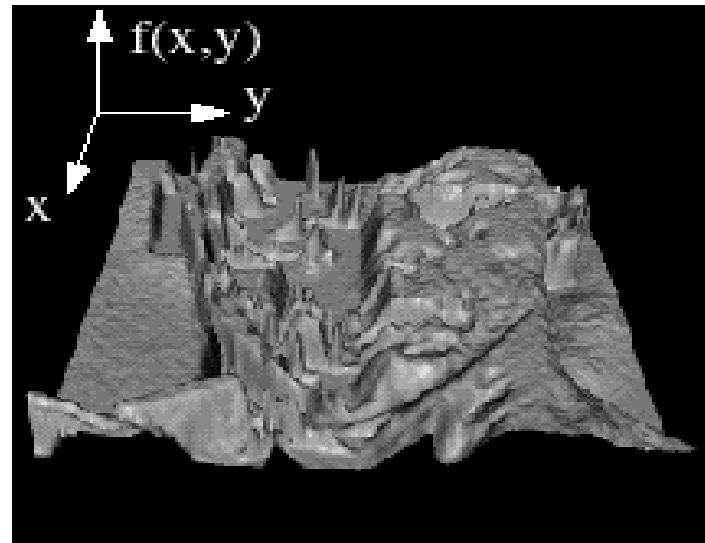
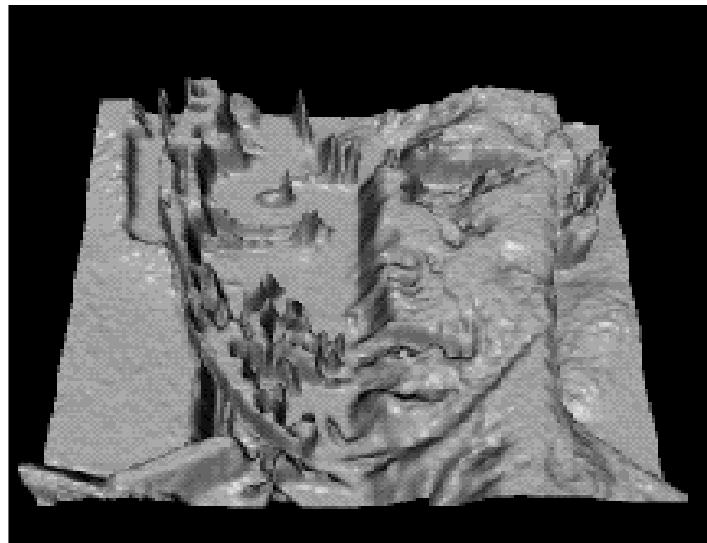
# Filters

CMP 719 – Computer Vision  
Pinar Duygulu  
Hacettepe University

# Today's topics

- **Image Formation**
- **Image filters in spatial domain**
  - Filter is a mathematical operation of a grid of numbers
  - Smoothing, sharpening, measuring texture
- **Image filters in the frequency domain**
  - Filtering is a way to modify the frequencies of images
  - Denoising, sampling, image compression
- **Templates and Image Pyramids**
  - Filtering is a way to match a template to the image
  - Detection, coarse-to-fine registration

# Images as functions



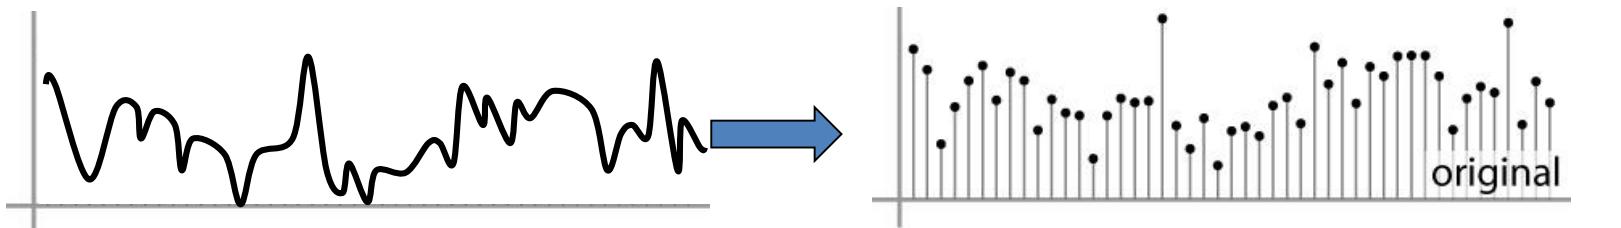
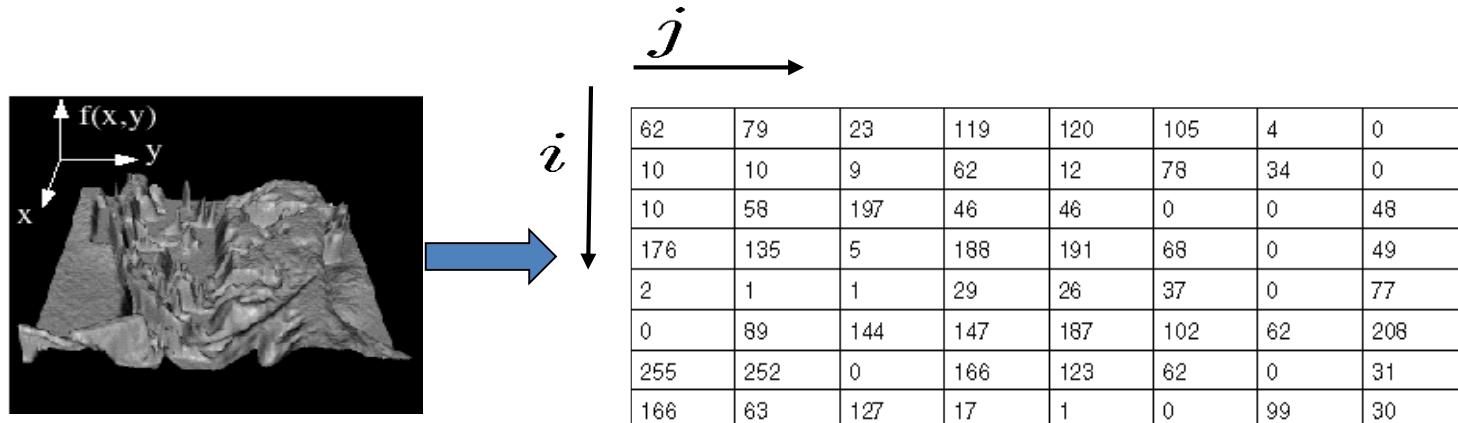
# Images as functions

- We can think of an image as a function,  $f$ , from  $\mathbf{R}^2$  to  $\mathbf{R}$ :
  - $f(x, y)$  gives the intensity at position  $(x, y)$
  - Realistically, we expect the image only to be defined over a rectangle, with a finite range:
    - $f: [a,b] \times [c,d] \rightarrow [0, 255]$
- A color image is just three functions pasted together. We can write this as a “vector-valued” function:

$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

# Digital images

- In computer vision we operate on **digital (discrete)** images:
  - Sample the 2D space on a regular grid
  - Quantize each sample (round to nearest integer)
- Image thus represented as a matrix of integer values.



# Images as discrete functions

- Cartesian Coordinates

$$f[n, m] = \begin{bmatrix} & & & \vdots & \\ \ddots & & & & \\ & f[-1, 1] & f[0, 1] & f[1, 1] & \\ \dots & f[-1, 0] & \underline{f[0, 0]} & f[1, 0] & \dots \\ & f[-1, -1] & f[0, -1] & f[1, -1] & \\ & \vdots & & \ddots & \end{bmatrix}$$

# Today's topics

- Image Formation
- Image filters in spatial domain
  - Filter is a mathematical operation of a grid of numbers
  - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
  - Filtering is a way to modify the frequencies of images
  - Denoising, sampling, image compression
- Templates and Image Pyramids
  - Filtering is a way to match a template to the image
  - Detection, coarse-to-fine registration

# Zebras vs. Dalmatians



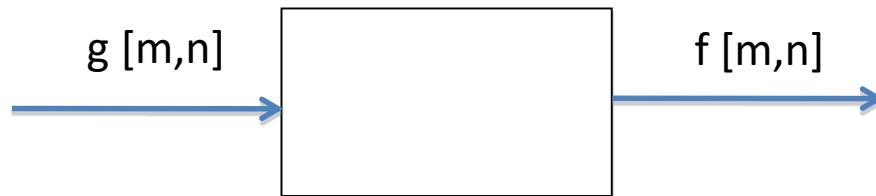
Both zebras and dalmatians have black and white pixels in about the same number

- if we shuffle the images point-wise processing is not affected

Need to measure properties relative to small *neighborhoods* of pixels

- find different image patterns

# Filtering

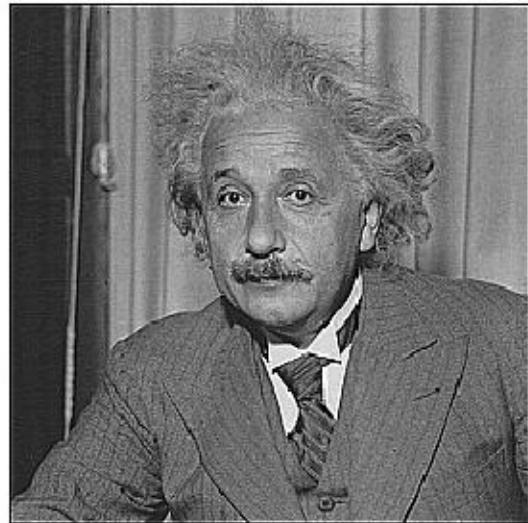
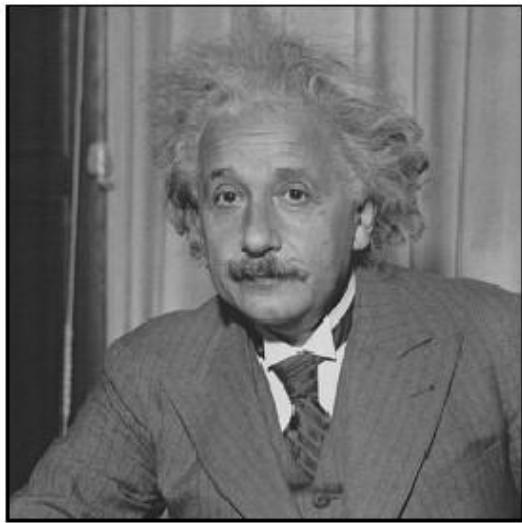


We want to remove unwanted sources of variation, and keep the information relevant for whatever task we need to solve

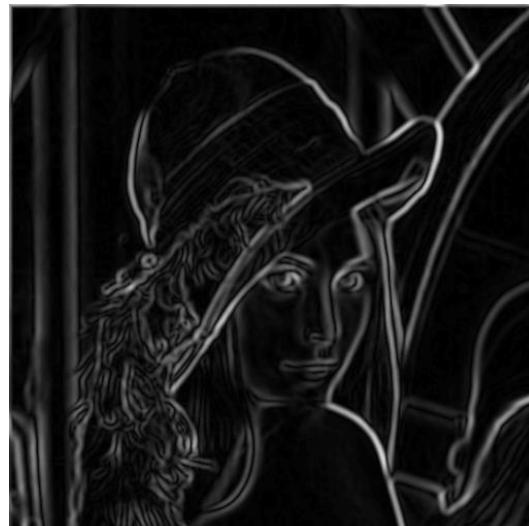


# Filters

- **Filtering:**
  - Form a new image whose pixels are a combination of original pixel values
  - compute function of local neighborhood at each position
- **Goals:**
- Extract useful information from the images
  - Features (textures, edges, corners, distinctive points, blobs...)
- Modify or enhance image properties:
  - super-resolution; in-painting; de-noising, resizing
- Detect patterns
  - Template matching



*Smooth/Sharpen Images...*



*Find edges...*



*Find waldo...*

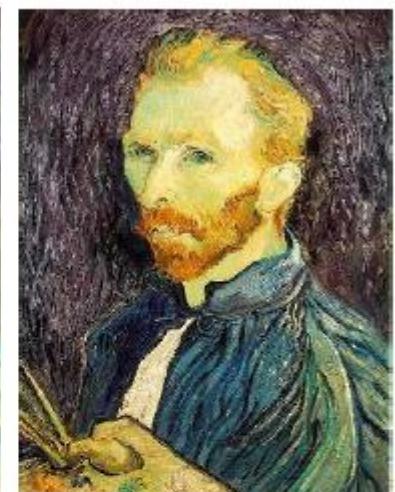
De-noising



Salt and pepper noise



Super-resolution



In-painting



Bertamio et al

# Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise

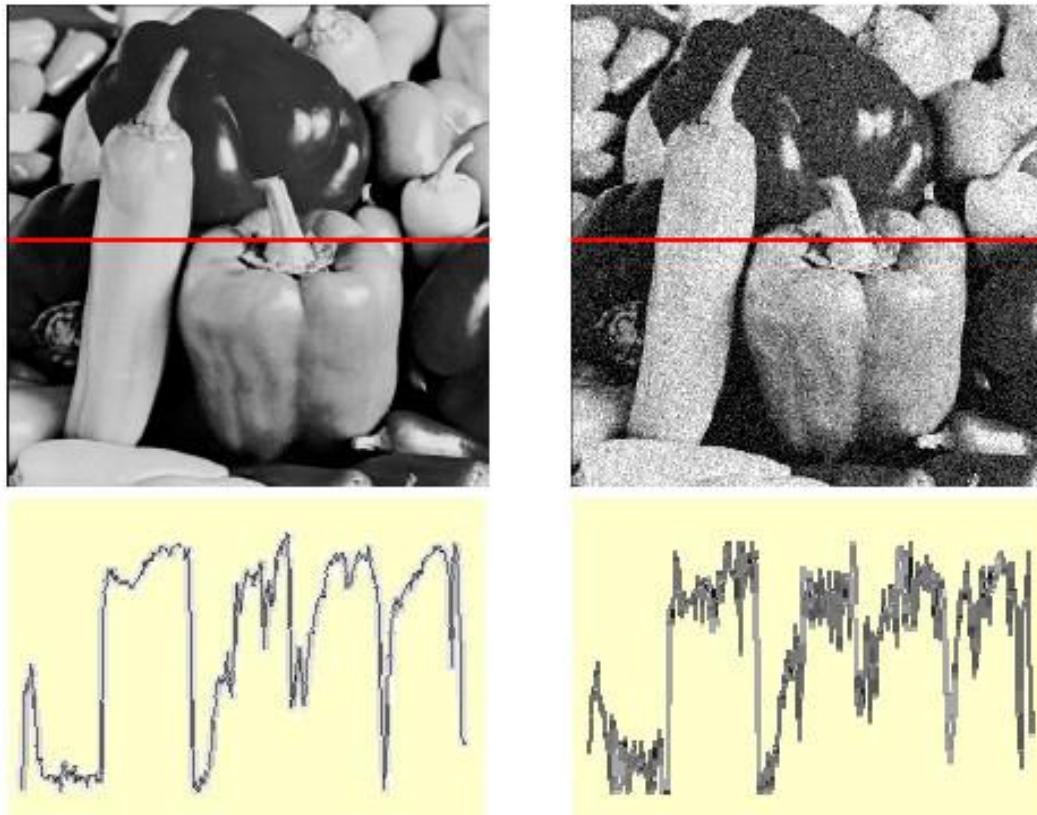


Impulse noise



Gaussian noise

# Gaussian noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

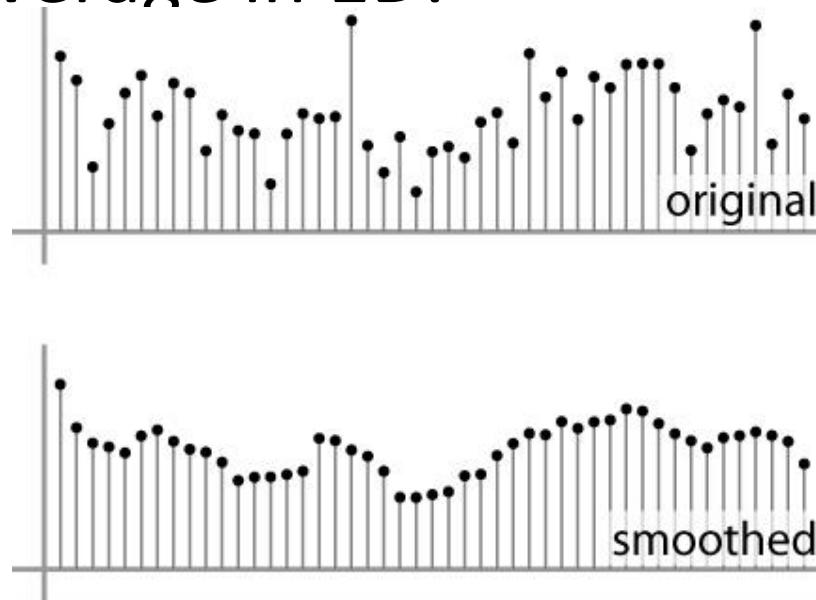
```
>> noise = randn(size(im)).*sigma;  
  
>> output = im + noise;
```

# First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
  - Expect pixels to be like their neighbors
  - Expect noise processes to be independent from pixel to pixel

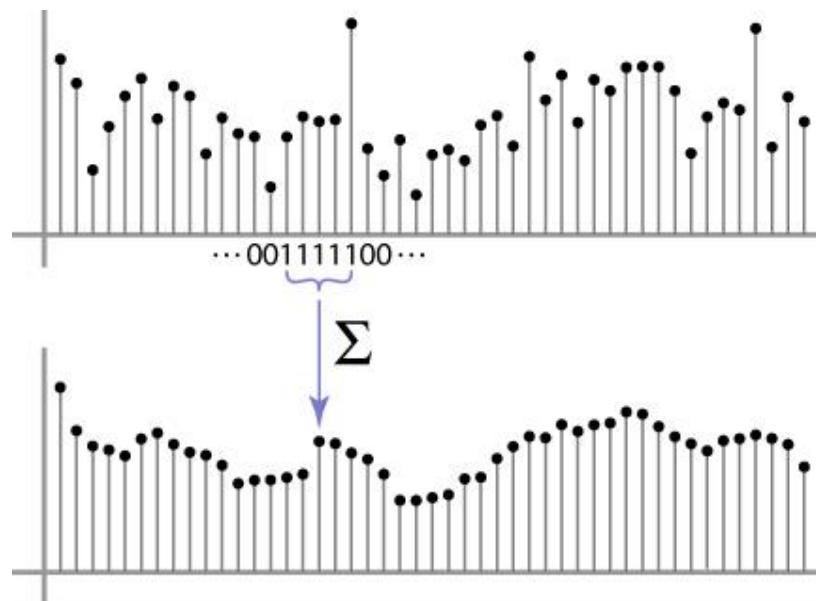
# First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



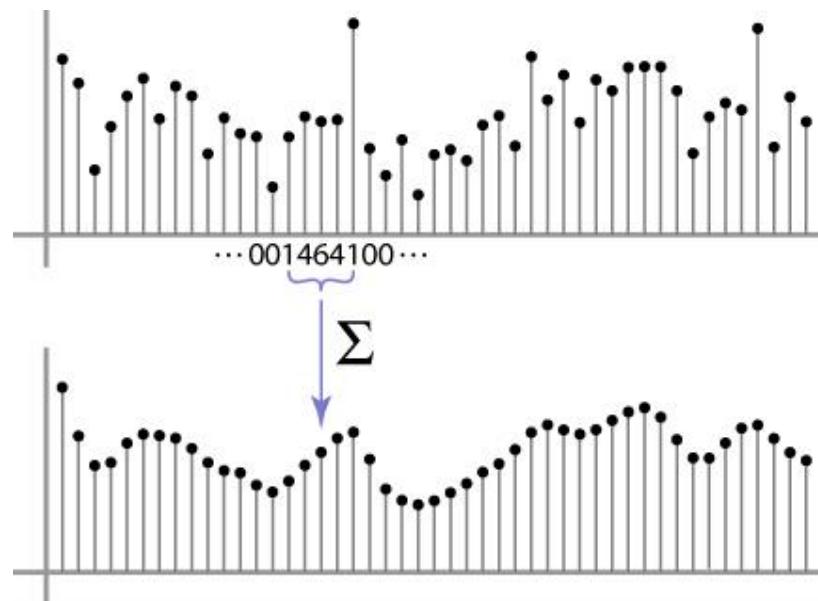
# Weighted Moving Average

- Can add weights to our moving average
- *Weights*  $[1, 1, 1, 1, 1] / 5$



# Weighted Moving Average

- Non-uniform weights  $[1, 4, 6, 4, 1] / 16$



# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

			0							

# Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10									

# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	0	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$


# Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$


# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$


# Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

# Correlation filtering

Say the averaging window size is  $2k+1 \times 2k+1$ :

$$G[i, j] = \frac{1}{(2k+1)^2} \underbrace{\sum_{u=-k}^k}_{\text{Attribute uniform weight}} \underbrace{\sum_{v=-k}^k}_{\text{Loop over all pixels in neighborhood around to each pixel}} F[i+u, j+v]$$

*Attribute uniform weight Loop over all pixels in neighborhood around to each pixel image pixel  $F[i,j]$*

Now generalize to allow different weights depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] \underbrace{F[i+u, j+v]}_{\text{Non-uniform weights}}$$

# Correlation filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called cross-correlation, denoted  $G = H \otimes F$

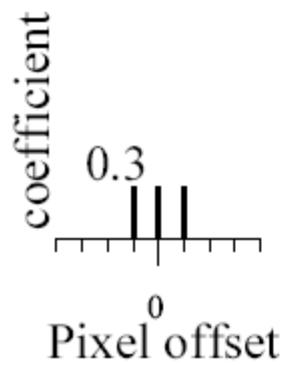
Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter “kernel” or “mask”  $H[u, v]$  is the prescription for the weights in the linear combination.

# Averaging Filter

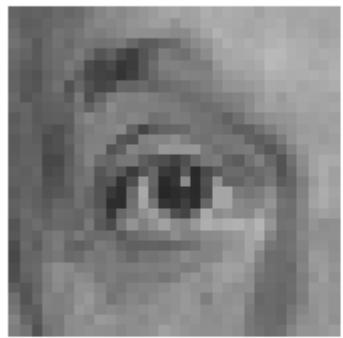


original

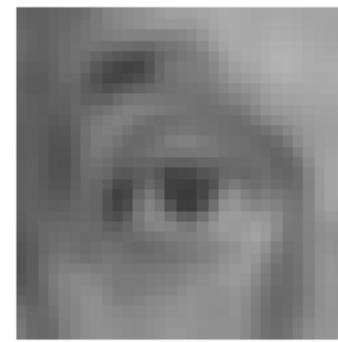
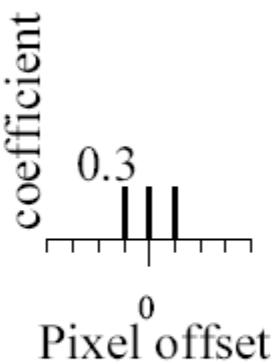


?

# Averaging Filter

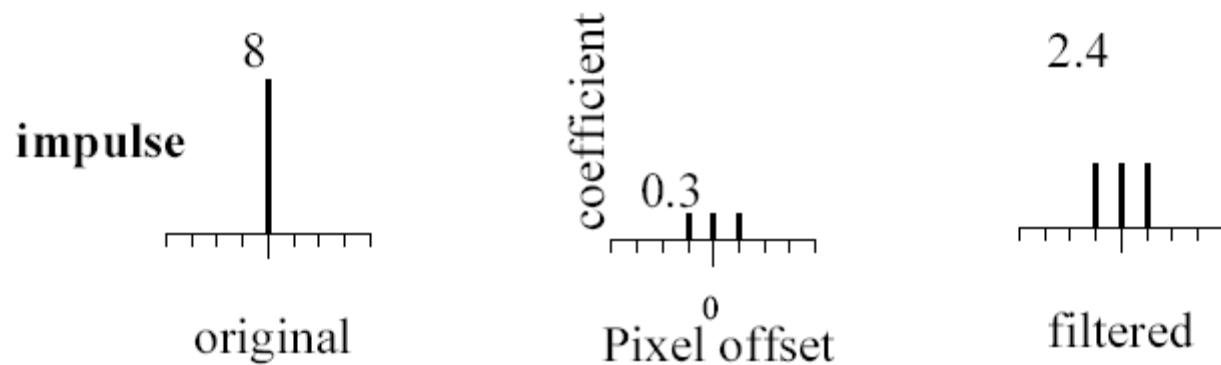


original

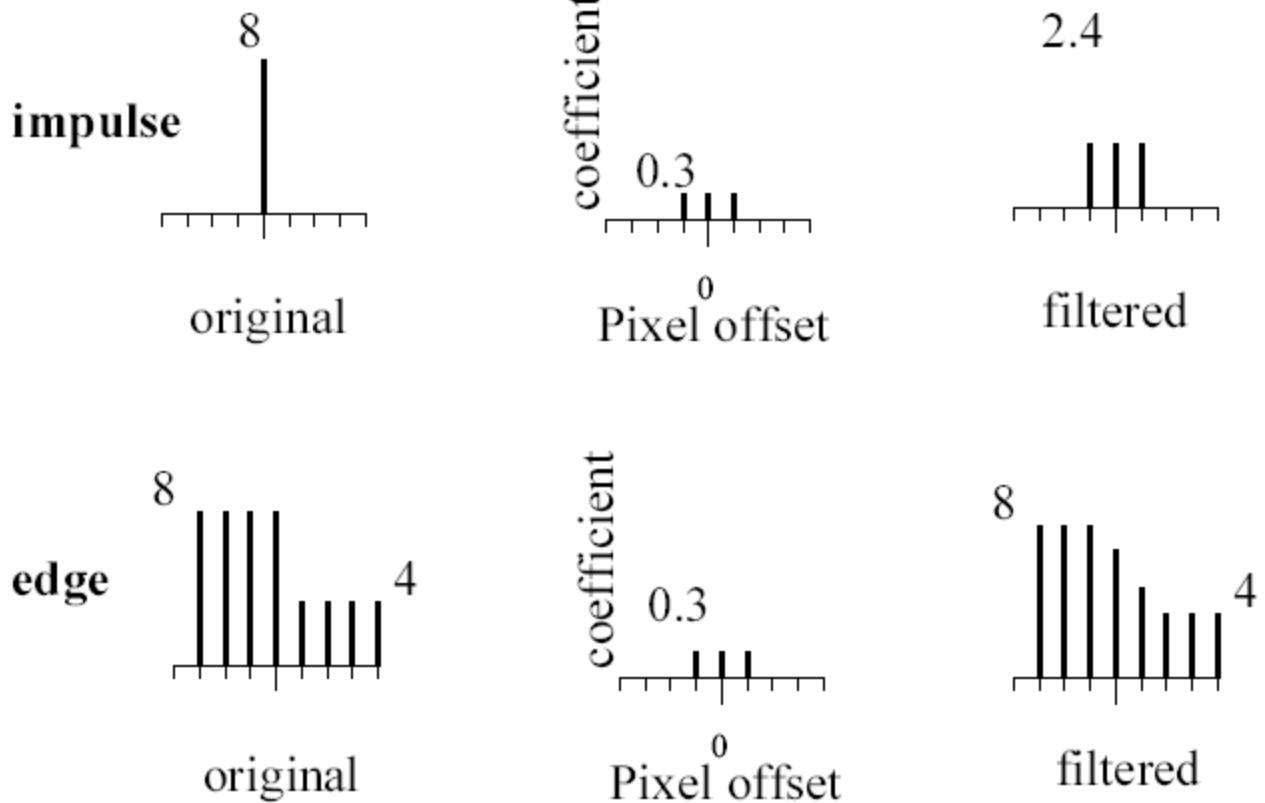


Blurred (filter applied in both dimensions).

# Averaging Filter



# Averaging Filter



# Averaging filter

- What values belong in the kernel  $H$  for the moving average example?

F[x, y]									
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



H[u, v]									
1	1	1	0	0	0	0	0	0	0
1	?	1	0	0	0	0	0	0	0
1	1	1	0	0	0	0	0	0	0

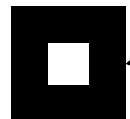
$$\frac{1}{9}$$

“box filter”

G[x, y]									
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0
0	10	20	30	30	0	0	0	0	0

$$G = H \otimes F$$

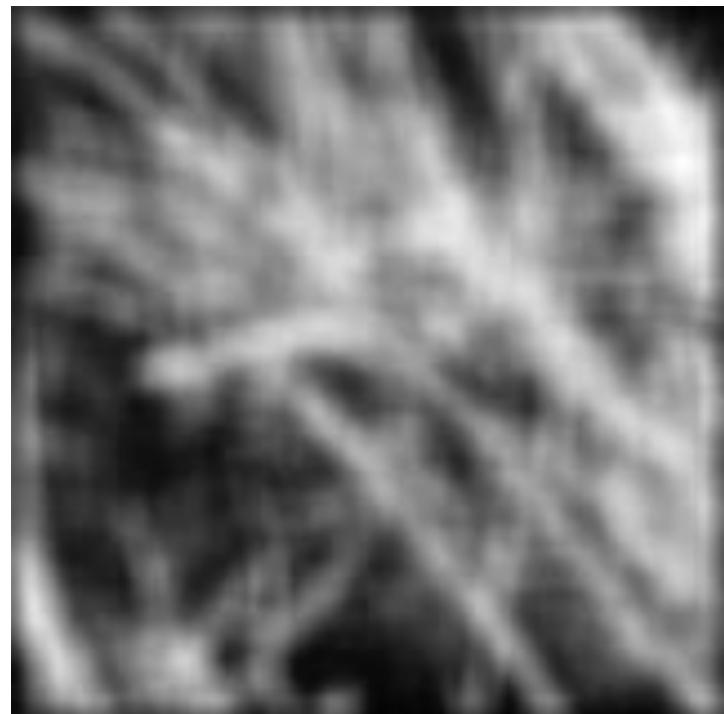
# Smoothing by averaging



depicts box filter:  
white = high value, black = low value

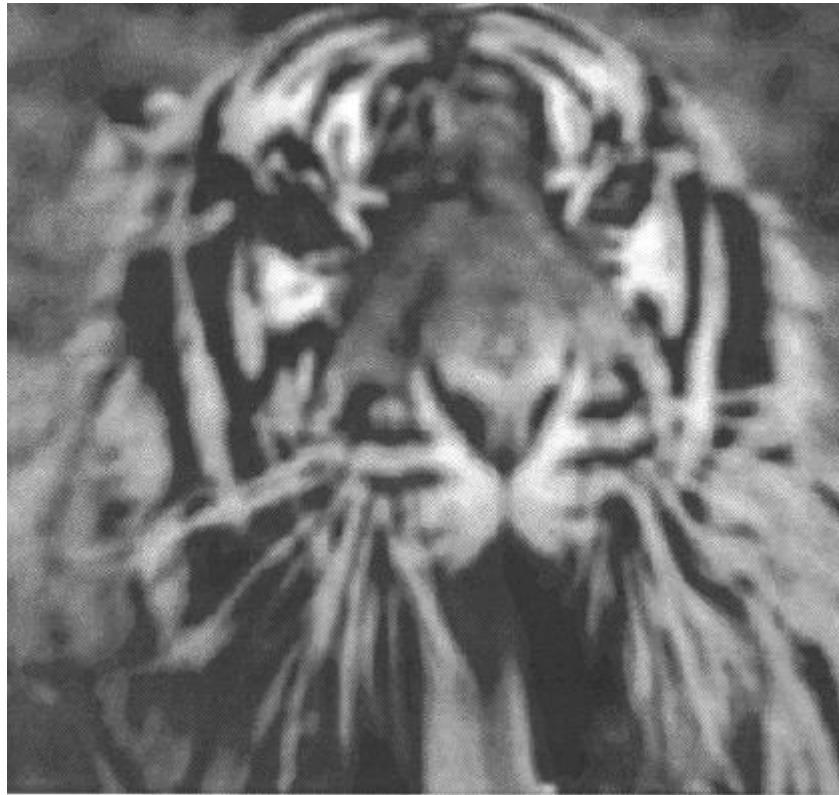


original



filtered

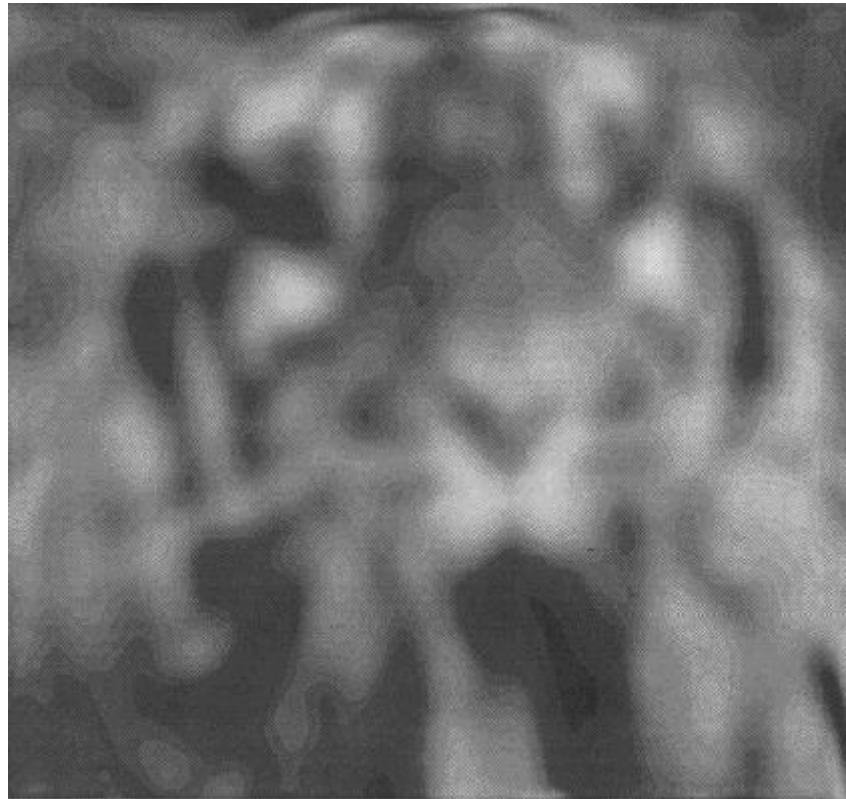
# Example



# Example



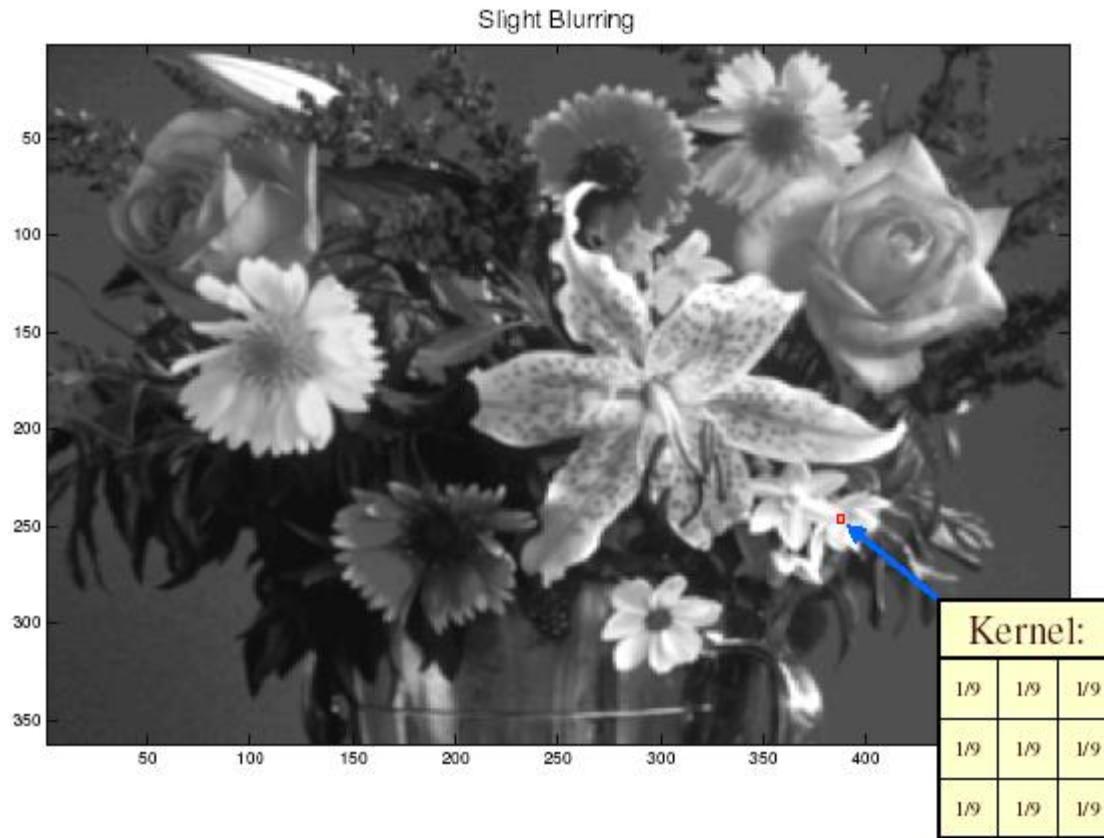
# Example



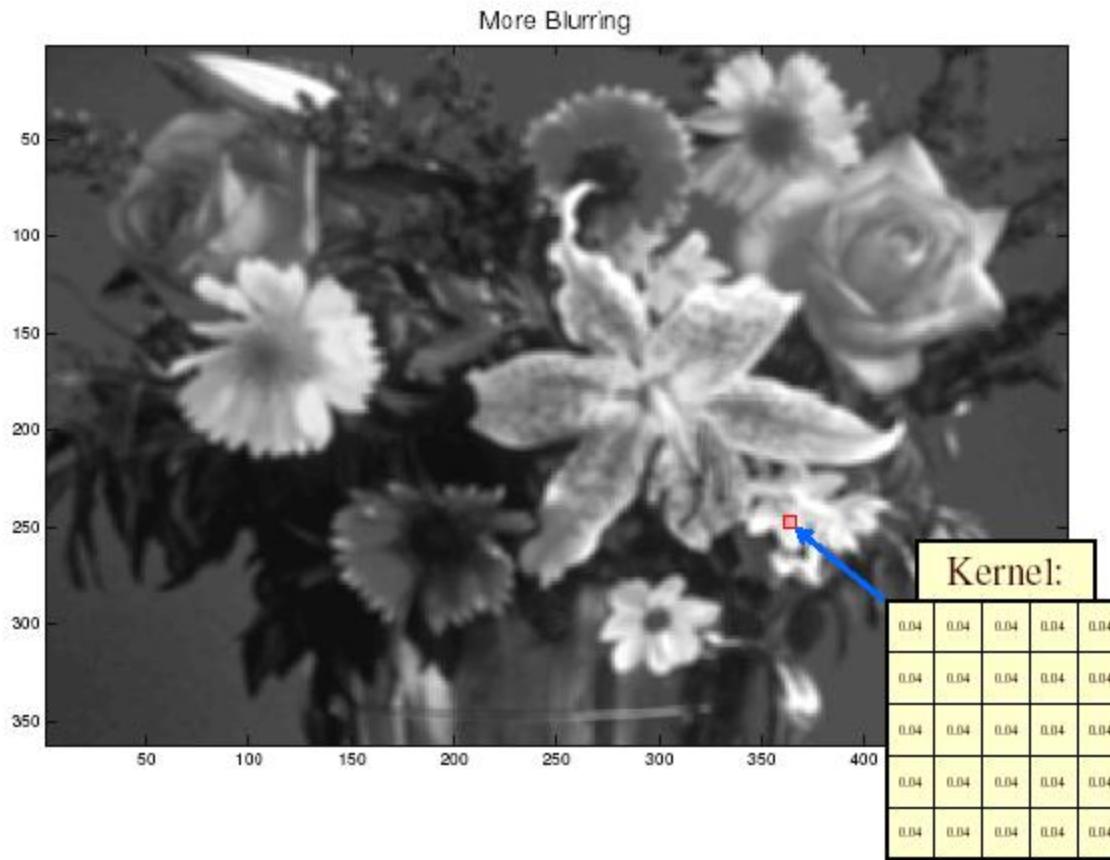
# Smoothing by averaging



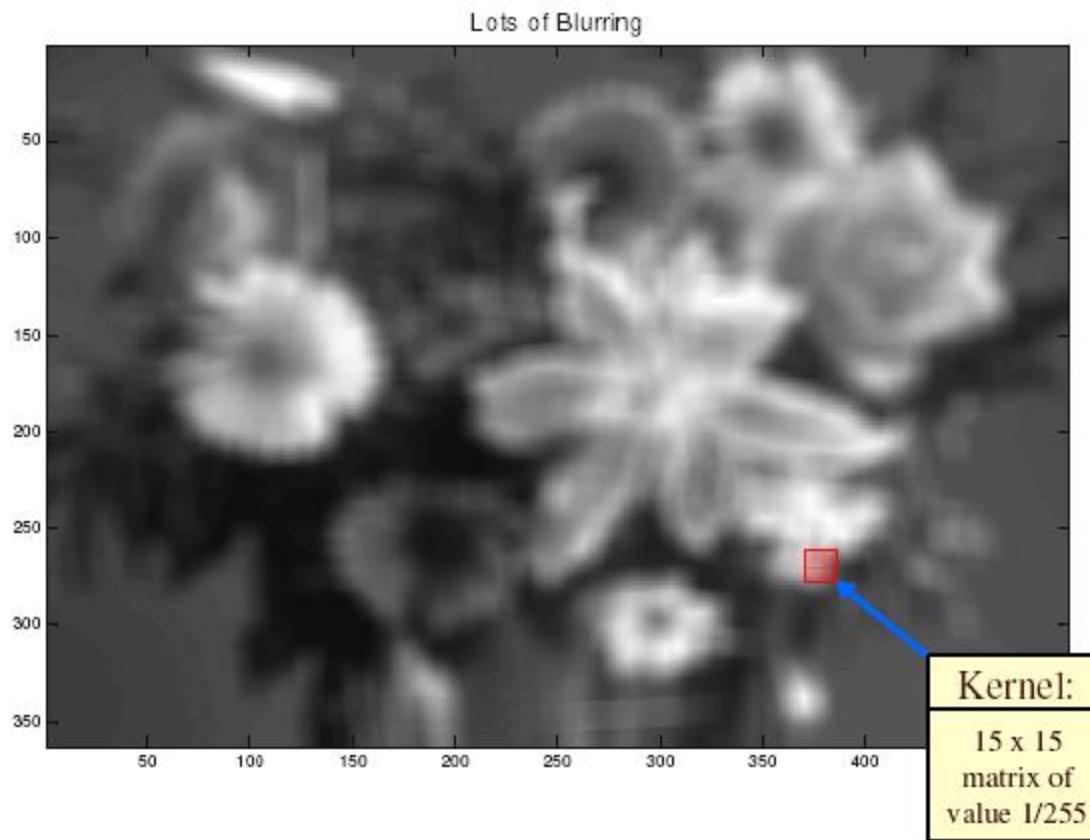
# Smoothing by averaging



# Smoothing by averaging

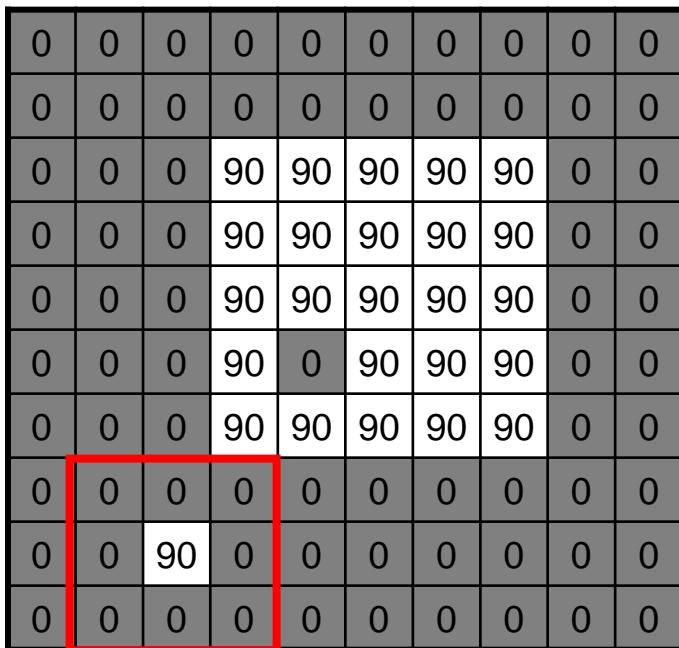


# Smoothing by averaging



# Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?



$F[x, y]$

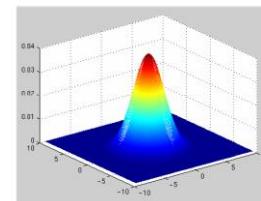
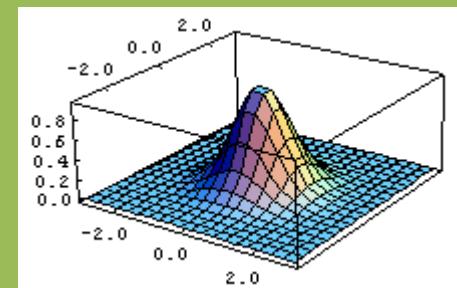
$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$

$H[u, v]$

A weighted average that weights pixels at its center much more strongly than its boundaries

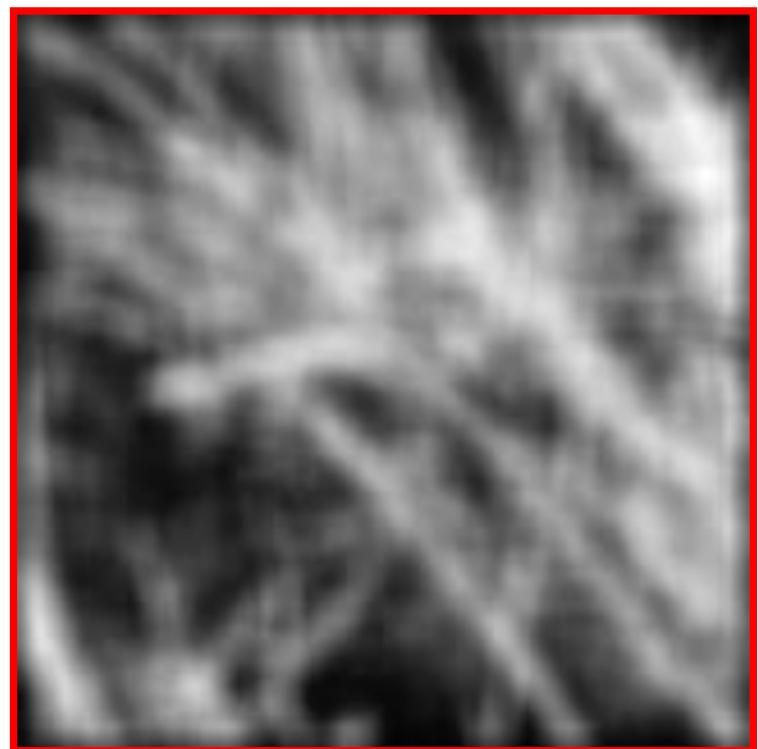
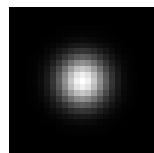
This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$

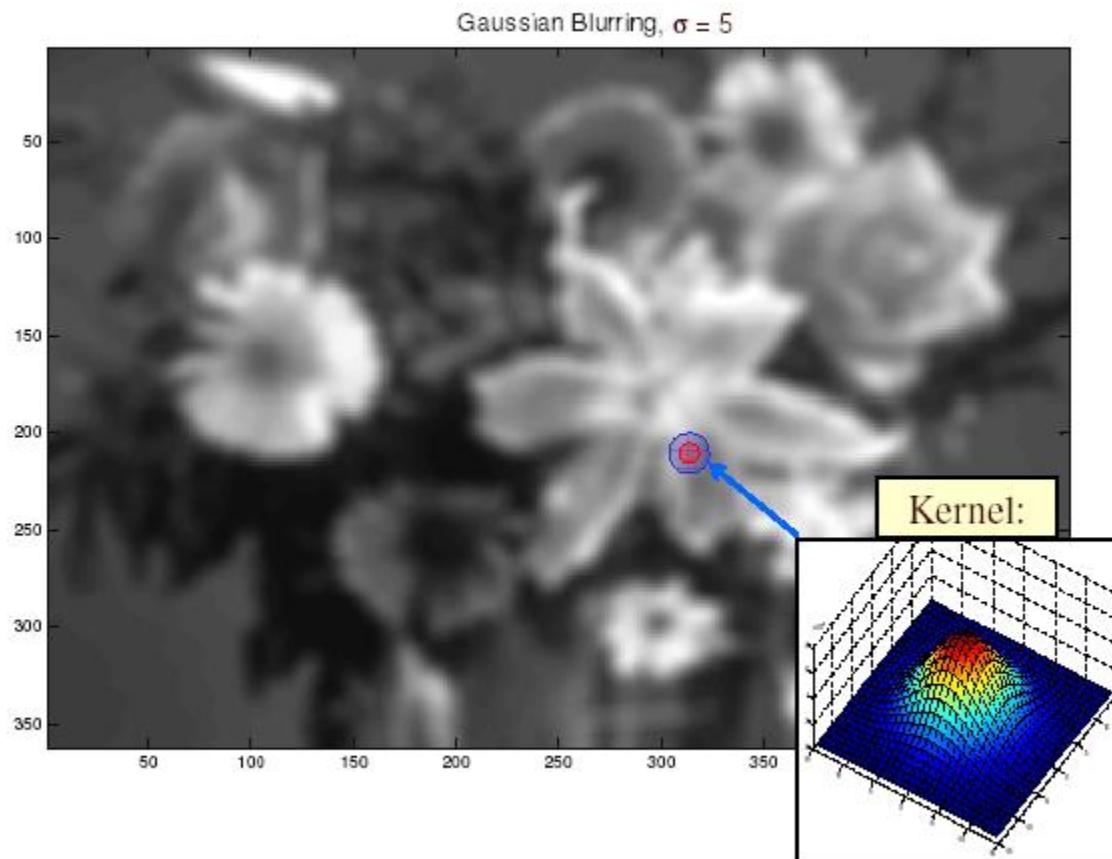


Source: S. Seitz

# Smoothing with a Gaussian



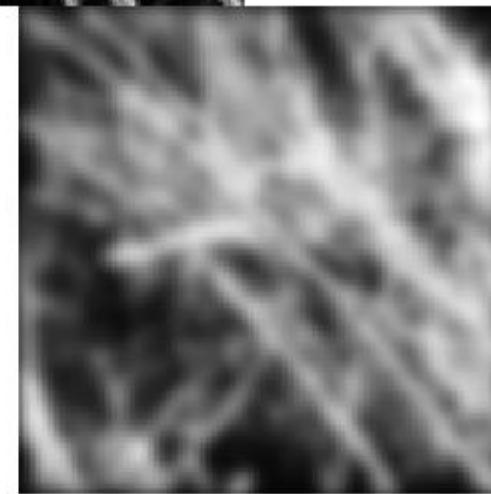
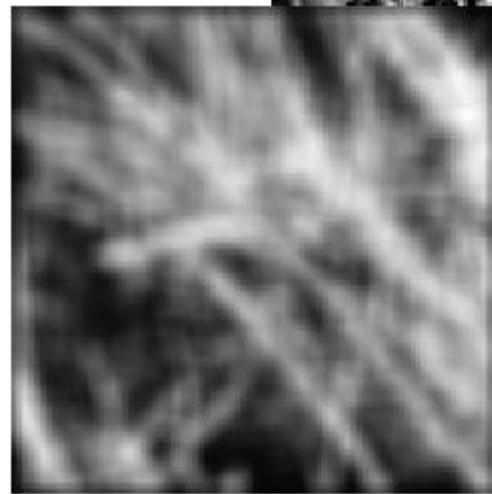
# Smoothing with a Gaussian



# Smoothing with a Gaussian

Result of blurring using  
a uniform local model

Produces a set of  
narrow vertical  
horizontal and vertical  
bars – ringing effect

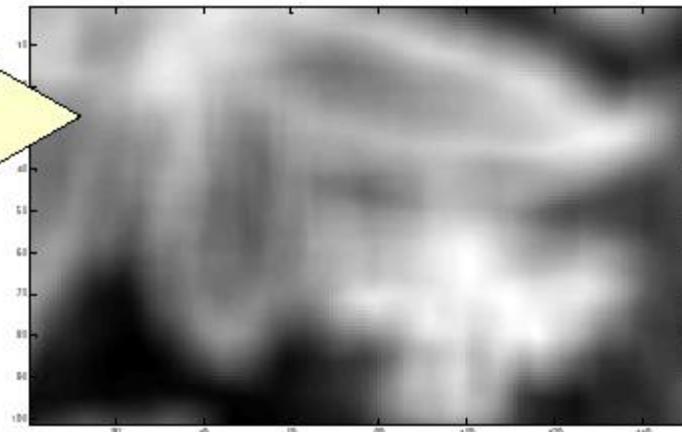
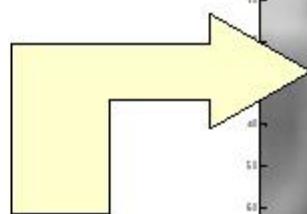


Result of blurring  
using a set of  
Gaussian weights

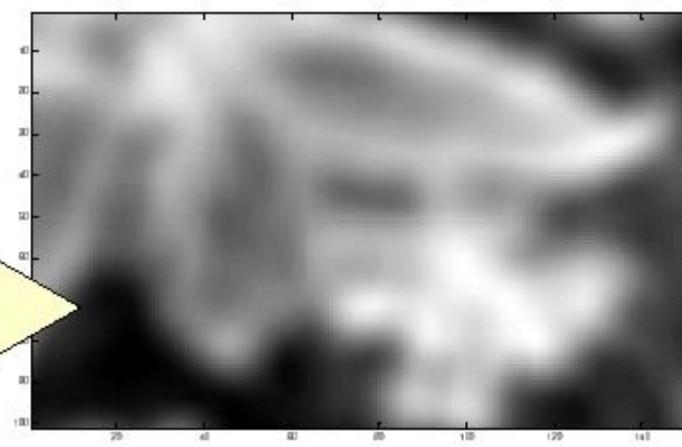
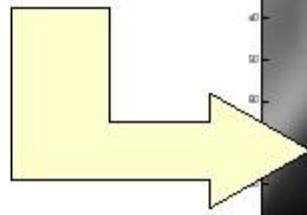


# Smoothing with a Gaussian

Simple  
Averaging

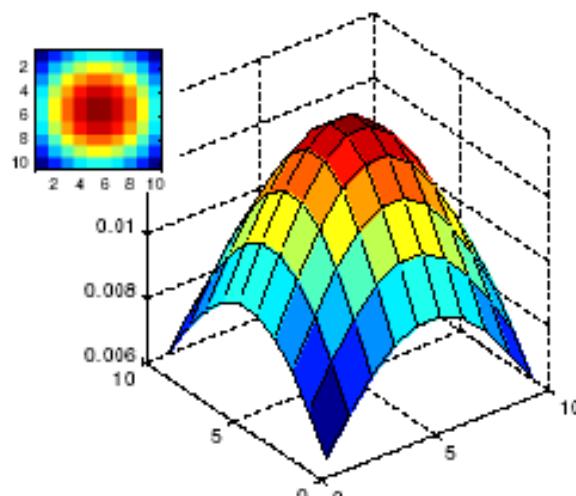


Gaussian  
Smoothing

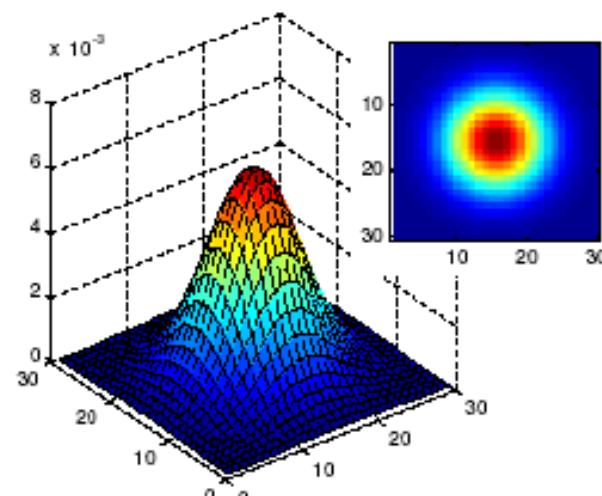


# Gaussian filters

- What parameters matter here?
- **Size** of kernel or mask
  - Note, Gaussian function has infinite support, but discrete filters use finite kernels



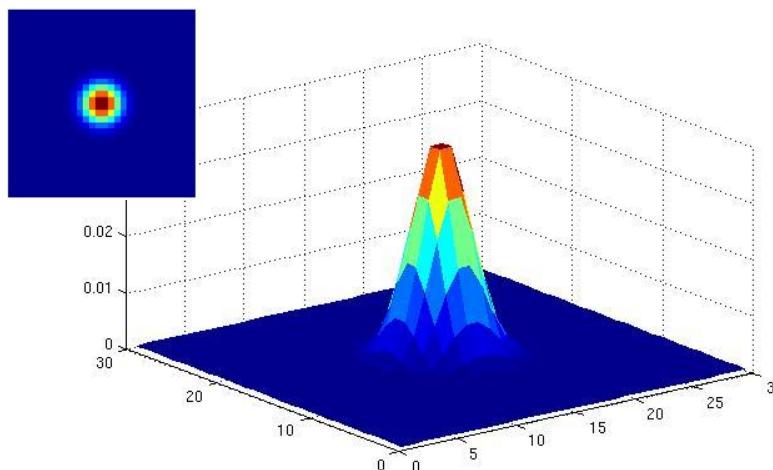
$\sigma = 5$  with 10  
x 10 kernel



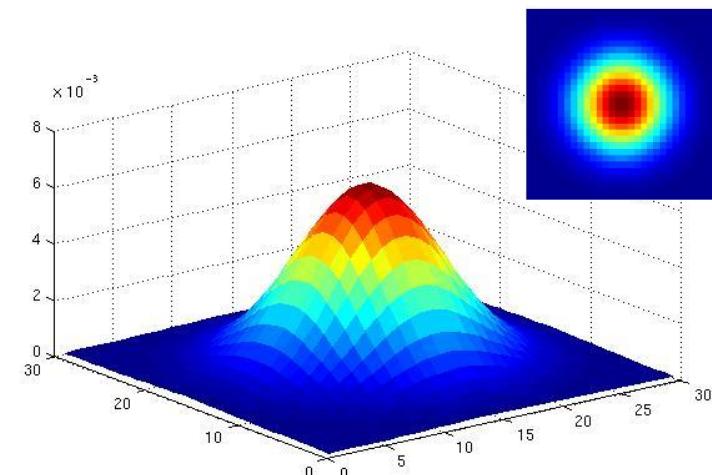
$\sigma = 5$  with 30  
x 30 kernel

# Gaussian filters

- What parameters matter here?
- **Variance of Gaussian:** determines extent of smoothing



$\sigma = 2$  with 30  
x 30 kernel



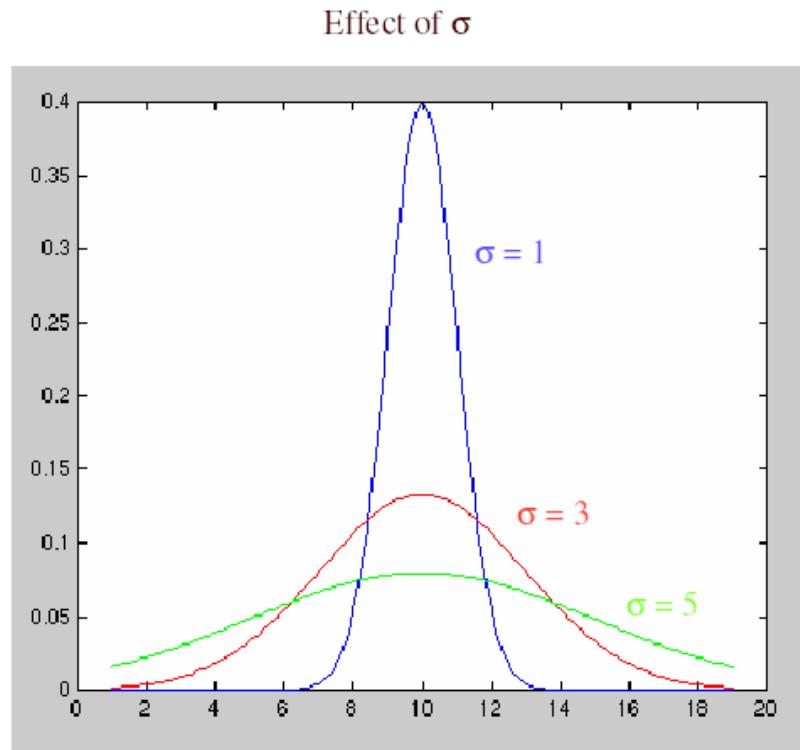
$\sigma = 5$  with 30  
x 30 kernel

# Smoothing with a Gaussian

If  $\sigma$  is small : the smoothing will have little effect

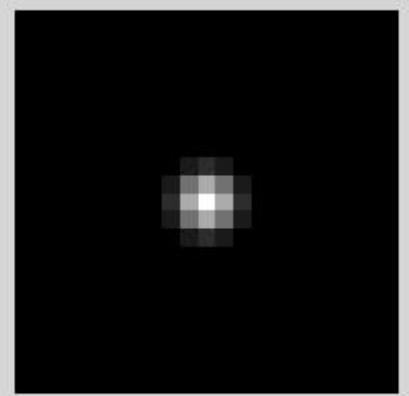
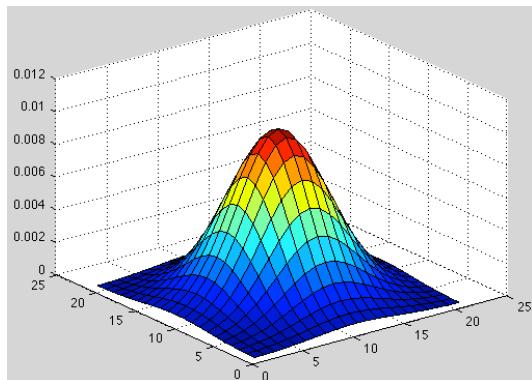
If  $\sigma$  is larger : neighboring pixels will have larger weights resulting in consensus of the neighbors

If  $\sigma$  is very large : details will disappear along with the noise



# Gaussian filter

$$G(x, y; \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$\sigma=1$



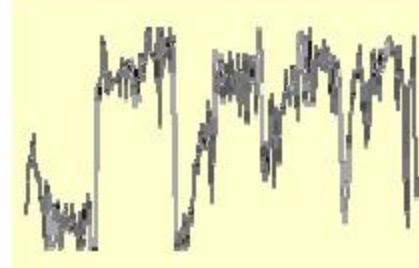
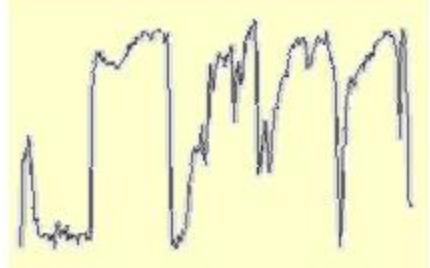
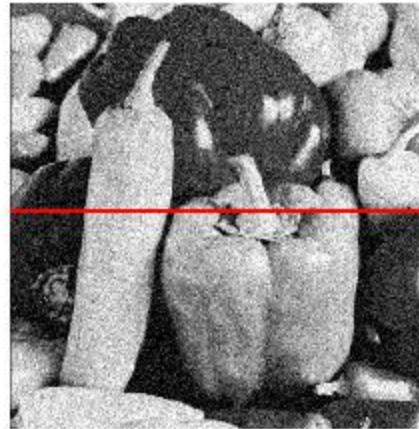
$\sigma=2$



$\sigma=4$

# Gaussian smoothing to remove noise

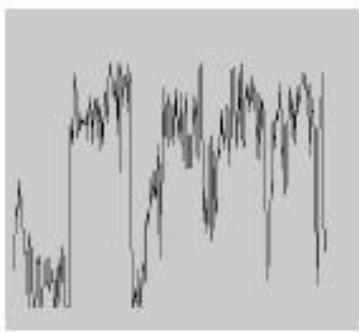
Image  
Noise



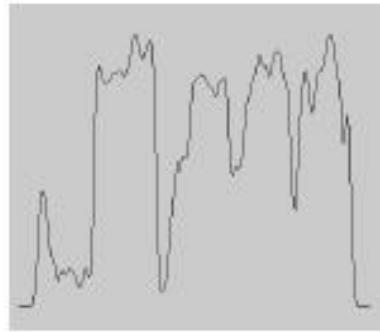
$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:  
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

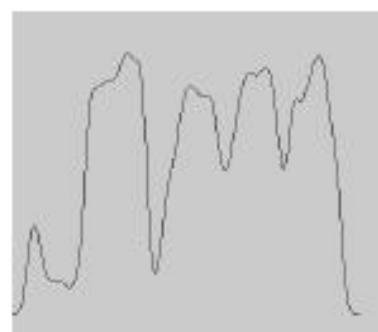
# Gaussian smoothing to remove noise



No smoothing



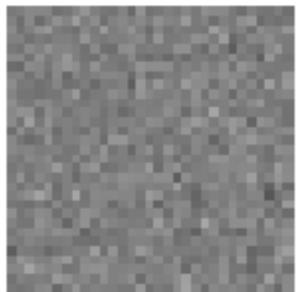
$\sigma = 2$



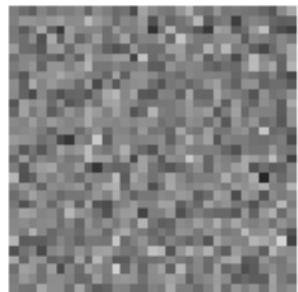
$\sigma = 4$

# Smoothing with a Gaussian

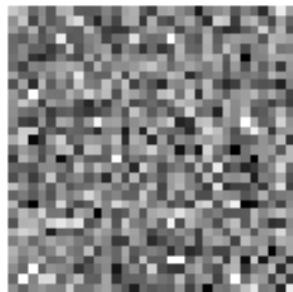
$\sigma=0.05$



$\sigma=0.1$



$\sigma=0.2$

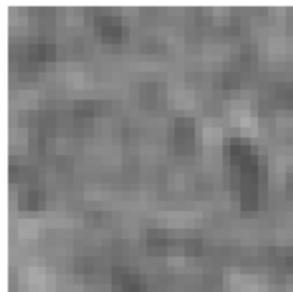
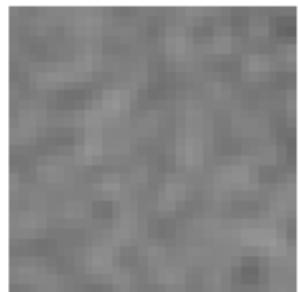
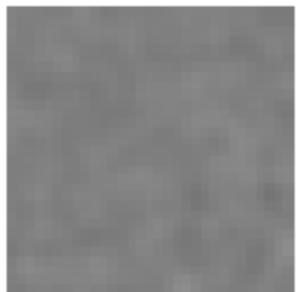


no  
smoothing

## The effects of smoothing

Each row shows smoothing with gaussians of different width; each column shows different realizations of an image of gaussian noise.

$\sigma=1$  pixel

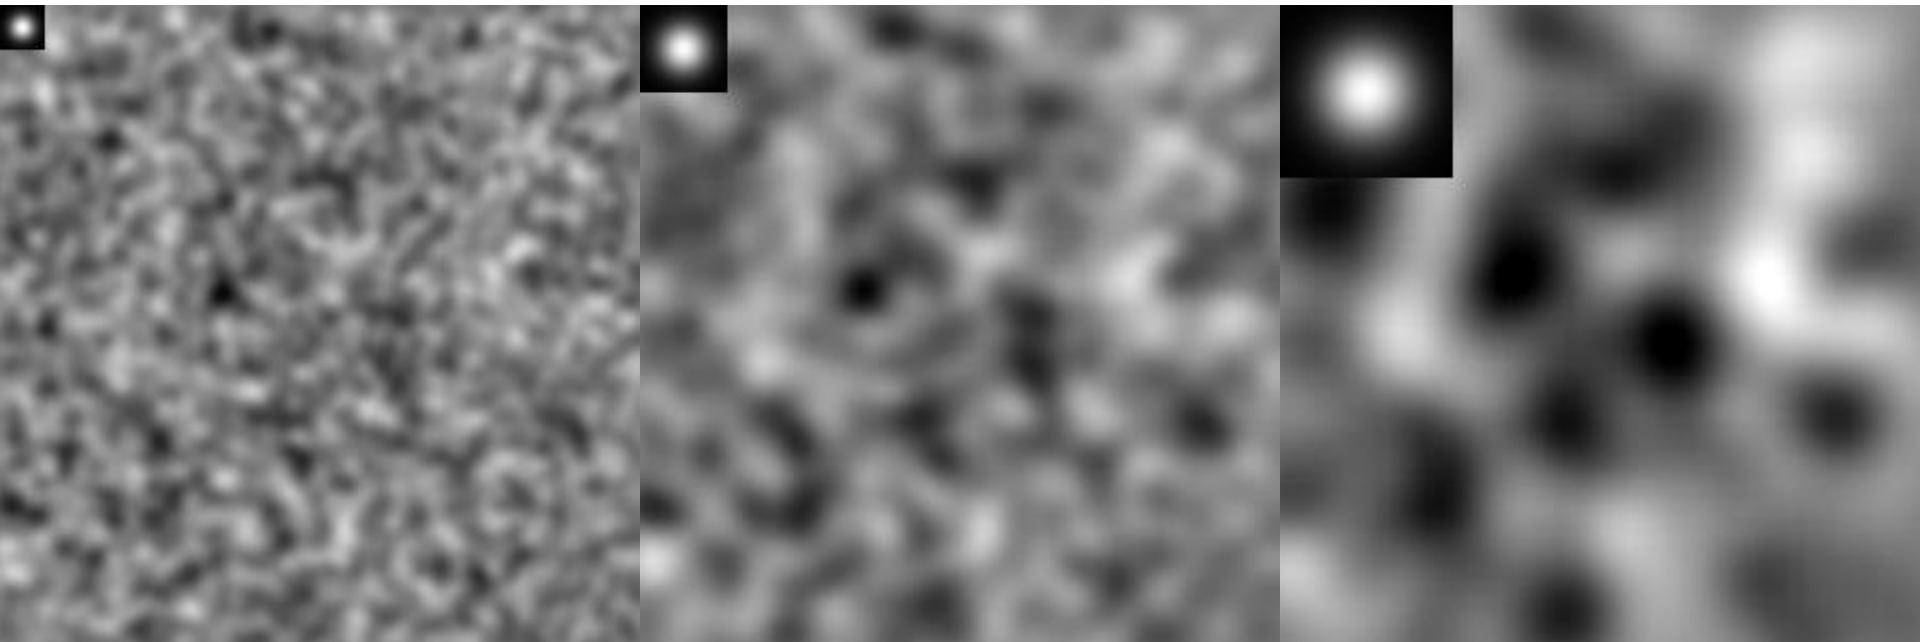


$\sigma=2$  pixels



# Smoothing with a Gaussian

- Filtered noise is sometimes useful
  - looks like some natural textures, can be used to simulate fire, etc.



# Gaussian kernel

$$g(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{(x^2 + y^2)}{2\sigma^2}\right)$$

$$\begin{matrix} 0.0751 & 0.1238 & 0.0751 \\ 0.1238 & 0.242 & 0.1238 \\ 0.0751 & 0.1238 & 0.0751 \end{matrix}$$

Gaussian is an approximation to the binomial distribution.

$$a_{nr} \equiv \frac{n!}{r!(n-r)!} \equiv \binom{n}{r}$$

Can approximate Gaussian using binomial coefficients.

1X3 filter: n=(3-1)=2, r=0,1,2

n = number of elements in the 1D filter minus 1

r = position of element in the filter kernel (0, 1, 2...)

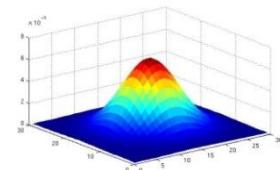
1	2	1
---	---	---

$$g = 1/4 \quad \boxed{1 \ 2 \ 1}$$

$$g' g = \begin{matrix} 0.0625 & 0.1250 & 0.0625 \\ 0.1250 & 0.2500 & 0.1250 \\ 0.0625 & 0.1250 & 0.0625 \end{matrix}$$

# Matlab

```
>> hsize = 10;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```



```
>> mesh(h);
```



```
>> imagesc(h);
```

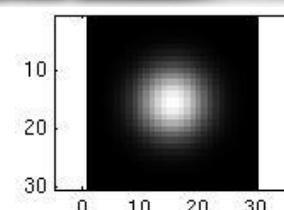
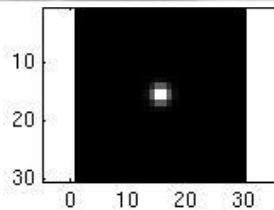
```
>> outim = imfilter(im, h);  
>> imshow(outim);
```



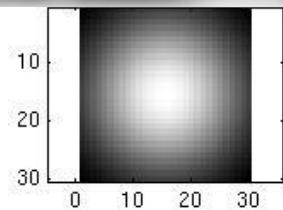
outim

# Smoothing with a Gaussian

Parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



...



```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

# Convolution

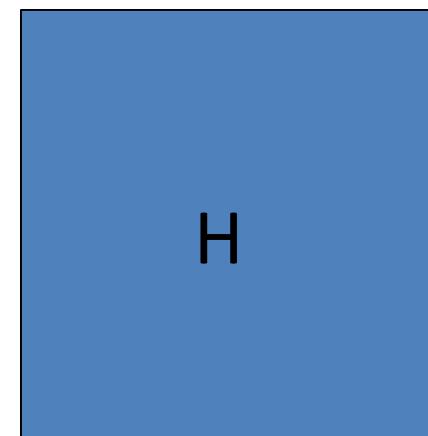
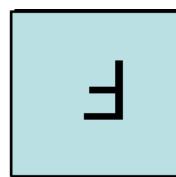
- Convolution:
  - Flip the filter in both dimensions (bottom to top, right to left)
  - Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$



*Notation for  
convolution  
operator*



# Convolution vs. correlation

Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

Cross-correlation

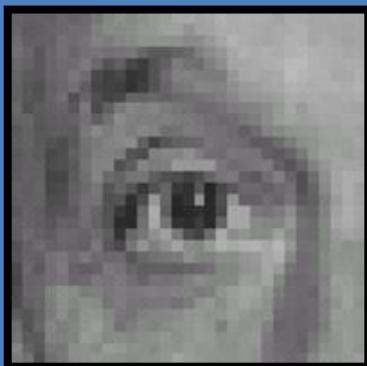
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$

For a Gaussian or box filter, how will the outputs differ?

If the input is an impulse signal, how will the outputs differ?

# Predict the filtered outputs



$$\begin{matrix} * & \begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix} & = ? \end{matrix}$$



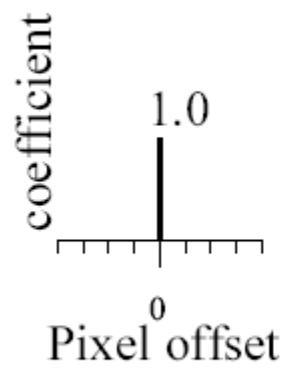
$$\begin{matrix} * & \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{matrix} & = ? \end{matrix}$$

$$\begin{matrix} * & \begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix} & - \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix} & = ? \end{matrix}$$

# Practice with linear filters



original

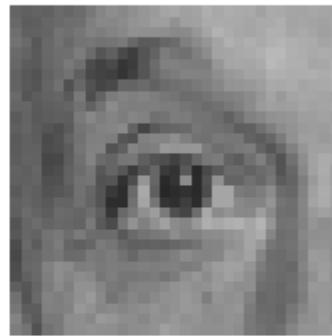
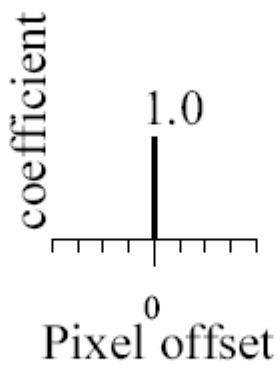


?

# Practice with linear filters



original



Filtered  
(no change)

# Practice with linear filters



Original

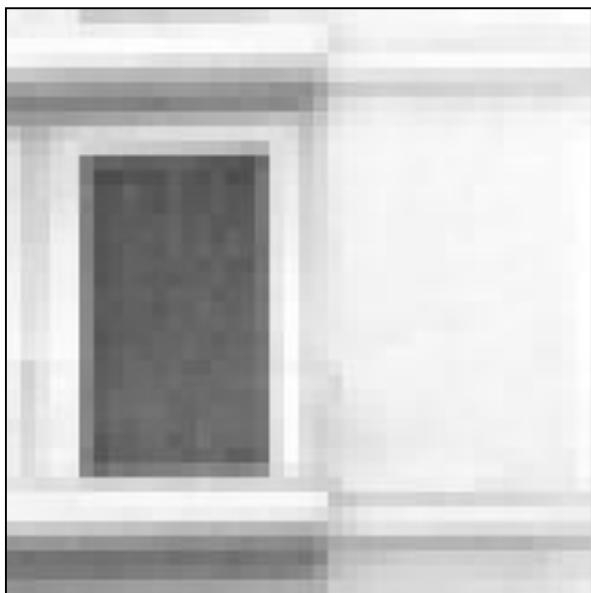
0	0	0
0	1	0
0	0	0



Filtered  
(no change)

# Impulse

$$f[m,n] = I \otimes g = \sum_{k,l} h[m-k, n-l] g[k, l]$$

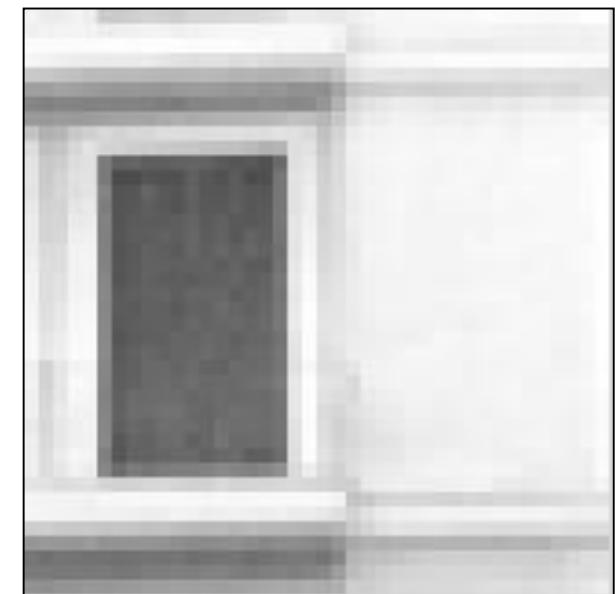


$g[m,n]$

$\otimes$

0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

$h[m,n]$



$f[m,n]$

# Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

# Practice with linear filters



Original

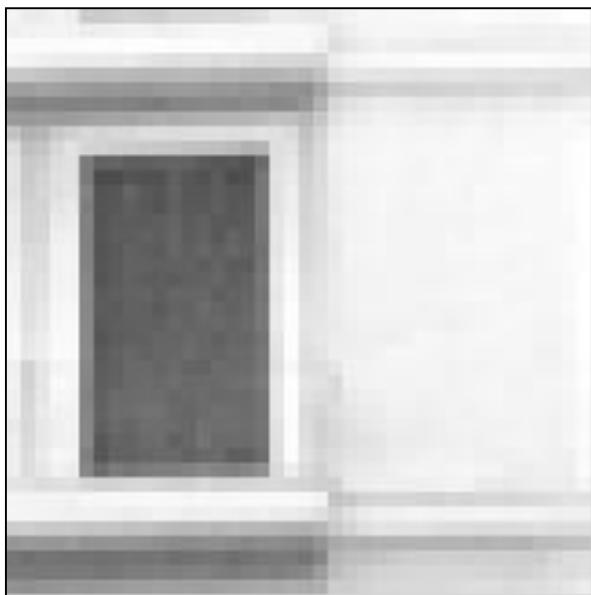
0	0	0
0	0	1
0	0	0



Shifted left  
by 1 pixel  
with  
correlation

# Shifts

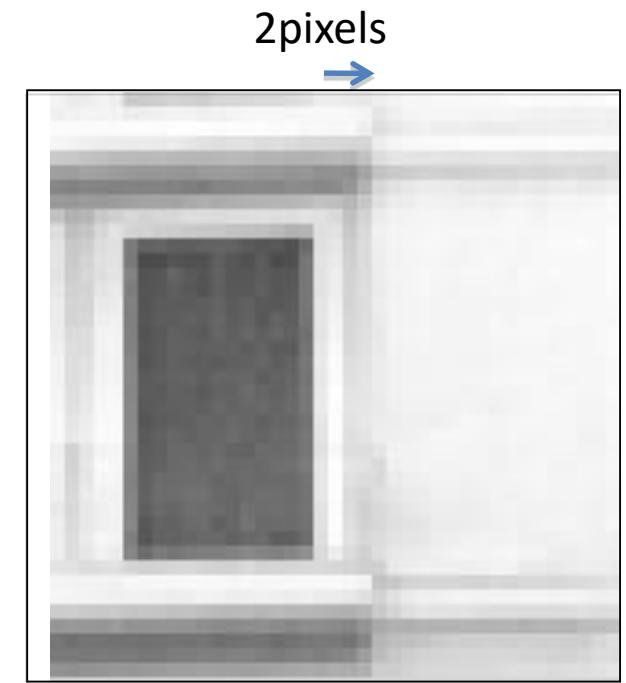
$$f[m, n] = I \otimes g = \sum_{k, l} h[m - k, n - l]g[k, l]$$



$\otimes$

0	0	0	0	0
0	0	0	0	0
0	0	0	0	1
0	0	0	0	0
0	0	0	0	0

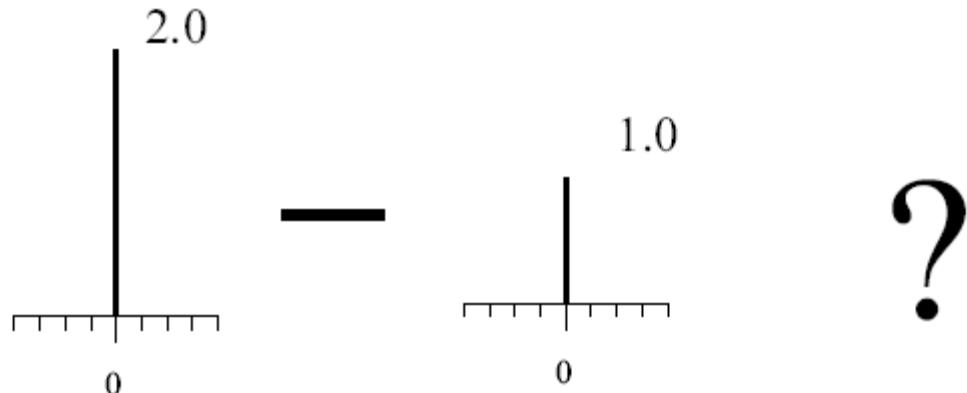
$h[m, n]$



$g[m, n]$

$f[m, n]$

# Practice with linear filters

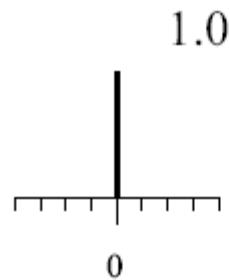
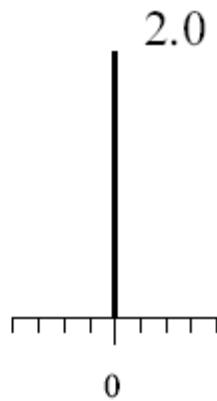


original

# Practice with linear filters



original

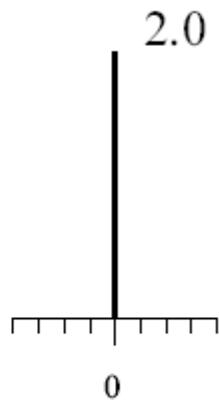


Filtered  
(no change)

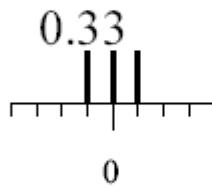
# Practice with linear filters



original



—

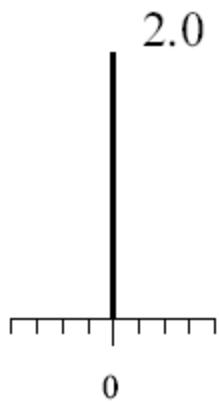


?

# Sharpening

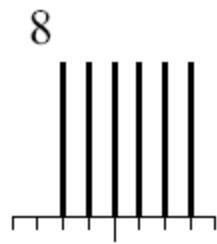


original

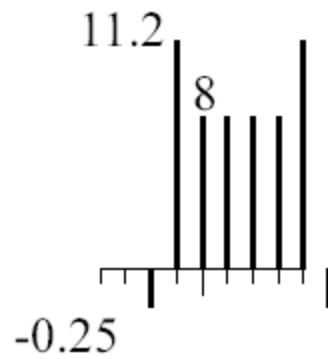
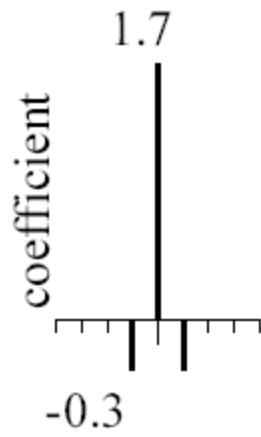


Sharpened  
original

# Sharpening



original



Sharpened  
(differences are  
accentuated; constant  
areas are left untouched).

# Practice with linear filters



Original

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

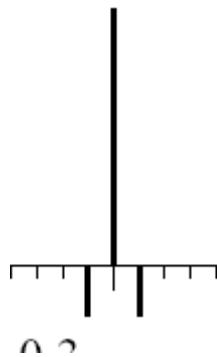
-

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

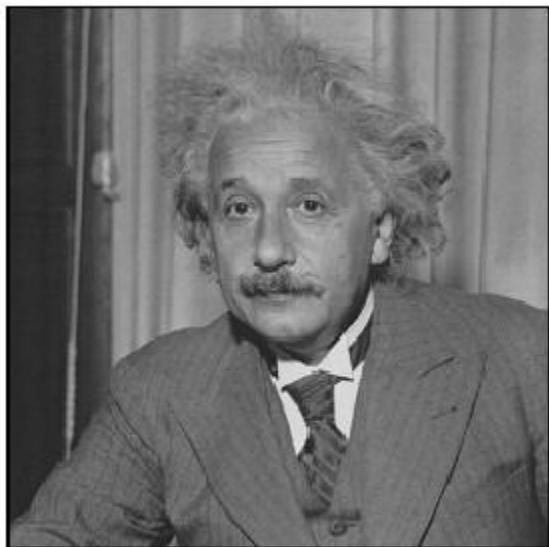


Sharpening filter

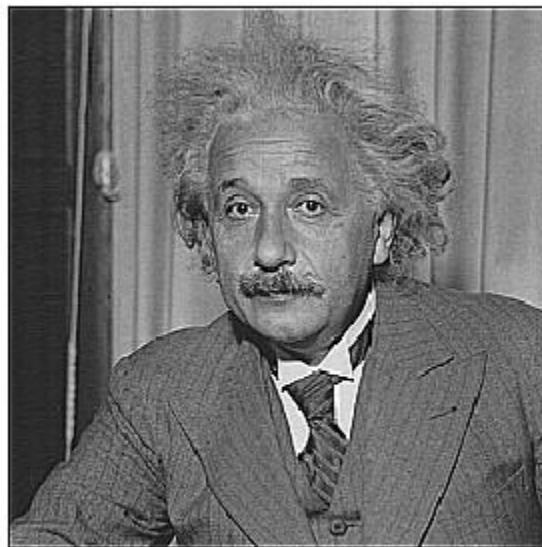
- Accentuates differences with local average



# Filtering examples: sharpening



before



after

# Rectangular filter



$g[m,n]$

$\otimes$



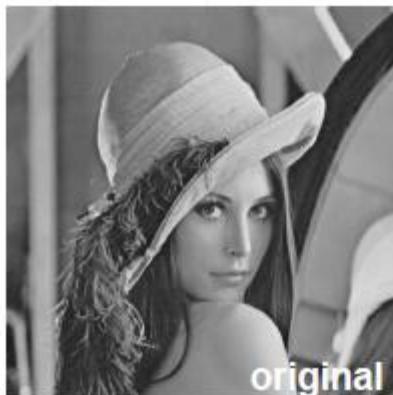
$h[m,n]$

=



$f[m,n]$

## What does blurring take away?



-



=



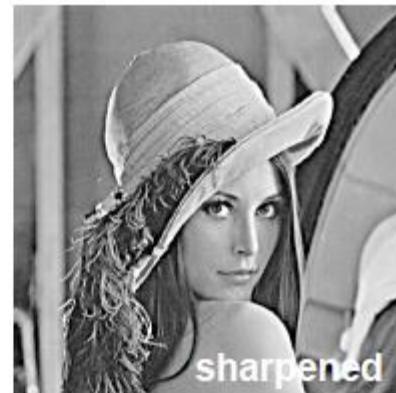
- Let's add it back:



+ a



=



# Rectangular filter



$g[m,n]$

$\otimes$             =

$h[m,n]$



$f[m,n]$

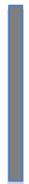
# Rectangular filter



$g[m,n]$

$\otimes$

$h[m,n]$

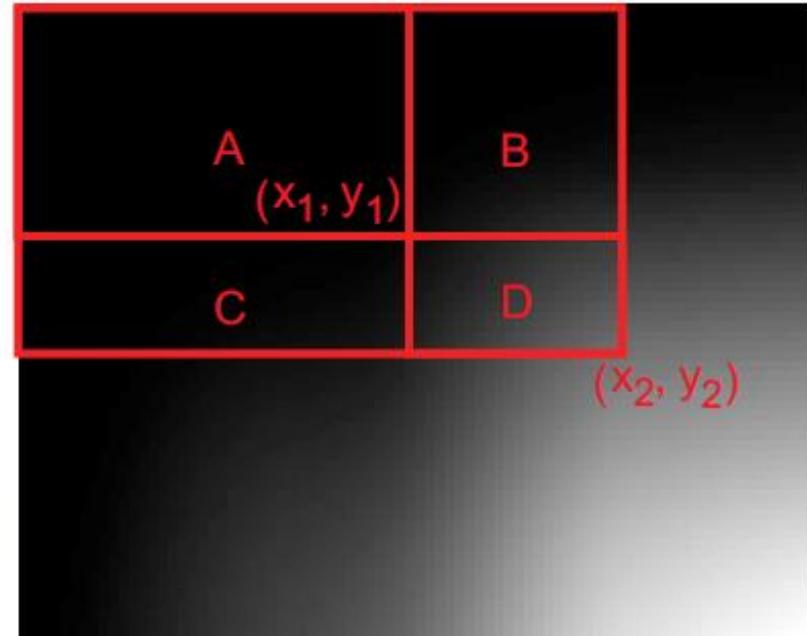


=



$f[m,n]$

# Integral image



# Shift invariant linear system

- **Shift invariant:**
  - Operator behaves the same everywhere, i.e. the value of the output depends on the pattern in the image neighborhood, not the position of the neighborhood.
- **Linear:**
  - Superposition:  $h * (f_1 + f_2) = (h * f_1) + (h * f_2)$
  - Scaling:  $h * (k f) = k (h * f)$

# Properties of convolution

- Linear & shift invariant

- Commutative:

$$f * g = g * f$$

- Associative

$$(f * g) * h = f * (g * h)$$

- Identity:

unit impulse  $e = [..., 0, 0, 1, 0, 0, ...]$ .  $f * e = f$

- Differentiation:

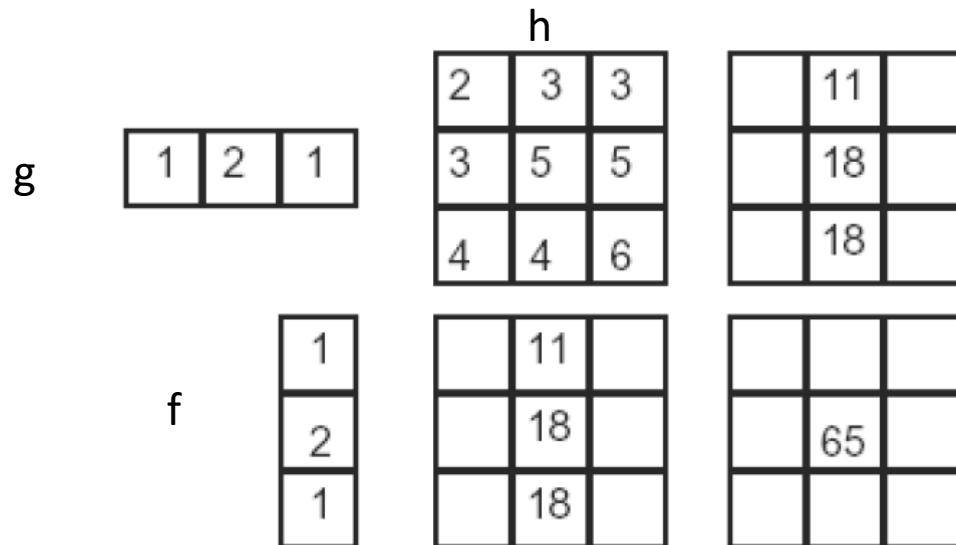
$$\frac{\partial}{\partial x} (f * g) = \frac{\partial f}{\partial x} * g$$

# Separability

- In some cases, filter is separable, and we can factor into two steps:
  - Convolve all rows
  - Convolve all columns

# Separability

- In some cases, filter is separable, and we can factor into two steps: e.g.,



What is the computational complexity advantage for a separable filter of size  $k \times k$ , in terms of number of operations per output pixel?

$$\begin{array}{c|c|c|c|c} 1 & \times & [1 \ 2 \ 1] & = & \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} & = 2 + 6 + 3 = 11 \\ \hline 2 & & & & \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} & = 6 + 20 + 10 = 36 \\ \hline 1 & & & & & = 4 + 8 + 6 = 18 \\ \hline & & & & & \hline & & & & & 65 \end{array}$$

# Advantages of separability

First convolve the image with a one dimensional horizontal filter

Then convolve the result of the first convolution with a one dimensional vertical filter

For a  $k \times k$  Gaussian filter, 2D convolution requires  $k^2$  operations per pixel

But using the separable filters, we reduce this to  $2k$  operations per pixel.

# Seperable Gaussian

$$g(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp(-x^2/(2\sigma^2))$$

$$g(y) = \frac{1}{\sqrt{2\pi}\sigma} \exp(-y^2/(2\sigma^2))$$

Product?

$$g(x, y) = \frac{1}{2\pi\sigma^2} \exp(-(x^2 + y^2)/(2\sigma^2))$$

# Advantages of Gaussians

- Convolution of a Gaussian with itself is another Gaussian
  - so we can first smooth an image with a small Gaussian
  - then, we convolve that smoothed image with another small Gaussian and the result is equivalent to smoother the original image with a larger Gaussian.
- If we smooth an image with a Gaussian having sd  $\sigma$  twice, then we get the same result as smoothing the image with a Gaussian having standard deviation  $(2\sigma)^{1/2}$

# Effect of smoothing filters

5x5

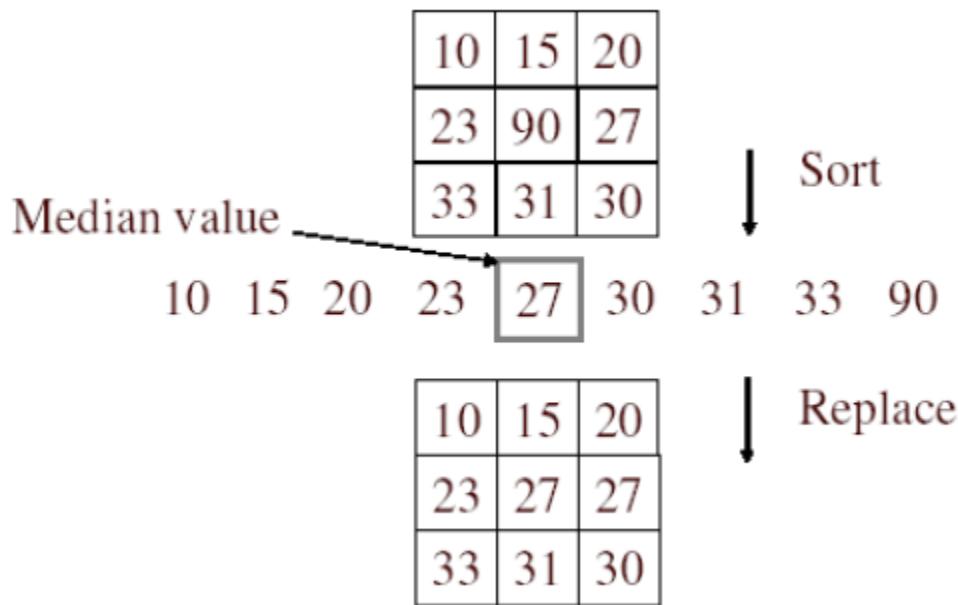


Additive Gaussian noise



Salt and pepper noise

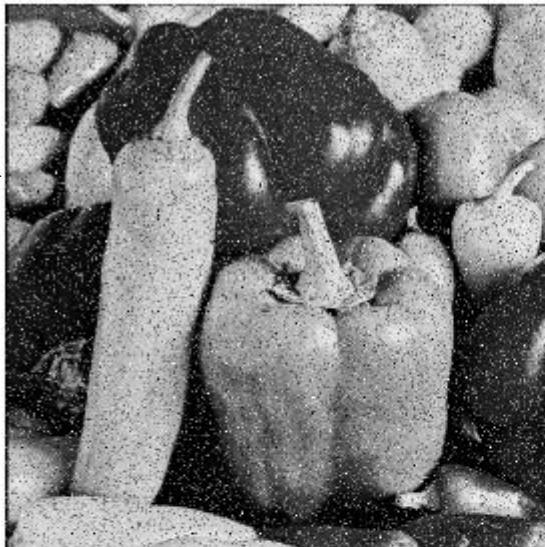
# Median filter



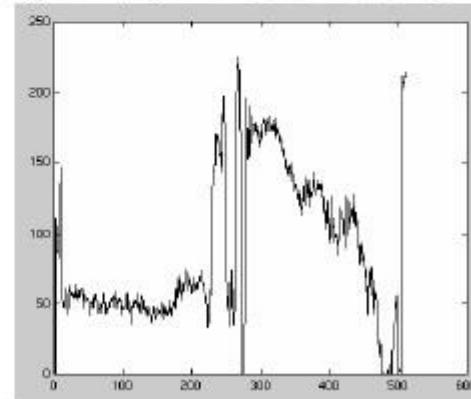
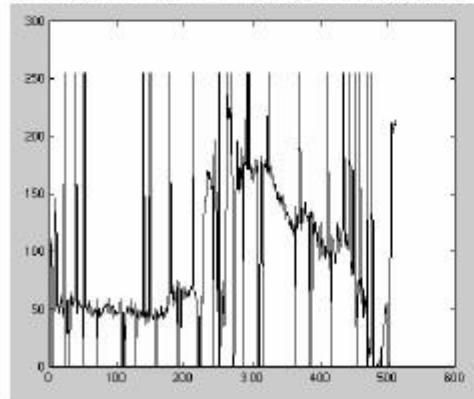
- No new pixel values introduced
- Removes spikes: good for impulse, salt & pepper noise

# Median filter

Salt and  
pepper noise



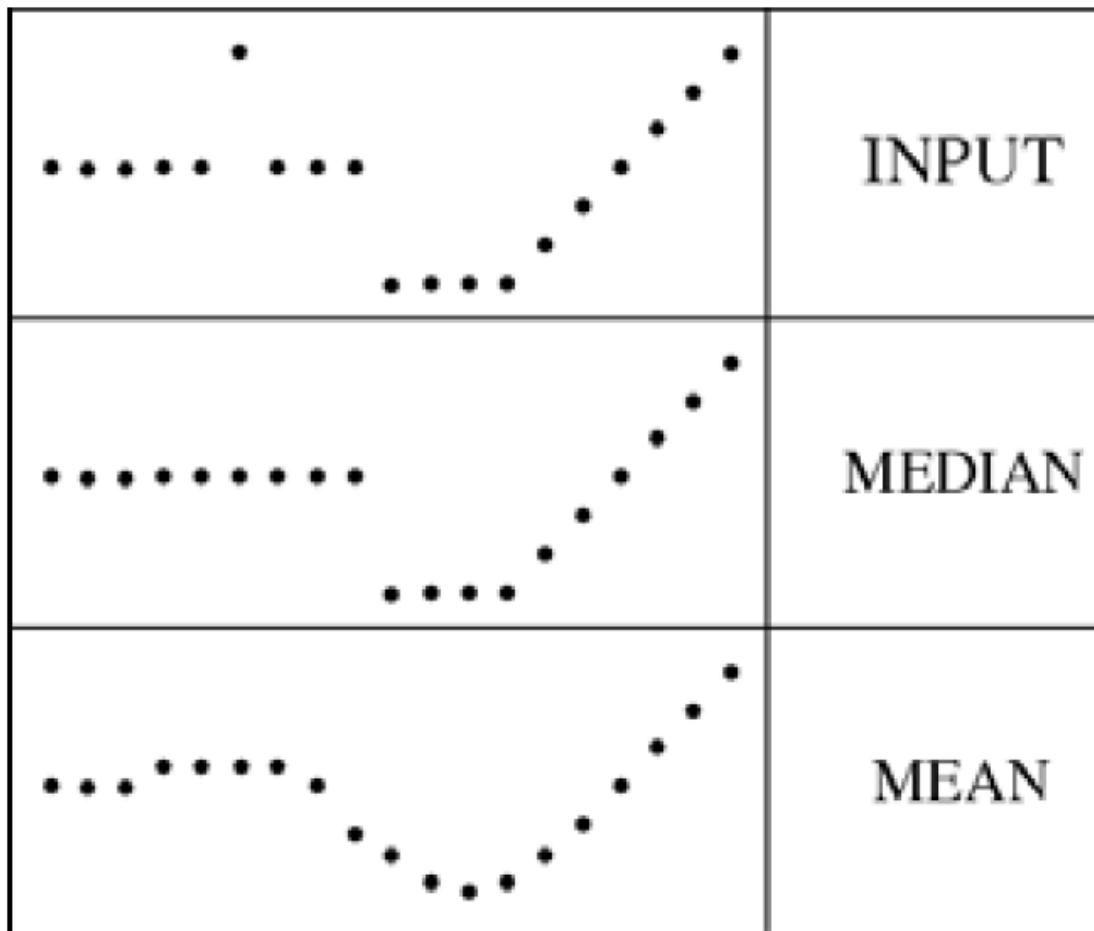
Median  
filtered



Plots of a row of the image

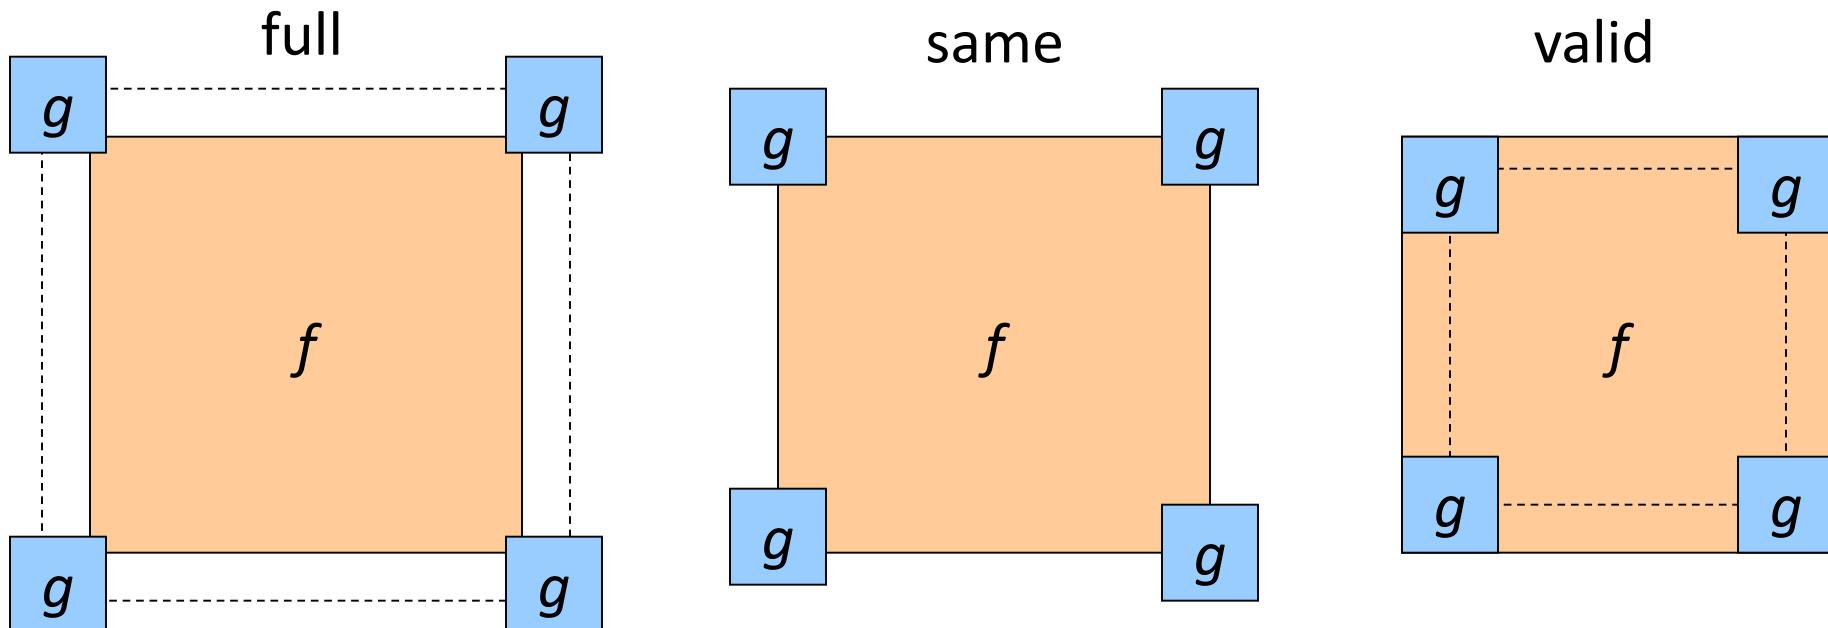
# Median filter

- Median filter is edge preserving



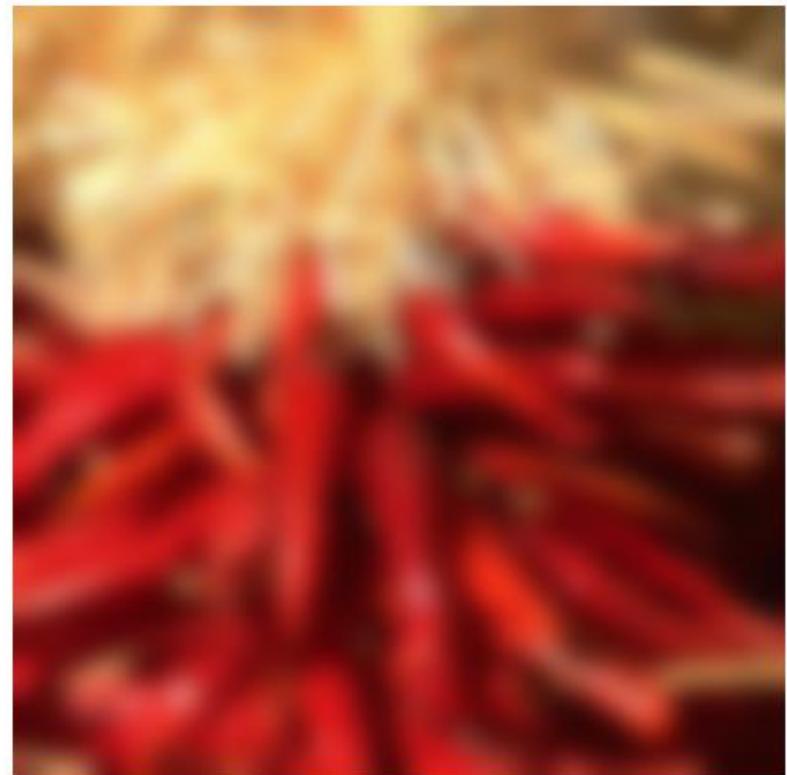
# Boundary issues

- What is the size of the output?
- MATLAB: `filter2(g, f, shape)`
  - `shape = 'full'`: output size is sum of sizes of  $f$  and  $g$
  - `shape = 'same'`: output size is same as  $f$
  - `shape = 'valid'`: output size is difference of sizes of  $f$  and  $g$



# Boundary issues

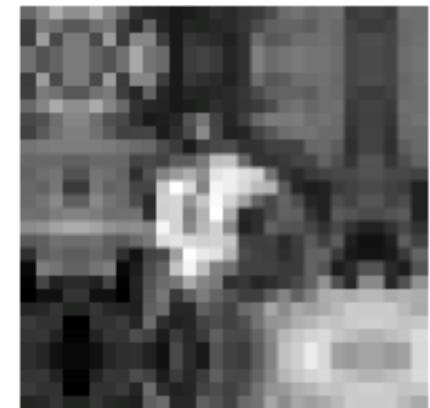
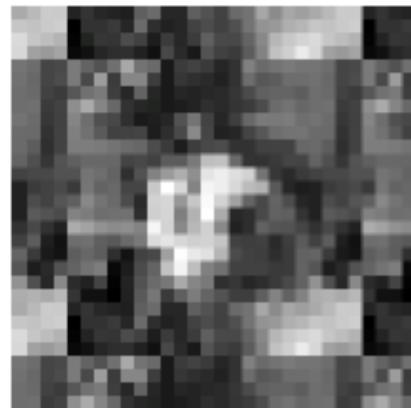
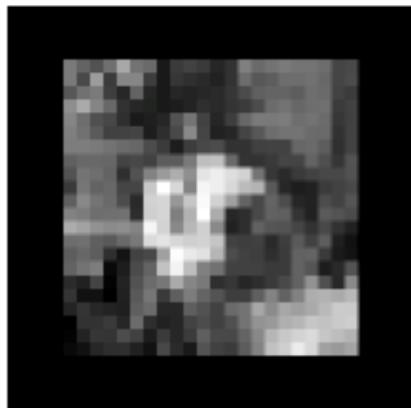
- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge



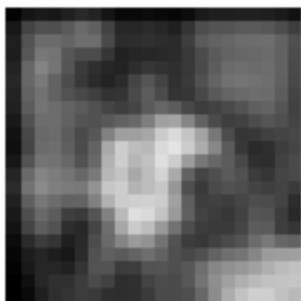
# Boundary issues

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods (MATLAB):
    - clip filter (black): `imfilter(f, g, 0)`
    - wrap around: `imfilter(f, g, 'circular')`
    - copy edge: `imfilter(f, g, 'replicate')`
    - reflect across edge: `imfilter(f, g, 'symmetric')`

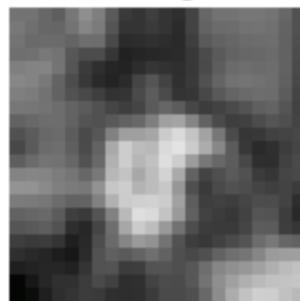
# Borders



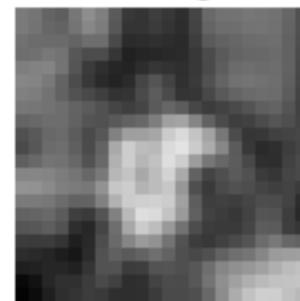
zero



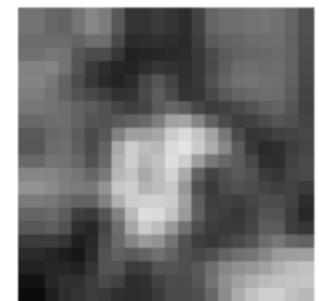
wrap



clamp



mirror



blurred: zero

normalized zero

clamp

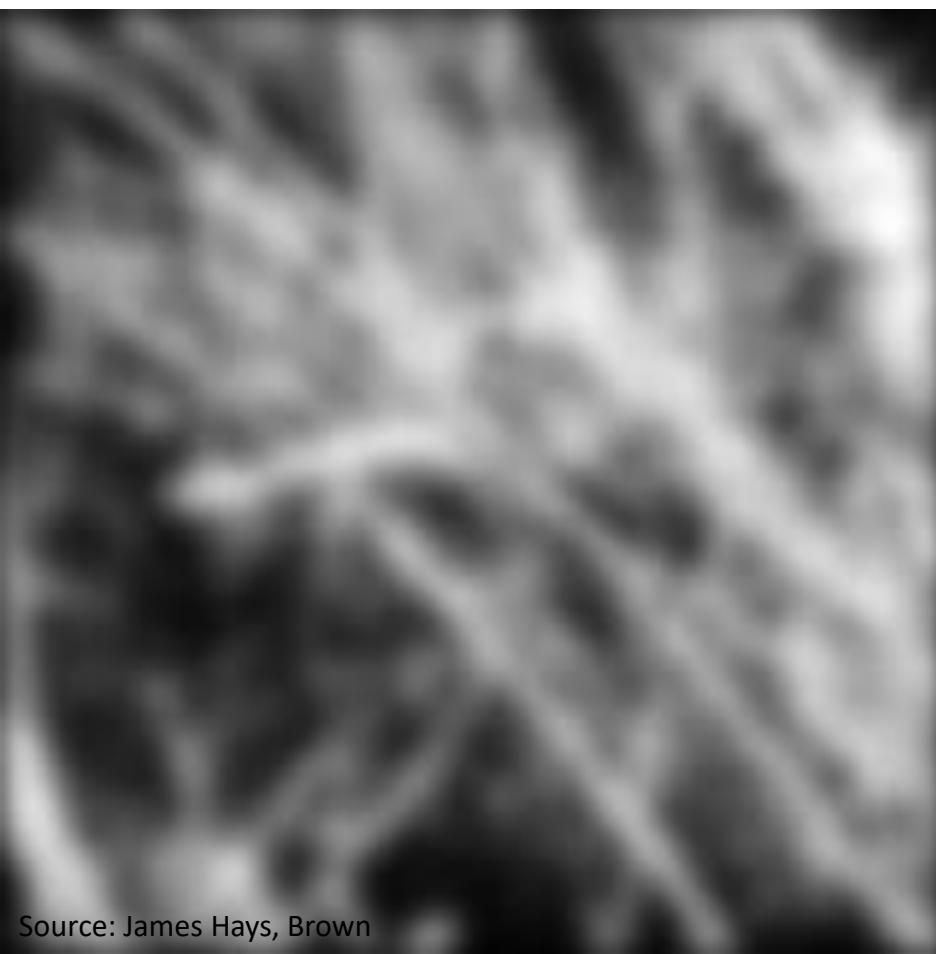
mirror

# Today's topics

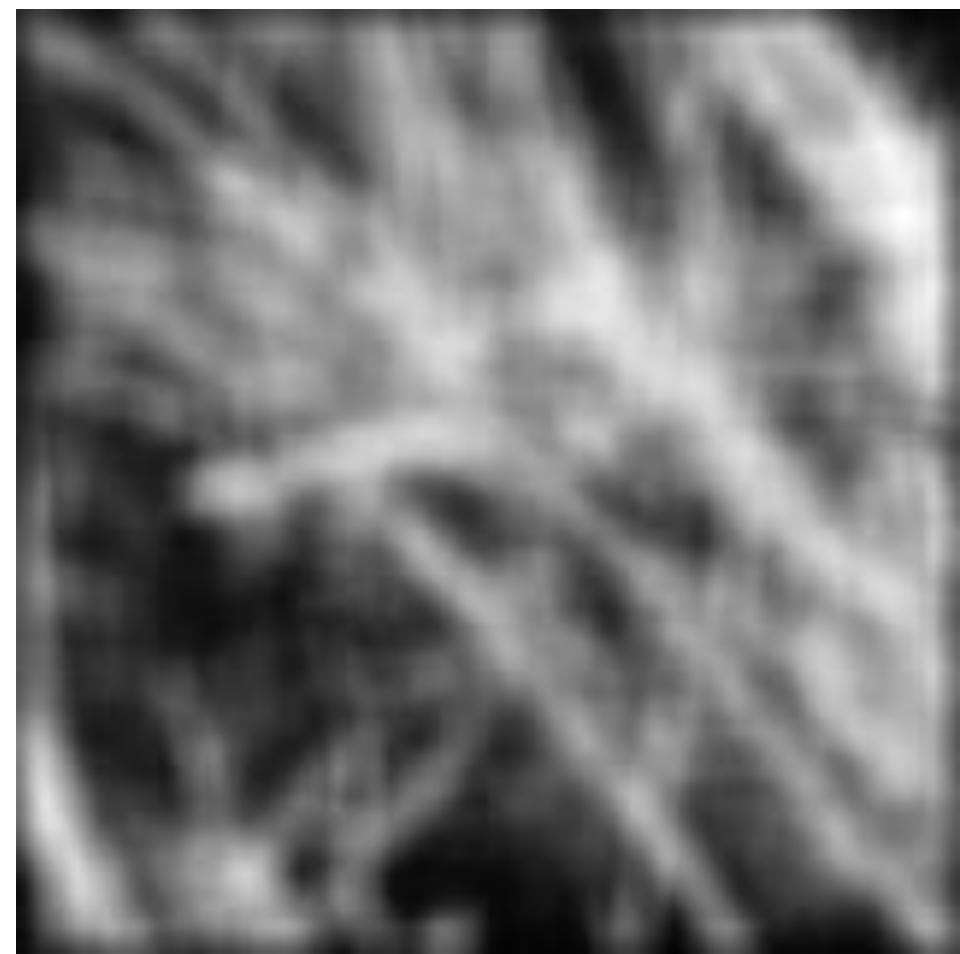
- Image Formation
- Image filters in spatial domain
  - Filter is a mathematical operation of a grid of numbers
  - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
  - Filtering is a way to modify the frequencies of images
  - Denoising, sampling, image compression
- Templates and Image Pyramids
  - Filtering is a way to match a template to the image
  - Detection, coarse-to-fine registration

# Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

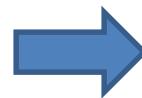
Gaussian



Box filter



# Why does a lower resolution image still make sense to us? What do we lose?



# Jean Baptiste Joseph Fourier (1768-

1830)

had crazy idea (1807):

*Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.*

*...the manner in which the author arrives at these equations is not exempt of difficulties and...his analysis to integrate them still leaves something to be desired on the score of generality and even rigour.*

- Don't believe it?
  - Neither did Lagrange, Laplace, Poisson and other big wigs
  - Not translated into English until 1878!
- But it's (mostly) true!
  - called Fourier Series
  - there are some subtle restrictions

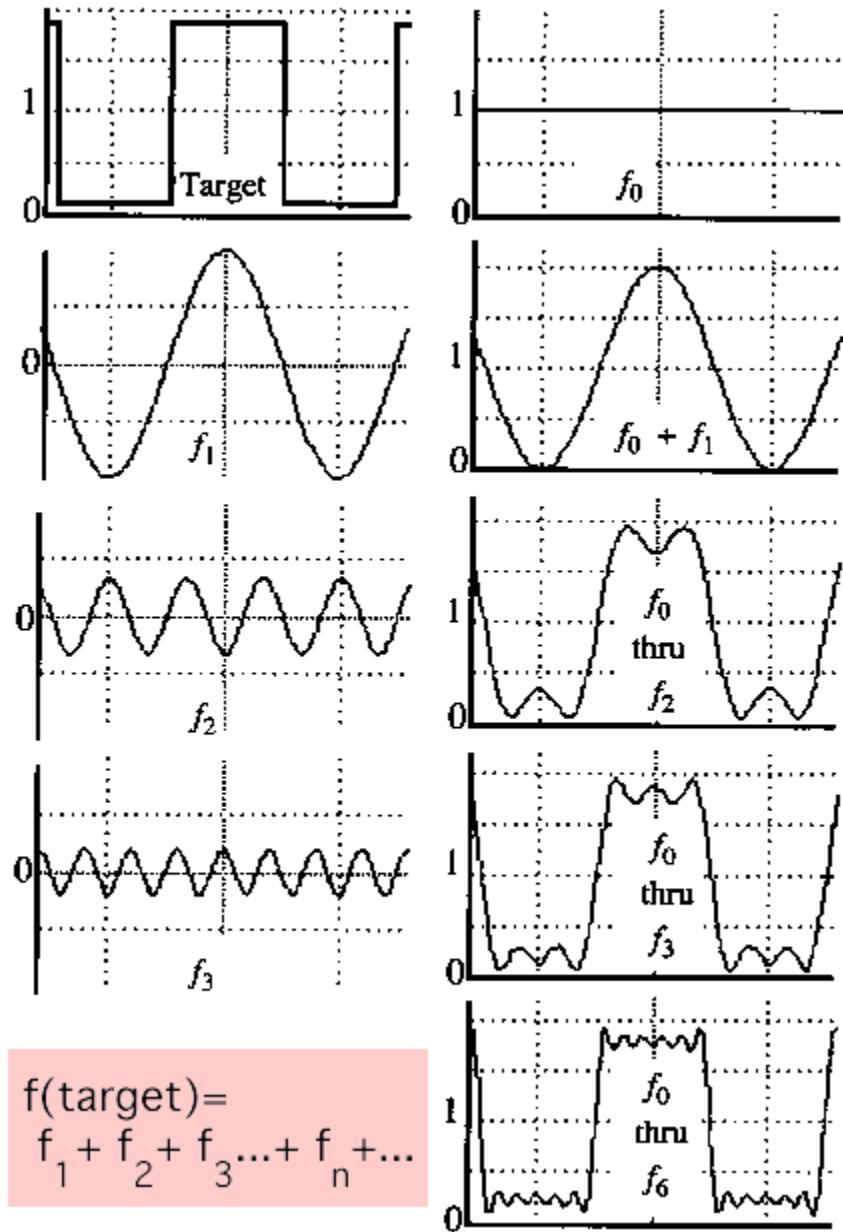


# A sum of sines

Our building block:

$$A \sin(\omega x + \phi)$$

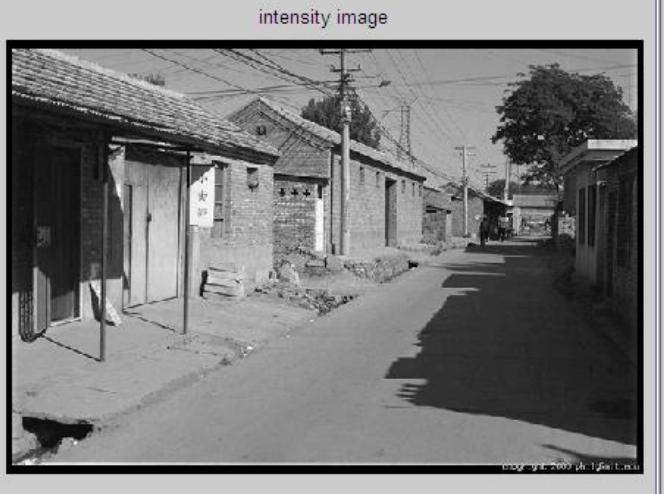
Add enough of them to get any signal  $f(x)$  you want!



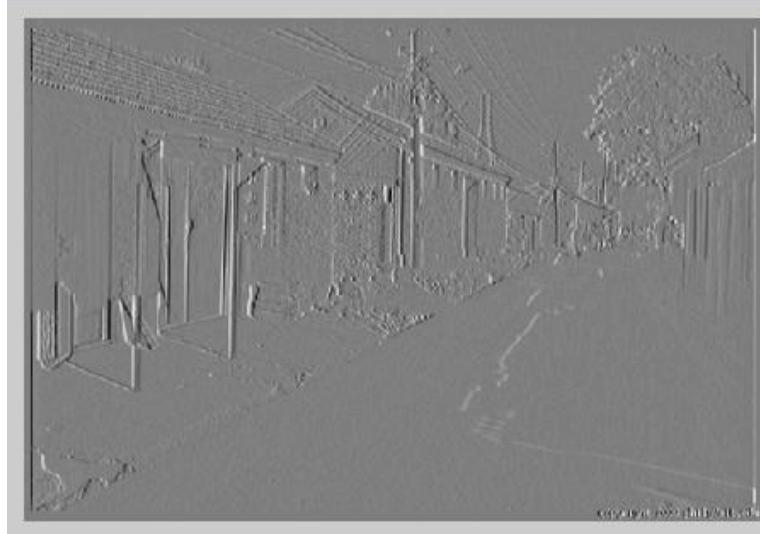
1	0	-1
2	0	-2
1	0	-1

# Filtering in spatial dom

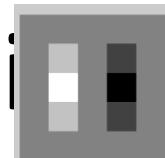
intensity image



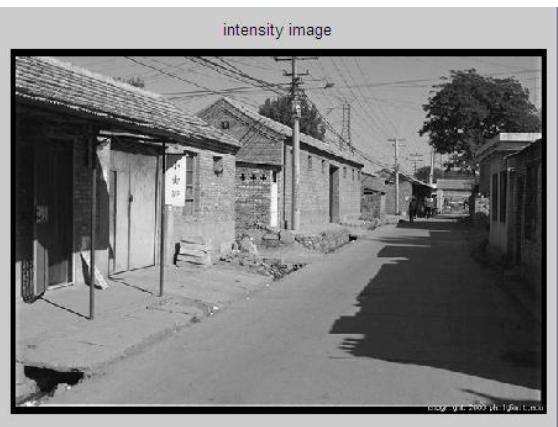
$$* \begin{matrix} \text{dark gray} & \text{white} \\ \text{black} & \text{white} \end{matrix} =$$



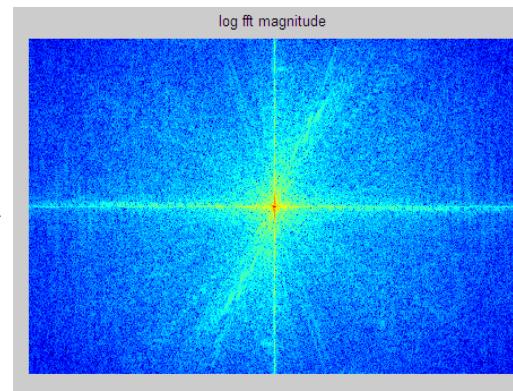
# Filtering in frequency domain



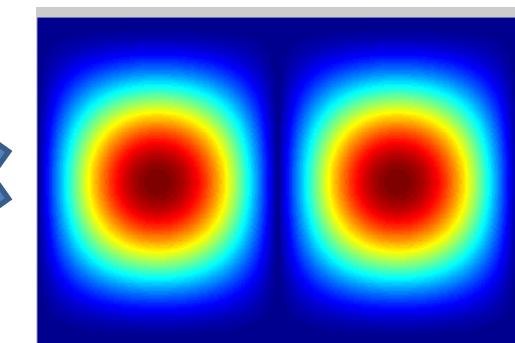
FFT  
↓



FFT  
→

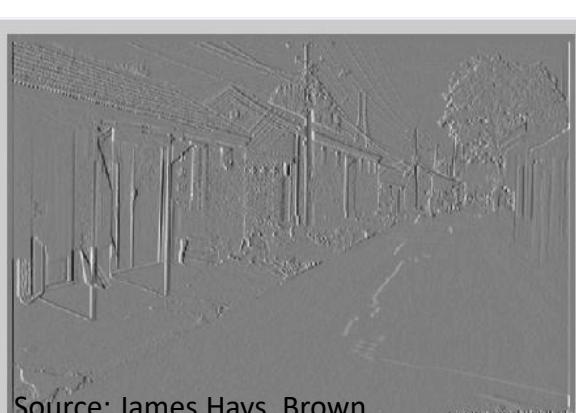
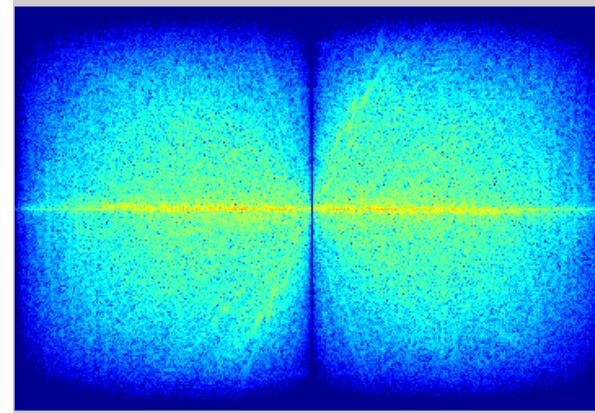


×



||

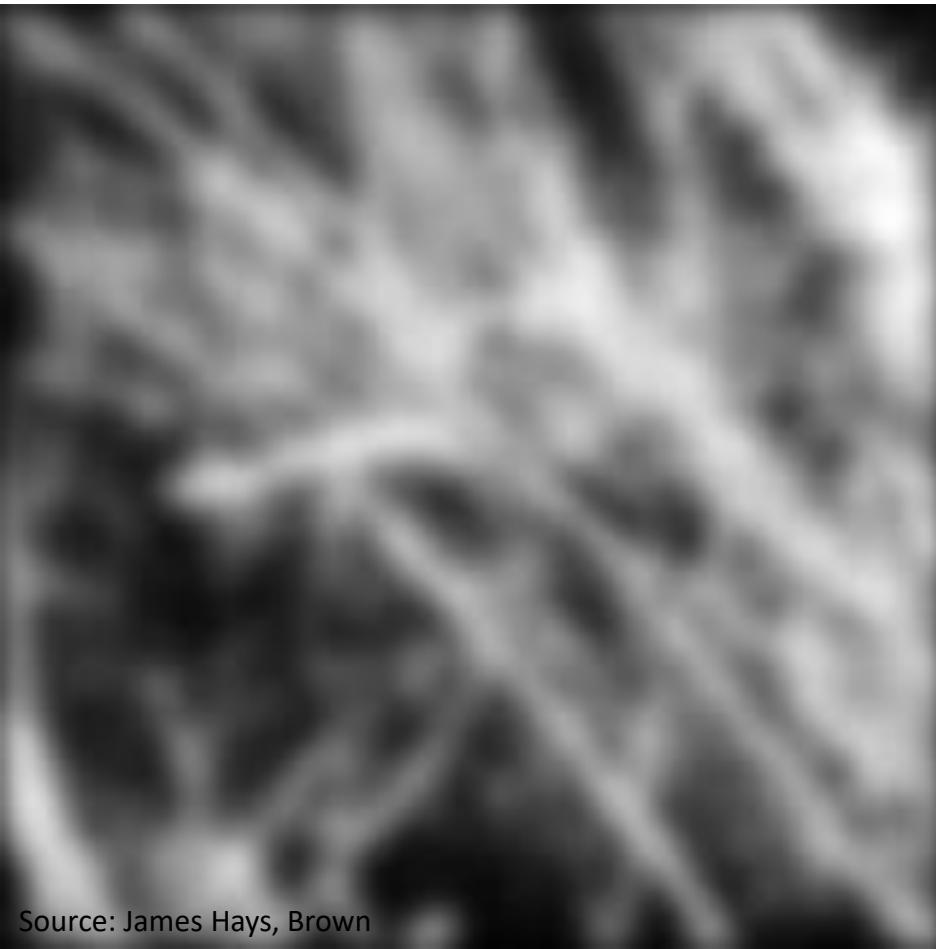
Inverse FFT  
←



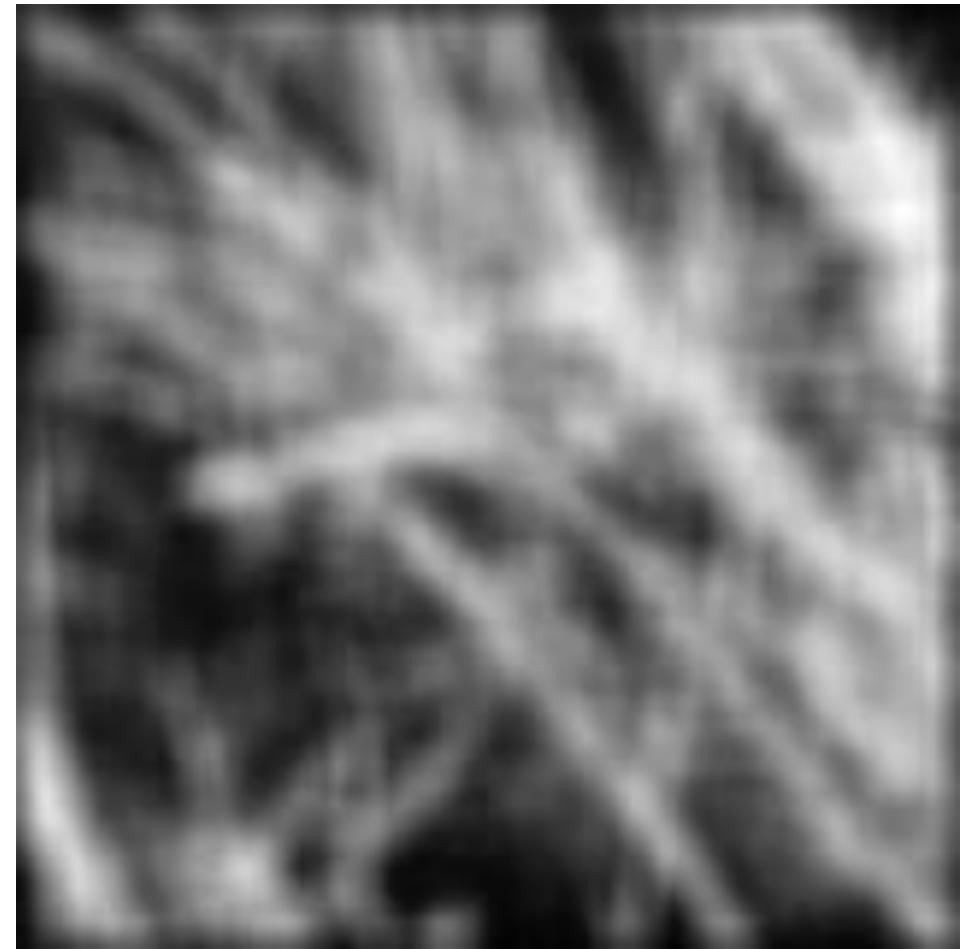
# Filtering

Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

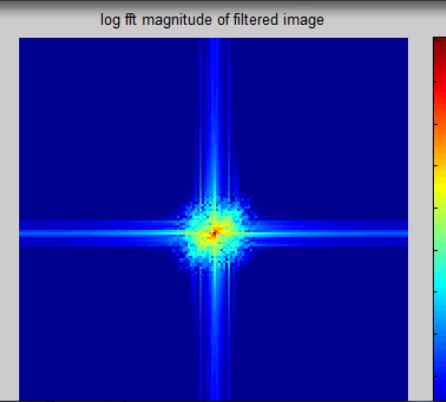
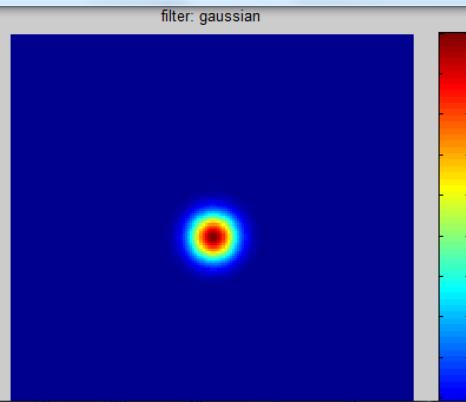
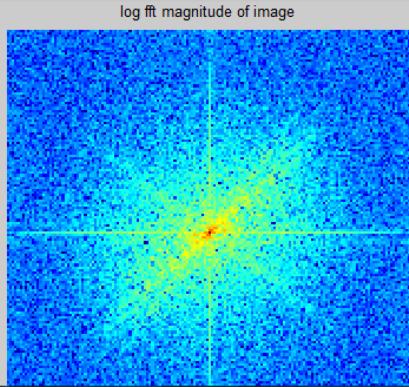
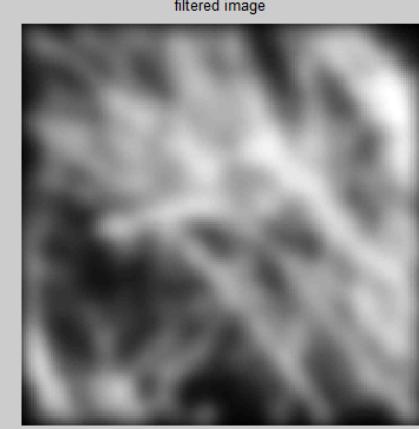
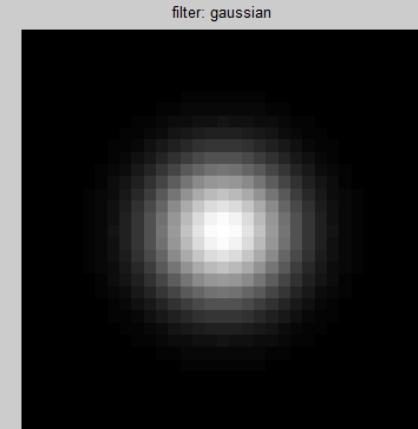
Gaussian



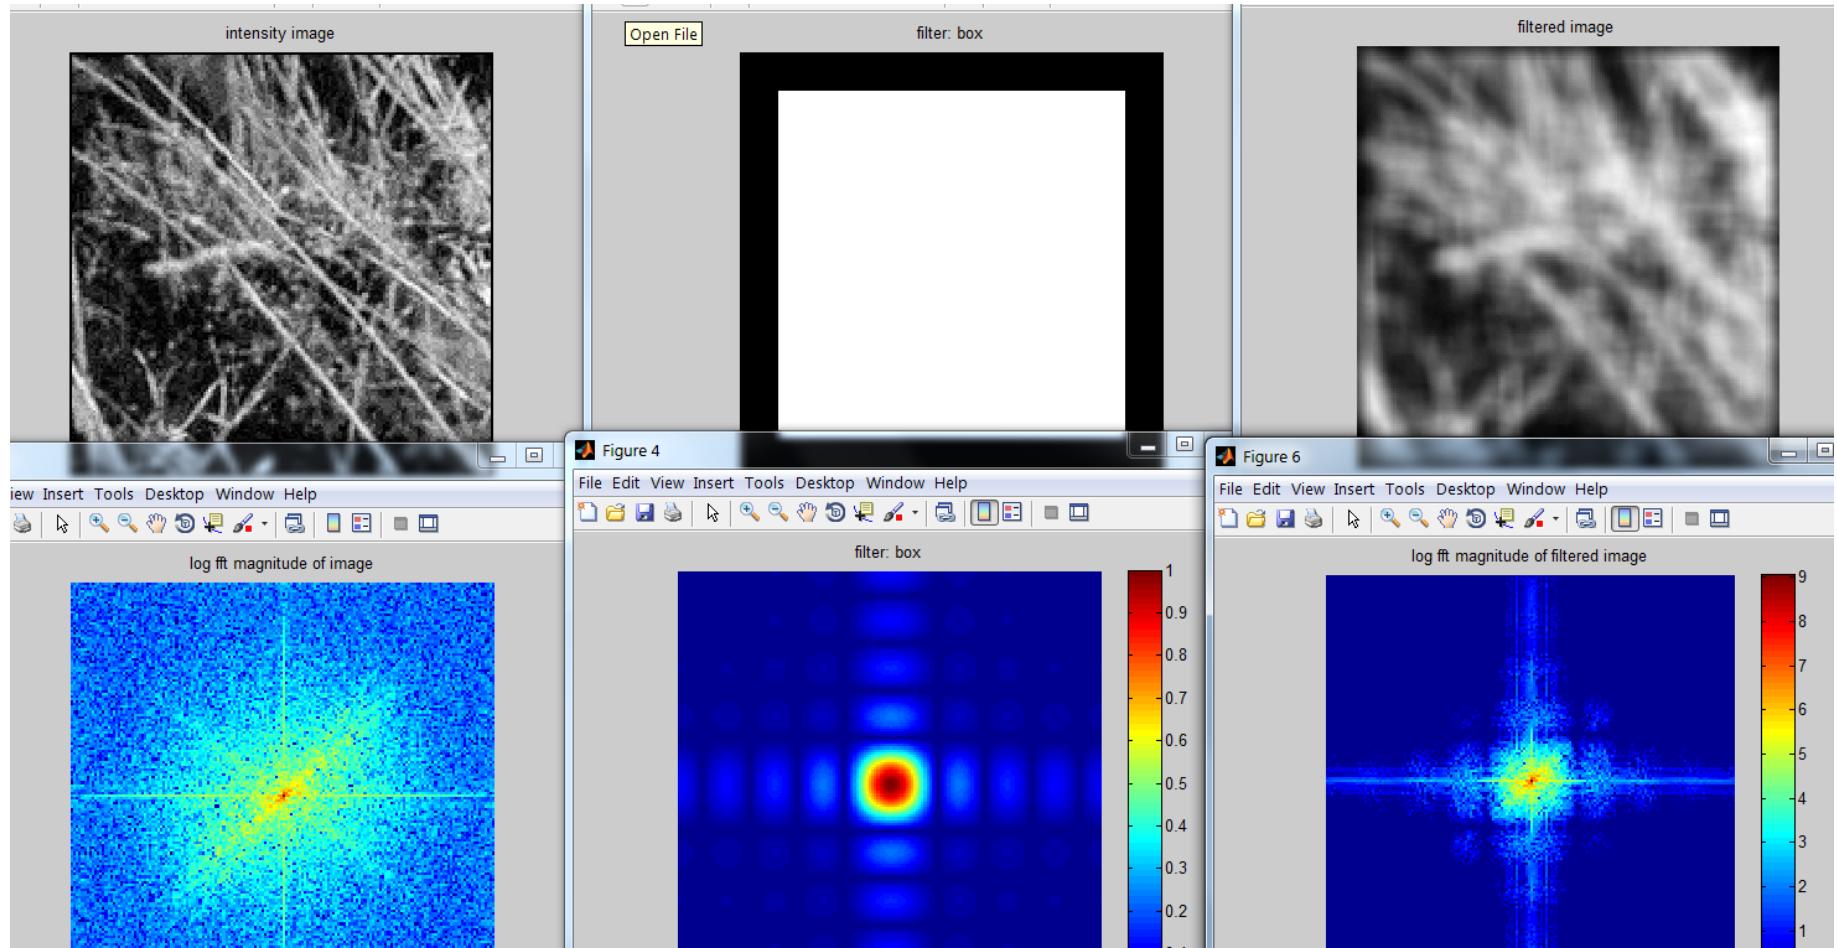
Box filter



# Gaussian

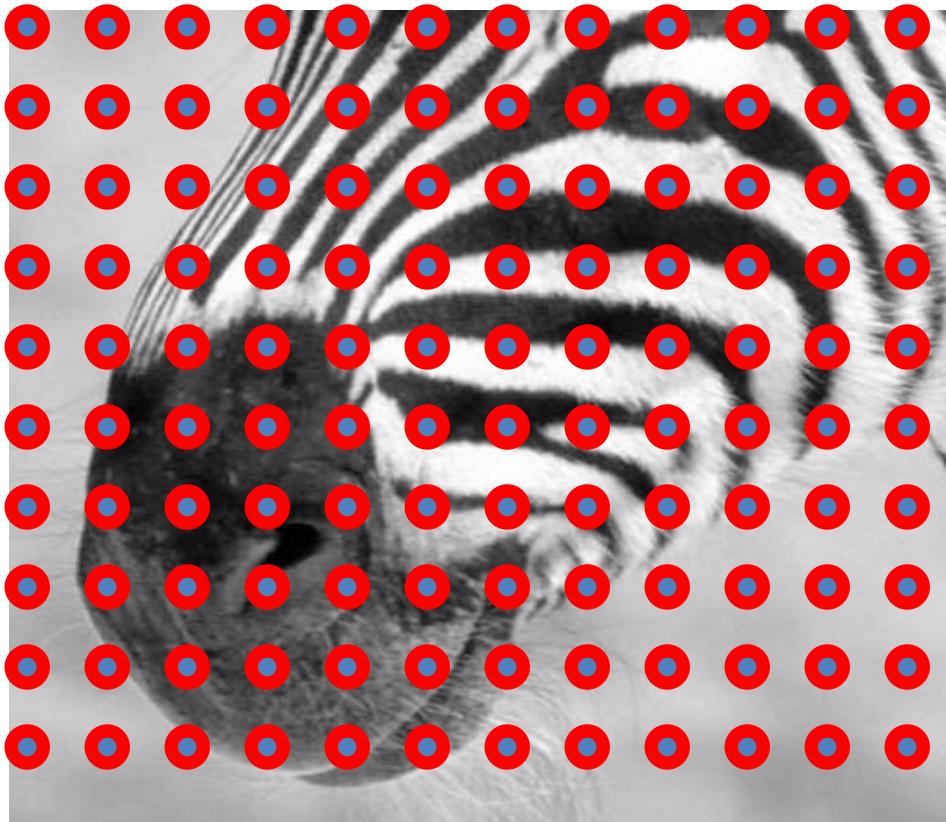


# Box Filter



Source: James Hays, Brown

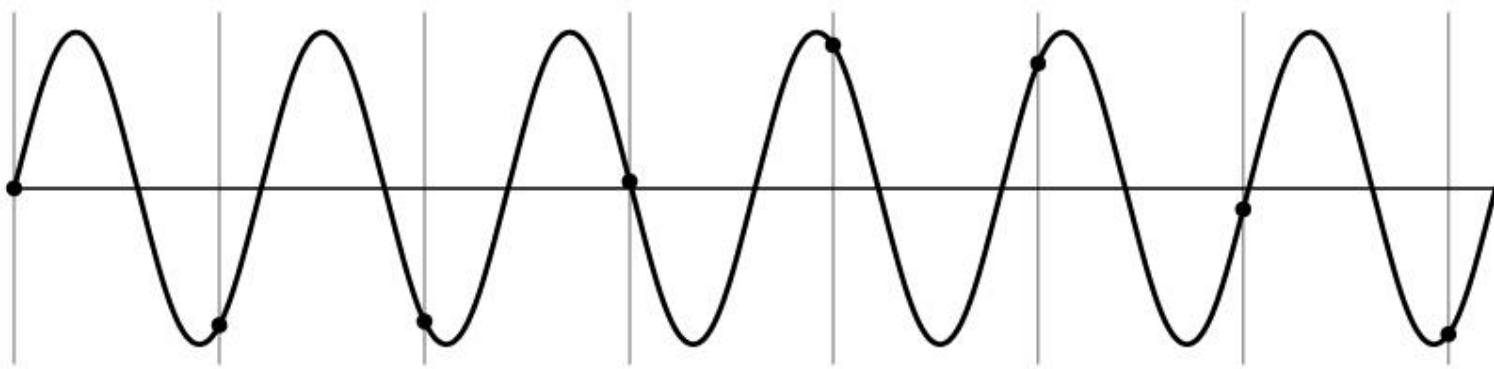
# Subsampling by a factor of 2



Throw away every other row and column  
to create a  $1/2$  size image

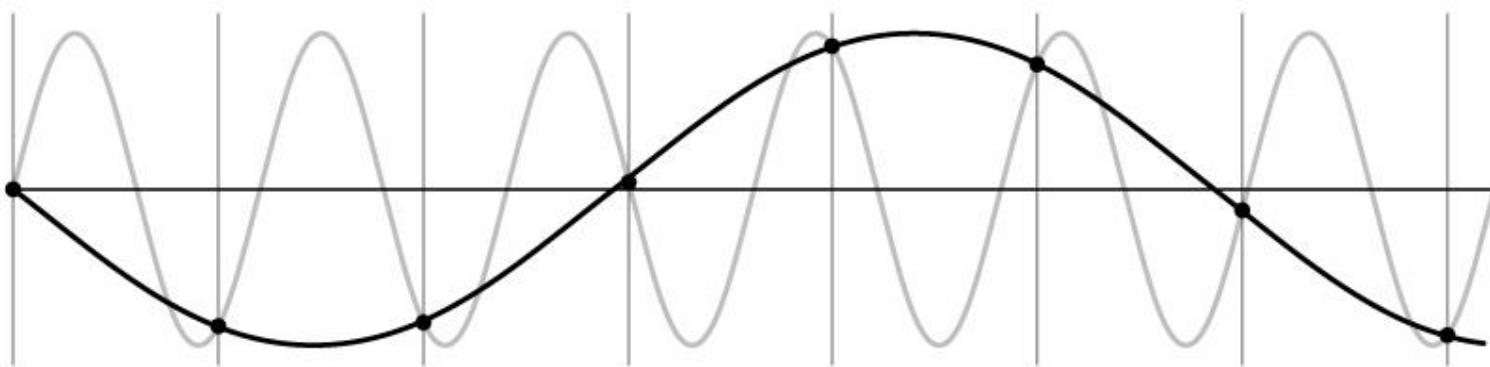
# Aliasing problem

- 1D example (sinewave):



# Aliasing problem

- 1D example (sinewave):



# Subsampling without pre-filtering



1/2

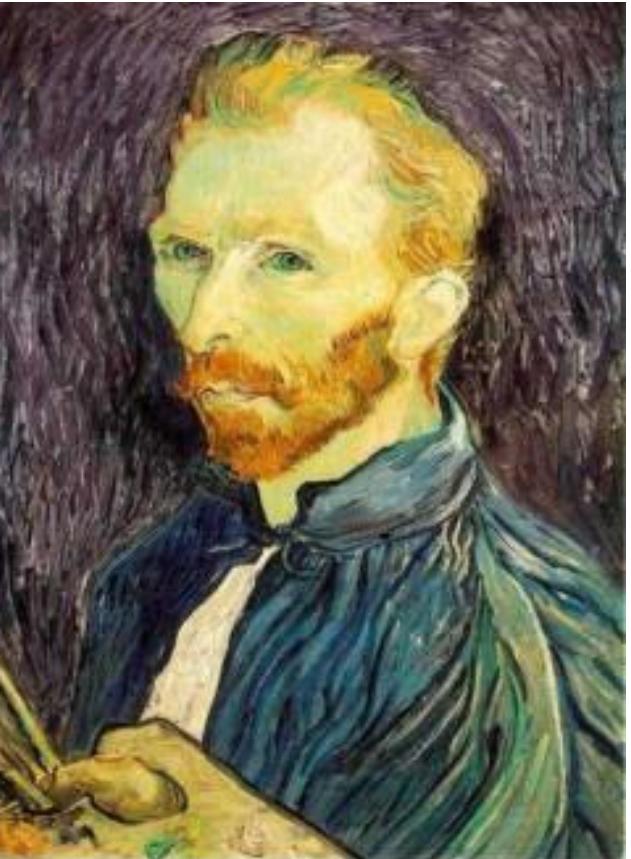


1/4 (2x zoom)



1/8 (4x zoom)

# Subsampling with Gaussian pre-filtering



Gaussian 1/2



G 1/4

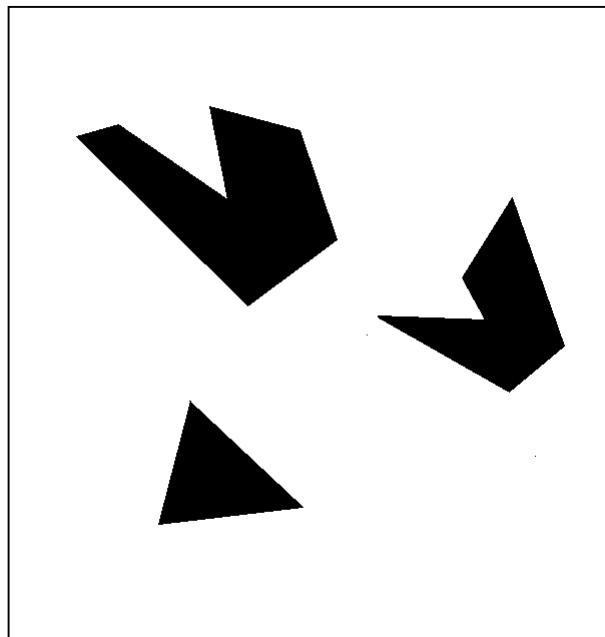


G 1/8

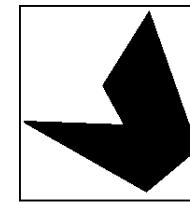
# Today's topics

- Image Formation
- Image filters in spatial domain
  - Filter is a mathematical operation of a grid of numbers
  - Smoothing, sharpening, measuring texture
- Image filters in the frequency domain
  - Filtering is a way to modify the frequencies of images
  - Denoising, sampling, image compression
- Templates and Image Pyramids
  - Filtering is a way to match a template to the image
  - Detection, coarse-to-fine registration

# Template matching



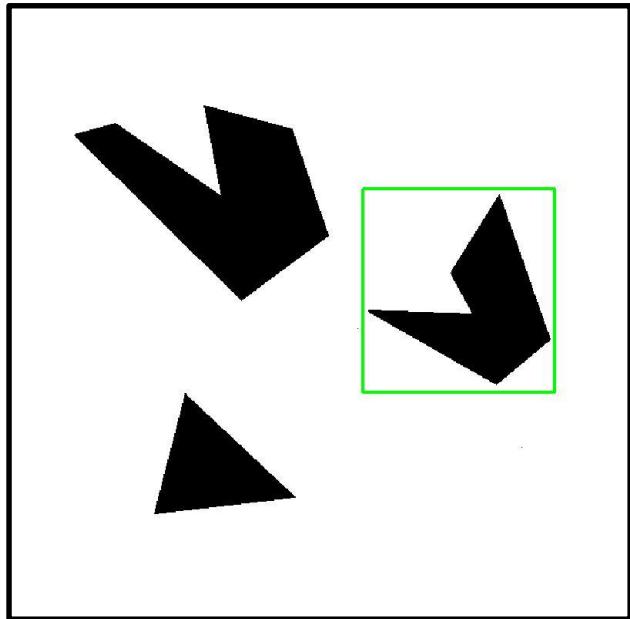
Scene



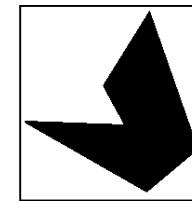
Template (mask)

A toy example

# Template matching

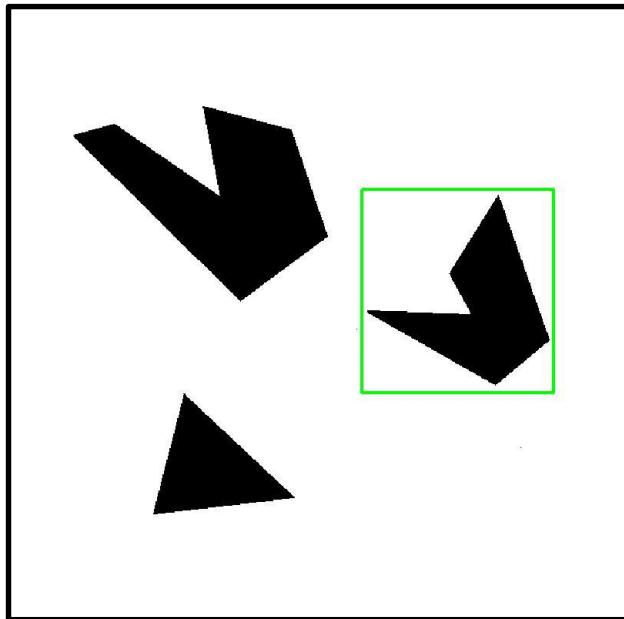


Detected template

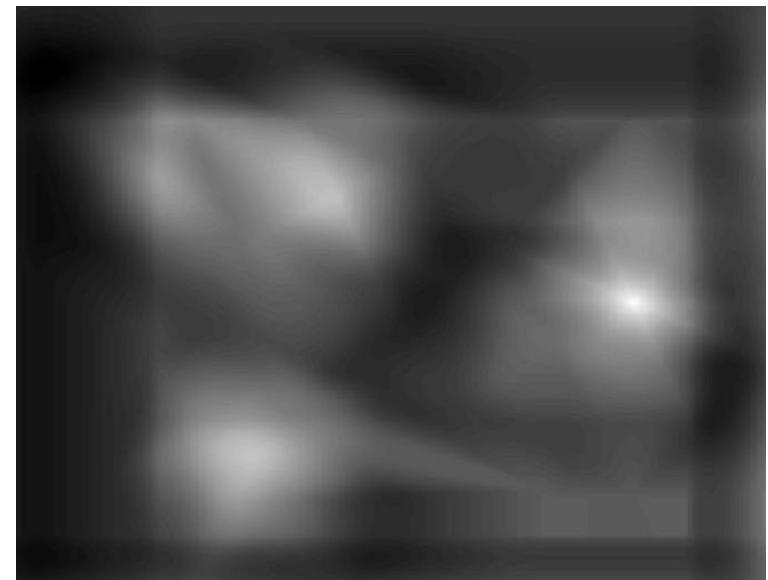


Template

# Template matching



Detected template



Correlation map

# Where's Waldo?

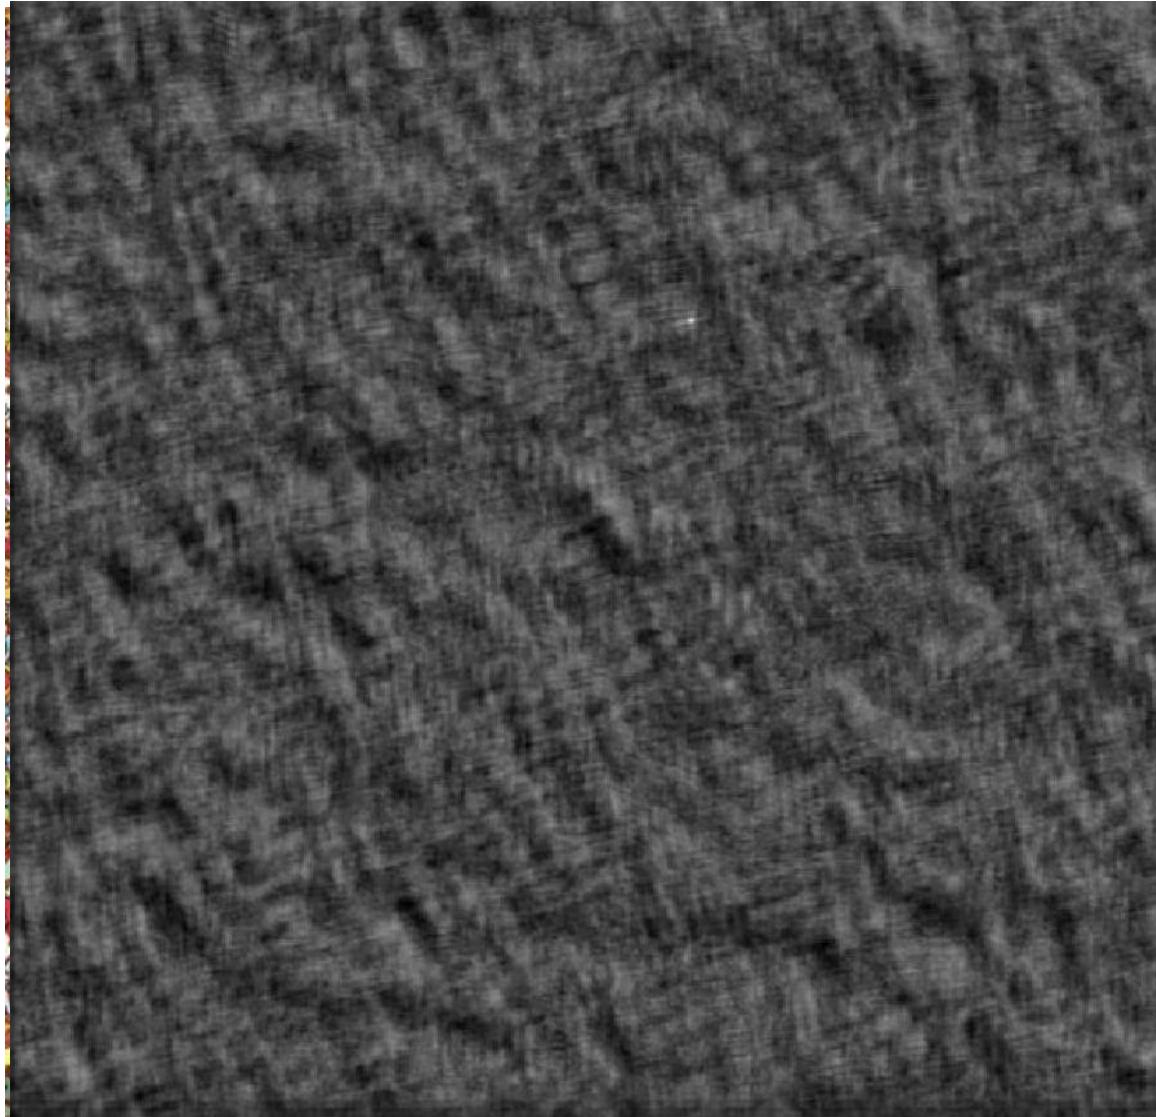


Scene



Template

# Where's Waldo?

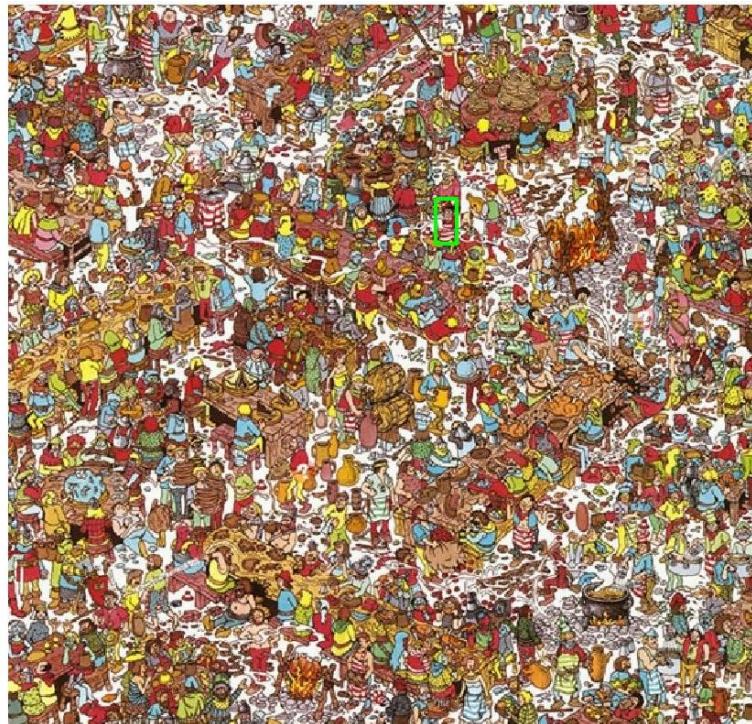


Scene

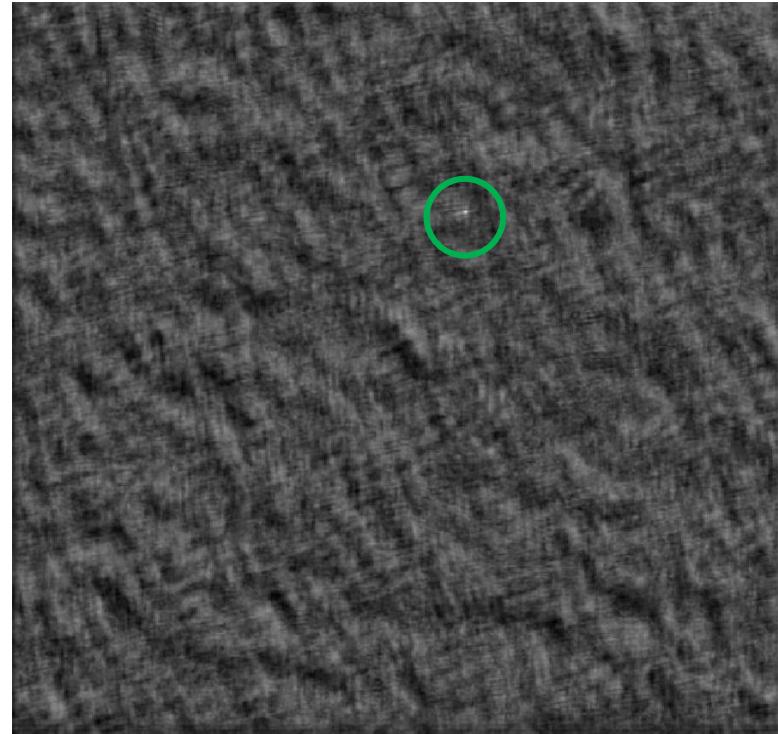


Template

# Where's Waldo?



Detected template



Correlation map

# Template matching



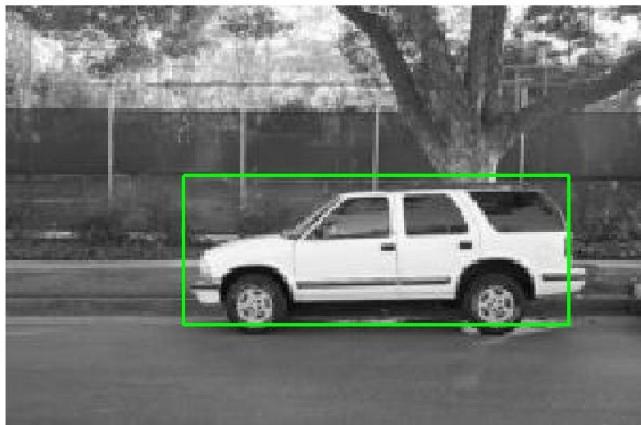
Scene



Template

What if the template is not identical to some subimage in the scene?

# Template matching



Detected template



Template

Match can be meaningful, if scale, orientation, and general appearance is right.

# Application

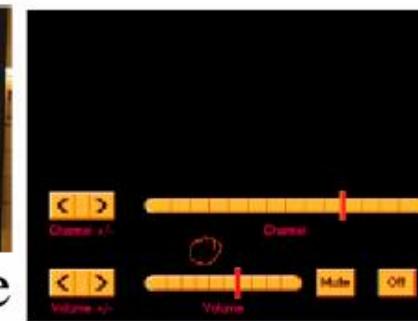
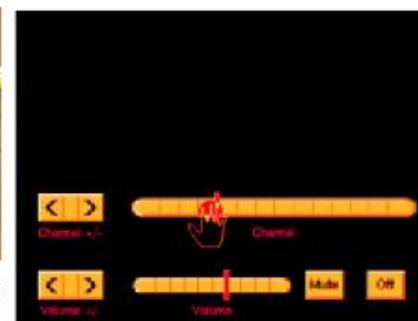
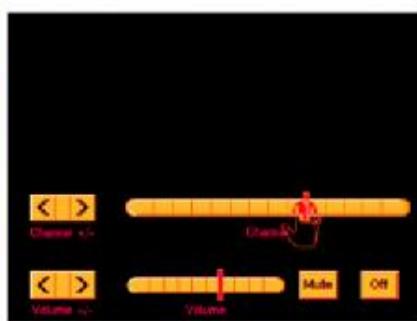
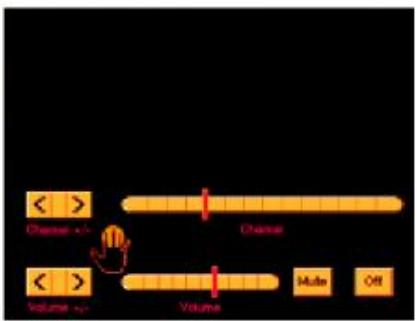
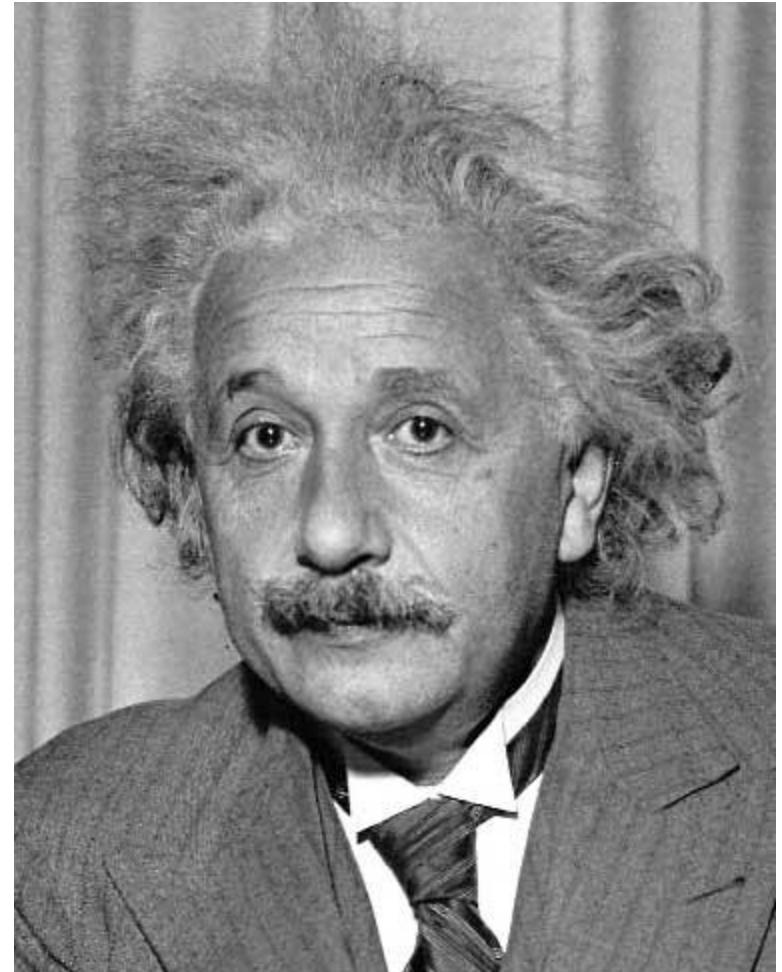


Figure from "Computer Vision for Interactive Computer Graphics," W.Freeman et al, IEEE Computer Graphics and Applications, 1998  
copyright 1998, IEEE

# Template matching

- Goal: find  in image
- Main challenge: What is a good similarity or distance measure between two patches?
  - Correlation
  - Zero-mean correlation
  - Sum Square Difference
  - Normalized Cross Correlation



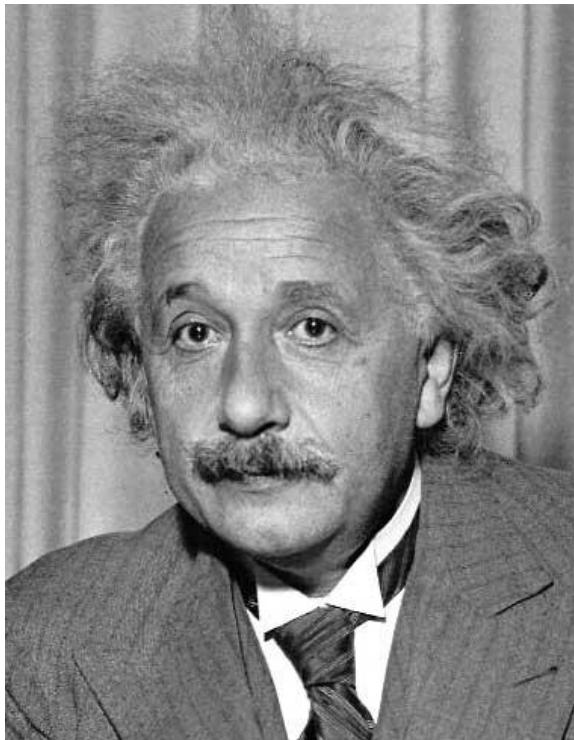
# Matching with filters

- Goal: find  in image
- Method 0: filter the image with eye patch

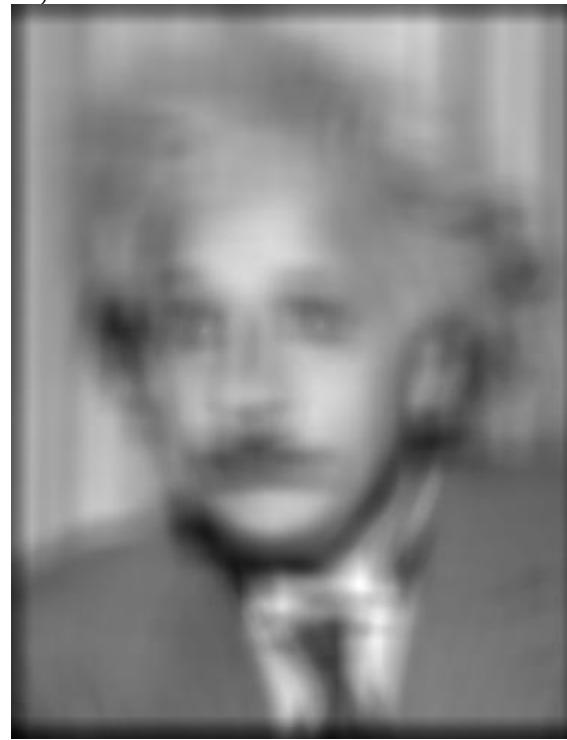
$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$



f = image  
g = filter



Input



Filtered Image

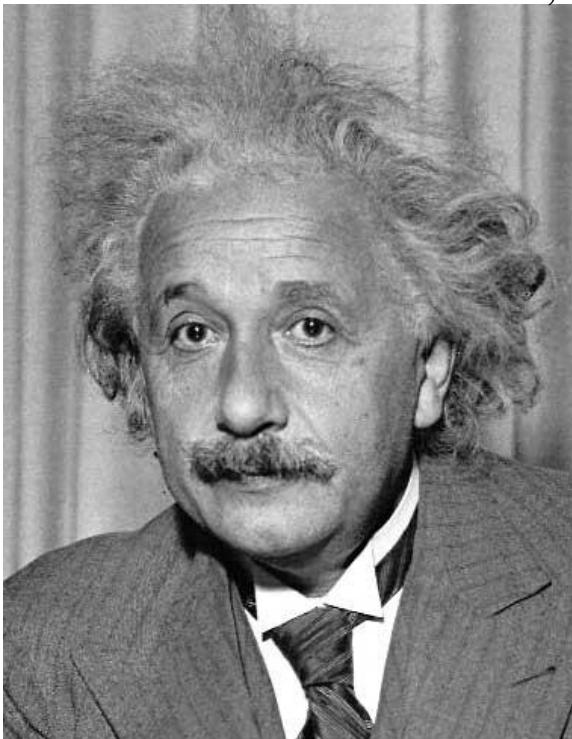
What went wrong?

response is stronger  
for higher intensity

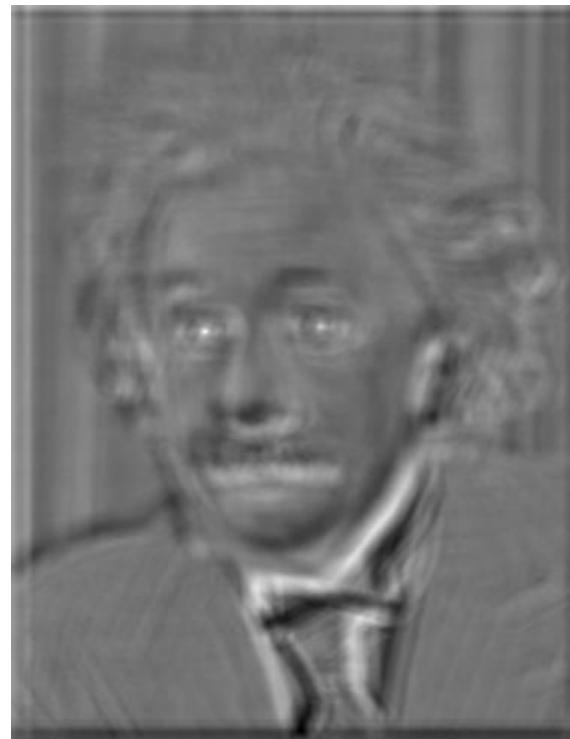
# Matching with filters

- Goal: find  in image
- Method 1: filter the image with zero-mean eye

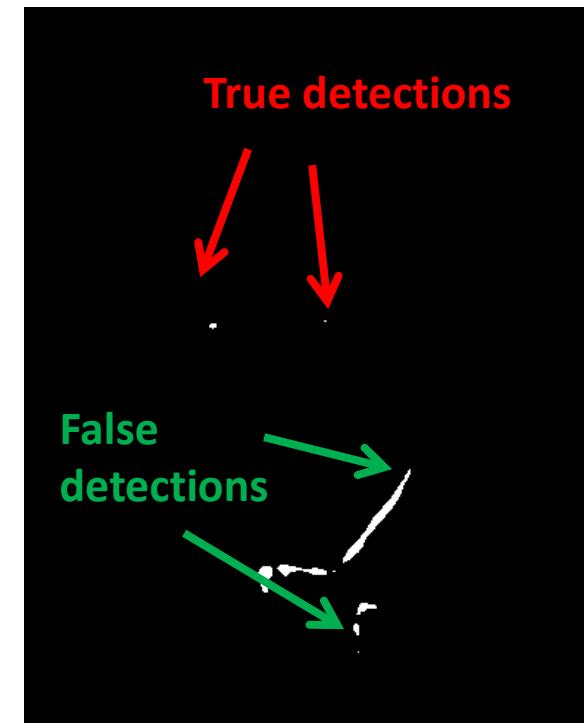
$$h[m,n] = \sum_{k,l} (f[k,l] - \bar{f}) \underbrace{(g[m+k, n+l])}_{\text{mean of } f}$$



Input



Filtered Image (scaled)



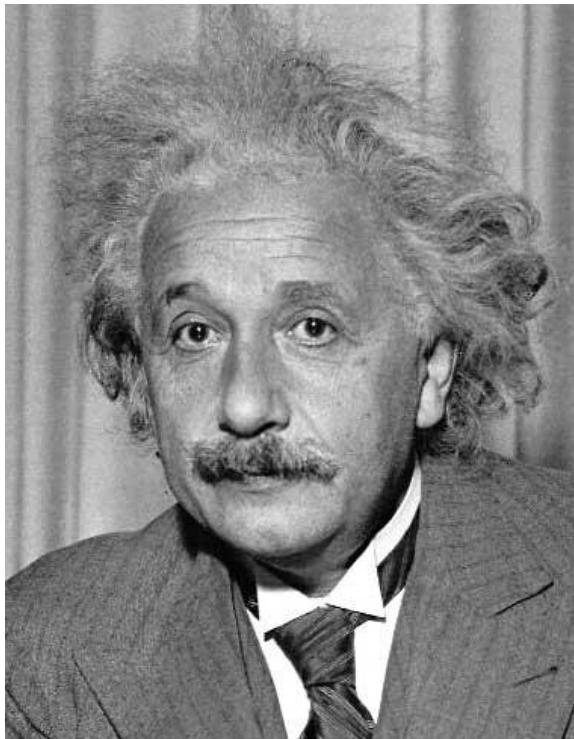
Thresholded Image

# Matching with filters

- Goal: find  in image

- Method 2: SSD

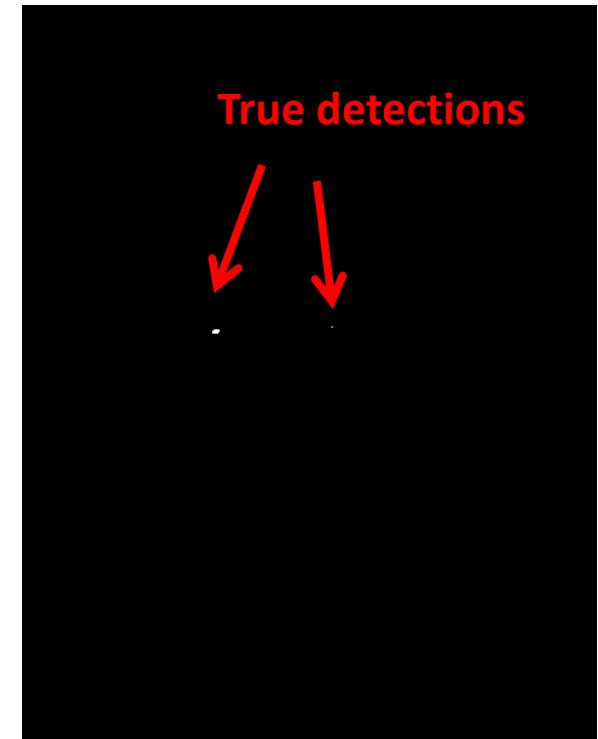
$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1 -  $\sqrt{\text{SSD}}$



Thresholded Image

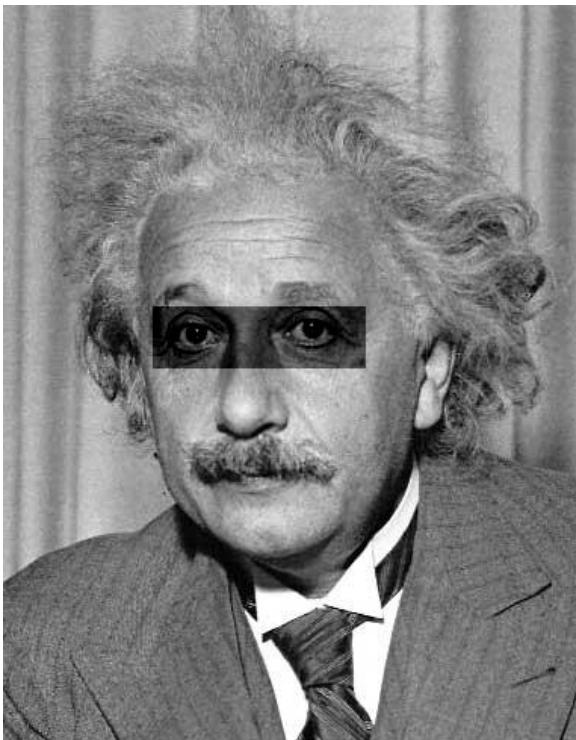
Slide: Hoiem

# Matching with filters

- Goal: find  in image

- Method 2: SSD

$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1- sqrt(SSD)

What's the potential downside  
of SSD?

SSD is sensitive to  
average intensity

# Matching with filters

- Goal: find  in image
- Method 3: Normalized cross-correlation

$$h[m,n] = \frac{\sum_{k,l} (g[k,l] - \bar{g})(f[m-k, n-l] - \bar{f}_{m,n})}{\left( \sum_{k,l} (g[k,l] - \bar{g})^2 \sum_{k,l} (f[m-k, n-l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

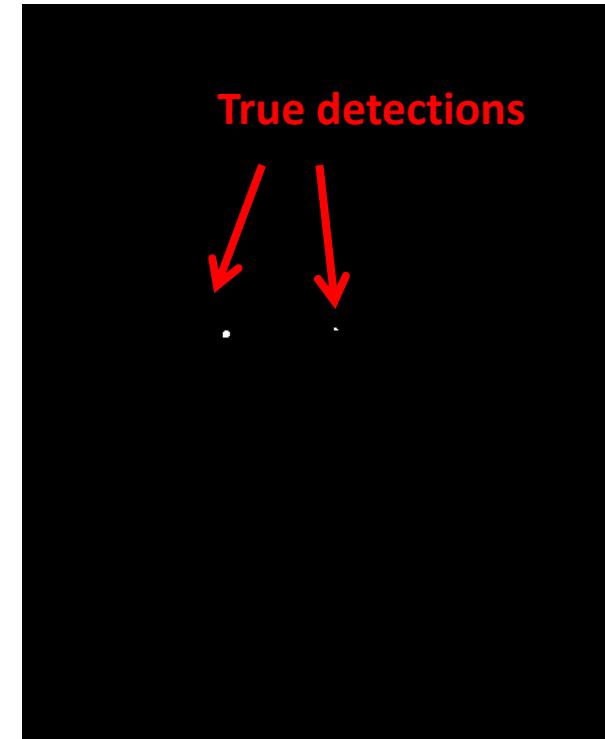
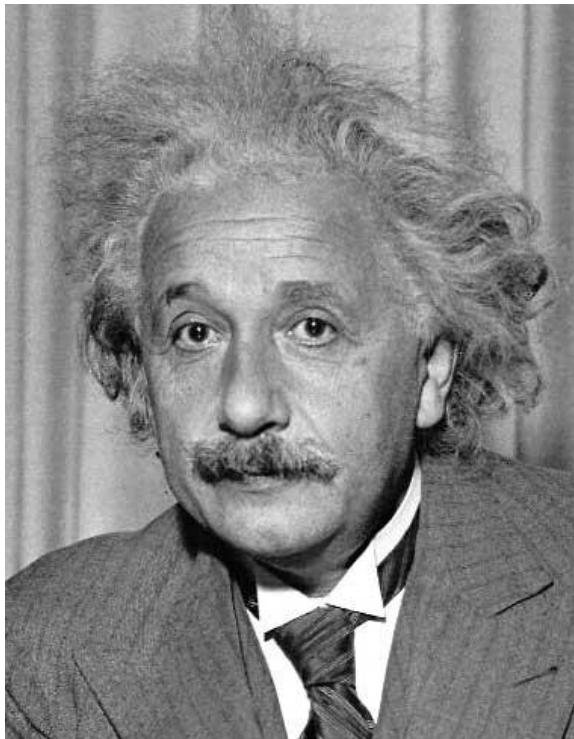
mean template                                  mean image patch

$\downarrow$      $\downarrow$

Matlab: `normxcorr2(template, im)`

# Matching with filters

- Goal: find  in image
- Method 3: Normalized cross-correlation



Input

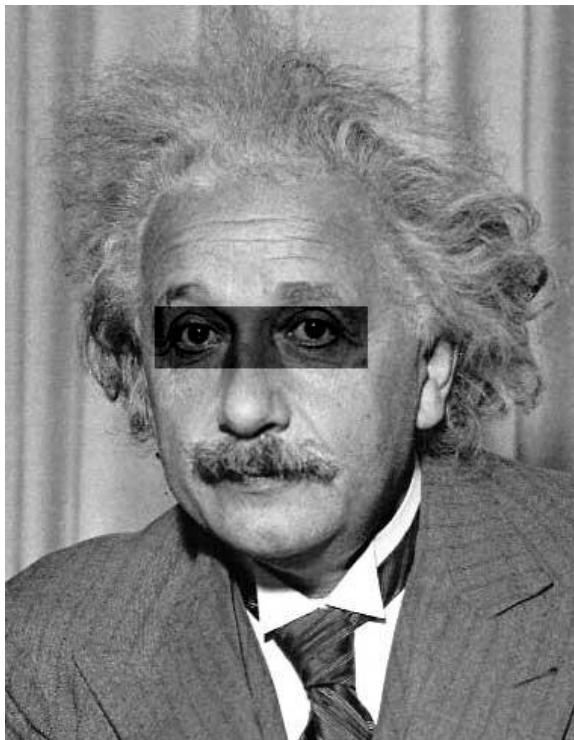
Normalized X-Correlation

Thresholded Image

Slide: Hoiem

# Matching with filters

- Goal: find  in image
- Method 3: Normalized cross-correlation



Input

Normalized X-Correlation

Thresholded Image

Slide: Hoiem

# Q: What is the best method to use?

A: Depends

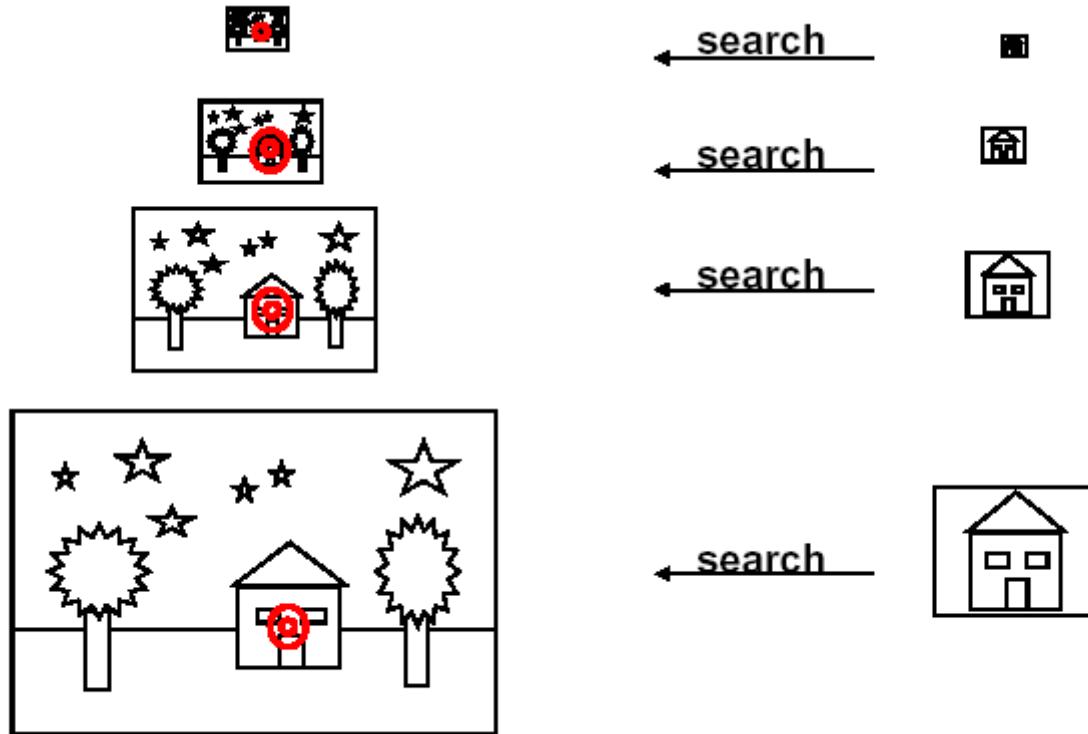
- SSD: faster, sensitive to overall intensity
- Normalized cross-correlation: slower, invariant to local average intensity and contrast

# Q: What if we want to find larger or smaller eyes?

Motivation for studying scale.

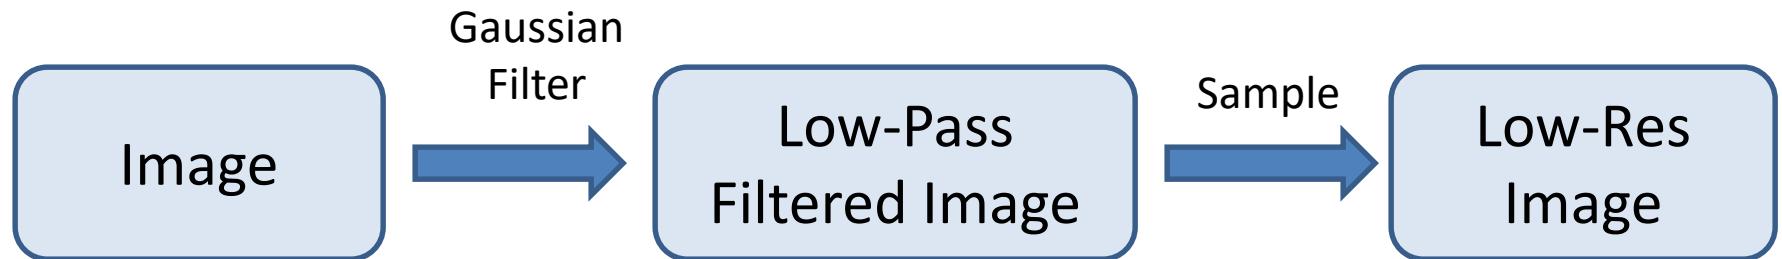


# A: Image Pyramid



Irani & Basri

# Review of Sampling

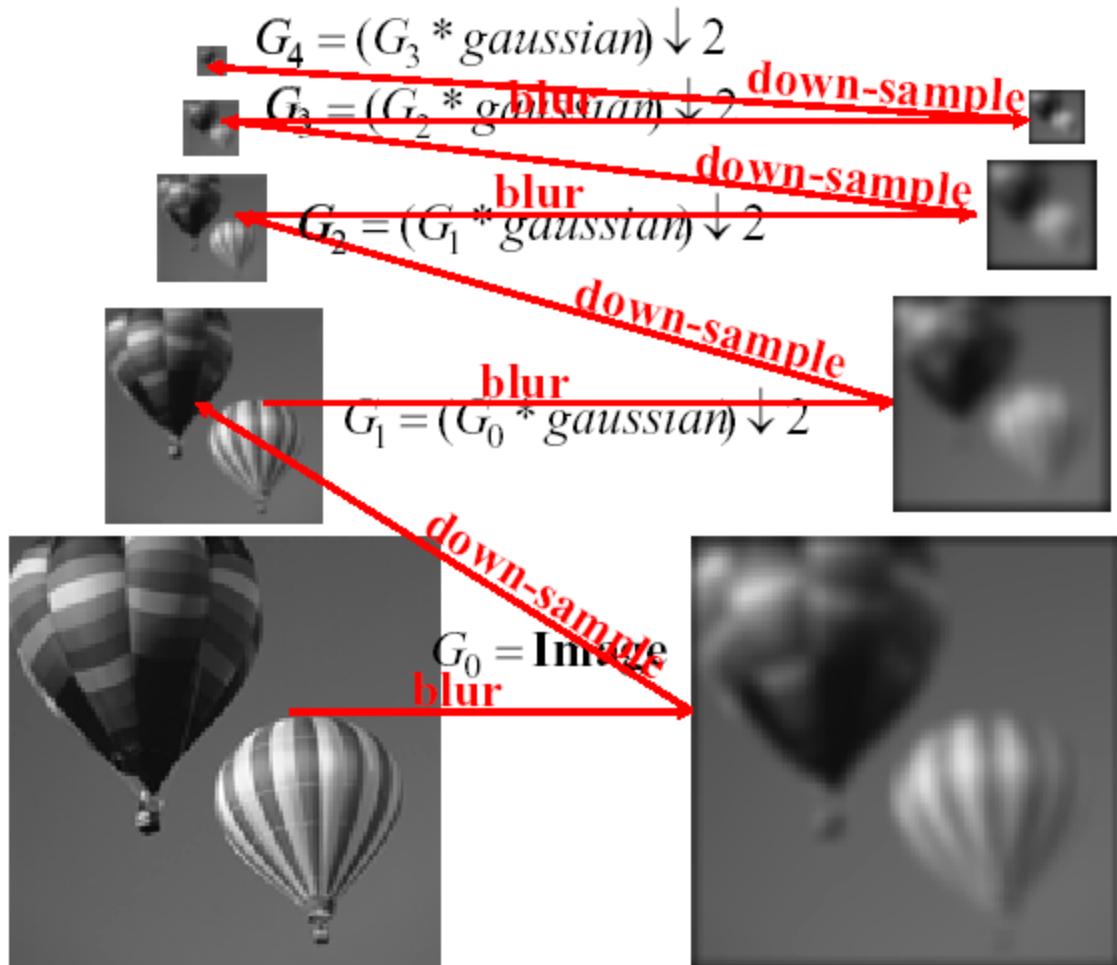


# Gaussian Pyramid

Low resolution

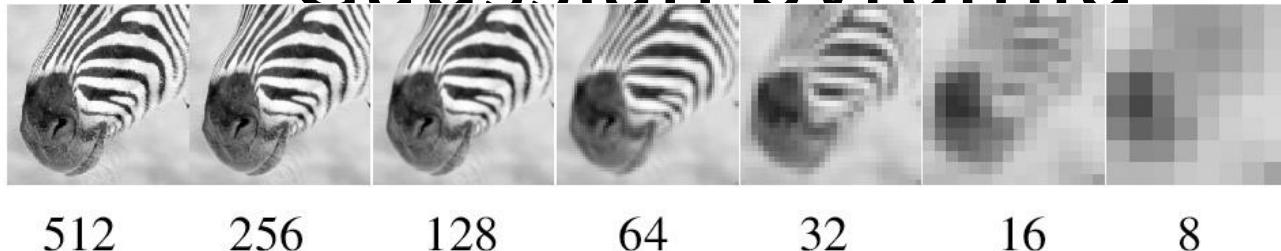


High resolution



Irani & Basri

# Gaussian pyramid



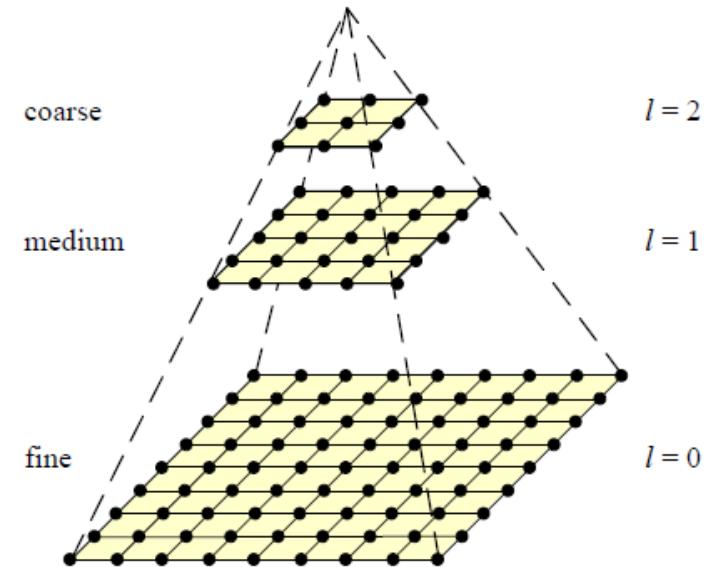
# Template Matching with Image Pyramids

Input: Image, Template

1. Match template at current scale
2. Downsample image
3. Repeat 1-2 until image is very small
4. Take responses above some threshold, perhaps with non-maxima suppression

# Coarse-to-fine Image Registration

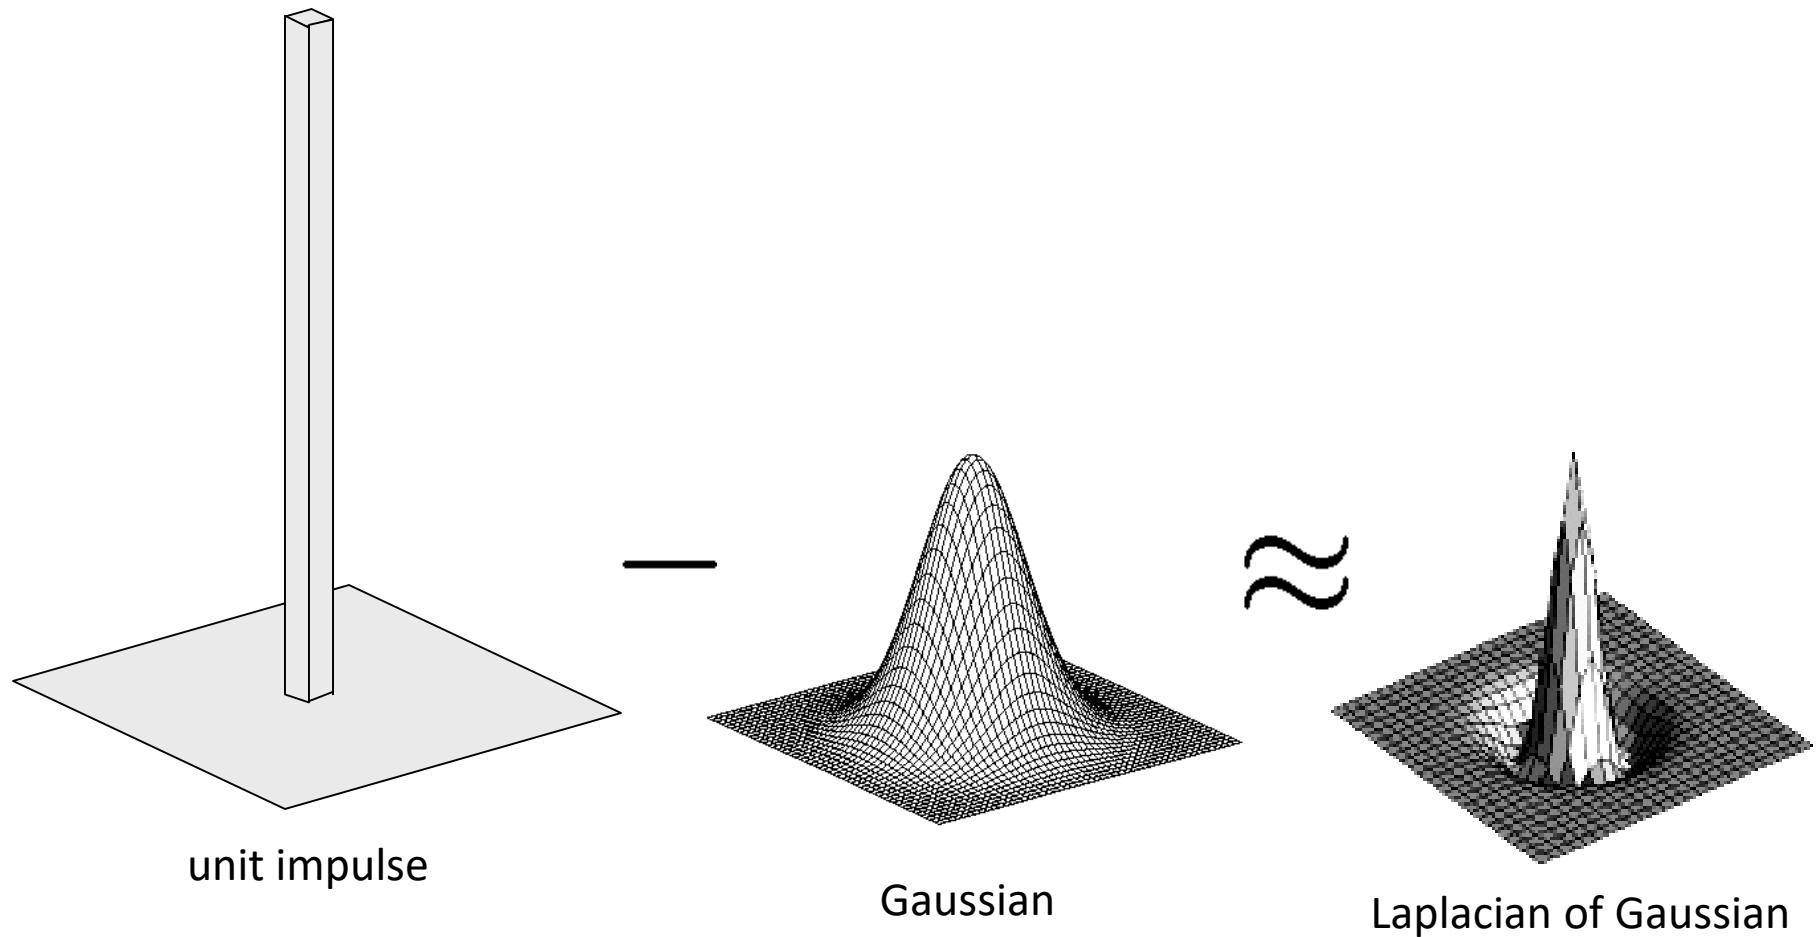
1. Compute Gaussian pyramid
2. Align with coarse pyramid
3. Successively align with finer pyramids
  - Search smaller range



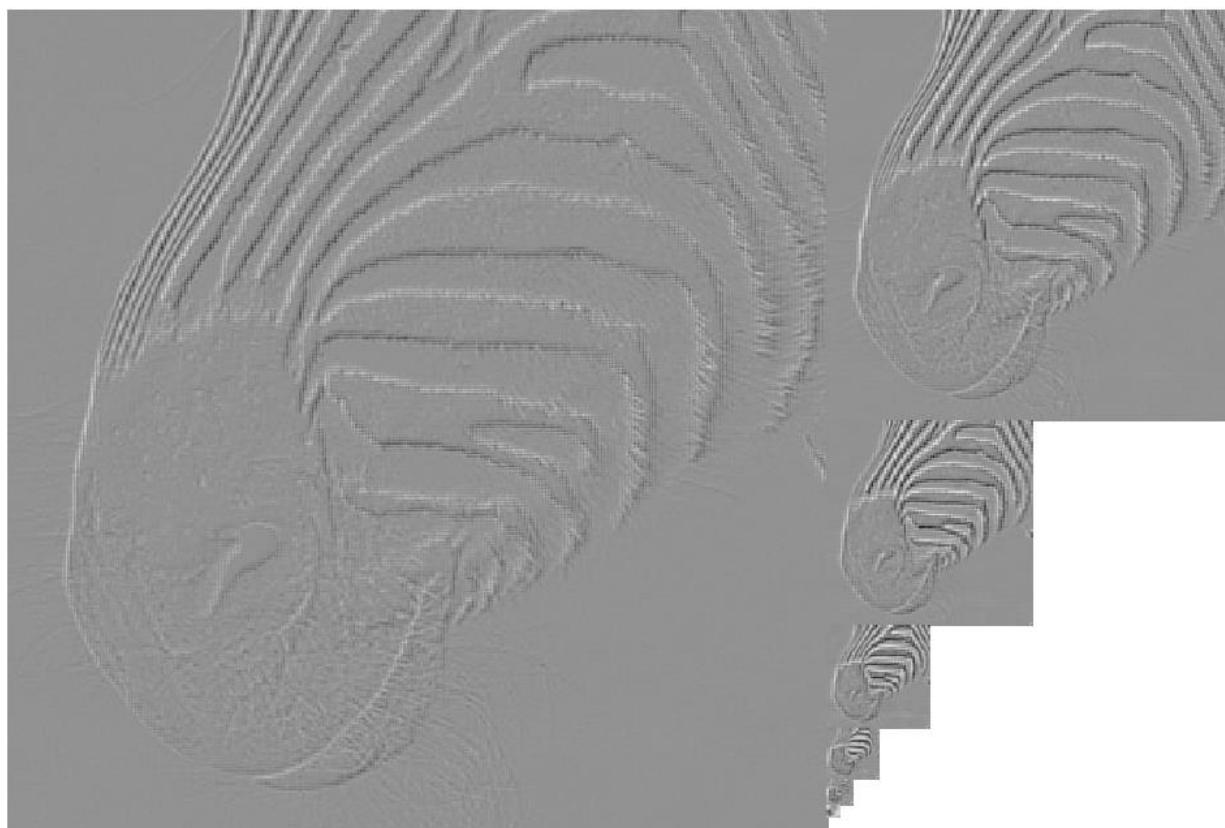
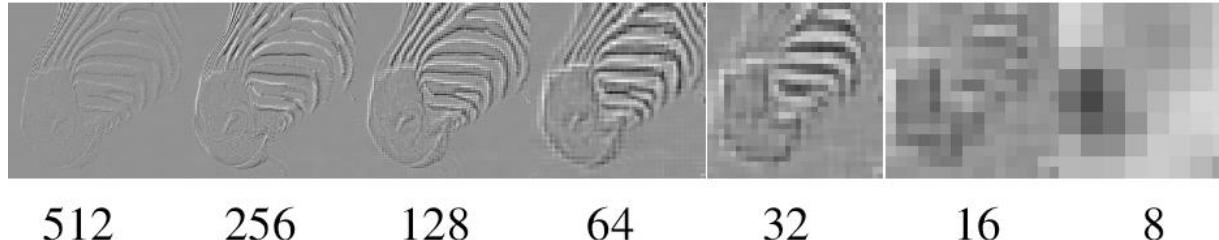
Why is this faster?

Are we guaranteed to get the same result?

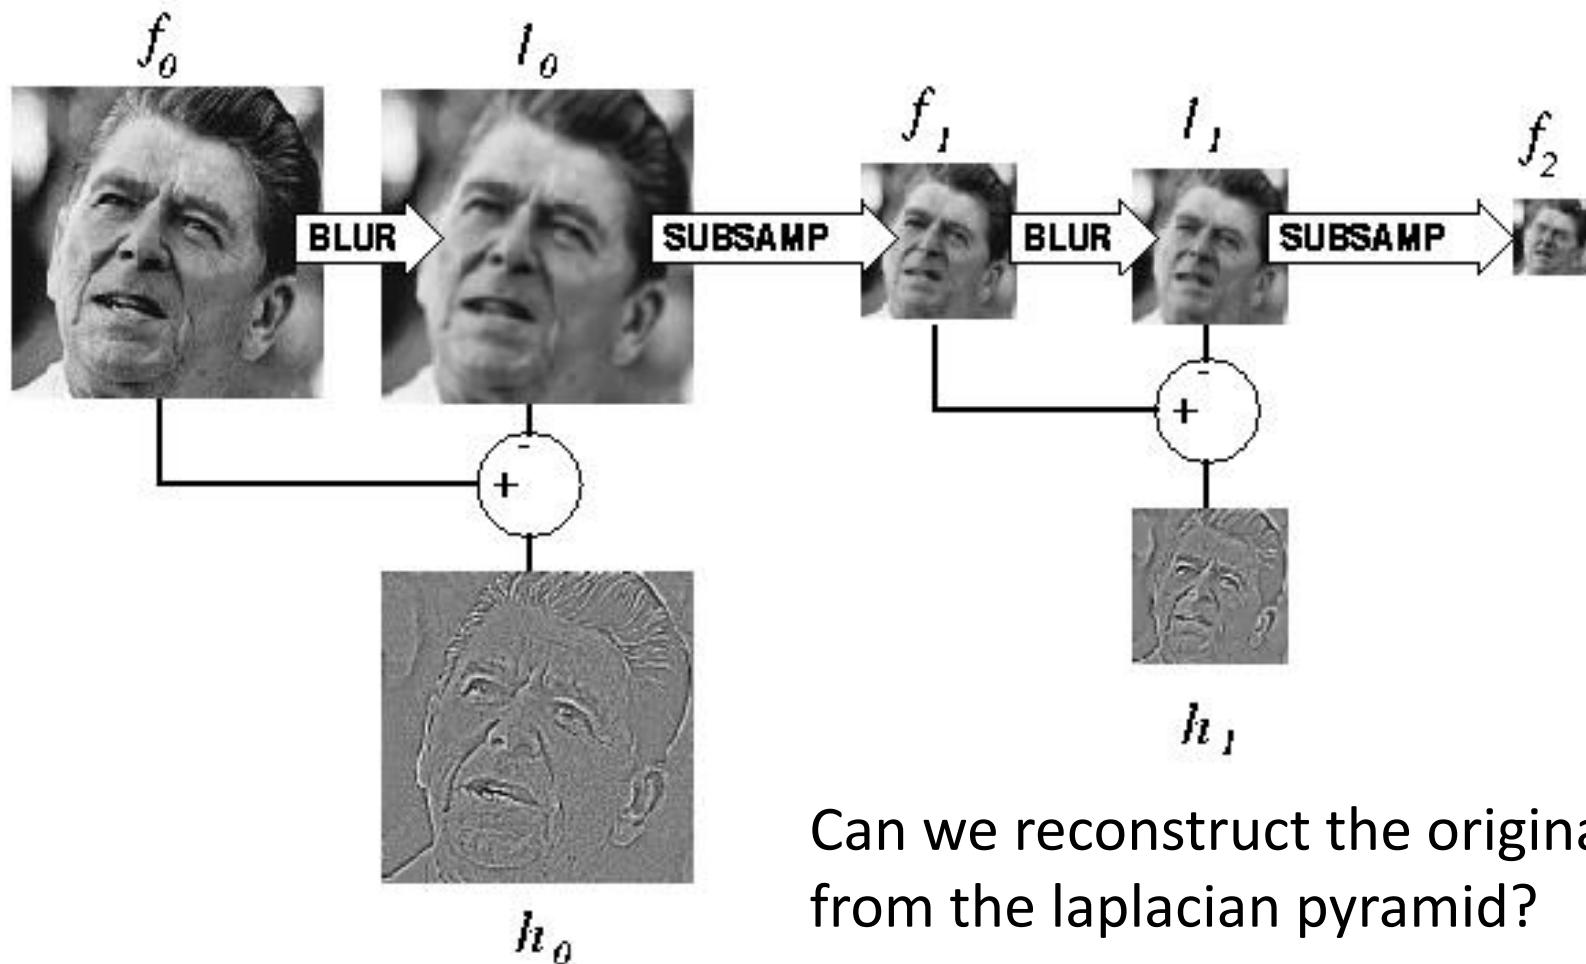
# Laplacian filter



# Laplacian pyramid

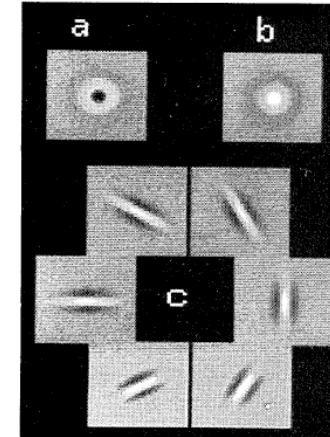
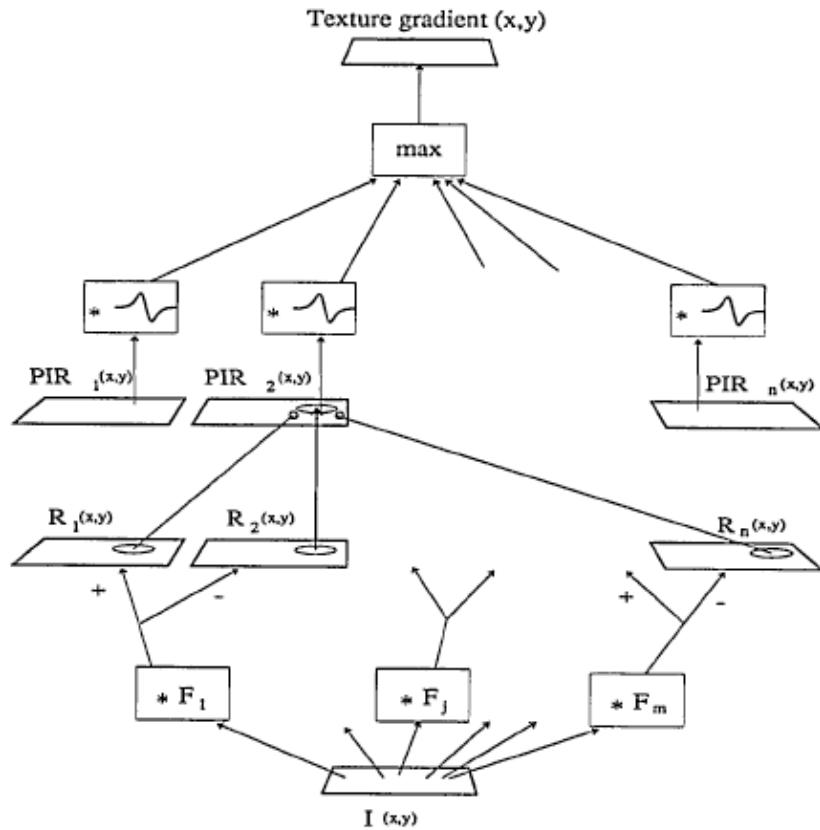


# Computing Gaussian/Laplacian Pyramid



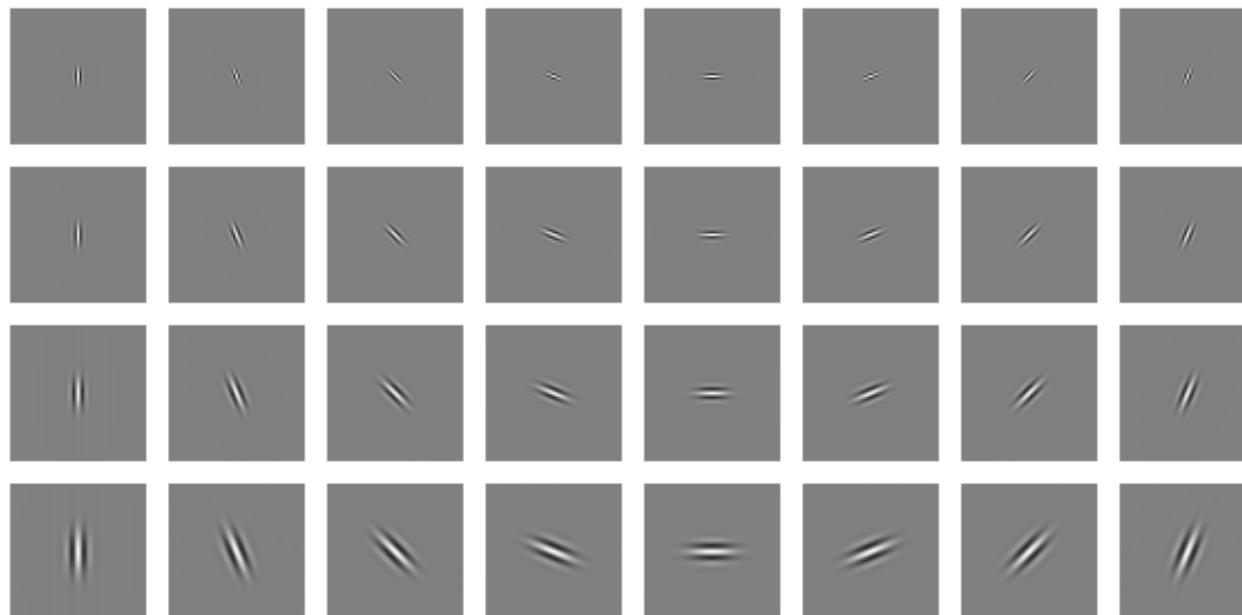
Can we reconstruct the original  
from the laplacian pyramid?

# Texture segmentation



# Clues from Human Perception

- Early processing in humans filters for various orientations and scales of frequency
- Perceptual cues in the mid-high frequencies dominate perception
- When we see an image from far away, we are effectively subsampling it



Early Visual Processing: Multi-scale edge and blob filters