

Grouping and Segmentation

CMP719– Computer Vision

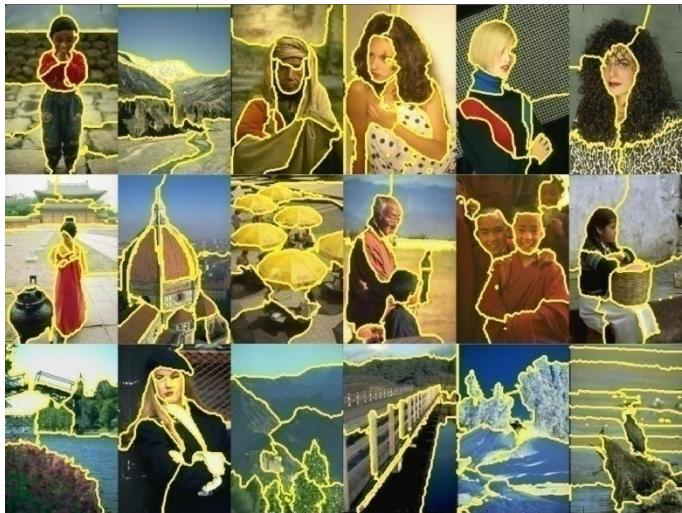
Pinar Duygulu

Hacettepe University

Grouping in vision

- Goals:
 - Gather features that belong together
 - Obtain an intermediate representation that compactly describes key image or video parts

Examples of grouping in vision

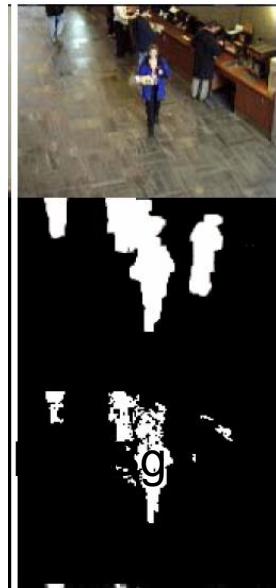


[Figure by J. Shi]

Determine image regions

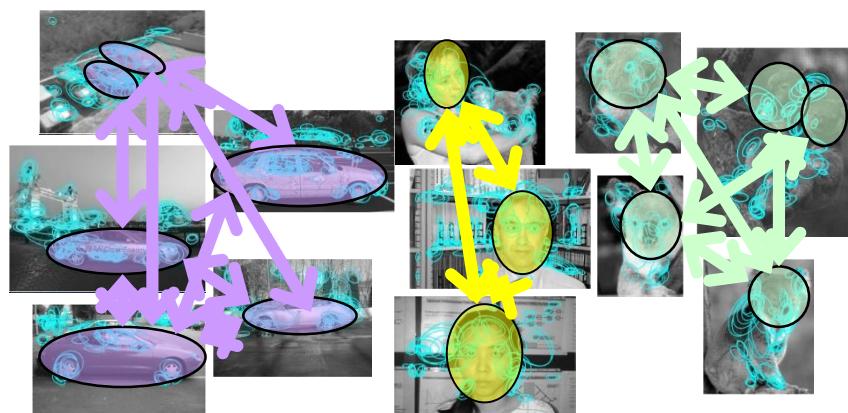


Group video frames into shots



[Figure by Wang & Suter]

Figure-ground



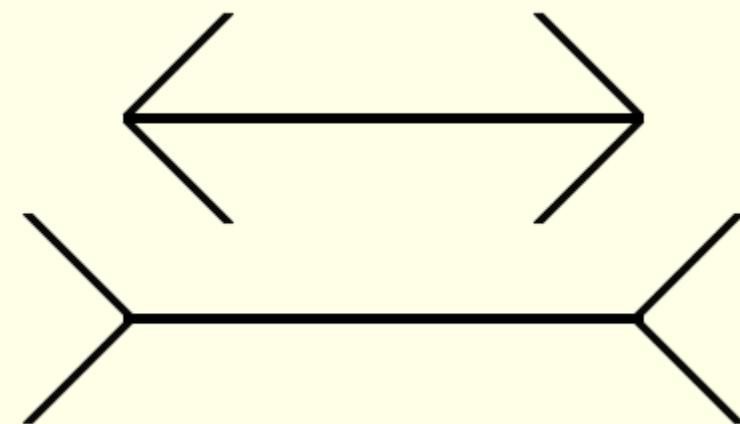
[Figure by Grauman & Darrell]

Object-level grouping

Grouping in vision

- Goals:
 - Gather features that belong together
 - Obtain an intermediate representation that compactly describes key image (video) parts
- Top down vs. bottom up segmentation
 - Top down: pixels belong together because they are from the same object
 - Bottom up: pixels belong together because they look similar
- Hard to measure success
 - What is interesting depends on the app.

Muller-Lyer illusion



What things should be grouped?
What cues indicate groups?

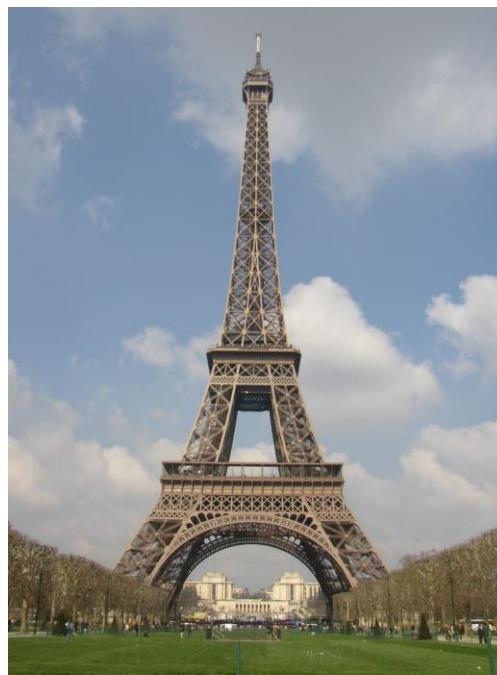
Gestalt

- Gestalt: whole or group
 - Whole is greater than sum of its parts
 - Relationships among parts can yield new properties/features
- Psychologists identified series of factors that predispose set of elements to be grouped (by human visual system)

Similarity



Symmetry



Common fate



Image credit: Arthus-Bertrand (via F. Durand)

Proximity



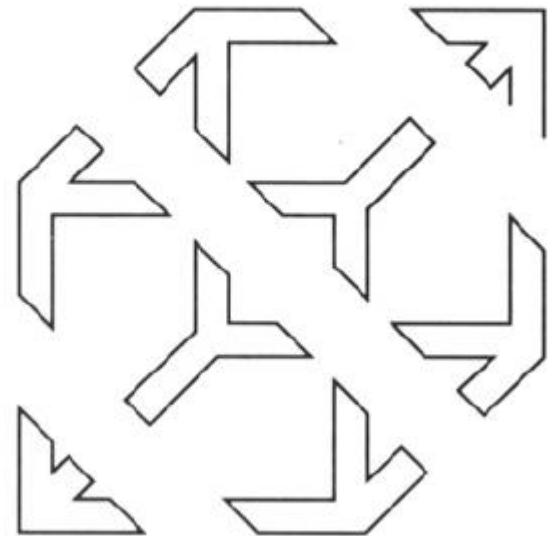
A “simple” segmentation problem

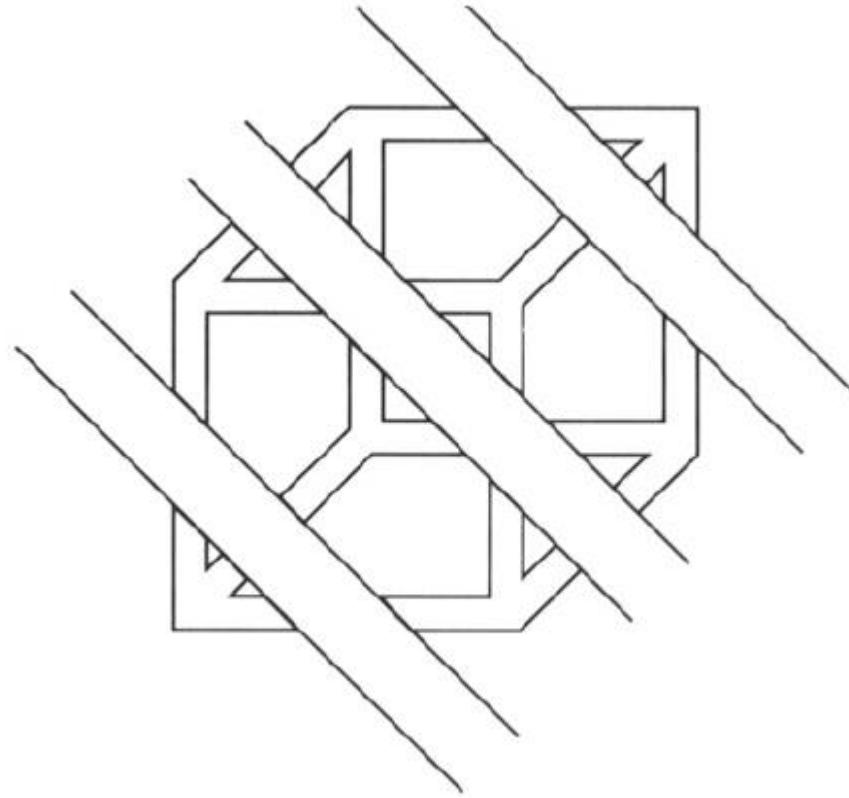


It can get a lot harder

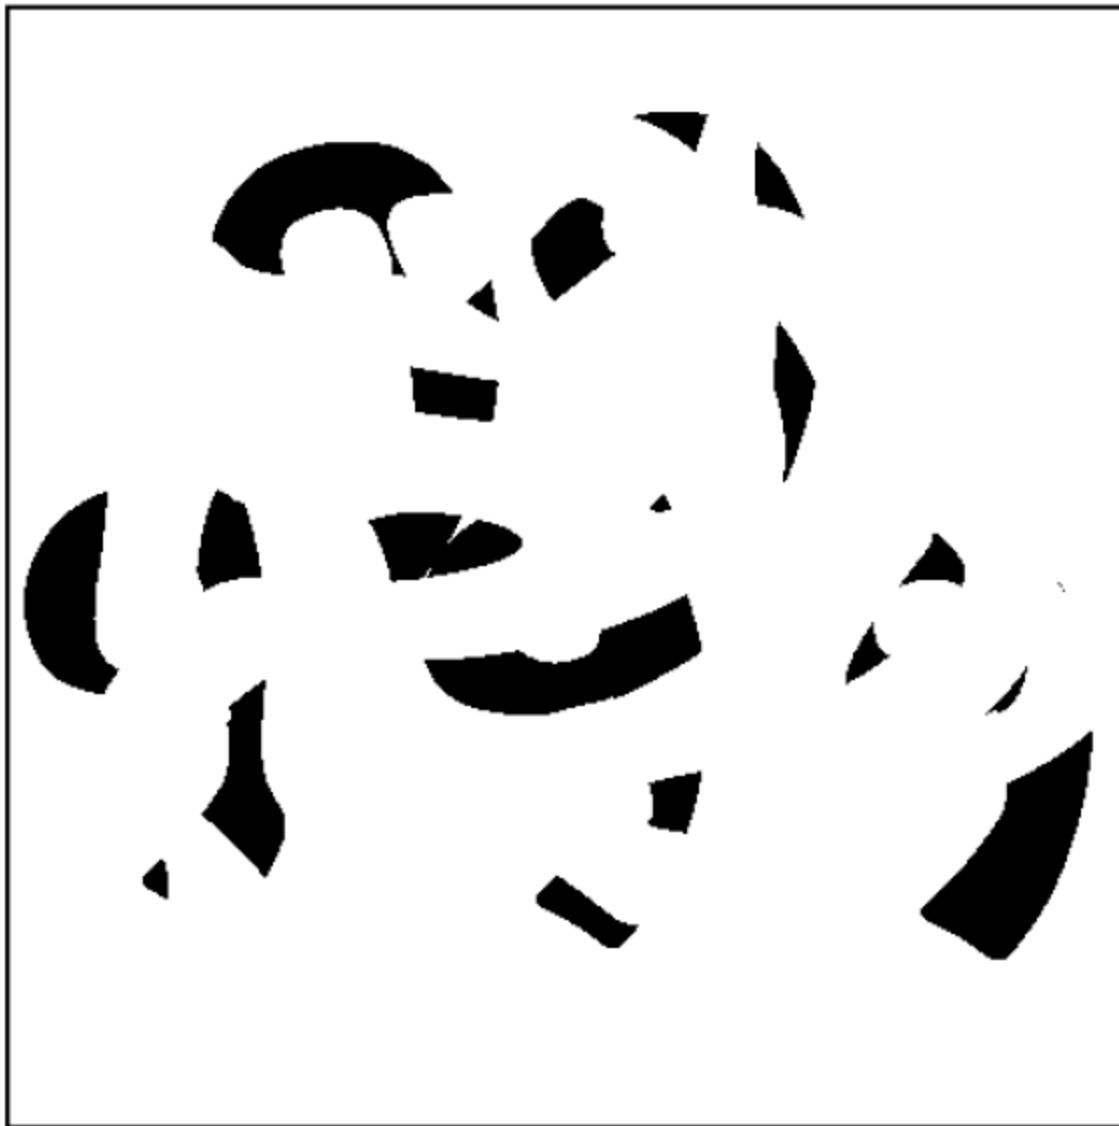


Brady, M. J., & Kersten, D. (2003). Bootstrapped learning of novel objects. *J Vis*, 3(6), 413-422





Continuity, explanation by occlusion



D. Forsyth





Magritte, 1957

Groupings by Invisible Completions

A



B



C



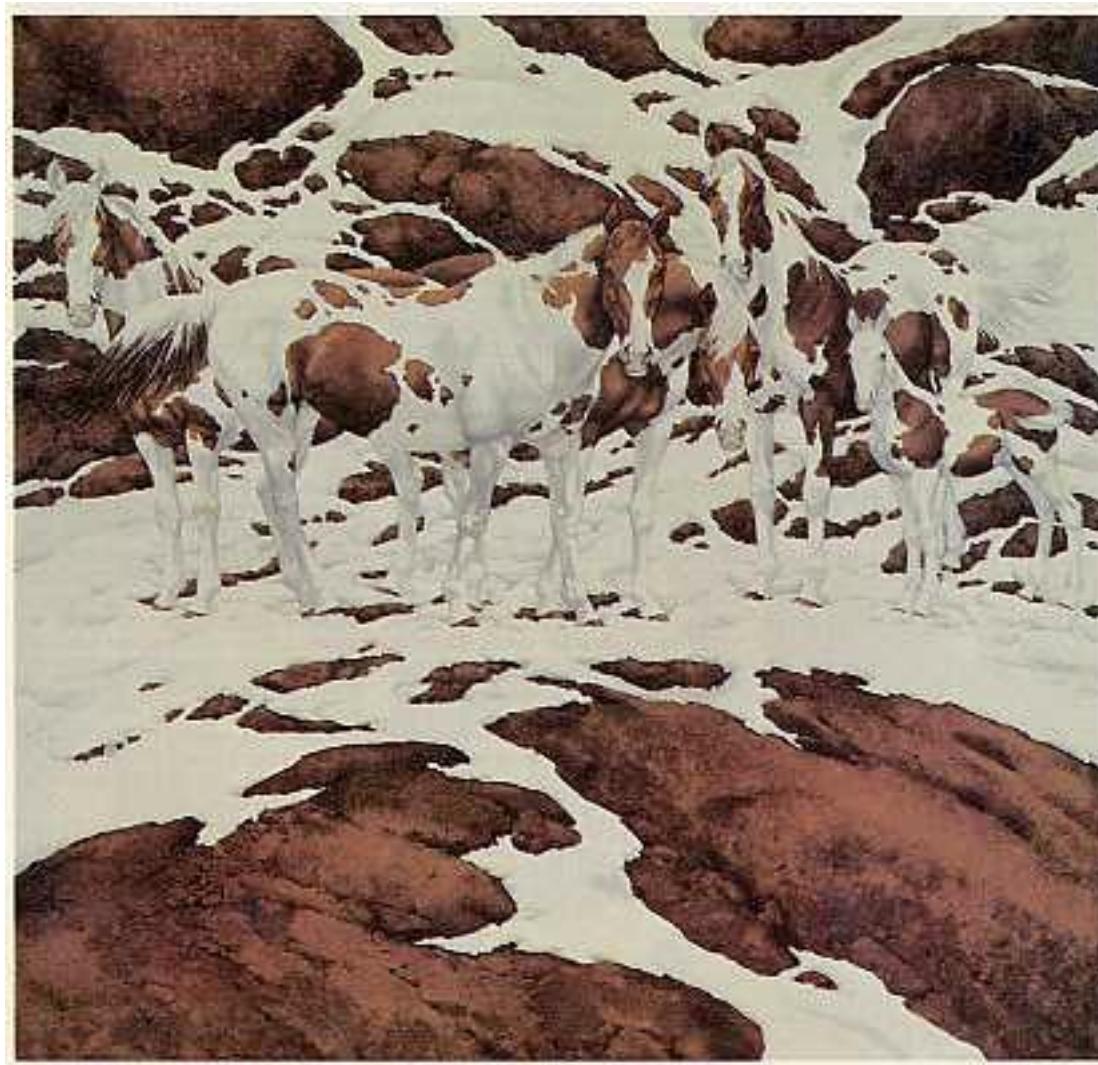
D



* Images from Steve Lehar's Gestalt papers



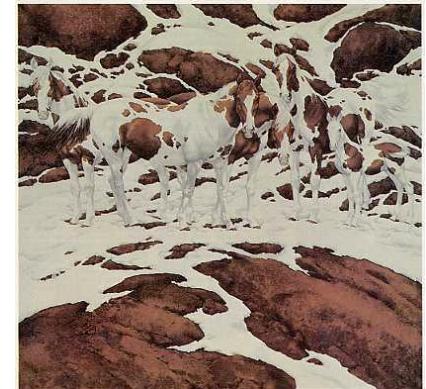
1970s: R. C. James



2000s: Bev Doolittle

Perceptual organization

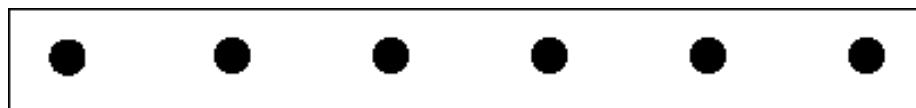
“...the processes by which the bits and pieces of visual information that are available in the retinal image are structured into the larger units of perceived objects and their interrelations”



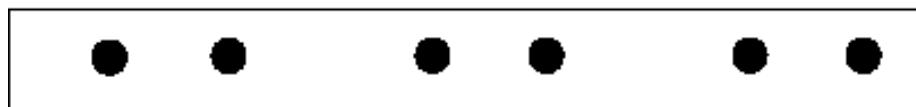
Stephen E. Palmer, *Vision Science*,
1999

Gestalt

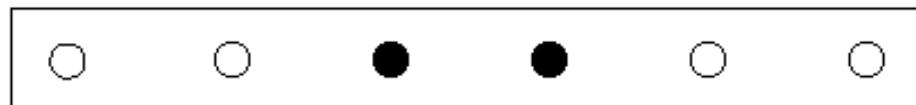
- Gestalt: whole or group
 - Whole is greater than sum of its parts
 - Relationships among parts can yield new properties/features
- Psychologists identified series of factors that predispose set of elements to be grouped (by human visual system)
- Inspiring observations/explanations; challenge remains how to best map to algorithms.



Not grouped



Proximity



Similarity



Similarity

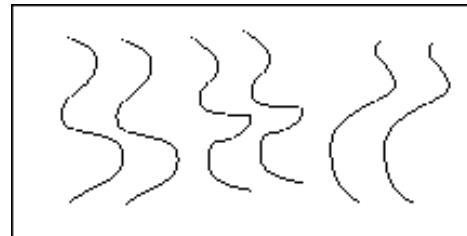


Common Fate

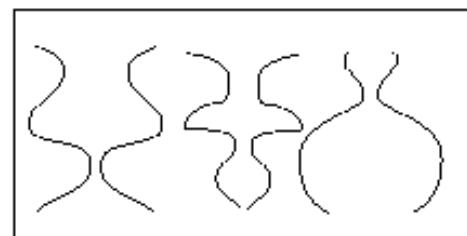


Common Region

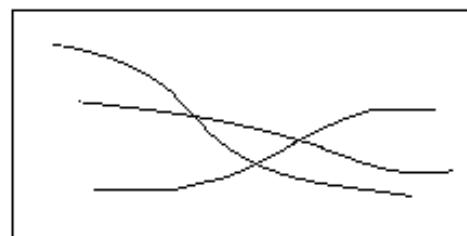




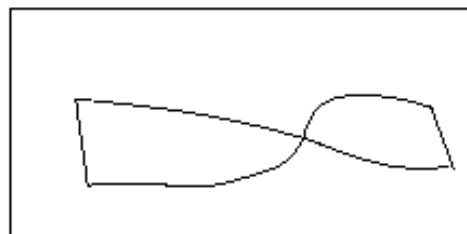
Parallelism



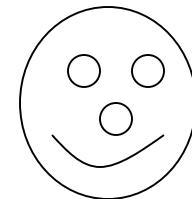
Symmetry



Continuity

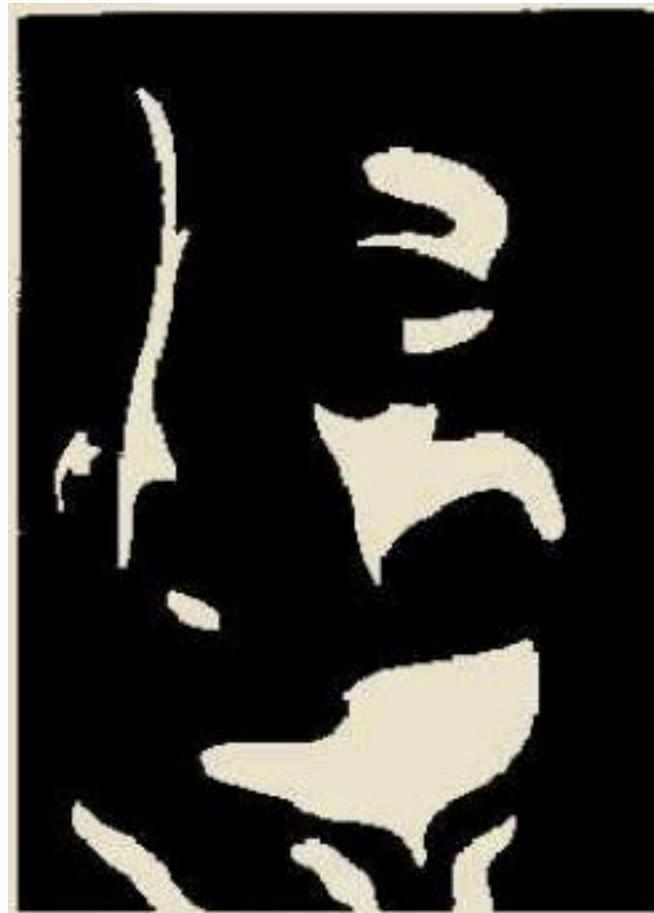


Closure



Familiar configuration

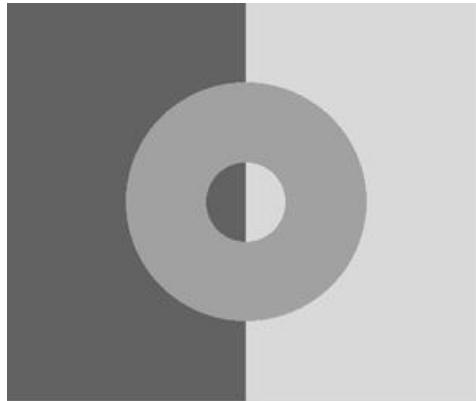
Familiarity



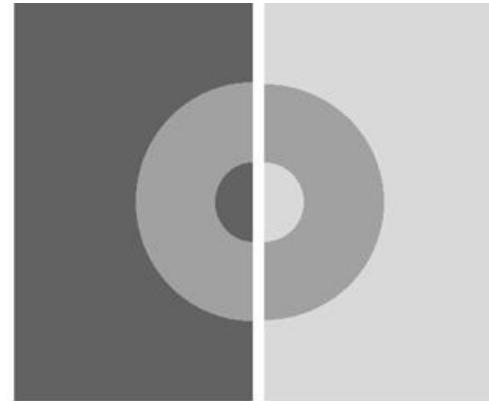
Familiarity



Influences of grouping



a



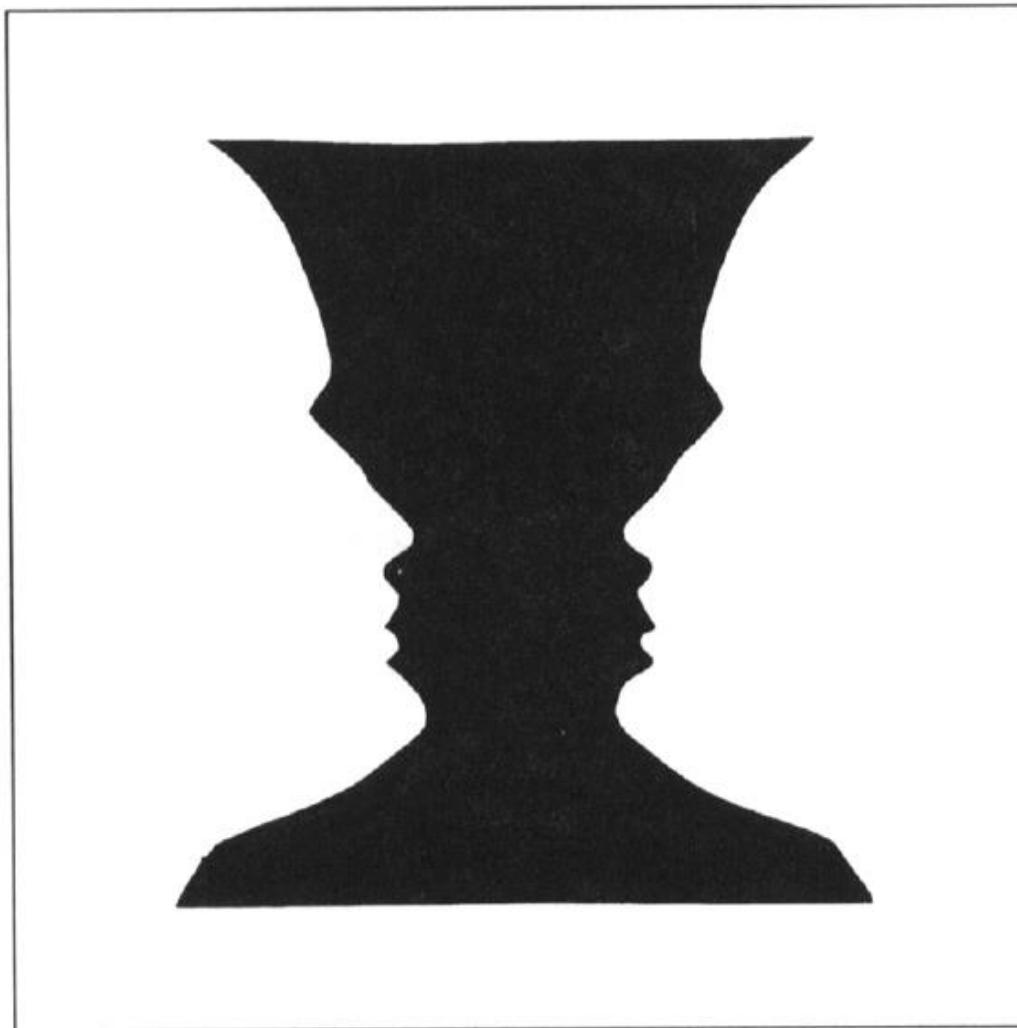
b



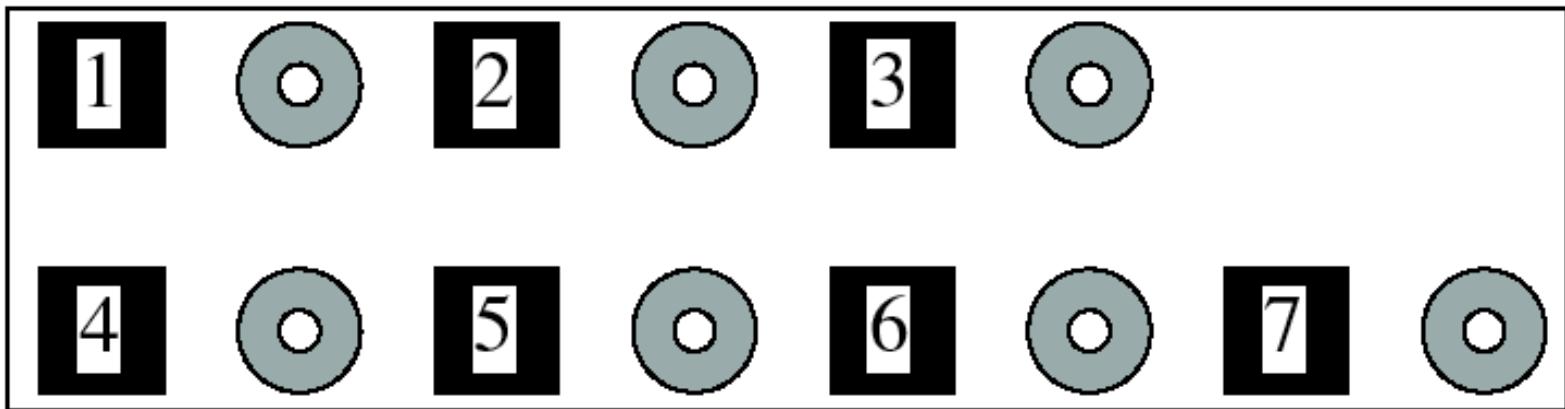
C

Grouping influences other perceptual mechanisms such as lightness perception

Figure-ground

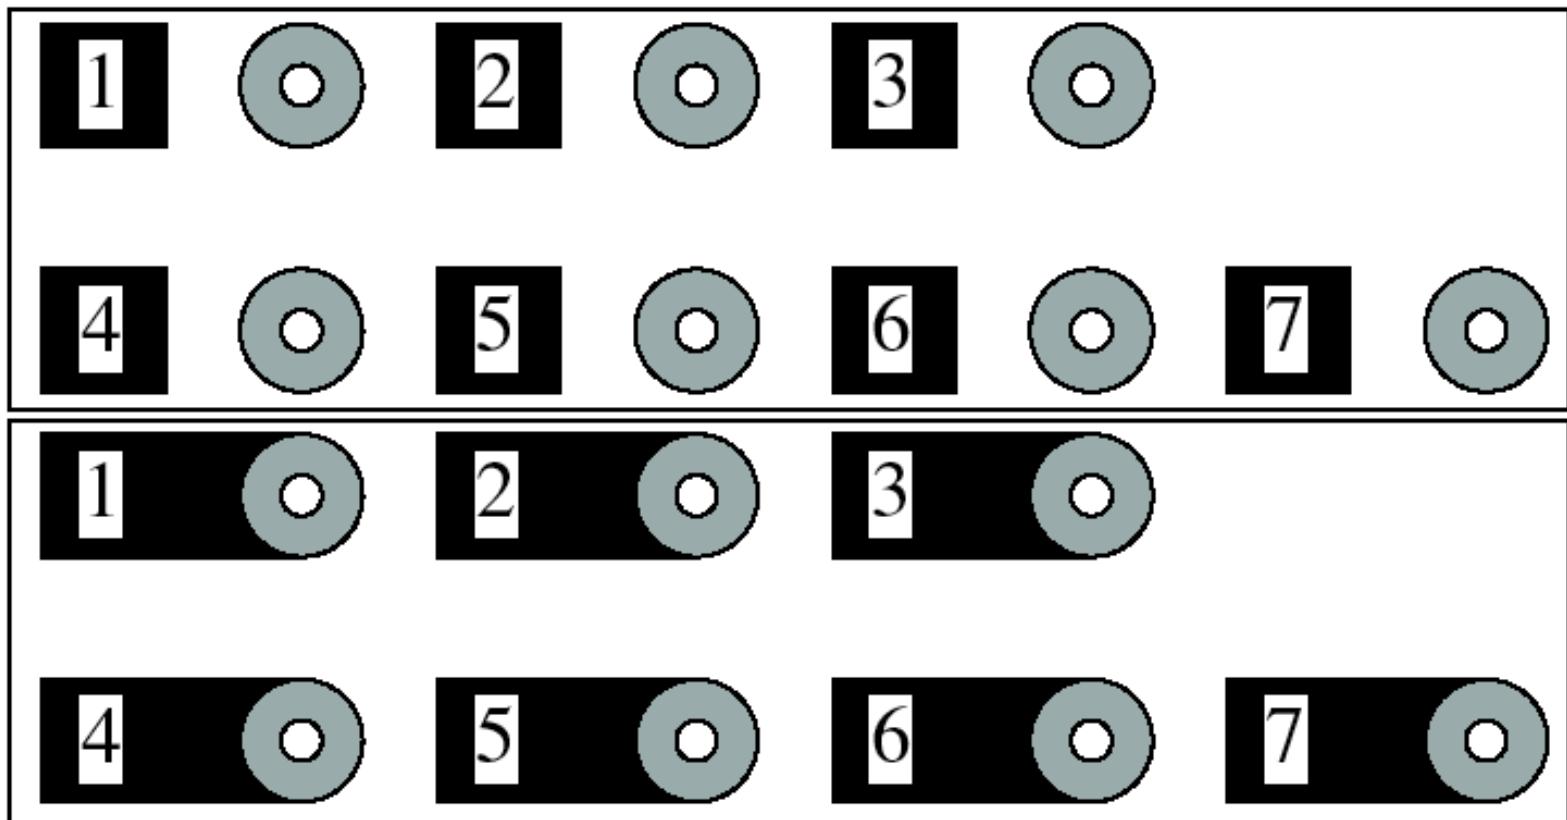


Grouping phenomena in real life



Forsyth & Ponce, Figure 14.7

Grouping phenomena in real life



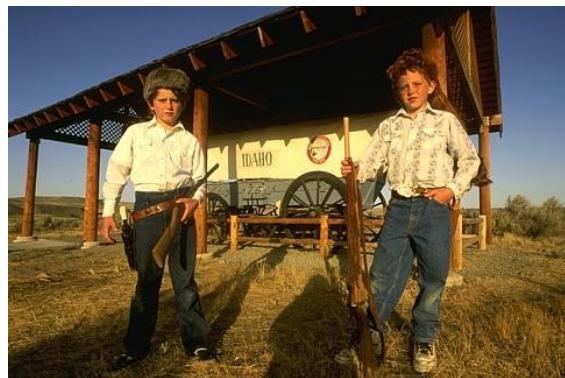
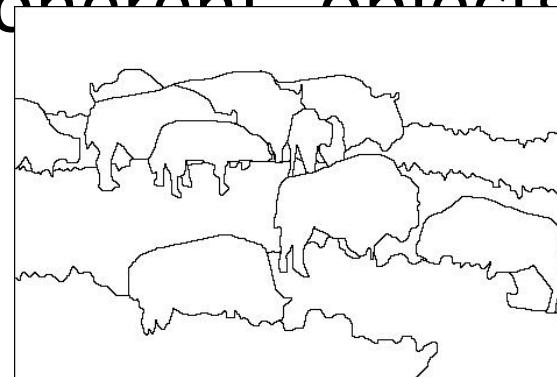
Forsyth & Ponce, Figure 14.7

The goals of segmentation

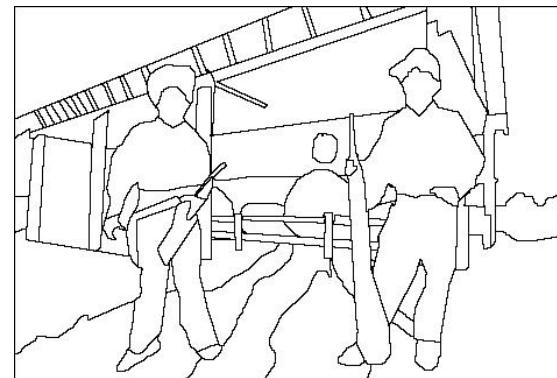
- Separate image into coherent “objects”



image



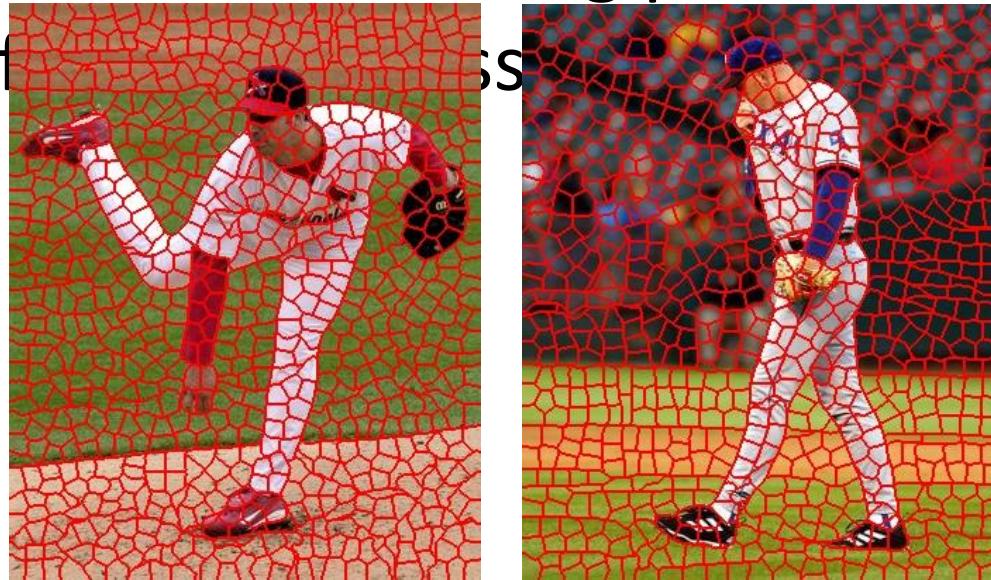
human segmentation



The goals of segmentation

- Separate image into coherent “objects”
- Group together similar-looking pixels for efficiency of processing

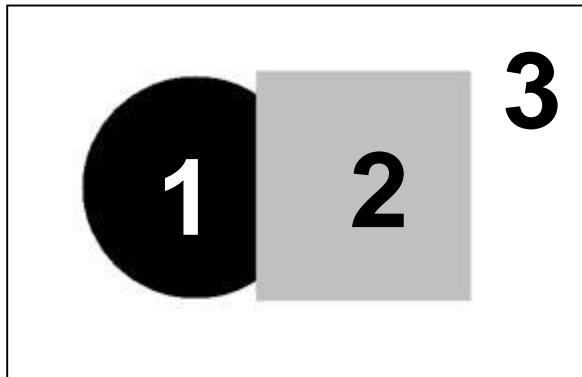
“superpixels”



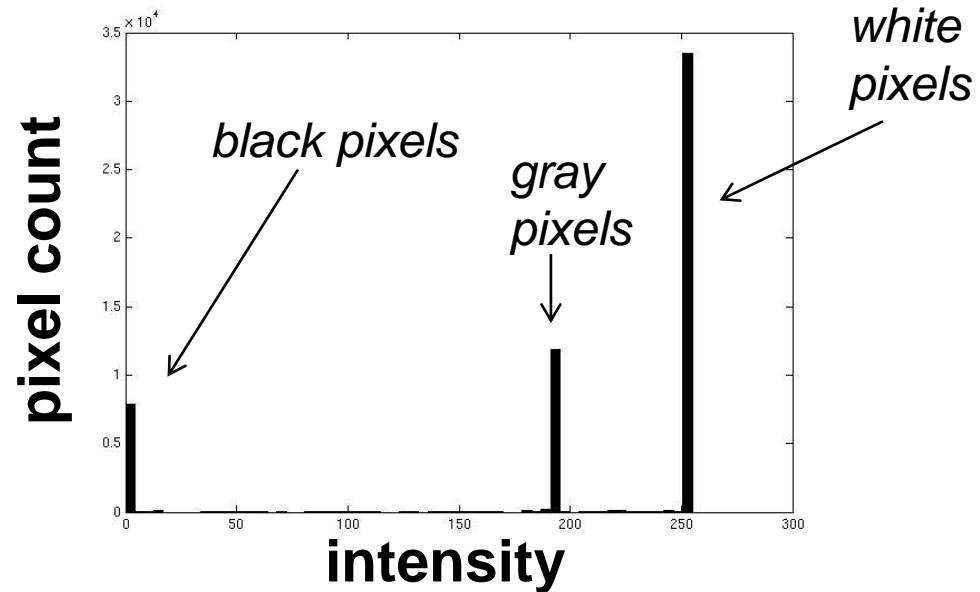
X. Ren and J. Malik. [Learning a classification model for segmentation](#). ICCV 2003.

Source: Lana Lazebnik

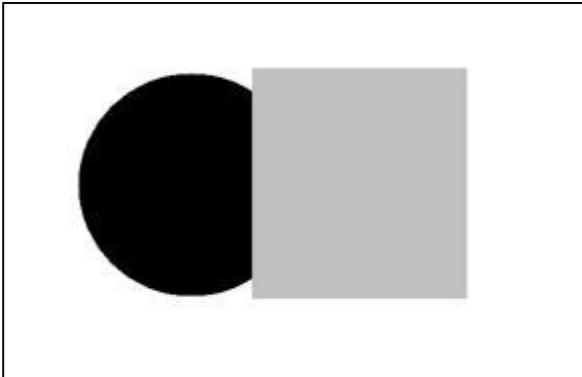
Image segmentation: toy example



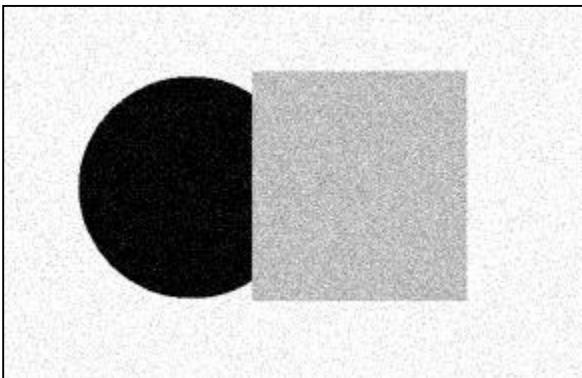
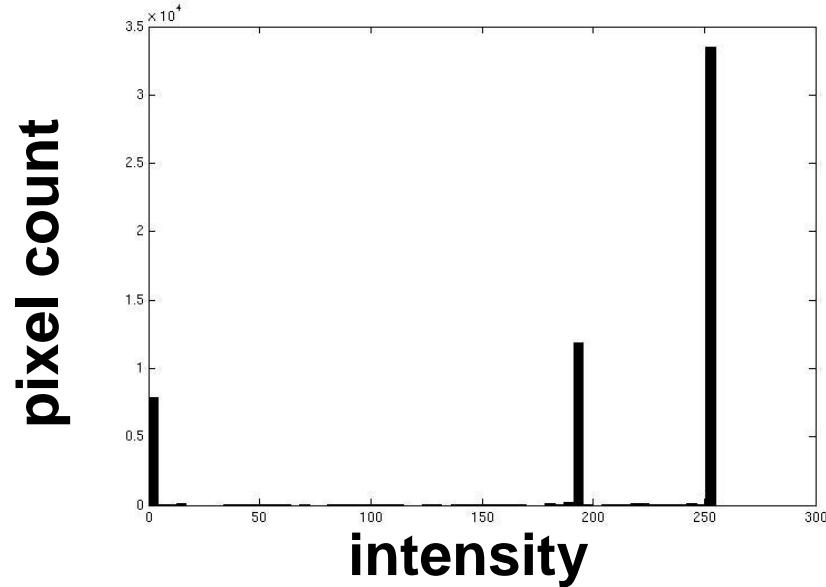
input image



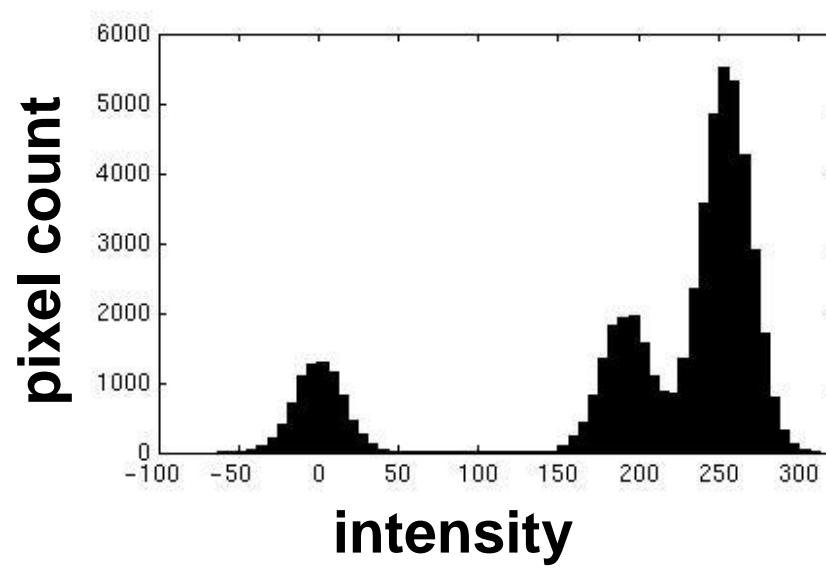
- These intensities define the three groups.
- We could label every pixel in the image according to which of these primary intensities it is.
 - i.e., *segment* the image based on the intensity feature.
- What if the image isn't quite so simple?

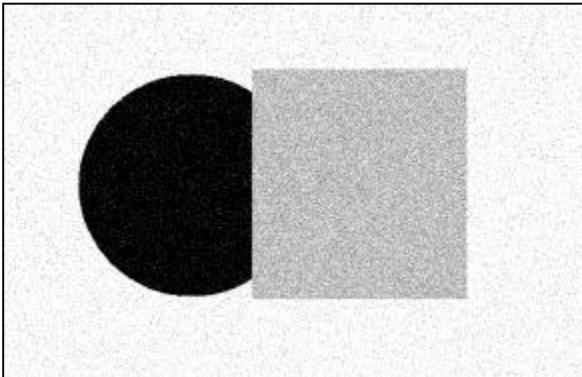


input image

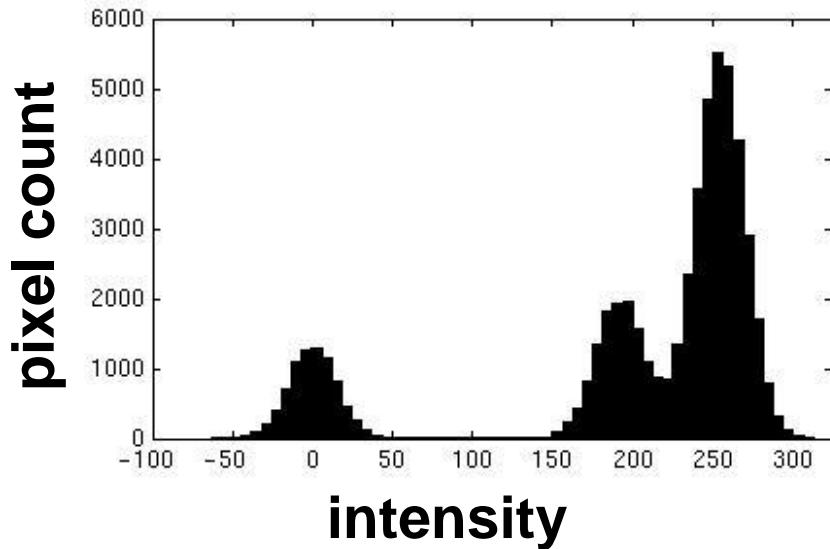


input image

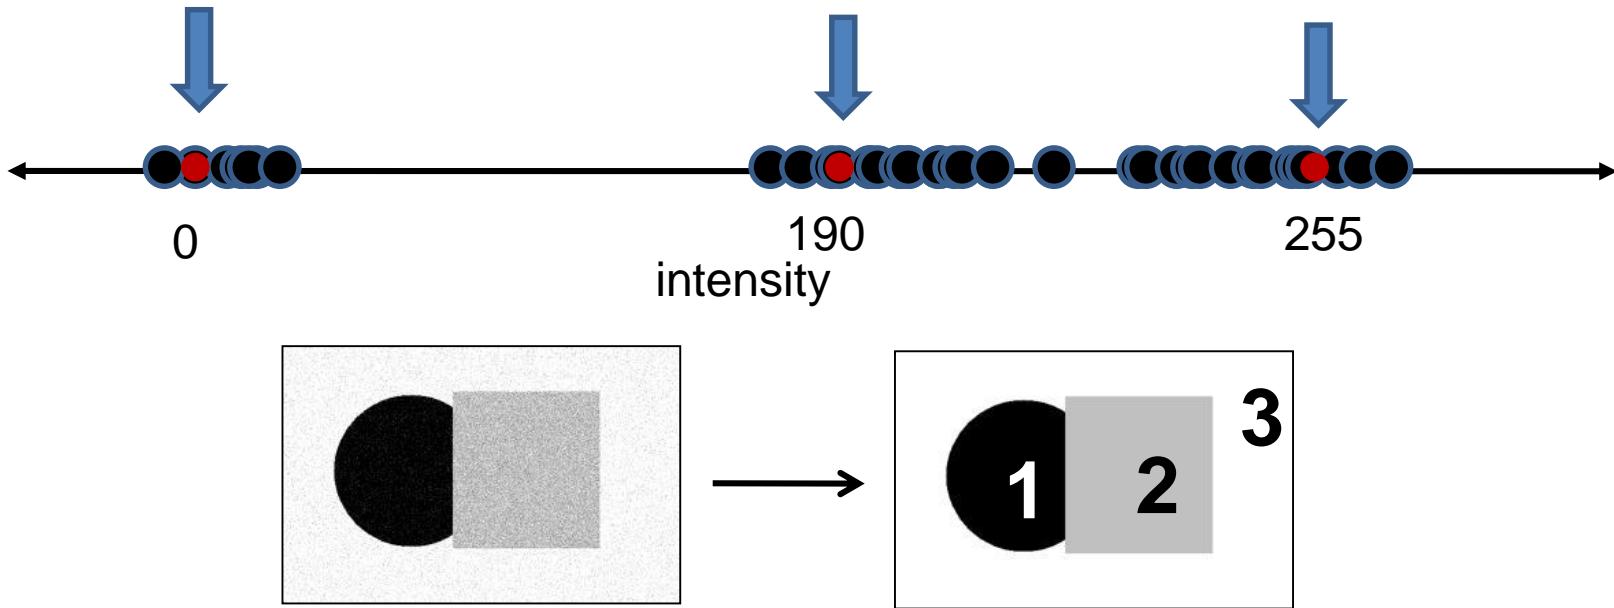




input image



- Now how to determine the three main intensities that define our groups?
- We need to ***cluster***.

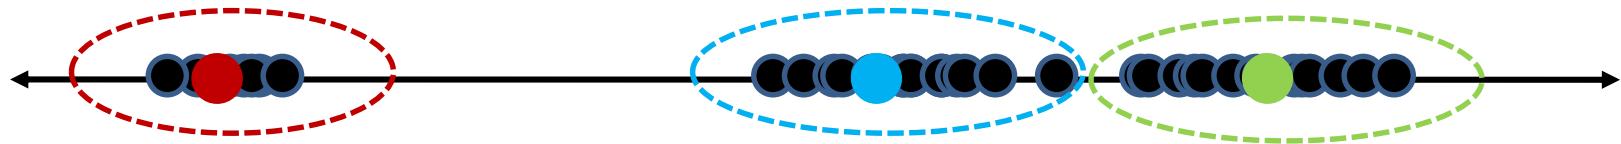


- Goal: choose three “centers” as the **representative** intensities, and label every pixel according to which of these centers it is nearest to.
- Best cluster centers are those that minimize SSD between all points and their nearest cluster center c_i :

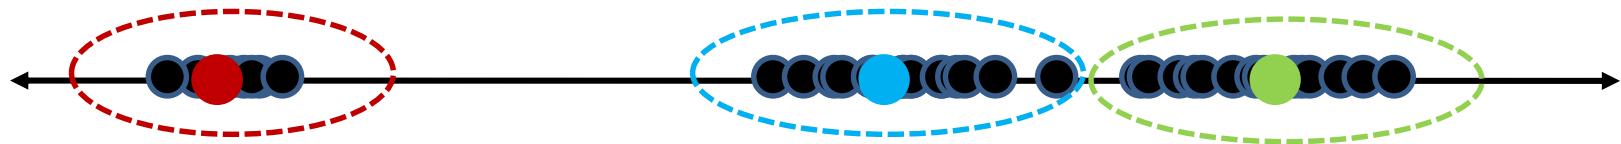
$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

Clustering

- With this objective, it is a “chicken and egg” problem:
 - If we knew the **cluster centers**, we could allocate points to groups by assigning each to its closest center.



- If we knew the **group memberships**, we could get the centers by computing the mean per group.



K-means clustering

- Basic idea: randomly initialize the k cluster centers, and iterate between the two steps we just saw.
 1. Randomly initialize the cluster centers, c_1, \dots, c_K
 2. Given cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
 3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
 4. If c_i have changed, repeat Step 2



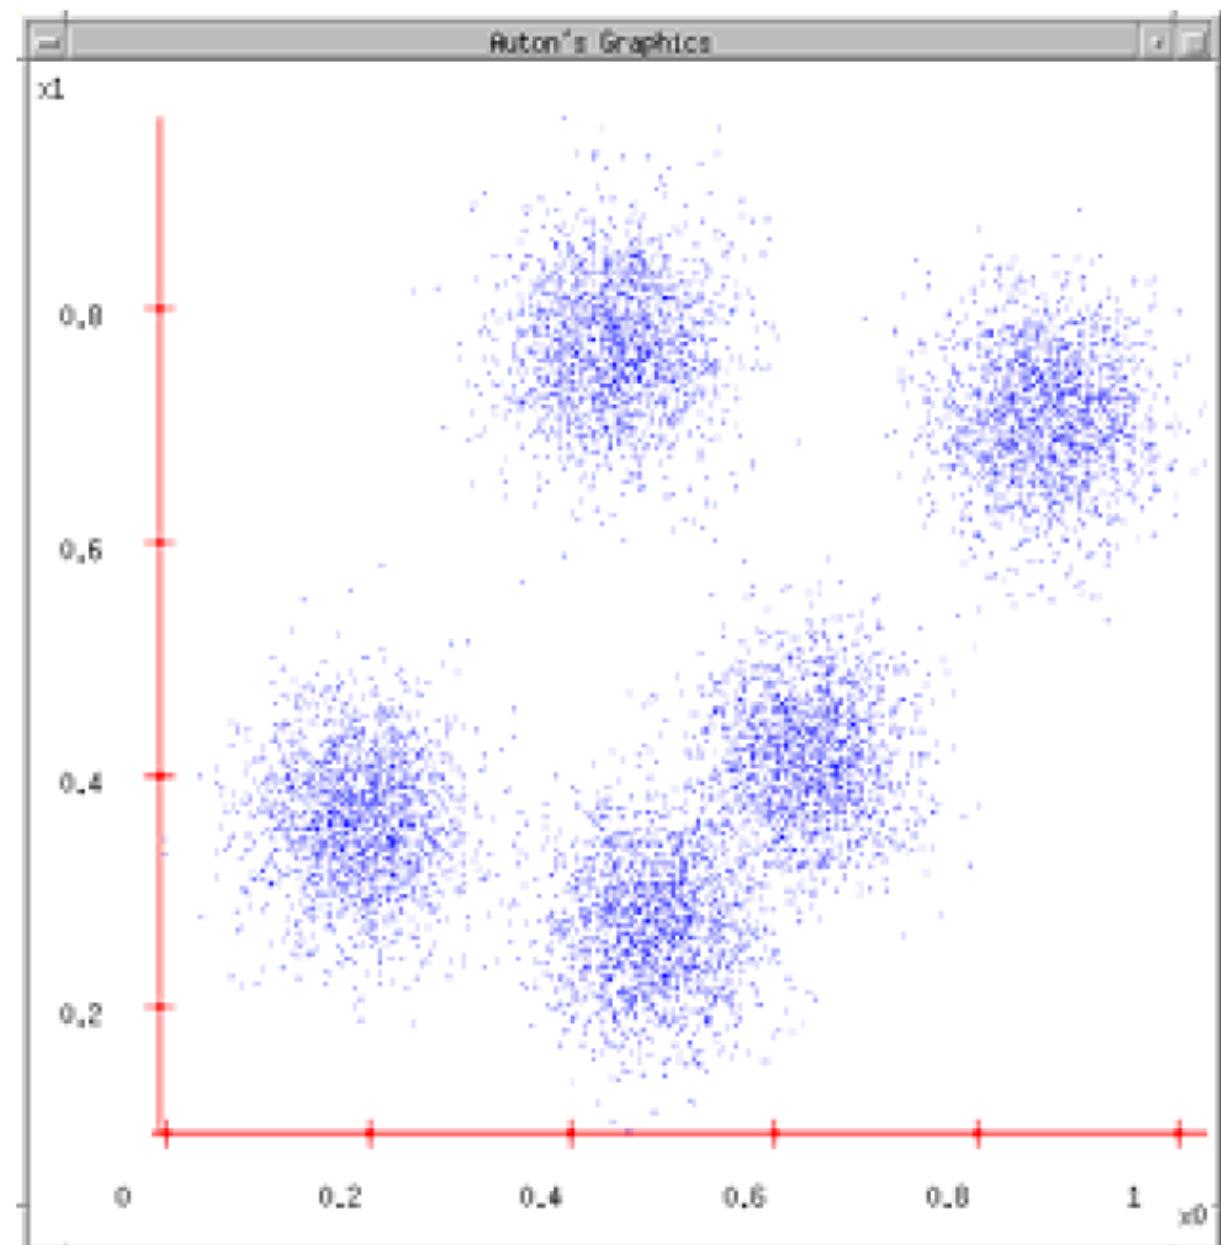
Properties

- Will always converge to *some* solution
- Can be a “local minimum”
 - does not always find the global minimum of objective function:

$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$

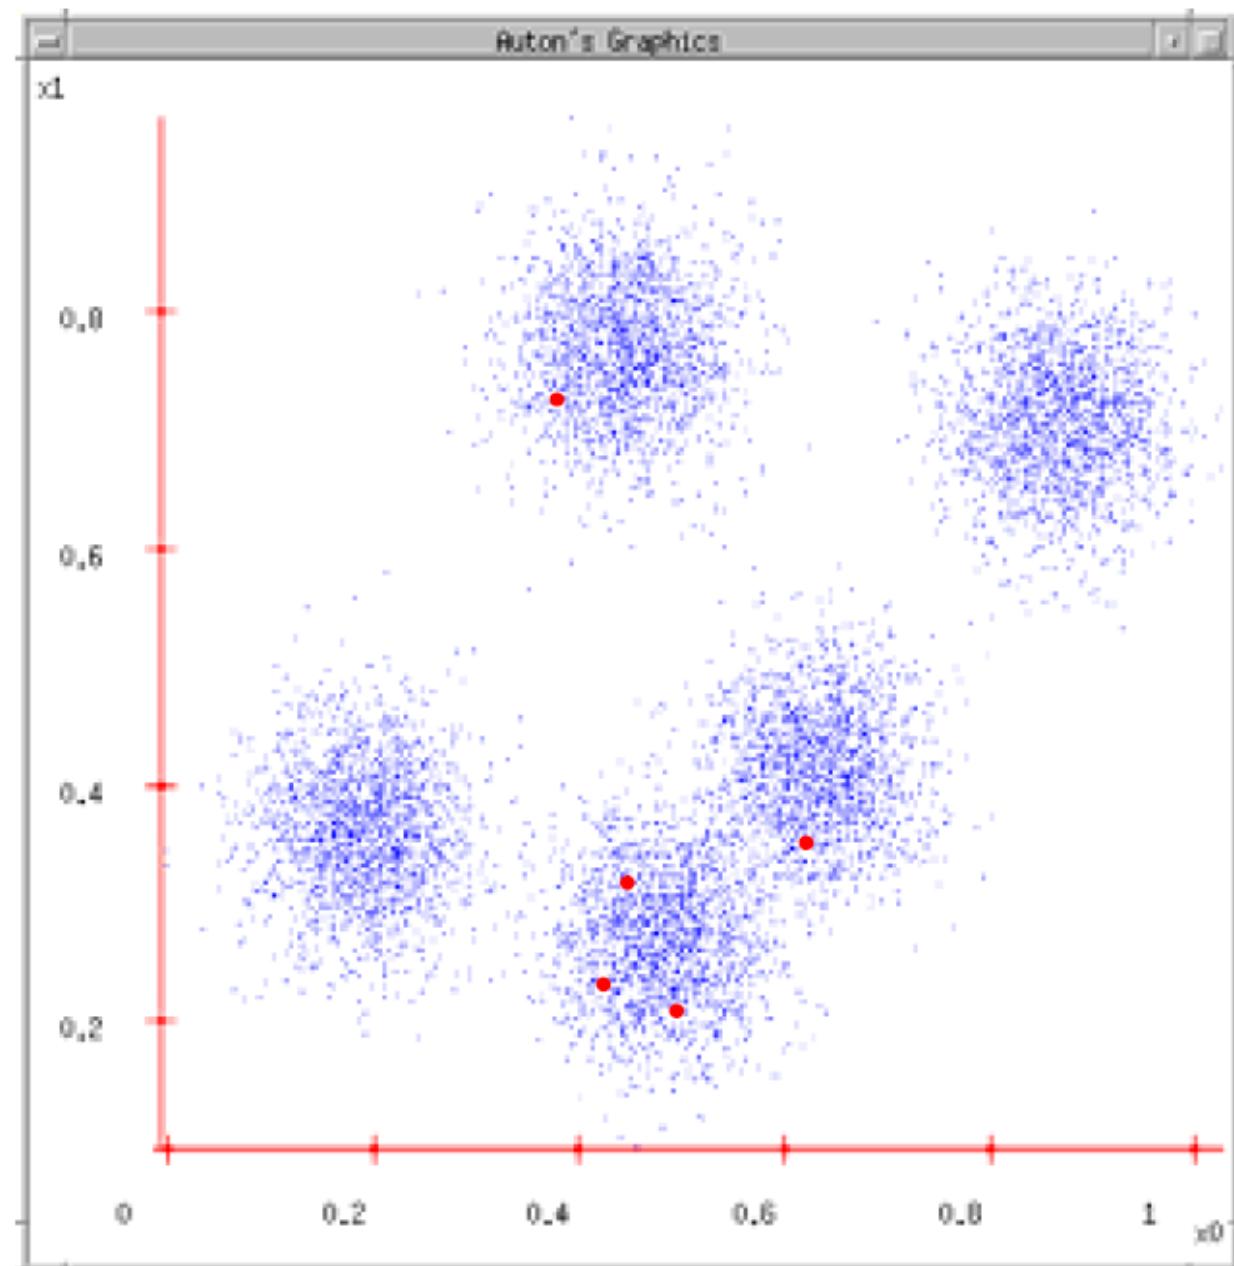
K-means

1. Ask user how many clusters they'd like.
(e.g. k=5)



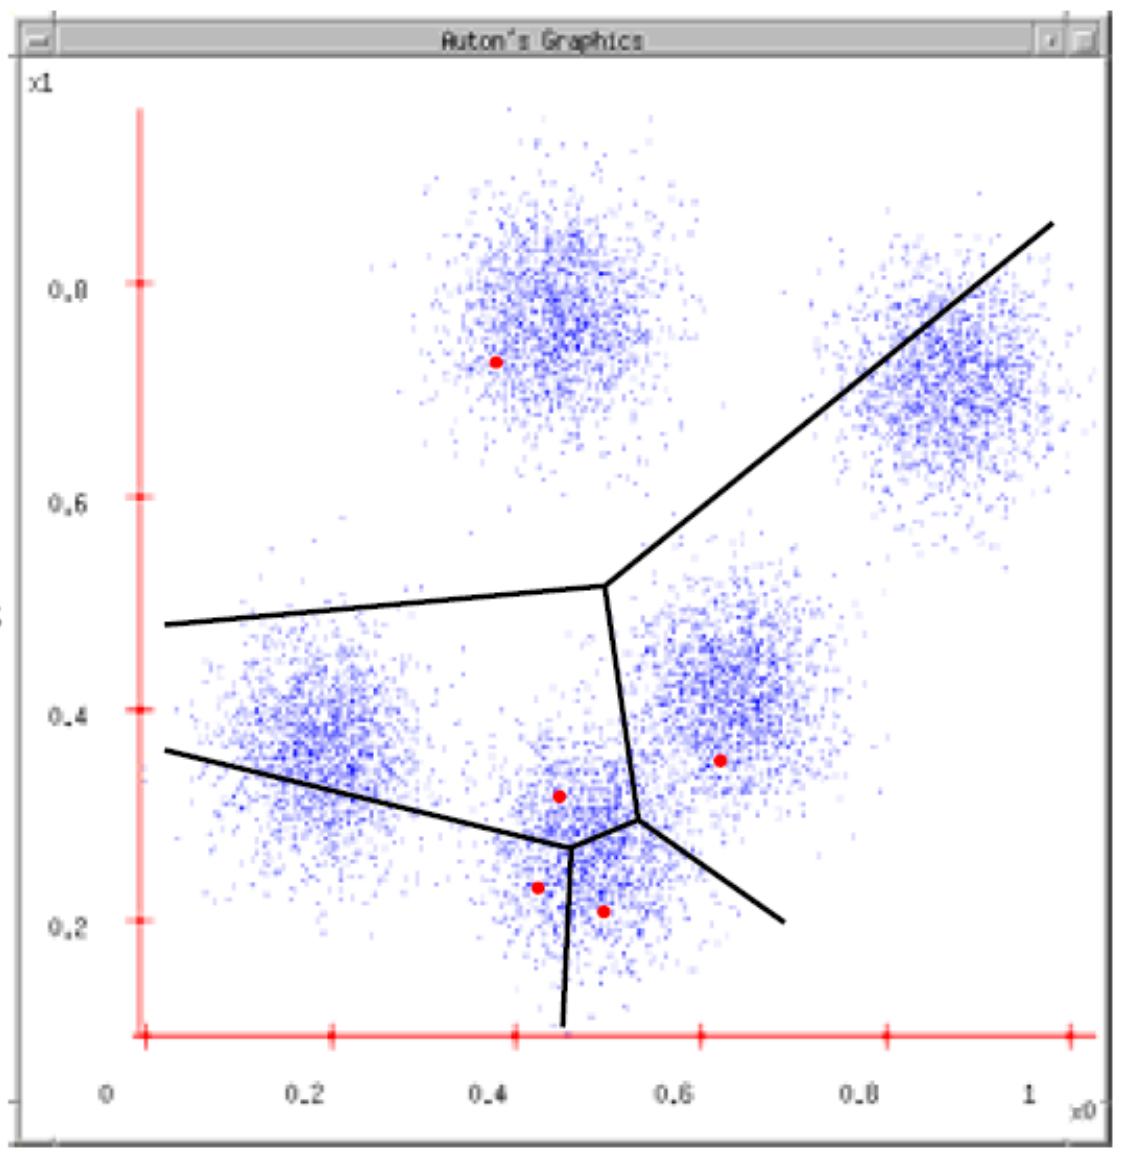
K-means

1. Ask user how many clusters they'd like.
(e.g. k=5)
2. Randomly guess k cluster Center locations



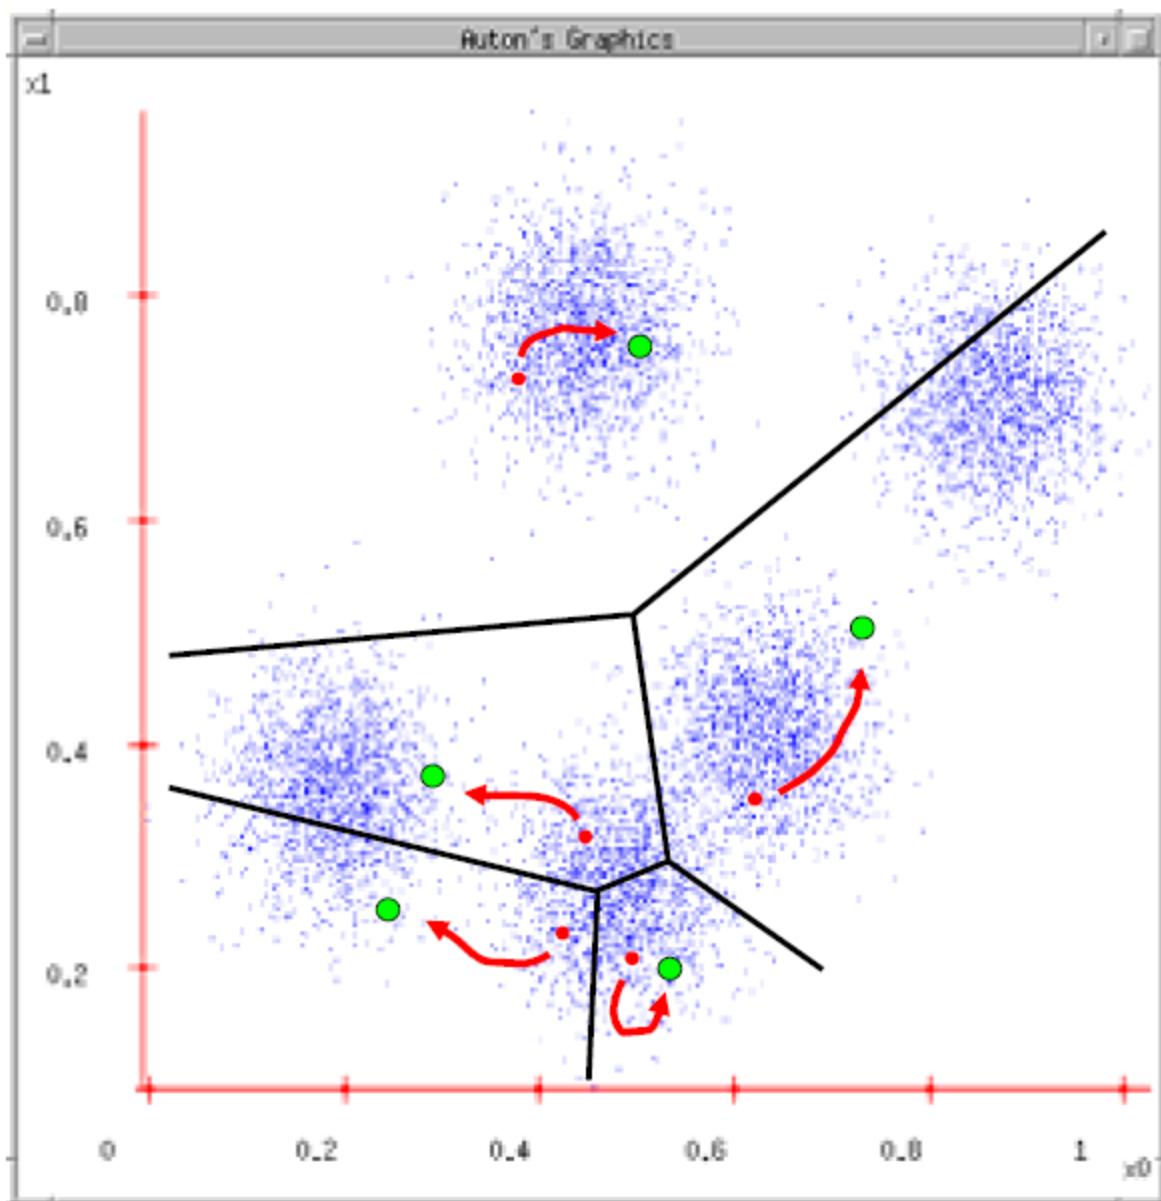
K-means

1. Ask user how many clusters they'd like.
(e.g. k=5)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to. (Thus each Center "owns" a set of datapoints)



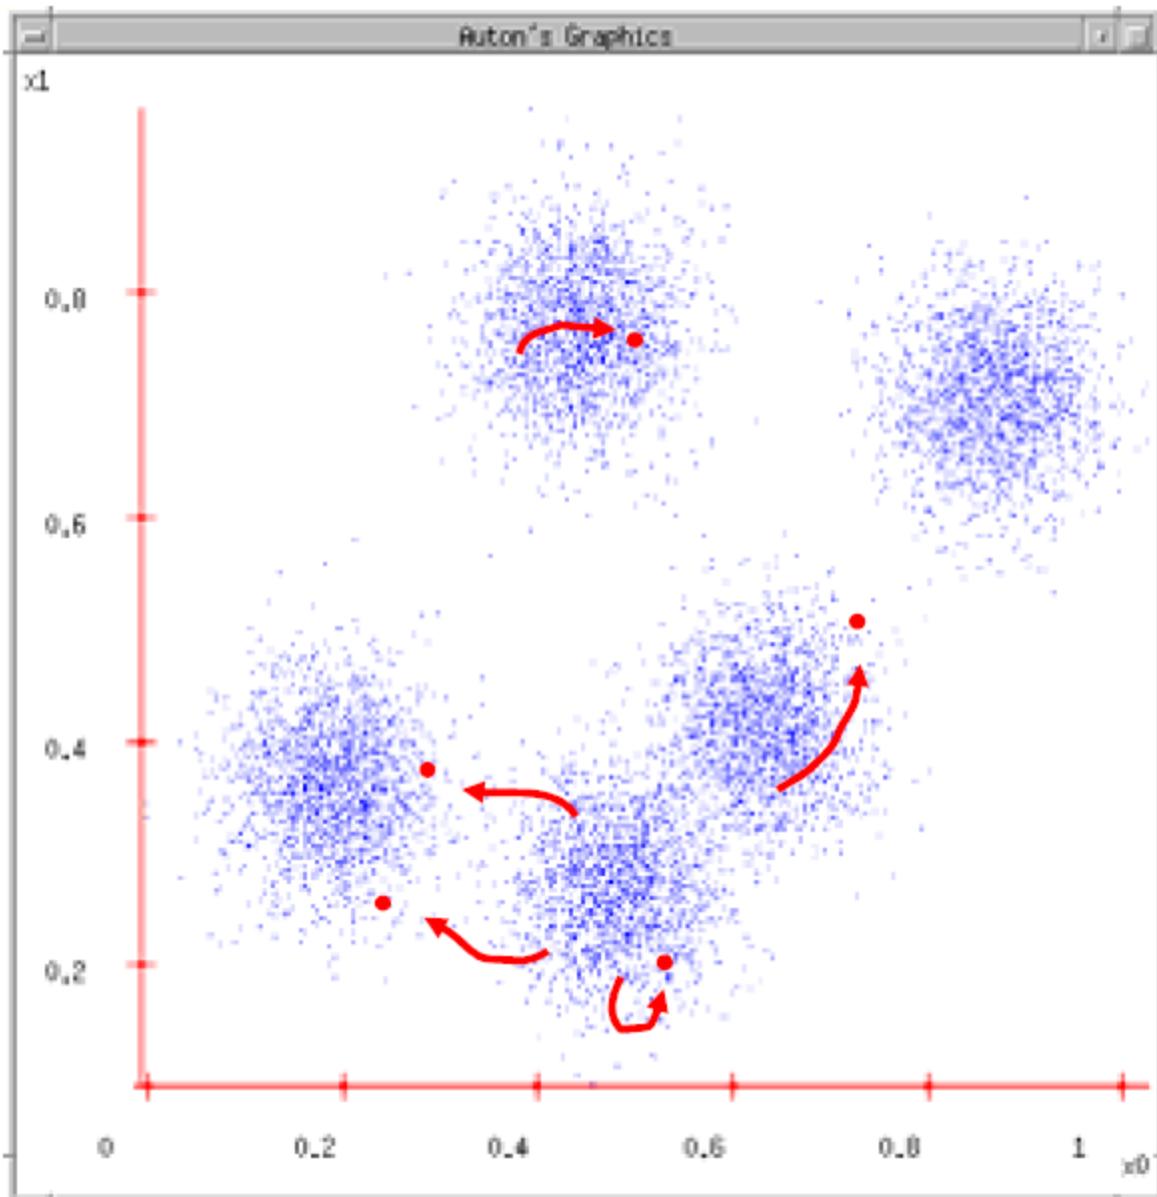
K-means

1. Ask user how many clusters they'd like.
(e.g. k=5)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns



K-means

1. Ask user how many clusters they'd like.
(e.g. $k=5$)
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns...
5. ...and jumps there
6. ...Repeat until terminated!



K-means clustering

- Java demo:

<http://kovan.ceng.metu.edu.tr/~maya/kmeans/index.html>

http://home.dei.polimi.it/matteucc/Clustering/tutorial_html/AppletKM.html

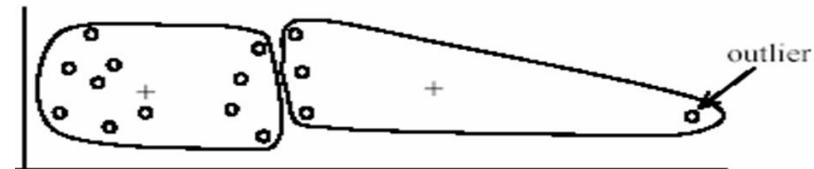
K-means: pros and cons

Pros

- Simple, fast to compute
- Converges to local minimum of within-cluster squared error

Cons/issues

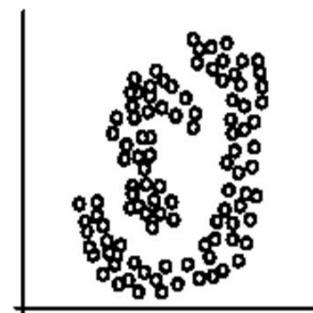
- Setting k ?
- Sensitive to initial centers
- Sensitive to outliers
- Detects spherical clusters
- Assuming means can be computed



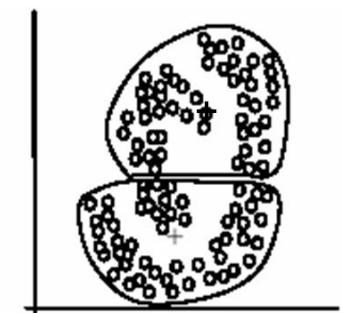
(A): Undesirable clusters



(B): Ideal clusters



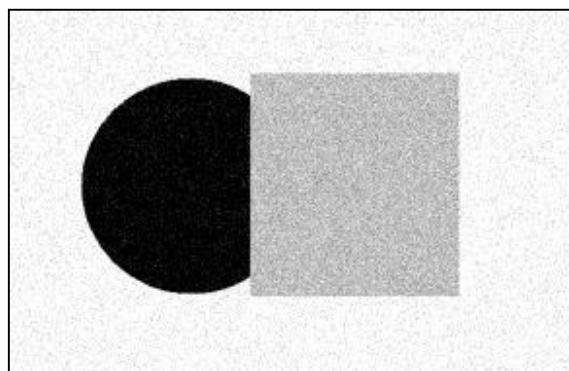
(A): Two natural clusters



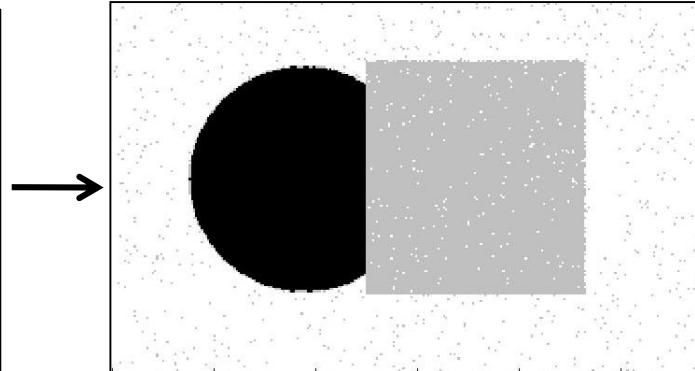
(B): k -means clusters

An aside: Smoothing out cluster assignments

- Assigning a cluster label per pixel may yield outliers:

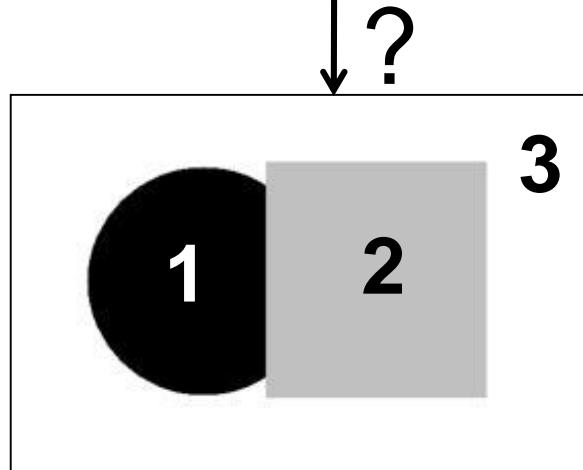


original



labeled by cluster center's
intensity

- How to ensure they are spatially smooth?



Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based
on **intensity** similarity



Feature space: intensity value (1-d)

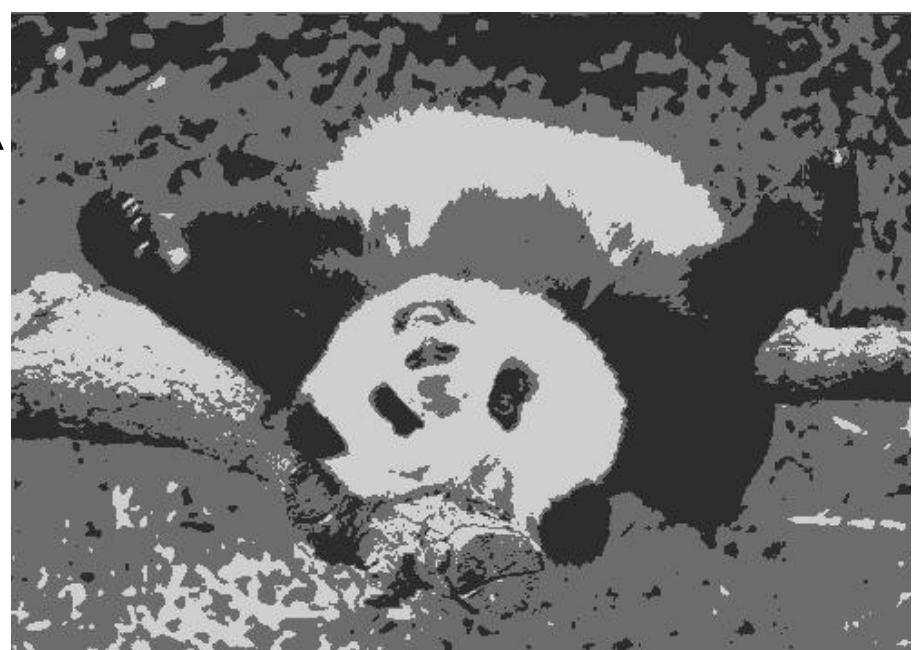


K=2



*quantization of the feature space;
segmentation label map*

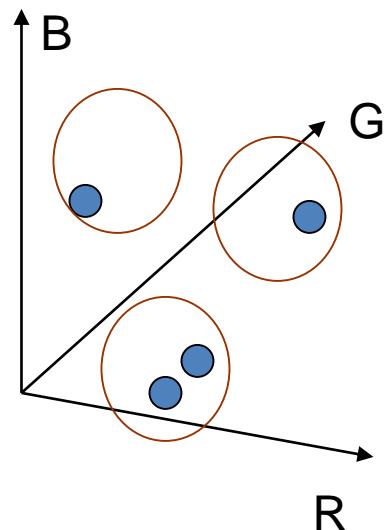
K=3



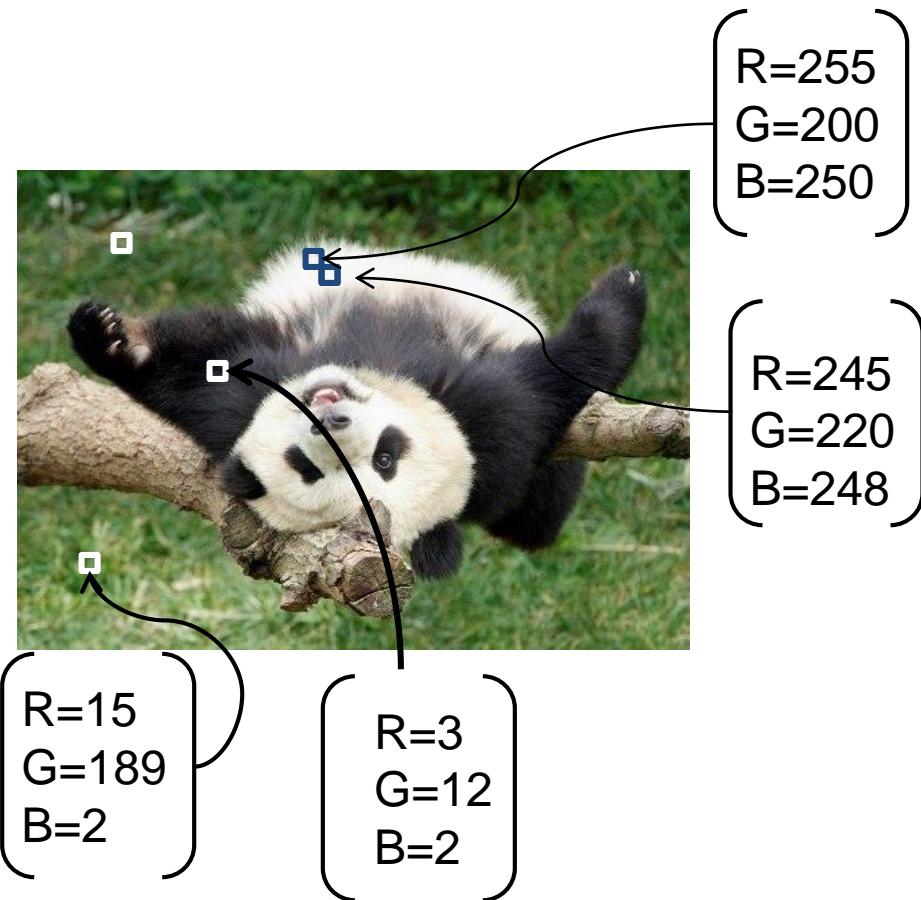
Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based
on **color** similarity



Feature space: color value (3-d)



Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based
on **intensity** similarity



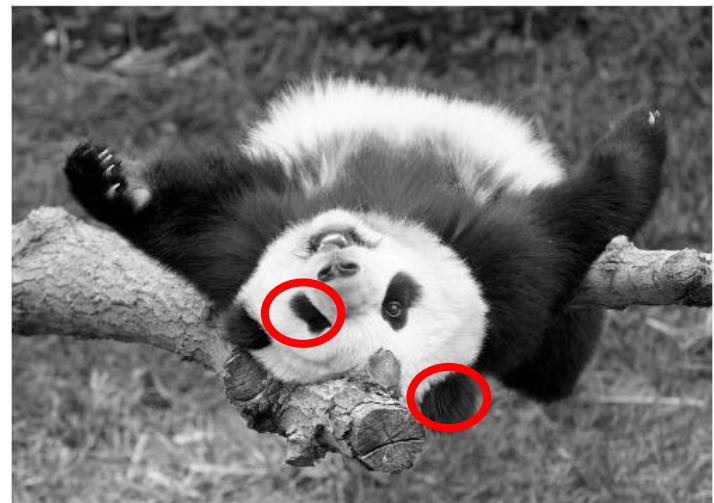
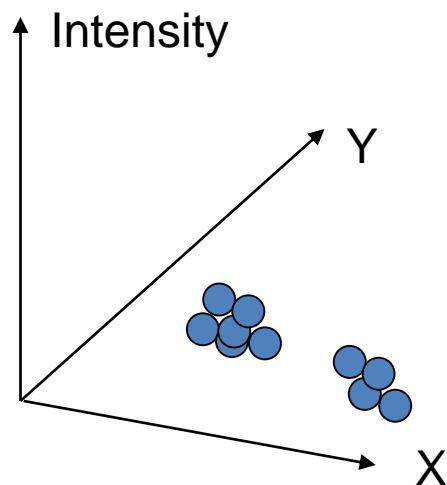
Clusters based on intensity
similarity don't have to be spatially
coherent.



Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on
intensity+position similarity



Both regions are black, but if we also include **position (x,y)**, then we could group the two into distinct segments; way to encode both similarity & proximity.

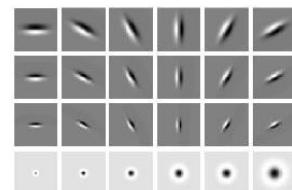
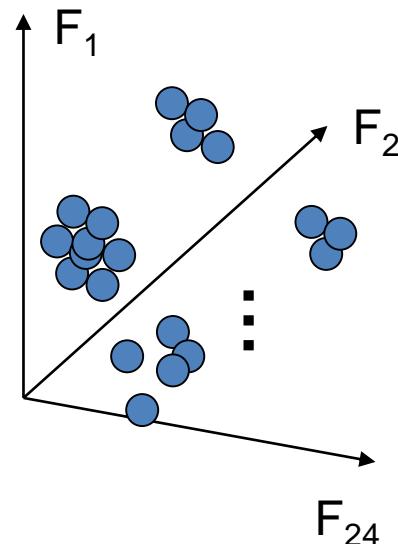
- Segmentation as clustering
- Color, Brightness, position alone are not enough to distinguish all regions...



Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

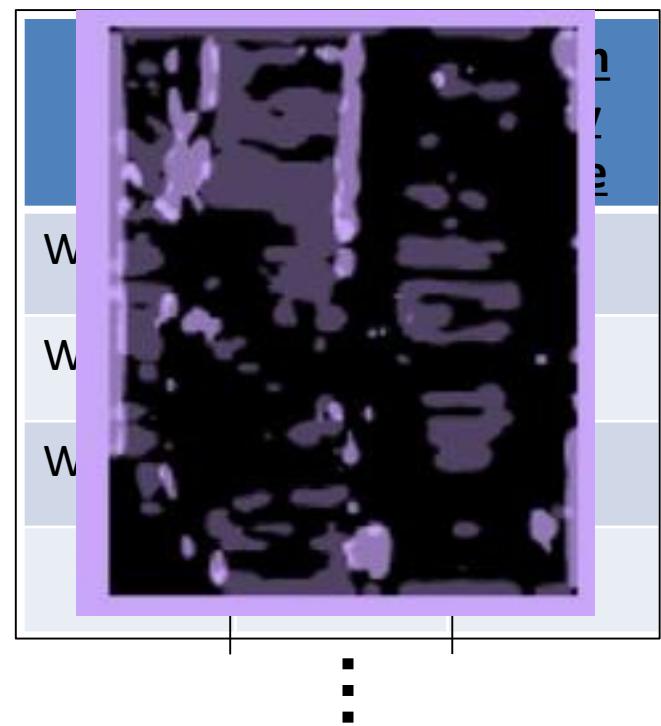
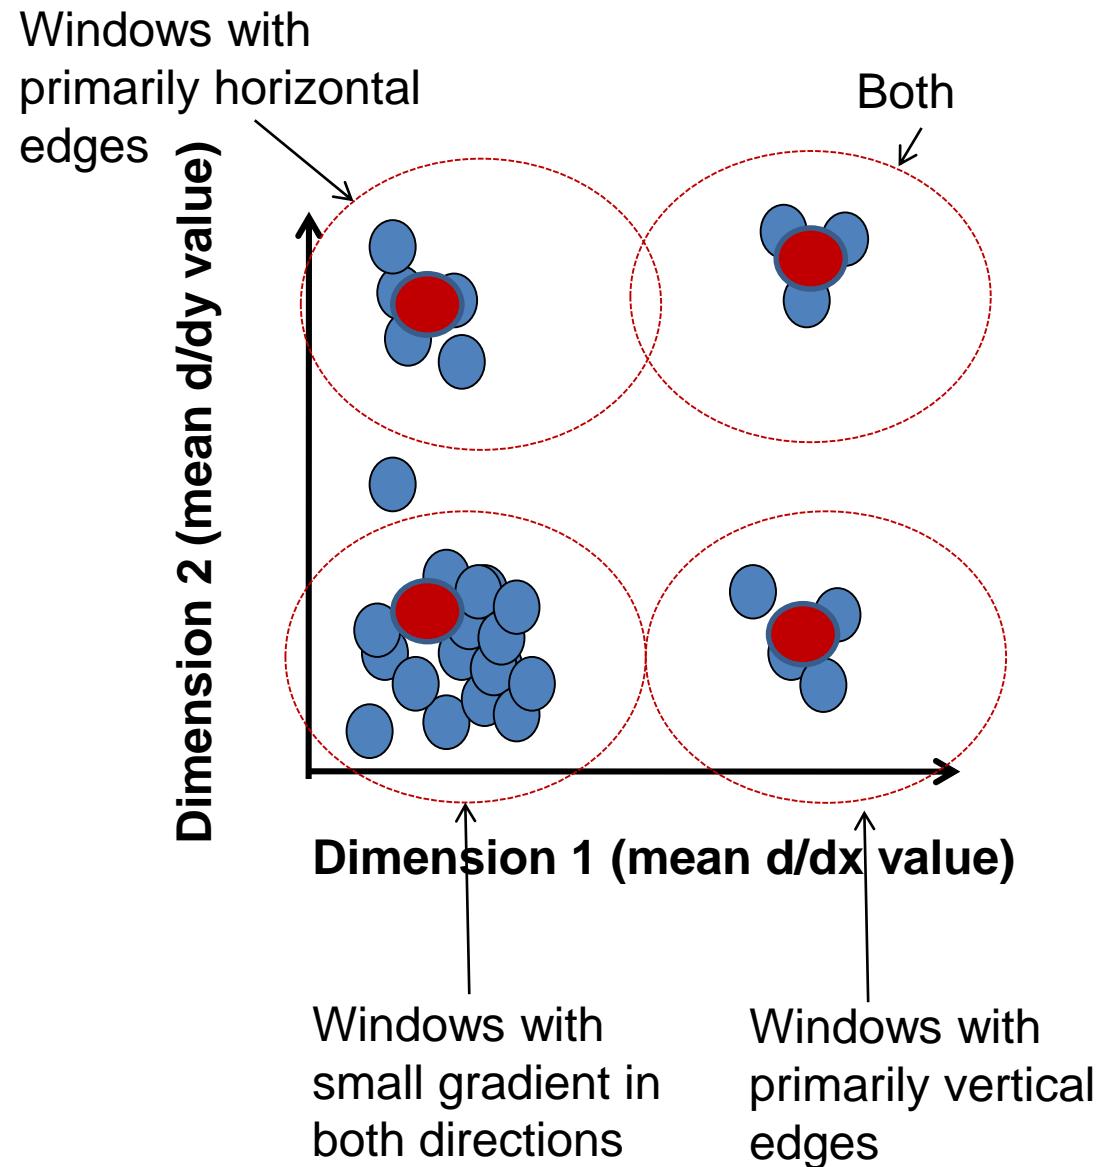
Grouping pixels based
on **texture** similarity



Filter bank
of 24 filters

Feature space: filter bank responses (e.g., 24-d)

Recall: texture representation example



statistics to summarize patterns in small windows

Segmentation with texture features

- Find “textons” by **clustering** vectors of filter bank outputs
- Describe texture in a window based on *texton histogram*

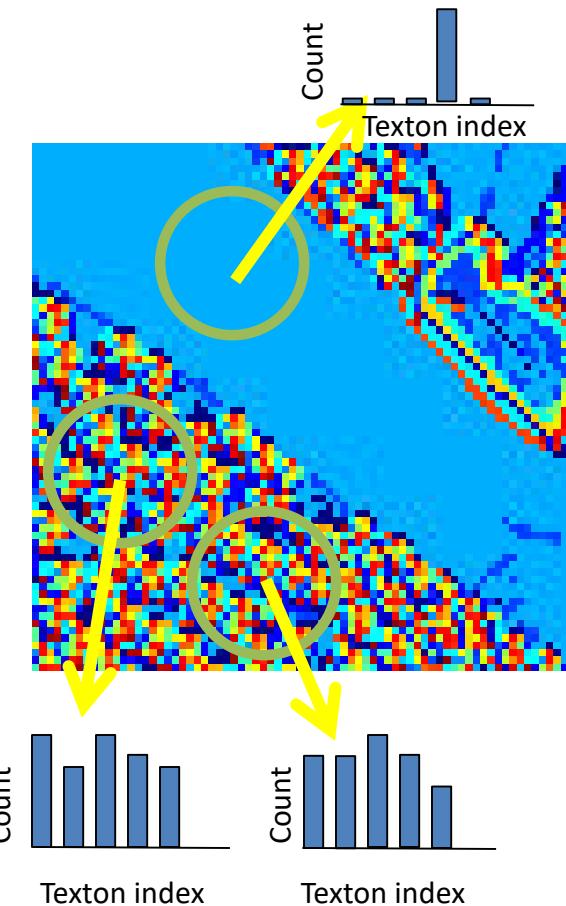
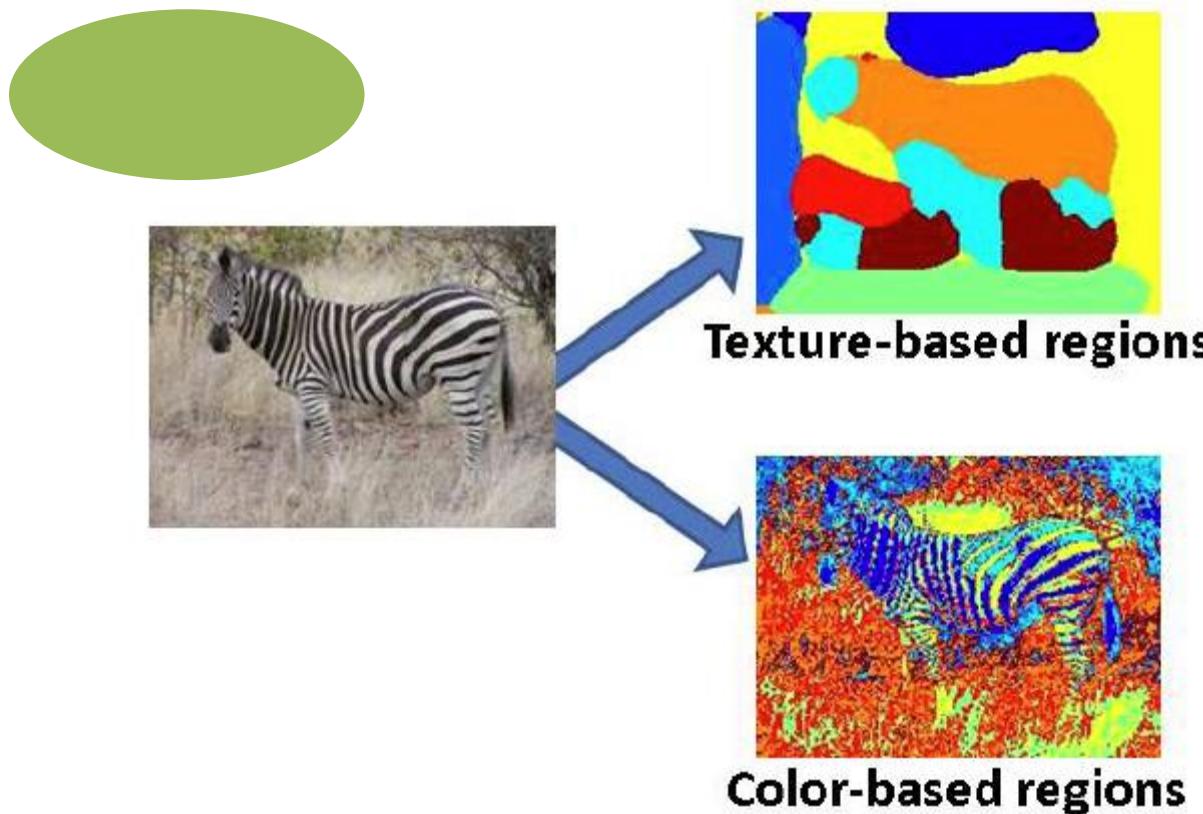


Image segmentation example



Pixel properties vs. neighborhood properties

query



query

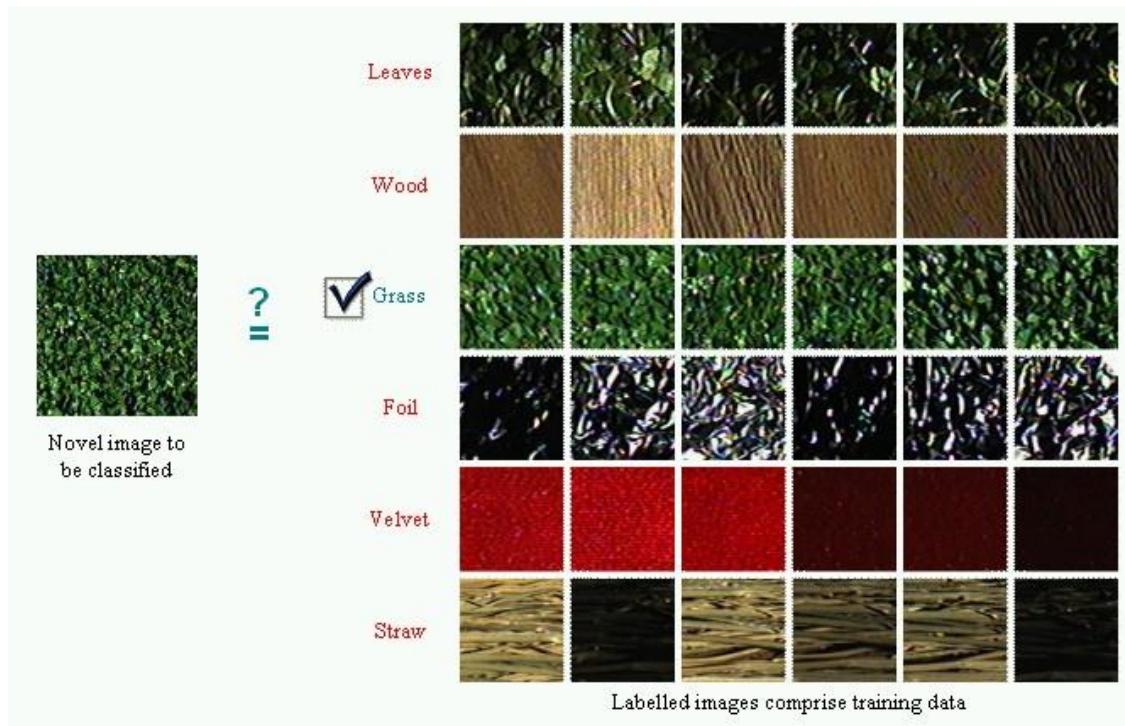


These look very similar in terms of their color distributions (histograms).

How would their *texture* distributions compare?

Material classification example

For an image of a single texture, we can classify it according to its global (image-wide) texton histogram.

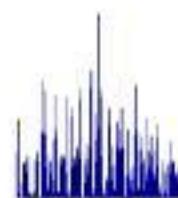


Material classification example

Nearest neighbor classification: label the input according to the nearest known example's label.



Novel Image

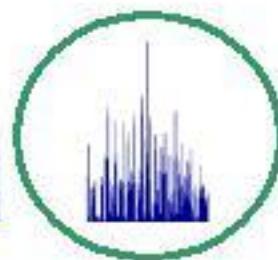


Model

$$\chi^2 =$$



Plastic



Grass

$$\chi^2(h_i, h_j) = \frac{1}{2} \sum_{k=1}^K \frac{[h_i(k) - h_j(k)]^2}{h_i(k) + h_j(k)}$$

Outline

- What are grouping problems in vision?
- Inspiration from human perception
 - Gestalt properties
- Bottom-up segmentation via clustering
 - Algorithms:
 - Mode finding and mean shift: k-means, mean-shift
 - Graph-based: normalized cuts
 - Features: color, texture, ...
 - Quantization for texture summaries

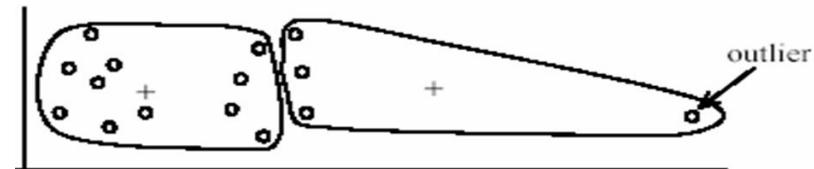
K-means: pros and cons

Pros

- Simple, fast to compute
- Converges to local minimum of within-cluster squared error

Cons/issues

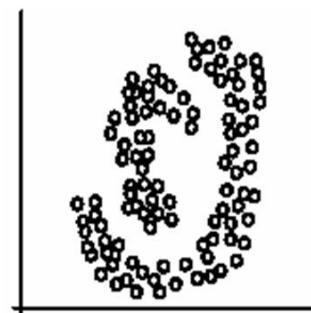
- Setting k?
- Sensitive to initial centers
- Sensitive to outliers
- Detects spherical clusters
- Assuming means can be computed



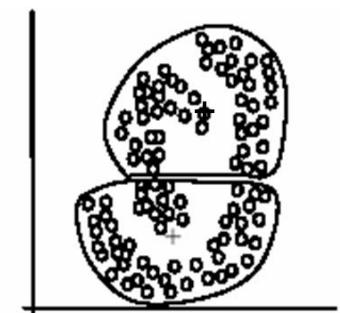
(A): Undesirable clusters



(B): Ideal clusters



(A): Two natural clusters



(B): k -means clusters

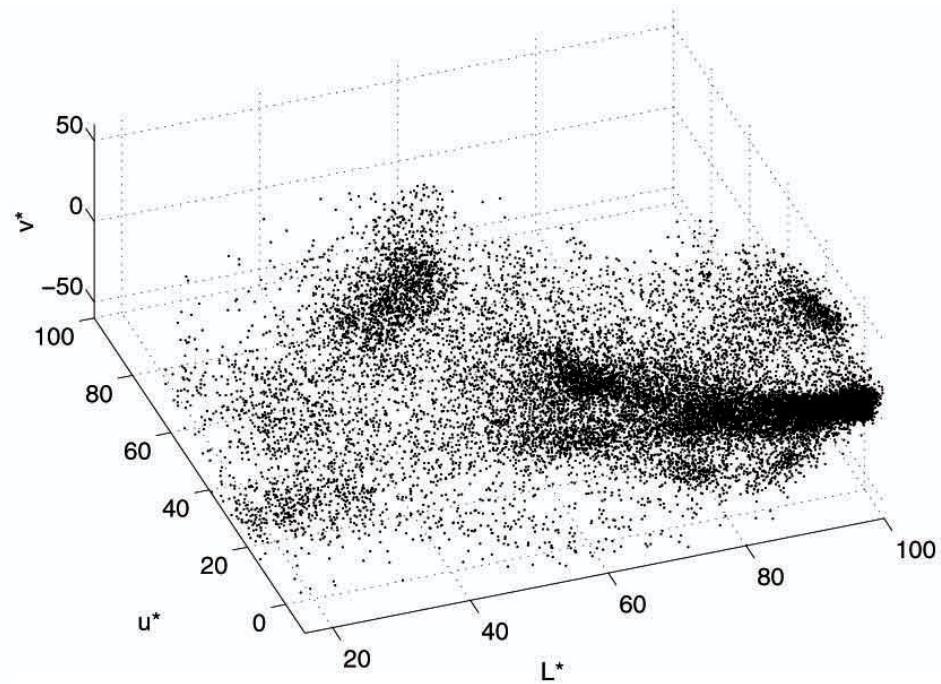
Mean shift algorithm

- The mean shift algorithm seeks *modes* or local maxima of density in the feature space

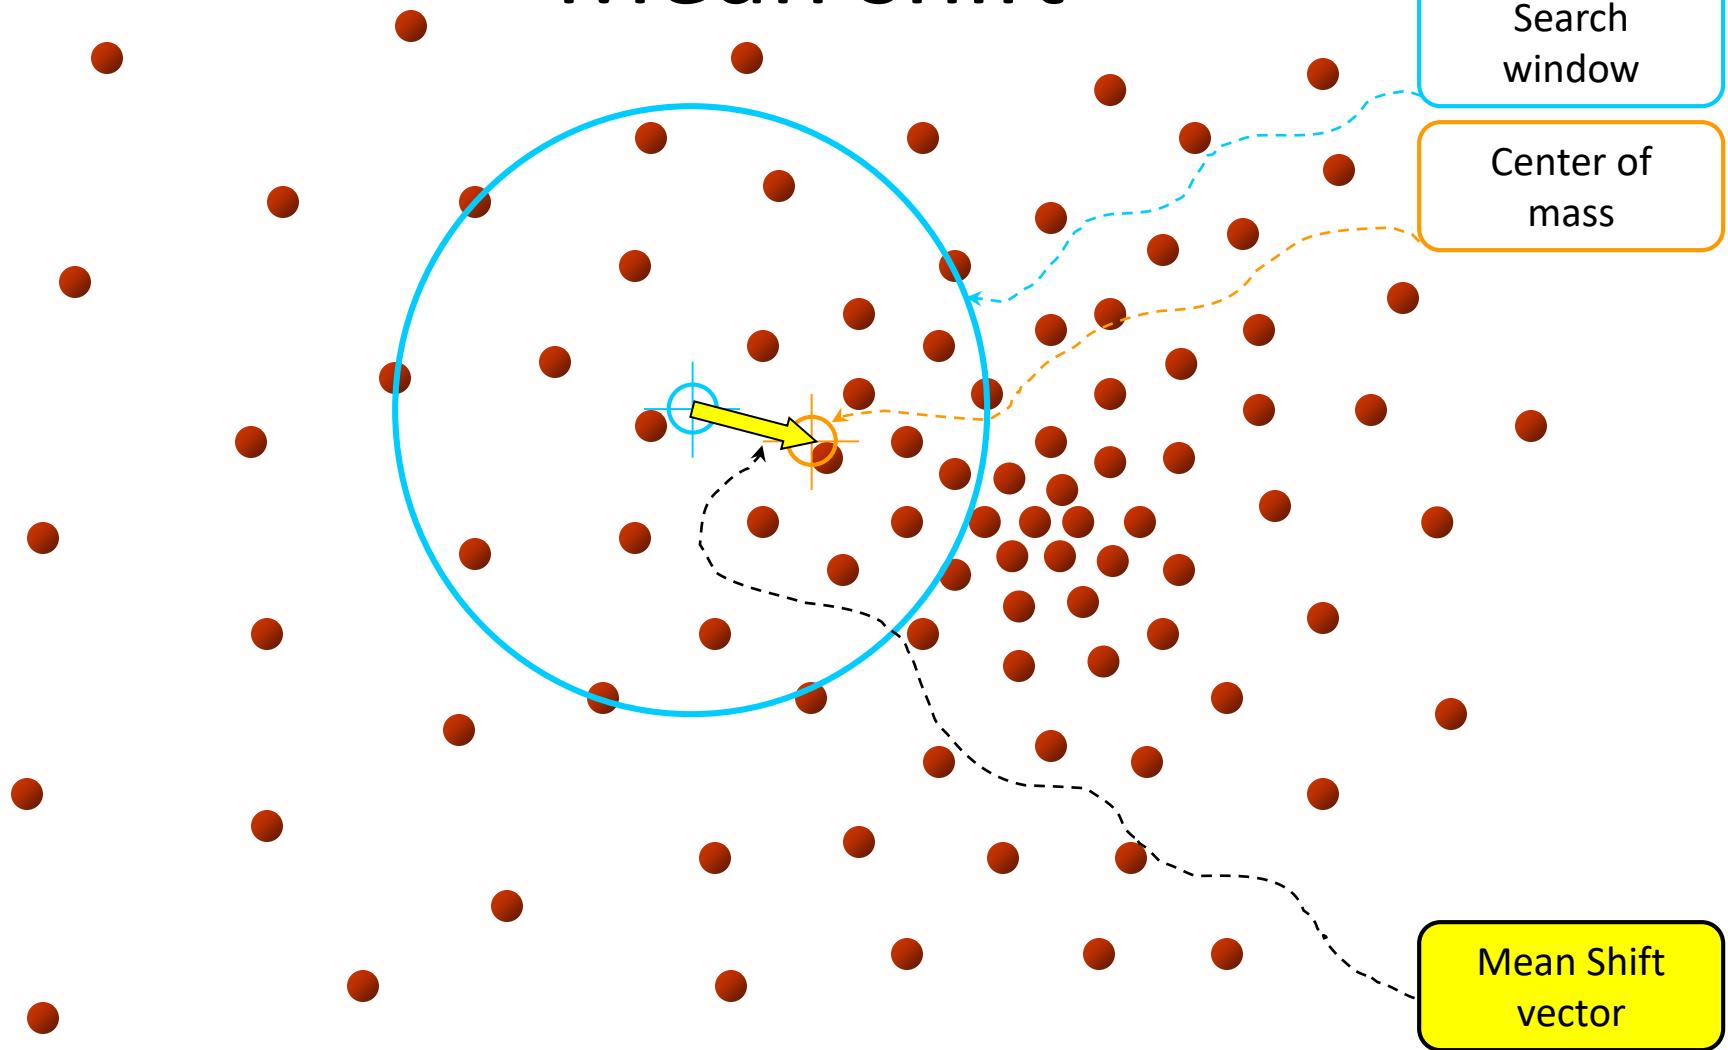
image



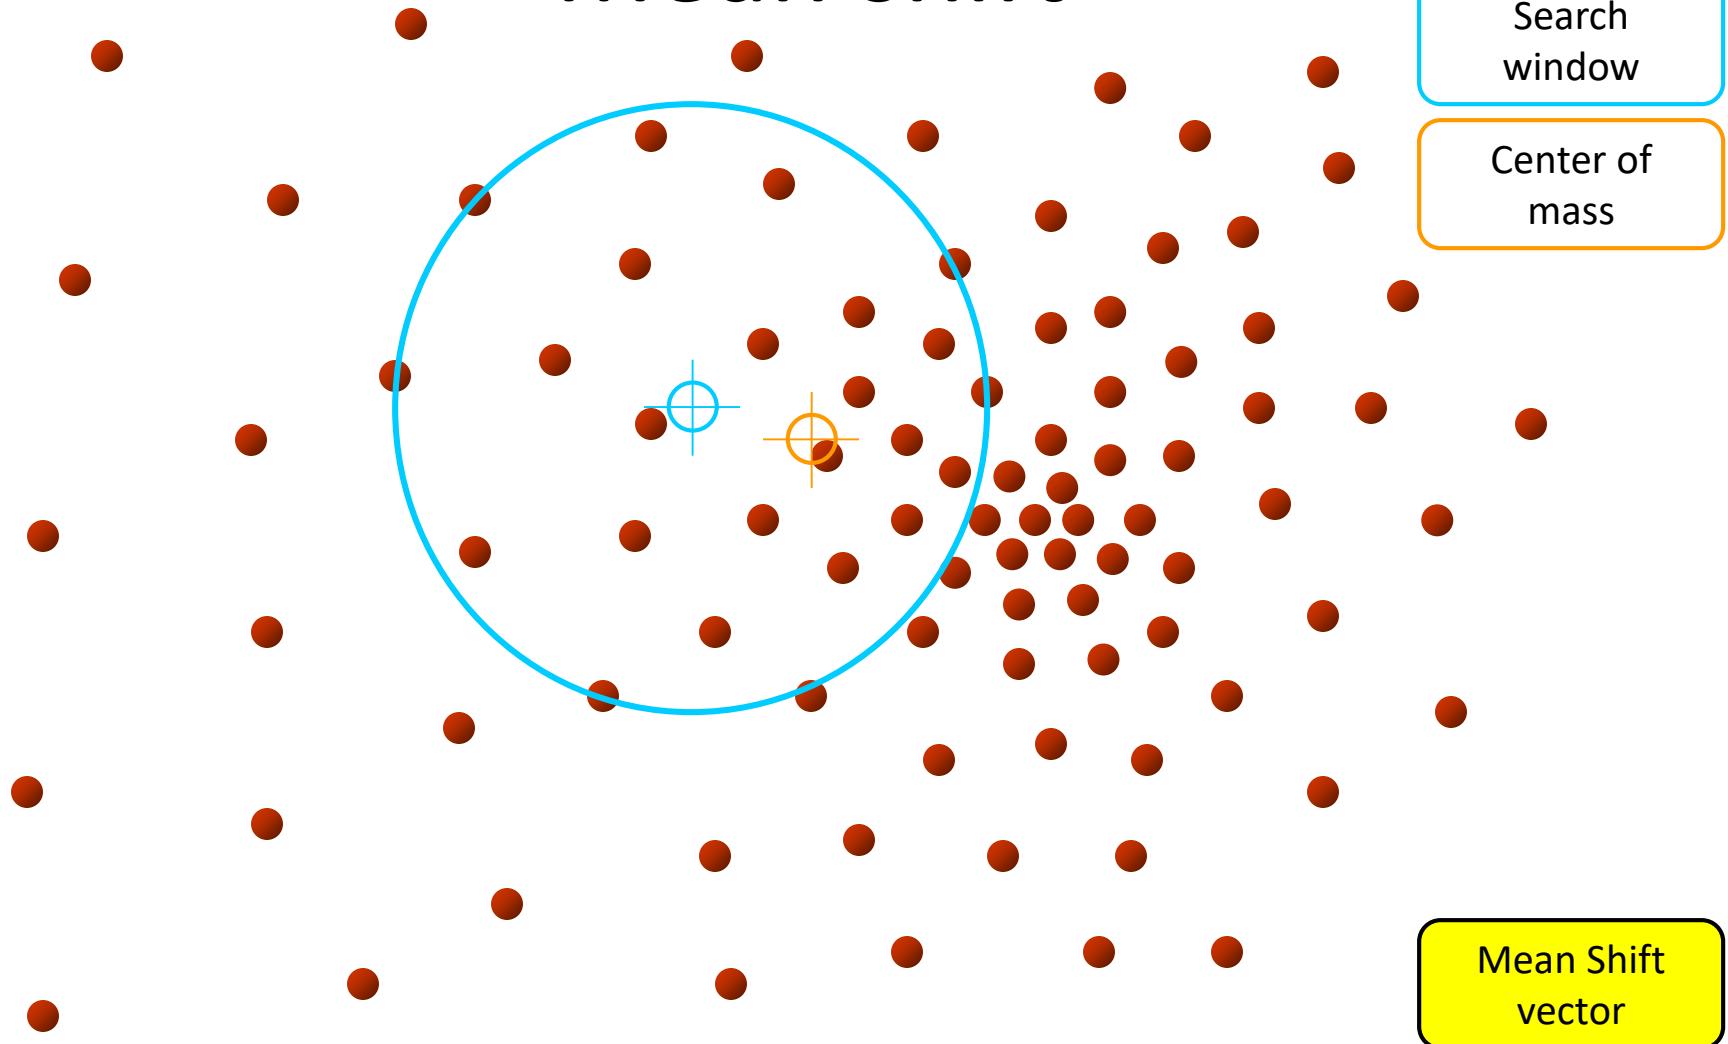
Feature space
($L^*u^*v^*$ color values)



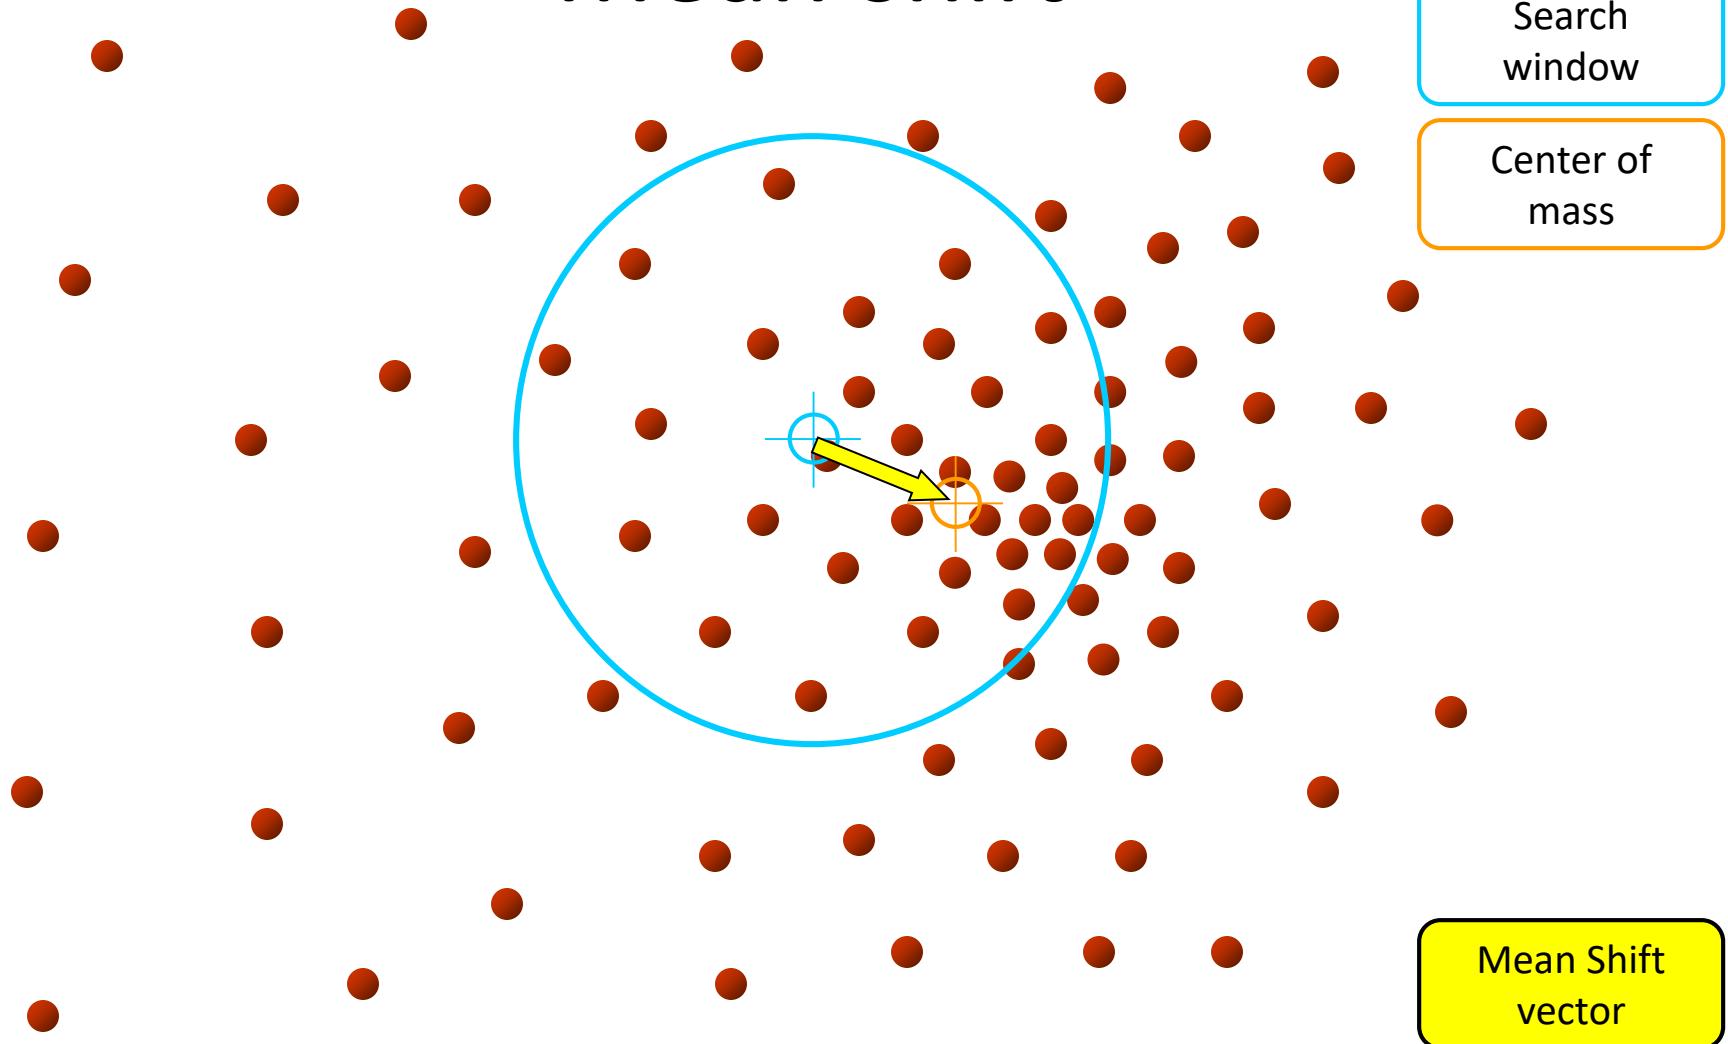
Mean shift



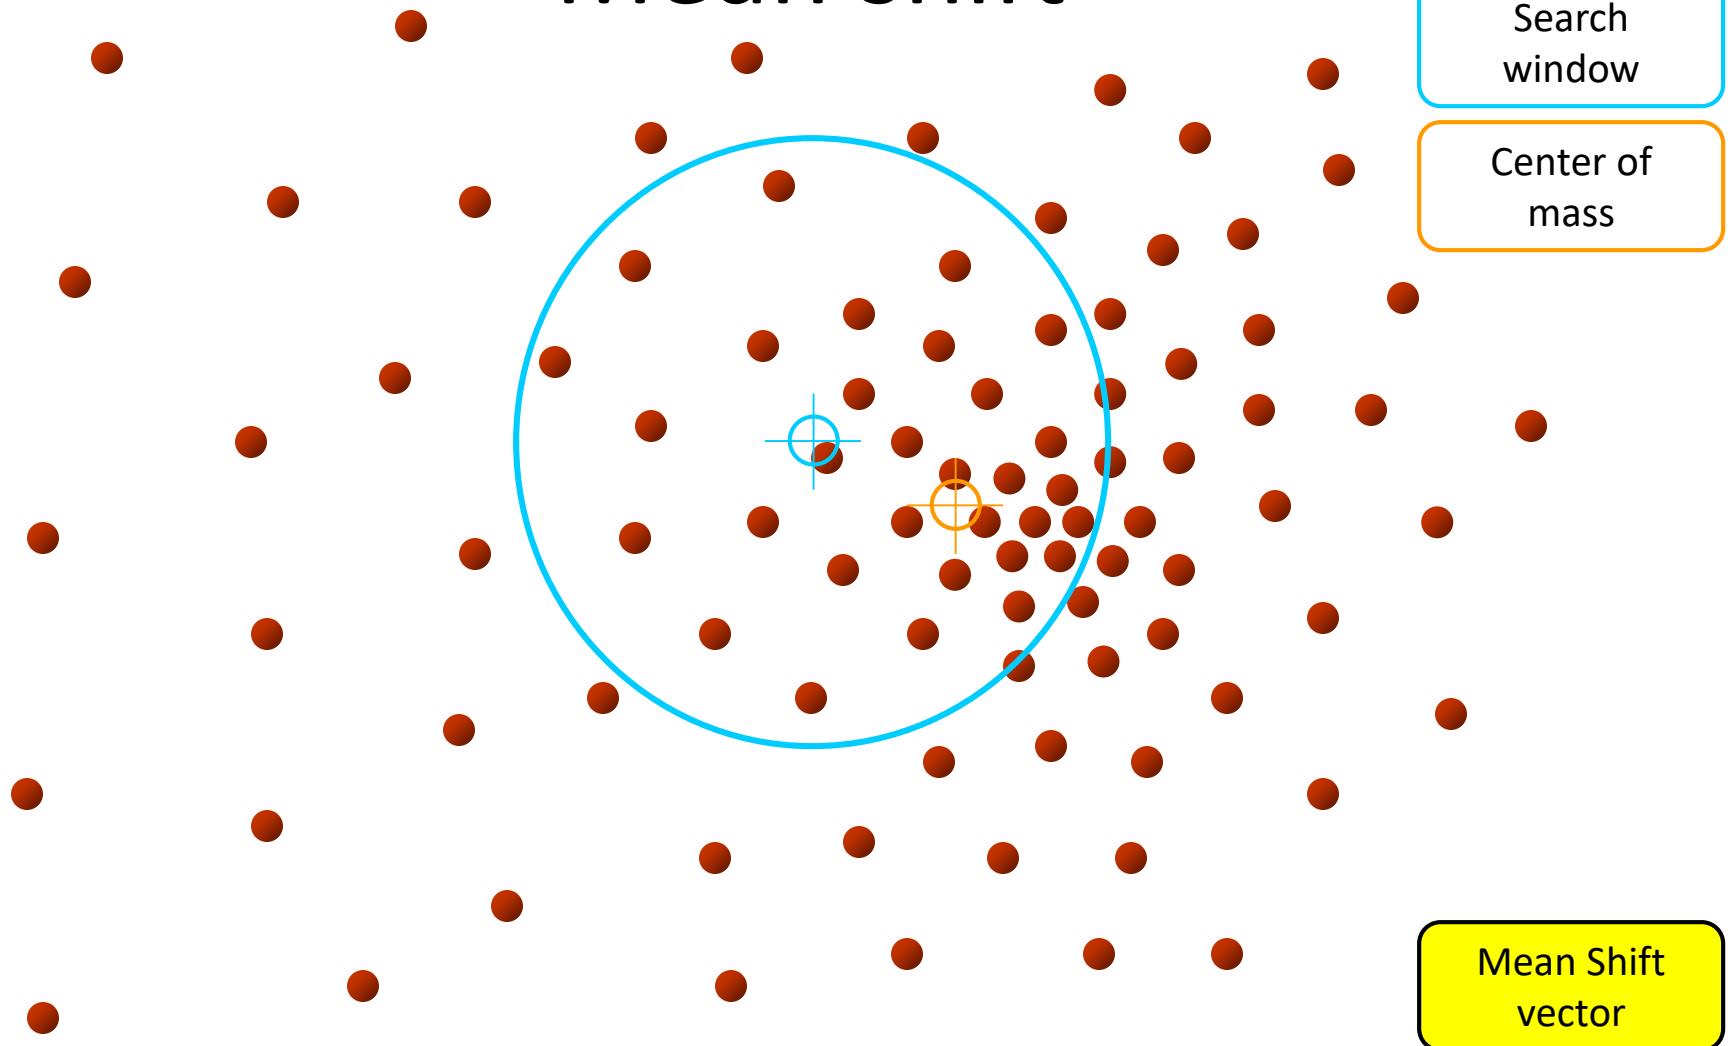
Mean shift



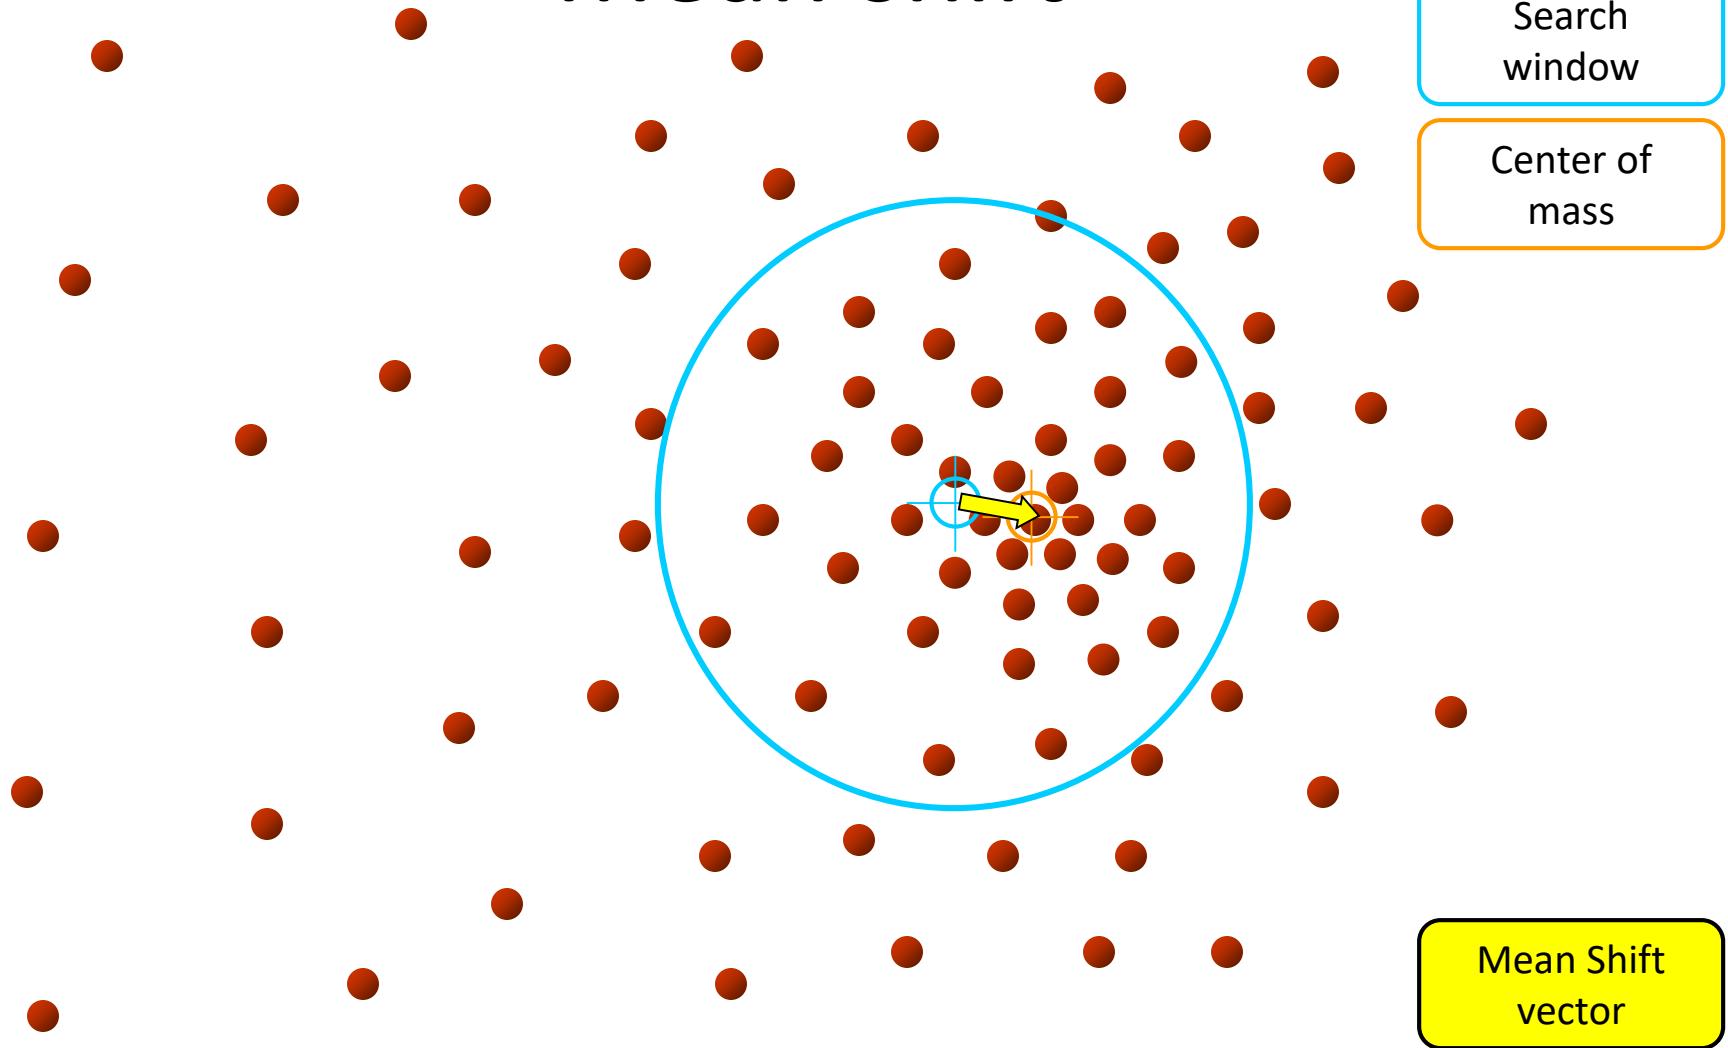
Mean shift



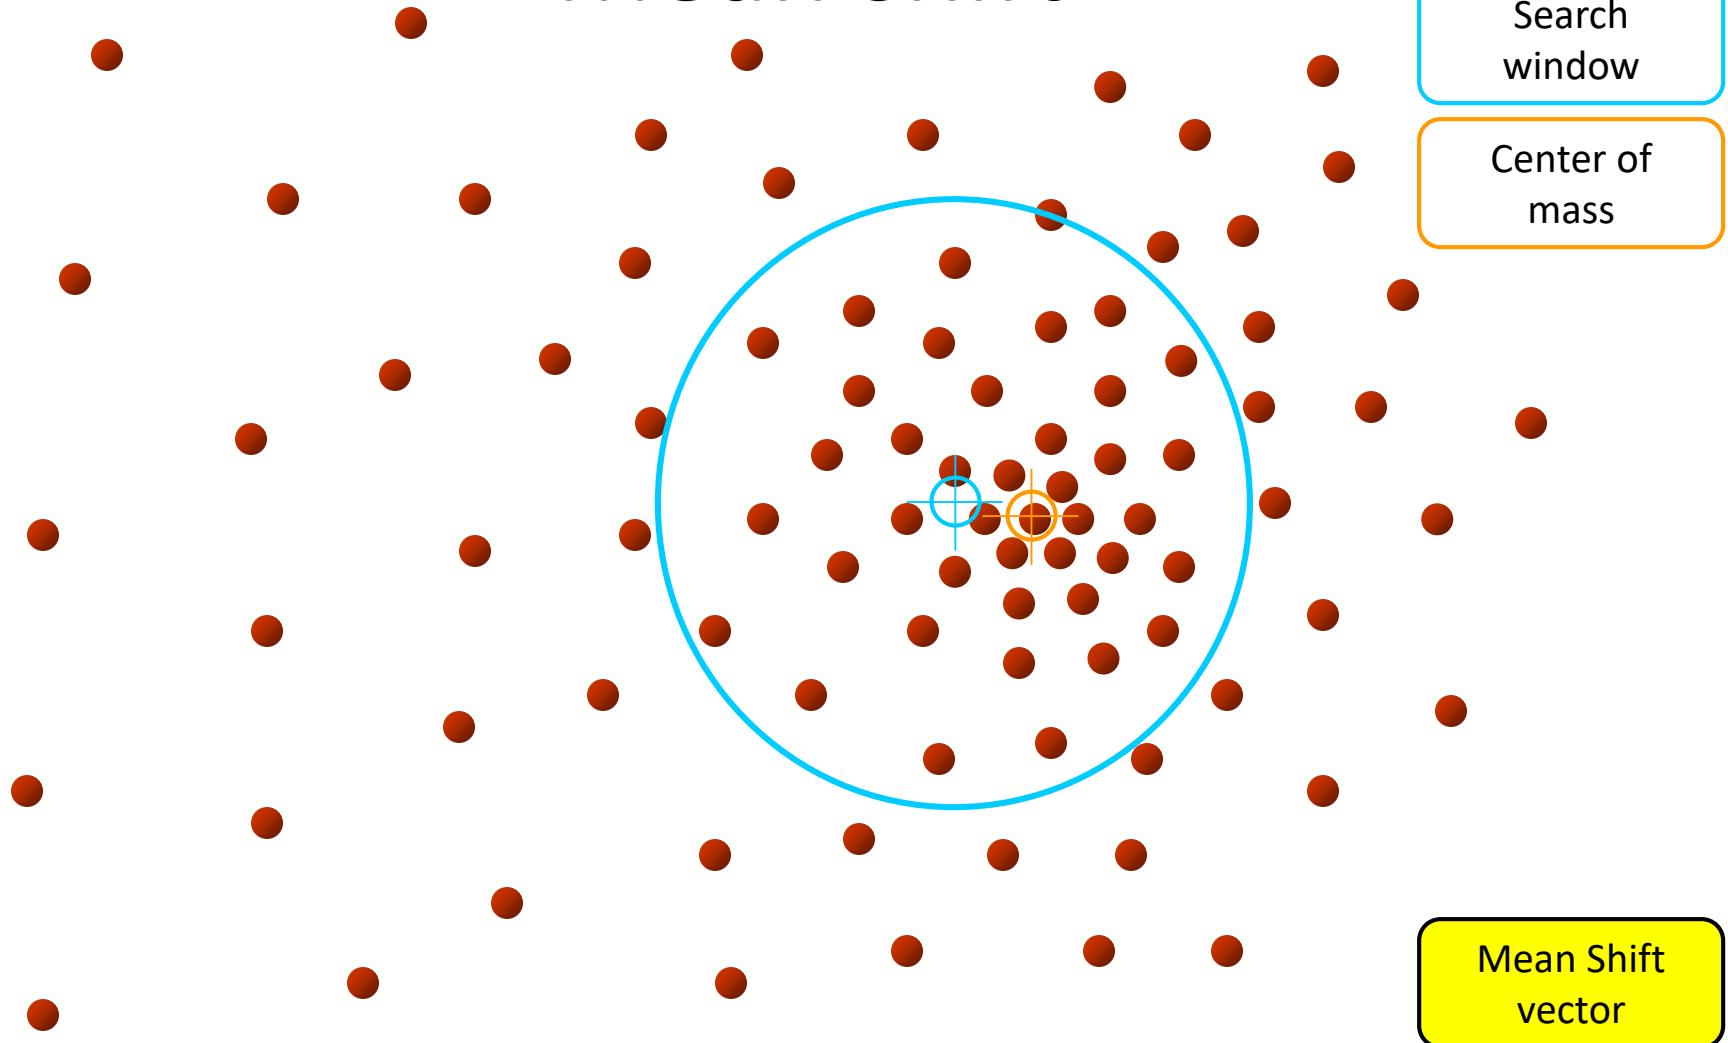
Mean shift



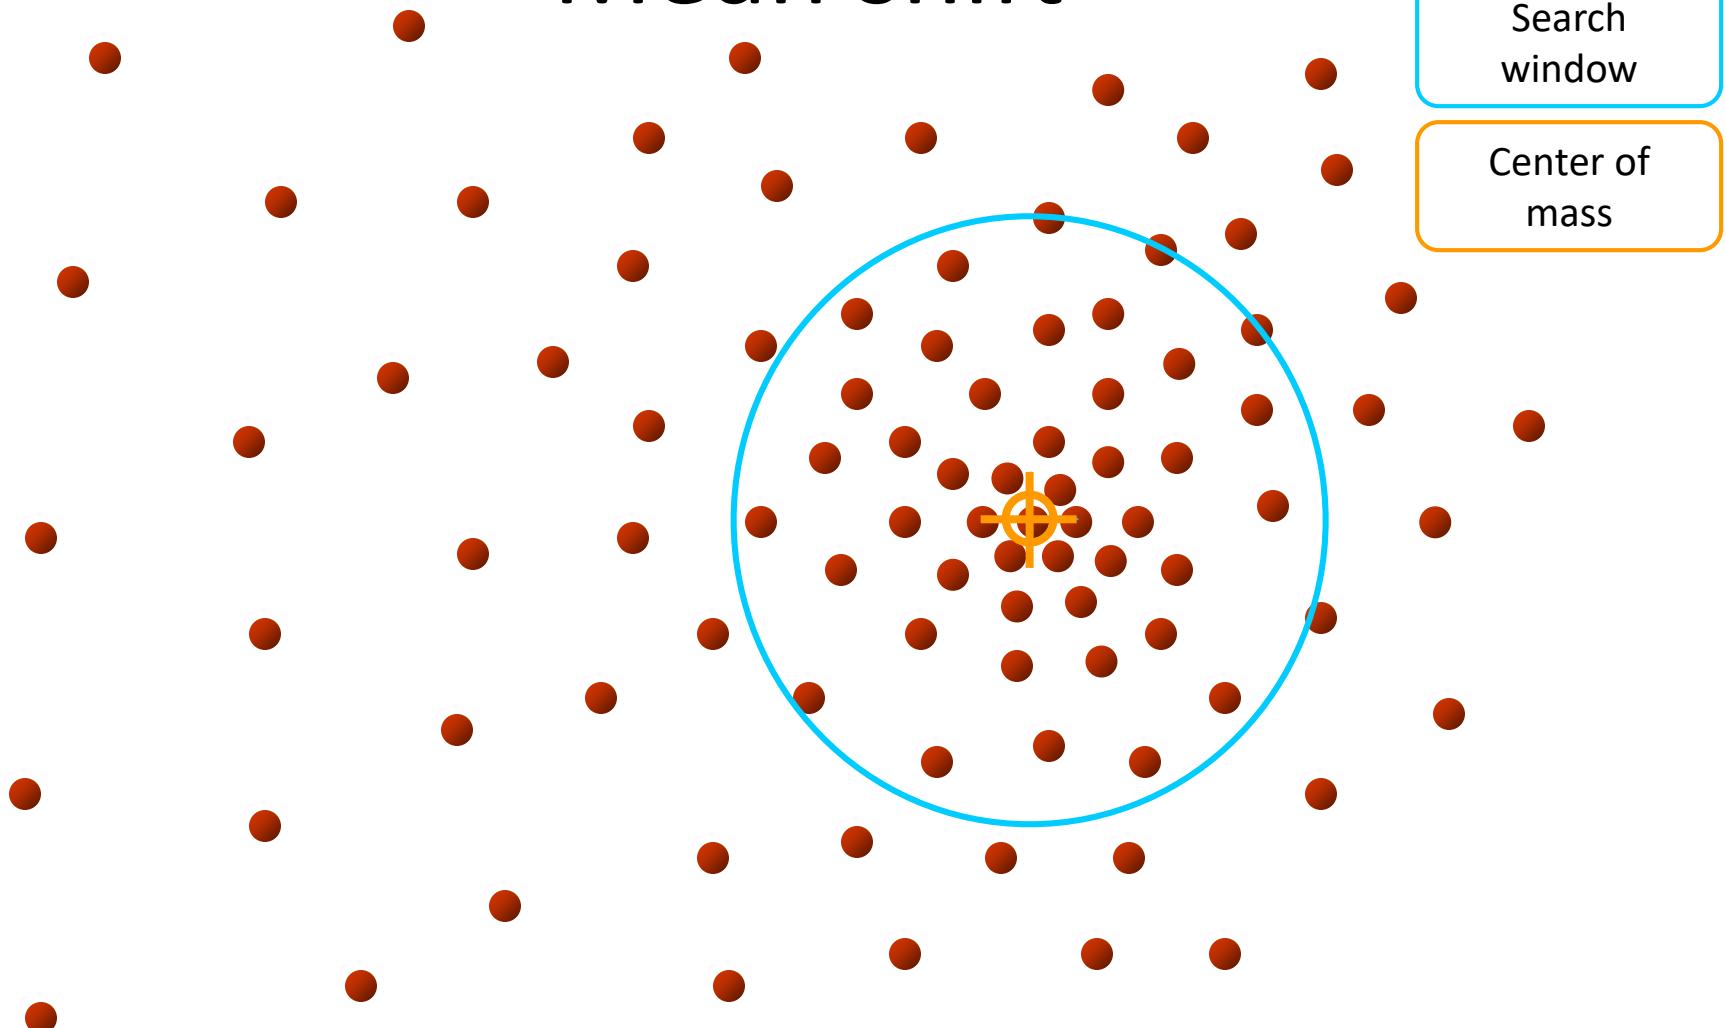
Mean shift



Mean shift



Mean shift

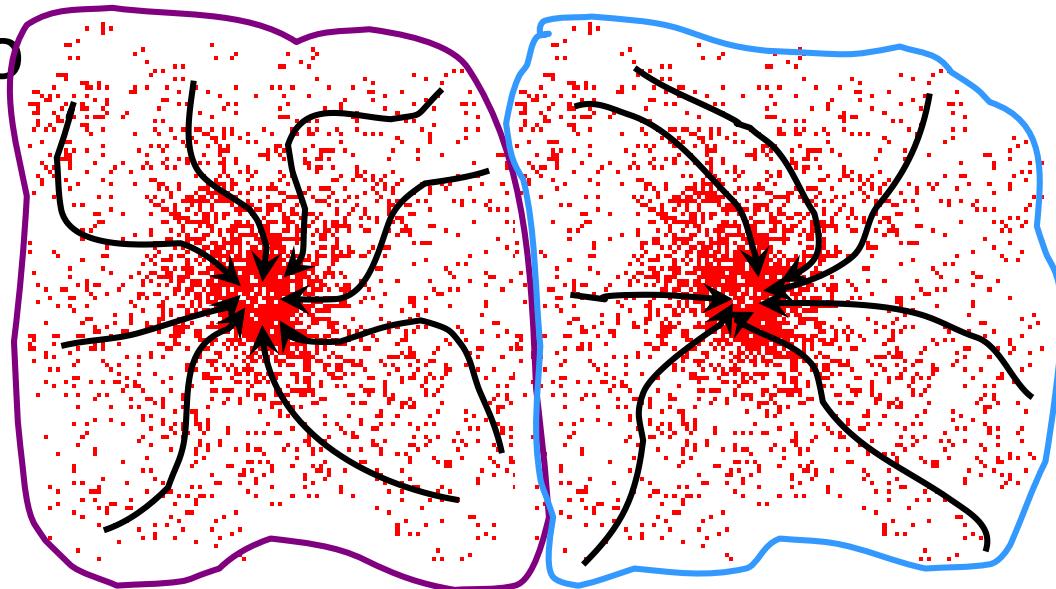


Search
window

Center of
mass

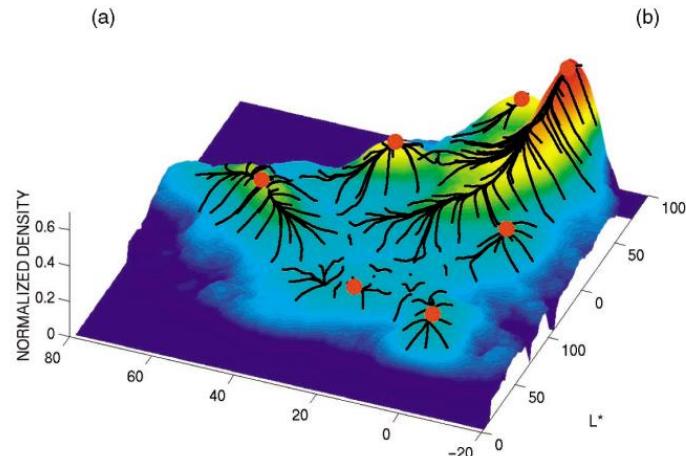
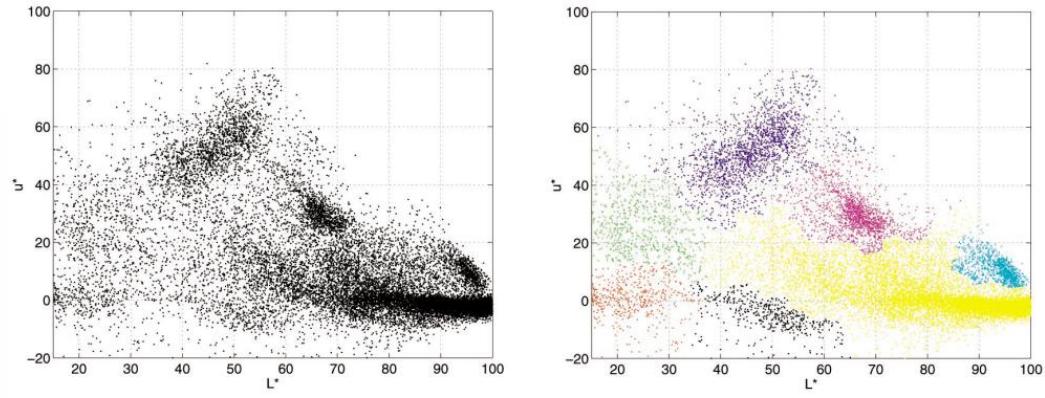
Mean shift clustering

- Cluster: all data points in the attraction basin of a mode
- Attraction basin: the region for which all trajectories converge to a mode



Mean shift clustering/segmentation

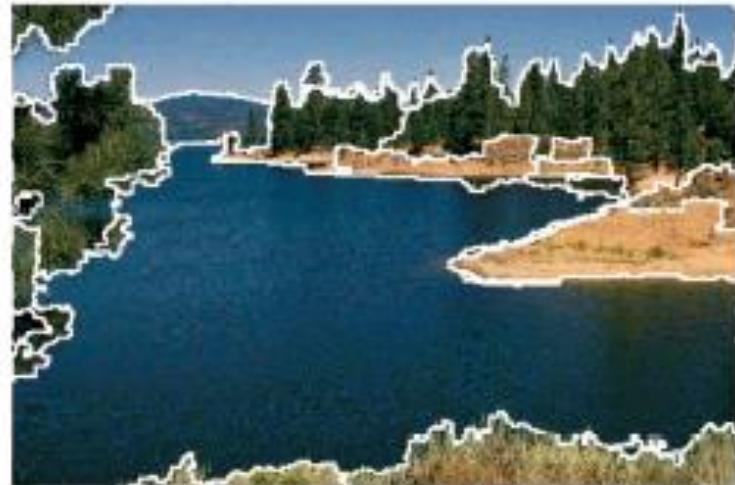
- Find features (color, gradients, texture, etc)
- Initialize windows at individual feature points
- Perform mean shift for each window until convergence
- Merge windows that end up near the same “peak” or mode



Mean shift segmentation results



Mean shift segmentation results



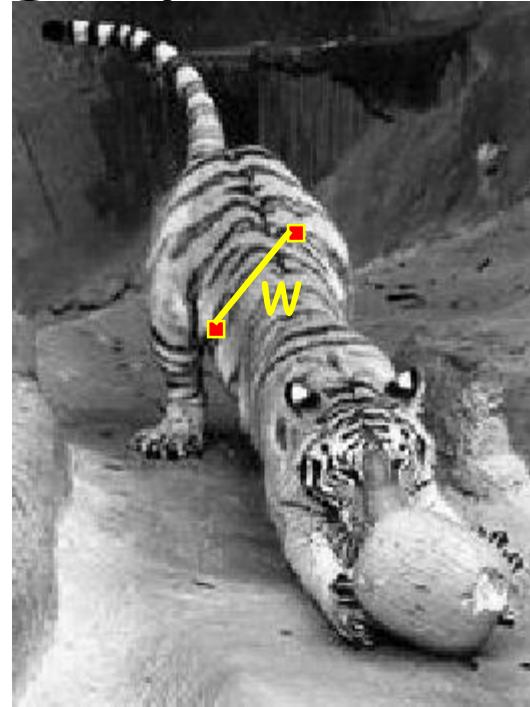
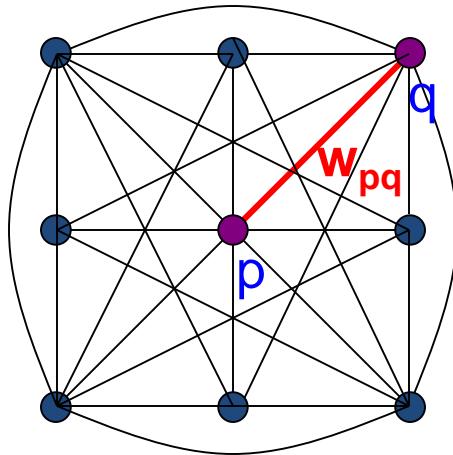
Mean shift

- Pros:
 - Does not assume shape on clusters
 - One parameter choice (window size)
 - Generic technique
 - Find multiple modes
- Cons:
 - Selection of window size
 - Does not scale well with dimension of feature space

Outline

- What are grouping problems in vision?
- Inspiration from human perception
 - Gestalt properties
- Bottom-up segmentation via clustering
 - Algorithms:
 - Mode finding and mean shift: k-means, mean-shift
 - Graph-based: normalized cuts
 - Features: color, texture, ...
 - Quantization for texture summaries

Images as graphs



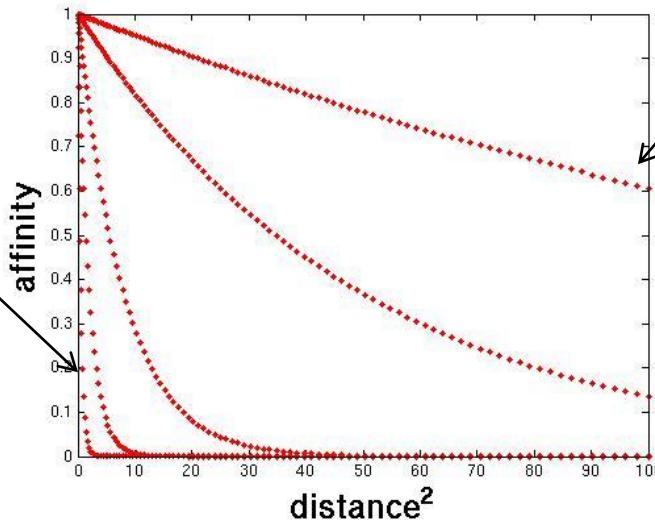
- *Fully-connected graph*
 - node (vertex) for every pixel
 - link between *every* pair of pixels, p, q
 - affinity weight w_{pq} for each link (edge)
 - w_{pq} measures *similarity*
 - similarity is *inversely proportional* to difference (in color and position...)

Measuring affinity

- One possibility:

$$aff(x, y) = \exp\left\{-\left(\frac{1}{2\sigma_d^2}\right)(\|x - y\|^2)\right\}$$

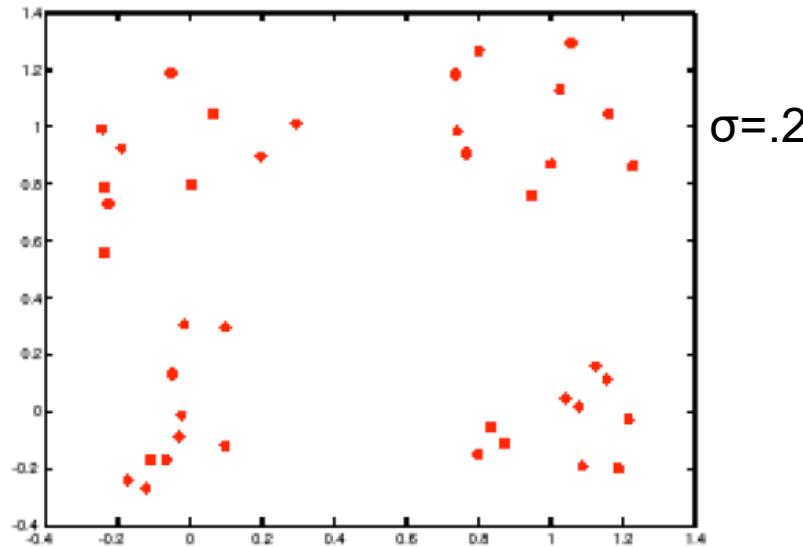
Small sigma:
group only
nearby points



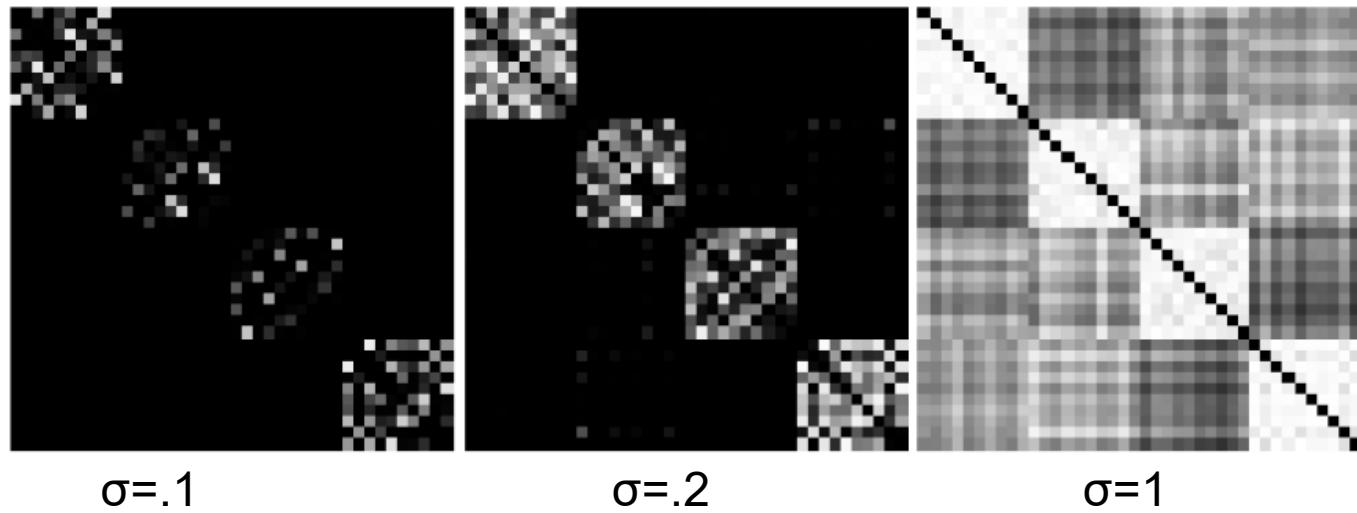
Large sigma:
group distant
points

Measuring affinity

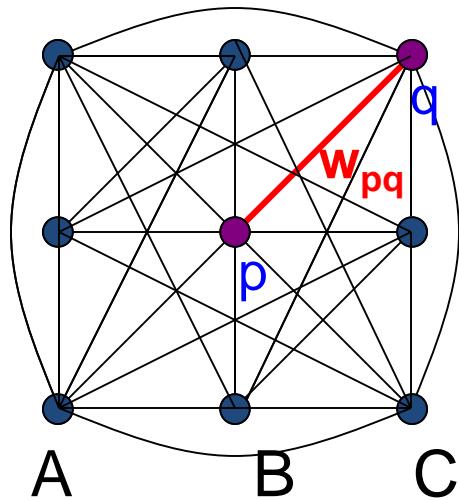
Data points



Affinity
matrices

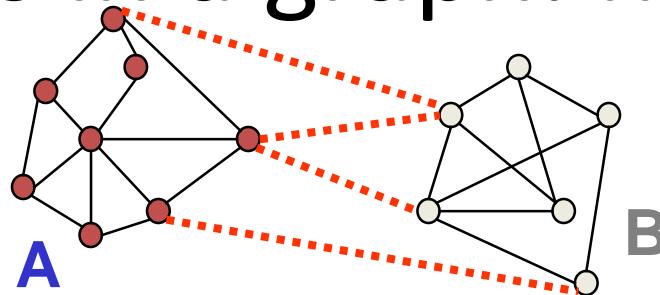


Segmentation by Graph Cuts



- Break Graph into Segments
 - Want to delete links that cross **between** segments
 - Easiest to break links that have low similarity (low weight)
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

Cuts in a graph: Min cut



- Link Cut
 - set of links whose removal makes a graph disconnected
 - cost of a cut:

$$cut(A, B) = \sum_{p \in A, q \in B} w_{p,q}$$

Find minimum cut

- gives you a segmentation
- fast algorithms exist for doing this

Minimum cut

- Problem with minimum cut:
Weight of cut proportional to number of edges in the cut;
tends to produce small, isolated components.

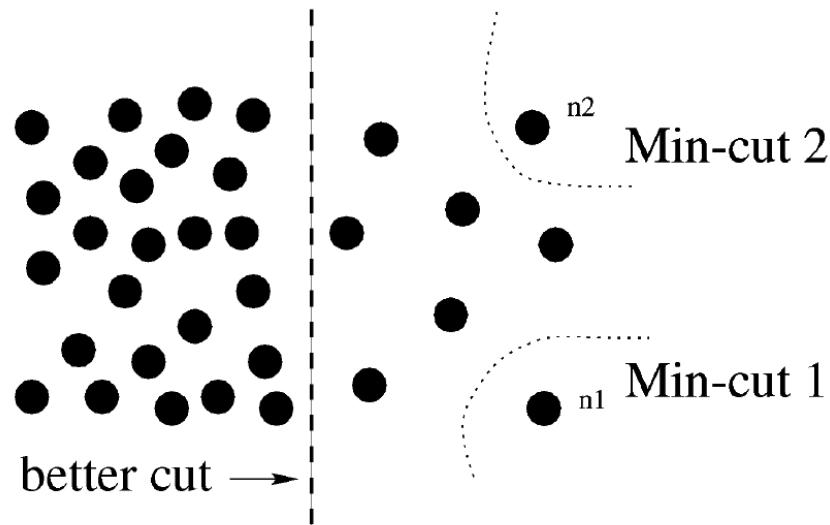
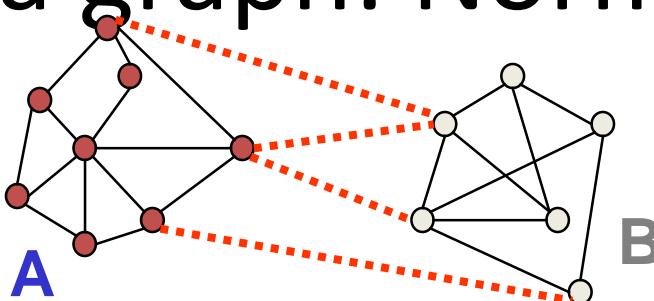


Fig. 1. A case where minimum cut gives a bad partition.

Cuts in a graph: Normalized cut



Normalized Cut

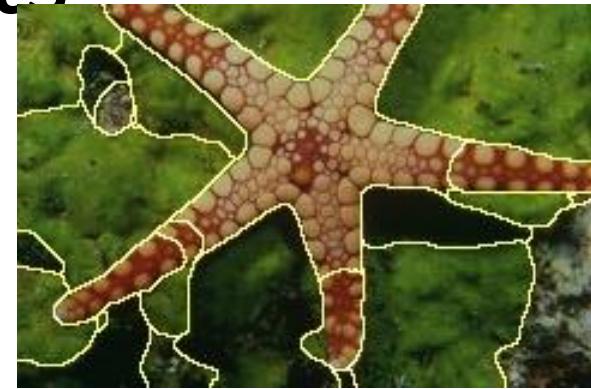
- fix bias of Min Cut by **normalizing** for size of segments:

$$Ncut(A, B) = \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)}$$

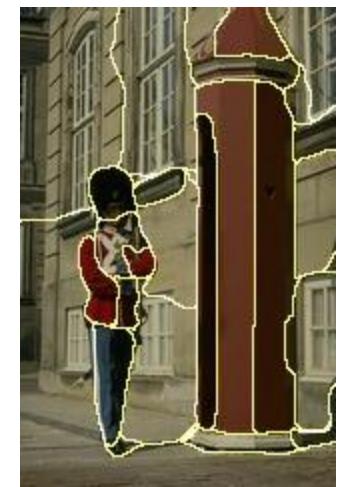
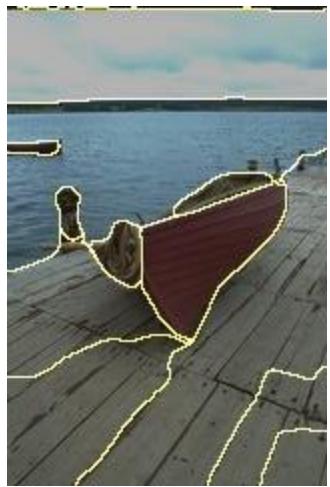
$assoc(A, V)$ = sum of weights of all edges that touch A

- Ncut value small when we get two clusters with many edges with high weights, and few edges of low weight between them
- Approximate solution for minimizing the Ncut value : generalized eigenvalue problem.

Example results



Results: Berkeley Segmentation Engine



<http://www.cs.berkeley.edu/~fowlkes/BSE/>

Normalized cuts: pros and cons

Pros:

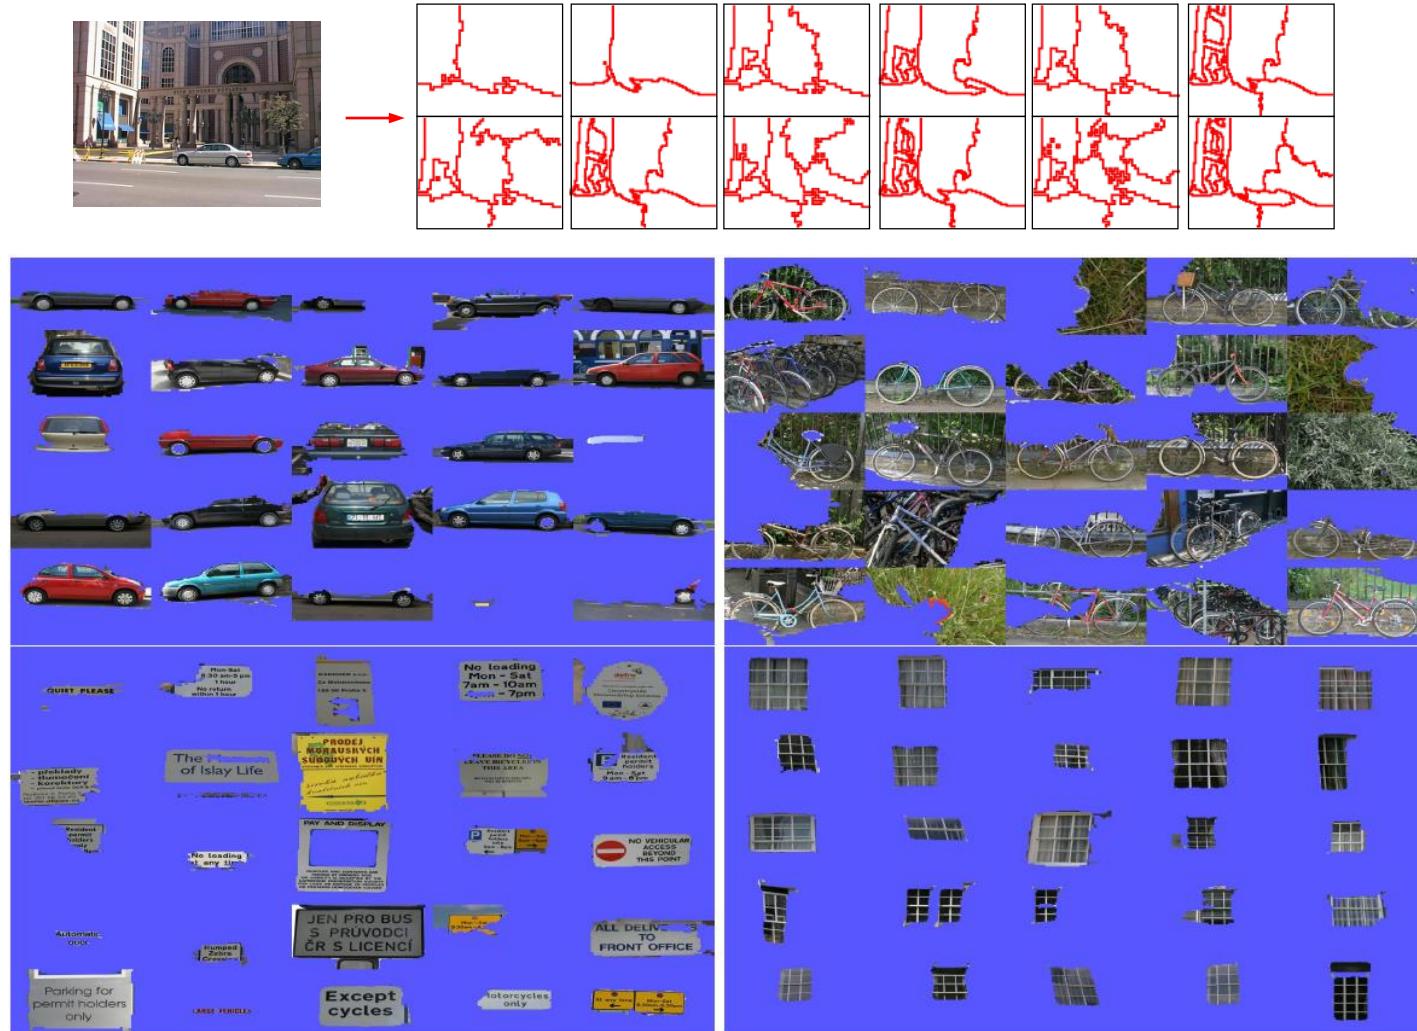
- Generic framework, flexible to choice of function that computes weights (“affinities”) between nodes
- Does not require model of the data distribution

Cons:

- Time complexity can be high
 - Dense, highly connected graphs → many affinity computations
 - Solving eigenvalue problem
- Preference for balanced partitions

Segments as primitives for recognition

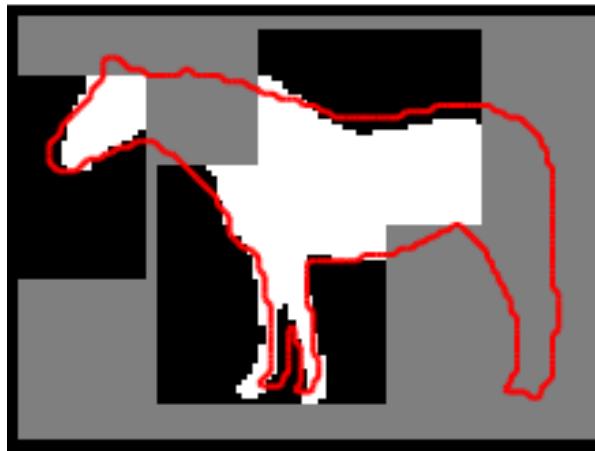
Multiple segmentations



- B. Russell et al., [“Using Multiple Segmentations to Discover Objects and their Extent in Image Collections,” CVPR 2006](#)

Slide credit: Lana Lazebnik

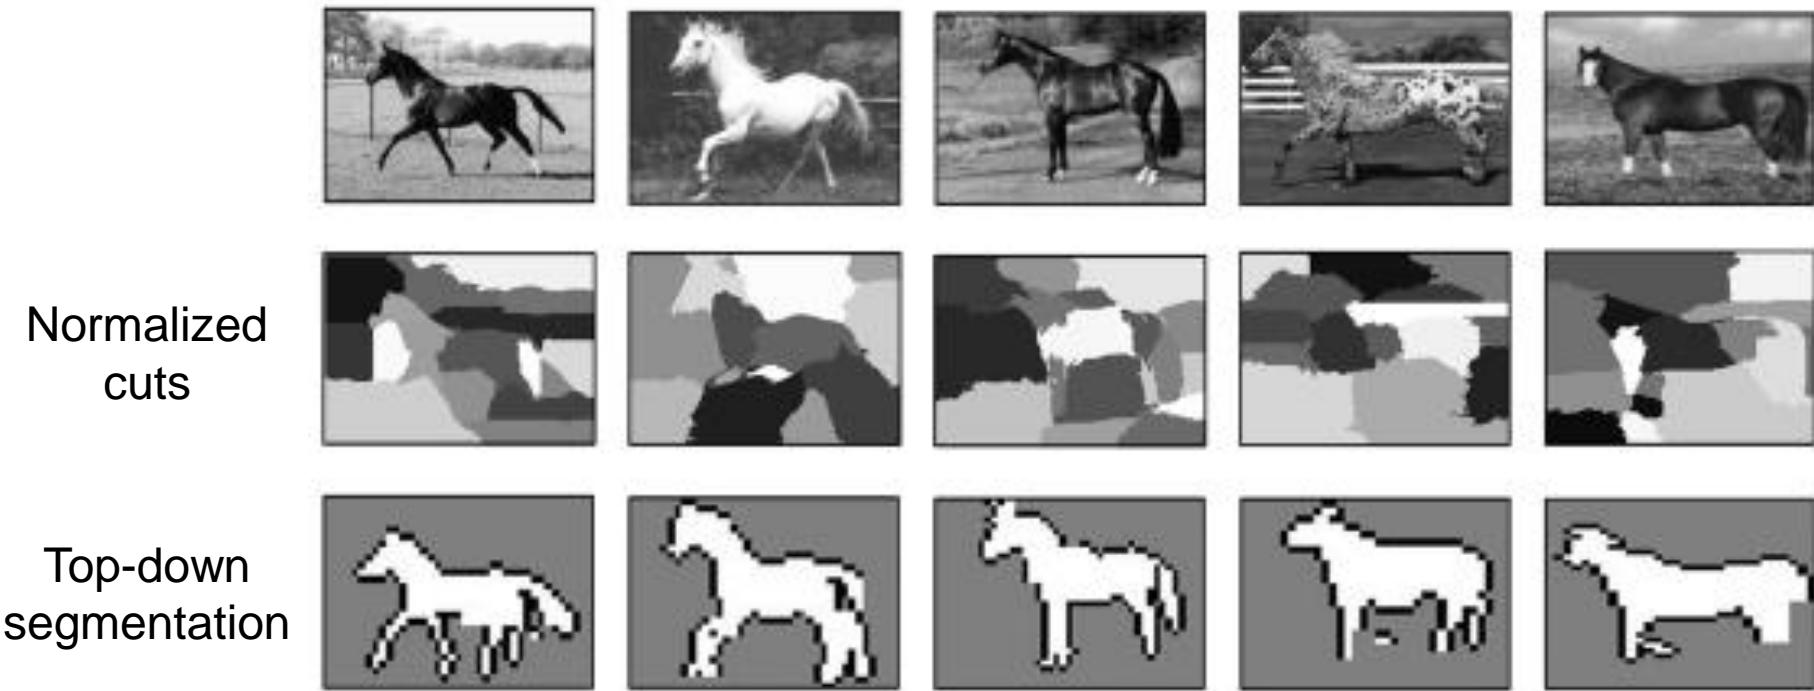
Top-down segmentation



E. Borenstein and S. Ullman, “[Class-specific, top-down segmentation,](#)” ECCV 2002

A. Levin and Y. Weiss, “[Learning to Combine Bottom-Up and Top-Down Segmentation,](#)” ECCV 2006.

Top-down segmentation



- E. Borenstein and S. Ullman, “[Class-specific, top-down segmentation](#),” ECCV 2002
- A. Levin and Y. Weiss, “[Learning to Combine Bottom-Up and Top-Down Segmentation](#),” ECCV 2006.

Motion segmentation



Input sequence

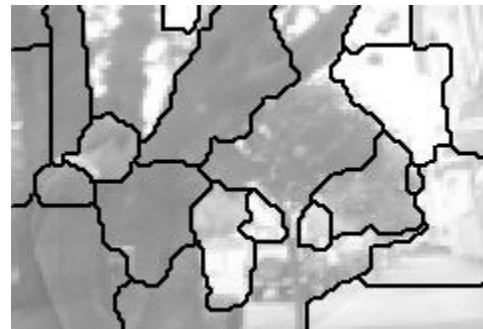


Image Segmentation



Motion Segmentation



Input sequence



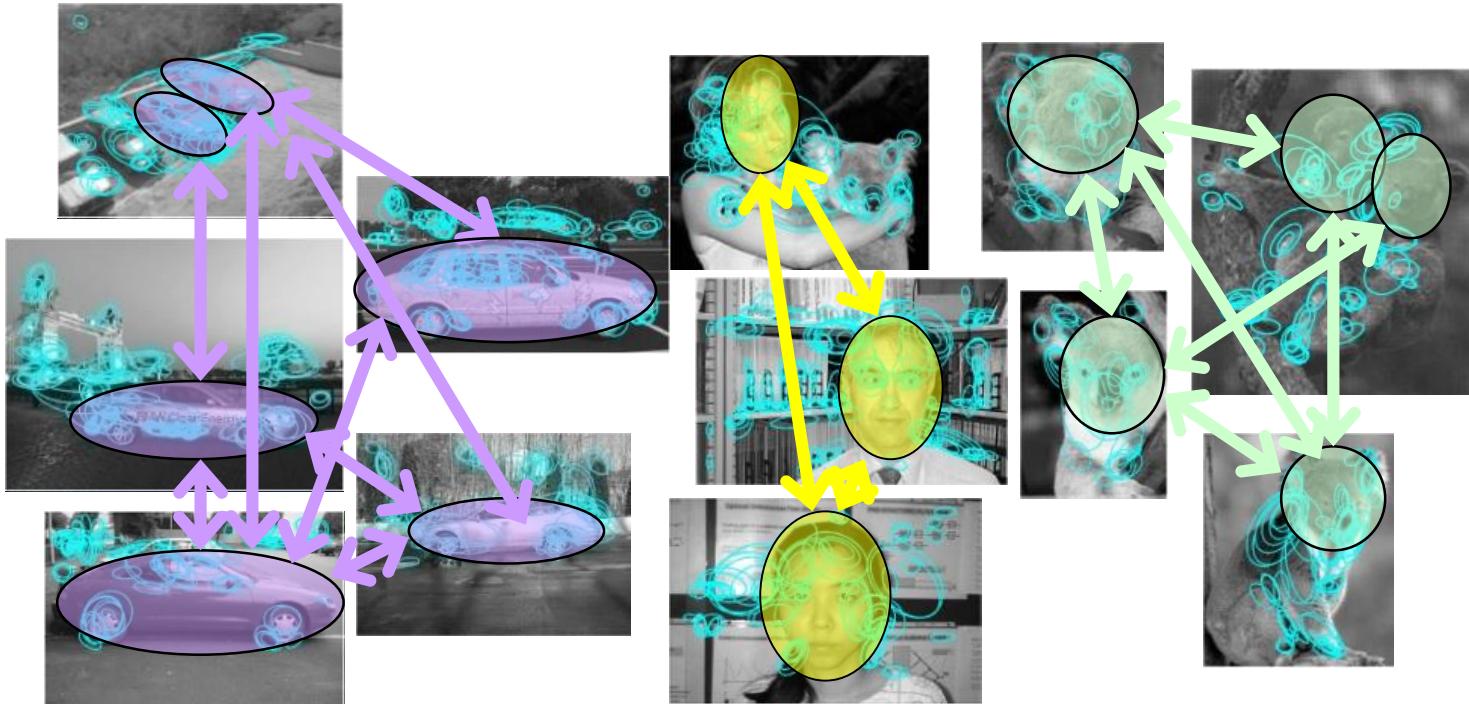
Image Segmentation



Motion Segmentation

A.Barbu, S.C. Zhu. Generalizing Swendsen-Wang to sampling arbitrary posterior probabilities,
IEEE Trans. PAMI, August 2005.
Kristen Grauman

Image grouping



K. Grauman & T. Darrell, Unsupervised Learning of Categories from Sets of Partially Matching Image Features, CVPR 2006.
Kristen Grauman