

Implementing Subprograms

BBM 301 – Programming Languages

The General Semantics of Calls and Returns

- The subprogram call and return operations of a language are together called its *subprogram linkage*
- General semantics of subprogram calls
 - Parameter passing methods
 - Stack-dynamic allocation of local variables
 - Save the execution status of calling program
 - Transfer of control and arrange for the return
 - If subprogram nesting is supported, access to nonlocal variables must be arranged

The General Semantics of Calls and Returns

- General semantics of subprogram returns:
 - Out mode and inout mode parameters must have their values returned
 - Deallocation of stack-dynamic locals
 - Restore the execution status
 - Return control to the caller

Implementing “Simple” Subprograms

- Simple Programs:
 - Subprograms cannot be nested
 - All local variables are static
 - Ex: Early versions of Fortran
- Call Semantics:
 - Save the execution status of the caller (caller/callee)
 - Pass the parameters (caller)
 - Pass the return address to the subprogram (caller)
 - Transfer control to the subprogram (caller)

Implementing “Simple” Subprograms

- Return Semantics:
 - If pass-by-value-result or out mode parameters are used, move the current values of those parameters to their corresponding actual parameters (callee)
 - If it is a function, move the functional value to a place the caller can get it (callee)
 - Restore the execution status of the caller (caller/callee)
 - Transfer control back to the caller (callee)

Storage required for call/return actions

- Required storage:
 - Status information about the caller
 - Parameters
 - return address
 - return value for functions
- These along with local variables and the subprogram code, form the complete collection of information a subprogram needs to execute and then return to the caller

Implementing “Simple” Subprograms: Parts

- Two separate parts: the actual code (constant) and the non-code part (local variables and data that can change)
- The format, or layout, of the non-code part of an executing subprogram is called an *activation record*
- The form of an activation record is static.
- An *activation record instance (ARI)* is a concrete example of an activation record (the collection of data for a particular subprogram activation)

An Activation Record for “Simple” Subprograms

Because languages with simple subprograms do not support recursion, there can be only one active version of a given subprogram at a time

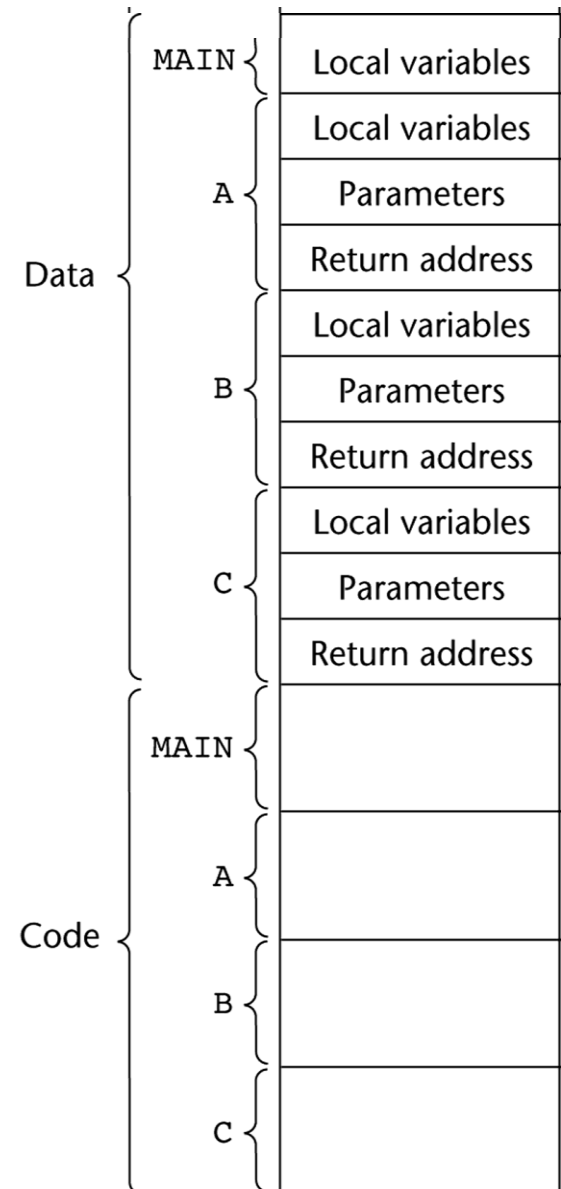
Local variables
Parameters
Return address

Therefore, there can be only a single instance of the activation record for a subprogram

Since activation record instance of a simple subprogram has fixed size it can be statically allocated. It could be also attached to the code part

Code and Activation Records of a Program with “Simple” Subprograms

- Note that code could be attached to ARIs
- Also, the four program units could be compiled at different times
- Linker put the compiled parts together when it is called for the main program



Implementing Subprograms with Stack-Dynamic Local Variables

- Advantage of Stack-Dynamic Local Variables
 - Support for recursion
- More complex activation record
 - The compiler must generate code to cause implicit allocation and deallocation of local variables
 - Recursion must be supported (adds the possibility of multiple simultaneous activations of a subprogram)
 - Each activation needs its own activation record and the number of activations is limited only by the memory size of the machine.

Typical Activation Record for a Language with Stack-Dynamic Local Variables

Local variables
Parameters
Dynamic link
Return address

↑
Stack top

Return address: pointer to the code segment of the caller and an offset address in that code segment of the instruction following the call

Dynamic link: pointer to the base of the activation record instance of the caller

In static scoped languages this link is used to provide traceback information when a run-time error occurs.

In dynamic-scoped languages, the dynamic link is used to access non-local variables

Because the return address, dynamic link, and parameters are placed in the activation record instance by the caller, these entries must appear first.

Implementing Subprograms with Stack-Dynamic Local Variables

- In most languages,
 - The format of the activation record is known at compile time
 - In many cases, the size is also known, because all the local data are of fixed size.
 - In some languages, like Ada, the size of a local array can depend on the value of an actual parameter, so in those cases the format is static, but the size can be dynamic.

Implementing Subprograms with Stack-Dynamic Local Variables: Activation Record

- An activation record instance is dynamically created when a subprogram is called
- Because the call and return semantics specify that the subprogram last called is the first to complete, it is reasonable to use a stack, so
 - Activation record instances reside on the run-time stack
- Every subprogram activation (recursive or non-recursive) creates a new instance of an activation record on the stack.
- The *Environment Pointer (EP)* must be maintained by the run-time system. It always points at the base of the activation record instance of the currently executing program unit.

An Example: C Function

```
void sub(float total, int part)
{
    int list[5];
    float sum;
    ...
}
```

Local	sum
Local	list [4]
Local	list [3]
Local	list [2]
Local	list [1]
Local	list [0]
Parameter	part
Parameter	total
Dynamic link	
Return address	

Revised Semantic Call/Return Actions

- Create an activation record instance
- Save the execution status of the current program unit
- Compute and pass the parameters
- Pass the return address of the callee
- Transfer the control to the callee

Revised Semantic Call/Return Actions

Before (Prologue)

- Save the old EP to the stack as the dynamic link and create the new value
- Allocate local variables

After (Epilogue)

- If pass-by-value-result or out mode parameters are used, move the current values of those parameters to their corresponding actual parameters
- If it is a function, move the functional value to a place the caller can get it
- Restore the stack pointer by setting it to the value of the current EP minus one and set the EP to the old dynamic link
- Restore the execution status of the caller
- Transfer control back to the caller


An Example Without Recursion

```
void fun1(float r) {  
    int s, t;  
    ...  
    fun2(s);  
    ...  
}  
void fun2(int x) {  
    int y;  
    ...  
    fun3(y);  
    ...  
}  
void fun3(int q) {  
    ...  
}  
void main() {  
    float p;  
    ...  
    fun1(p);  
    ...  
}
```

main calls fun1
fun1 calls fun2
fun2 calls fun3

An Example Without Recursion

main calls fun1
fun1 calls fun2
fun2 calls fun3



ARI = activation record instance

Dynamic Chain and Local Offset

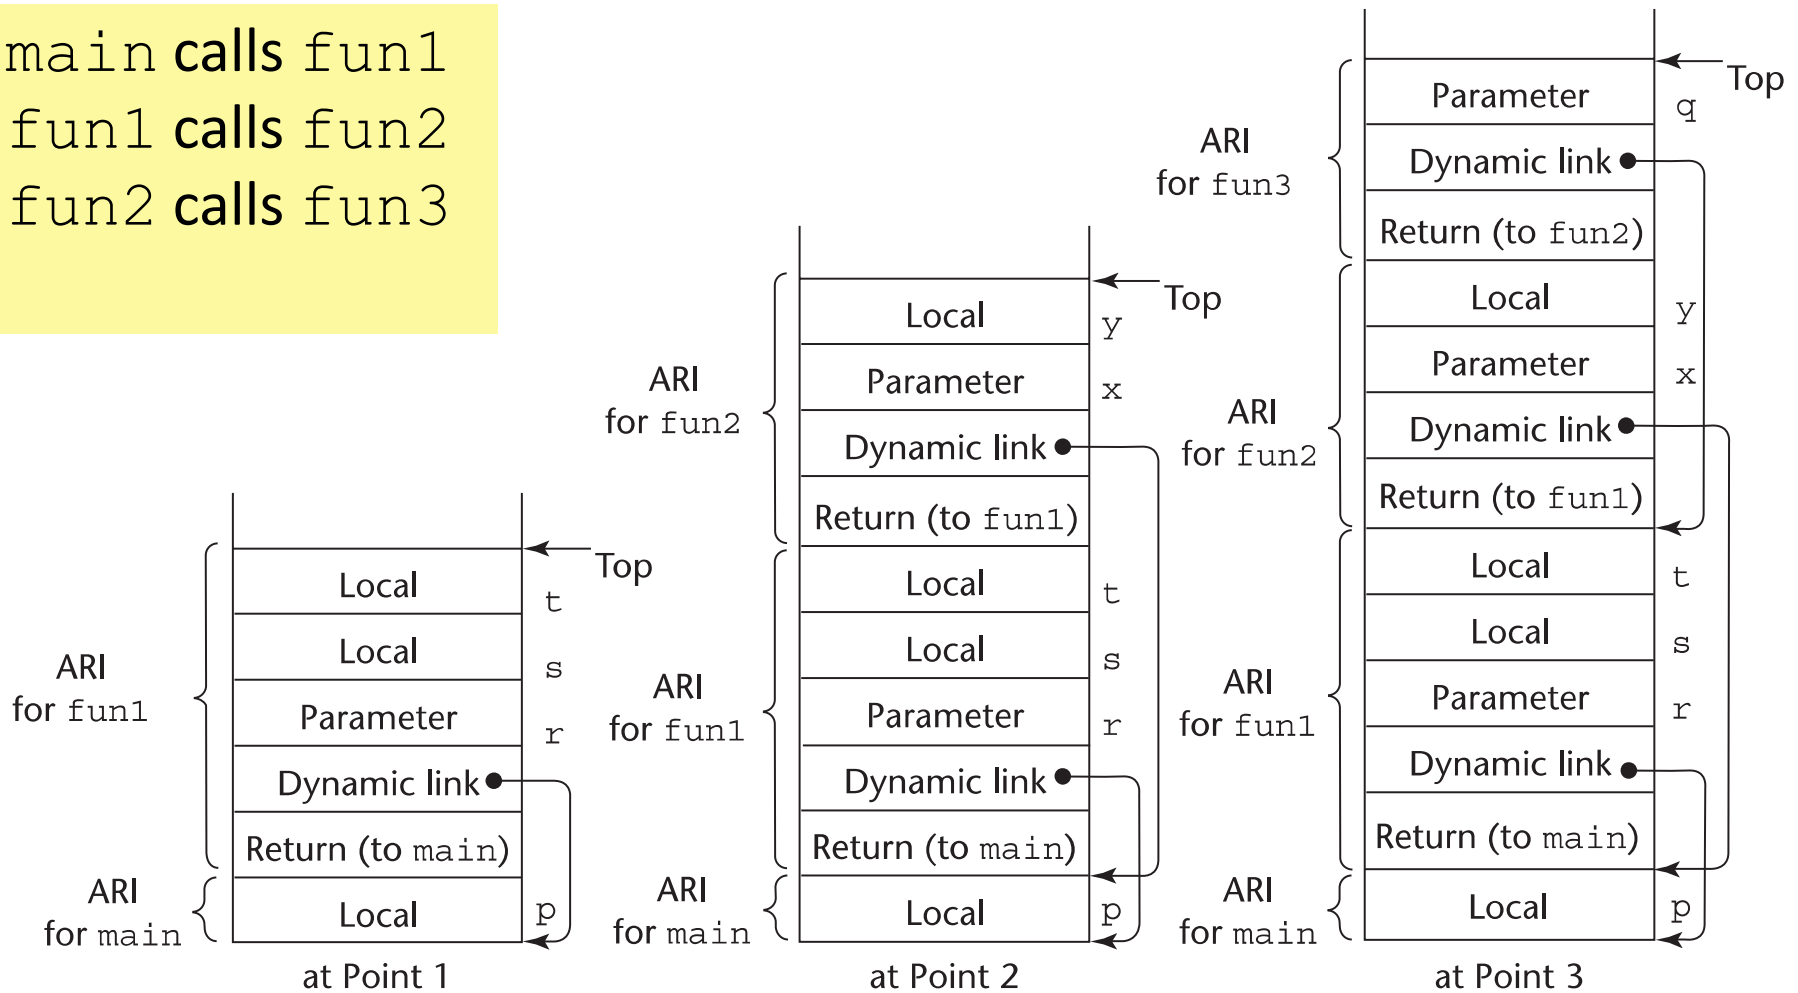
- *Dynamic chain (call chain)*: The collection of dynamic links in the stack at a given time
- Local variables can be accessed by their offset from the beginning of the activation record, whose address is in the EP. This offset is called the *local_offset*
- The local_offset of a local variable can be determined by the compiler at compile time
 - How?

Local offset

- The `local_offset` of a variable in an activation record can be determined at compile time, using the order, types, and sizes of variables.
- For simplicity, we assume that all variables take one position in the activation record.

An Example Without Recursion

main calls fun1
fun1 calls fun2
fun2 calls fun3




ARI = activation record instance

In fun1, the local_offset of s is 3; for t it is 4.
In fun2, the local_offset of y is 3.

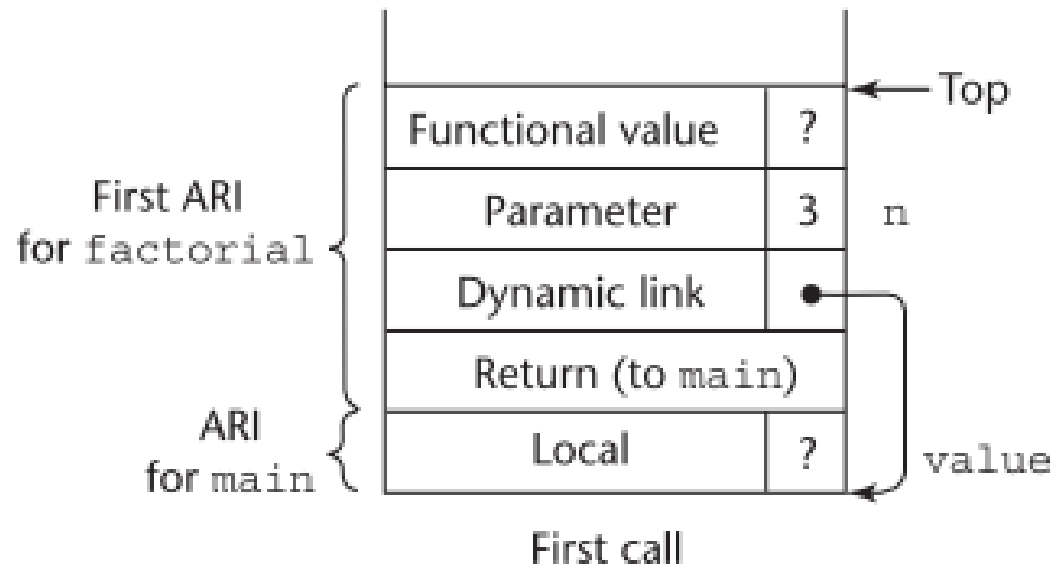
Recursion

```
int factorial (int n) {  
    <-----1  
    if (n <= 1)  
        return 1;  
    else return (n * factorial(n - 1));  
    <-----2  
}  
void main() {  
    int value;  
    value = factorial(3);  
    <-----3  
}
```

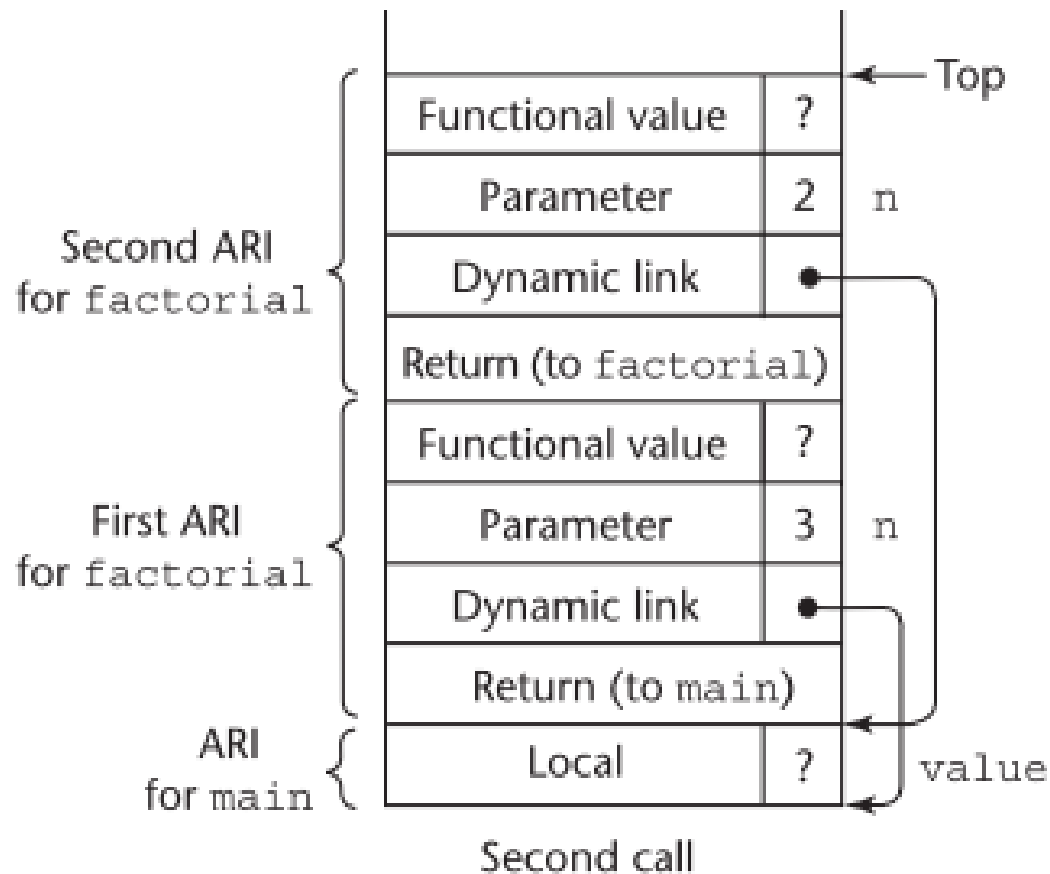
Activation Record for factorial



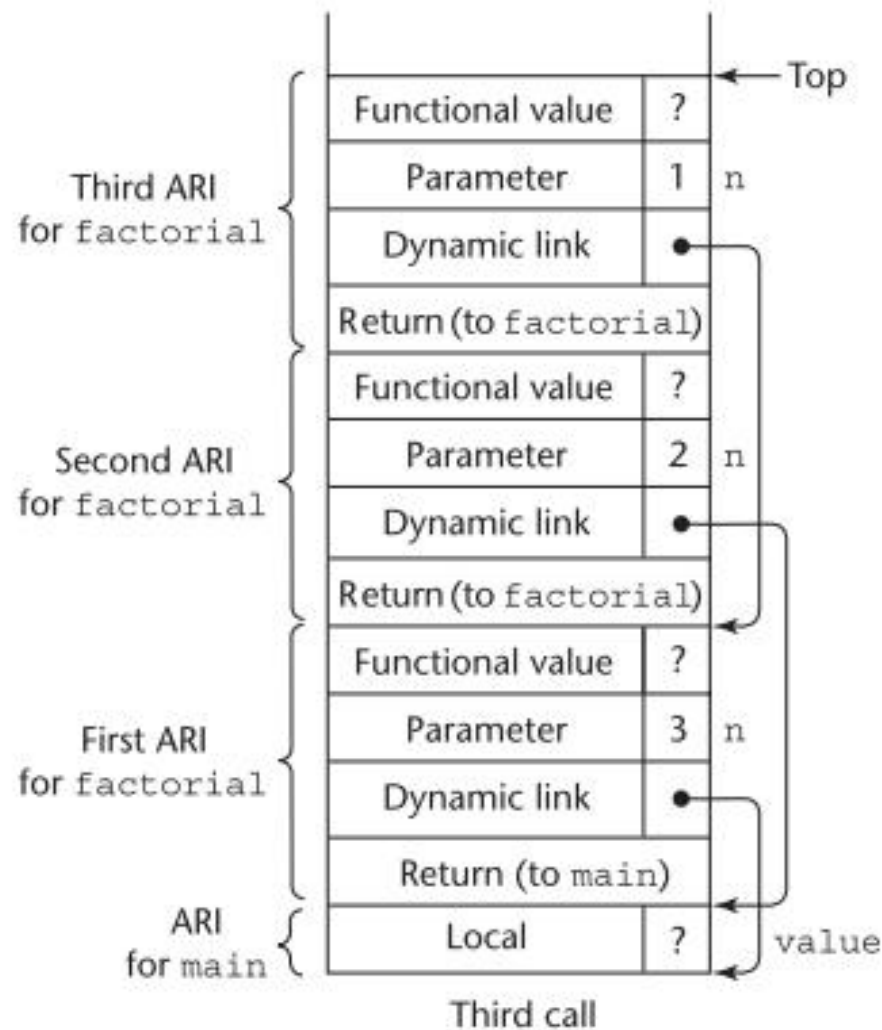
Stack contents during execution of main and factorial



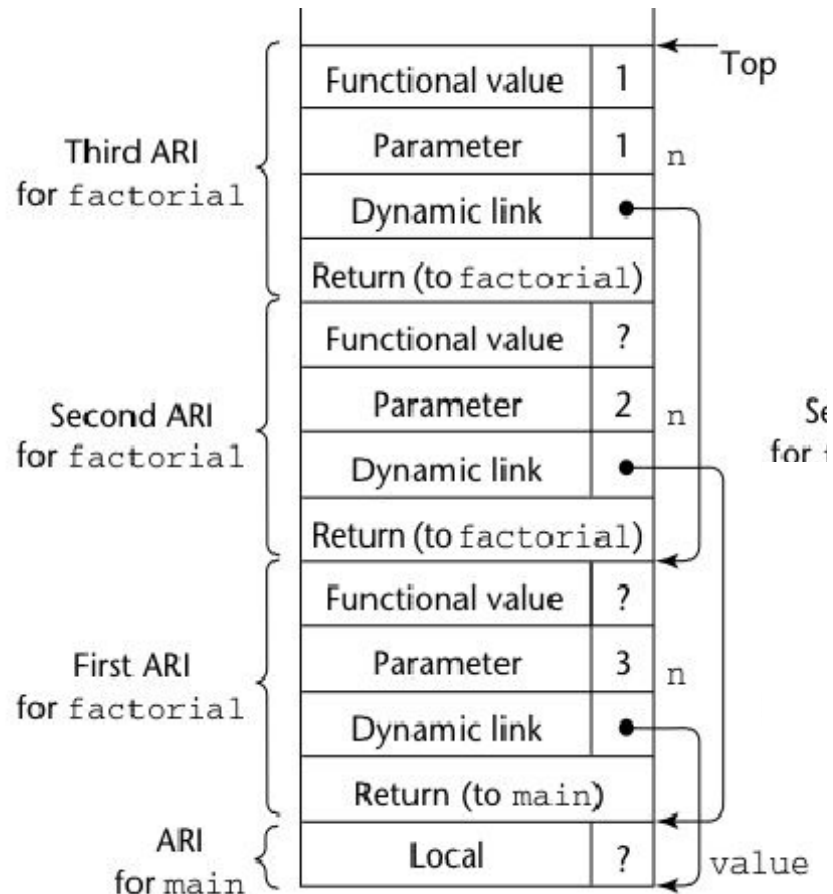
Stack contents during execution of main and factorial



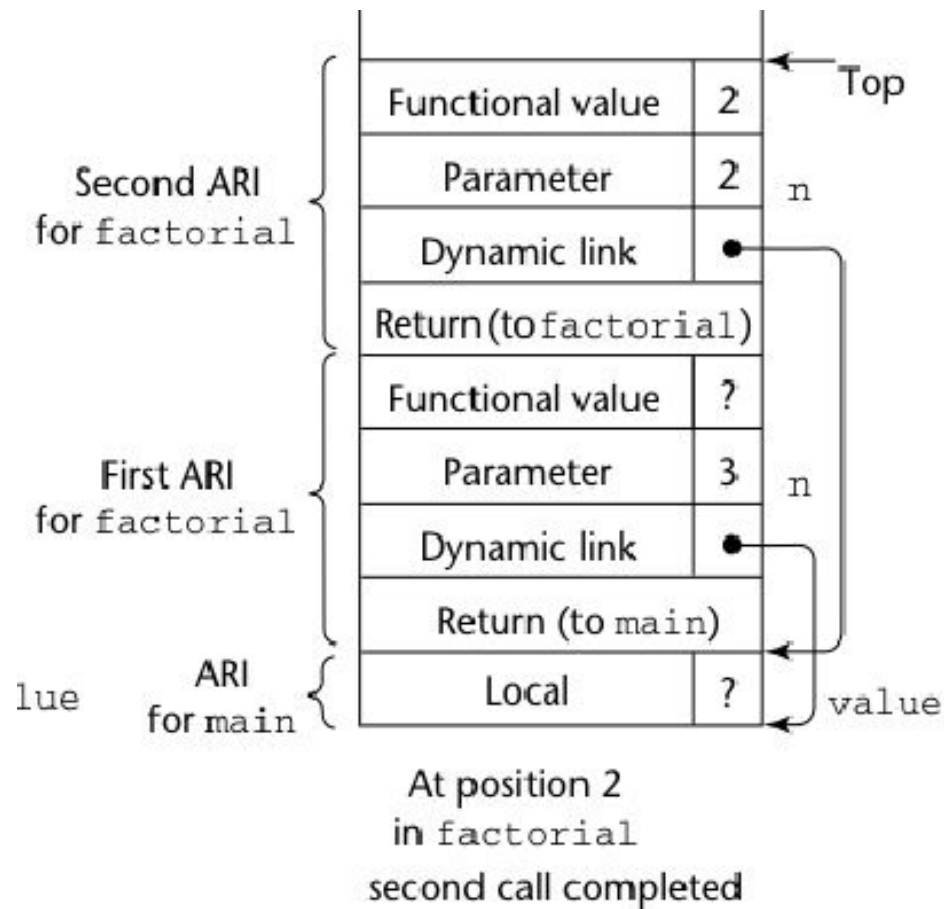
Stack contents during execution of main and factorial



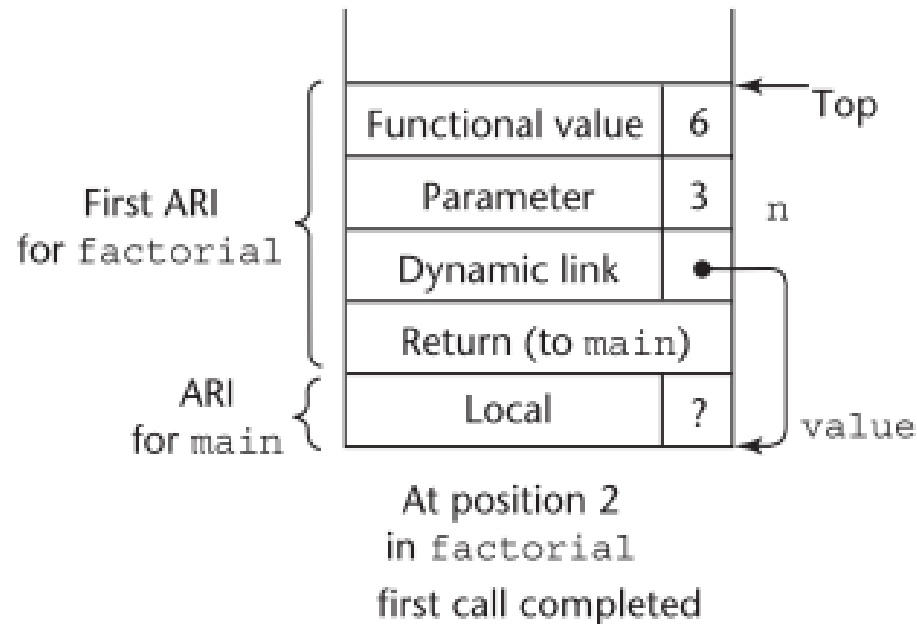
Stack contents during execution of main and factorial




Stack contents during execution of main and factorial



Stack contents during execution of main and factorial



Stack contents during execution of main and factorial



Nested Subprograms

- Some non-C-based static-scoped languages (e.g., Fortran 95, Ada, Python, JavaScript, Ruby, and Lua) use stack-dynamic local variables and allow subprograms to be nested
- All variables that can be non-locally accessed reside in some activation record instance in the stack
- The process of locating a non-local reference:
 1. Find the correct activation record instance
 2. Determine the correct offset within that activation record instance

Locating a Non-local Reference

- Finding the offset is easy
- Finding the correct activation record instance
 - Static semantic rules guarantee that all non-local variables that can be referenced have been allocated in some activation record instance that is on the stack when the reference is made

Implementing Static Scoping

- The *static link* in an activation record instance for subprogram A points to the bottom of one of the activation record instances of A's static parent
- A *static chain* is a chain of static links that connects certain activation record instances
- The static chain from an activation record instance connects it to all of its static ancestors
- *Static_depth* is an integer associated with a static scope whose value is the depth of nesting => indicates how deeply it is nested in the outermost scope

Static Scoping (continued)

- A program unit that is not nested inside any other unit has a `static_depth` of 0.
- A subprogram A is defined in a nonnested program unit, its `static_depth` is 1.
- If subprogram A contains the definition of a nested subprogram B , then B 's `static_depth` is 2.
- The *chain_offset* or *nesting_depth* of a nonlocal reference is the difference between the `static_depth` of the reference and `static_depth` of the procedure containing its declaration
- A reference to a variable can be represented by the pair:
(chain_offset, local_offset),
where `local_offset` is the offset in the activation record of the variable being referenced

```
# Global scope
```

```
...
```

```
def f1():
```

```
    def f2():
```

```
        def f3():
```

```
            ...
```

```
        # end of f3
```

```
    ...
```

```
    # end of f2
```

```
...
```

```
# end of f1
```

Static depth of f1 = 0

Static depth of f2 = 1

Static depth of f3 = 2

Example Ada Program*

```
procedure Main_2 is
  X : Integer;
  procedure Bigsub is
    A, B, C : Integer;
    procedure Sub1 is
      A, D : Integer;
      begin -- of Sub1
        A := B + C;  <-----1
      end; -- of Sub1
    procedure Sub2(X : Integer) is
      B, E : Integer;
      procedure Sub3 is
        C, E : Integer;
        begin -- of Sub3
          Sub1;
          E := B + A;  <-----2
        end; -- of Sub3
      begin -- of Sub2
        Sub3;
        A := D + E;  <-----3
      end; -- of Sub2 }
    begin -- of Bigsub
      Sub2(7);
    end; -- of Bigsub
  begin
    Bigsub;
  end; of Main_2 }
```

Example Ada Program (continued)

- Call sequence for `Main_2`

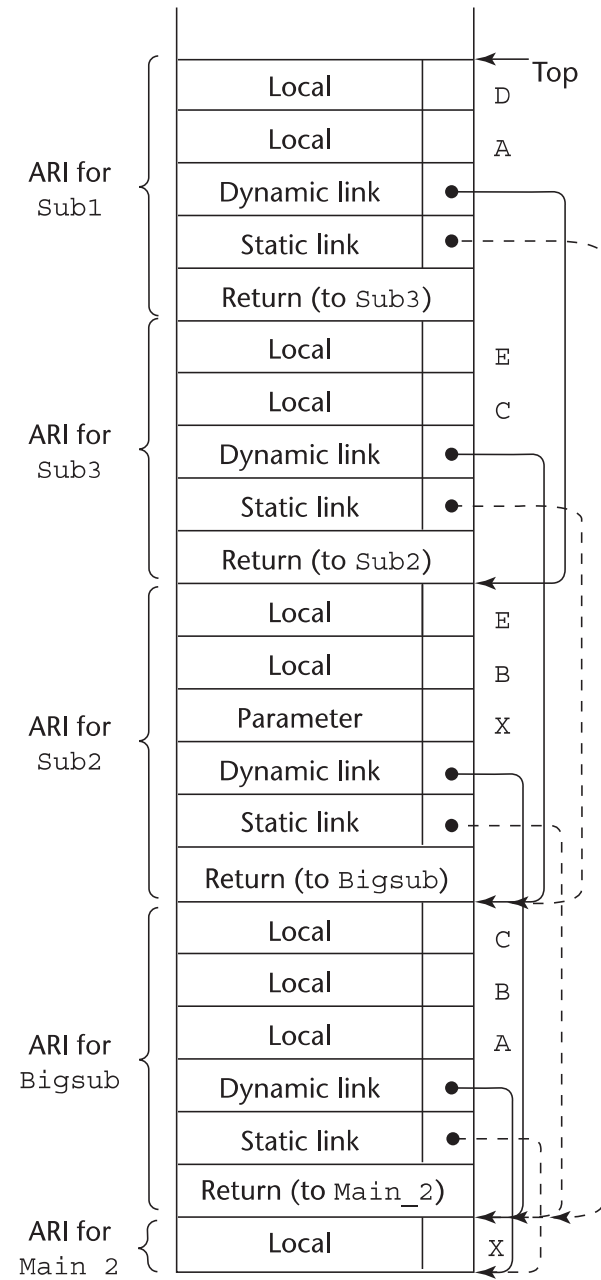
`Main_2` **calls** `Bigsb`

`Bigsb` **calls** `Sub2`

`Sub2` **calls** `Sub3`

`Sub3` **calls** `Sub1`

Stack Contents at Position 1



Example Ada Program*

```
procedure Main_2 is
  X : Integer;
  procedure Bigsub is
    A, B, C : Integer;
    procedure Sub1 is
      A, D : Integer;
      begin -- of Sub1
        A := B + C; <-----1
      end; -- of Sub1
    procedure Sub2(X : Integer) is
      B, E : Integer;
      procedure Sub3 is
        C, E : Integer;
        begin -- of Sub3
          Sub1;
          E := B + A; <-----2
        end; -- of Sub3
      begin -- of Sub2
        Sub3;
        A := D + E; <-----3
      end; -- of Sub2 }
    begin -- of Bigsub
      Sub2(7);
    end; -- of Bigsub
  begin
    Bigsub;
  end; of Main_2 }
```

Variable	Chain offset	Local offset
A at position 1		
B at position 1		
C at position 1		
E at position 2		
B at position 2		
A at position 2		
A at position 3		
D at position 3		
E at position 3		

Evaluation of Static Chains

- Problems:
 1. A nonlocal reference is slow if the nesting depth is large
 2. Time-critical code is difficult:
 - a. Costs of nonlocal references are difficult to determine
 - b. Code changes can change the nesting depth, and therefore the cost

Blocks

- Recall that blocks are user-specified local scopes for variables

- An example in C

```
{int temp;  
  temp = list [upper];  
  list [upper] = list [lower];  
  list [lower] = temp  
}
```

- The lifetime of `temp` in the above example begins when control enters the block
- An advantage of using a local variable like `temp` is that it cannot interfere with any other variable with the same name

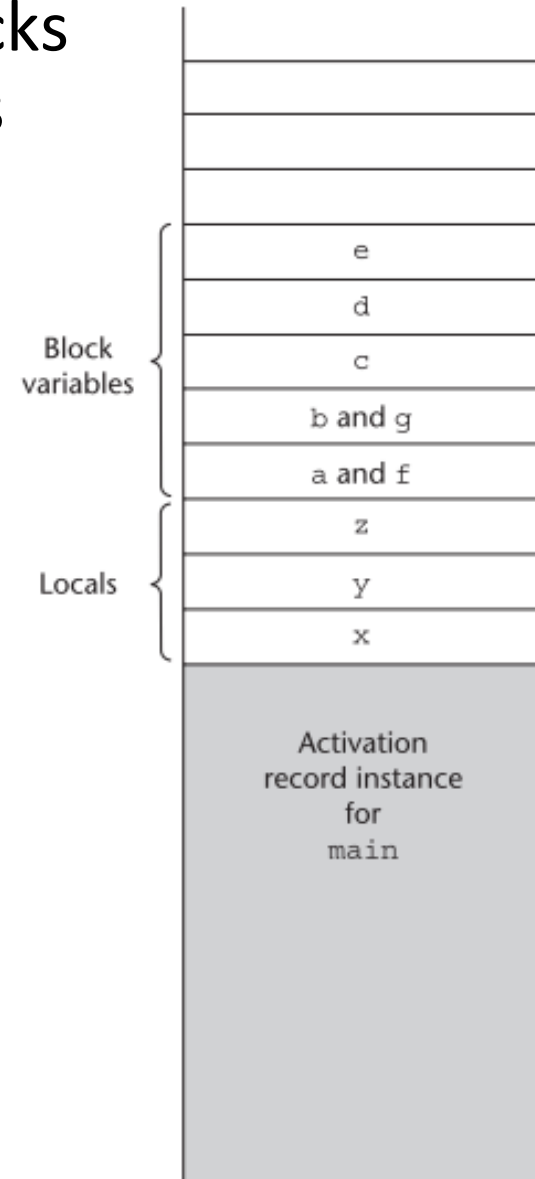
Implementing Blocks

- Two Methods:
 1. Treat blocks as parameter-less subprograms that are always called from the same location
 - Every block has an activation record; an instance is created every time the block is executed
 2. Since the maximum storage required for a block can be statically determined, this amount of space can be allocated after the local variables in the activation record

Implementing Blocks

- Block variable storage when blocks are not treated as parameterless procedures

```
void main() {  
    int x, y, z;  
    while ( ... ) {  
        int a, b, c;  
        ...  
        while ( ... ) {  
            int d, e;  
            ...  
        }  
    }  
    while ( ... ) {  
        int f, g;  
        ...  
    }  
    ...  
}
```




Implementing Dynamic Scoping

- *Deep Access*: non-local references are found by searching the activation record instances on the dynamic chain
 - Length of the chain cannot be statically determined
 - Every activation record instance must have variable names
- *Shallow Access*: put locals in a central place
 - One stack for each variable name
 - Central table with an entry for each variable name

Using Shallow Access to Implement Dynamic Scoping

```
void sub3() {  
    int x, z;  
    x = u + v;  
    ...  
}  
void sub2() {  
    int w, x;  
    ...  
}  
void sub1() {  
    int v, w;  
    ...  
}  
void main() {  
    int v, u;  
    ...  
}
```



(The names in the stack cells indicate the program units of the variable declaration.)

main calls sub1
sub1 calls sub1
sub1 calls sub2
sub2 calls sub3

Summary

- Subprogram linkage semantics requires many action by the implementation
- Simple subprograms have relatively basic actions
- Stack-dynamic languages are more complex
- Subprograms with stack-dynamic local variables and nested subprograms have two components
 - actual code
 - activation record

Summary (continued)

- Activation record instances contain formal parameters and local variables among other things
- Static chains are the primary method of implementing accesses to non-local variables in static-scoped languages with nested subprograms
- Access to non-local variables in dynamic-scoped languages can be implemented by use of the dynamic chain or thru some central variable table method