

Hashing

BBM371 Data Management

Motivation

- Consider the problem of searching an array for a given value
 - If the array is not sorted, the search requires $O(n)$ time
 - If the value isn't there, we need to search all n elements
 - If the value is there, we search $n/2$ elements on average
 - If the array is sorted, we can do a binary search
 - A binary search requires $O(\log n)$ time
 - About equally fast whether the element is found or not
 - It doesn't seem like we could do much better
 - How about an $O(1)$, that is, constant time search?
 - We can do it *if* the array is organized in a particular way


Hashing

- Suppose we were to come up with a “magic function” that, given a value to search for, would tell us exactly where in the array to look
 - If it’s in that location, it’s in the array
 - If it’s not in that location, it’s not in the array
- This function would have no other purpose
- If we look at the function’s inputs and outputs, they probably won’t “make sense”
- This function is called a **hash function** because it “makes hash” of its inputs

Hash Function

- Hash function h :
 - Mapping from U to the slots of a hash table $T[0..m-1]$.
 $h : U \rightarrow \{0, 1, \dots, m-1\}$
- With arrays, key k maps to slot $A[k]$.
- With hash tables, key k maps or “hashes” to slot $T[h[k]]$.
- $h[k]$ is the *hash value* of key k .

Hashing (cont'd)



Finding the Hash Function

- How can we come up with this magic function?
- In general, we cannot--there is no such magic function ☹
 - In a few specific cases, where all the possible values are known in advance, it has been possible to compute a perfect hash function
- What is the next best thing?
 - A perfect hash function would tell us exactly where to look
 - In general, the best we can do is a function that tells us where to *start* looking!

Example Hash Function

- Map a key k into one of the m slots by taking the remainder of k divided by m (Division Method). That is,

$$h(k) = k \bmod m$$

$h(\text{ssn}) = \text{ssn} \bmod 100$ (i.e., the last two digits)

e.g., if $\text{ssn} = 10123411$ then $h(10123411) = 11$


Keys as Natural Numbers

- Hash functions assume that the keys are natural numbers.
- When they are not, have to interpret them as natural numbers.
- Example: Interpret a character string as an integer expressed in some radix notation. Suppose the string is CLRS:
 - ASCII values: C=67, L=76, R=82, S=83.
 - There are 128 basic ASCII values.
 - So, $CLRS = 67 \cdot 128^3 + 76 \cdot 128^2 + 82 \cdot 128^1 + 83 \cdot 128^0$
141,764,947.

Collision

- When two values hash to the same array location, this is called a **collision**
- Collisions are normally treated as “first come, first served”—the first value that hashes to the location gets it
- We have to find something to do with the second and subsequent values that hash to this same location


Collisions



Handling Collisions


- **Chaining**

- Store all elements that hash to the same slot in a linked list.
- Store a pointer to the head of the linked list in the hash table slot.




- **Open Addressing**


- All elements stored in hash table itself.
- When collisions occur, use a systematic (consistent) procedure to store elements in free slots of the table.



Chaining



Chaining (cont'd)



Dictionary Operations with Chaining

- Chained-Hash-Insert (T, x)

- Insert x at the head of list $T[h(\text{key}[x])]$.
- Worst-case complexity – $O(1)$.

- Chained-Hash-Delete (T, x)


- Delete x from the list $T[h(\text{key}[x])]$.
- Worst-case complexity – proportional to length of list with singly-linked lists. $O(1)$ with doubly-linked lists.

- Chained-Hash-Search (T, k)

- Search an element with key k in list $T[h(k)]$.
- Worst-case complexity – proportional to length of list.

Analysis of Hashing with Chaining: Worst Case

- How long does it take to search for an element with a given key?
- Worst case:
 - All n keys hash to the same slot
 - Worst-case time to search is $\Theta(n)$, plus time to compute the hash function



Open Addressing

- If collision occurs, open addressing scheme **probes** for some other empty (or open) location in which to place the item.
- The sequence of locations that we examine is called the **probe sequence**.
- There are different open-addressing schemes:
 - *Linear Probing*
 - *Quadratic Probing*
 - *Double Hashing*

Linear Probing

- In linear probing, we search the hash table sequentially starting from the original hash location.
 - If a location is occupied, we check the next location
 - We wrap around from the last table location to the first table location if necessary.

Linear Probing: Example

- Table Size is 11 (0..10)
- Hash Function: $h(x) = x \bmod 11$
- Insert keys:
 - $20 \bmod 11 = 9$
 - $30 \bmod 11 = 8$
 - $2 \bmod 11 = 2$
 - $13 \bmod 11 = 2 \rightarrow 2+1=3$
 - $25 \bmod 11 = 3 \rightarrow 3+1=4$
 - $24 \bmod 11 = 2 \rightarrow 2+1, 2+2, 2+3=5$
 - $10 \bmod 11 = 10$
 - $9 \bmod 11 = 9 \rightarrow 9+1, 9+2 \bmod 11 = 0$

0	9
1	
2	2
3	13
4	25
5	24
6	
7	
8	30
9	20
10	10

Linear Probing: Clustering Problem

- One of the problems with linear probing is that table items tend to **cluster** together in the hash table.
 - This means that the table contains groups of consecutively occupied locations.
- This phenomenon is called **primary clustering**.
- Clusters can get close to one another, and merge into a larger cluster.
 - Thus, the one part of the table might be quite dense, even though another part has relatively few items.
- Primary clustering causes long probe searches and therefore decreases the overall efficiency.

Quadratic Probing

- Primary clustering problem can be almost eliminated if we use **quadratic probing** scheme.
- In quadratic probing,
 - We start from the original hash location i
 - If a location is occupied, we check the locations $i+1^2$, $i+2^2$, $i+3^2$, $i+4^2$. . .
 - We wrap around from the last table location to the first table location if necessary.

Quadratic Probing: Example

- Table Size is 11 (0..10)
- Hash Function: $h(x) = x \bmod 11$
- Insert keys:
 - $20 \bmod 11 = 9$
 - $30 \bmod 11 = 8$
 - $2 \bmod 11 = 2$
 - $13 \bmod 11 = 2 \rightarrow 2+1^2=3$
 - $25 \bmod 11 = 3 \rightarrow 3+1^2=4$
 - $24 \bmod 11 = 2 \rightarrow 2+1^2, 2+2^2=6$
 - $10 \bmod 11 = 10$
 - $9 \bmod 11 = 9 \rightarrow 9+1^2, 9+2^2 \bmod 11, 9+3^2 \bmod 11 = 7$

0	
1	
2	2
3	13
4	25
5	
6	24
7	9
8	30
9	20
10	10

Double Hashing

- Double hashing also reduces clustering.
- In linear probing and quadratic probing , the probe sequences are independent from the key.
- We can select increments used during probing using a second hash function. The second hash function h_2 should be:

$$h_2(\text{key}) \neq 0$$

$$h_2 \neq h_1$$

- We first probe the location $h_1(\text{key})$
 - If the location is occupied, we probe the location $h_1(\text{key})+h_2(\text{key})$, $h_1(\text{key})+(2*h_2(\text{key}))$, ...

Double Hashing: Example

- Table Size is 11 (0..10)
- Hash Functions: $h_1(x) = x \bmod 11$
 $h_2(x) = 7 - (x \bmod 7)$
- Insert keys:
 - $58 \bmod 11 = 3$
 - $14 \bmod 11 = 3 \rightarrow 3+7=10$
 - $91 \bmod 11 = 3 \rightarrow 3+7, 3+2*7 \bmod 11=6$

0	
1	
2	
3	58
4	
5	
6	91
7	
8	
9	
10	14

Open Addressing: Retrieval & Deletion

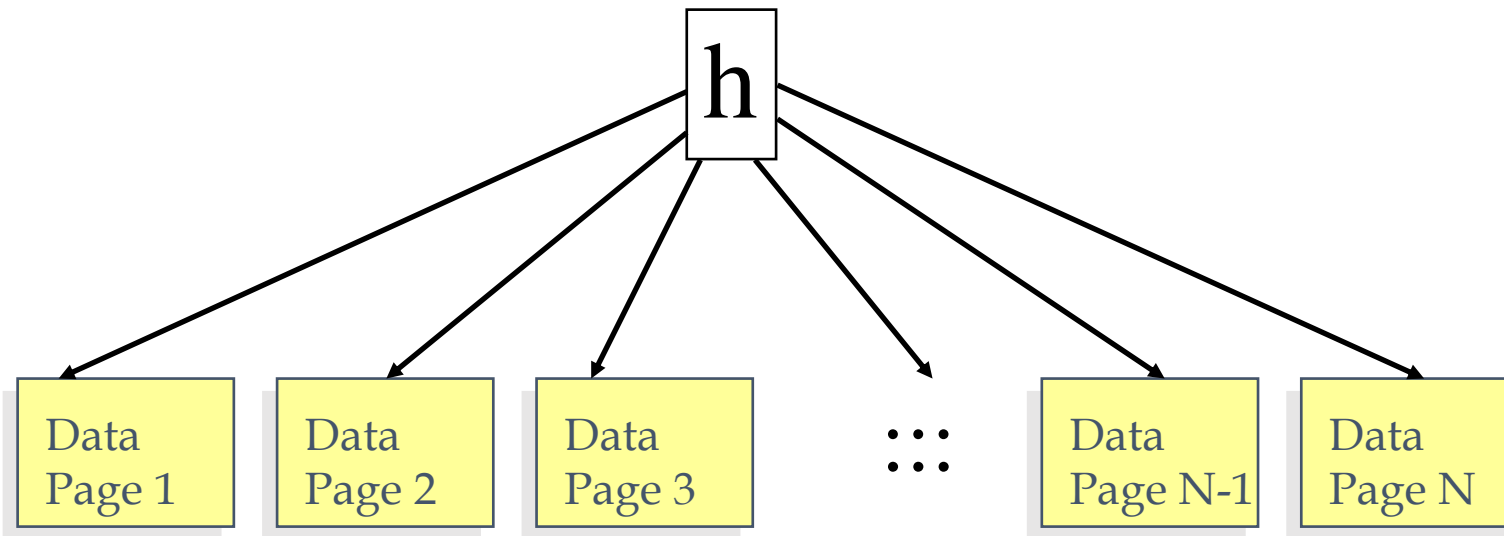
- In open addressing, to find an item with a given key:
 - We probe the locations (same as insertion) until we find the desired item or we reach to an empty location.
- Deletions in open addressing cause complications:
 - We CANNOT simply delete an item from the hash table because this new empty (deleted locations) cause to stop prematurely (incorrectly) indicating a failure during a retrieval.
 - ***Solution:*** We have to have three kinds of locations in a hash table: ***Occupied, Empty, Deleted.***
 - A deleted location will be treated as an occupied location during retrieval and insertion.

Hashing Techniques

- Static Hashing
- Dynamic Hashing
 - Linear
 - Extendable

Static Hashing

- With hash based indexing, we assume that we have a function h , which tells us where to place any given record.
- E.g., **$\text{page_number} = h(\text{value}) \bmod N$** , N should be prime

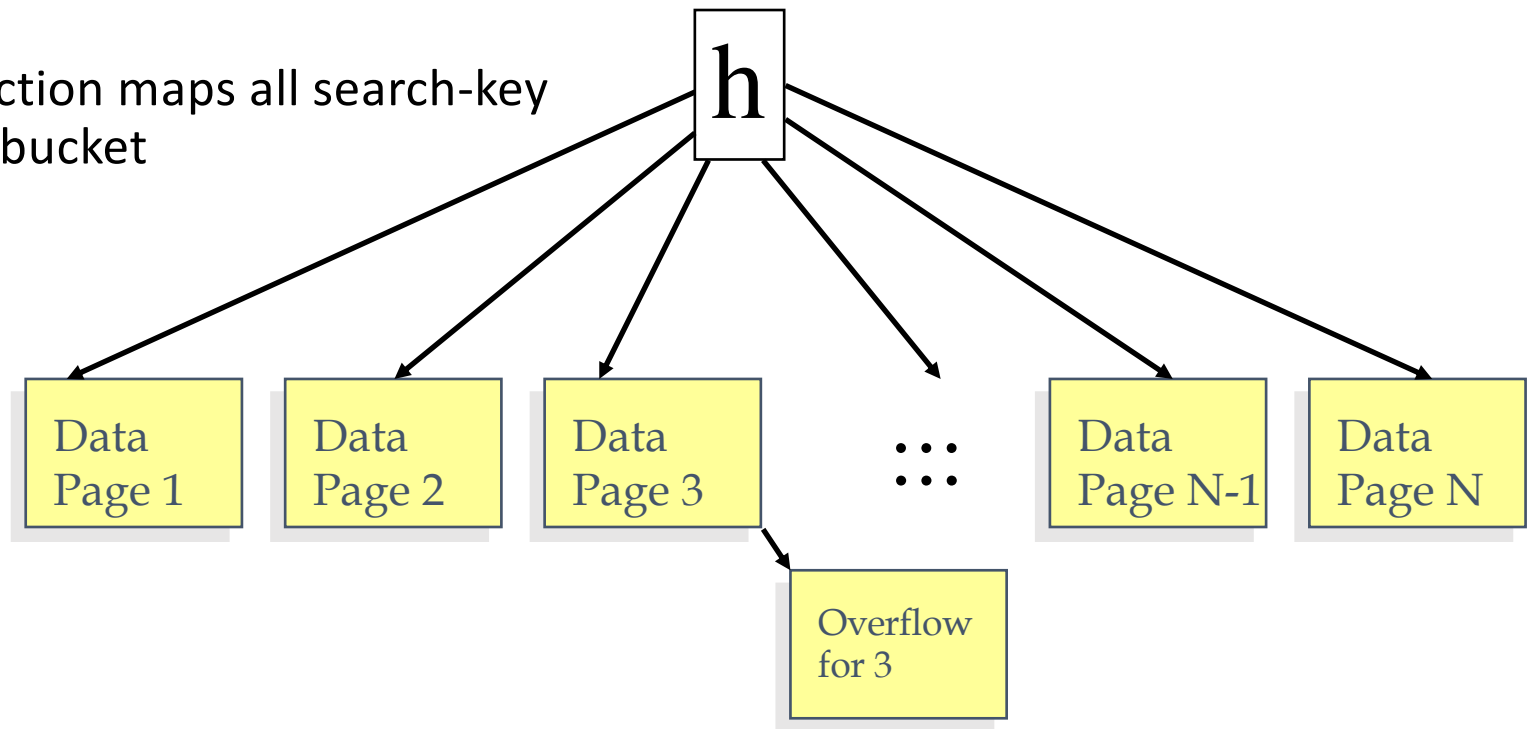


Static Hashing

- A **bucket** is a unit of storage containing one or more records (a bucket is typically a disk block).
- In a **hash file organization** we obtain the bucket of a record directly from its search-key value using a **hash function**.
- Hash function h is a function from the set of all search-key values K to the set of all bucket addresses B .
- Hash function is used to locate records for access, insertion as well as deletion.
- Records with different search-key values may be mapped to the same bucket; thus entire bucket has to be searched sequentially to locate a record.

Static Hashing: Overflow

- Insertion may cause overflow. The solution is to create chains of overflow pages.
- The worst hash function maps all search-key values to the same bucket



Deficiencies of Static Hashing

- In static hashing, function h maps search-key values to a fixed set of B bucket addresses.
 - Databases grow with time. If initial number of buckets is too small, performance will degrade due to too much overflows.
 - If file size at some point in the future is anticipated and number of buckets allocated accordingly, significant amount of space will be wasted initially.
 - If database shrinks, again space will be wasted.
 - One option is periodic re-organization of the file with a new hash function, but it is very expensive.
- These problems can be avoided by using techniques that allow the number of buckets to be modified dynamically.

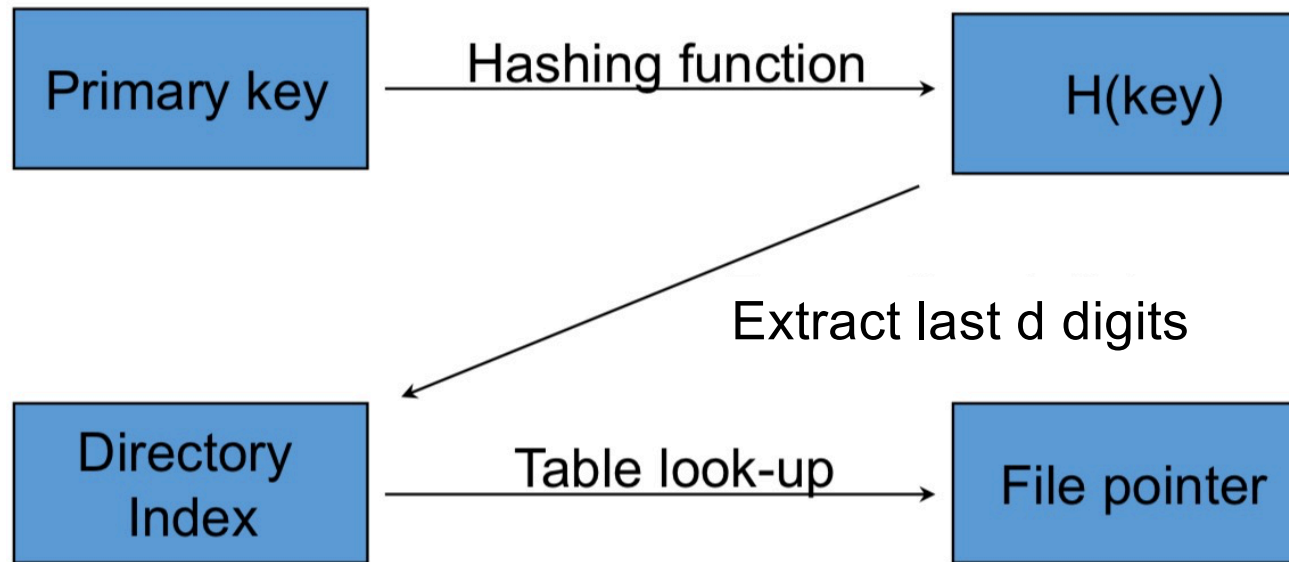
Dynamic Hashing

- Dynamic = Changing number of Buckets B dynamically
- Two methods
 - Extendible (or Extensible) Hashing: Grow B by doubling it
 - Linear Hashing: Grow B by incrementing it by 1
- To save storage space both methods can choose to shrink B dynamically
- Must avoid oscillations when removes and additions are both common.

Extendible Hashing

- **Idea:** Use *directory of pointers to buckets*
- double # of buckets B by *doubling the directory*, splitting just the bucket that overflowed!
- Directory much smaller than file, so doubling it is much cheaper.
- Only one page of data entries is split. *No overflow blocks.*
- Trick lies in how hash function is adjusted!

Extendible Hashing Overview

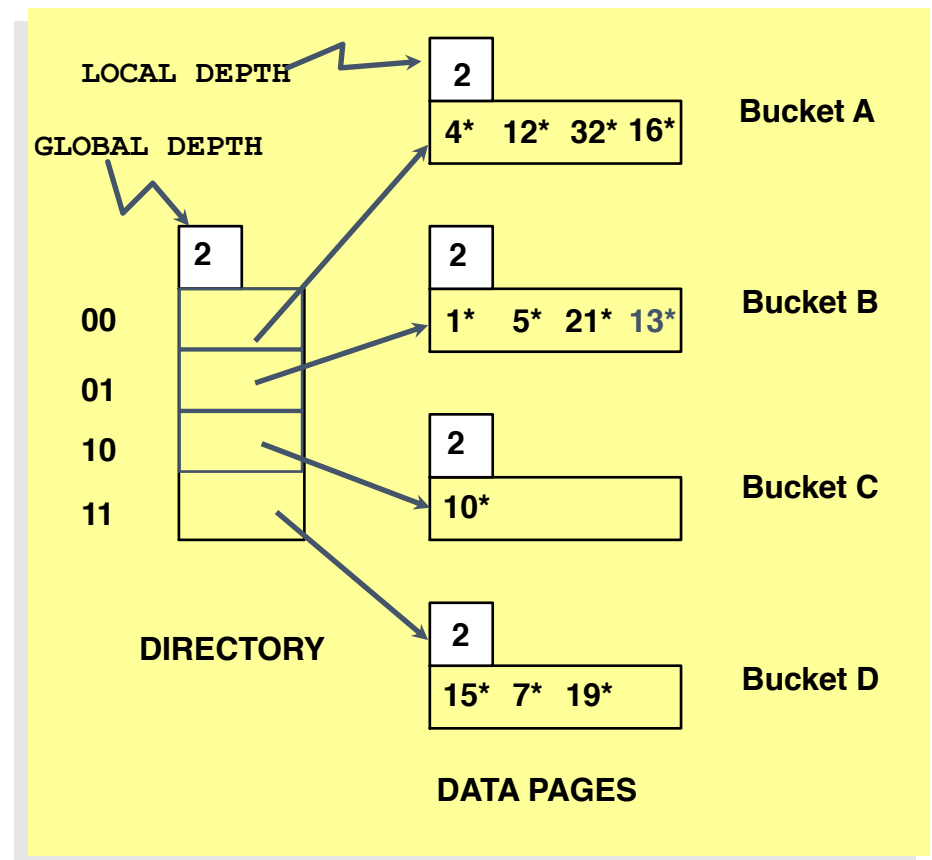


Extendible Hashing

- $h(k)$ maps keys to a fixed address space
- File pointers point to blocks of records known as buckets,
 - where an entire bucket is read by one physical data transfer, buckets may be added to or removed from the file dynamically
- The (last) d bits are used as an index in a directory array containing 2^d entries, which usually resides in primary memory
- The value d , the directory size (2^d), and the number of buckets change automatically as the file expands and contracts

Example

- Directory is array of size 4.
- To find bucket for r , take last '*global depth*' # bits of $h(r)$; we denote r by $h(r)$.
 - If $h(r) = 5 = \text{binary } 101$, it is in bucket pointed to by 01 .
- *Global depth of directory*: Max # of bits needed to tell which bucket an entry belongs to.
- *Local depth of a bucket*: # of bits used to determine if an entry belongs to this bucket.



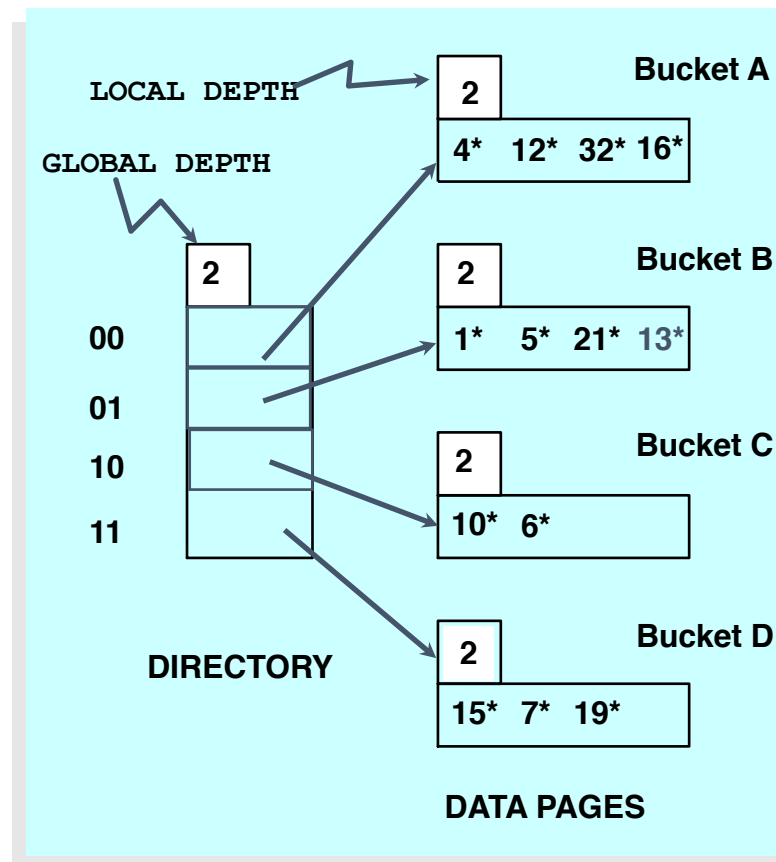
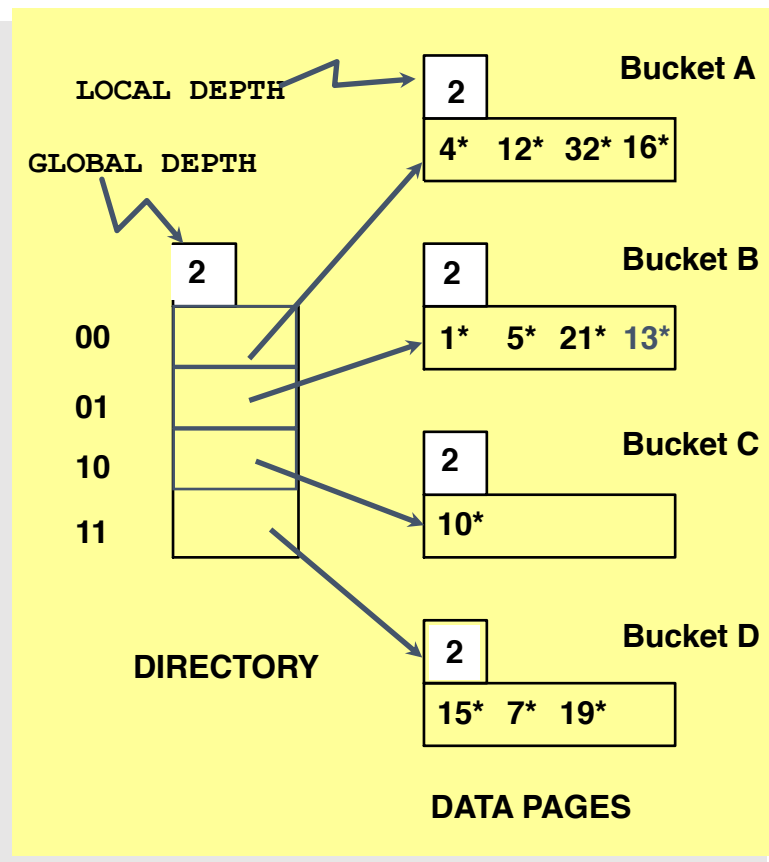
Insert an Item

- Locate the bucket
 - If there is space in bucket, insert the item.
 - If bucket is full, split it (*allocate new page, re-distribute*).
- *If necessary*, double the directory.
 - If insert causes *local depth* to become $>$ *global depth*
 - directory is doubled by *copying it over* and 'fixing' pointer to split image page.

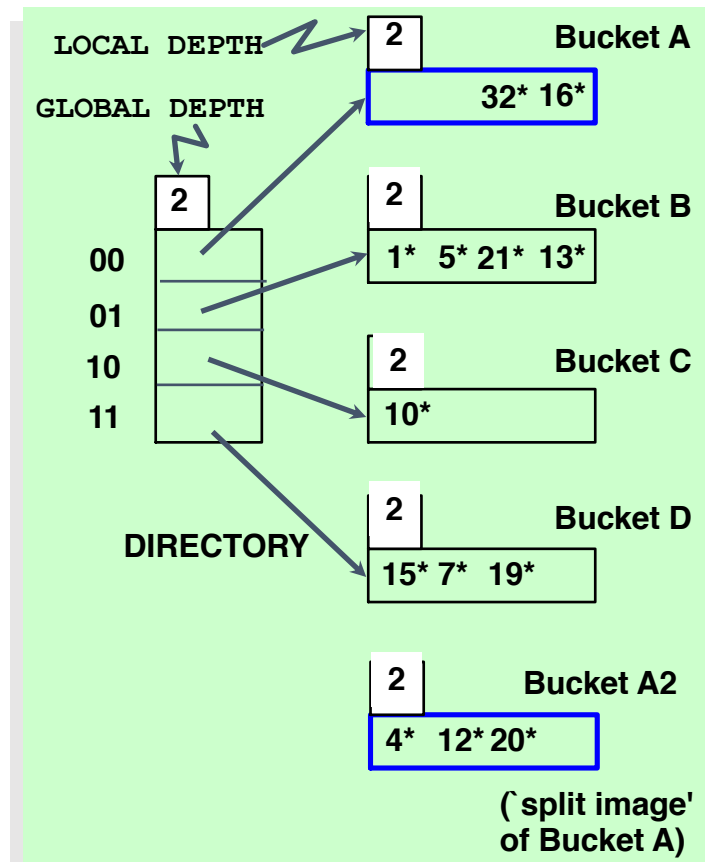
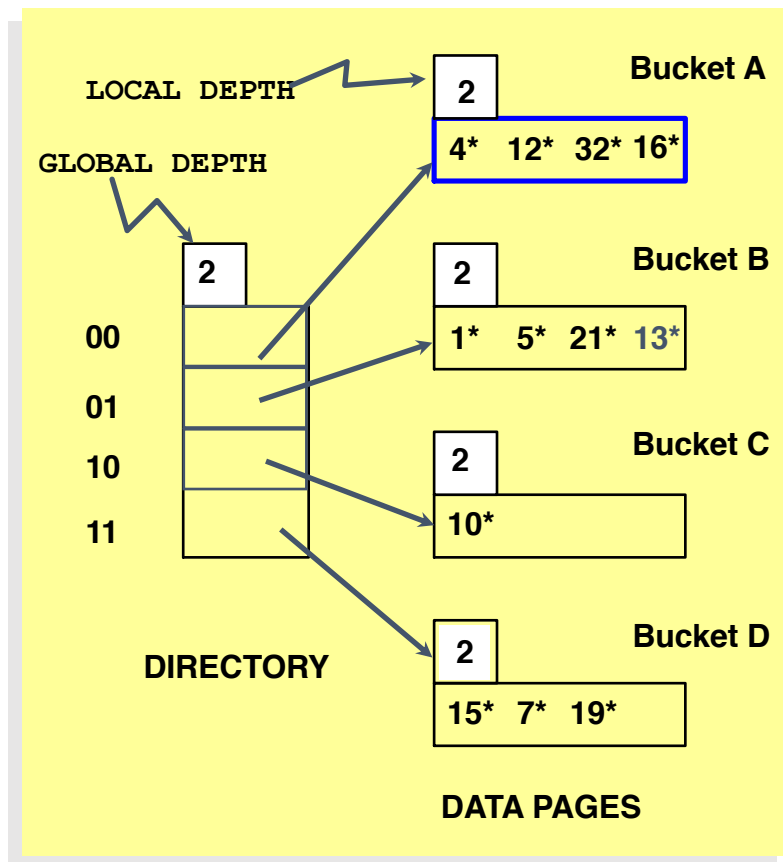
Insert $h(r) = 6$ (The Easy Case)

6 = binary 00110

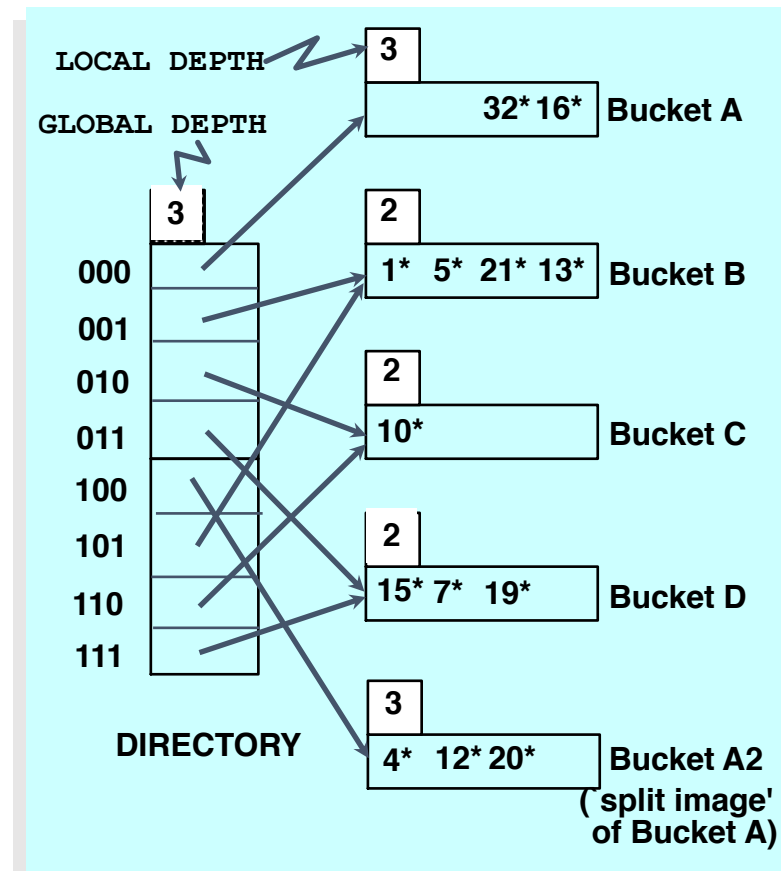
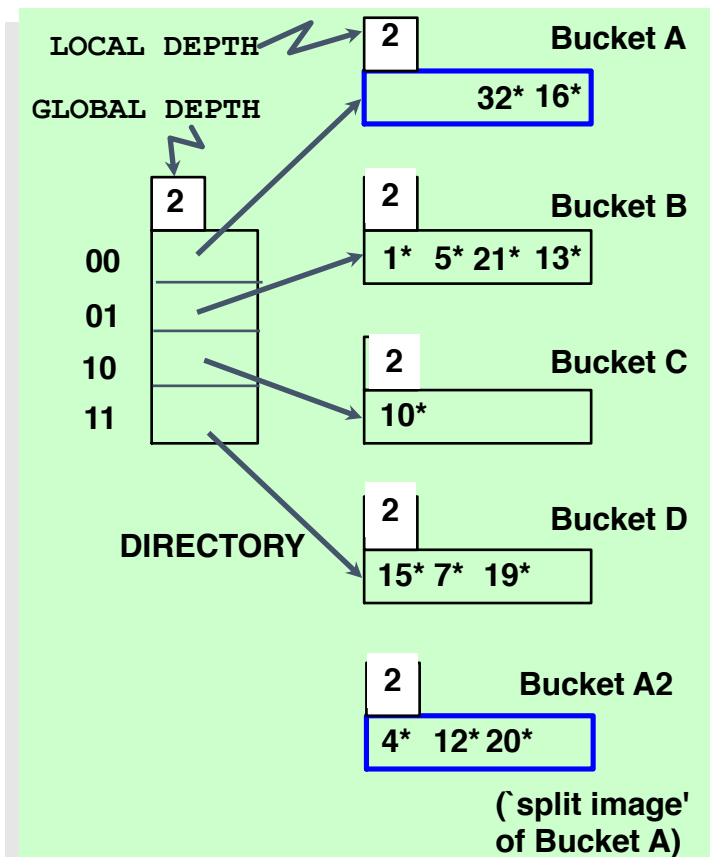
6 = binary 00110



Insert $h(r) = 20$ (Causes Doubling) 20 = binary 10100



Insert $h(r) = 20$ (Causes Doubling) $20 = \text{binary } 10100$



Comments on Extendible Hashing

- If directory fits in memory, equality search answered with one disk access; else two.
 - 100MB file, 100 bytes/rec, 4K pages contains 1,000,000 records (as data entries) and 25,000 directory elements; chances are high that directory will fit in memory.
 - Directory grows in spurts, and, if the distribution *of hash values* is skewed, directory can grow large.
- **Delete**: If removal of data entry makes bucket empty, can be merged with 'split image'. If each directory element points to same bucket as its split image, we can halve directory (this is rare in practice).

Linear Hashing

- This is another dynamic hashing scheme, an alternative to Extendible Hashing.
- LH handles the problem of long overflow chains without using a directory.
- Idea: Use a family of hash functions h_0, h_1, h_2, \dots
 - $h_i(\text{key}) = h(\text{key}) \bmod(2^i N)$; N = initial # buckets
 - h is some hash function (range is *not* 0 to $N-1$)
 - If $N = 2^{d_0}$, for some d_0 , h_i consists of applying h and looking at the last d_i bits, where $d_i = d_0 + i$.
 - h_{i+1} doubles the range of h_i (similar to directory doubling)

Linear Hashing: Bucket Split

- When the first overflow occurs (it can occur in any bucket), bucket 0, which is pointed by p , is split (rehashed) into two buckets:
 - The original bucket 0 and a new bucket m .
- A new empty page is also added in the overflown bucket to accommodate the overflow.
- The search values originally mapped into bucket 0 (using function h_0) are now distributed between buckets 0 and m using a new hashing function h_1 .





Linear Hashing: Insertion

- Locate bucket to insert
- If bucket to insert into is full:
 - Add overflow page and insert data entry.
 - (*Maybe*) Split *Next* bucket and increment *Next*.
- A split occurs in case of overflow
- Since buckets are split round-robin, long overflow chains don't develop!

Linear Hashing: Background

Insert 20





We have seen what it means to split a bucket...



Before

After

We have seen what it means to add an overflow page to a bucket...



Before


After


Triggering Splits

- A split performed whenever a bucket overflow occurs is an **uncontrolled** split.
- Let l denote the Linear Hash- ing scheme's load factor, i.e., $l = S/b$ where S is the total number of records and b is the number of buckets used.
- The load factor achieved by uncontrolled splits is usually between 50–70%, depending on the page size and the search value distribution.
- In practice, higher storage utilization is achieved if a split is triggered not by an overflow, but when the load factor l becomes greater than some upper threshold, which is called **controlled** split.

Overview of Splitting as Rounds

- Splits occur in a round robin fashion, i.e., as rounds.






Bucket to be split

Next


h_{Level}



Bucket to be split

Next

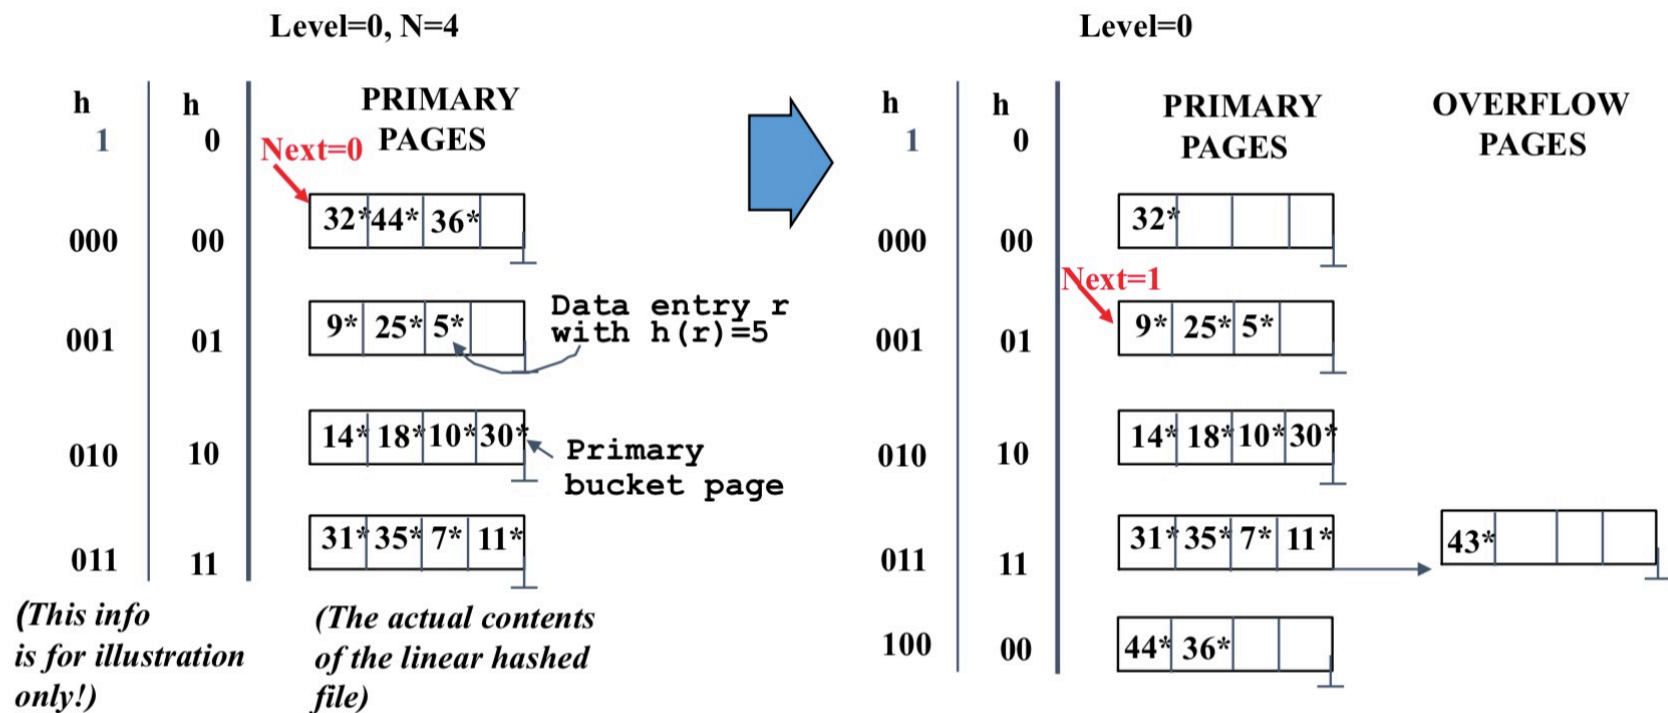
h_{Level}



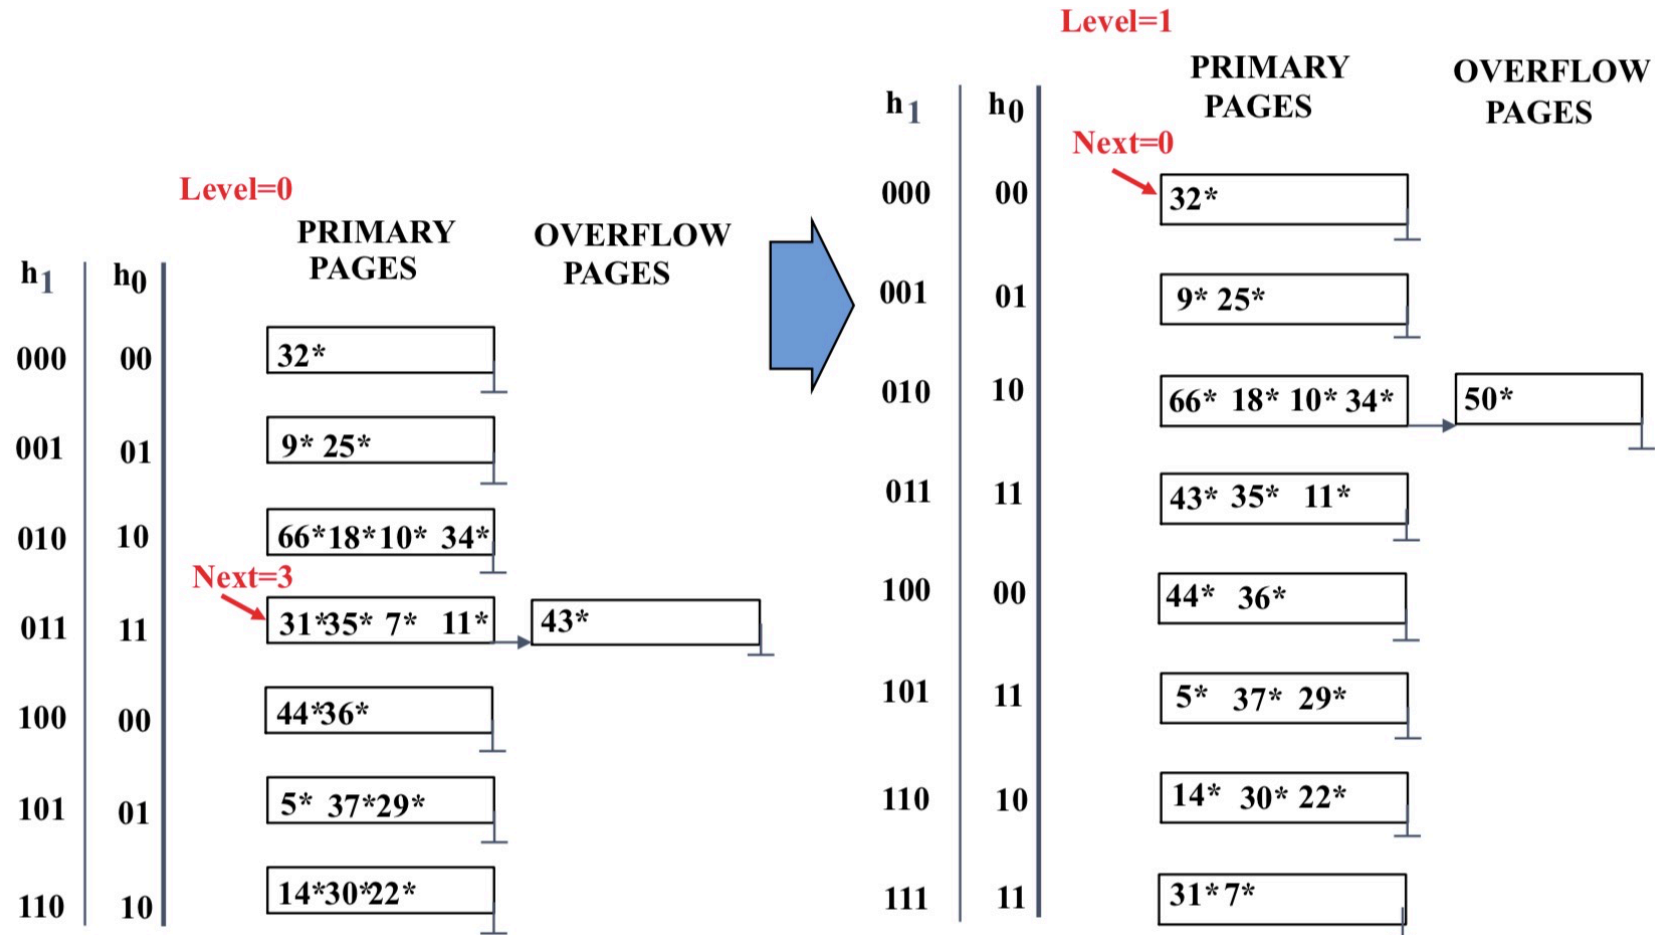
Linear Hashing: Example

On **split**, $h_{\text{Level}+1}$ is used to **re-distribute** entries.

insert 43 → adds overflow → triggers split → increments Next



Example: End of Round



Summary

- Hash-based indexes: best for equality searches, cannot support range searches.
- Static Hashing can lead to long overflow chains.
- Extendible Hashing avoids overflow pages by splitting a full bucket when a new data entry is to be added to it. (*Duplicates may require overflow pages.*)
- Directory to keep track of buckets, doubles periodically.
- Can get large with skewed data; additional I/O if this does not fit in main memory.