

Today

- Image formation
- Color
- Filtering

Image formation

- How are objects in the world captured in an image?
 - Mapping between image and world coordinates

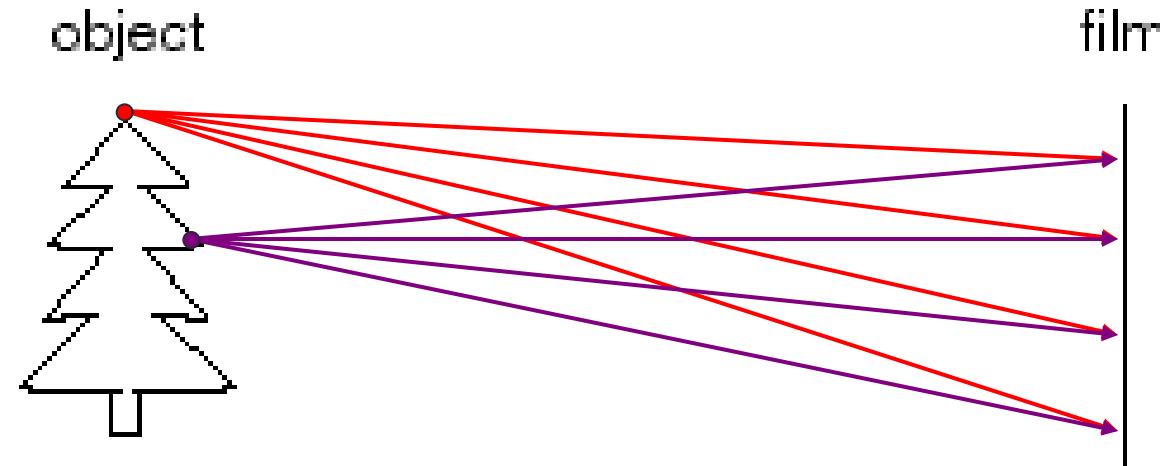
Physical parameters of image formation

- Photometric
 - Surfaces' reflectance properties
 - Direction and intensity of light reaching sensor
- Geometric
 - Type of projection
 - Pinhole camera model
 - Projective geometry
 - Camera pose
- Optical
 - Sensor's lens type
 - focal length, field of view, aperture, shutter speed

What do you need to make a camera from scratch?



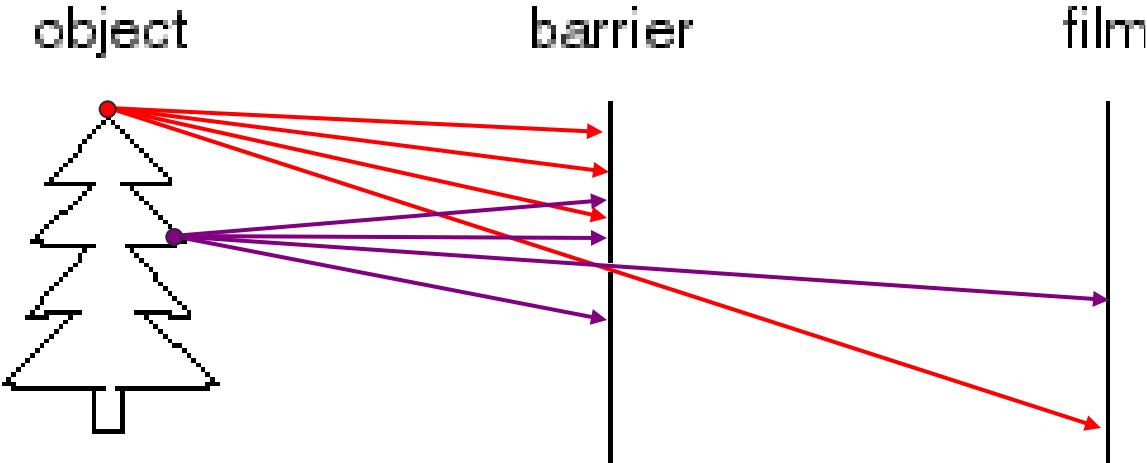
Image formation



Let's design a camera

- Idea 1: put a piece of film in front of an object
- Do we get a reasonable image?

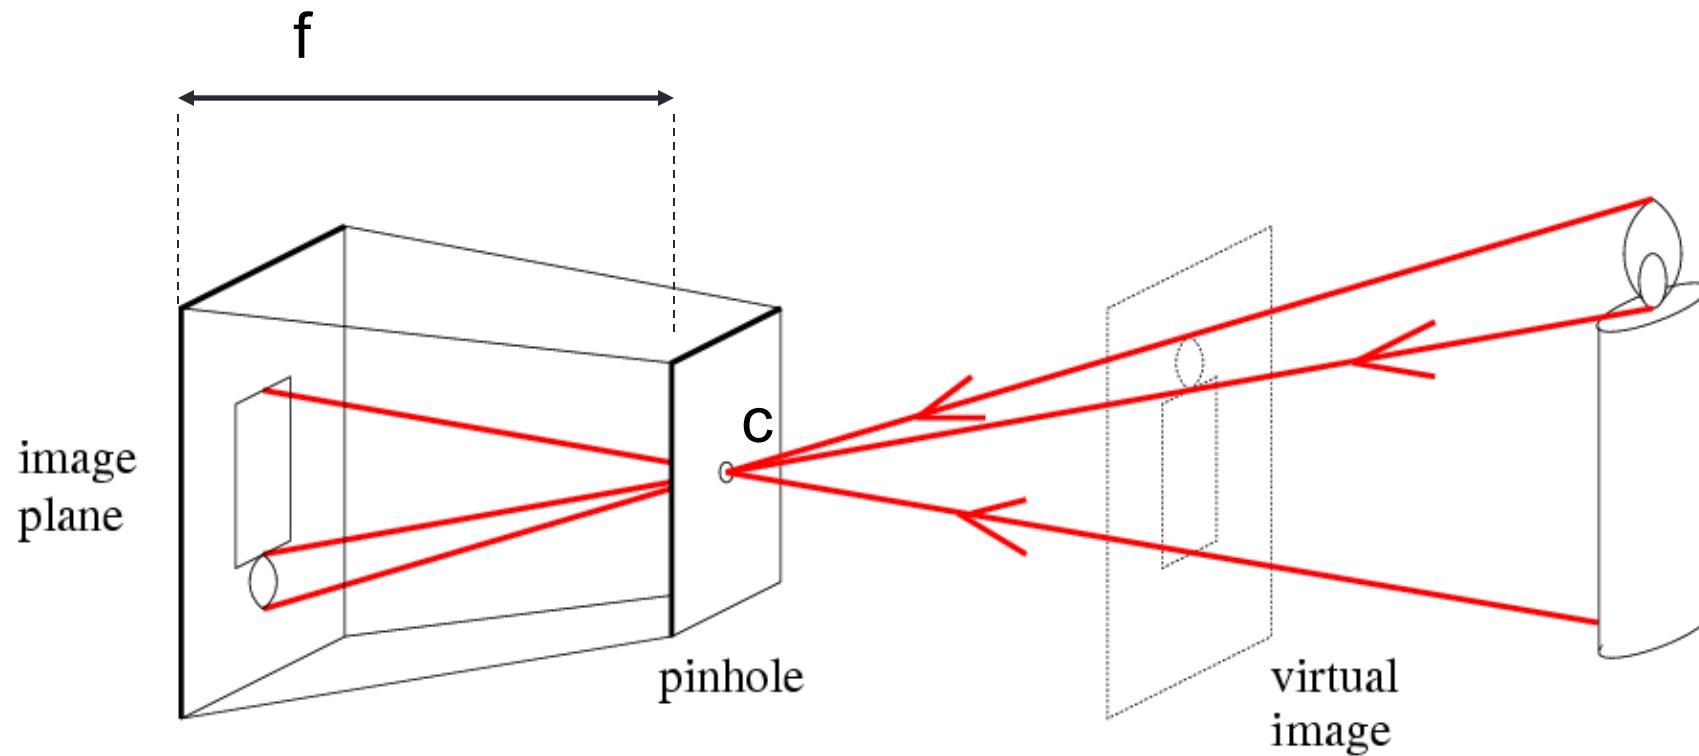
Pinhole camera



Idea 2: add a barrier to block off most of the rays

- This reduces blurring
- The opening known as the **aperture**

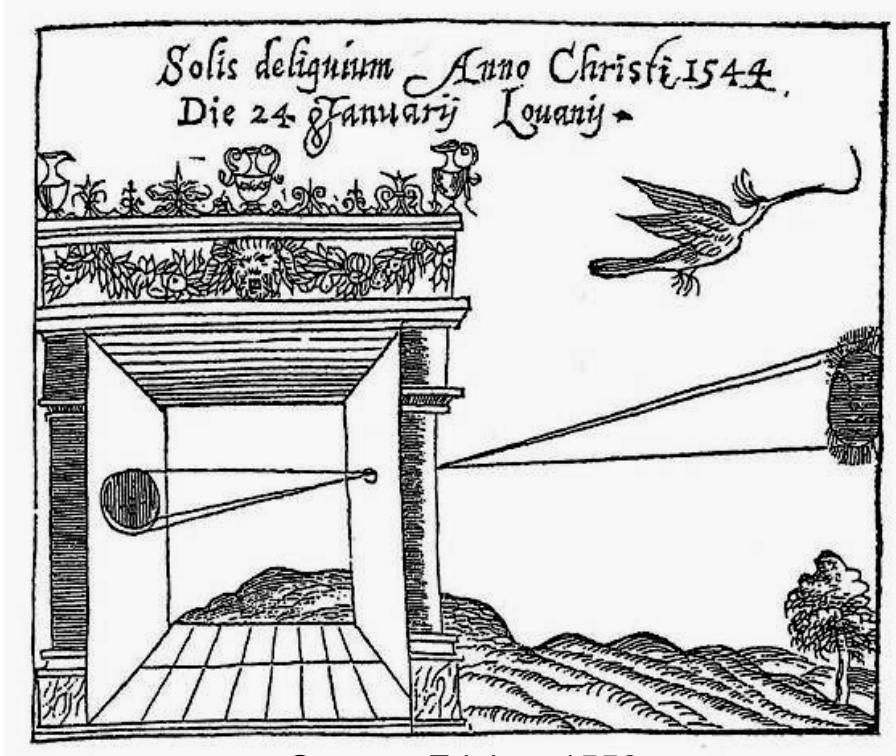
Pinhole camera



f = focal length

c = center of the camera

Camera Obscura



- Basic principle known to Mozi (470-390 BCE), Aristotle (384-322 BCE)
- Drawing aid for artists: described by Leonardo da Vinci (1452-1519)

Camera Obscura



After scouting rooms and reserving one for at least a day, Morell masks the windows except for the aperture. He controls three elements: the size of the hole, with a smaller one yielding a sharper but dimmer image; the length of the exposure, usually eight hours; and the distance from the hole to the surface on which the outside image falls and which he will photograph. He used 4 x 5 and 8 x 10 view cameras and lenses ranging from 75 to 150 mm.

After he's done inside, it gets harder. "I leave the room and I am constantly checking the weather, I'm hoping the maid reads my note not to come in, I'm worrying that the sun will hit the plastic masking and it will fall down, or that I didn't trigger the lens."

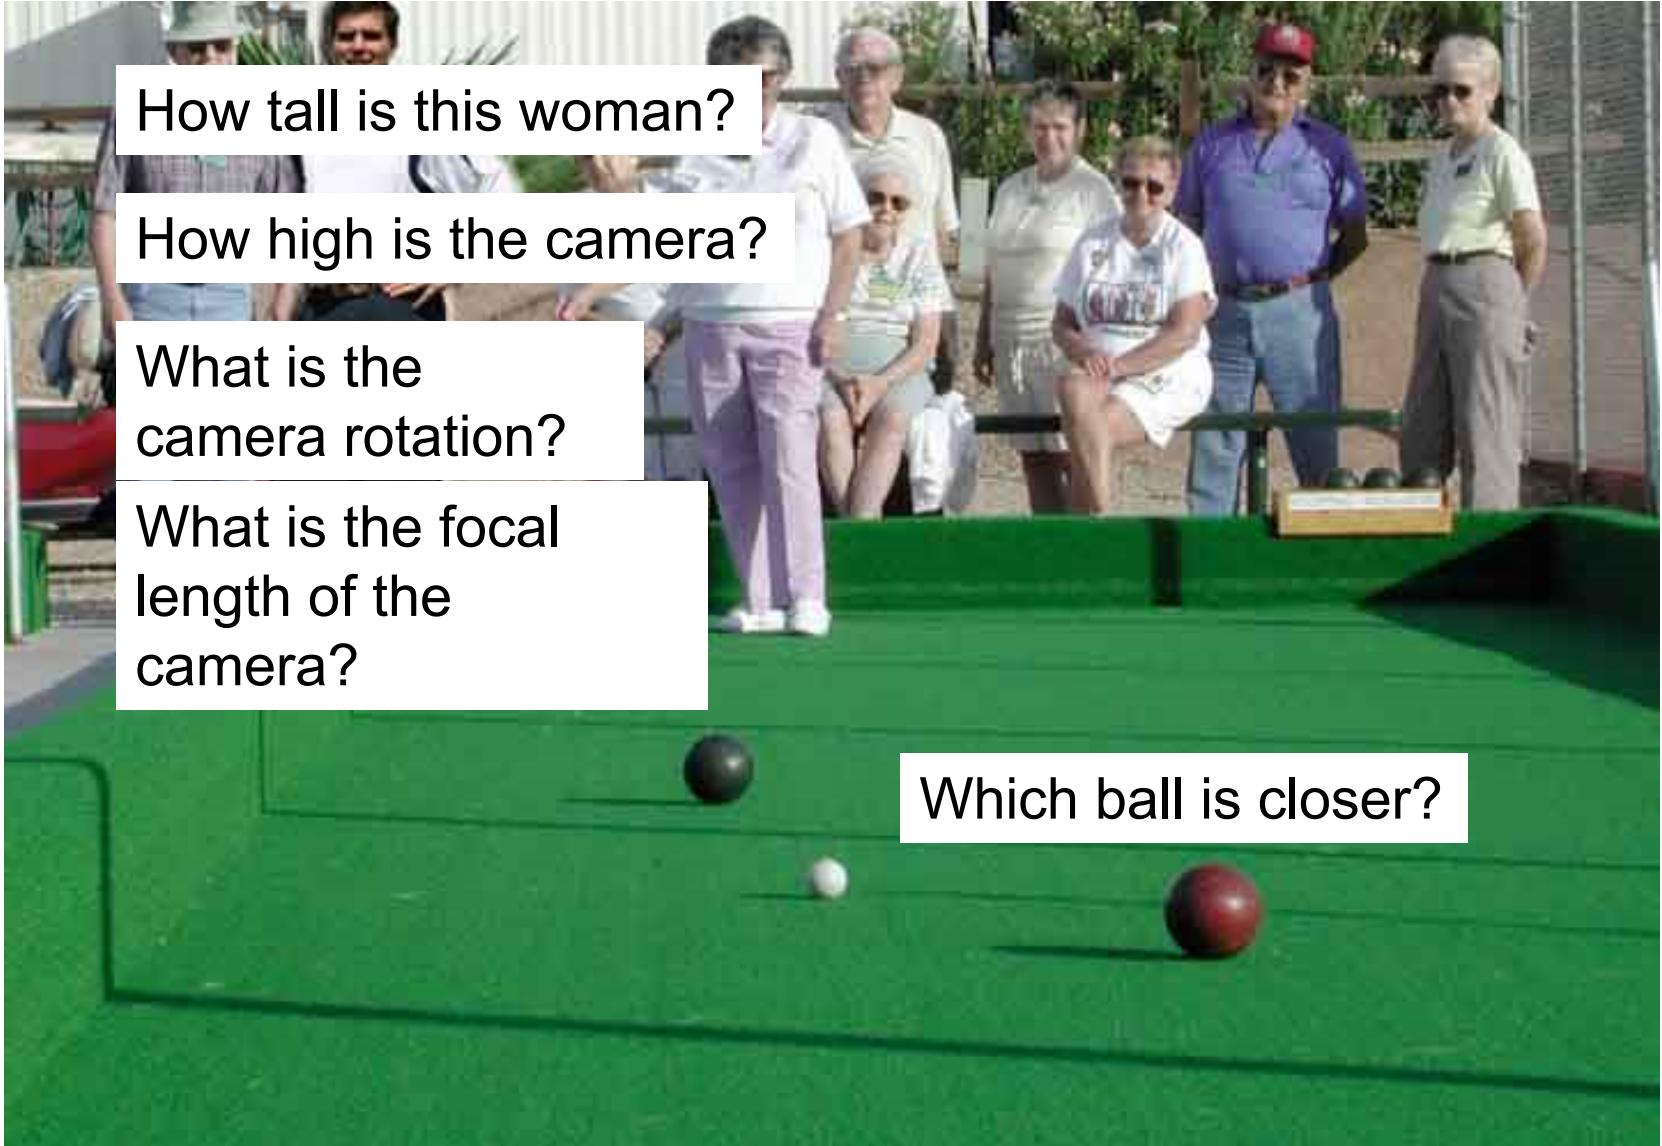
From *Grand Images Through a Tiny Opening*, **Photo District News**, February 2005

Abelardo Morell, Camera Obscura Image of Manhattan View Looking South in Large Room, 1996

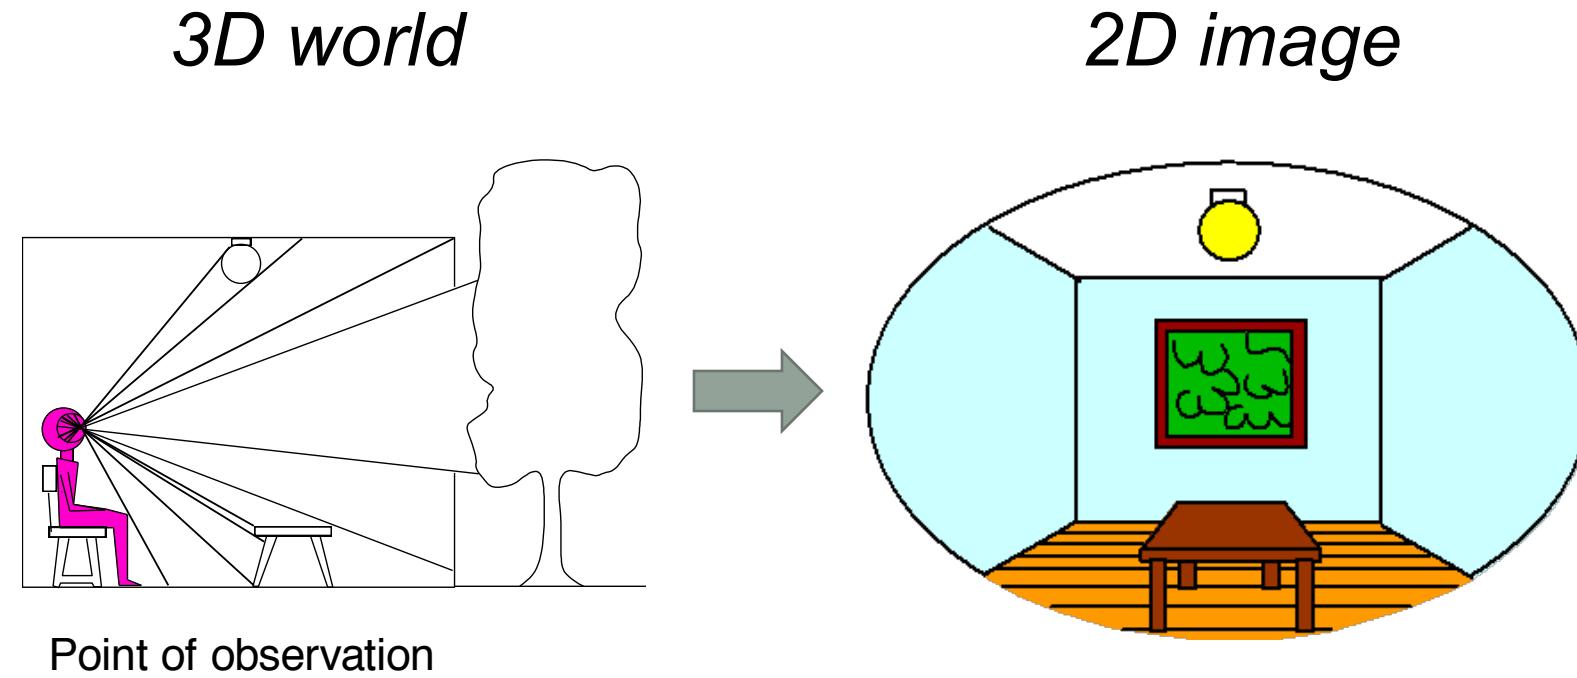


Abelardomorell.com

Camera and World Geometry



Dimensionality Reduction Machine (3D to 2D)



Projection can be tricky...



Projection can be tricky...

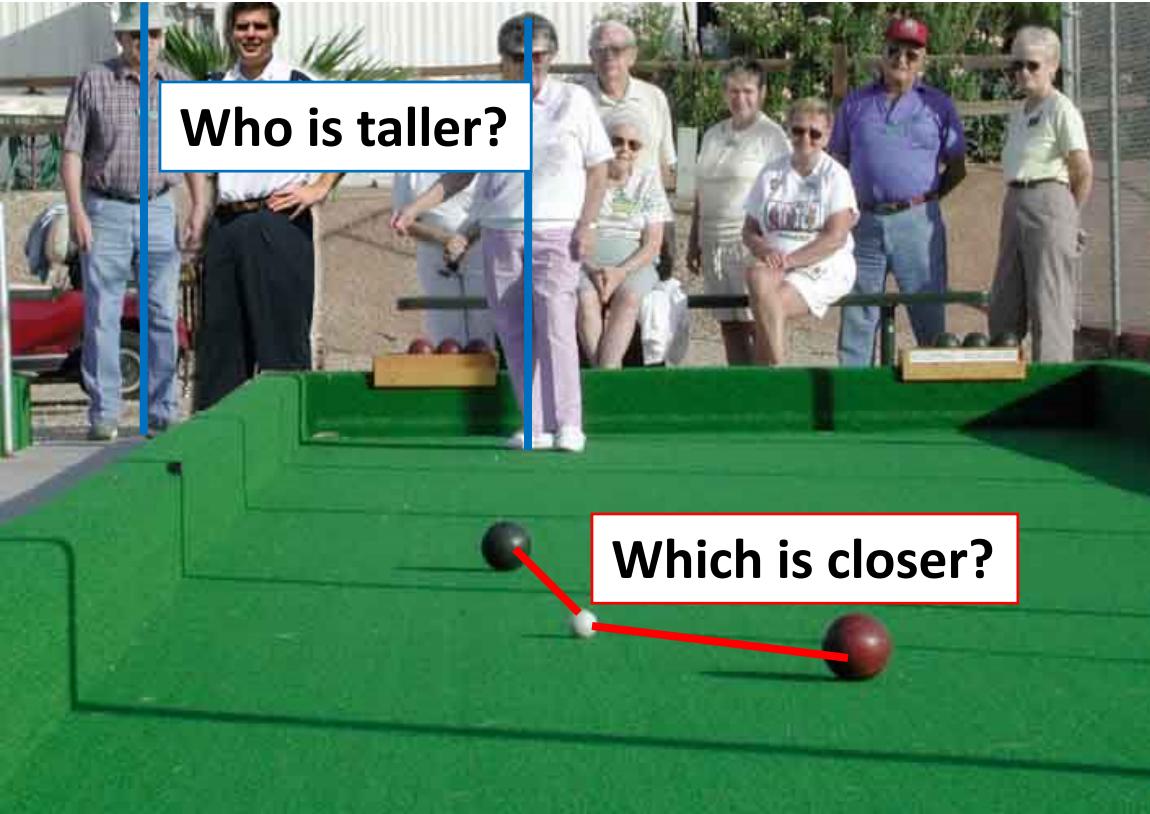


CoolOpticalIllusions.com

Projective Geometry

What is lost?

- Length



Length and area are not preserved

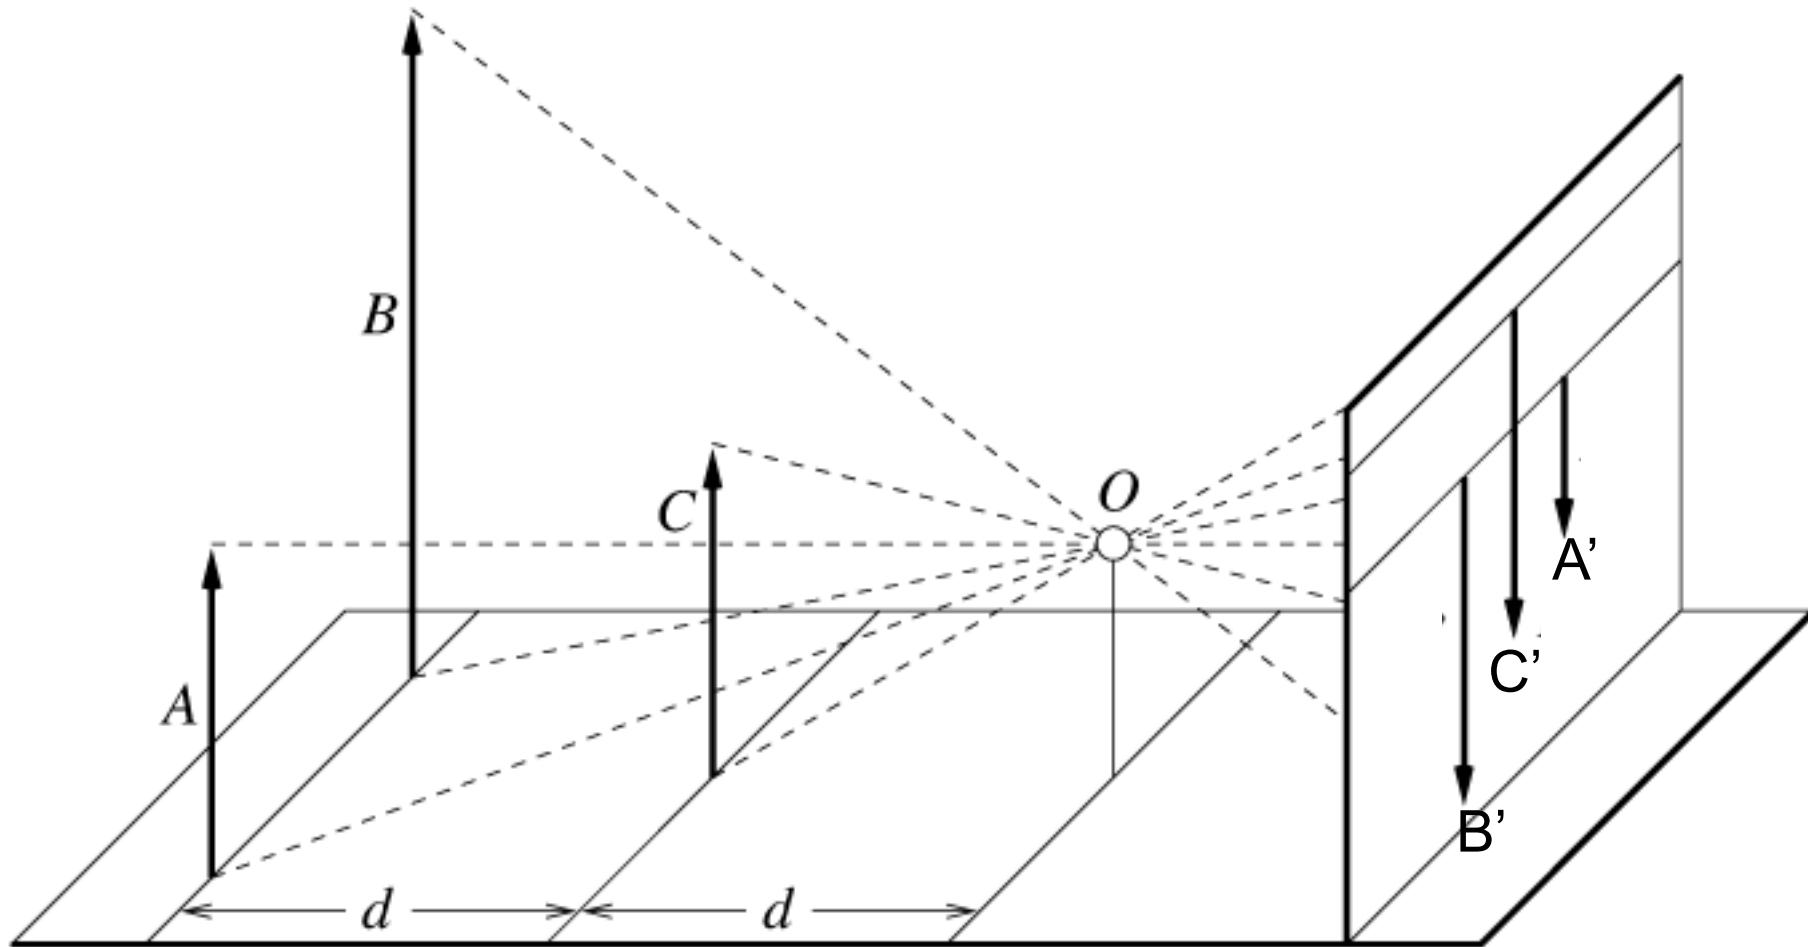
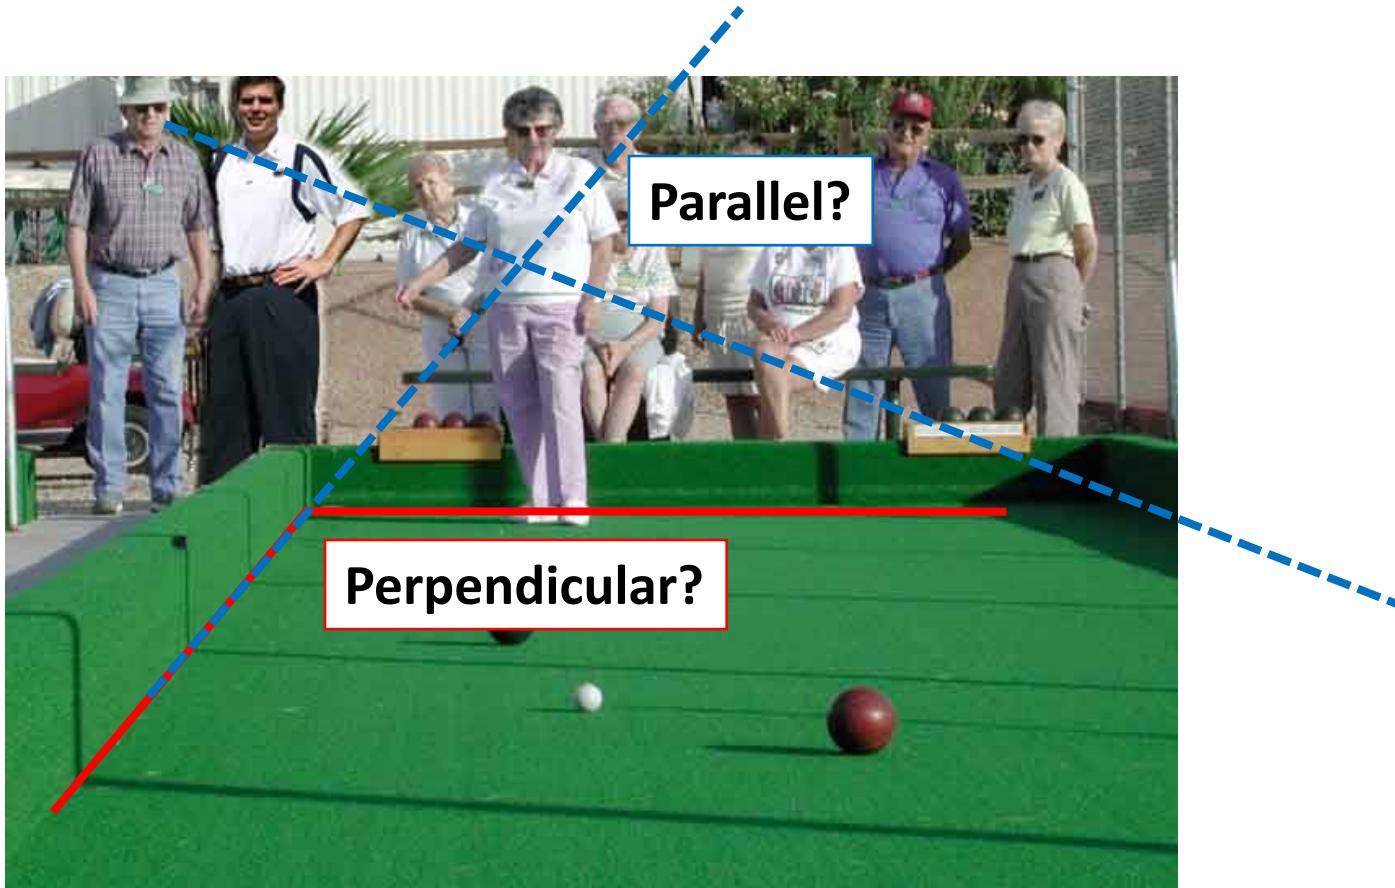


Figure by David Forsyth

Projective Geometry

What is lost?

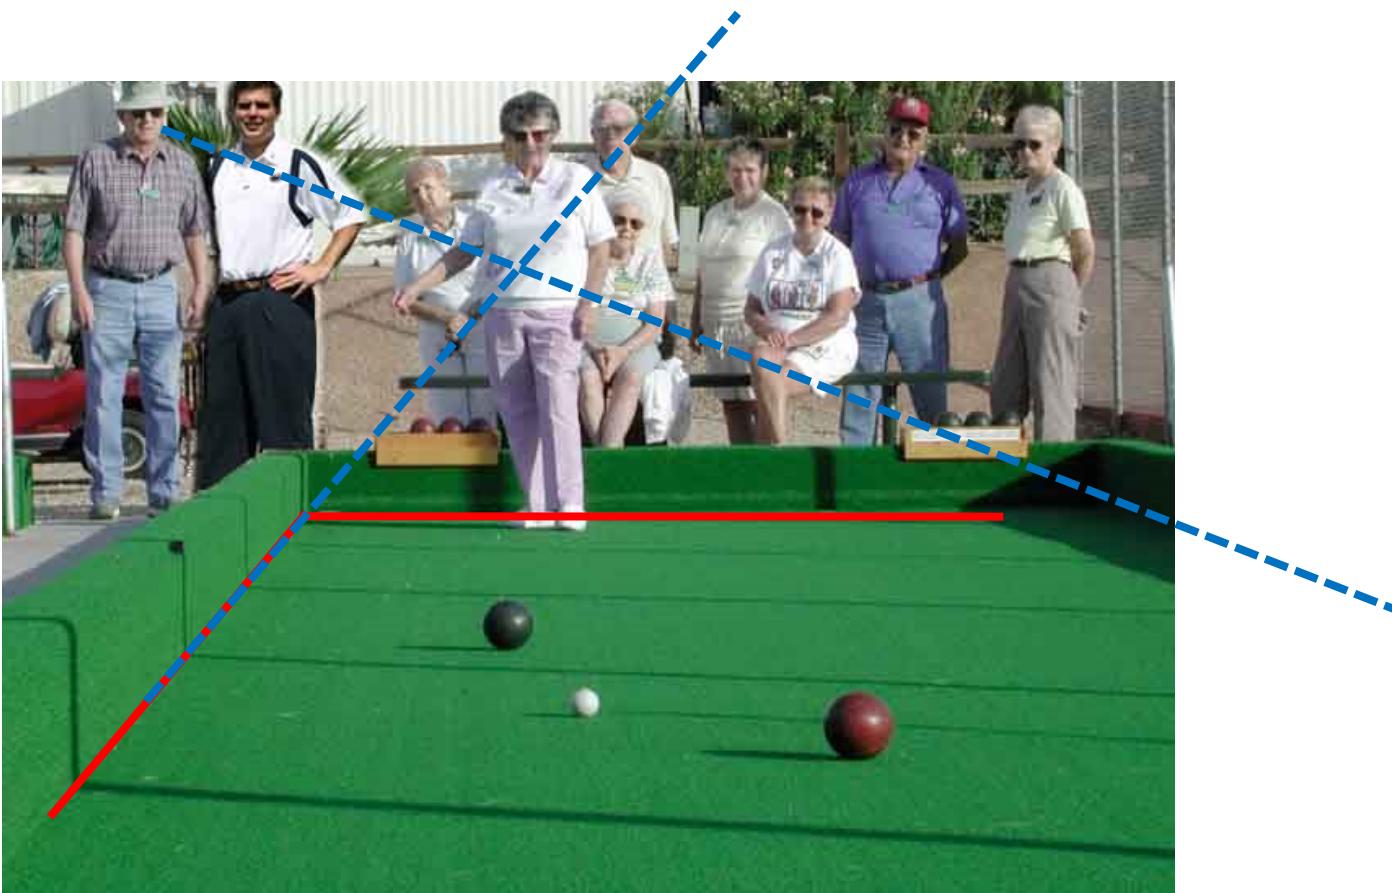
- Length
- Angles



Projective Geometry

What is preserved?

- Straight lines are still straight



Some Facts About Projection



3D lines project to 2D lines

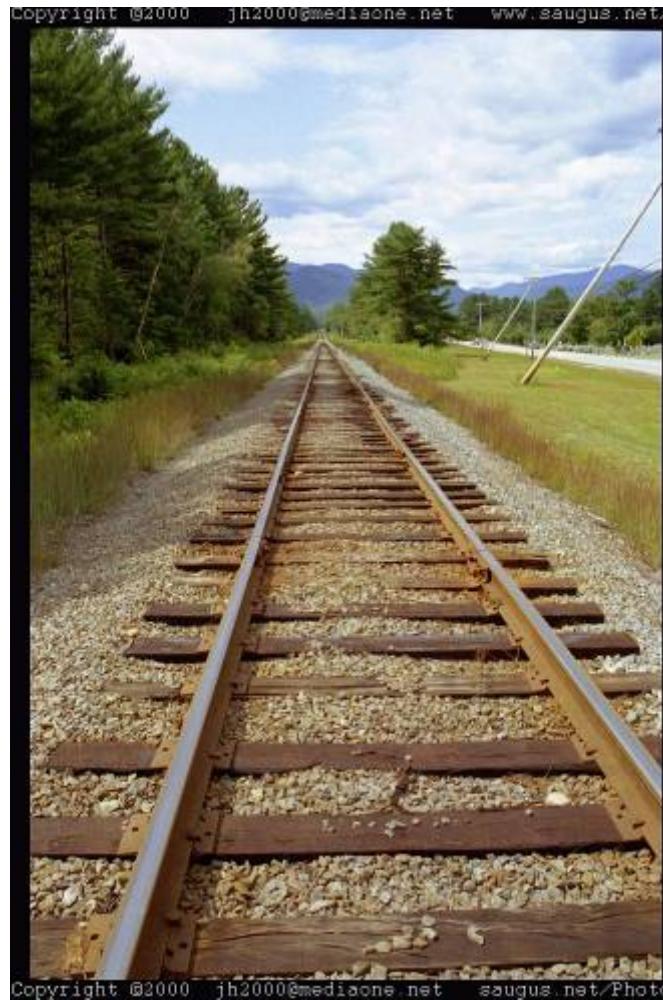
The projection of any 3D parallel lines converge at a vanishing point

Distant objects are smaller

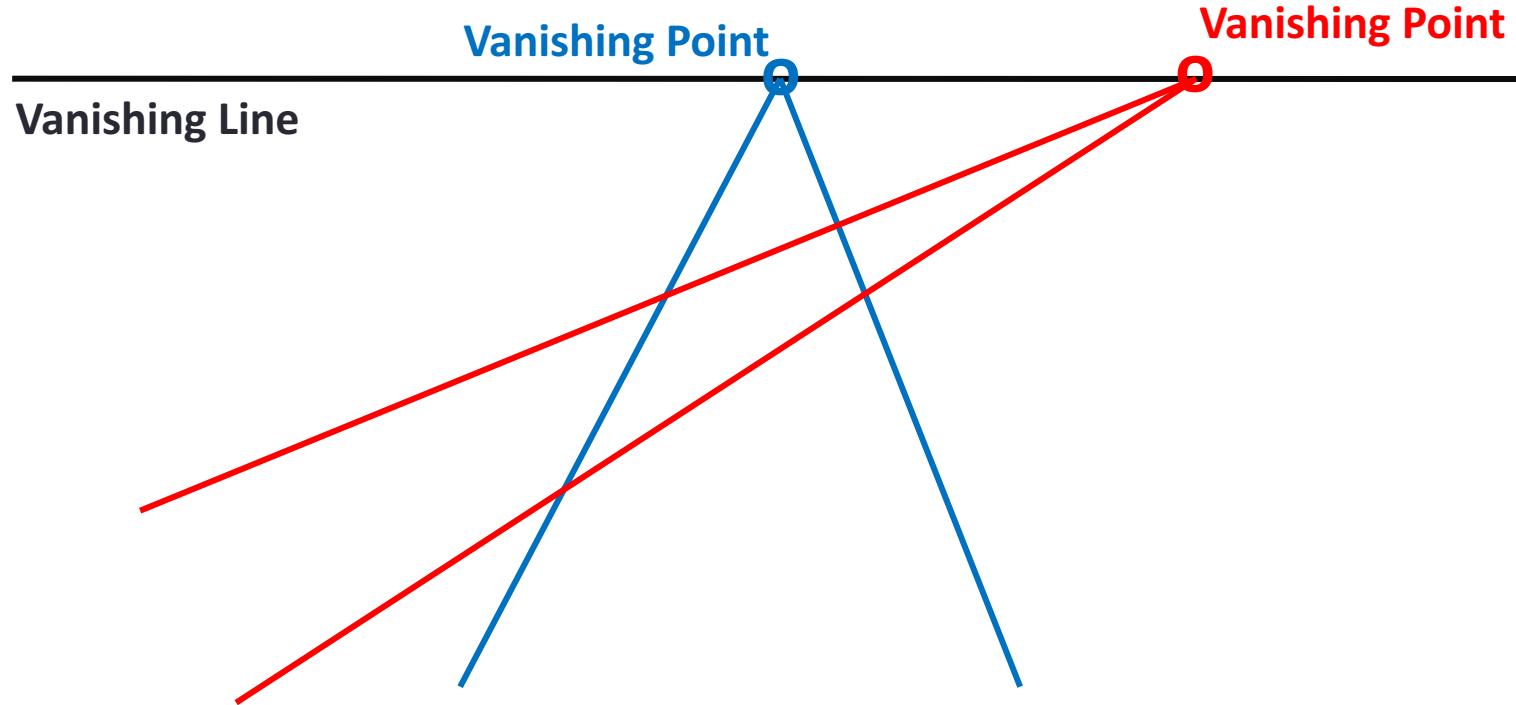


Vanishing points and lines

Parallel lines in the world intersect in the image at a “vanishing point”

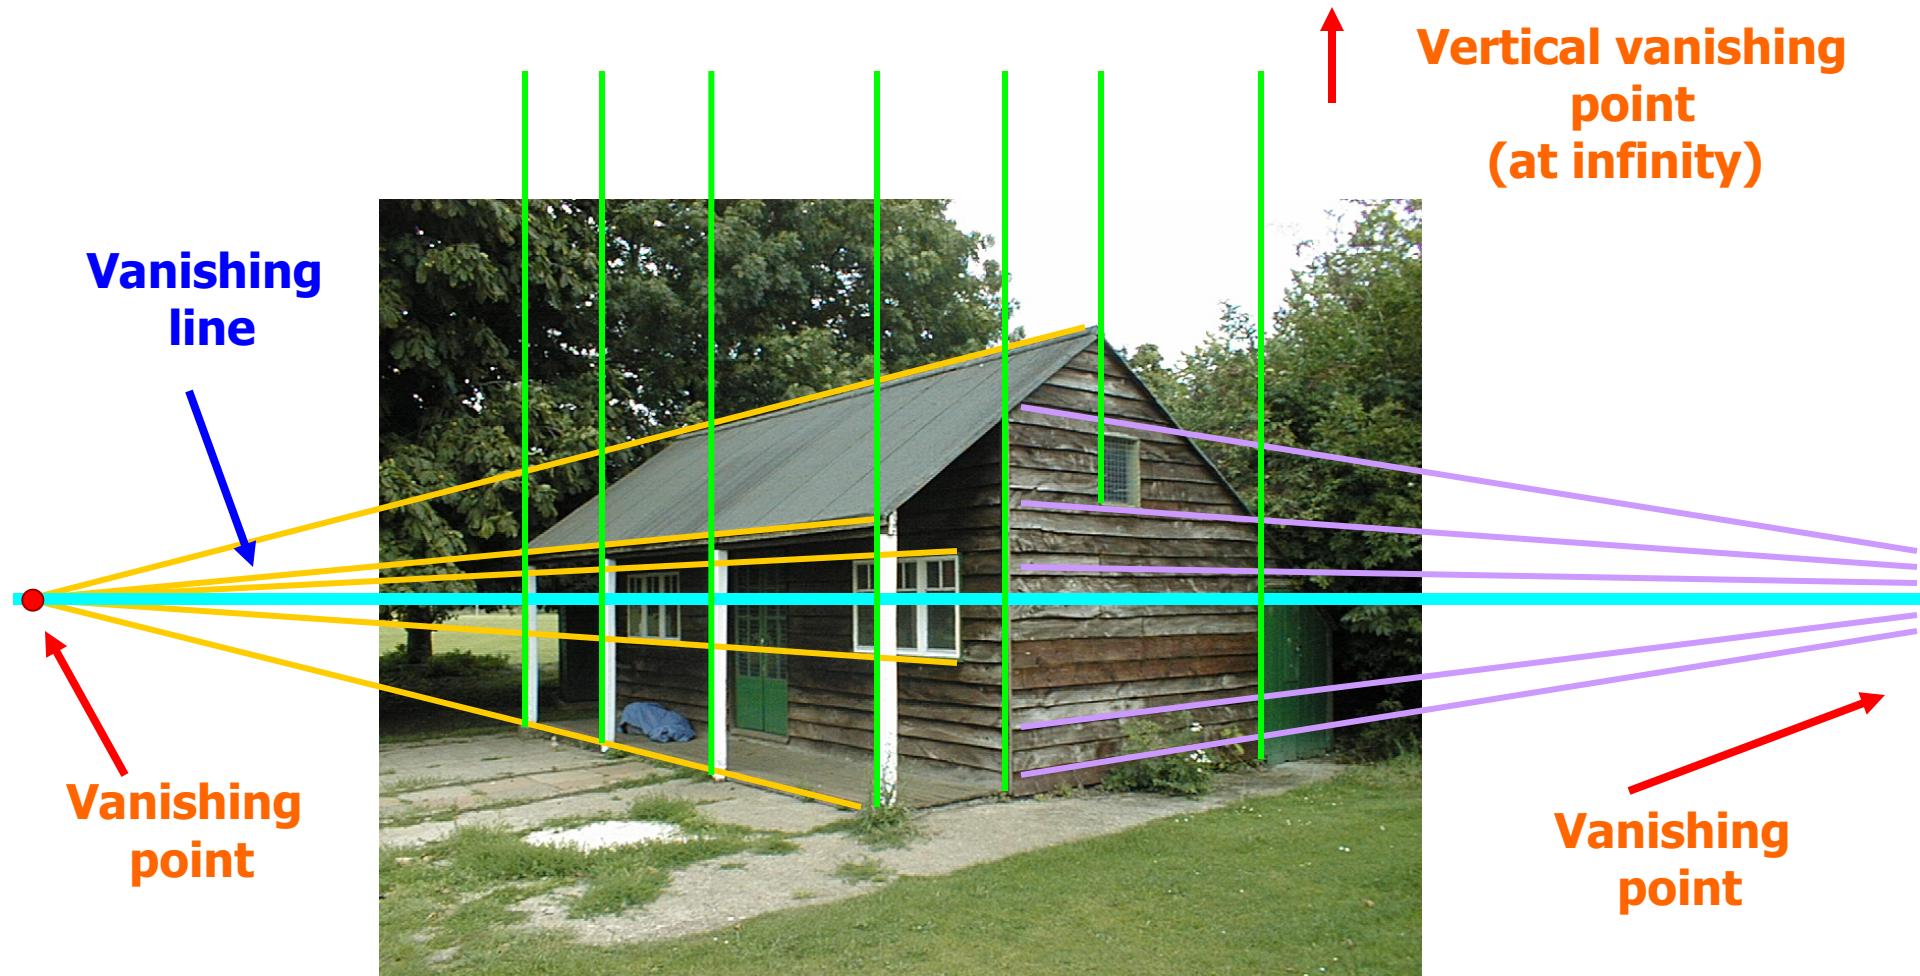


Vanishing points



- The projections of parallel 3D lines intersect at a **vanishing point**
- Vanishing point \leftrightarrow 3D direction of a line
- Not all lines that intersect are parallel

Vanishing points

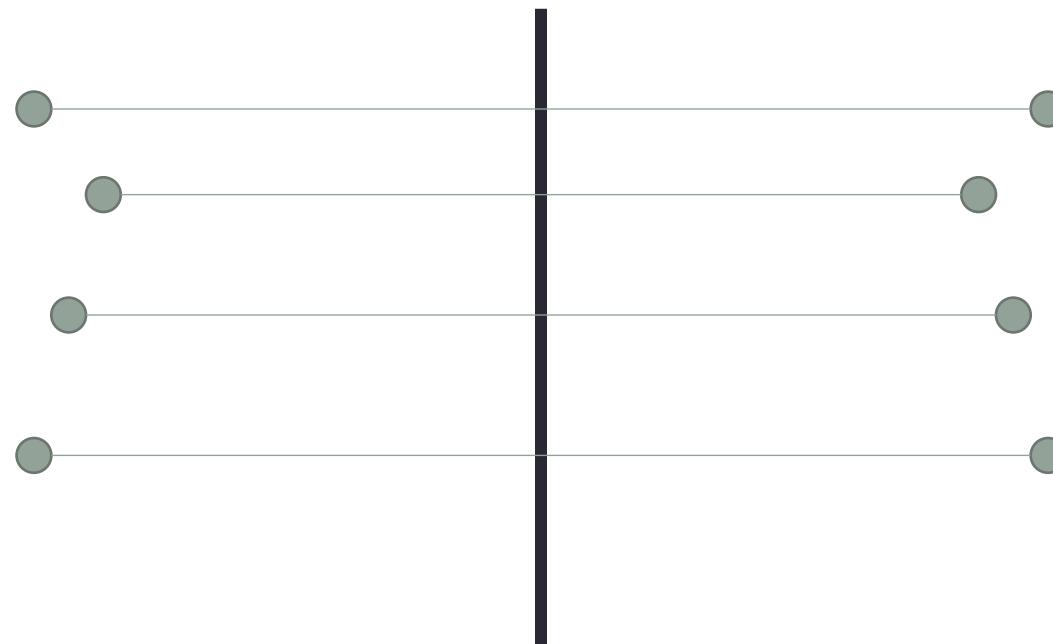


What's This Useful For?



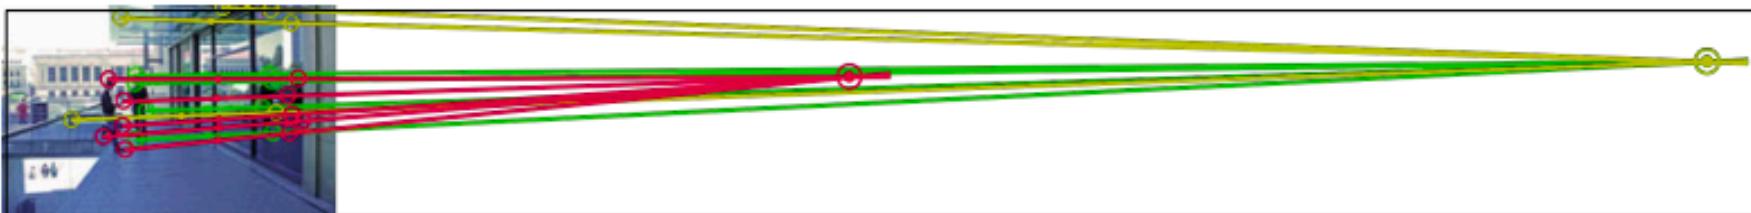
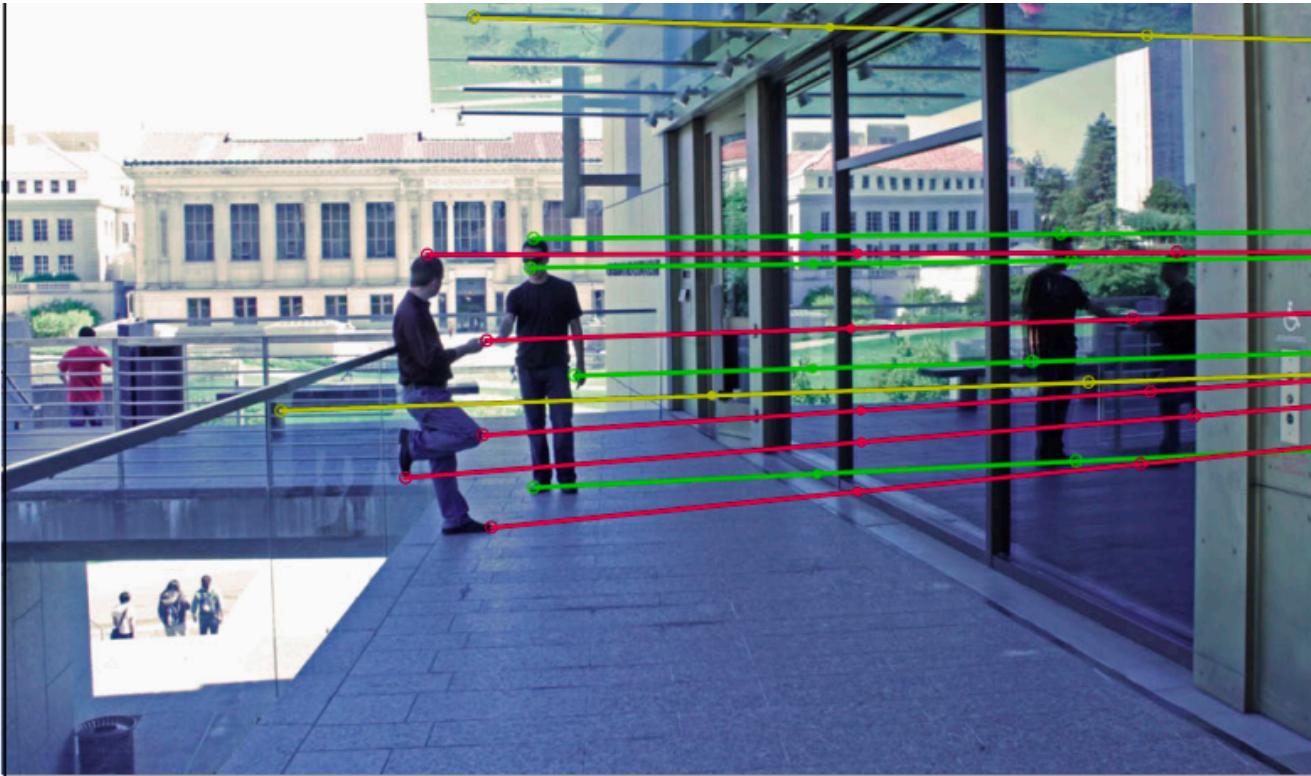
O'Brien and Farid, SIGGRAPH 2012, Exposing Photo Manipulation with Inconsistent Reflections

What's This Useful For?

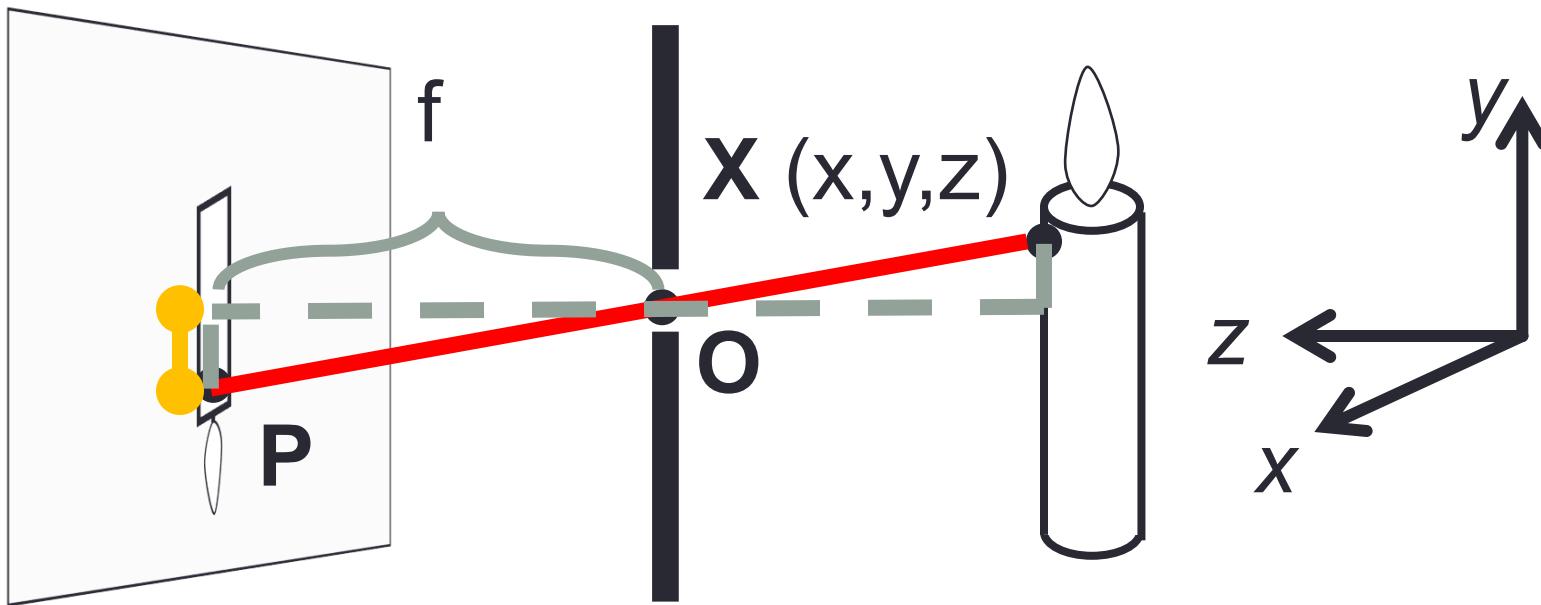


Mirror

What's This Useful For?

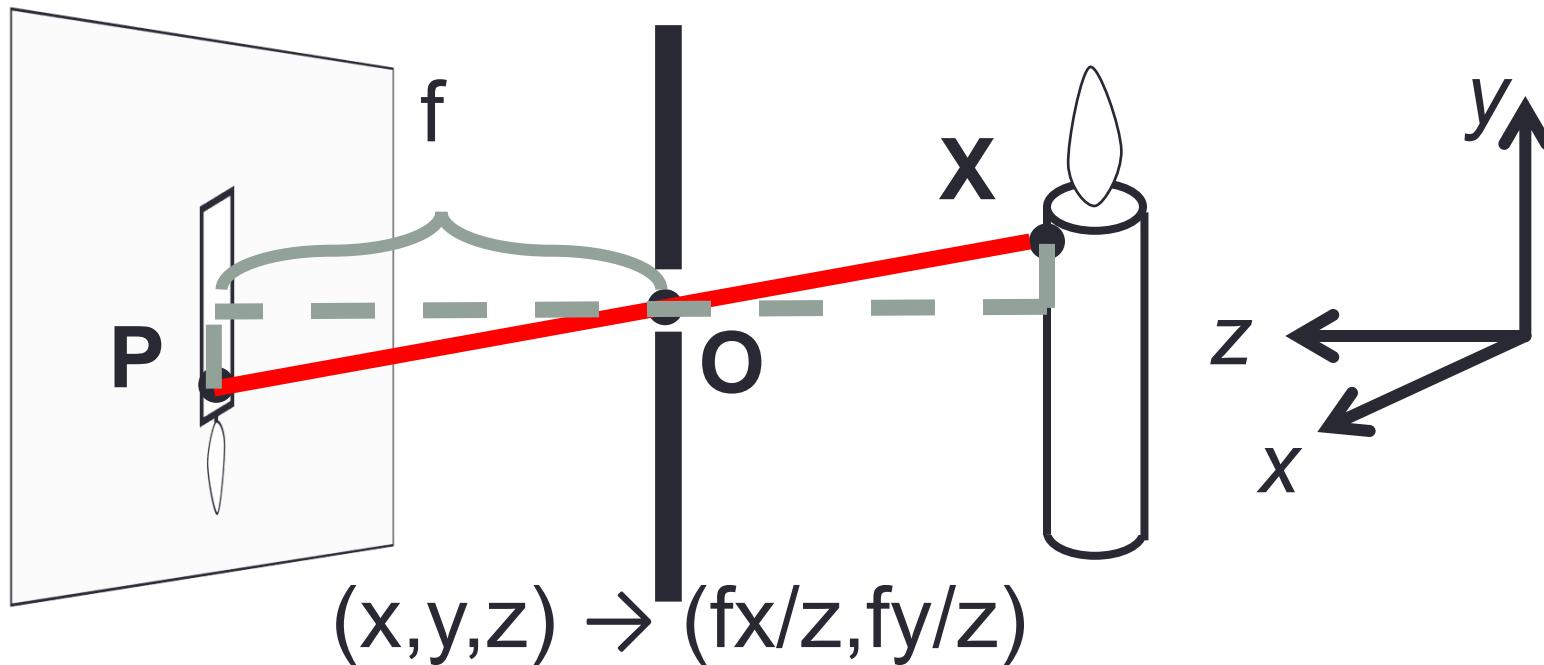


Projection Equations



Coordinate system: \mathbf{O} is origin, XY in image, Z sticks out.
 XY is image plane, Z is optical axis.
 (x, y, z) projects to $(fx/z, fy/z)$ via similar triangles

Projection Equation



is this linear?

Nope: division by z is non-linear
(and risks division by 0)

Homogeneous Coordinates (2D)

Trick: add a dimension!

This also clears up lots of nasty special cases

Physical
Point

$$\begin{bmatrix} x \\ y \end{bmatrix}$$

→
Concat
 $w=1$

Homogeneous
Point

$$\begin{bmatrix} u \\ v \\ w \end{bmatrix}$$

→
Divide
by w

Physical
Point

$$\begin{bmatrix} u/w \\ v/w \end{bmatrix}$$

What if $w = 0$?

Benefits of Homogeneous Coords

General equation of 2D line:

$$ax + by + c = 0$$

Homogeneous Coordinates

$$\mathbf{l}^T \mathbf{p} = 0, \quad \mathbf{l} = \begin{bmatrix} a \\ b \\ c \end{bmatrix}, \mathbf{p} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Homogeneous coordinates

Conversion

Converting to *homogeneous* coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image
coordinates

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

homogeneous scene
coordinates

Converting *from* homogeneous coordinates

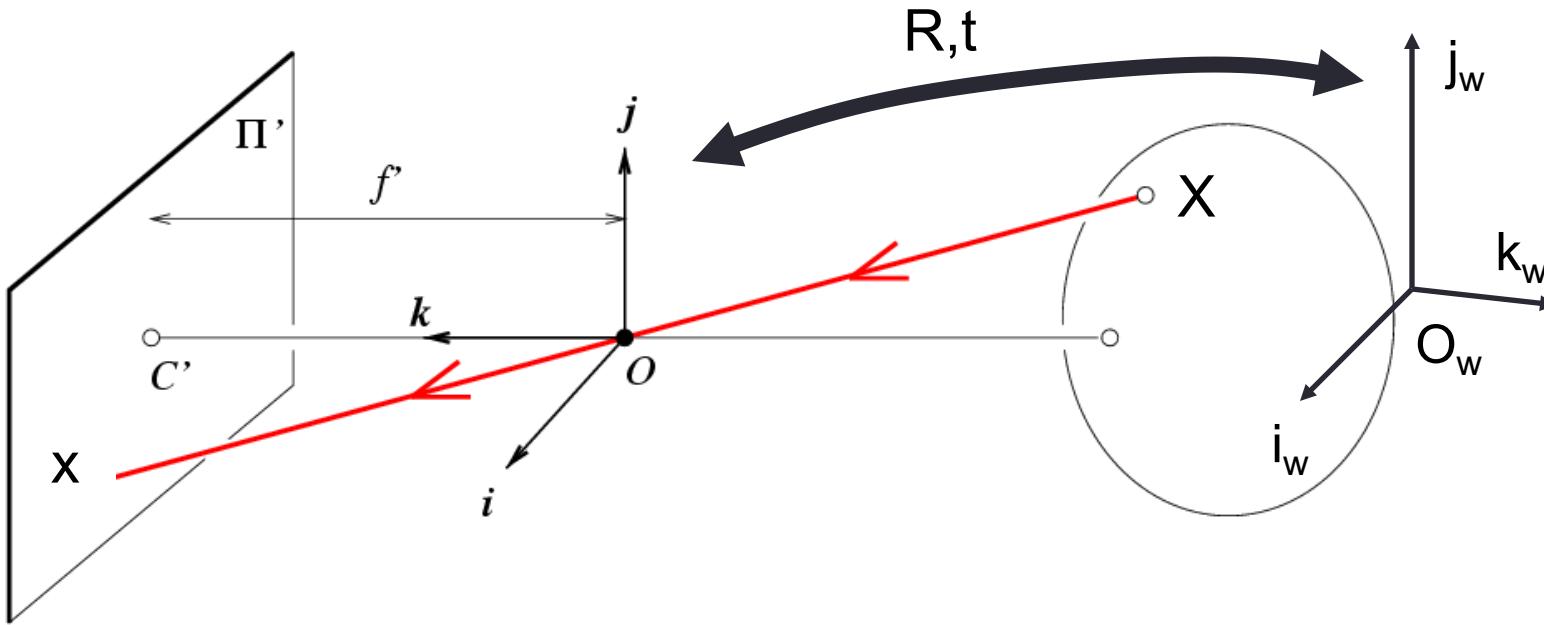
$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

$$\begin{bmatrix} x \\ y \\ z \\ w \end{bmatrix} \Rightarrow (x/w, y/w, z/w)$$

Benefits of Homogeneous Coords

- Lines (3D) and points (2D → 3D) are now the same dimension.
- Use the *cross* (\times) and *dot product* for:
 - Intersection of lines l and m : $l \times m$
 - Line through two points p and q : $p \times q$
 - Point p on line l : $l^T p$
- Parallel lines, vertical lines become easy (compared to $y=mx+b$)

Projection matrix



$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$

\mathbf{x} : Image Coordinates: $(u, v, 1)$

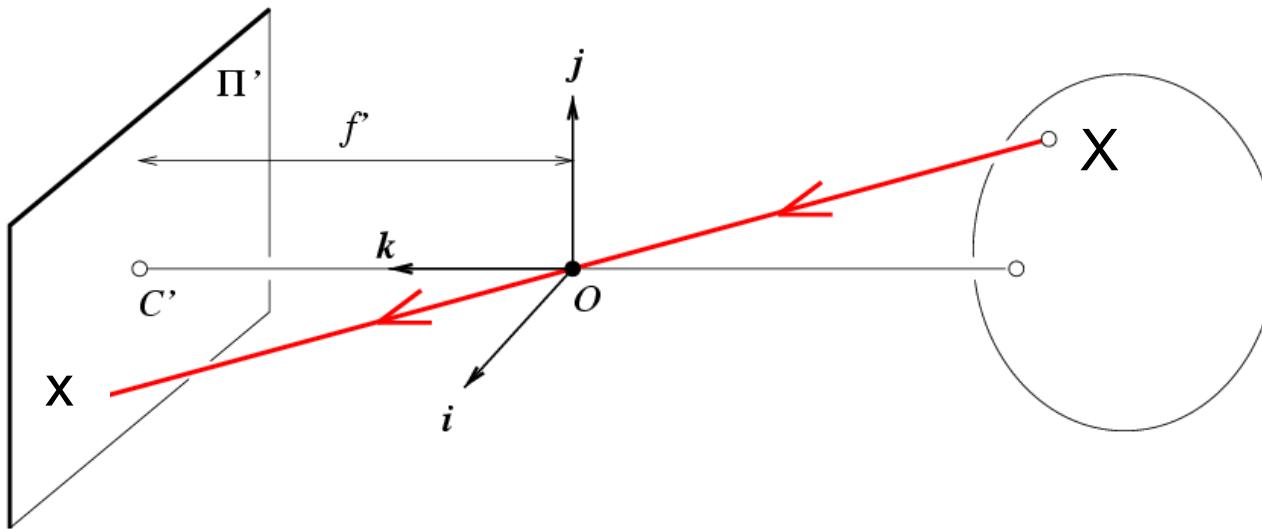
\mathbf{K} : Intrinsic Matrix (3x3)

\mathbf{R} : Rotation (3x3)

\mathbf{t} : Translation (3x1)

\mathbf{X} : World Coordinates: $(X, Y, Z, 1)$

Projection matrix



Intrinsic Assumptions

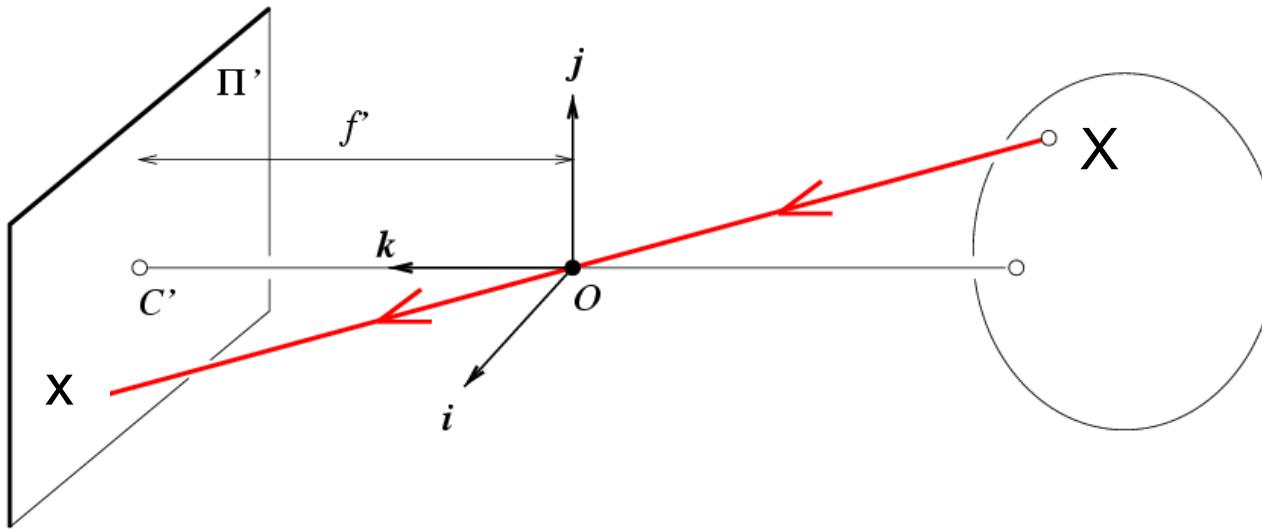
- Unit aspect ratio
- Optical center at $(0,0)$
- No skew

Extrinsic Assumptions

- No rotation
- Camera at $(0,0,0)$

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Projection matrix



Intrinsic Assumptions

- Unit aspect ratio
- Optical center at $(0,0)$
- No skew

Extrinsic Assumptions

- No rotation
- Camera at $(0,0,0)$

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Remove assumption: known optical center

Intrinsic Assumptions

- Unit aspect ratio
- No skew

Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & u_0 \\ 0 & f & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Remove assumption: square pixels

Intrinsic Assumptions

- No skew

Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Remove assumption: non-skewed pixels

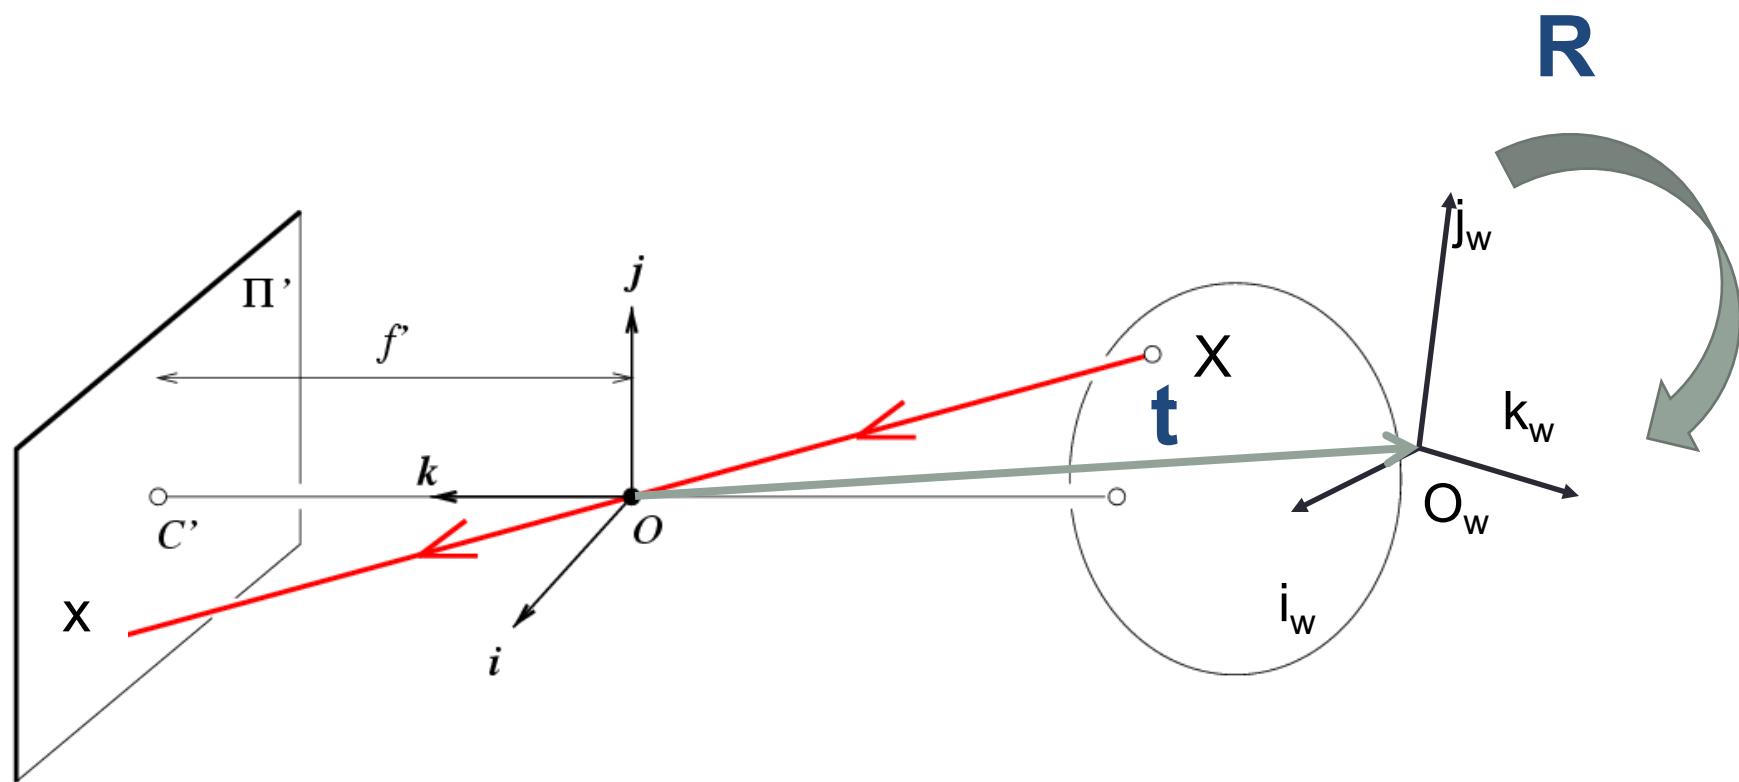
Intrinsic Assumptions Extrinsic Assumptions

- No rotation
- Camera at (0,0,0)

$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Note: different books use different notation for parameters

Oriented and Translated Camera



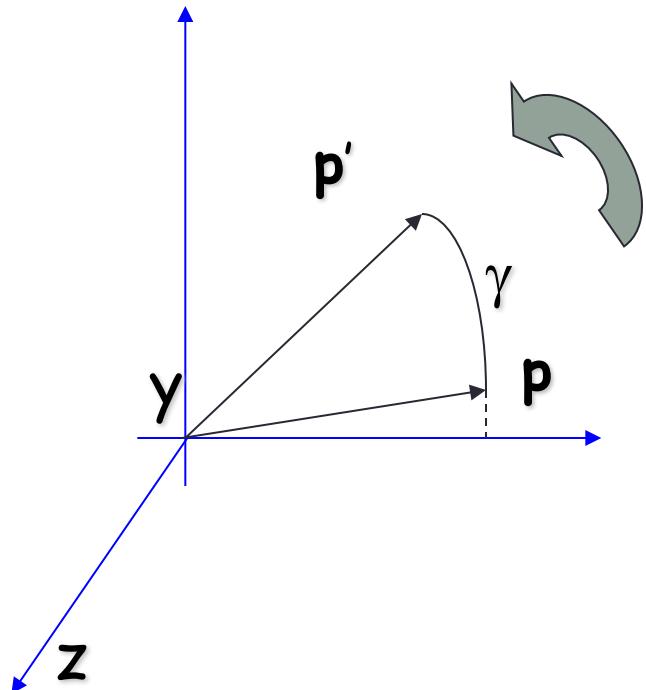
Allow camera translation

Intrinsic Assumptions Extrinsic Assumptions
• No rotation

$$\mathbf{x} = \mathbf{K}[\mathbf{I} \quad \mathbf{t}] \mathbf{X} \rightarrow w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & 0 & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

3D Rotation of Points

Rotation around the coordinate axes, counter-clockwise:



$$R_x(\alpha) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \alpha & -\sin \alpha \\ 0 & \sin \alpha & \cos \alpha \end{bmatrix}$$

$$R_y(\beta) = \begin{bmatrix} \cos \beta & 0 & \sin \beta \\ 0 & 1 & 0 \\ -\sin \beta & 0 & \cos \beta \end{bmatrix}$$

$$R_z(\gamma) = \begin{bmatrix} \cos \gamma & -\sin \gamma & 0 \\ \sin \gamma & \cos \gamma & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Allow camera rotation

$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$



$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Degrees of freedom

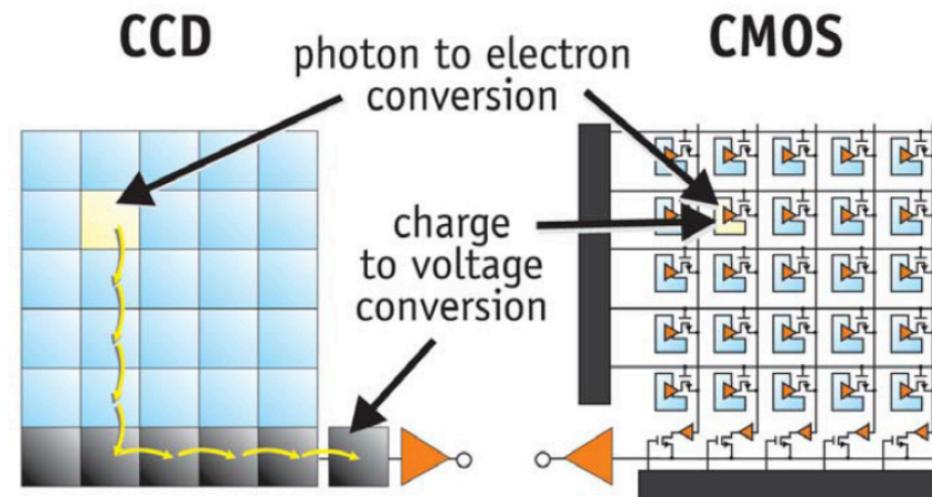
$$\mathbf{x} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{X}$$



$$w \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} \alpha & s & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

5 6

From Photon to Photo

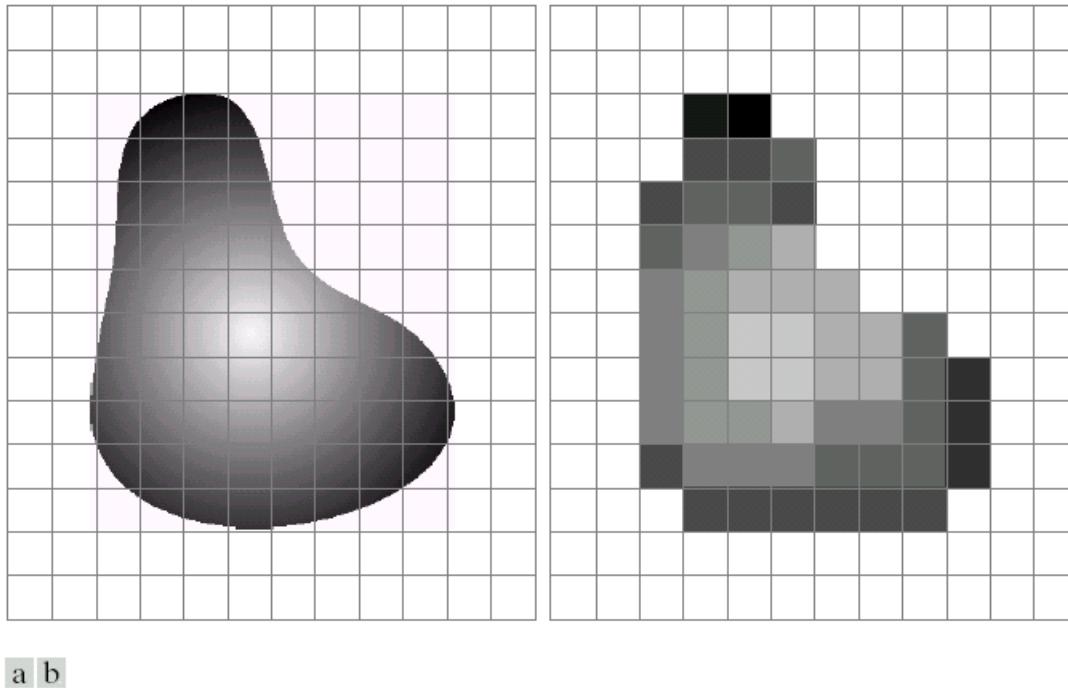


CCDs move photogenerated charge from pixel to pixel and convert it to voltage at an output node. CMOS imagers convert charge to voltage inside each pixel.

A digital camera replaces film with a sensor array

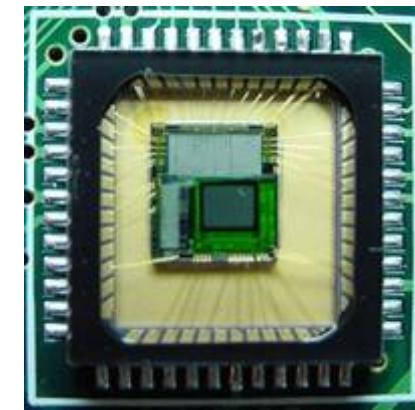
- Each cell in the array is light-sensitive diode that converts photons to electrons
- Two common types
 - **Charge Coupled Device (CCD)**
 - **Complementary metal oxide semiconductor (CMOS)**
- <http://electronics.howstuffworks.com/digital-camera.htm>

Digital Images



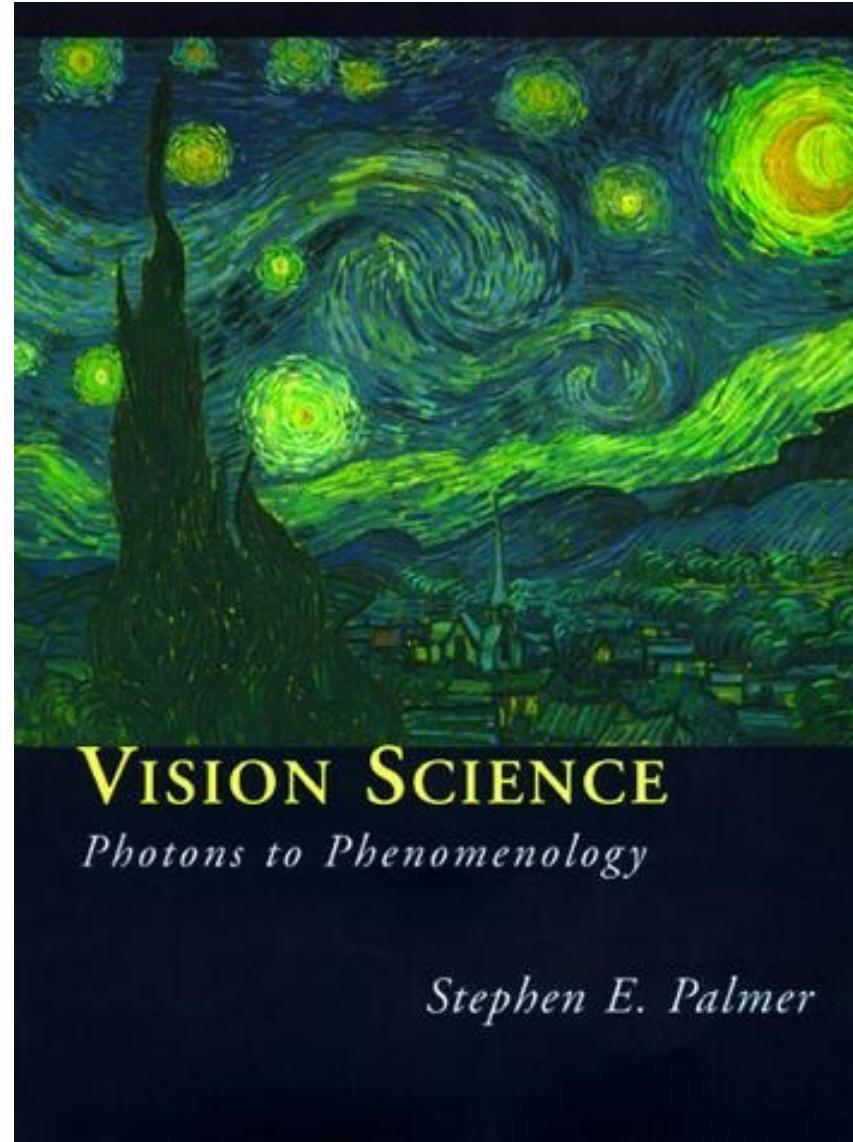
a b

FIGURE 2.17 (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.



What is color?

- The result of interaction between physical light in the environment and our visual system.
- A *psychological property* of our visual experiences when we look at objects and lights, *not a physical property* of those objects or lights.

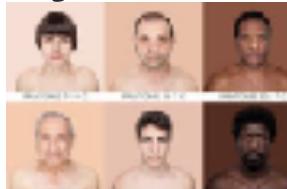


Why does a visual system need color?

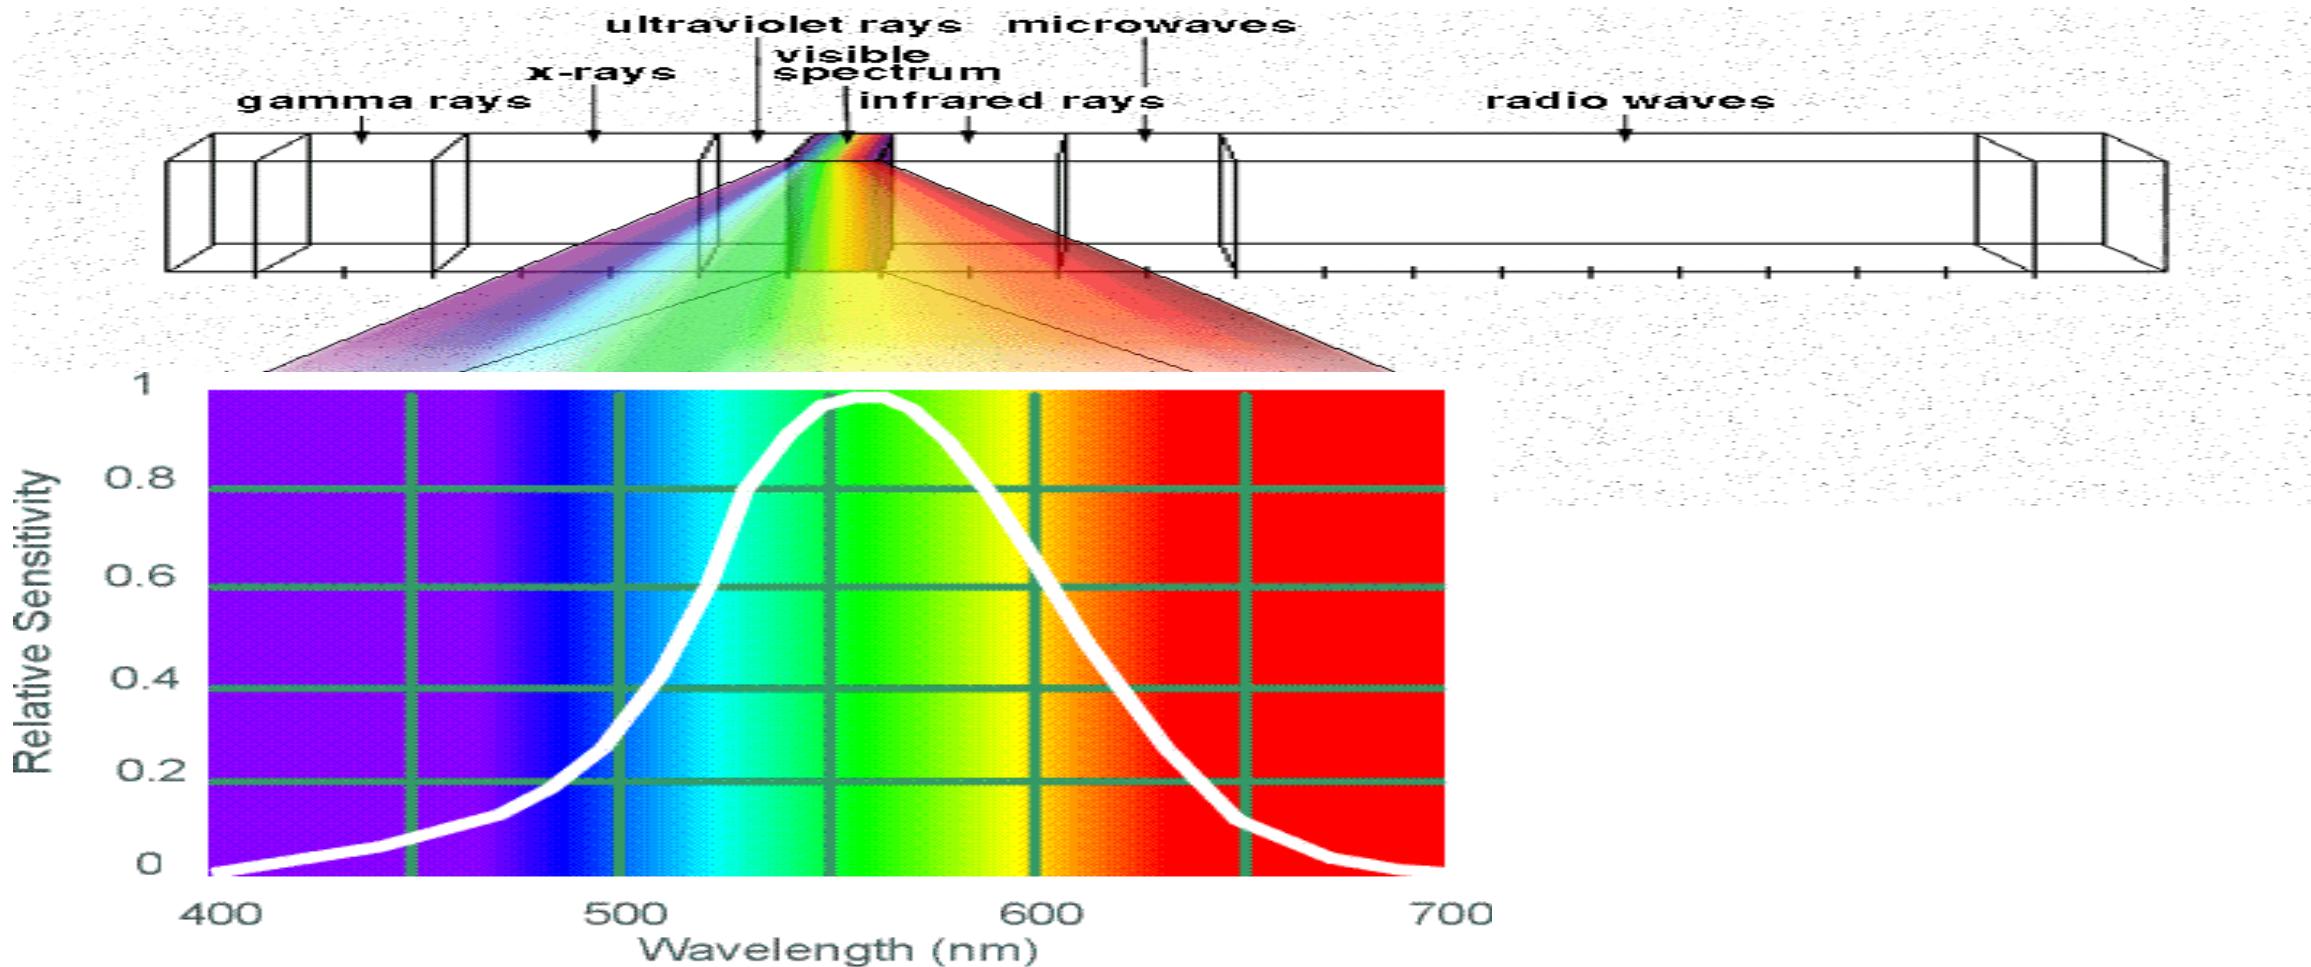


An incomplete list

- To tell what food is edible.
- To distinguish material changes from shading changes.
- To group parts of one object together in a scene.
- To find people's skin.
- Check whether a person's appearance looks normal/healthy.



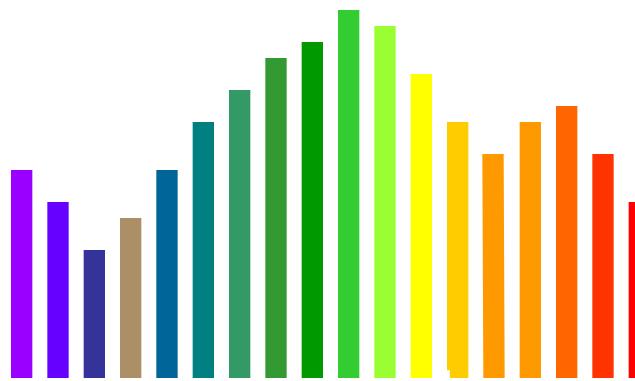
Electromagnetic Spectrum



Human Luminance Sensitivity Function

The Physics of Light

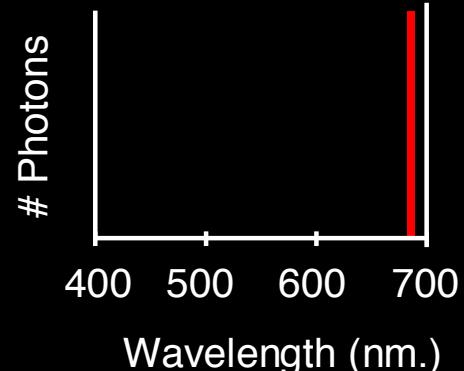
Any patch of light can be completely described physically by its spectrum: the number of photons (per time unit) at each wavelength 400 - 700 nm.



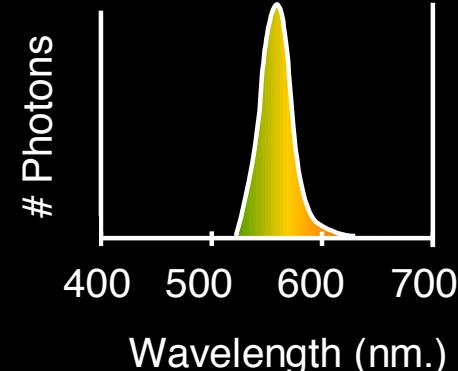
The Physics of Light

Some examples of the spectra of light sources

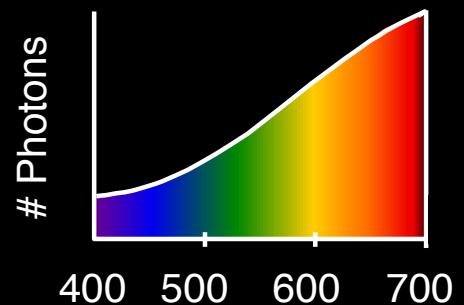
A. Ruby Laser



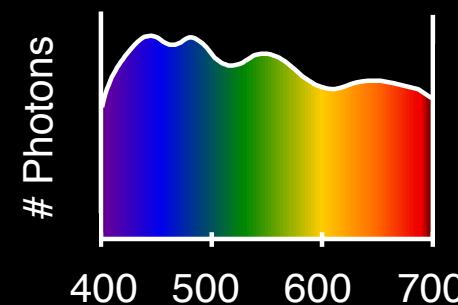
B. Gallium Phosphide Crystal



C. Tungsten Lightbulb

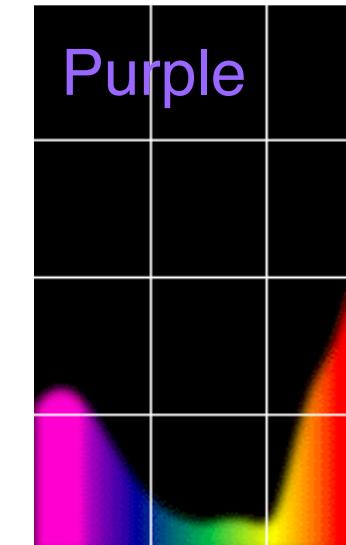
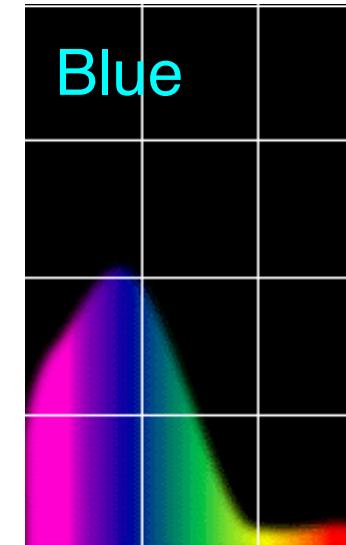
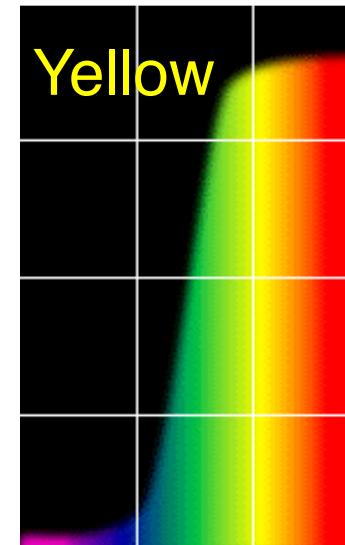
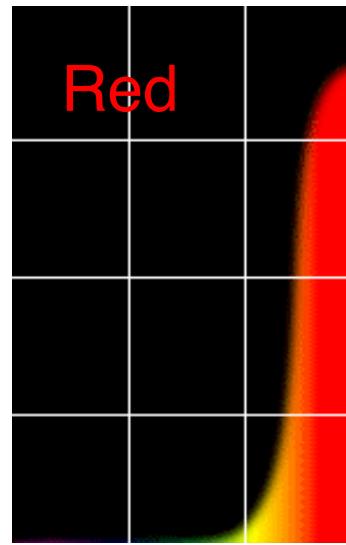


D. Normal Daylight



The Physics of Light

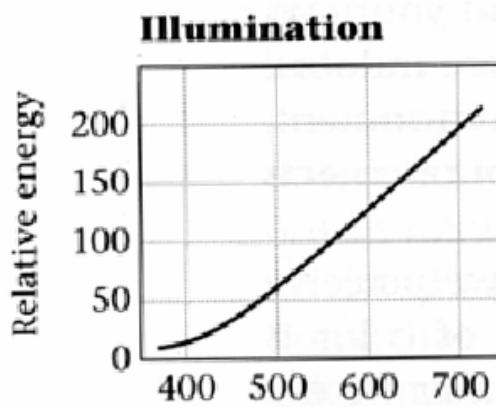
Some examples of the reflectance spectra of surfaces



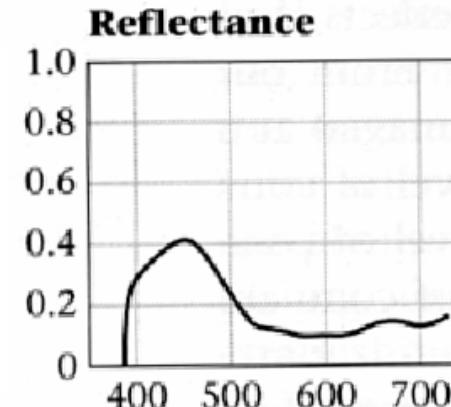
Interaction of light and surfaces



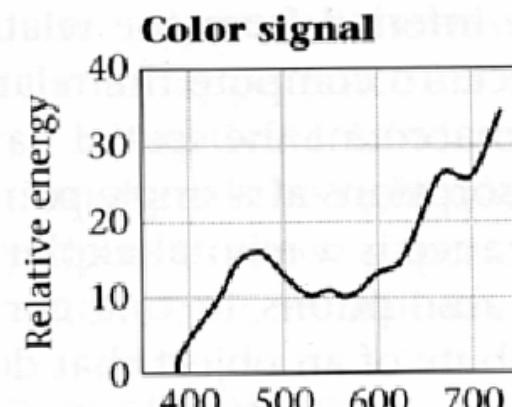
- Reflected color is the result of interaction of light source spectrum with surface reflectance
- Spectral radiometry
 - All definitions and units are now “per unit wavelength”
 - All terms are now “spectral”



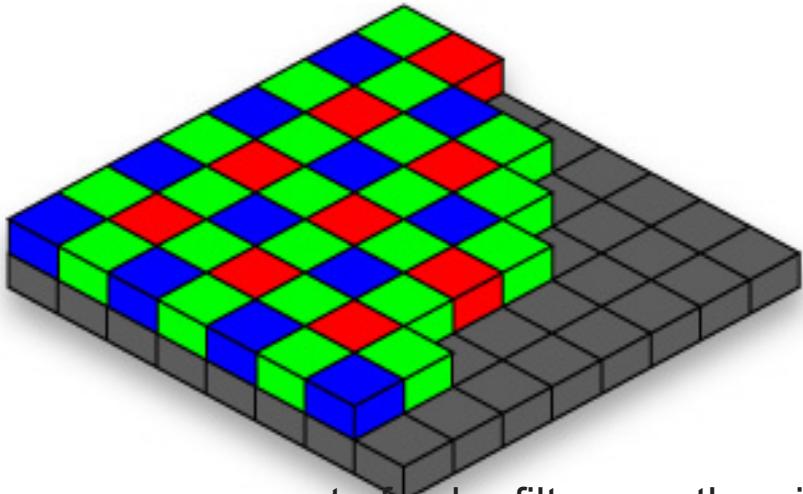
• *



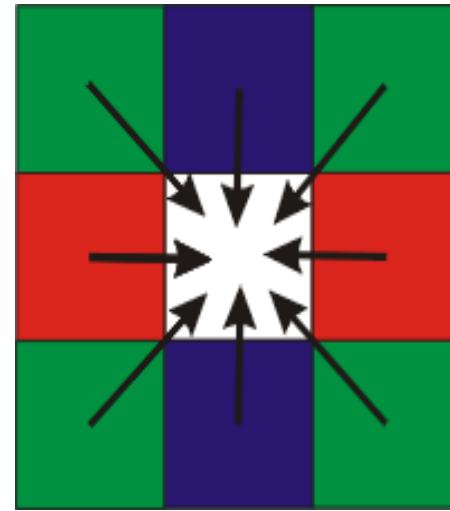
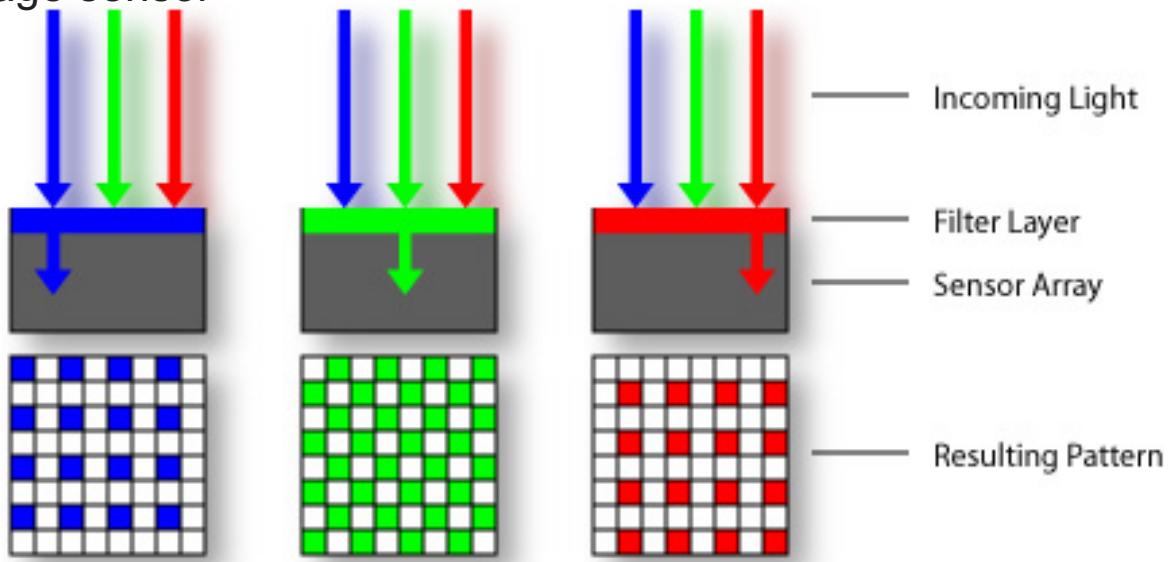
=



Practical Color Sensing: Bayer Grid

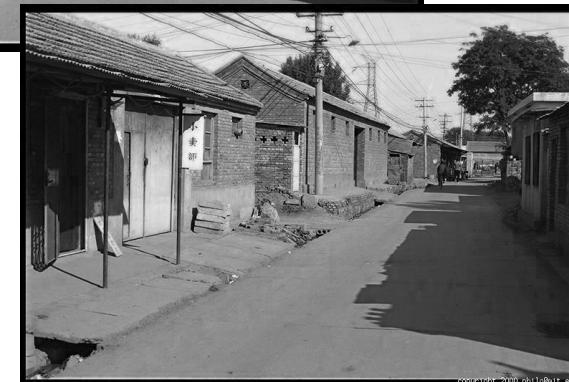
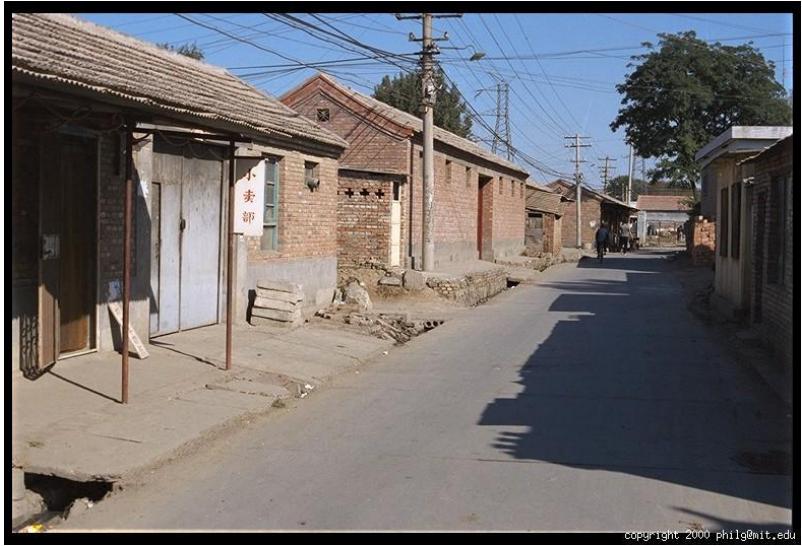


The Bayer arrangement of color filters on the pixel array of an image sensor



- Estimate RGB at 'G' cells from neighboring values

Color Image



Images in Matlab

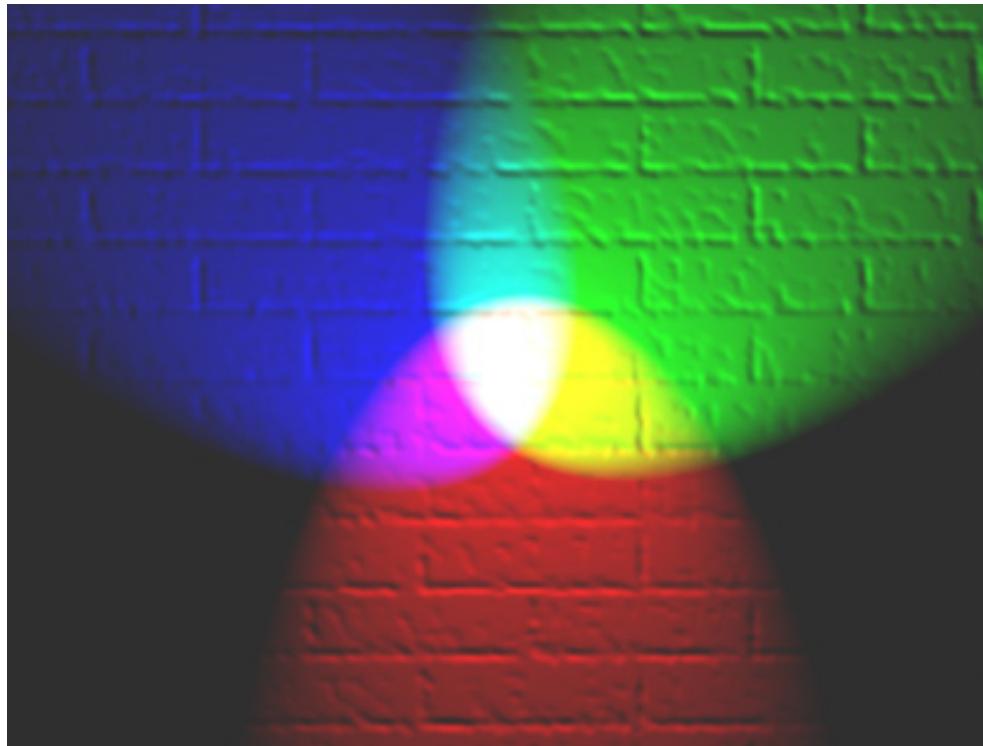
- Images represented as a matrix
- Suppose we have a NxM RGB image called “im”
 - $\text{im}(1,1,1)$ = top-left pixel value in R-channel
 - $\text{im}(y, x, b)$ = y pixels down, x pixels to right in the bth channel
 - $\text{im}(N, M, 3)$ = bottom-right pixel in B-channel

The diagram illustrates a 12x12 matrix representing an RGB image. A vertical arrow labeled "row" points downwards, and a horizontal arrow labeled "column" points to the right. The matrix is divided into three 4x4 sub-matrices labeled R (Red), G (Green), and B (Blue) with red borders. The matrix values are as follows:

0.92	0.93	0.94	0.97	0.62	0.37	0.85	0.97	0.93	0.92	0.99	
0.95	0.89	0.82	0.89	0.56	0.31	0.75	0.92	0.81	0.95	0.91	
0.89	0.72	0.51	0.55	0.51	0.42	0.57	0.41	0.49	0.91	0.92	
0.96	0.95	0.88	0.94	0.56	0.46	0.91	0.87	0.90	0.97	0.95	
0.71	0.81	0.81	0.87	0.57	0.37	0.80	0.88	0.89	0.79	0.85	
0.49	0.62	0.60	0.58	0.50	0.60	0.58	0.50	0.61	0.45	0.33	
0.86	0.84	0.74	0.58	0.51	0.39	0.73	0.92	0.91	0.49	0.74	
0.96	0.67	0.54	0.85	0.48	0.37	0.88	0.90	0.94	0.82	0.93	
0.69	0.49	0.56	0.66	0.43	0.42	0.77	0.73	0.71	0.90	0.99	
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	
0.85	0.75	0.55	0.55	0.55	0.45	0.75	0.72	0.77	0.75	0.71	
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	
0.85	0.75	0.55	0.55	0.55	0.45	0.75	0.72	0.77	0.75	0.71	
0.79	0.73	0.90	0.67	0.33	0.61	0.69	0.79	0.73	0.93	0.97	
0.91	0.94	0.89	0.49	0.41	0.78	0.78	0.77	0.89	0.99	0.93	

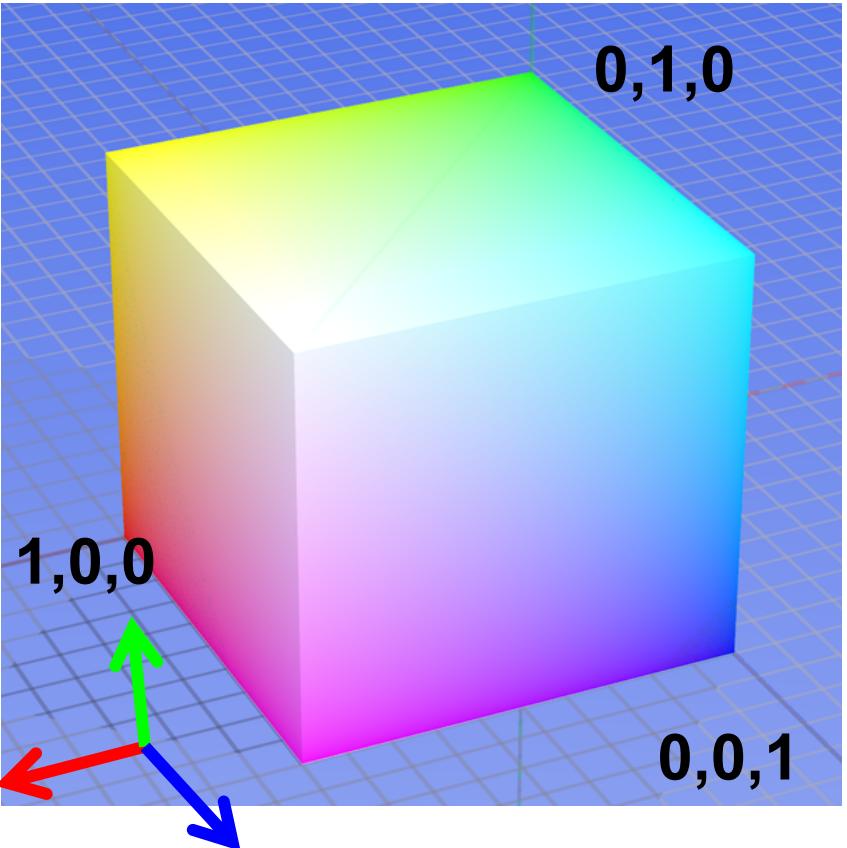
Color spaces

- How can we represent color?



Color spaces: RGB

Default color space



Some drawbacks

- Strongly correlated channels
- Non-perceptual



R
(G=0,B=0)



G
(R=0,B=0)

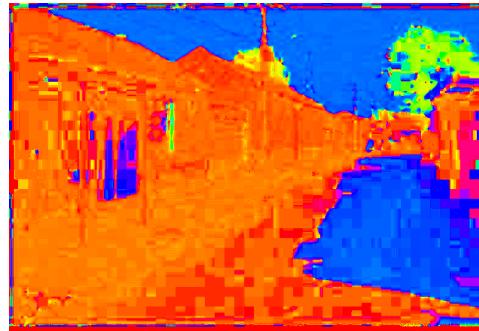
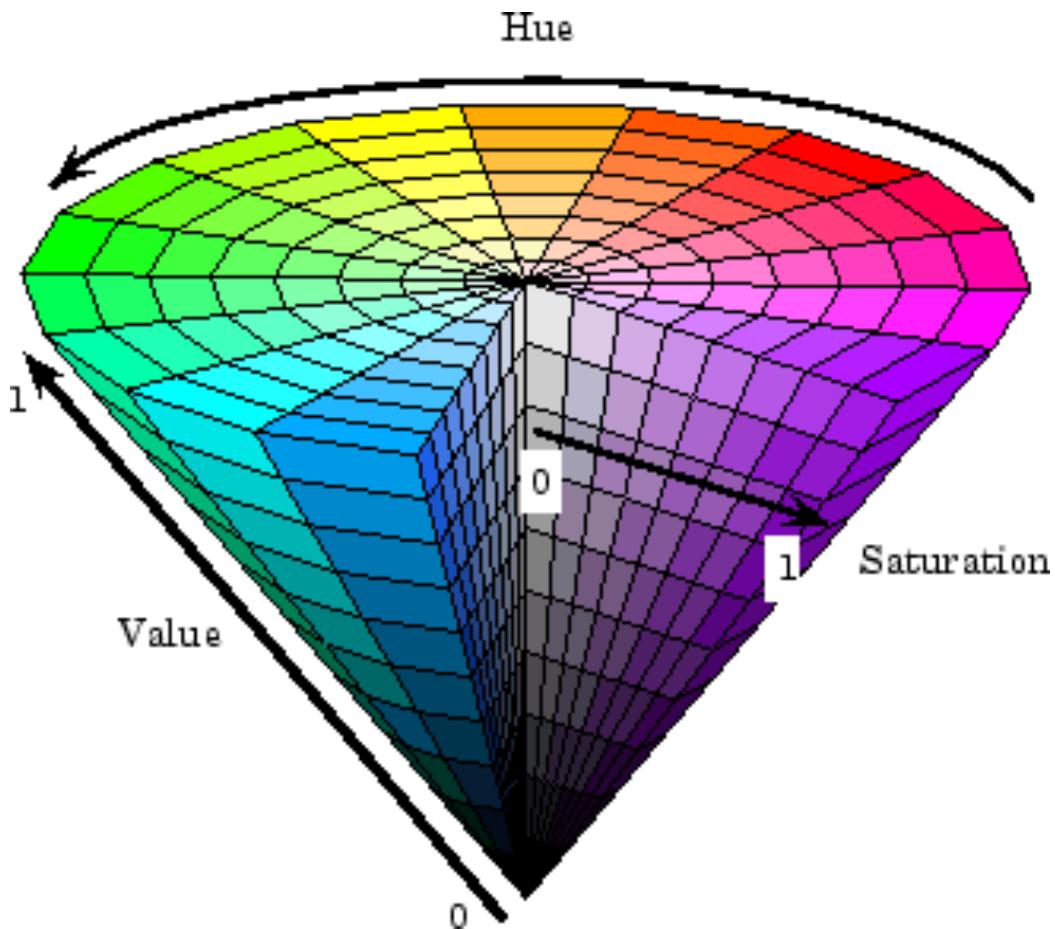


B
(R=0,G=0)

Color spaces: HSV



Intuitive color space



H
($S=1, V=1$)



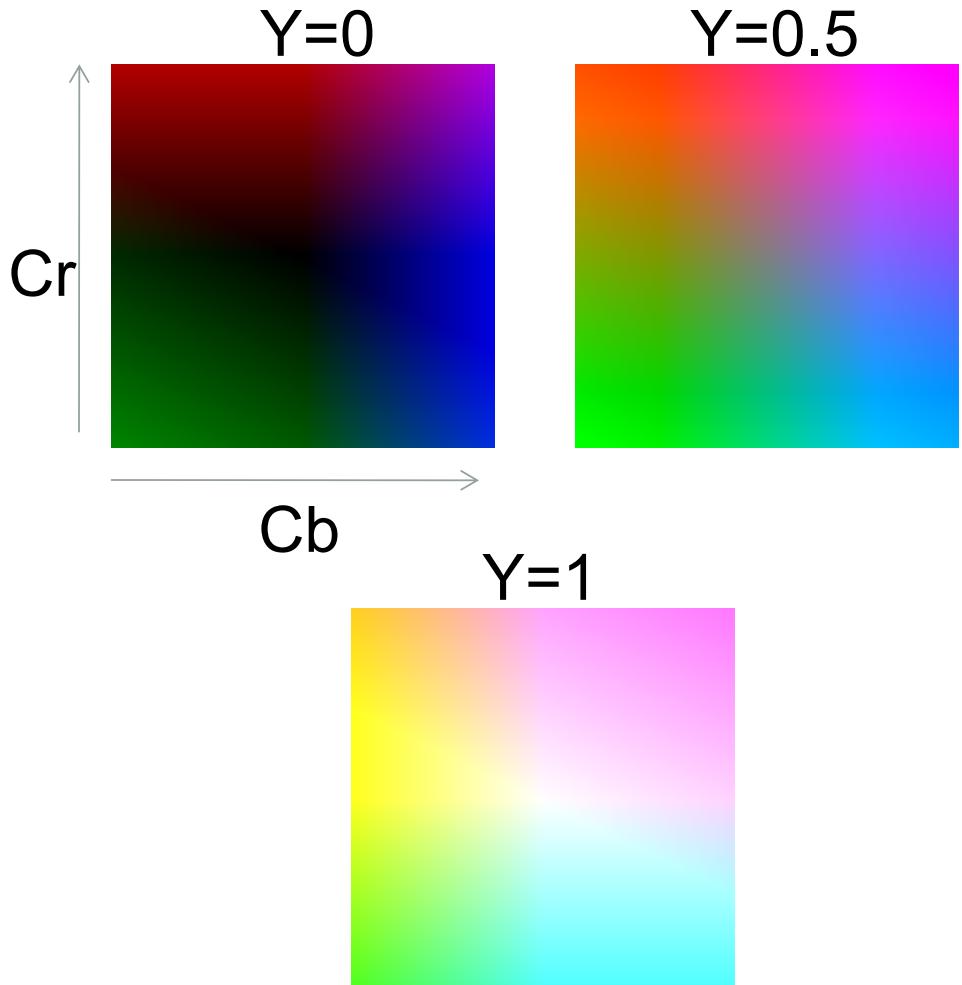
S
($H=1, V=1$)



V
($H=1, S=0$)

Color spaces: YCbCr

Fast to compute, good for compression, used by TV



Y
($Cb=0.5, Cr=0.5$)



Cb
($Y=0.5, Cr=0.5$)

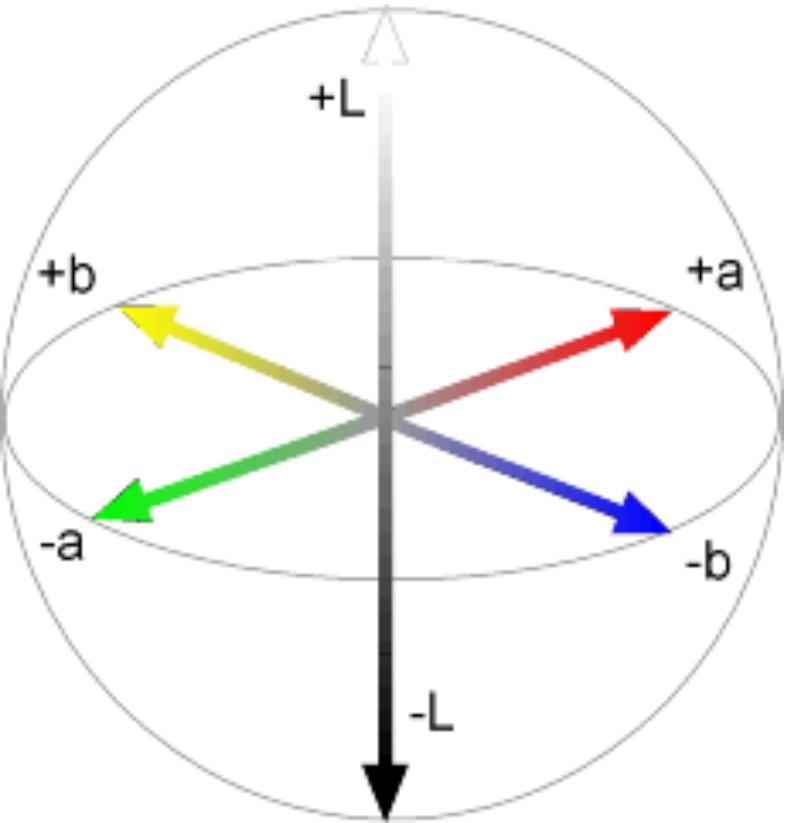


Cr
($Y=0.5, Cb=0.5$)

Color spaces: L*a*b*



“Perceptually uniform”* color space



L
($a=0, b=0$)



a
($L=65, b=0$)



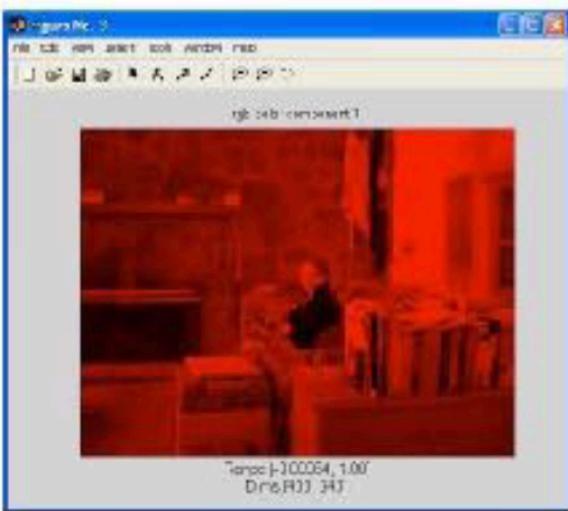
b
($L=65, a=0$)

luminance, chrominance color components: Y, I, Q

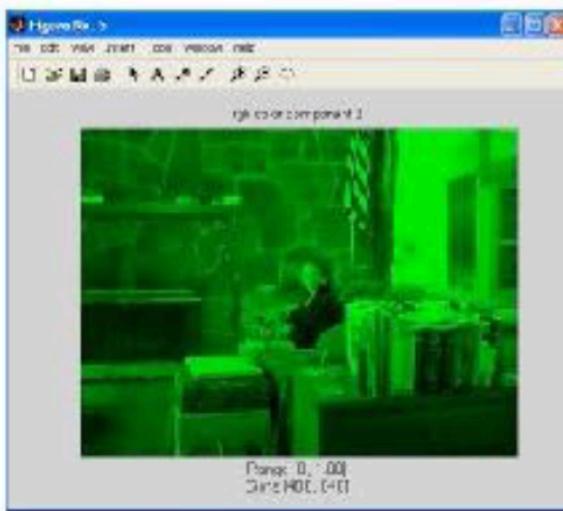
$$\begin{pmatrix} Y \\ I \\ Q \end{pmatrix} = \begin{pmatrix} 0.299 & 0.587 & 0.114 \\ 0.596 & -0.274 & -0.322 \\ 0.211 & -0.523 & 0.312 \end{pmatrix} \begin{pmatrix} R \\ G \\ B \end{pmatrix}$$



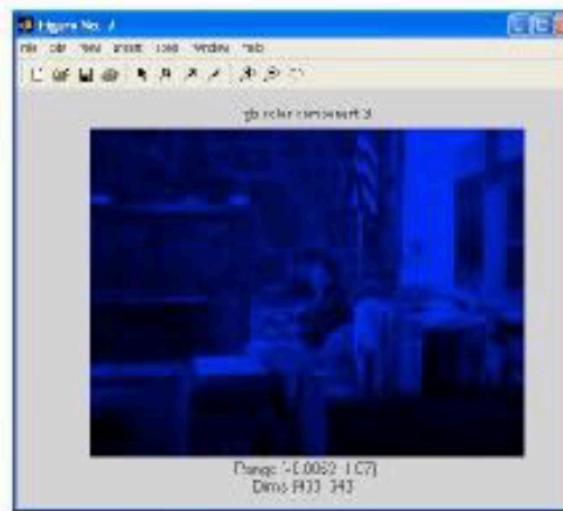
YIQ - RGB



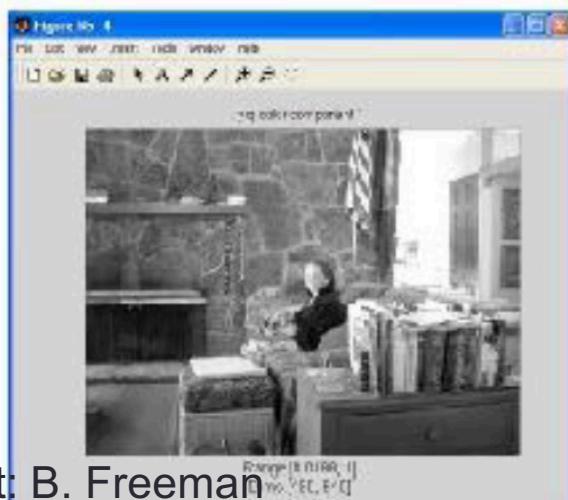
R



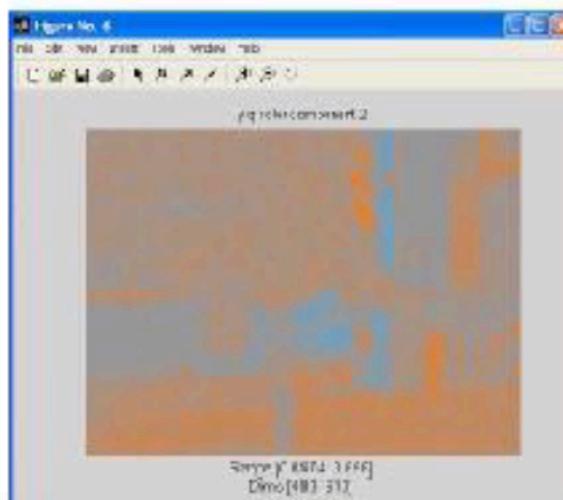
G



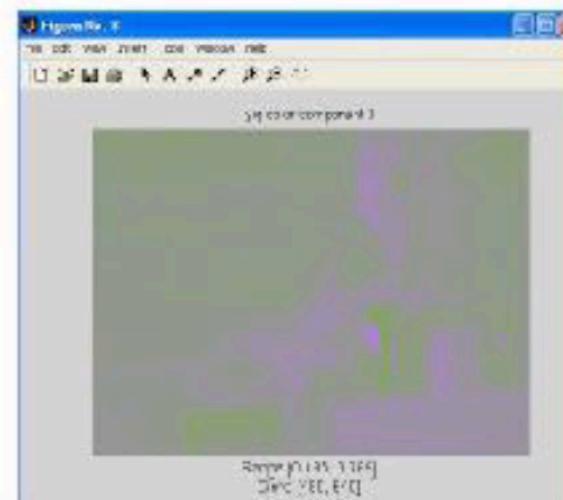
B



I



Q



S

Most information in intensity



Only color shown – constant intensity

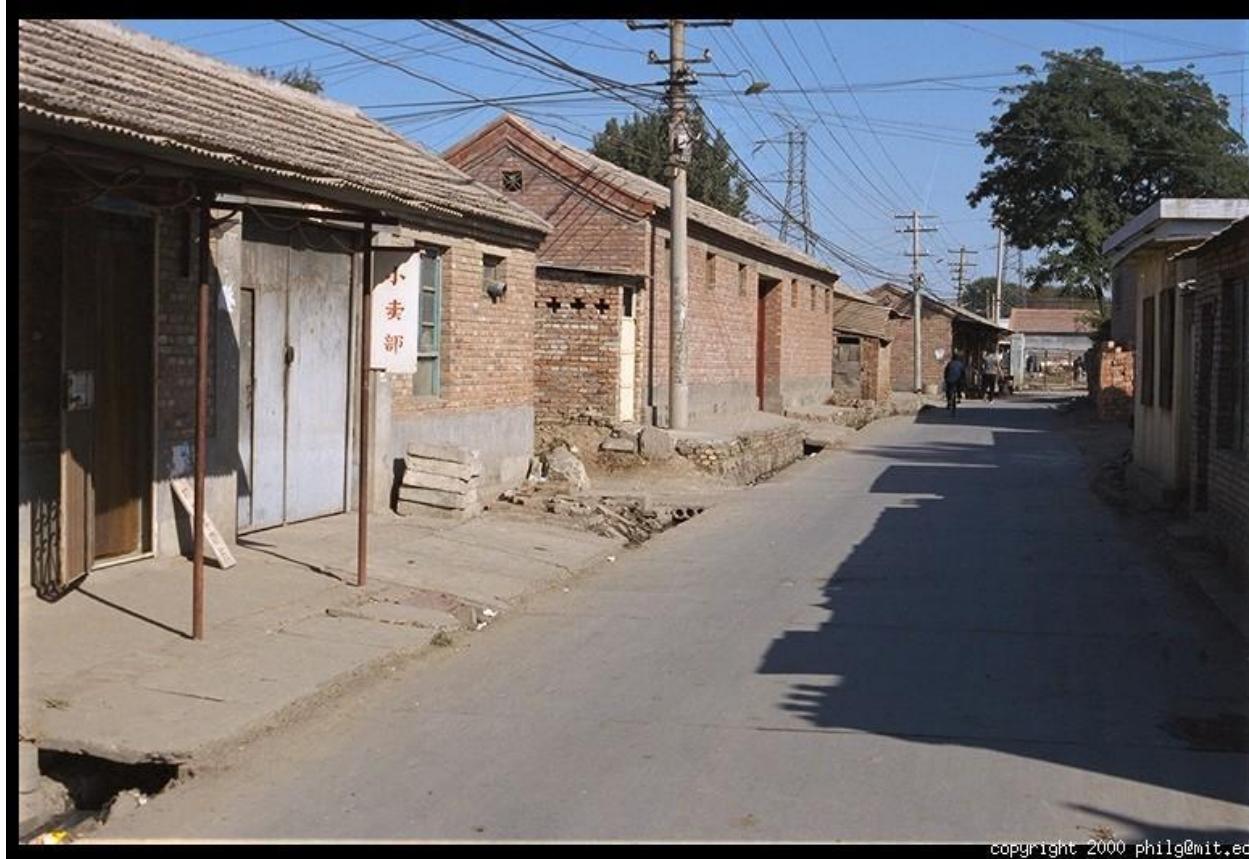
Most information in intensity



copyright 2000 philg@mit.edu

Only intensity shown – constant color

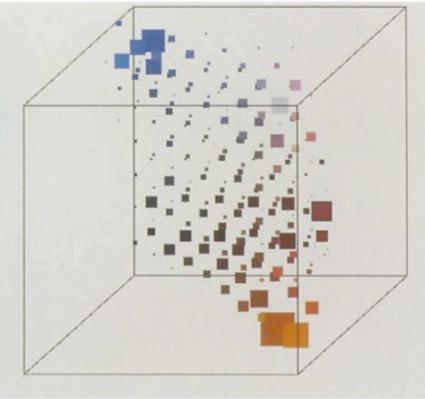
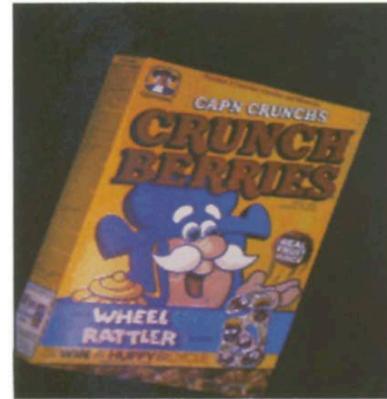
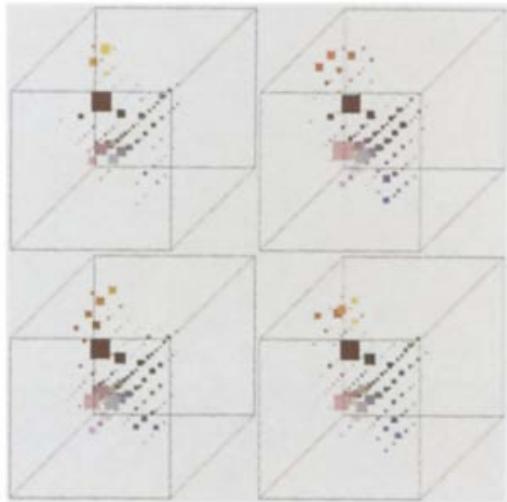
Most information in intensity



Original image

Uses of color in computer vision

Color histograms for indexing and retrieval



Swain and Ballard, [Color Indexing](#), IJCV 1991.

Uses of color in computer vision

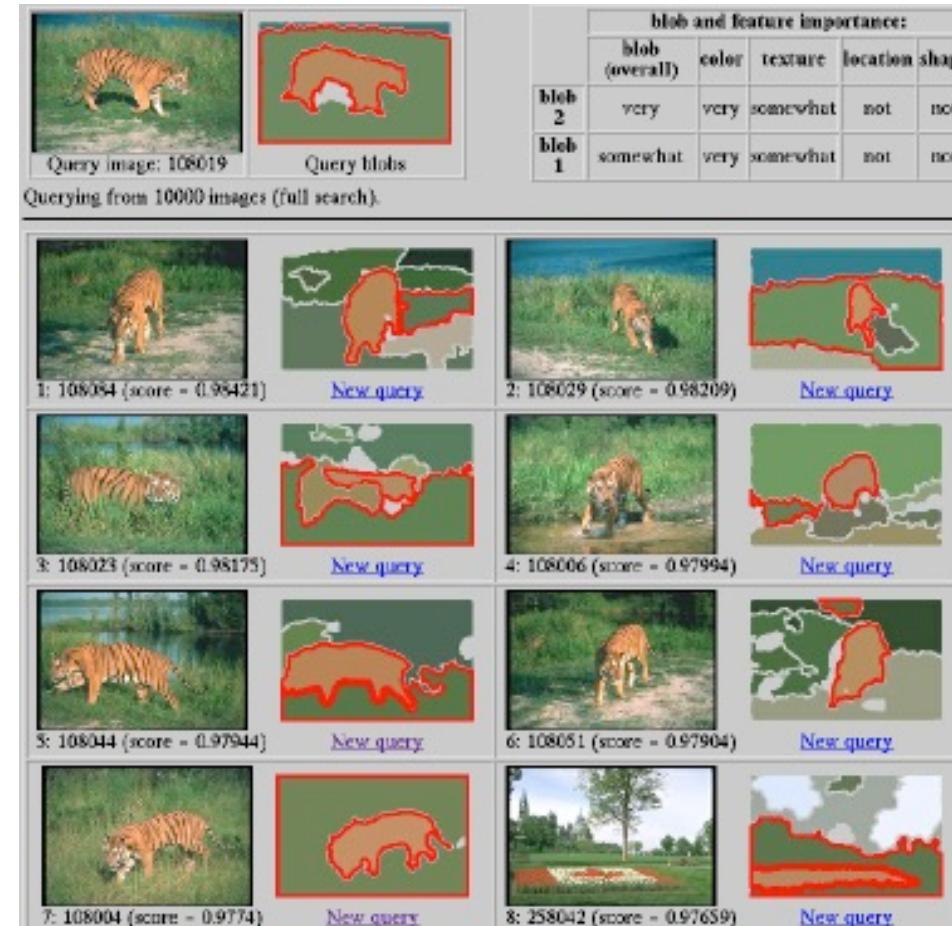
Skin detection



M. Jones and J. Rehg, [Statistical Color Models with Application to Skin Detection](#), IJCV
2002.

Uses of color in computer vision

Image segmentation and retrieval



C. Carson, S. Belongie, H. Greenspan, and Ji. Malik, Blobworld: Image segmentation using Expectation-Maximization and its application to image querying, ICVIS 1999.

Uses of color in computer vision

Building appearance models for tracking

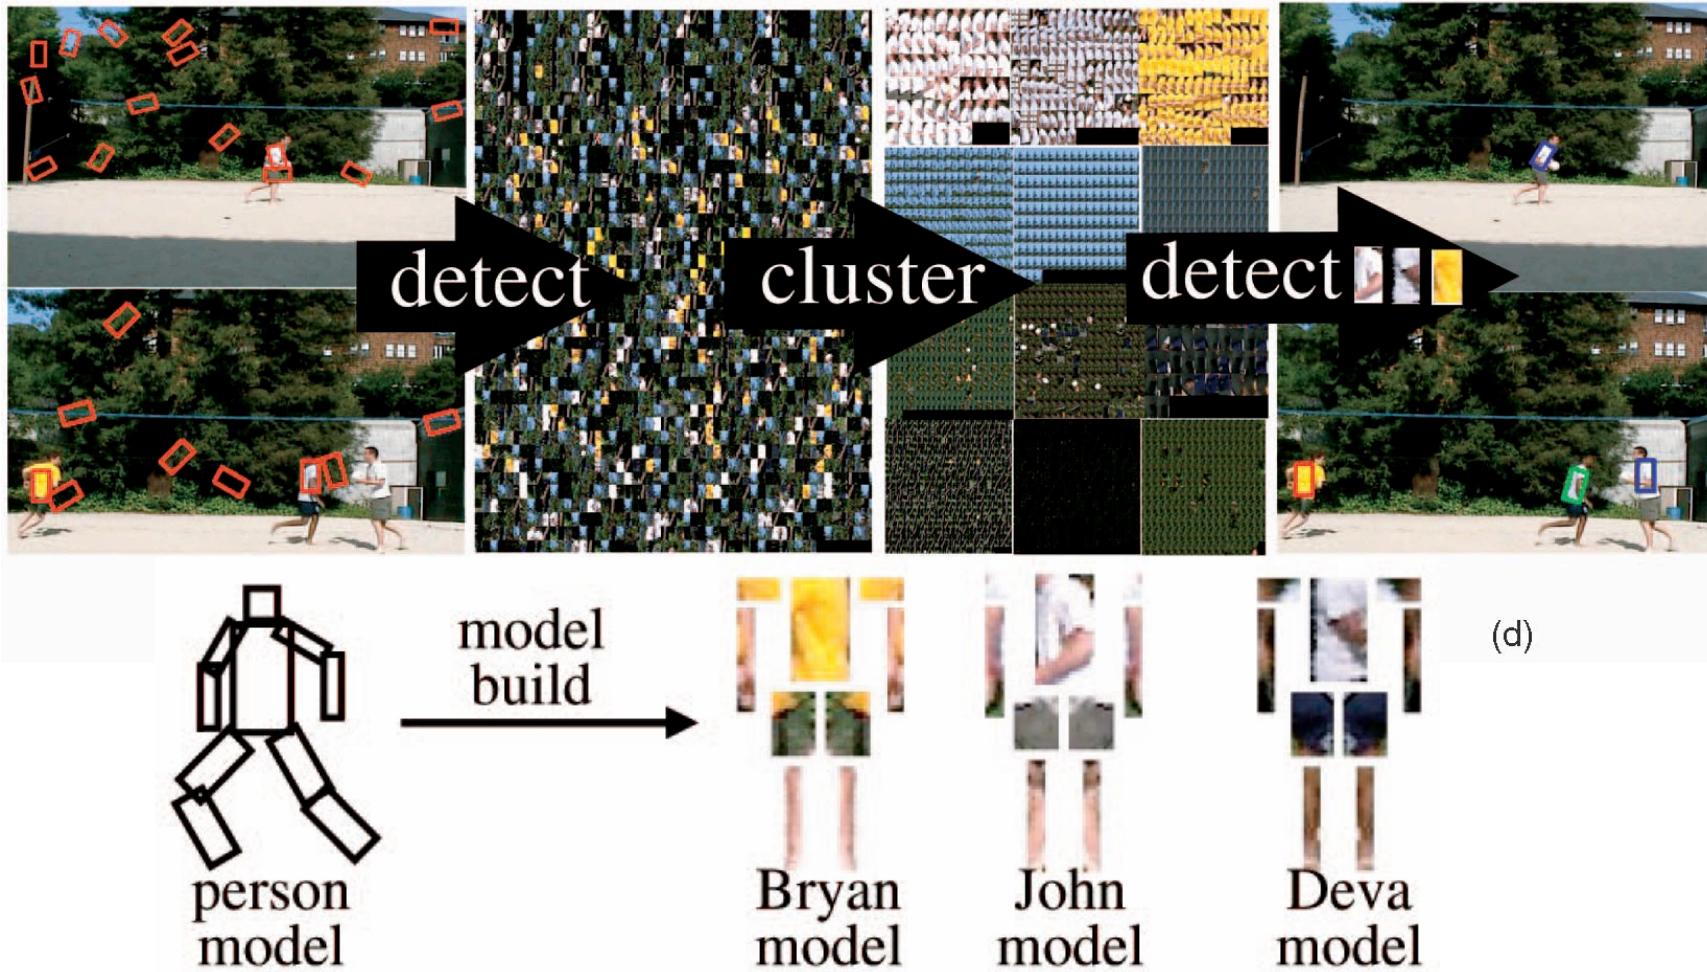
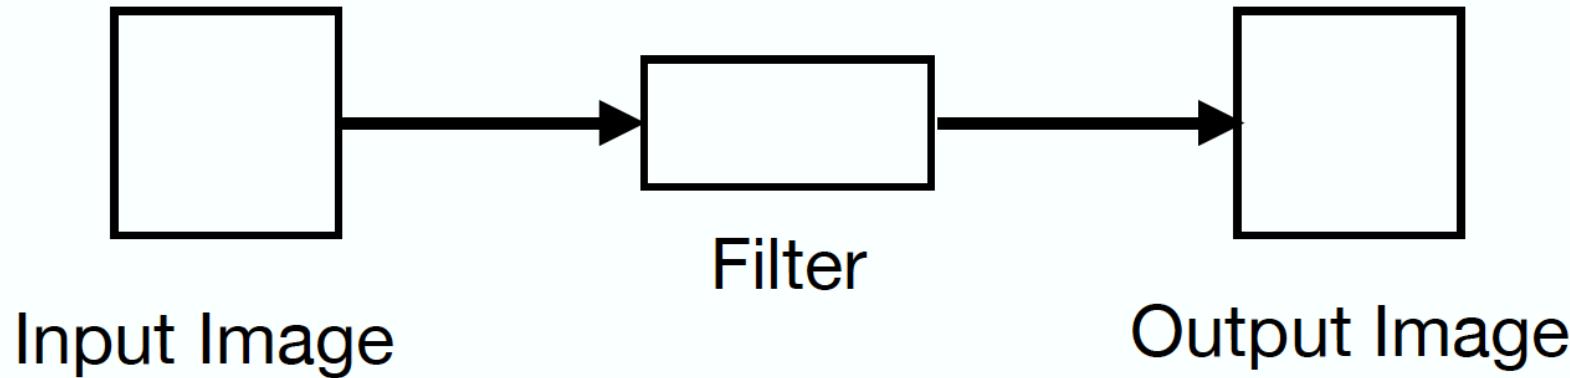


Image Filtering

BBM416

Image Filtering



We want to remove unwanted sources of variation, and keep the information relevant for whatever task we need to solve

Image filtering

- Image filtering: for each pixel, compute function of local neighborhood and output a new value
 - Function specified by a “filter” or mask saying how to combine values from neighbors.
 - Same function applied at each position
 - Output and input image are typically the same size

Image filtering

- Linear filtering: function is a weighted sum/difference of pixel values
- Uses of filtering:
 - Enhance images
 - Denoise, smooth, increase contrast, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching
 - Convolutional Neural Networks

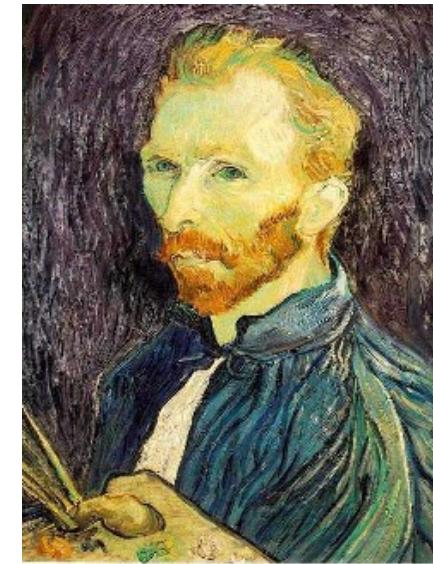
De-noising



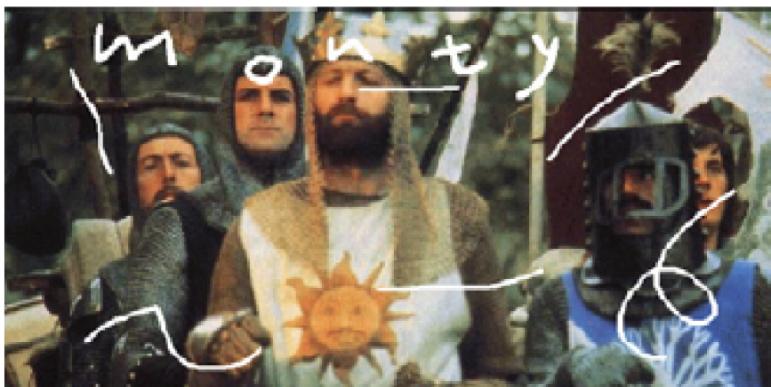
Salt and pepper noise



Super-resolution



In-painting



Bertamio et al.

Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Take lots of images and average them!
What's the next best thing?

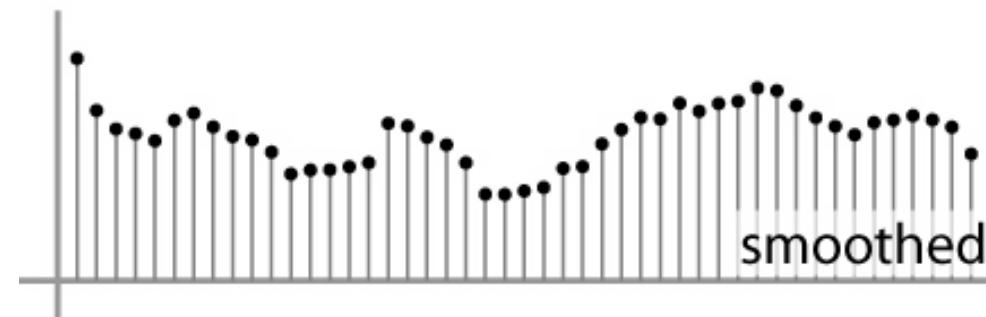
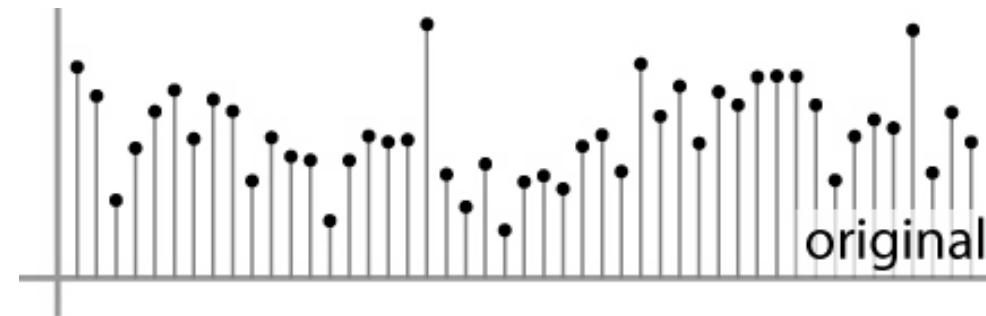
Source: S. Seitz

First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
 - Expect pixels to be like their neighbors
 - Expect noise processes to be independent from pixel to pixel

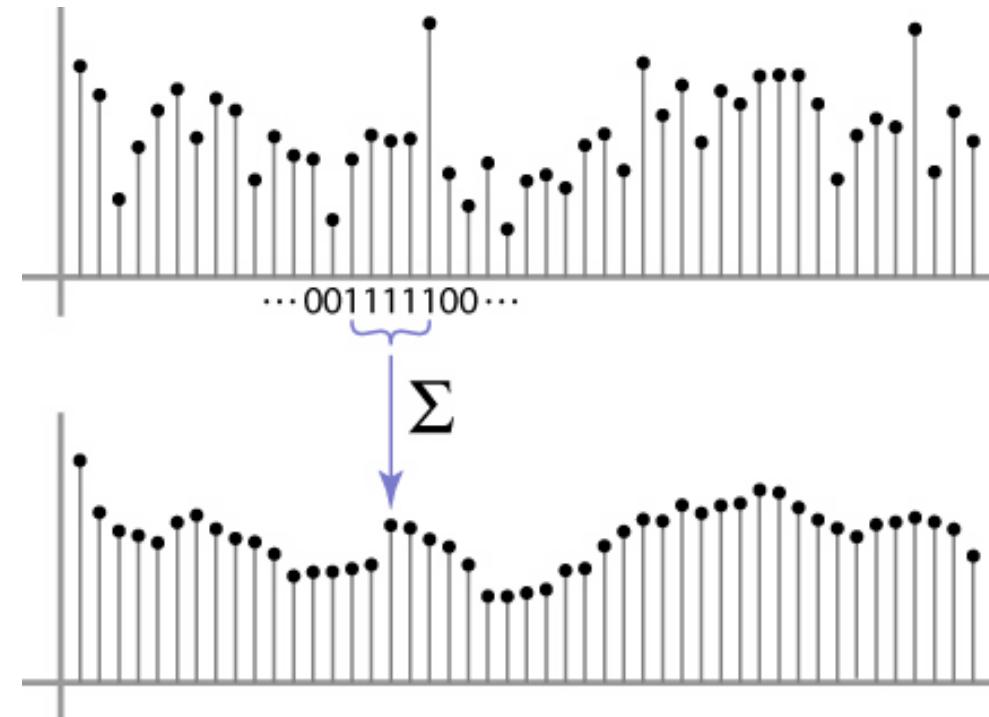
First attempt at a solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



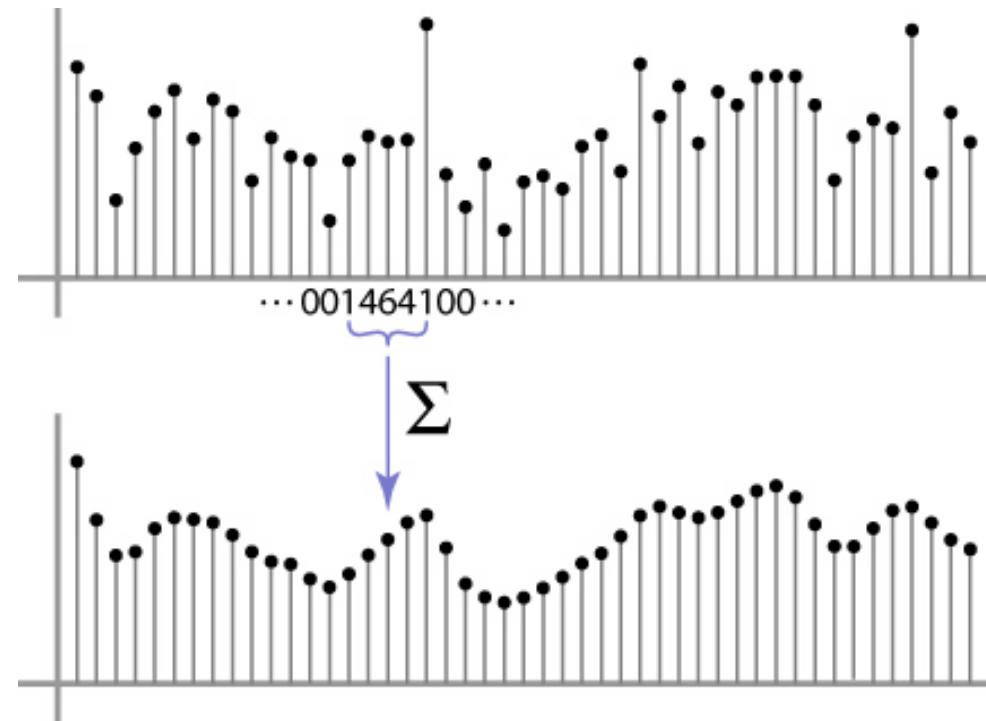
Weighted Moving Average

- Can add weights to our moving average
- $Weights [1, 1, 1, 1, 1] / 5$



Weighted Moving Average

- Non-uniform weights $[1, 4, 6, 4, 1] / 16$



Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10									

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

0	10	20								

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

Moving Average In 2D

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30				

Moving Average In 2D

$F[x, y]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

Averaging filter

- What values belong in the kernel H for the moving average example?

The diagram illustrates the convolution operation $G = F * H$. The input image $F[x, y]$ is a 10x10 grid. The filter $H[u, v]$ is a 3x3 grid with a central element marked with a question mark. The output image $G[x, y]$ is a 7x7 grid. A red box highlights the central 3x3 subgrid of F where the value 90 is highlighted. A red square highlights the value 30 in the corresponding position of G .

$$G = H \otimes F$$

Linear, shift-invariant filters

- A fundamental kind of image processing operation
- Linear: preserves summation and scalar multiplication

$$f(I_1 + I_2) = f(I_1) + f(I_2)$$

$$f(aI) = af(I)$$

- Shift-invariant: commutes with shifting the image

$$f(s_{\Delta x, \Delta y}(I)) = s_{\Delta x, \Delta y}(f(I))$$

(i.e. filter does the same thing at different places in the image)

- Surprisingly, all such operations can be computed with one simple algorithm: **convolution**

Convolution

- Same moving average operation, expressed mathematically:

$$c[i] = \frac{1}{2r + 1} \sum_{j=i-r}^{i+r} a[j].$$

Convolution

- Simple averaging:

$$c[i] = \frac{1}{2r+1} \sum_{j=i-r}^{i+r} a[j].$$

every sample gets the same weight

- Convolution: same idea but with **weighted** average

$$(a \star b)[i] = \sum_j a[j]b[i-j].$$

each sample gets its own weight (normally zero far away)

- Visually: reflect b and slide it over so that $b[0]$ lines up with $a[j]$
- This is all convolution is: it is a **moving weighted average**

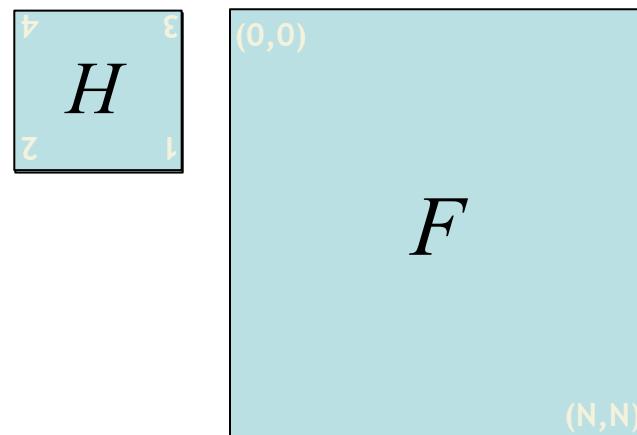
Convolution

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

$$G = H \star F$$

↑
Notation for convolution operator



Correlation filtering

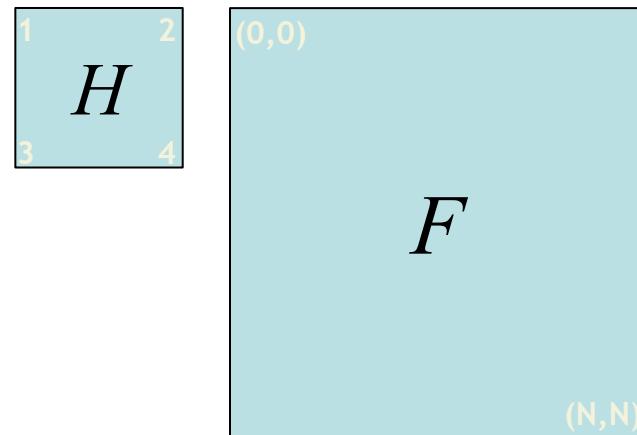
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called cross-correlation, denoted

$$G = H \otimes F$$

Filtering an image: replace each pixel
with a linear combination of its
neighbors.

Conceptually simpler, but not as nice as convolution



Convolution vs. correlation

Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$

For a Gaussian or box filter, how will the outputs differ?

If the input is an impulse signal, how will the outputs differ?

Correlation vs. Convolution

- Correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$

- Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

Matlab:
`filter2`
`imfilter`

Note the difference!

Matlab:
`conv2`

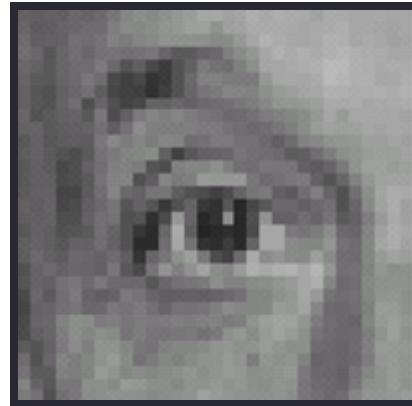
- Note $G = H \star F$

➤ If $H[-u, -v] = H[u, v]$ (if the filter is symmetric) then correlation \equiv convolution.

Properties of Convolution

- Linear and shift-invariant
- Commutative: $a * b = b * a$
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$
- Identity: unit impulse $e = [0, 0, 1, 0, 0]$,
 $a * e = a$

Practice with linear filters

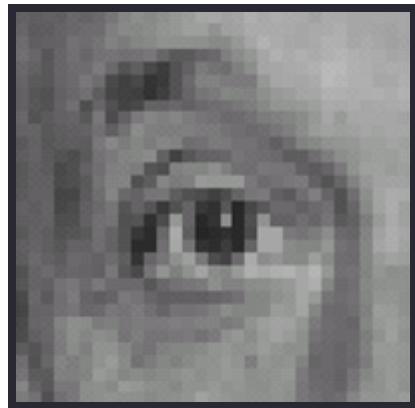


Original

0	0	0
0	1	0
0	0	0

?

Practice with linear filters



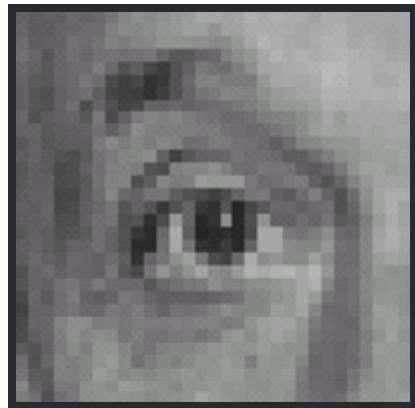
Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters

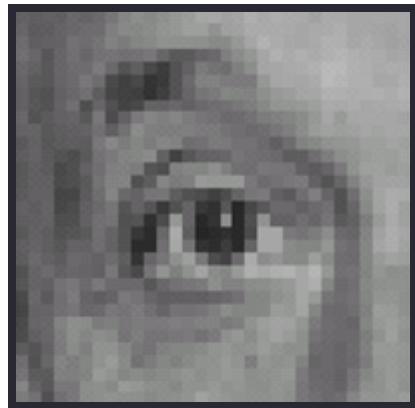


Original

0	0	0
0	0	1
0	0	0

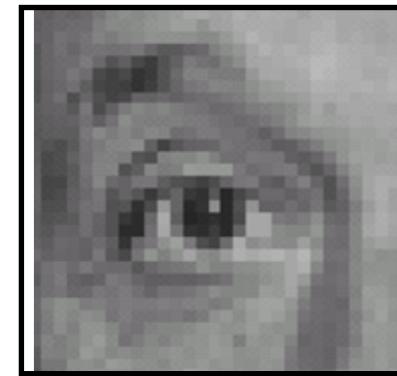
?

Practice with linear filters



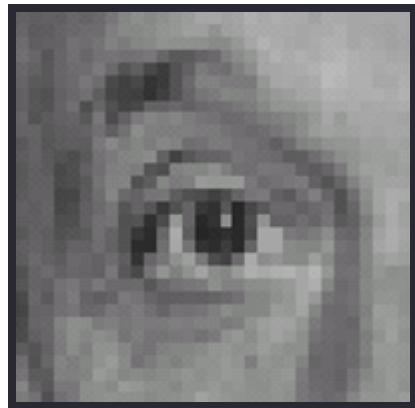
Original

0	0	0
0	0	1
0	0	0



Shifted right
by 1 pixel with
convolution

Practice with linear filters



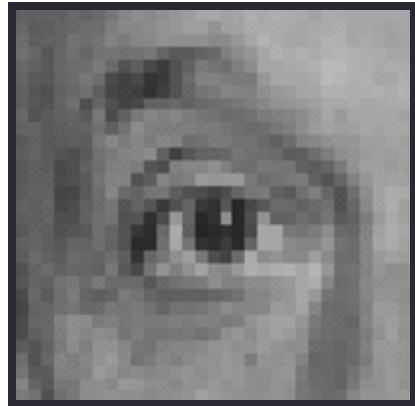
Original

0	0	0
0	0	1
0	0	0



Shifted left
by 1 pixel
with
correlation

Practice with linear filters

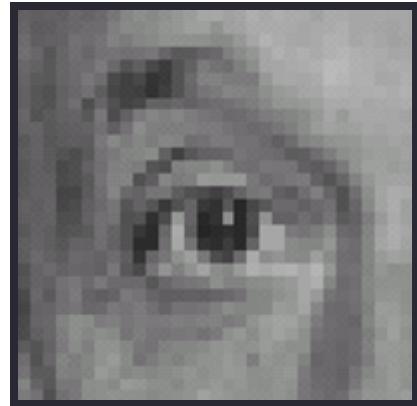


$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

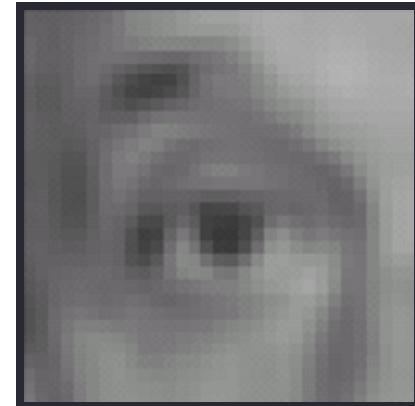
Original

Practice with linear filters



Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



Blur (with a
box filter)

Practice with linear filters



Original

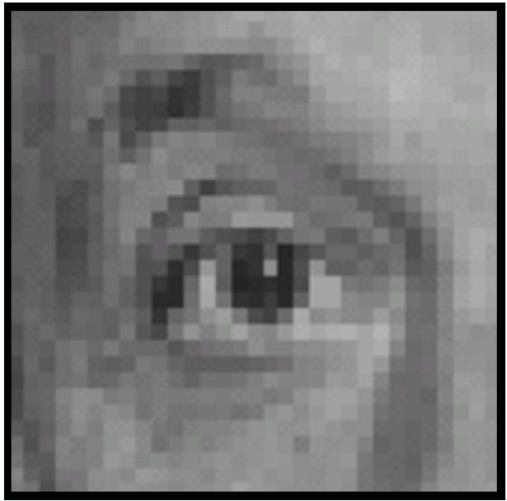
$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

-

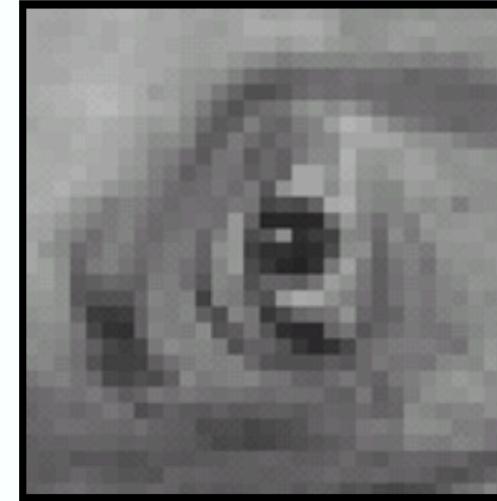
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

Practice with linear filters



?



Interactive tool for some filters: <https://setosa.io/ev/image-kernels/>

Practice with linear filters



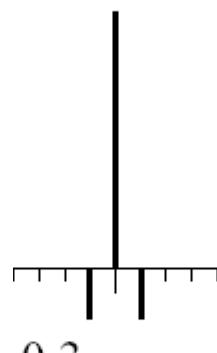
Original

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array}$$

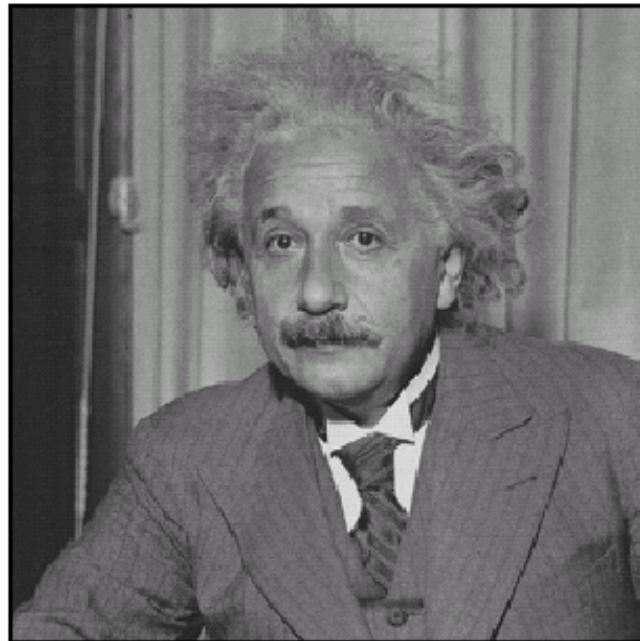
$$- \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



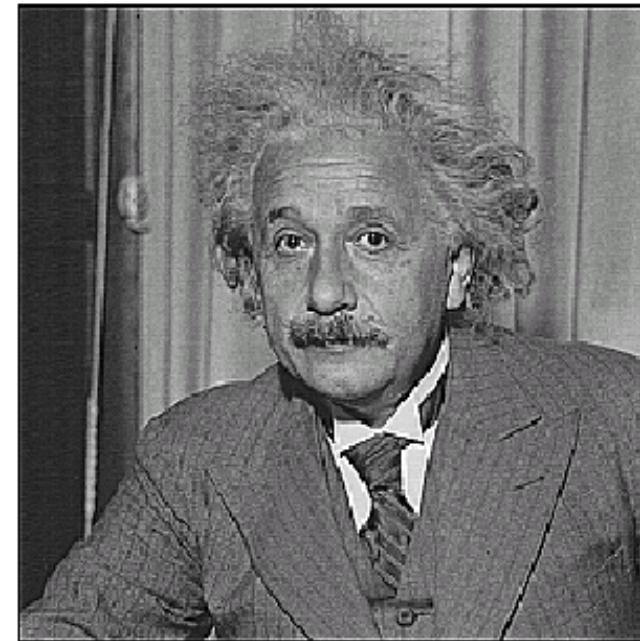
Sharpening filter
- Accentuates differences
with local average



Sharpening

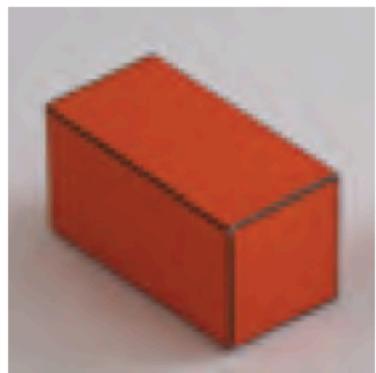


before

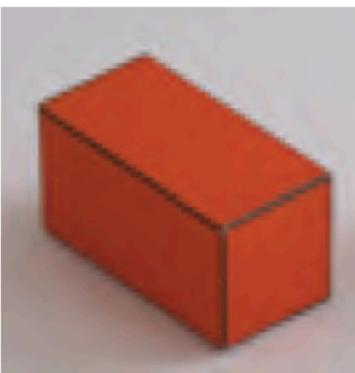


after

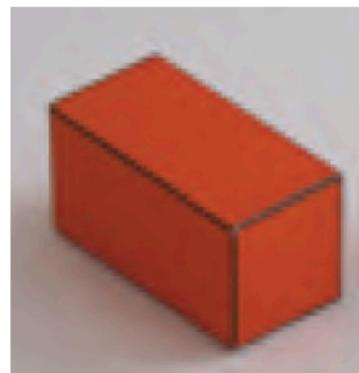
Practice with linear filters



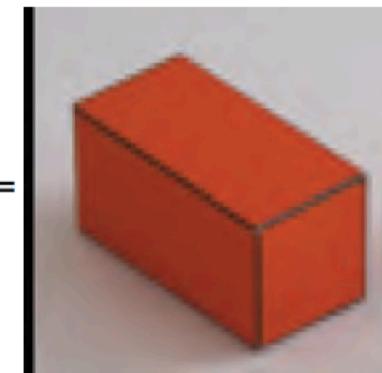
$$\odot \begin{array}{|c|c|c|c|c|} \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 1 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline \end{array} =$$



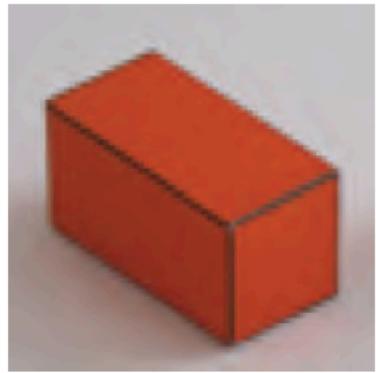
a)



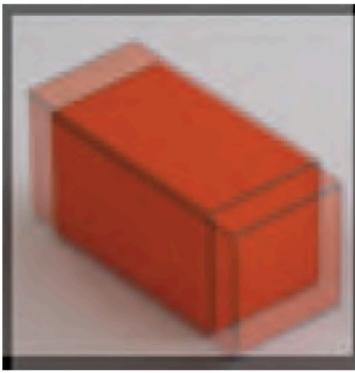
$$\odot \begin{array}{|c|c|c|c|c|} \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 1 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline \end{array} =$$



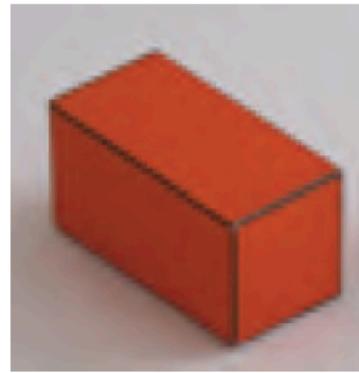
b)



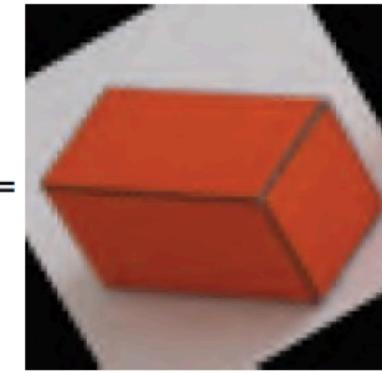
$$\odot \begin{array}{|c|c|c|c|c|} \hline .5 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & .5 \\ \hline \end{array} =$$



c)



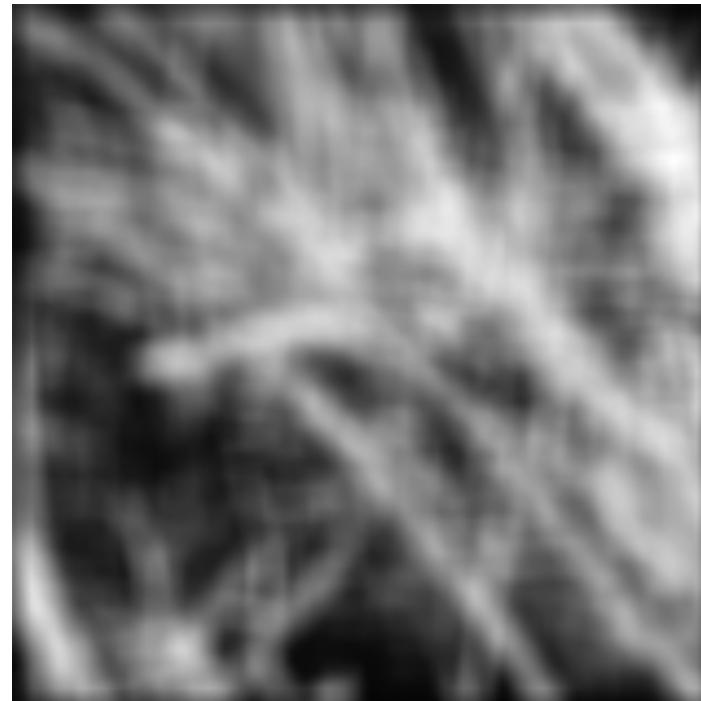
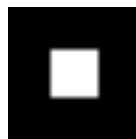
$$\odot \boxed{?} =$$



d)

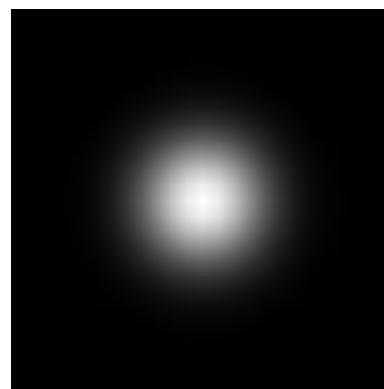
Smoothing with box filter revisited

- What's wrong with this picture?
- What's the solution?



Smoothing with box filter revisited

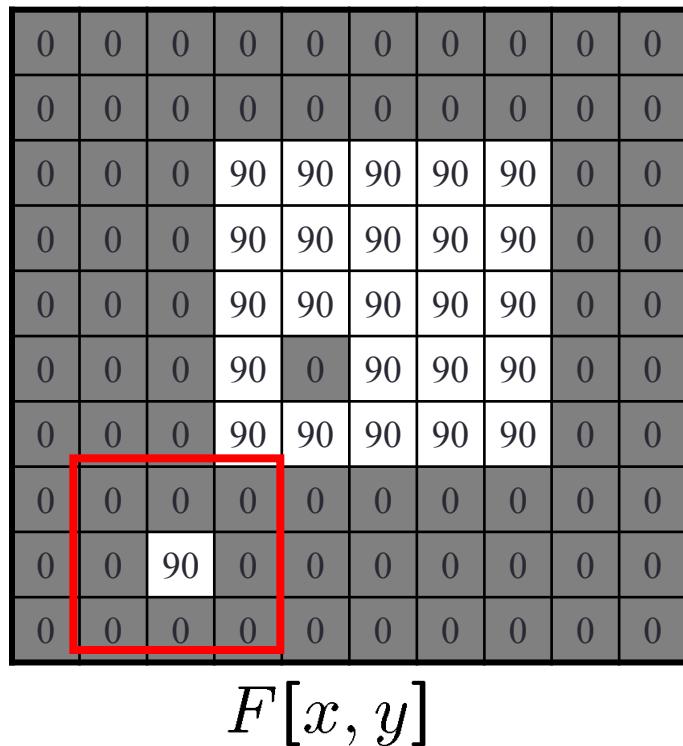
- What's wrong with this picture?
- What's the solution?
 - To eliminate edge effects, weight contribution of neighborhood pixels according to their closeness to the center



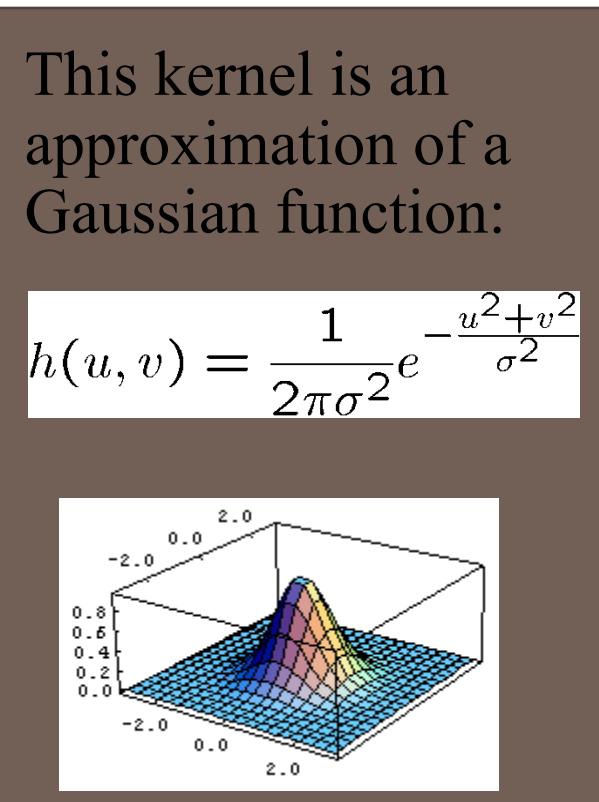
“fuzzy blob”

Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

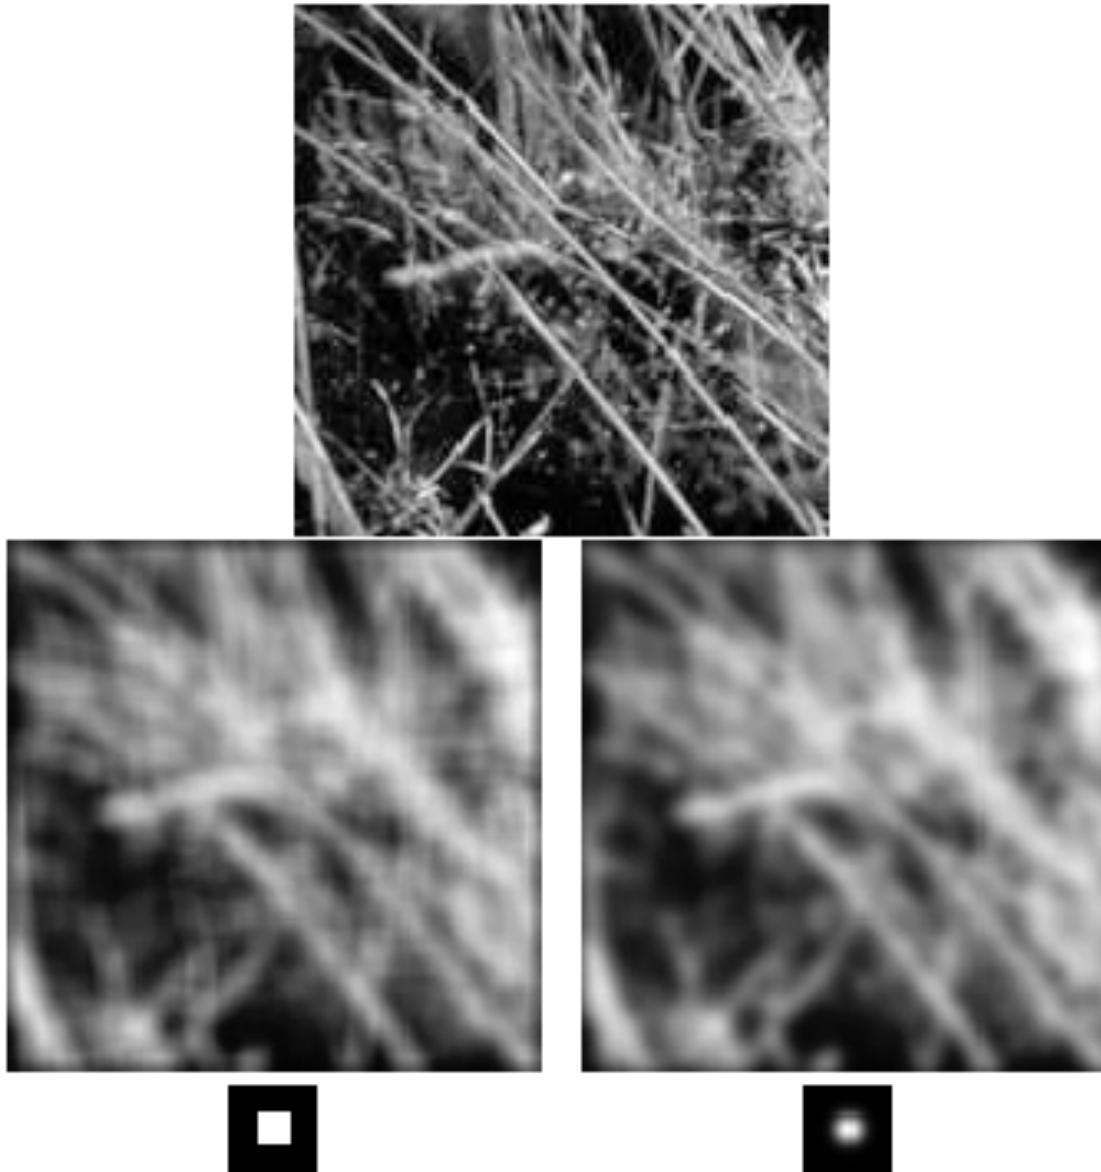


$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$
$$H[u, v]$$



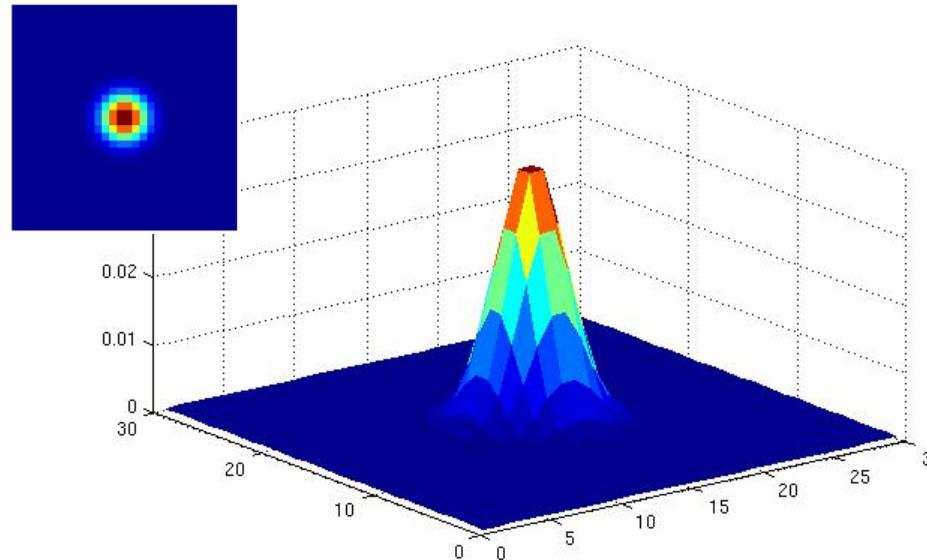
Removes high-frequency components from the image (it is a “low-pass filter”).

Gaussian vs. box filtering

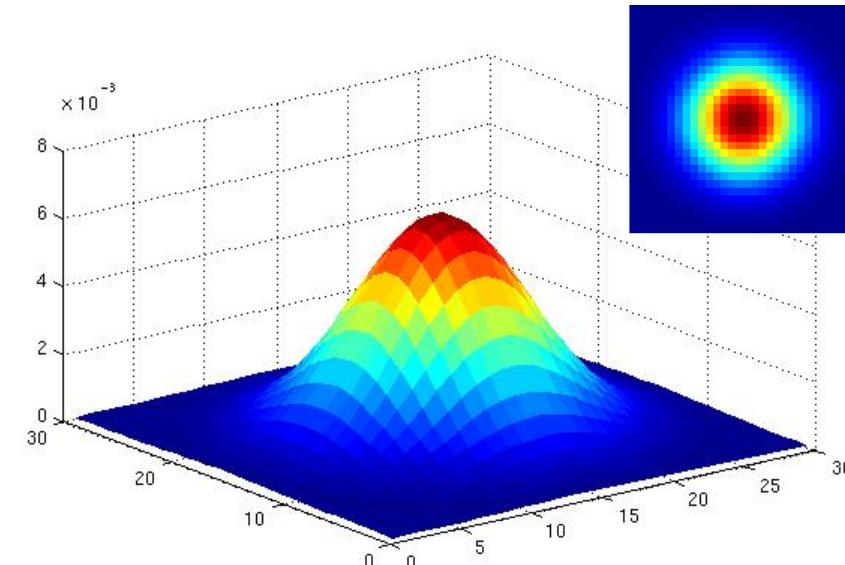


Gaussian Kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



$\sigma = 2$ with 30×30
kernel

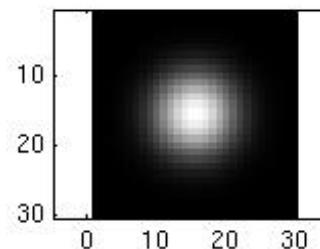
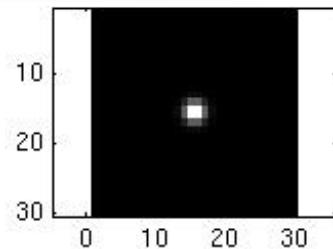


$\sigma = 5$ with 30×30
kernel

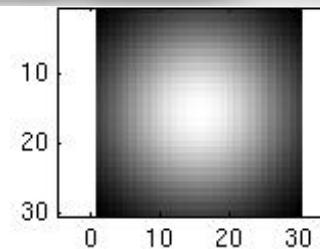
- Standard deviation σ : determines extent of smoothing

Smoothing with a Gaussian

Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.

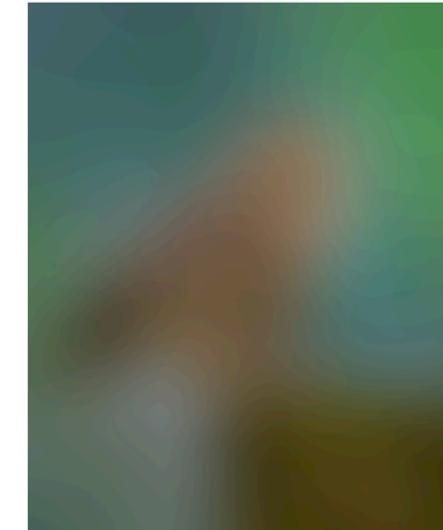
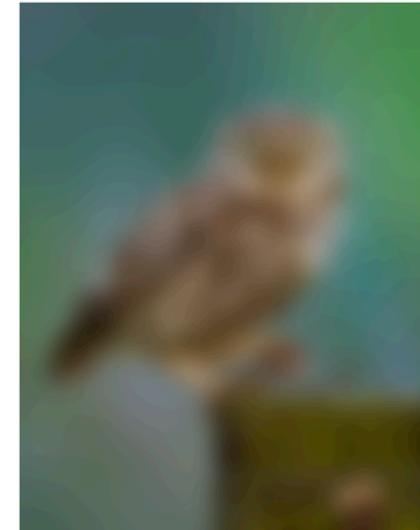


...

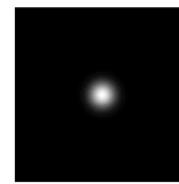


```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

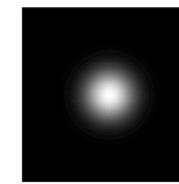
Gaussian Filters



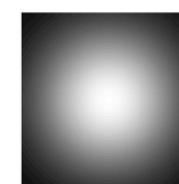
$\sigma = 1$ pixel



$\sigma = 5$ pixels



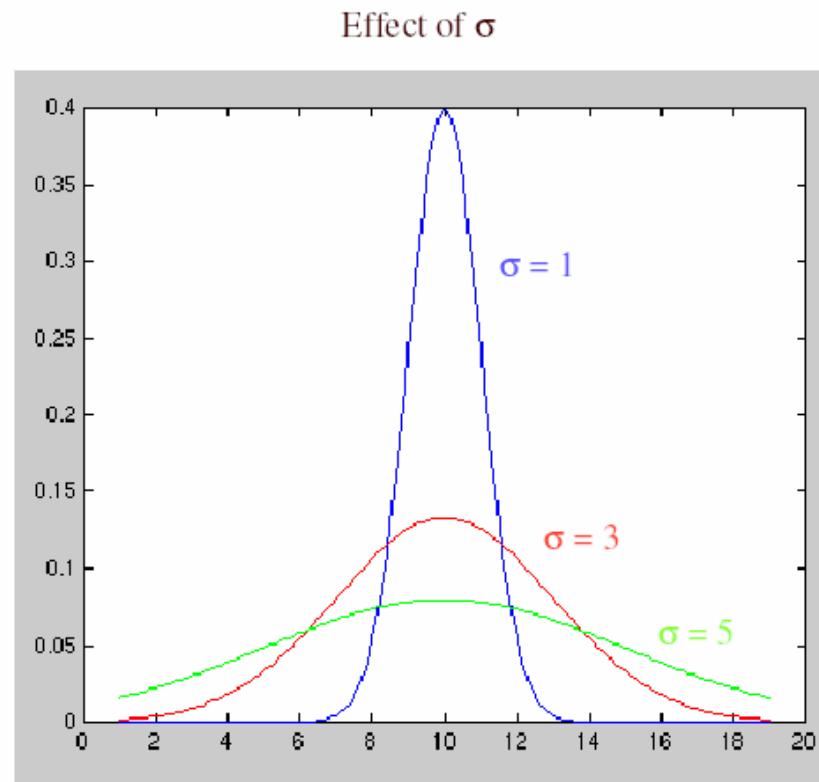
$\sigma = 10$ pixels



$\sigma = 30$ pixels

How big should the filter be?

- Values at edges should be near zero \leftarrow important!
- Rule of thumb for Gaussian: set filter half-width to about 3σ



Separable Filtering

$$a_2[i, j] = a_1[i]a_1[j]$$

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

=

0	0	0	0	0
0	0	0	0	0
1	4	6	4	1
0	0	0	0	0
0	0	0	0	0

*

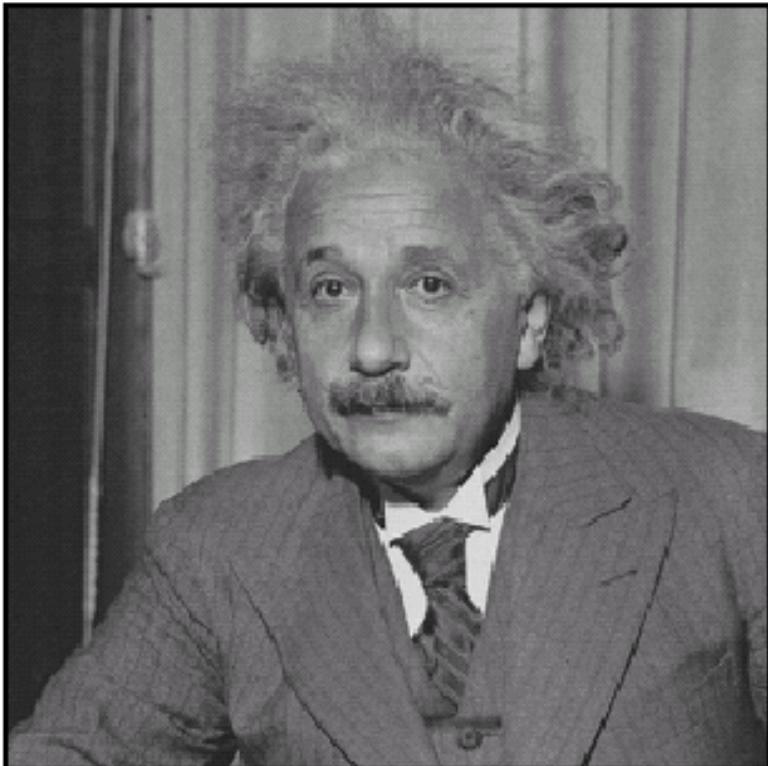
0	0	1	0	0
0	0	4	0	0
0	0	6	0	0
0	0	4	0	0
0	0	1	0	0

second, convolve with this

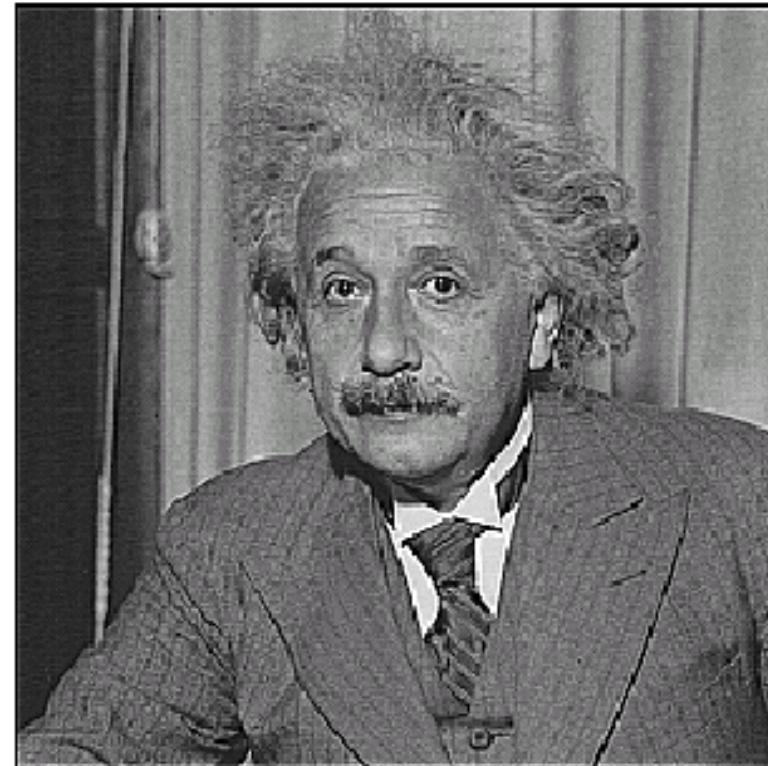
first, convolve with this

$$\sum_{i'} a_1[i'] \left(\sum_{j'} a_1[j'] b[i - i', j - j'] \right)$$

Sharpening revisited

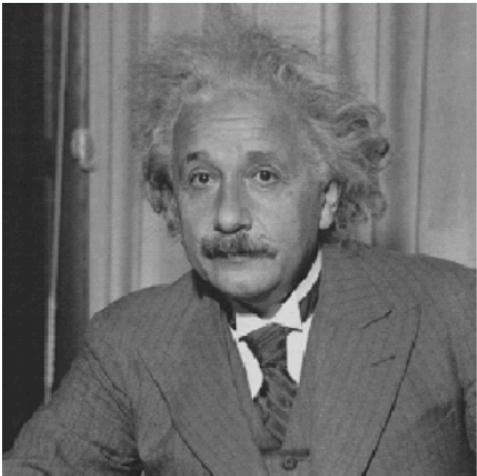


before

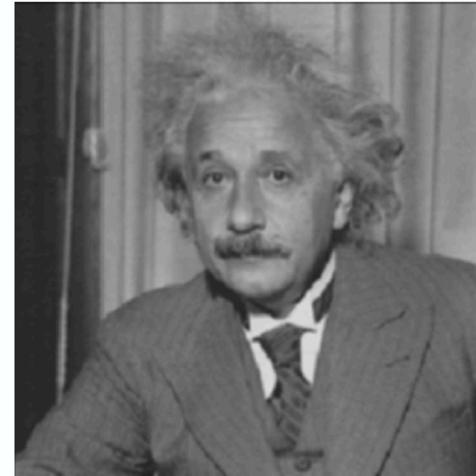


after

Sharpening revisited



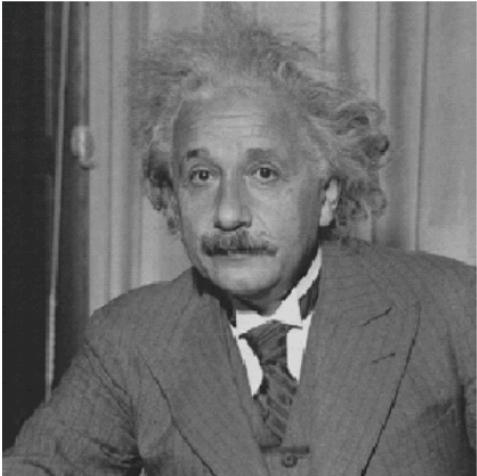
Image



Blurred



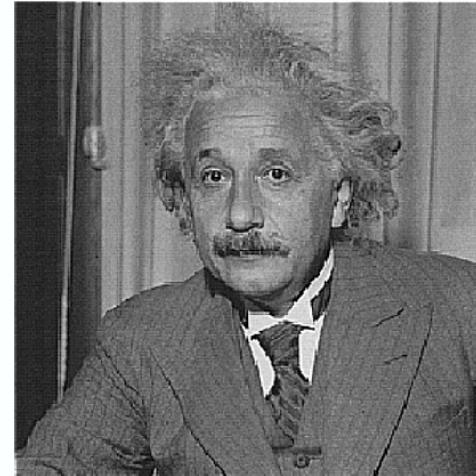
Detail



Image

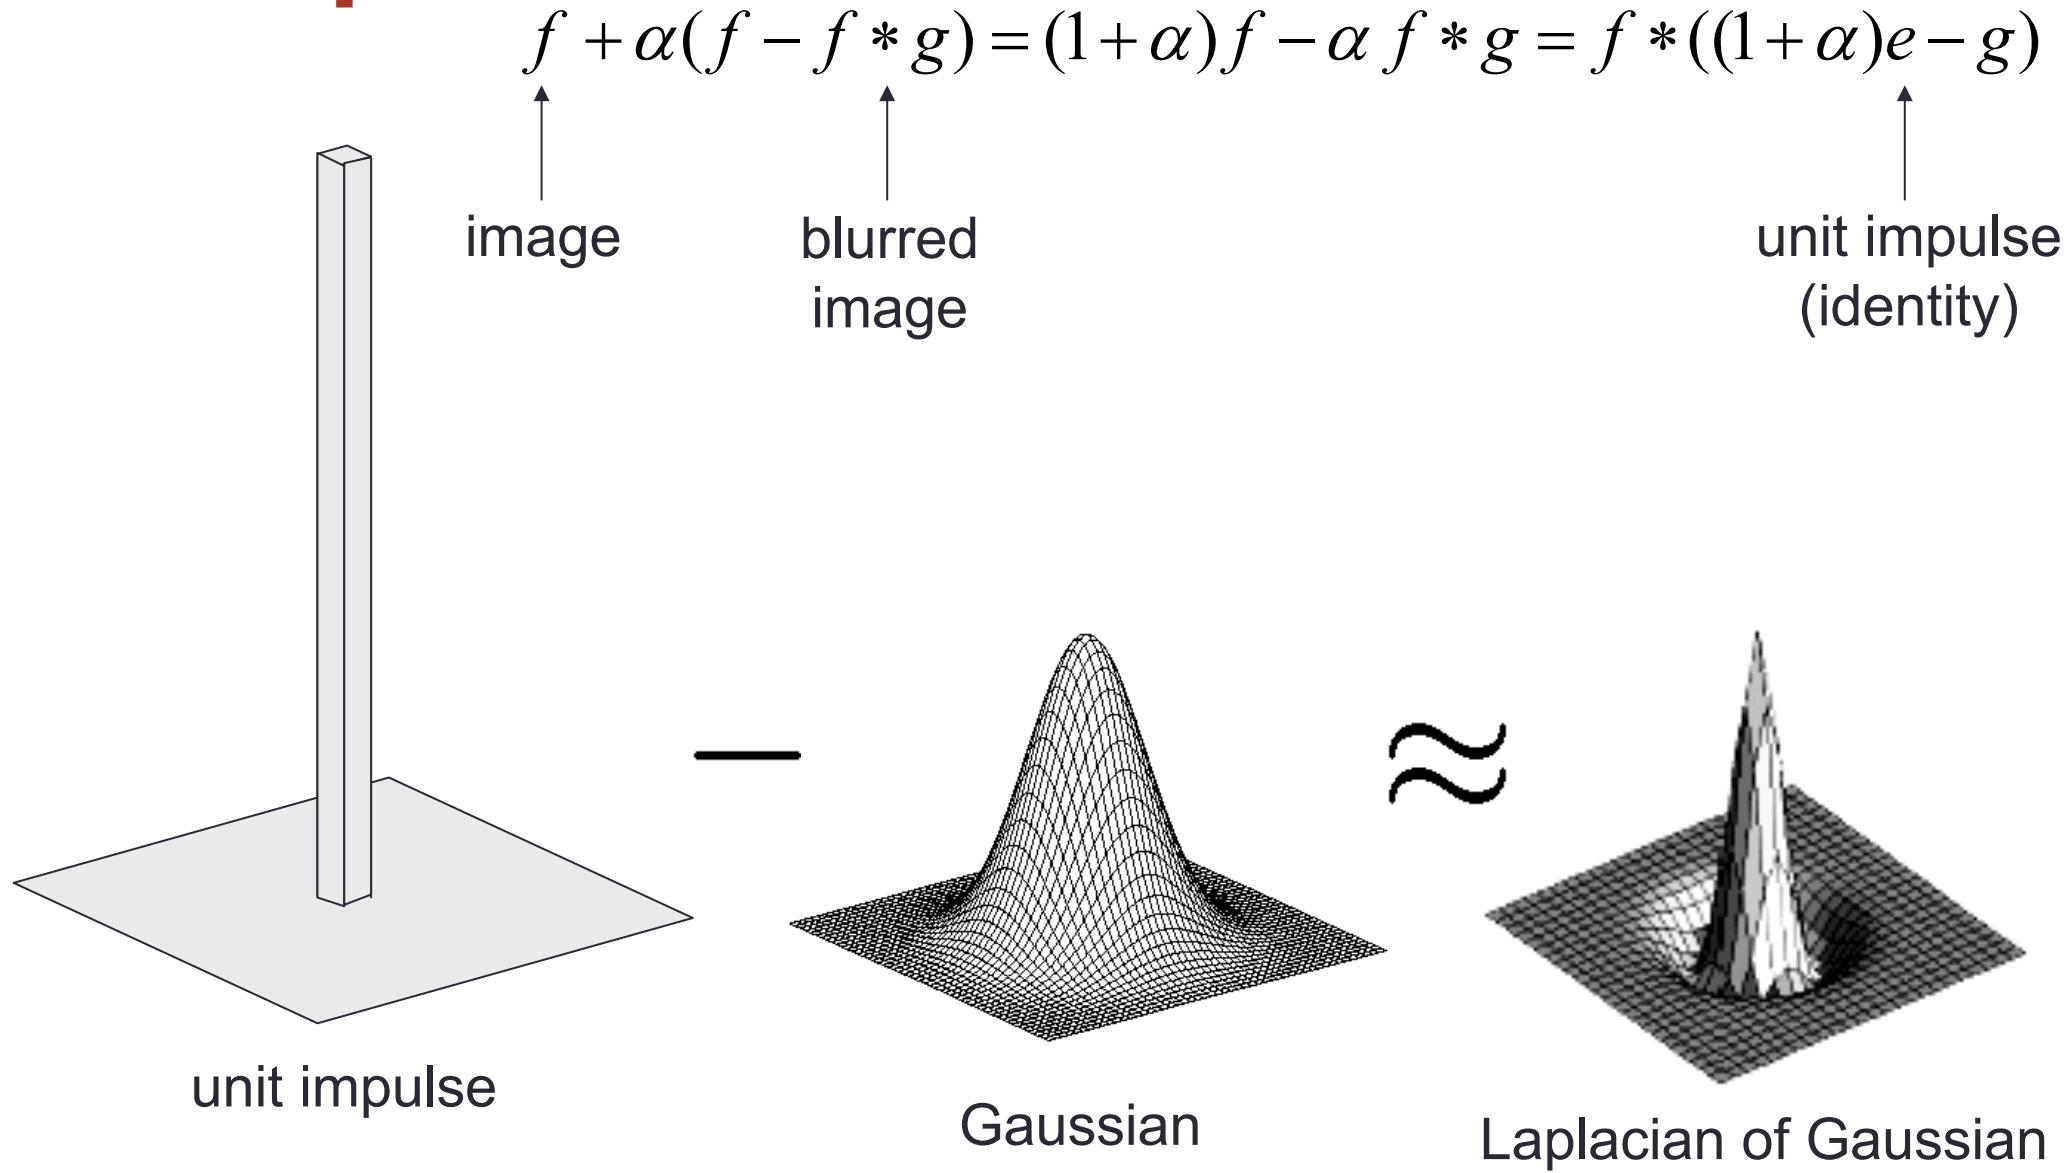


Detail



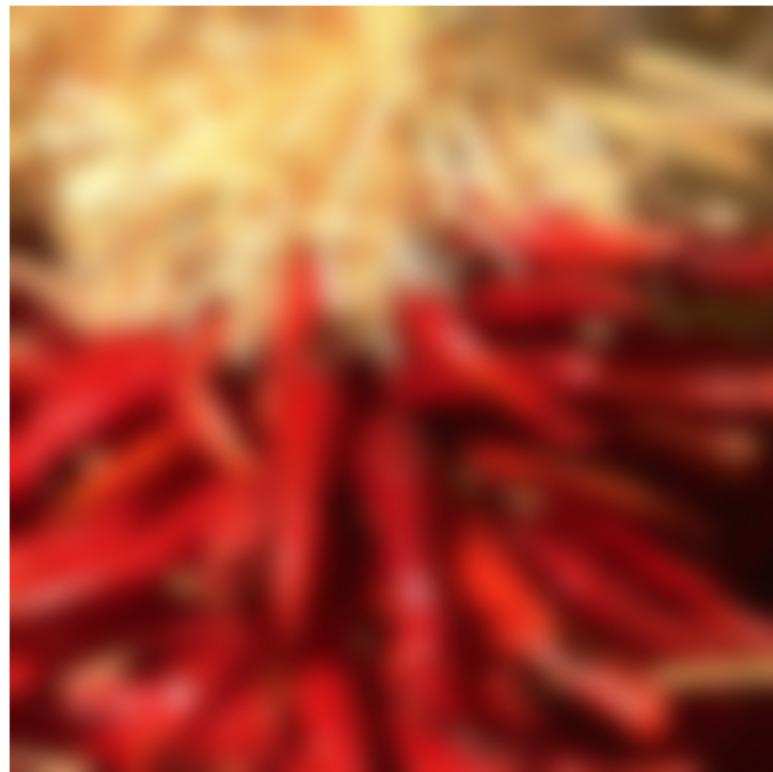
Sharpened

Unsharp mask filter



Boundary issues

- What about near the edge?
 - the filter window falls off the edge of the image
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge



Boundary issues

- What about near the edge?
 - the filter window falls off the edge of the image
 - need to extrapolate
 - methods (MATLAB):
 - clip filter (black): `imfilter(f, g, 0)`
 - wrap around: `imfilter(f, g, 'circular')`
 - copy edge: `imfilter(f, g, 'replicate')`
 - reflect across edge: `imfilter(f, g, 'symmetric')`

Common Types of Noise



Original



Salt and pepper noise



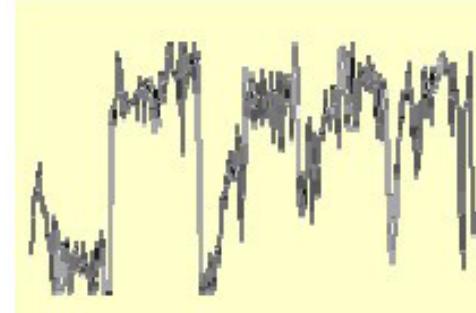
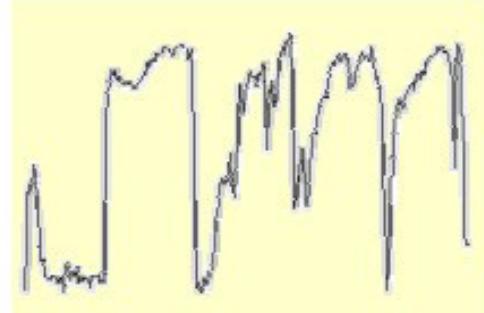
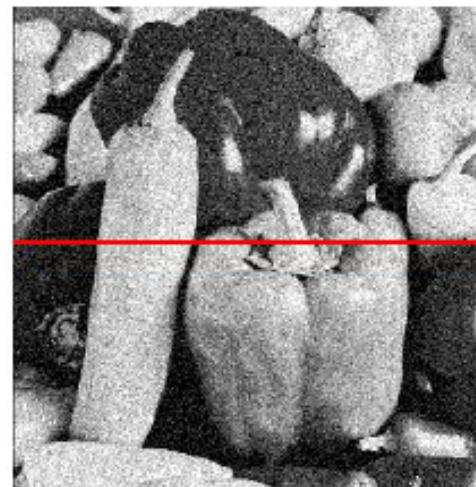
Impulse noise



Gaussian noise

- **Salt and pepper noise:** contains random occurrences of black and white pixels
- **Impulse noise:** contains random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution

Gaussian noise



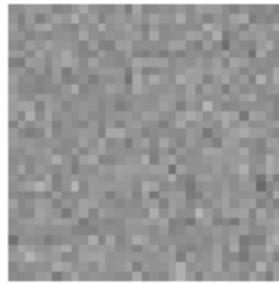
$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

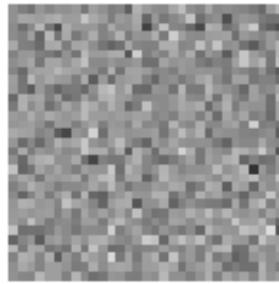
```
>> noise = randn(size(im)).*sigma;  
>> output = im + noise;
```

Reducing Gaussian noise

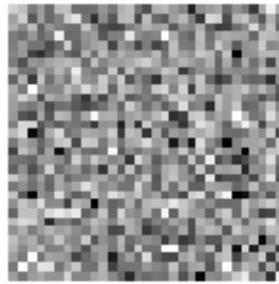
$\sigma=0.05$



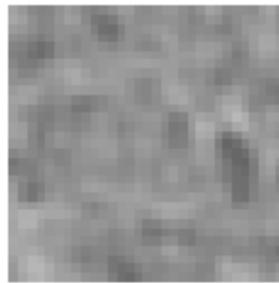
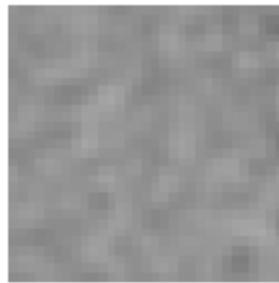
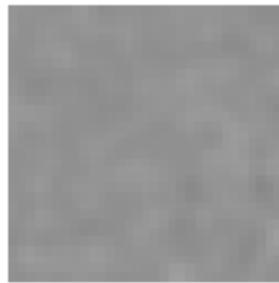
$\sigma=0.1$



$\sigma=0.2$



no
smoothing



$\sigma=1$ pixel



$\sigma=2$ pixels



Smoothing with larger standard deviations suppresses noise,
but also blurs the image

Reducing salt-and-pepper noise

3x3



5x5



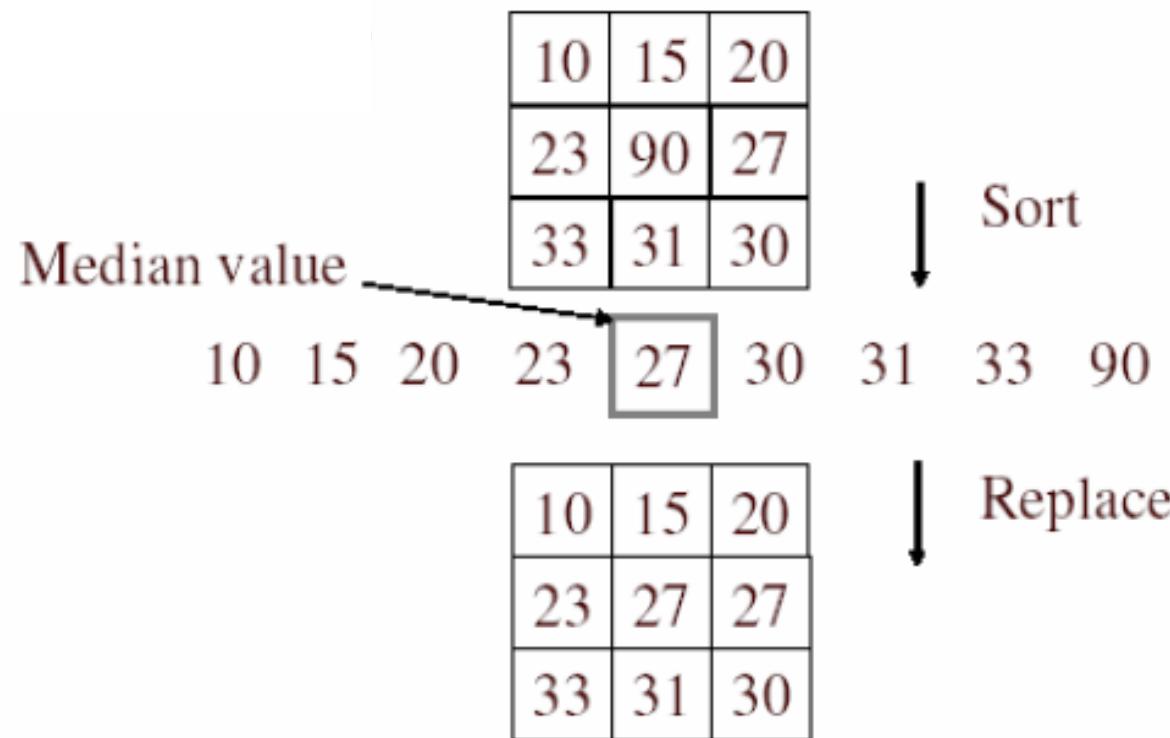
7x7



- What's wrong with the results?

Alternative idea: Median filtering

- A **median filter** operates over a window by selecting the median intensity in the window

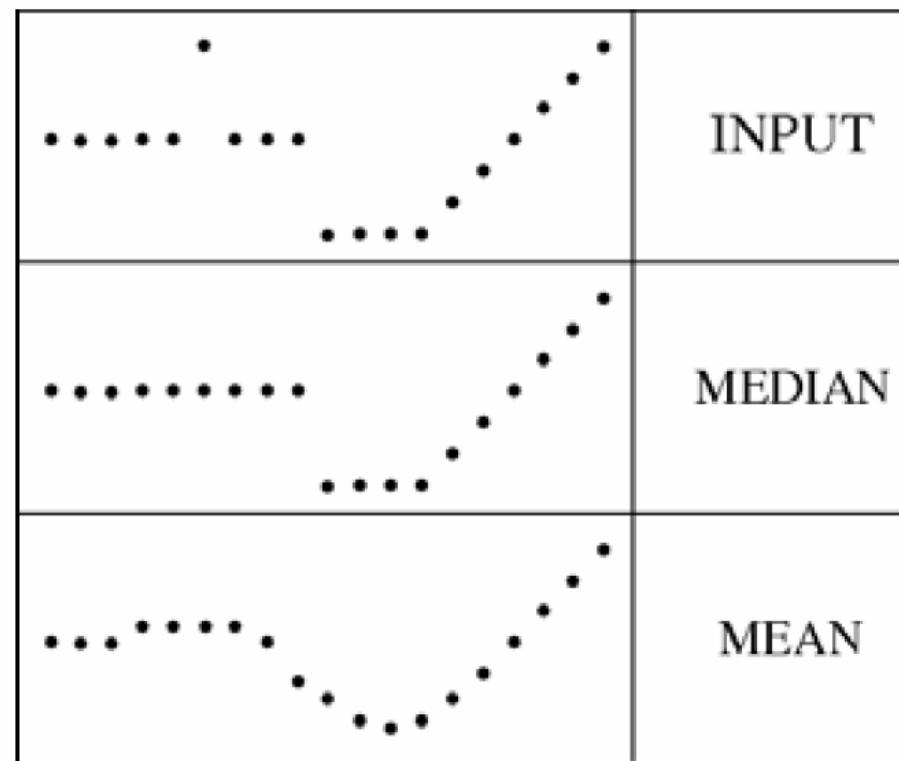


- Is median filtering linear?

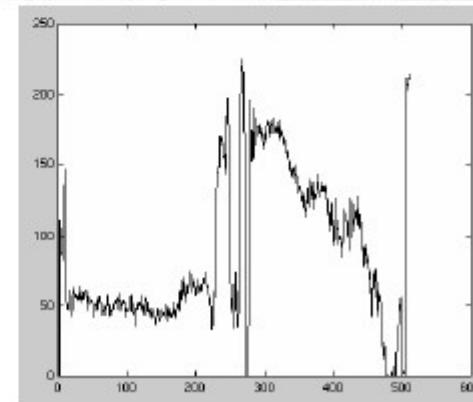
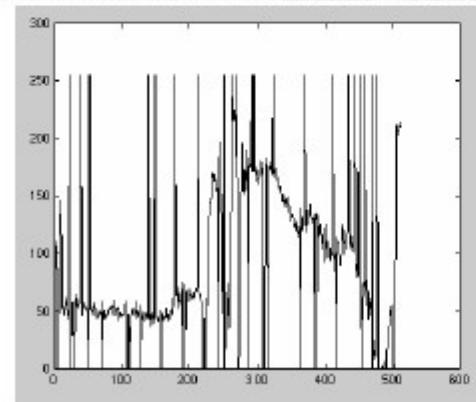
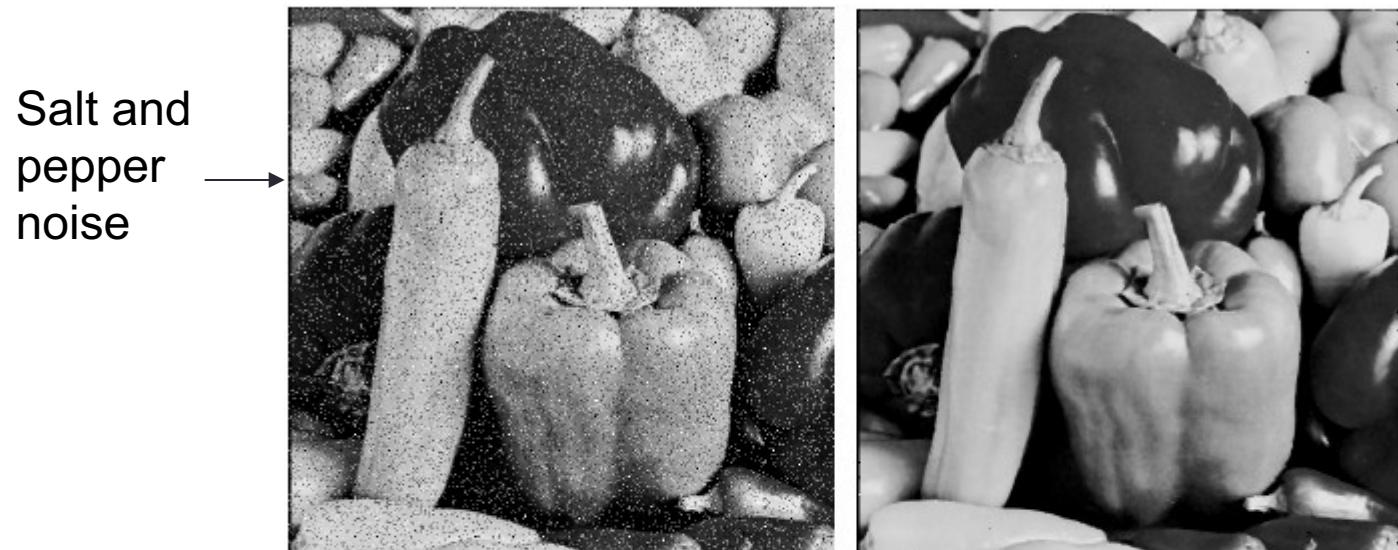
Median filter

- What advantage does median filtering have over Gaussian filtering?
 - Robustness to outliers
 - Edge preserving

filters have width 5 :

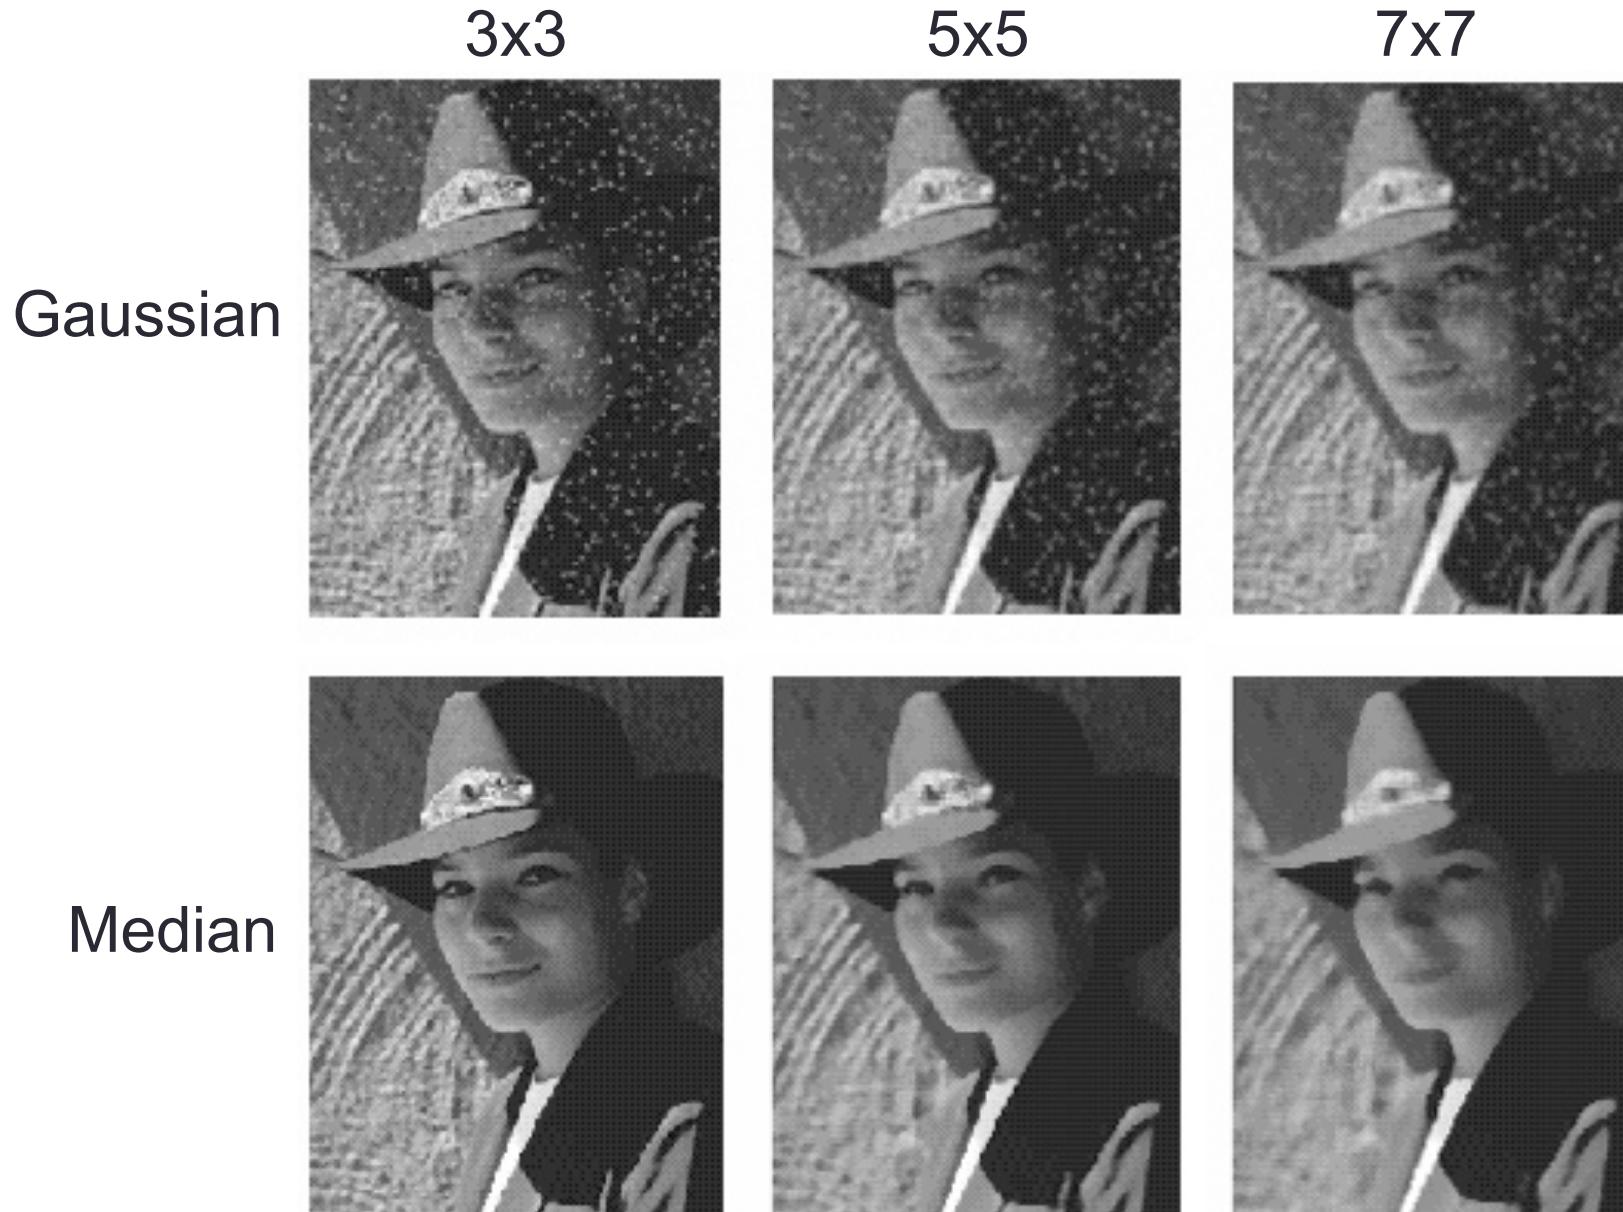


Median filter



Plots of a row of the image

Gaussian vs. median filtering



Readings

Szeliski Book
Chapter 2 and 3